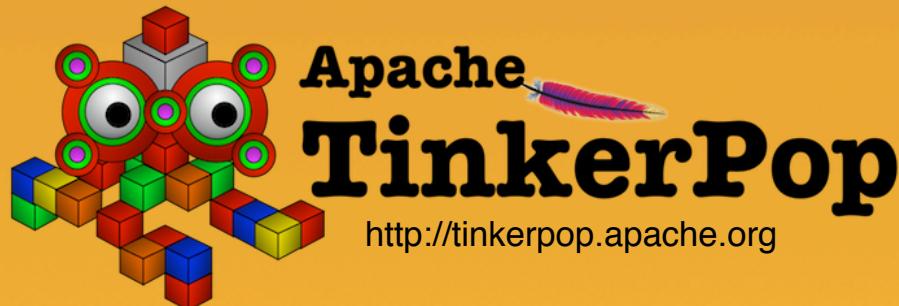


# Gremlin's Graph Traversal Machinery



Dr. Marko A. Rodriguez  
Director of Engineering at DataStax, Inc.  
Project Management Committee, Apache TinkerPop



CASSANDRA  
SUMMIT 2016

$$f : X \rightarrow X$$

The function f is a process that maps a structure of type X to a structure of type X.

$$f(x_1) = x_2$$

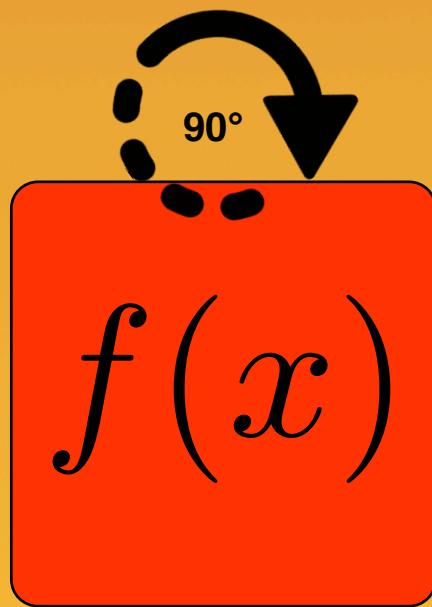
$$x_1 \in X$$

$$x_2 \in X$$

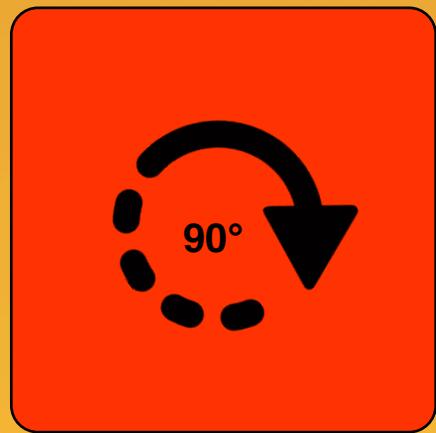
The function f maps the object x1 (from the set of X) to the object x2 (from the set of X).

$f(x)$ 

A **step** is a function.



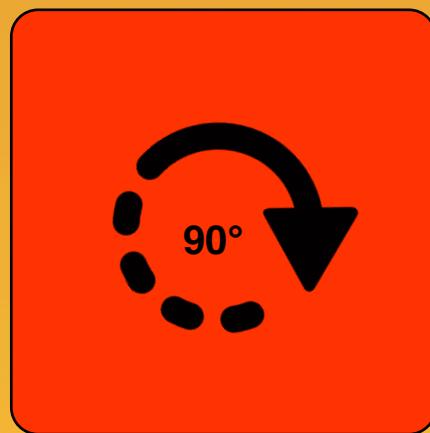
Assume that this step rotates an X by 90°.



The algorithm of the step is a “black box.”

```
class Traverser<V> {  
    V value;  
}
```

```
class Traverser<V> {  
    V value;  
}
```



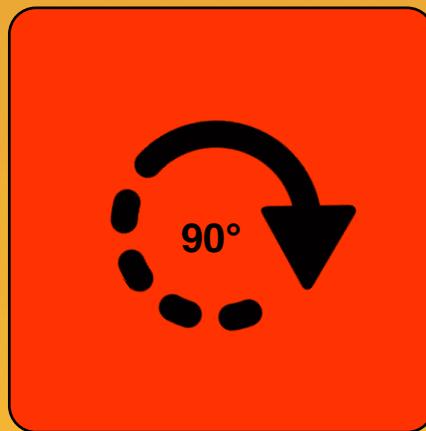
A **traverser** wraps a value of type V.

```
class Traverser<V> {  
    V value;  
}
```

Traverser<Integer>

```
class Traverser<V> {  
    V value;  
}
```

Traverser<Integer>



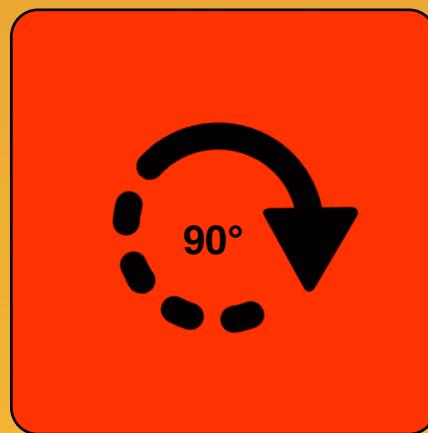
The step maps an integer traverser to an integer traverser.

```
class Traverser<V> {  
    V value;  
}
```

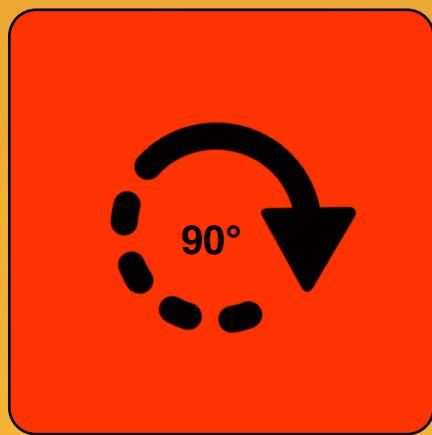
Traverser(0)

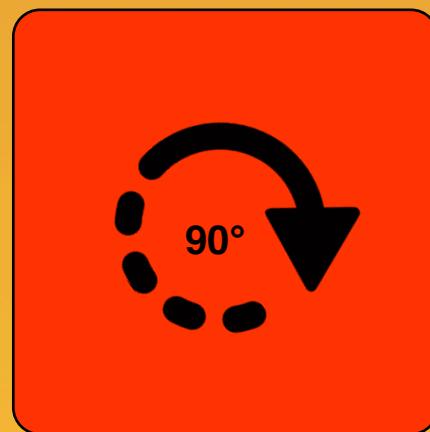
```
class Traverser<V> {  
    V value;  
}
```

Traverser(90)

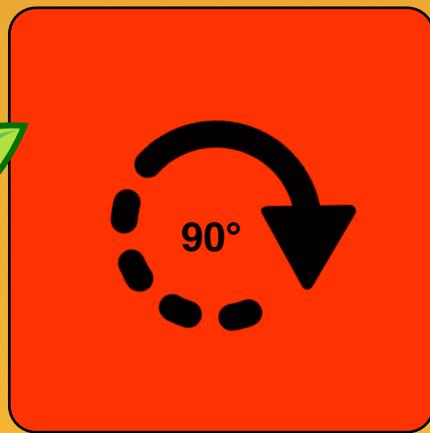
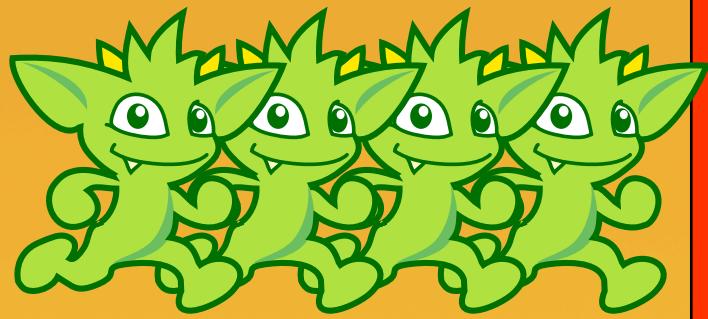


A traverser of with a rotation of  $0^\circ$  becomes a traverser with a rotation of  $90^\circ$ .

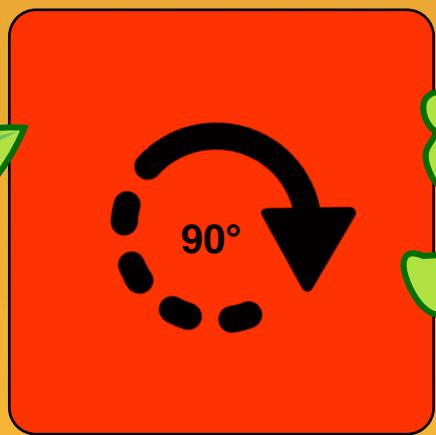


$T[N] \rightarrow T[N]$  $\in T[N]$   
 $\in T[N]$ 

Both the input and output traversers are of type `Traverser<Integer>`.



A stream of input traversers...



...yields a stream of output traversers.

```
class Traverser<V> {  
    V value;  
    long bulk;  
}
```

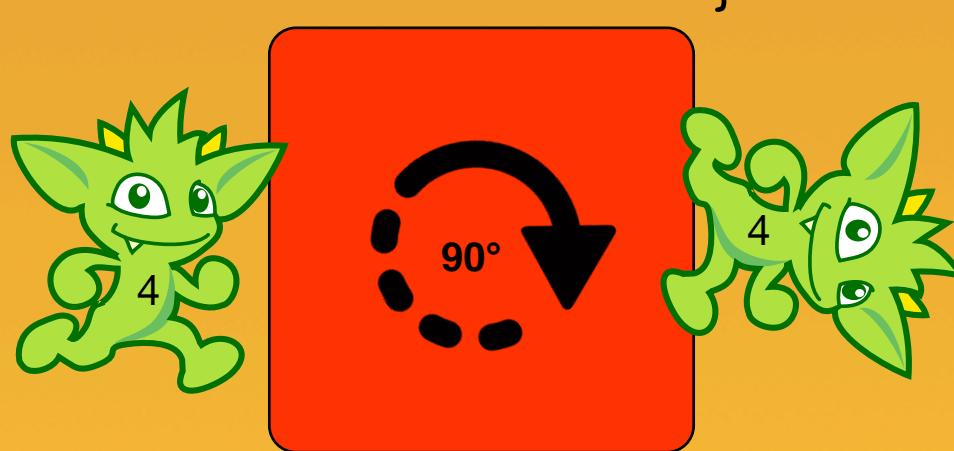
```
class Traverser<V> {  
    V value;  
    long bulk;  
}
```



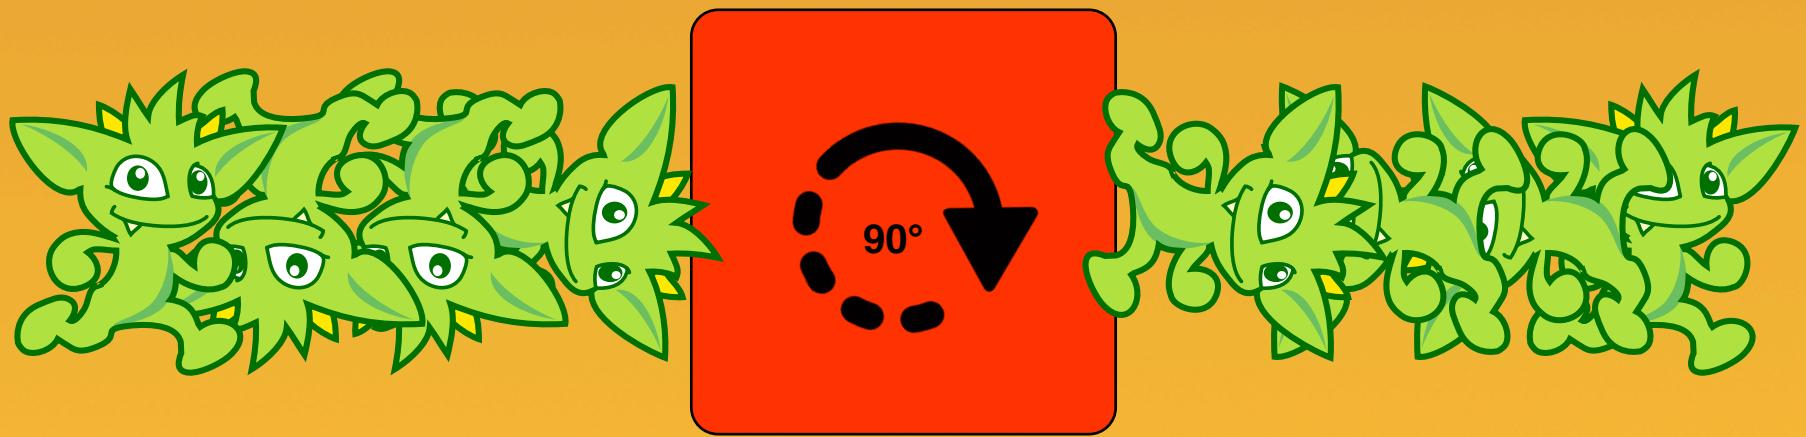
A traverser can have a **bulk** which denotes how many V values it represents.

```
class Traverser<V> {  
    V value;  
    long bulk;  
}
```

```
class Traverser<V> {  
    V value;  
    long bulk;  
}
```



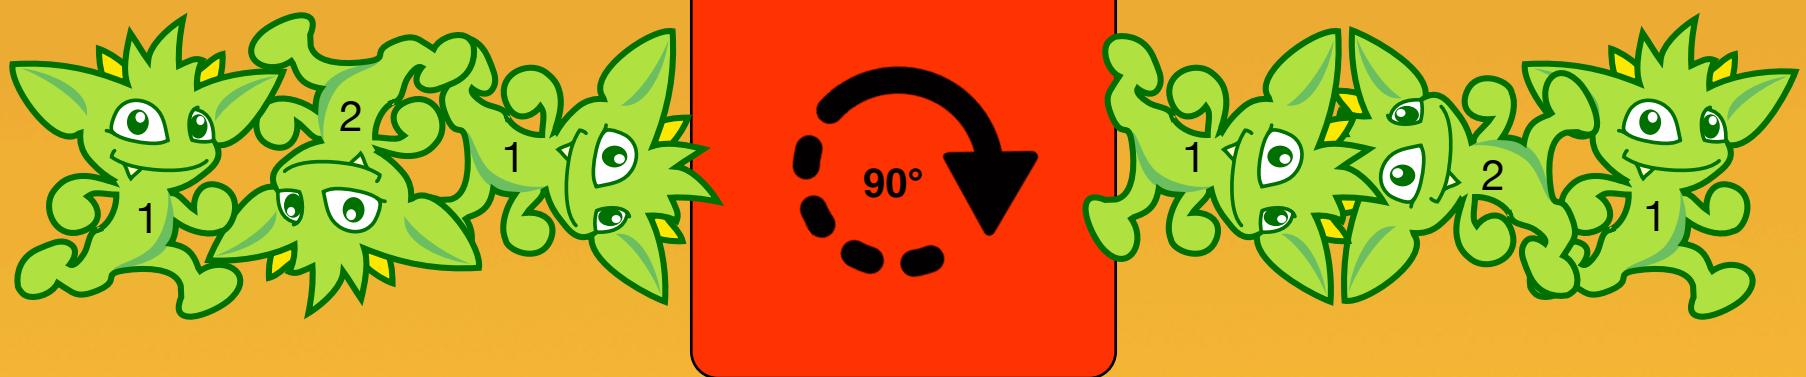
Bulking groups identical traversers to reduce the number of evaluations of a step.



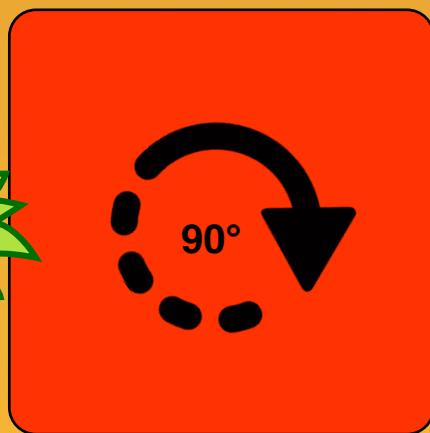
A variegated stream of input traversers yields a variegated stream of output traversers.

```
class Traverser<V> {  
    V value;  
    long bulk;  
}
```

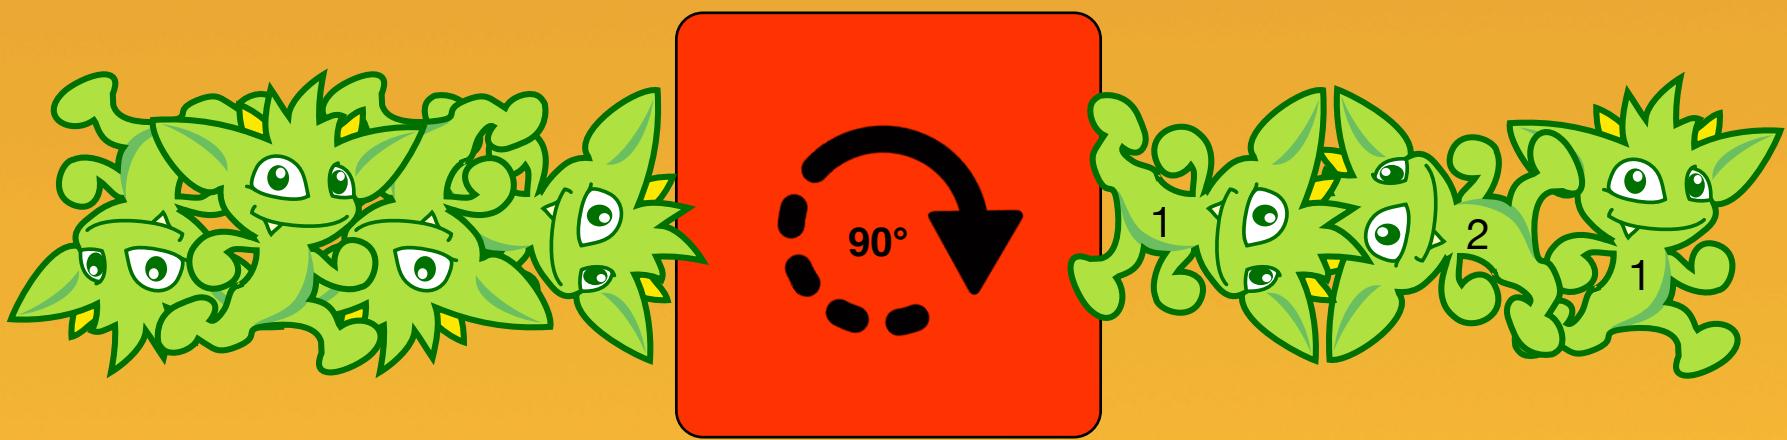
```
class Traverser<V> {  
    V value;  
    long bulk;  
}
```



Bulking can reduce the size of the stream.



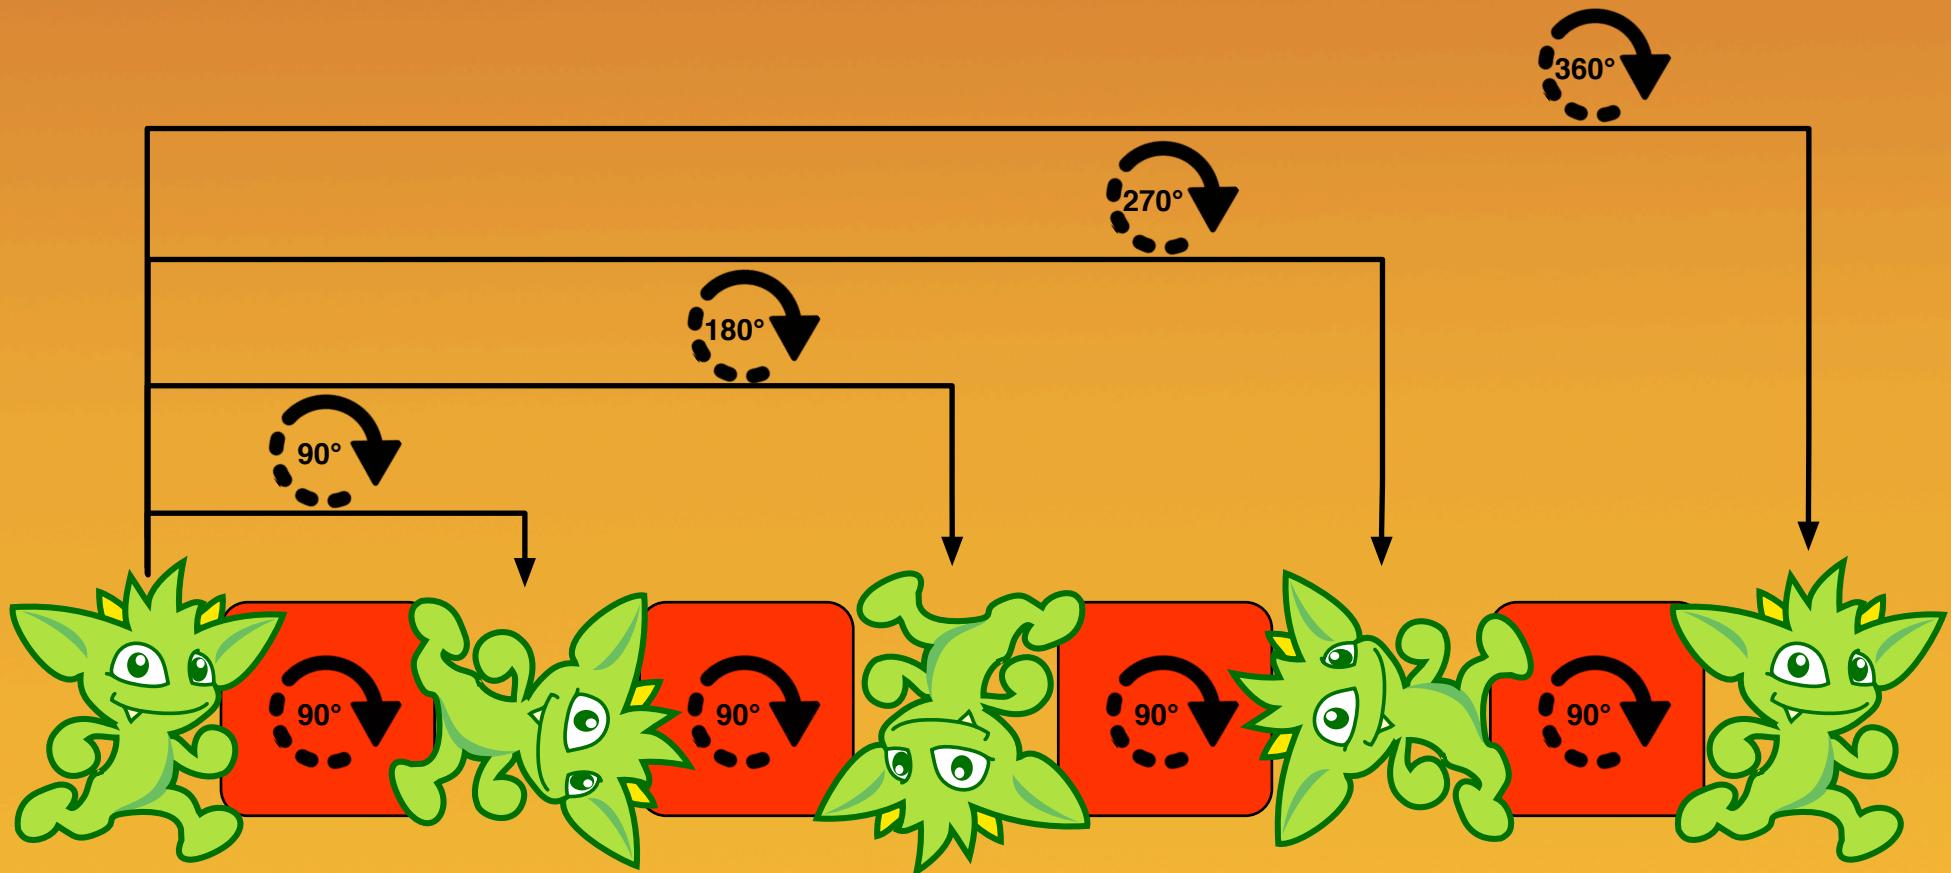
If the order of the stream does not matter...



total bulk = 4  
total count = 4

total bulk = 4  
total count = 3

...then reordering can increase the likelihood of bulking.



A **traversal** is a list of one or more steps.

$$f : X \rightarrow Y$$

$$h : Y \rightarrow Z$$

Different functions can yield different mappings between different domains and ranges.

$$Y \ y = f(x)$$

$$Z \ z = h(y)$$

The output of  $f$  can be the input to  $h$  because the range of  $f$  is the domain of  $h$  (i.e.  $Y$ ).

$$y = f(x)$$

$$z = h(y)$$

$$y=f(x)$$

$$z=h(y)$$

$$f(x)=y$$

$$h(y)=z$$

$$f(x)=y$$

$$h(y)=z$$

$$f(x)\not\equiv(y)=z$$

$$f(x) \; h \qquad = z$$

$$x^f_h=z$$

$$x \cdot f \cdot h = z$$

readable

$$x \cdot f \cdot h = z$$

≡

unreadable

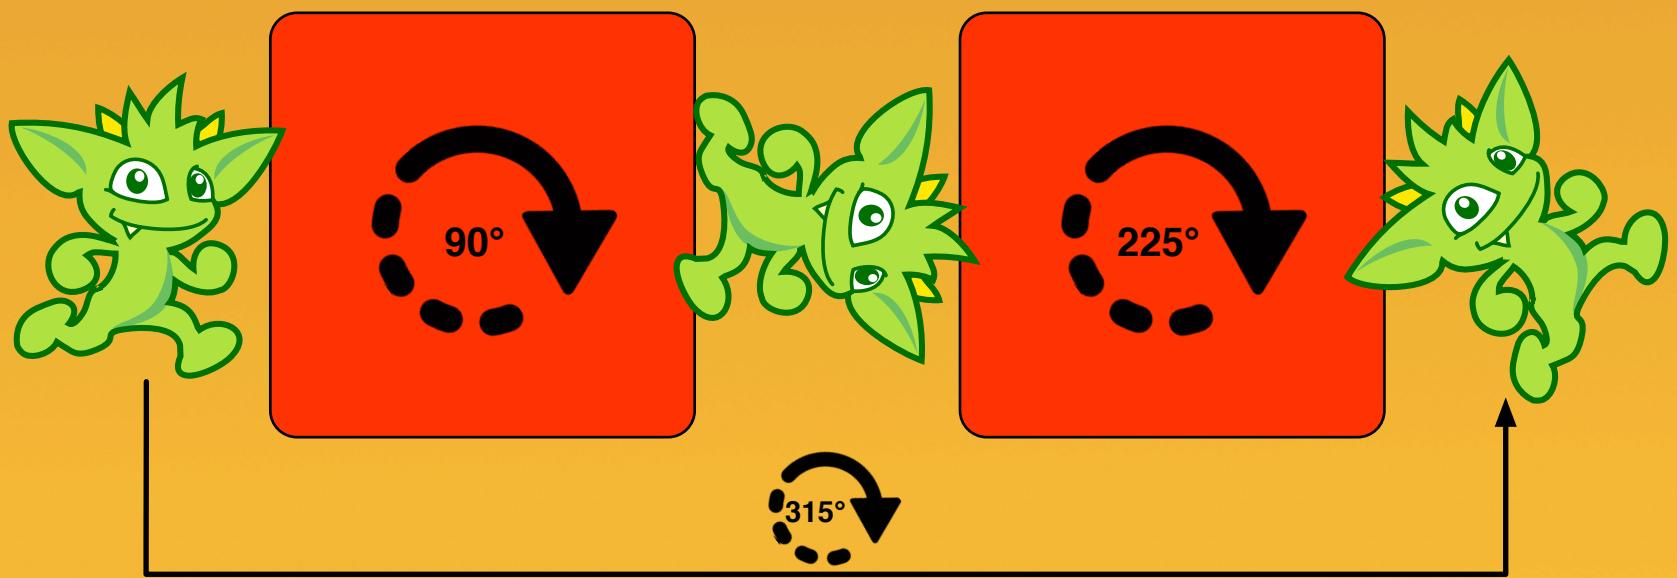
$$h(f(x)) = z$$

$$z = x \cdot f \cdot h$$

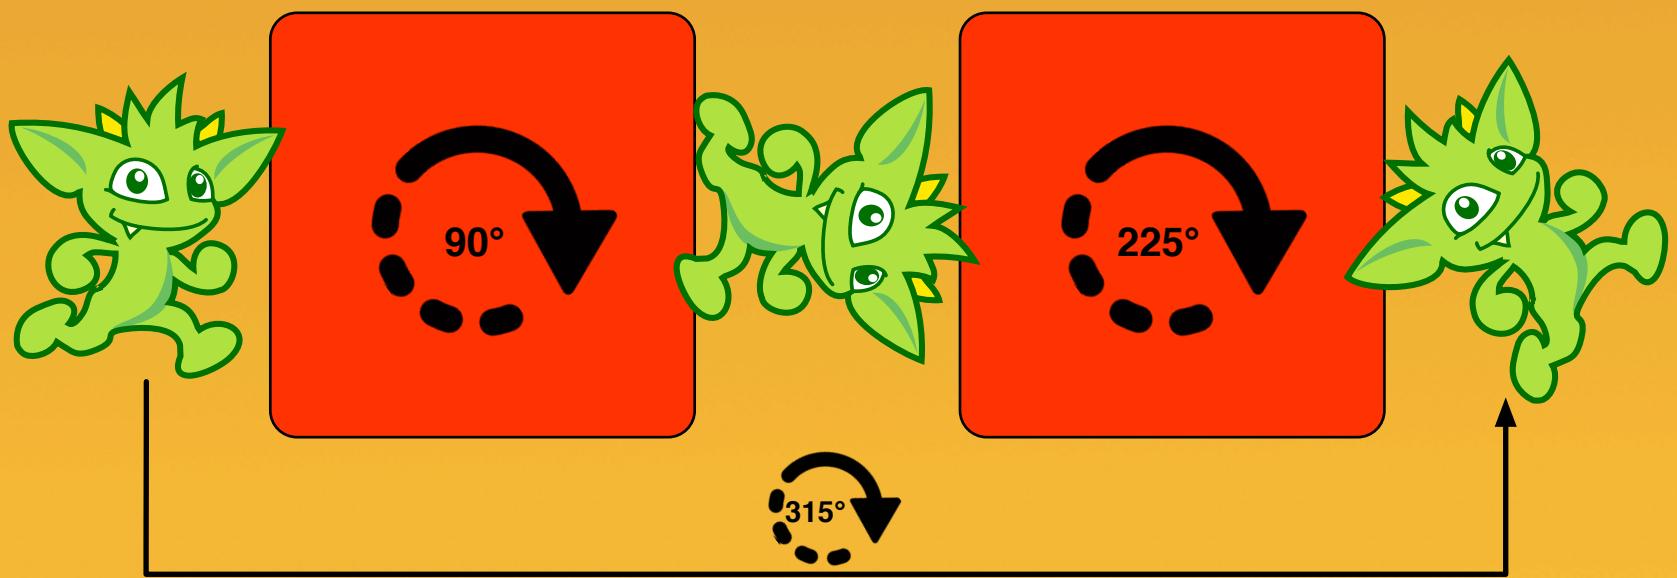
$$z = x \cdot f \cdot h$$

`z = x.f().h().next()`

head/start   stream/iterator/traversal

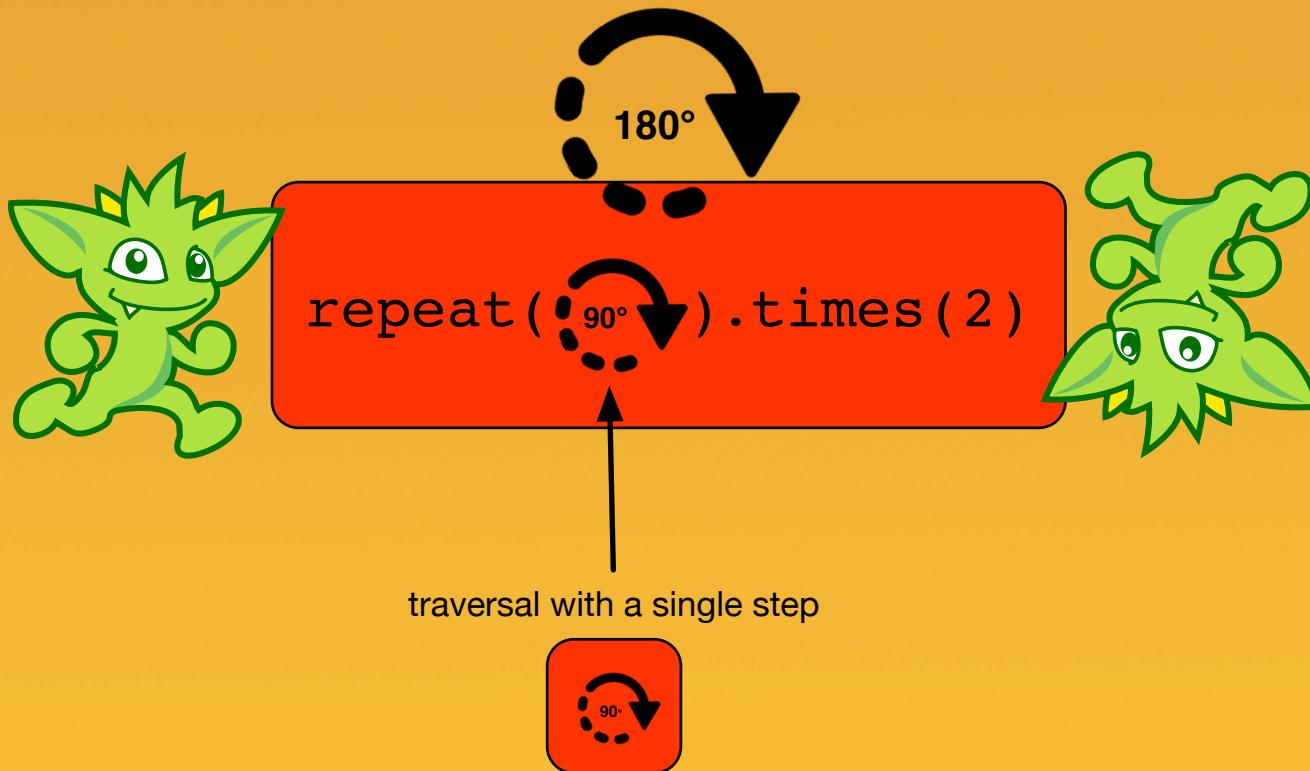


Different types of steps exist and their various compositions yield query expressivity.



 =  .  ( ) .  ( ) . next ( )

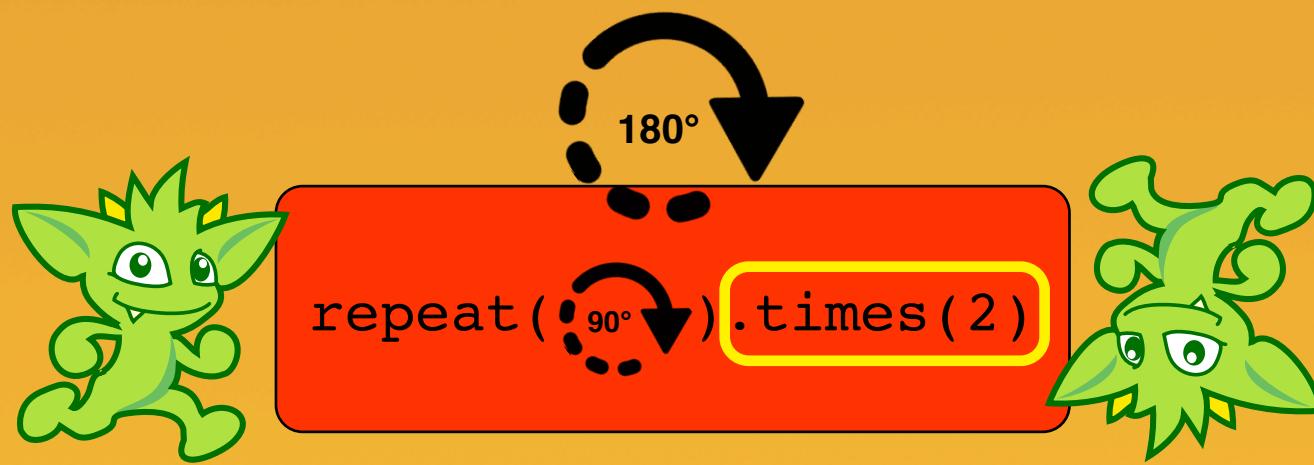
Steps process traverser streams/flows and their composition is a traversal/iterator.



**Anonymous traversals** can serve as step arguments.

`...select("a","b").by("name")`

`...groupCount().by(out().count())`

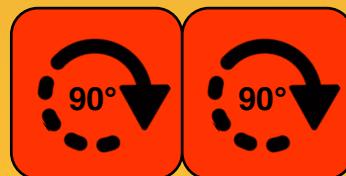
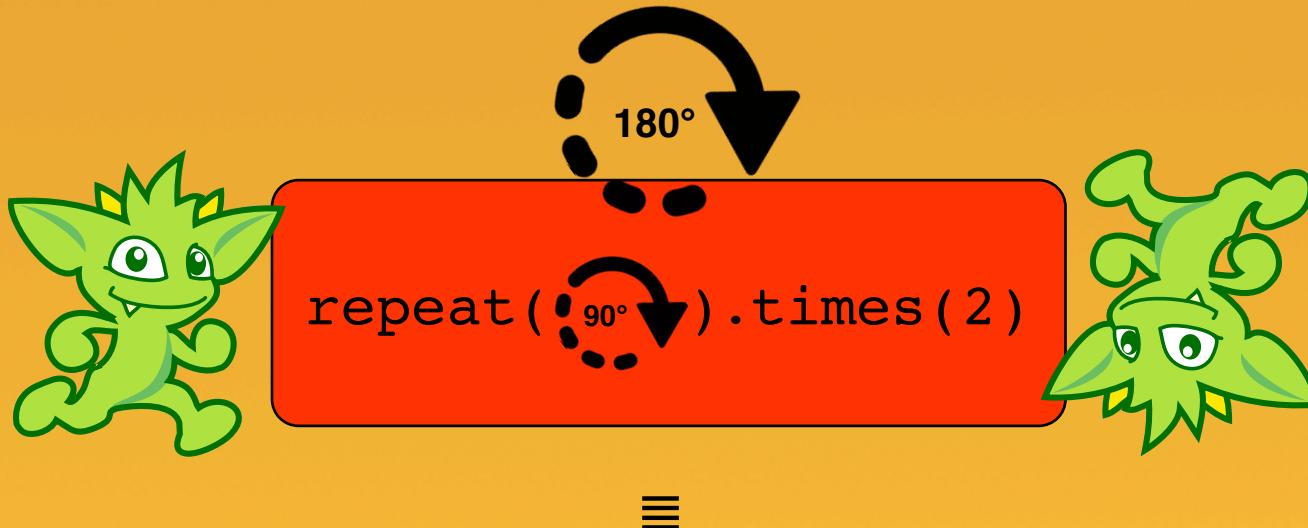


`...addE("knows").from("a").to("b")`

`...order().by("age", decr)`

Some functions serve as **step-modulators** instead of steps.

```
interface TraversalStrategy {  
    void apply(Traversal traversal);  
    Set<TraversalStrategy> applyPrior();  
    Set<TraversalStrategy> applyPost();  
}
```



During optimization, **traversal strategies** may rewrite a traversal to a more efficient, semantically equivalent form.

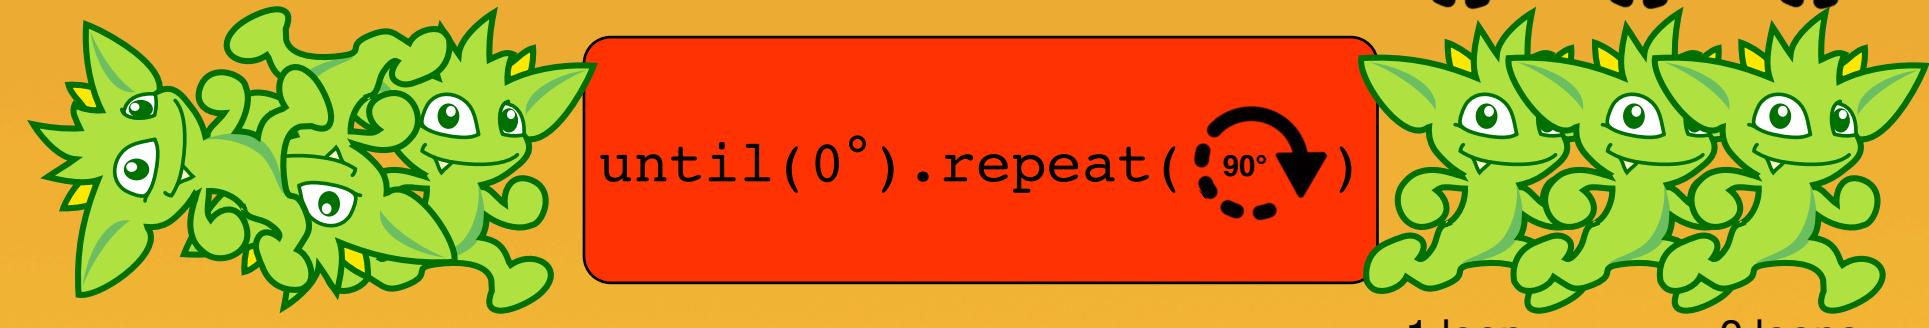


```
repeat(90°).until(0°)
```

Continuously process a traverser until some condition is met.



`repeat().until()` provides do/while-semantics.



≡

`while(x != !) do`



`until().repeat()` provides while/do-semantics.



```
until(0°).repeat(3)
```

3

Even if the traversers in the input stream can not be bulked, the output traversers may be able to be.

**map(x)**

one-to-one

**flatMap(x)**

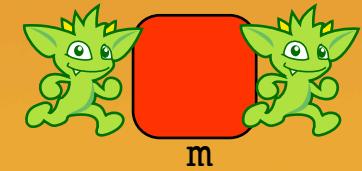
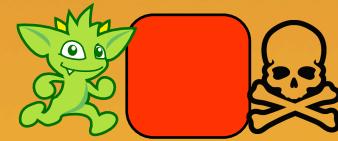
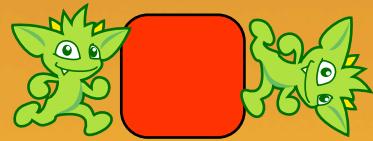
one-to-many

**filter(x)**

one-to-(one-or-none)

**sideEffect(x)**

one-to-same



**map(x)**

one-to-one

**flatMap(x)**

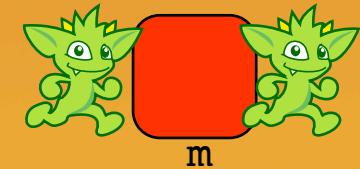
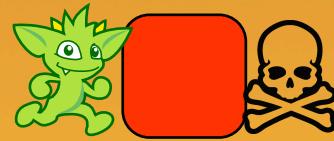
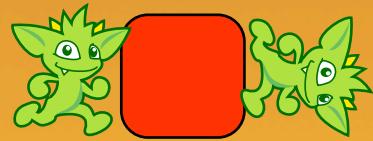
one-to-many

**filter(x)**

one-to-(one-or-none)

**sideEffect(x)**

one-to-same

`out("knows")``groupCount("m")``select("a","b")``has("name","gremlin")``values("age")``in("created")``store("m")``path()``where("a",eq("b"))``tree("m")``id()``and(x,y)``subgraph("m")``order()``match(x,y,z)``or(x,y)``label()``coin(0.5)``group("m")``mean()``properties()``sample(10)``sum()``outE("knows")``simplePath()``count()``v()``dedup()``groupCount()`

\* A collection of examples. Not the full step library.

**map(x)**

one-to-one

**flatMap(x)**

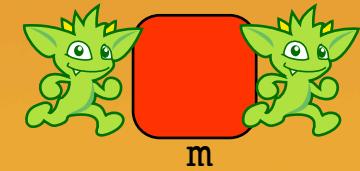
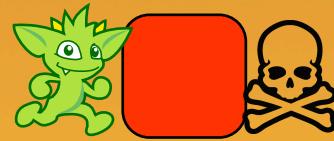
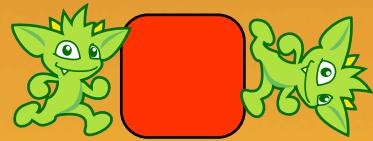
one-to-many

**filter(x)**

one-to-(one-or-none)

**sideEffect(x)**

one-to-same



`out("knows")`

`has("name", "gremlin")`

`values("age")`

`path()`

`label()`

`group("m")`

`outE("knows")`

`simplePath()`

`v()`

`groupCount()`

**map(x)**

one-to-one

**flatMap(x)**

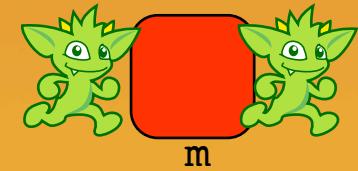
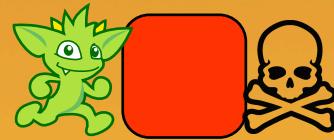
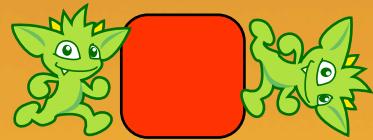
one-to-many

**filter(x)**

one-to-(one-or-none)

**sideEffect(x)**

one-to-same



$$T[V] \rightarrow T[V]^*$$

out("knows")

$$T[V \cup E] \rightarrow T[\mathbb{N}^+]$$

values("age")

path()

$$T[?] \rightarrow T[path]$$

$$T[V \cup E] \rightarrow \emptyset \cup T[V \cup E]$$

has("name", "gremlin")

$$\text{label()}$$

$$T[V \cup E] \rightarrow T[\text{string}]$$

$$T[V] \rightarrow T[E]^*$$

outE("knows")

$$T[?] \rightarrow T[\text{map}[?, \mathbb{N}^+]]$$

groupCount()

$$T[G] \rightarrow T[V]^*$$

v()

$$T[?] \rightarrow T[?]$$

group("m")

$$T[?] \rightarrow \emptyset \cup T[?]$$

simplePath()

```
g.V().has("name", "gremlin").  
out("knows").values("age").  
groupCount()
```

$T[V] \rightarrow T[V]^*$ $T[V \cup E] \rightarrow T[\mathbb{N}^+]$ $\text{values("age")}$	$\text{out("knows")}$ $T[V \cup E] \rightarrow \emptyset \cup T[V \cup E]$ $\text{has("name", "gremlin")}$
--	--

```
 $T[?] \rightarrow T[\text{map}[?, \mathbb{N}^+]]$      $T[G] \rightarrow T[V]^*$   
 $\text{v()}$   
 $\text{groupCount()}$ 
```

```

g.V().has("name", "gremlin").
    out("knows").values("age").
    groupCount()

```

	$v()$	$T[G] \rightarrow T[V]^*$
<code>has("name", "gremlin")</code>		$T[V \cup E] \rightarrow \emptyset \cup T[V \cup E]$
<code>out("knows")</code>		$T[V] \rightarrow T[V]^*$
<code>values("age")</code>		$T[V \cup E] \rightarrow T[\mathbb{N}^+]$
<code>groupCount()</code>		$T[?] \rightarrow T[\text{map}[?, \mathbb{N}^+]]$

Steps can be composed if their respective domains and ranges match.

```

g.V().has("name", "gremlin") .
  out("knows").values("age") .
  groupCount()

```

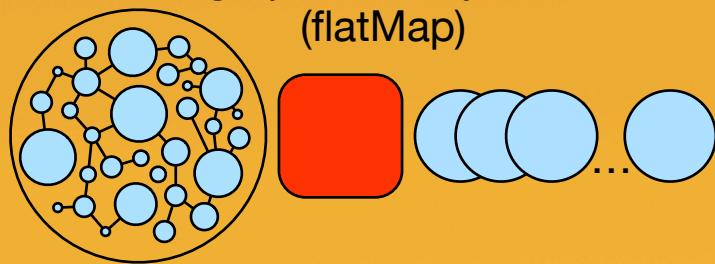
	v()	$T[G] \rightarrow T[V]^*$
has("name", "gremlin")		$T[V] \rightarrow \emptyset \cup T[V]$
out("knows")		$T[V] \rightarrow T[V]^*$
values("age")		$T[V] \rightarrow T[\mathbb{N}^+]$
groupCount()		$T[\mathbb{N}^+] \rightarrow T[\text{map}[\mathbb{N}^+, \mathbb{N}^+]]$

```
g.V().has("name", "gremlin").  
out("knows").values("age").  
groupCount()
```

What is the distribution of ages of the people that Gremlin knows?

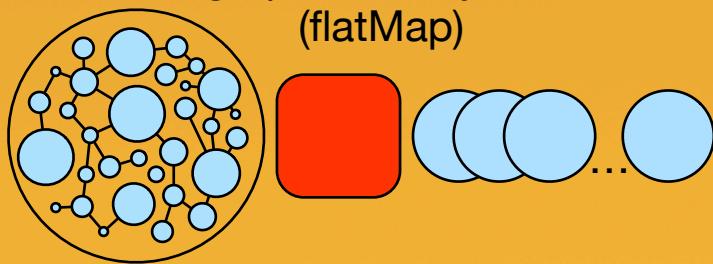
```
g.V().has("name", "gremlin").  
out("knows").values("age").  
groupCount()
```

one graph to many vertices  
(flatMap)

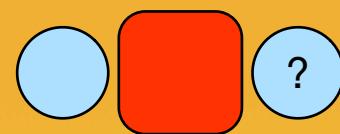


```
g.V().has("name", "gremlin").  
out("knows").values("age").  
groupCount()
```

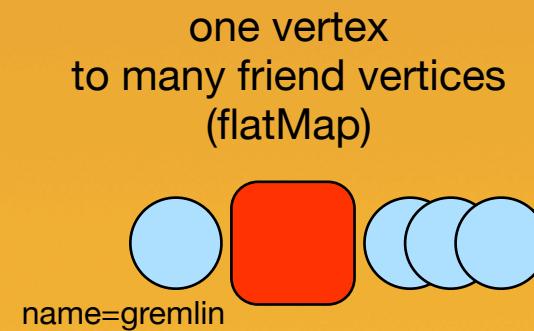
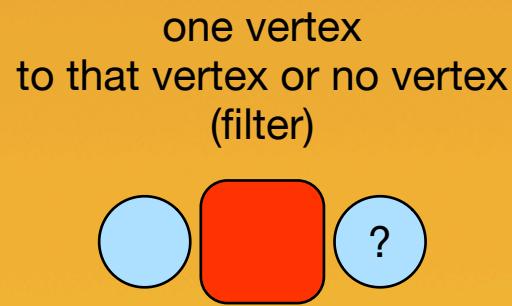
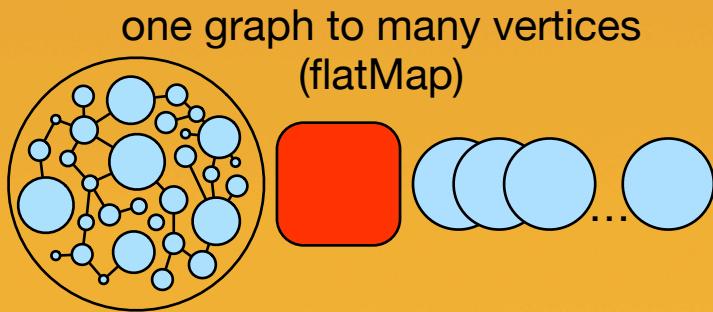
one graph to many vertices  
(flatMap)



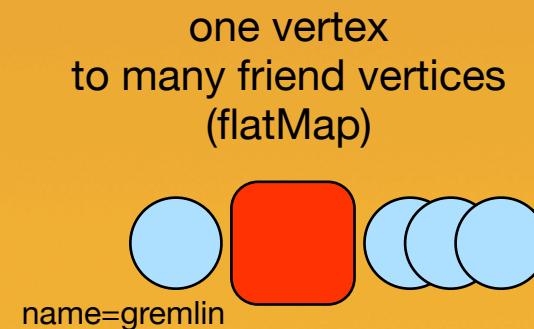
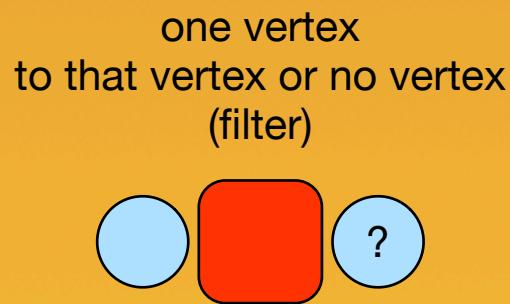
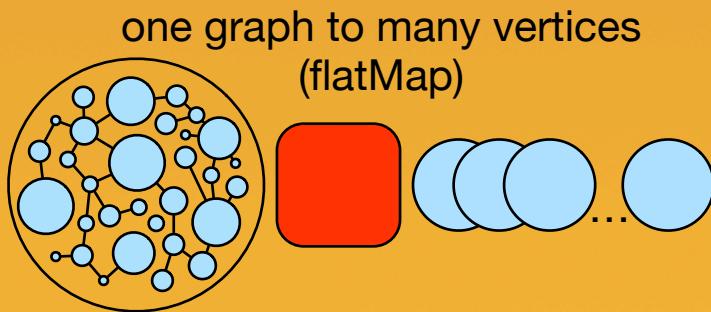
one vertex  
to that vertex or no vertex  
(filter)



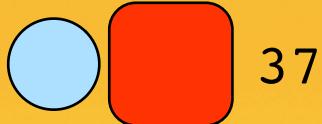
```
g.V().has("name", "gremlin").  
  out("knows").values("age").  
  groupCount()
```



```
g.V().has("name", "gremlin").  
out("knows").values("age").  
groupCount()
```

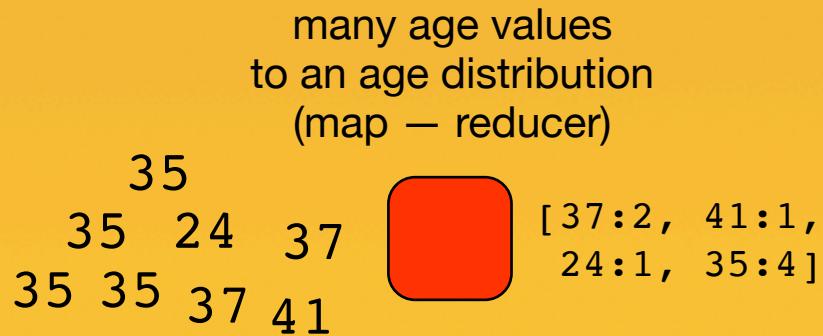
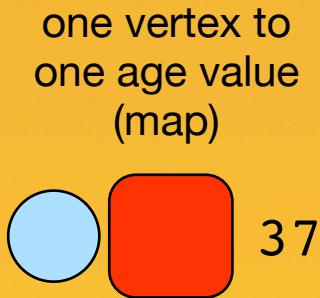
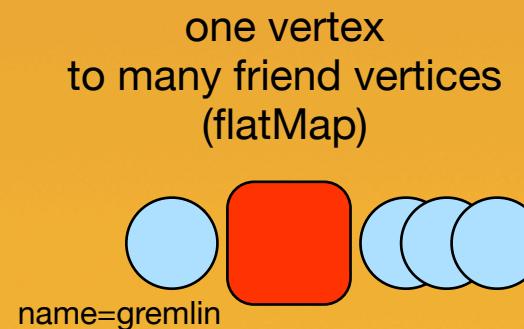
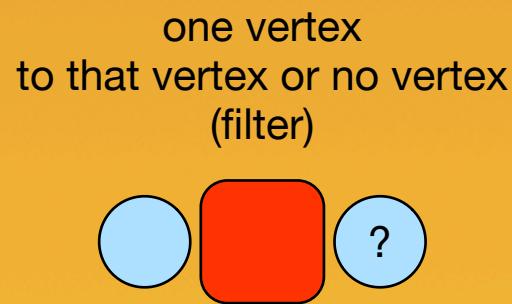
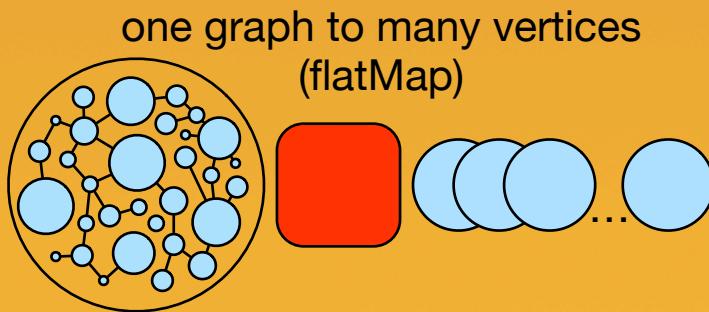


one vertex to  
one age value  
(map)

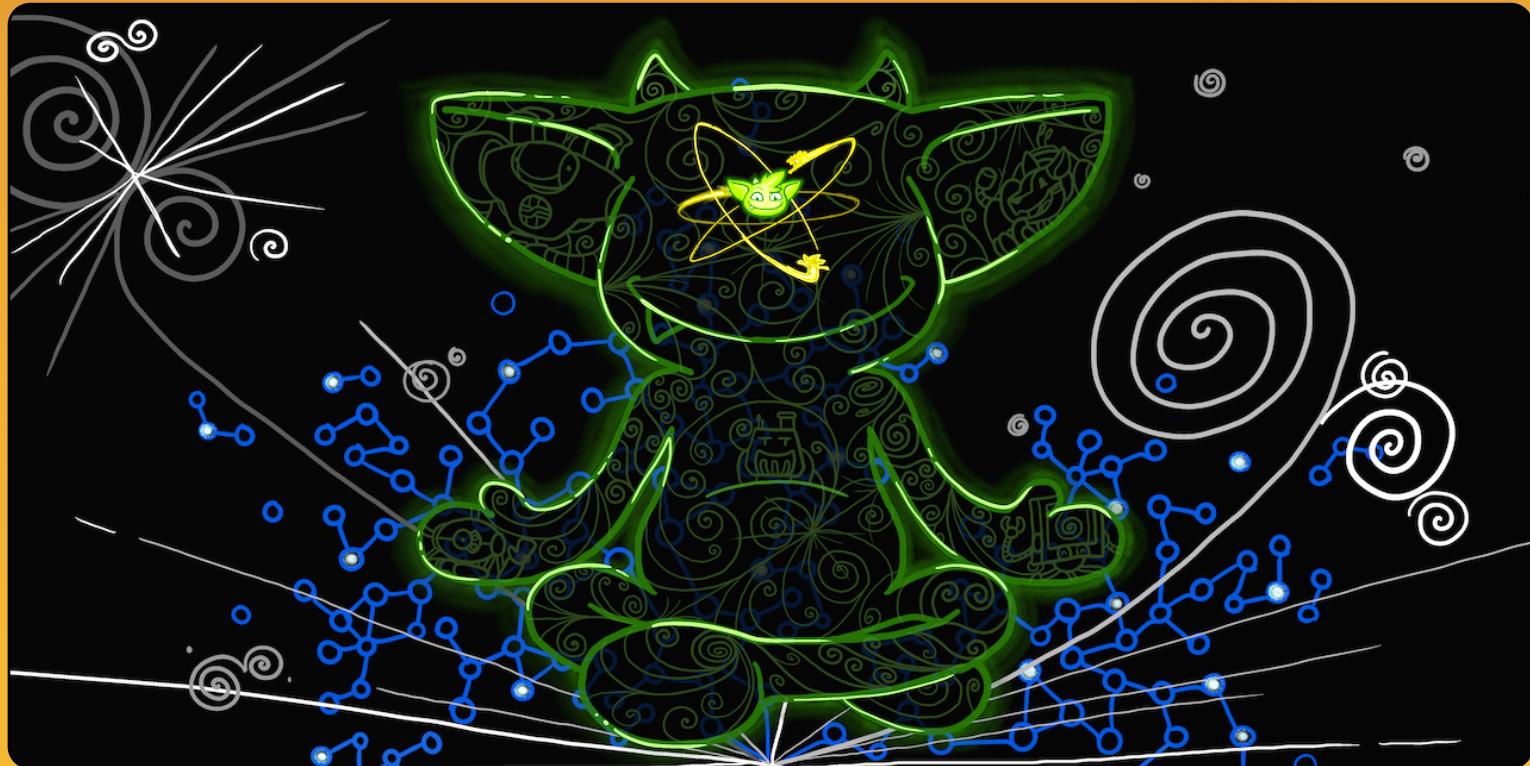


37

```
g.V().has("name", "gremlin").  
out("knows").values("age").  
groupCount()
```



# The Gremlin Traversal Language



The Gremlin Traversal Machine

# Gremlin Traversal Language

function composition

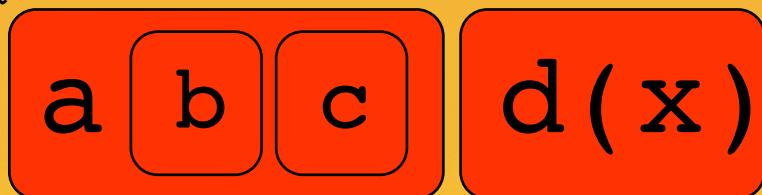


Traversal creation via  
step composition

fluent methods

`a().b().c()`

function nesting



method arguments

`a(b().c()).d(x)`

Step parameterization via  
traversal and constant nesting

Any language that supports function composition and function nesting can **host** Gremlin.

# Gremlin Traversal Machine

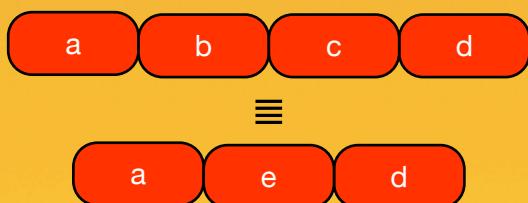
```
class Traverser<V> {  
    V value;  
    long bulk;  
}
```



```
class Step<S,E> {  
    Traverser<E> processNextStart();  
}
```

$$f(x)$$

```
class Traversal<S,E> implements Iterator<E> {  
    E next();  
    Traverser<E> nextTraverser();  
}
```



```
interface TraversalStrategy {  
    void apply(Traversal traversal);  
    Set<TraversalStrategy> applyPrior();  
    Set<TraversalStrategy> applyPost();  
}
```

The fundamental constructs of Gremlin's machinery.



Gremlin-Java

```
g.v(1).  
repeat(out("knows")).times(2).  
groupCount().by("age")
```

translates

```
[[v, 1]  
[repeat, [[out, knows]]]  
[times, 2]  
[groupCount]  
[by, age]]
```

compiles

Traversal

GraphStep

RepeatStep

GroupCountStep

Traversal  
Strategies

optimizes

GraphStep

VertexStep

VertexStep

GroupCountStep

executes



GraphStep

VertexStep

VertexStep

GroupCountStep



[29:2, 30:1,  
31:1, 35:10]



Gremlin-Python

```
g.v(1).  
repeat(out('knows')).times(2).  
groupCount().by('age')
```

Bytecode

```
[[v, 1]  
[repeat, [[out, knows]]]  
[times, 2]  
[groupCount]  
[by, age]]
```

Traversal

GraphStep

RepeatStep

GroupCountStep

Traversal  
Strategies

optimizes

GraphStep

VertexStep

VertexStep

GroupCountStep

executes



GraphStep

VertexStep

VertexStep

GroupCountStep



```
[29:2, 30:1,  
31:1, 35:10]
```

Language

Gremlin language variant

Gremlin-Python

```
g.v(1).  
repeat(out('knows')).times(2).  
groupCount().by('age')
```

Language agnostic bytecode

Bytecode

```
[[v, 1]  
[repeat, [[out, knows]]]  
[times, 2]  
[groupCount]  
[by, age]]
```

Execution engine assembly

Traversal

GraphStep

RepeatStep

GroupCountStep

VertexStep

VertexStep

Traversal  
Strategies

Execution engine optimization

GraphStep

VertexStep

VertexStep

GroupCountStep

Execution engine evaluation



GraphStep

VertexStep

VertexStep

GroupCountStep



```
[29:2, 30:1,  
31:1, 35:10]
```

translates

compiles

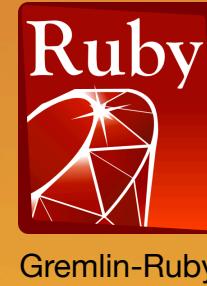
optimizes

executes

Language



Gremlin-Python



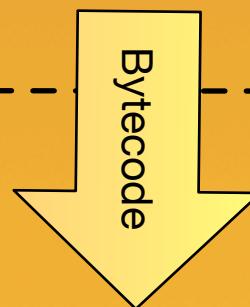
Gremlin-Ruby



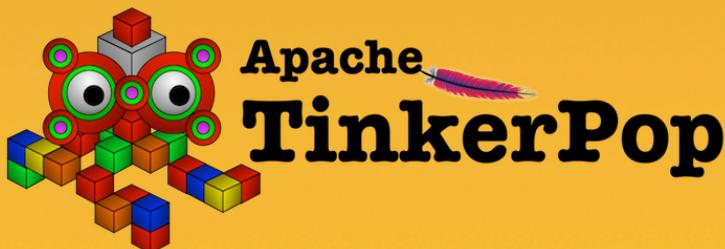
JavaScript  
Gremlin-JavaScript



Gremlin-Groovy



Clojure  
Gremlin-Clojure



Java-based Implementation

? ? ?

?-based Implementation

```
Python 2.7.2 (default, Oct 11 2012, 20:14:37)
[GCC 4.2.1 Compatible Apple Clang 4.0 (tags/Apple/clang-418.0.60)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
>>> from gremlin_python.structure.graph import Graph
>>> from gremlin_python.driver.driver_remote_connection import DriverRemoteConnection
```



CPython

Gremlin-Python

```
Python 2.7.2 (default, Oct 11 2012, 20:14:37)
[GCC 4.2.1 Compatible Apple Clang 4.0 (tags/Apple/clang-418.0.60)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
>>> from gremlin_python.structure.graph import Graph
>>> from gremlin_python.driver.driver_remote_connection import DriverRemoteConnection
>>> graph = Graph()
>>> g = graph.traversal().withRemote(DriverRemoteConnection('ws://localhost:8182','g'))
```

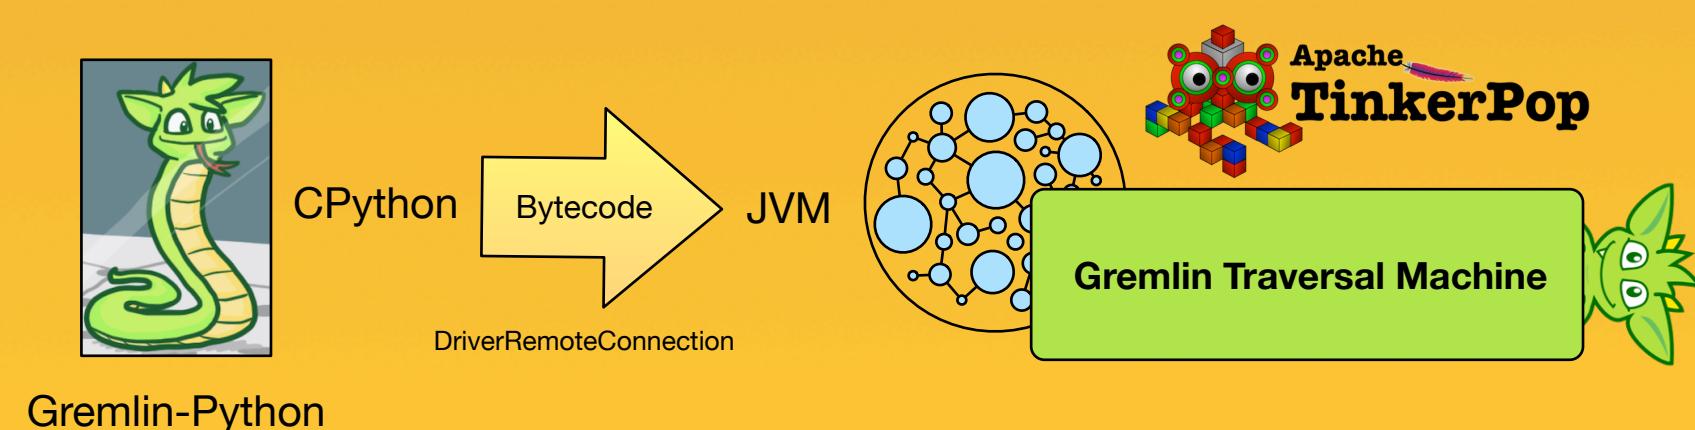


CPython

DriverRemoteConnection

Gremlin-Python

```
Python 2.7.2 (default, Oct 11 2012, 20:14:37)
[GCC 4.2.1 Compatible Apple Clang 4.0 (tags/Apple/clang-418.0.60)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
>>> from gremlin_python.structure.graph import Graph
>>> from gremlin_python.driver.driver_remote_connection import DriverRemoteConnection
>>> graph = Graph()
>>> g = graph.traversal().withRemote(DriverRemoteConnection('ws://localhost:8182','g'))
# nested traversal with Python slicing and attribute interception extensions
>>> g.V().hasLabel("person").repeat(both()).times(2).name[0:2].toList()
[u'marko', u'marko']
```



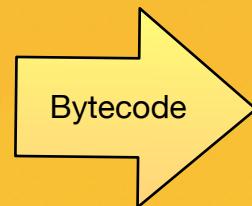
```

Python 2.7.2 (default, Oct 11 2012, 20:14:37)
[GCC 4.2.1 Compatible Apple Clang 4.0 (tags/Apple/clang-418.0.60)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
>>> from gremlin_python.structure.graph import Graph
>>> from gremlin_python.driver.driver_remote_connection import DriverRemoteConnection
>>> graph = Graph()
>>> g = graph.traversal().withRemote(DriverRemoteConnection('ws://localhost:8182','g'))
# nested traversal with Python slicing and attribute interception extensions
>>> g.V().hasLabel("person").repeat(both()).times(2).name[0:2].toList()
[u'marko', u'marko']
# a complex, nested multi-line traversal
>>> g.V().match(
...     as_("a").out("created").as_("b"), \
...     as_("b").in_("created").as_("c"), \
...     as_("a").out("knows").as_("c")). \
...     select("c"). \
...     union(in_("knows"),out("created")). \
...     name.toList()
[u'ripple', u'marko', u'lop']
>>>

```

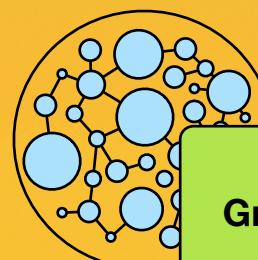


C<sub>Python</sub>



JVM

DriverRemoteConnection



Gremlin Traversal Machine



Gremlin-Python



**Gremlin**  
 $G = (V, E)$



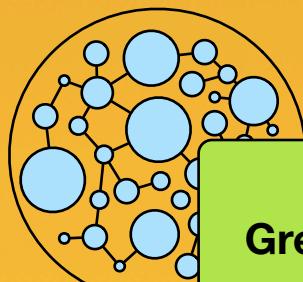
**SPARQL**



**SQL**



**neo4j**  
Cypher



Bytecode

Gremlin Traversal Machine

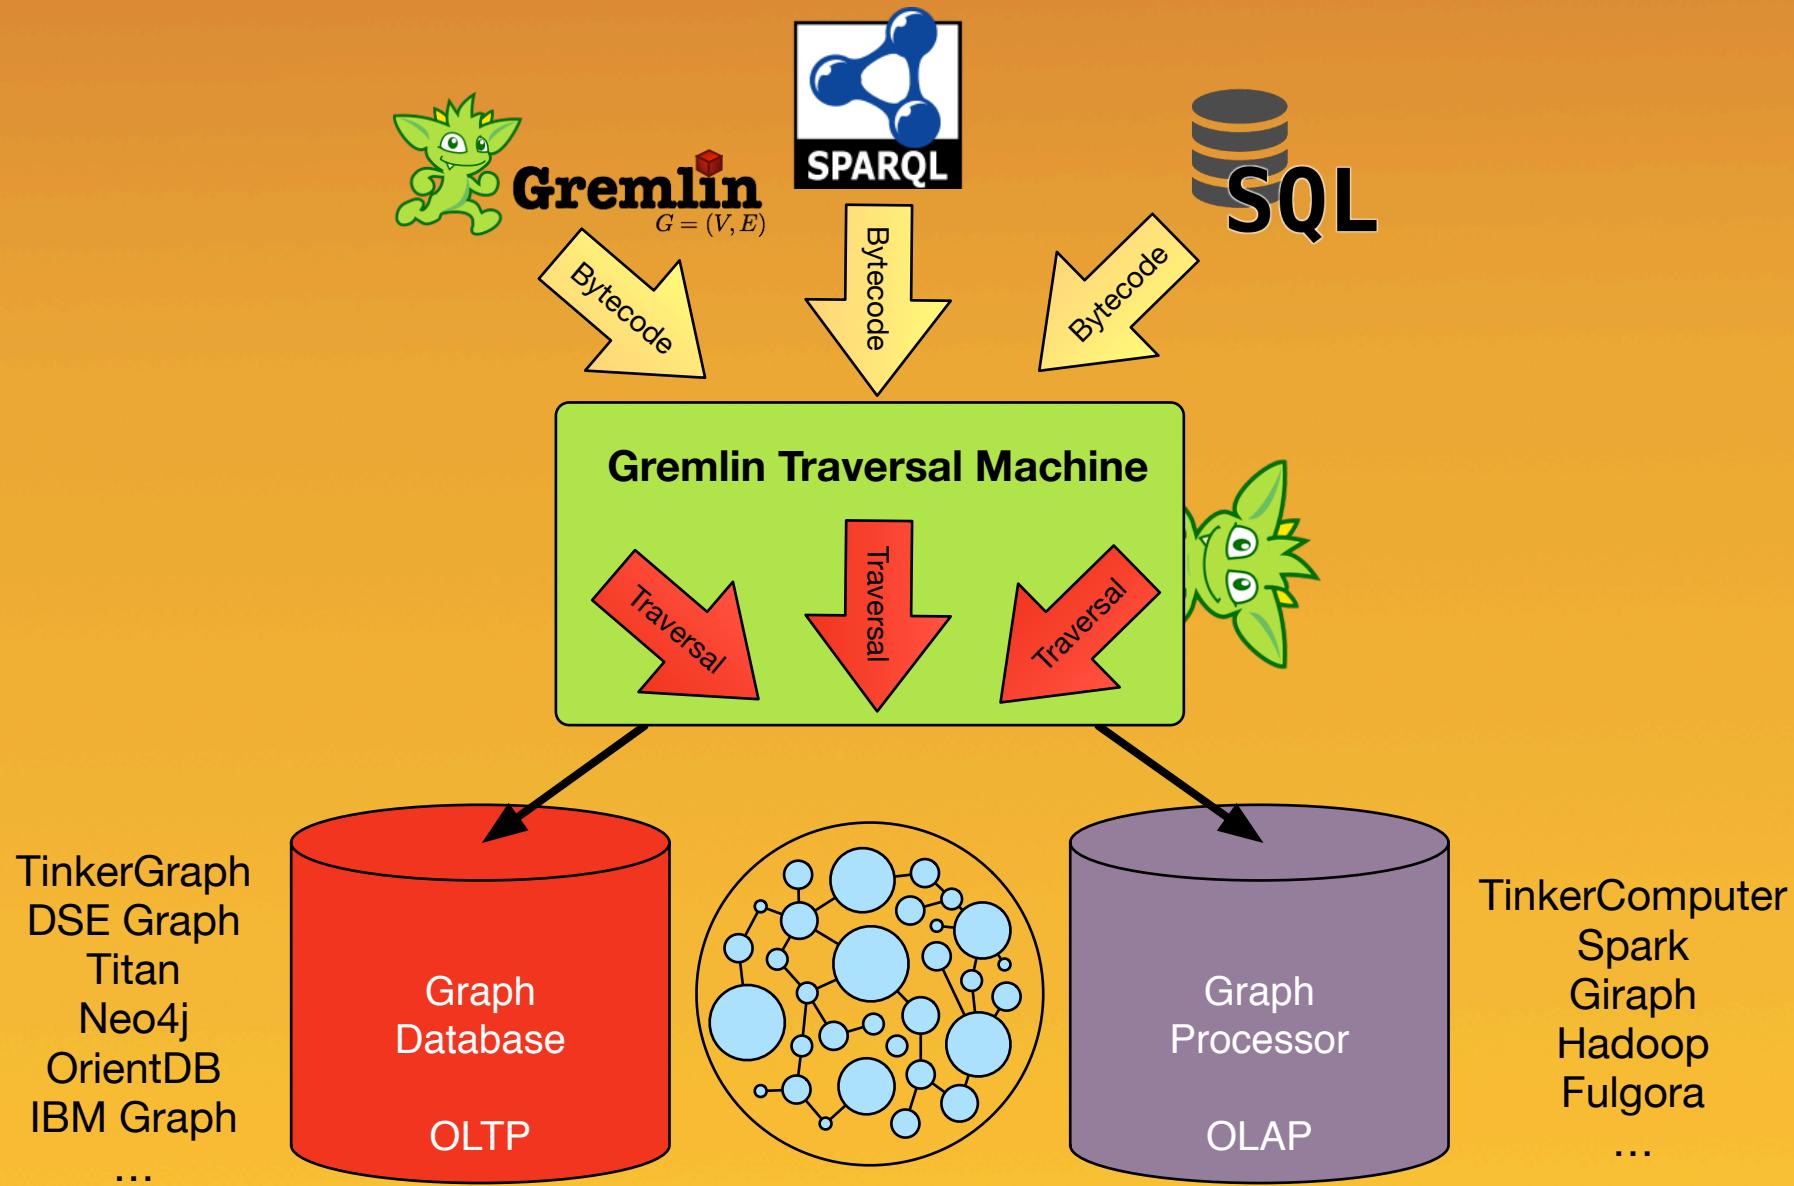


Distinct query languages (not only Gremlin language variants) can generate bytecode for evaluation by any OLTP/OLAP TinkerPop-enabled graph system.

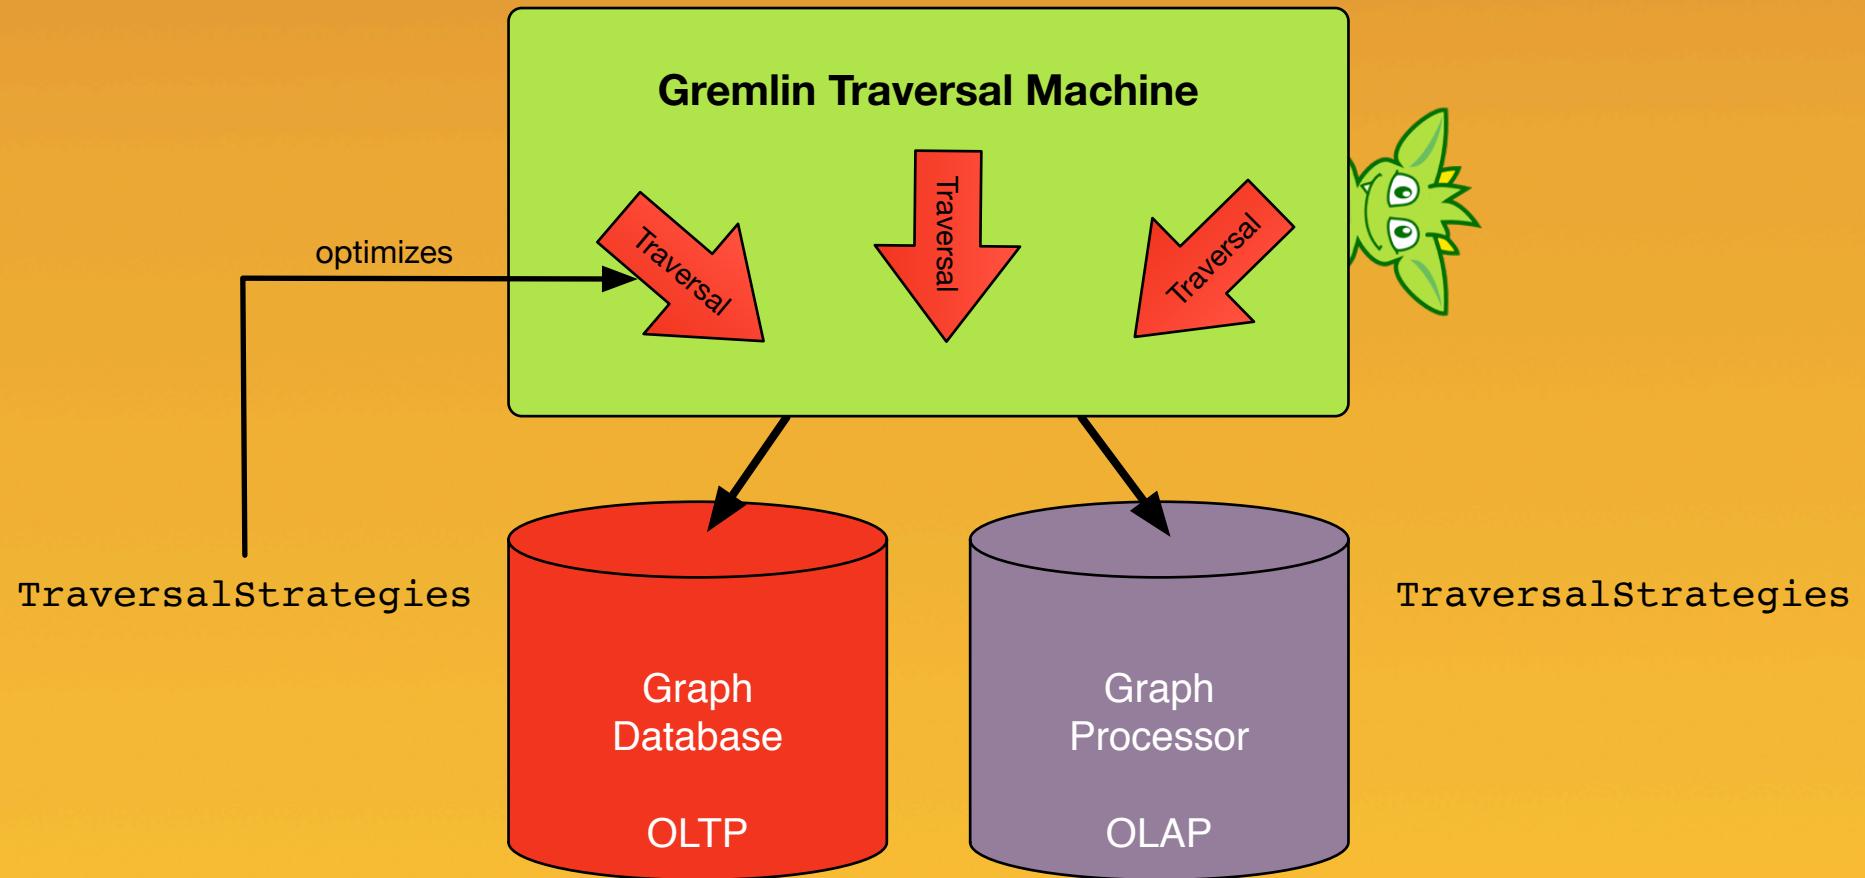


```
SELECT DISTINCT ?c
WHERE {
  ?a v:label "person" .
  ?a e:created ?b .
  ?b v:name ?c .
  ?a v:age ?d .
  FILTER (?d > 30)
}
```

```
[           ↓
[V],          Bytecode
[match,
 [[as, a], [hasLabel, person]],
 [[as, a], [out, created], [as, b]],
 [[as, a], [has, age, gt(30)]]],
 [select, b],
 [by, name],
 [dedup]
]
```



Gremlin traversals can be executed against OLTP graph databases and OLAP graph processors. That means that if, e.g., SPARQL compiles to bytecode, it can execute both OLTP and OLAP.



Graph system providers (OLTP/OLAP) typically have custom strategies registered with the Gremlin traversal machine that take advantage of unique, provider-specific features.

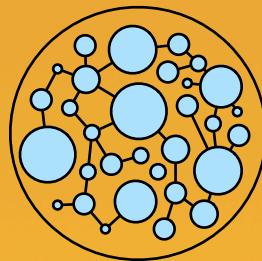
```
g.V().has("name", "gremlin").  
out("knows").values("age").  
groupCount()
```



DataStax  
Enterprise Graph



one graph to many vertices  
(flatMap)



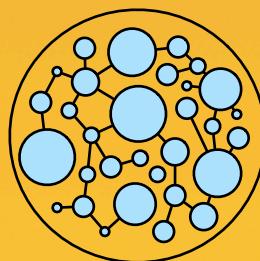
one vertex to that vertex or no vertex  
(filter)



GraphStepStrategy



one graph to many vertices using index lookup  
(flatMap)

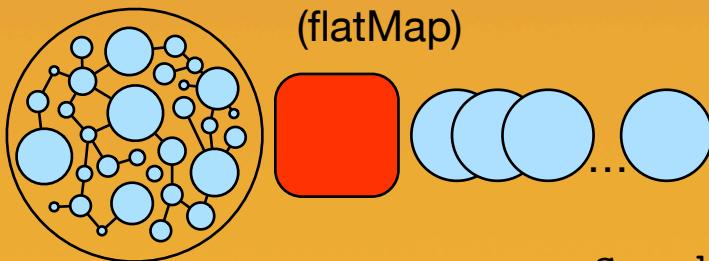


Most OLTP graph systems have a traversal strategy that combines [V,has\*]-sequences into a single global index-based flatMap-step.

`g.V().count()`



one graph to many vertices  
(flatMap)



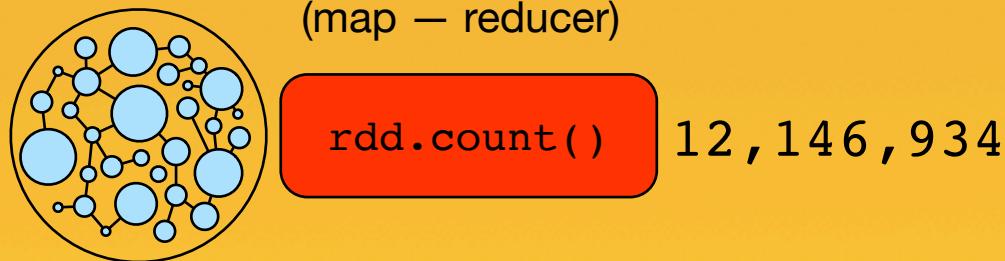
many vertices to long  
(map – reducer)



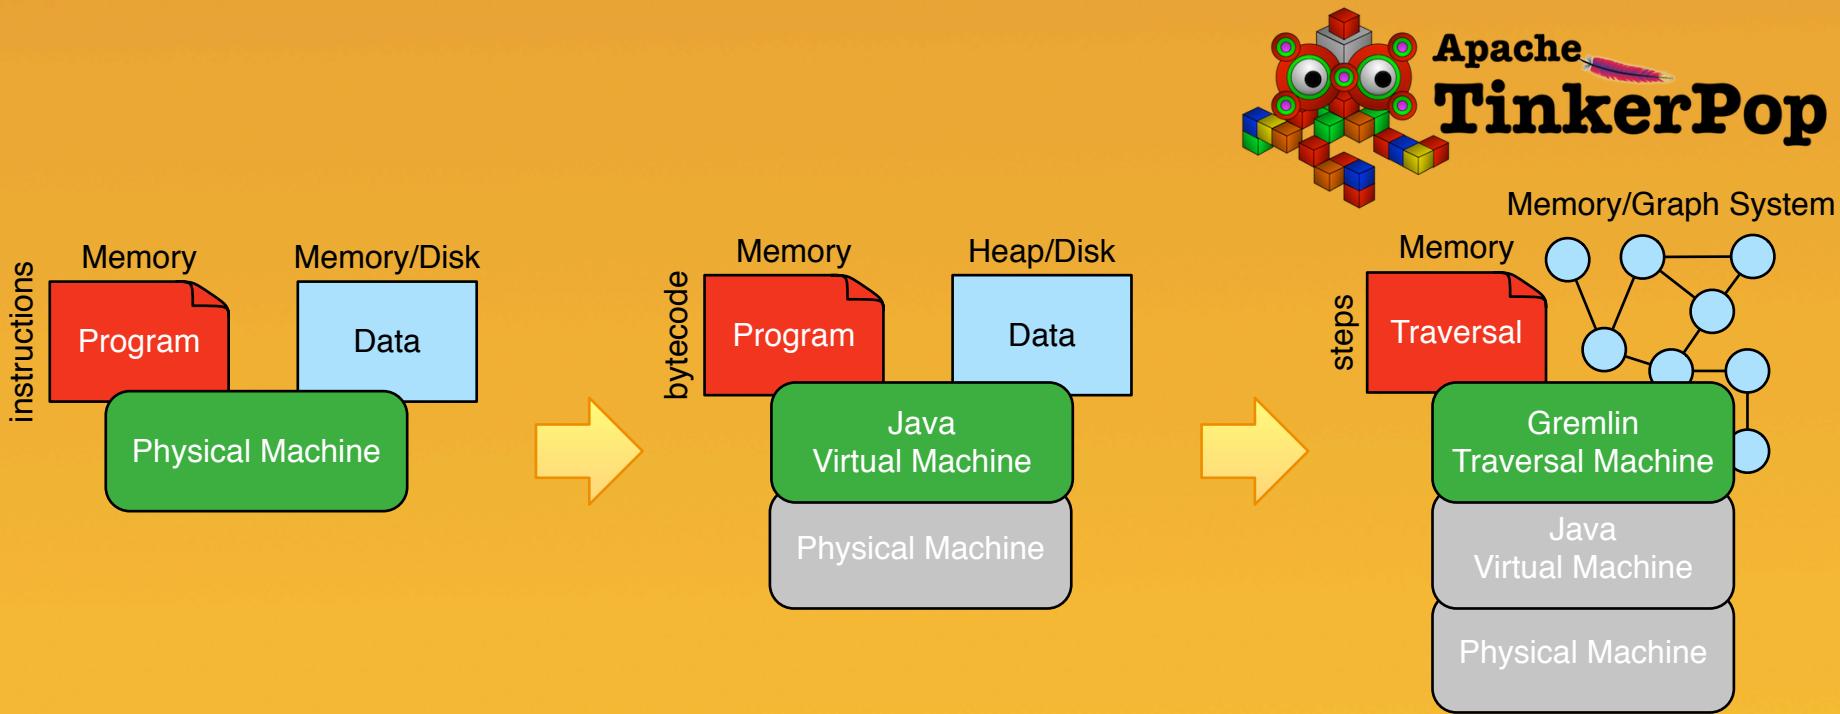
`SparkInterceptorStrategy`



one graph to long  
(map – reducer)



Most OLAP graph systems have a traversal strategy that bypasses Traversal semantics and implements reducers using the native API of the system.



From the physical computing machine to the Gremlin traversal machine.



# Stakeholders

**Application Developers**

**Language Providers**

Gremlin Language Variant  
Distinct Query Language

**Graph System Providers**

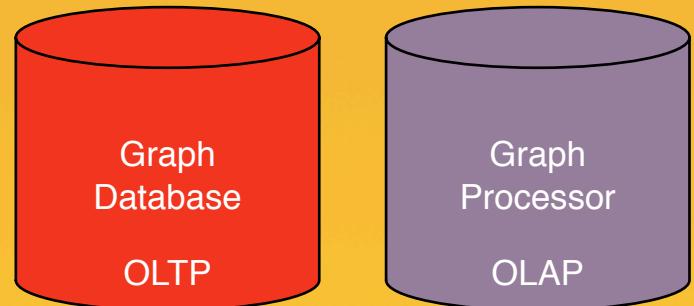
OLTP Provider  
OLAP Provider



# Stakeholders

## Application Developers

- + One query language for all OLTP/OLAP systems.



*Real-time and analytic queries are represented in Gremlin.*



# Stakeholders

## Application Developers



*Apache TinkerPop as the JDBC for graphs.*



# Stakeholders

## Application Developers

- + One query language for all OLTP/OLAP systems.
- + No vendor lock-in.
- + Gremlin is embedded in the developer's language.

SQL in Java

```
ResultSet result = statement.executeQuery(  
    "SELECT name FROM People \n" +  
    " ORDER BY age \n" +  
    " LIMIT 10")
```

vs.

Gremlin-Java

```
Iterator<String> result =  
g.V().hasLabel("person").  
order().by("age").  
limit(10).values("name")
```

No “fat strings.” The developer writes their graph database/processor queries in their native programming language.



# Stakeholders

## Language Providers

Gremlin Language Variant  
Distinct Query Language



Easy to generate bytecode.

```
GraphTraversal.getMethods()
    .findAll { GraphTraversal.class == it.returnType }
    .collect { it.name }
    .unique()
    .each {
        pythonClass.append(
            """ def ${it}(self, *args):
                self.bytecode.add_step("${it}", *args)
                return self
            """
        )
    }
}
```

*Gremlin-Python's source code is  
programmatically generated using Java reflection.*



# Stakeholders

## Language Providers

Gremlin Language Variant  
Distinct Query Language

 Easy to generate bytecode.

 Bytecode executes against  
TinkerPop-enabled systems.



*Language providers write a translator for the Gremlin traversal machine,  
not a particular graph database/processor.*

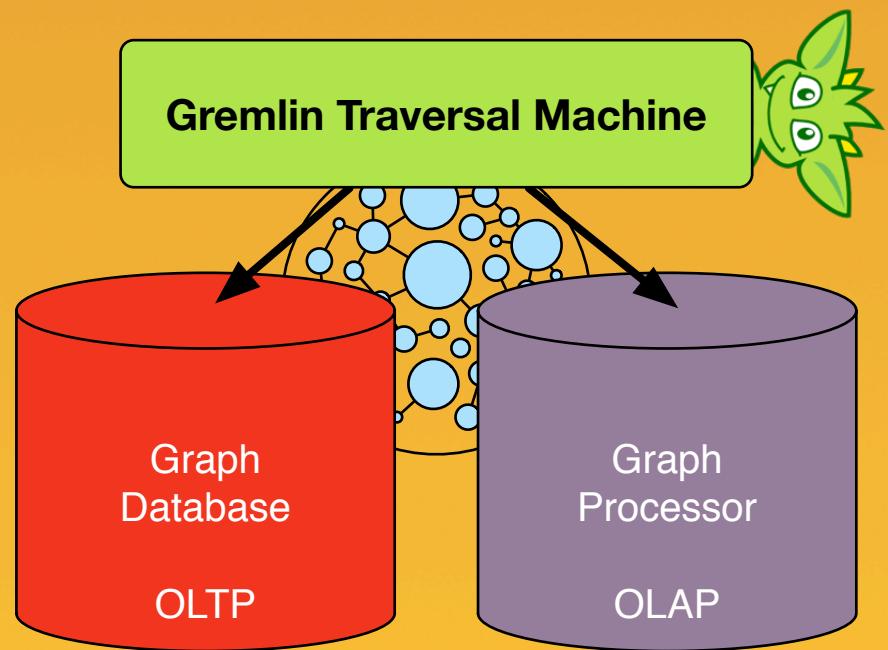


# Stakeholders

## Language Providers

Gremlin Language Variant  
Distinct Query Language

- + Easy to generate bytecode.
- + Bytecode executes against TinkerPop-enabled systems.
- + Provider can focus on design, not evaluation.



*The language designer does not have to concern themselves with OLTP or OLAP execution. They simply generate bytecode and the Gremlin traversal machine handles the rest.*

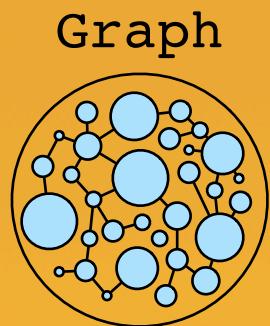


# Stakeholders

## Graph System Providers

OLTP Provider  
OLAP Provider

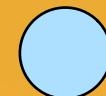
- + Easy to implement core interfaces.



Graph



Vertex



Transaction



Property  
key=value

TraversalStrategy





# Stakeholders

## Graph System Providers

OLTP Provider  
OLAP Provider

 Easy to implement core interfaces.

 Provider supports all provided languages.



*The provider automatically supports all query languages that have compilers that generate Gremlin bytecode.*



# Stakeholders

## Graph System Providers

OLTP Provider  
OLAP Provider

- + Easy to implement core interfaces.
- + Provider supports all provided languages.
- + OLTP providers can leverage existing OLAP systems.



DataStax  
Enterprise Graph

Spark

*DSE Graph leverages SparkGraphComputer for OLAP processing.*



# Stakeholders

## Application Developers

## Language Providers

## Graph System Providers

Gremlin Language Variant  
Distinct Query Language

OLTP Provider  
OLAP Provider

- |  |  |  |
|--|--|--|
|  One query language for all OLTP/OLAP systems.      |  Easy to generate bytecode.                           |  Easy to implement core interfaces.                   |
|  No vendor lock-in.                                 |  Bytecode executes against TinkerPop-enabled systems. |  Provider supports all provided languages.            |
|  Gremlin is embedded in the developer's language. |  Provider can focus on design, not evaluation.      |  OLTP providers can leverage existing OLAP systems. |

# Thank you.



<http://tinkerpop.apache.org>

<http://www.datastax.com/products/datastax-enterprise-graph>