

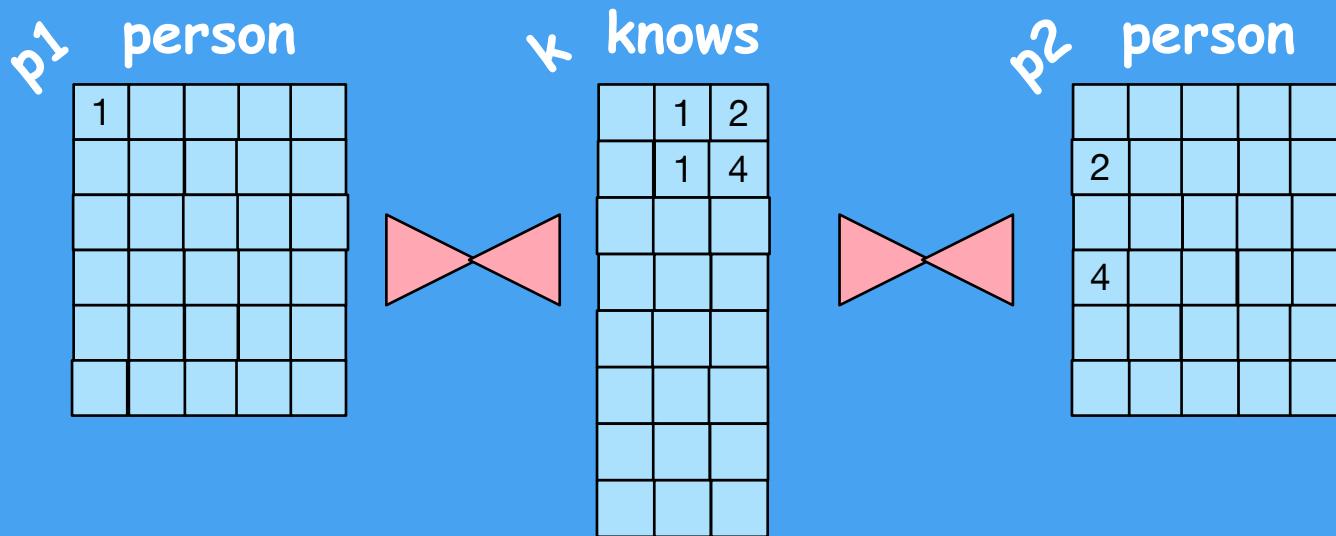
# Gremlin 101.3 on your FM Dial



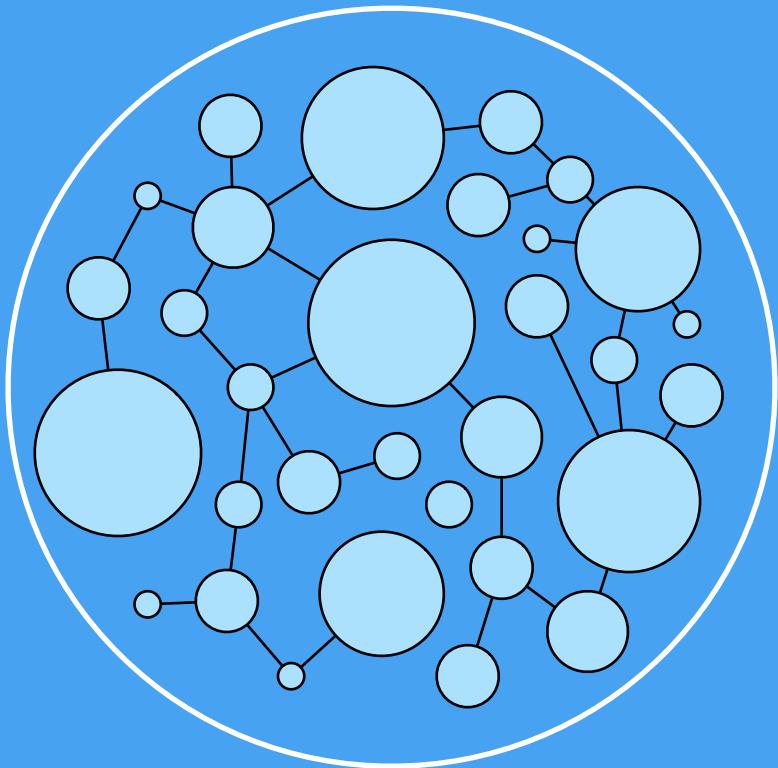
Marko A. Rodriguez :•:  
**DATASTAX**•:

Who are person 1's friends?

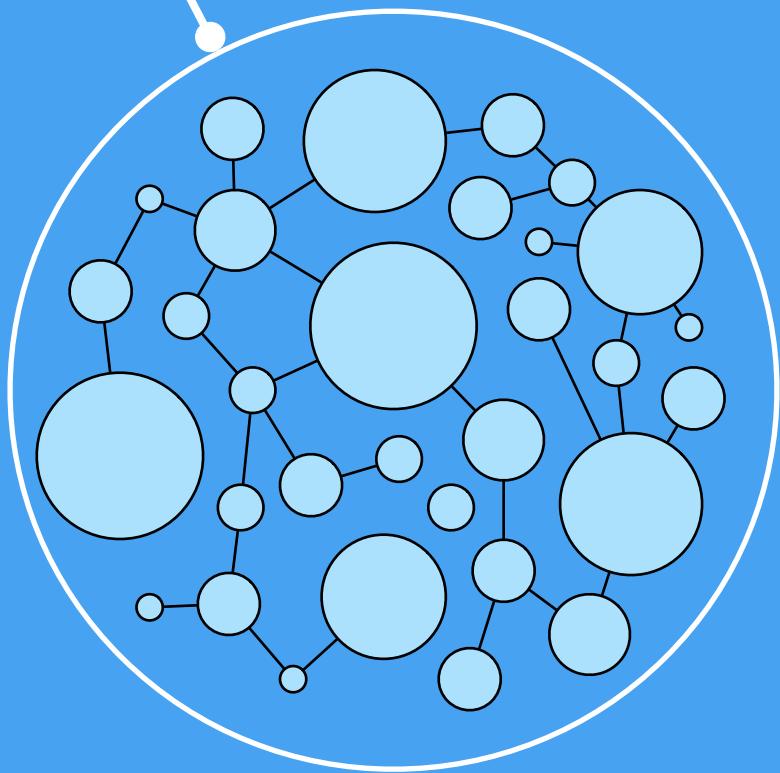
```
SELECT p2.*  
FROM person p1  
INNER JOIN knows k  
    ON k.out = p1.id  
INNER JOIN person p2  
    ON p2.id = k.in  
WHERE p1.id = 1
```



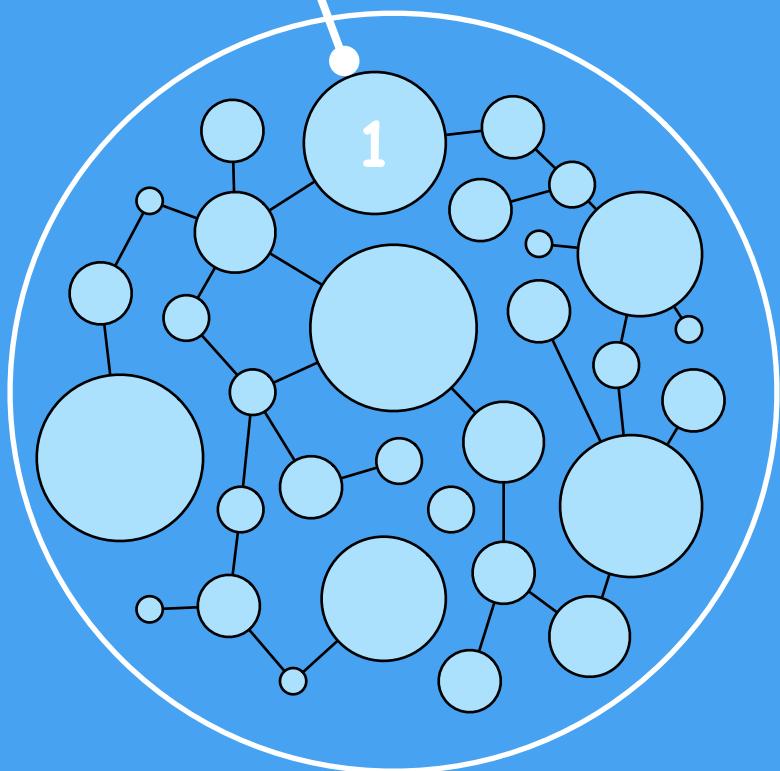
`g.V(1).out('knows')`



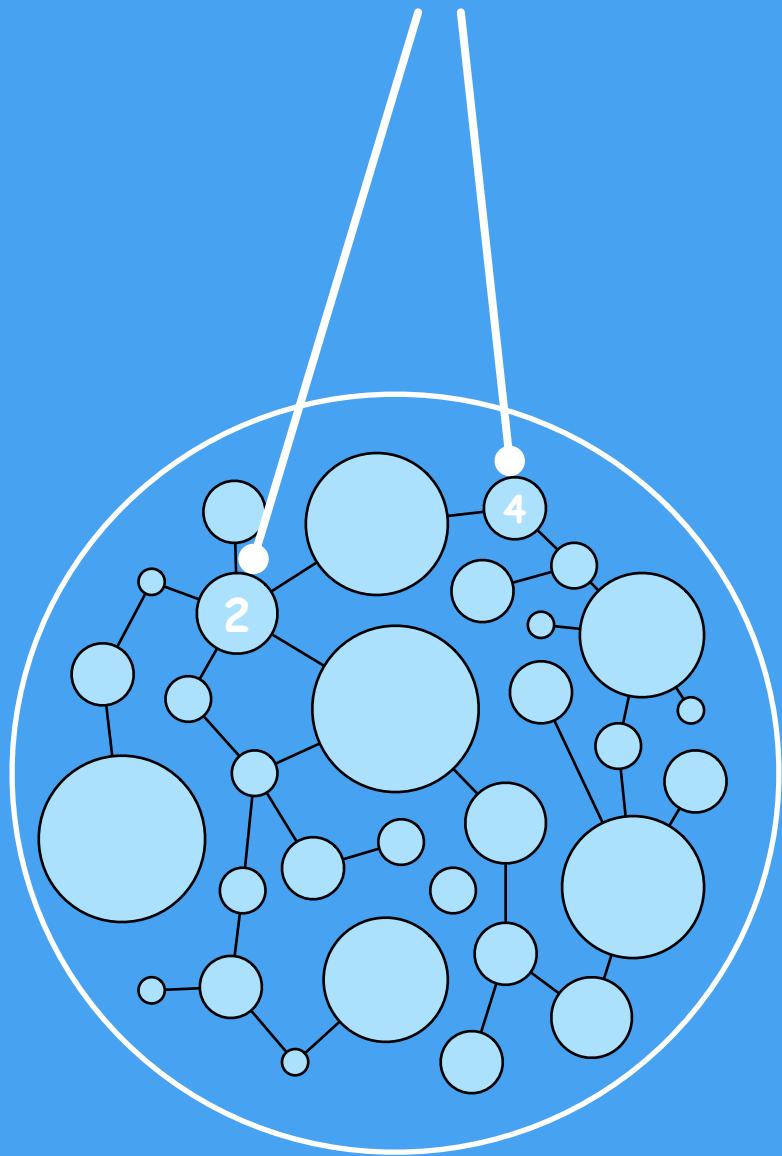
*g*



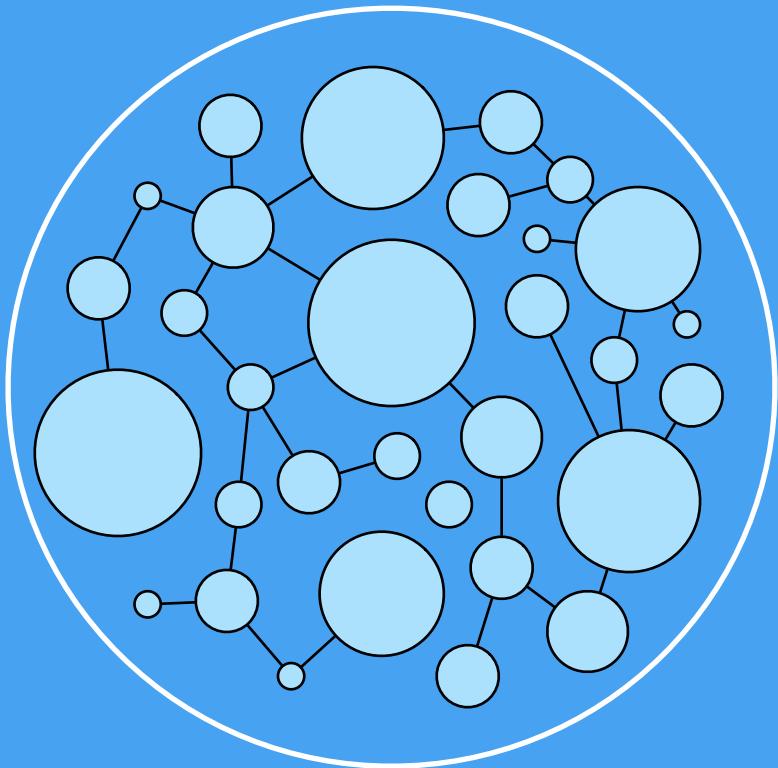
$g.V(1)$



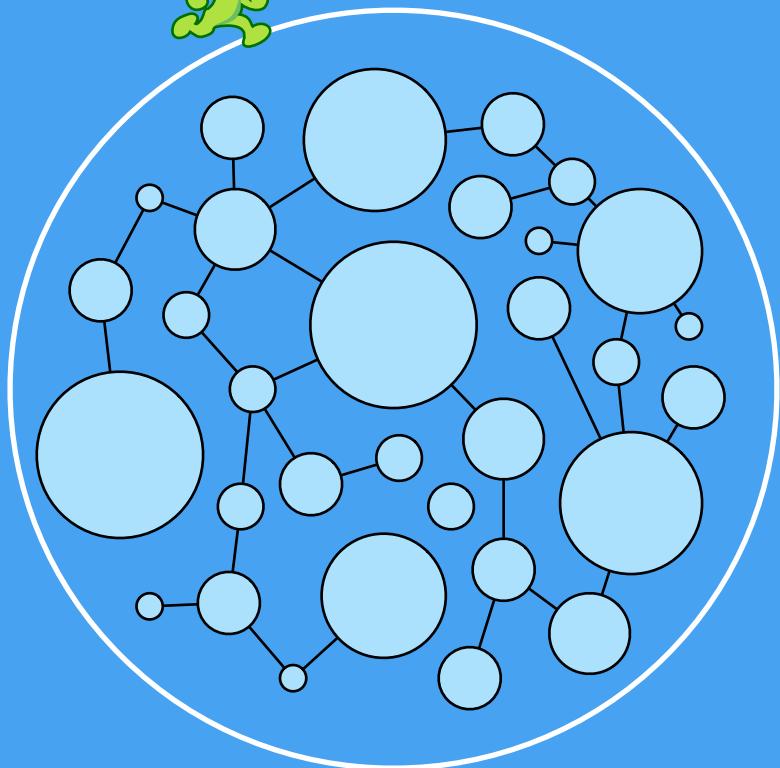
`g.V(1).out('knows')`



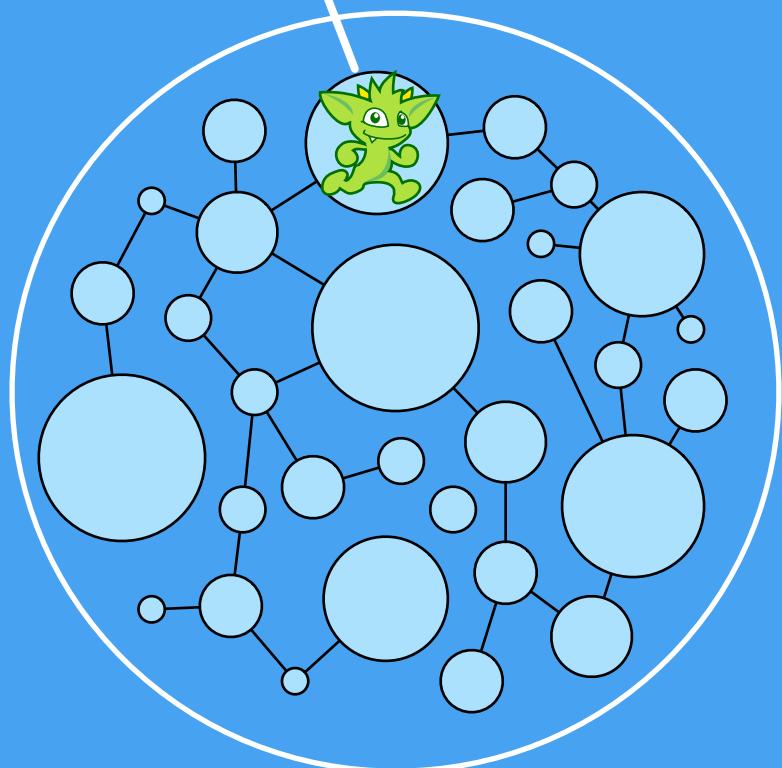
`g.V(1).out('knows')`



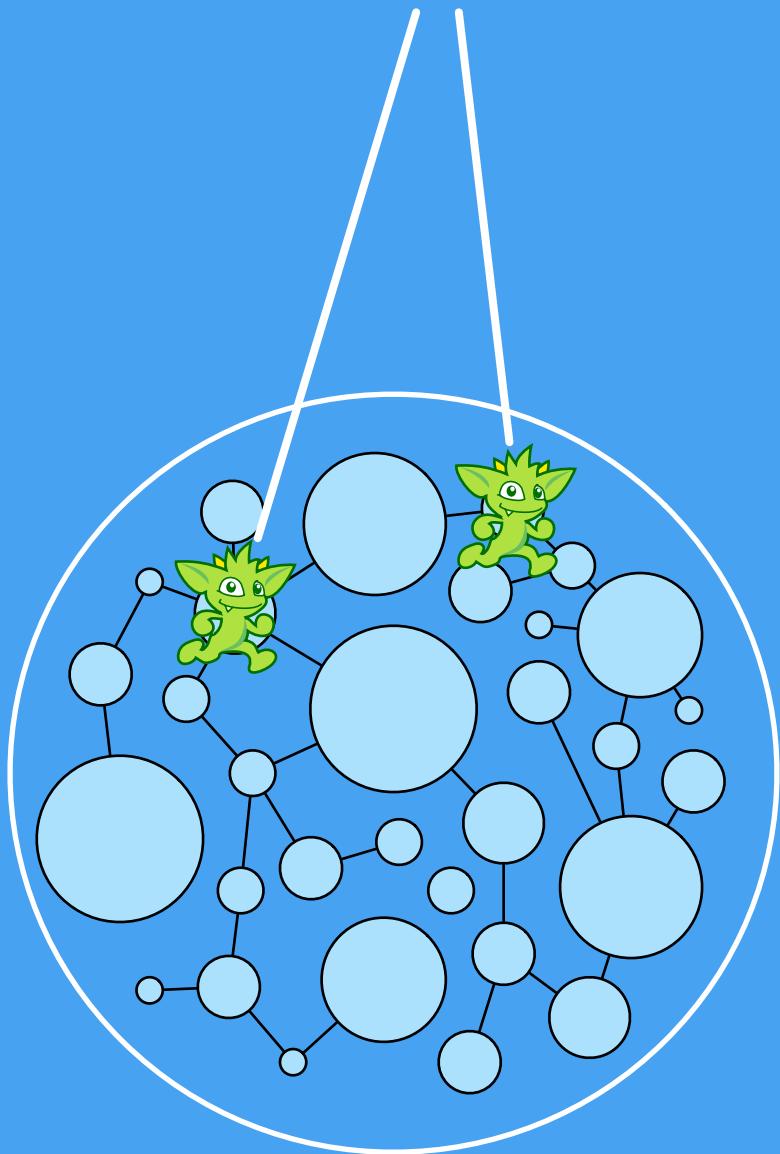
*g*



$g.V(1)$



`g.V(1).out('knows')`



$$G \xleftarrow{\mu} t \in T \xrightarrow{\psi} \Psi$$

Rodriguez, M.A., "The Gremlin Graph Traversal Machine and Language," ACM Proceedings of the 15th Symposium on Database Programming Languages, 2015.

<http://arxiv.org/abs/1508.03843>

**Traversers**

$$G \xleftarrow{\mu} t \in T \psi \xrightarrow{\Psi} \text{Traversal}$$

**Graph**

# Traversers

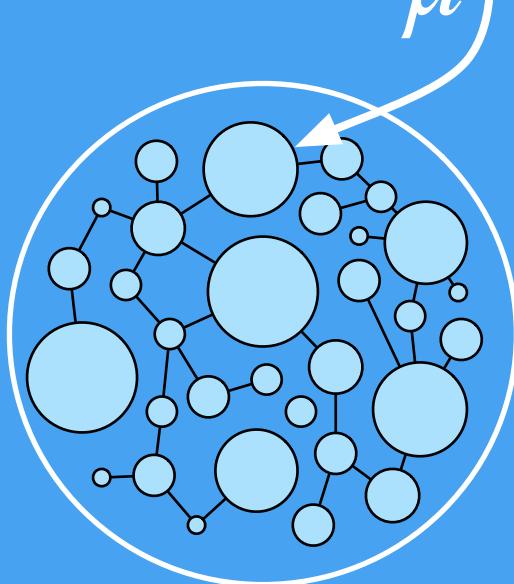
$G \xleftarrow[\text{object}]{\mu} t \in T \xrightarrow[\text{traverser}]{\psi} \Psi$   $\xrightarrow[\text{step}]{\Psi}$  Traversal

$\Psi$   
Traversal  
Steps, SideEffects

$T$   
Traversers  
Locations, Bulk, Sack, Path, Loops

$G$   
Graph  
Vertices, Edges, Properties

`g.V(1).out('knows')`



$\Psi$

The Traversal

```
g.V(1).out('knows')
```

```
g.V(1).out('knows')
```

compiles to

GraphStep

VertexStep

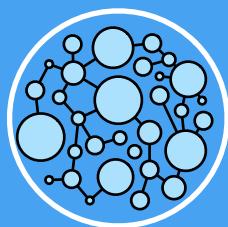
# `g.V(1).out('knows')`

**compiles to**

GraphStep

VertexStep

**executes as**



GraphStep

1

VertexStep

2

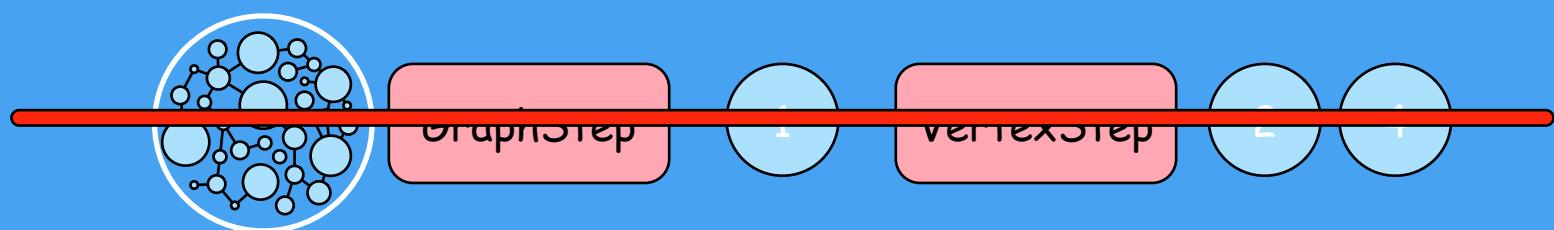
4

# `g.V(1).out('knows')`

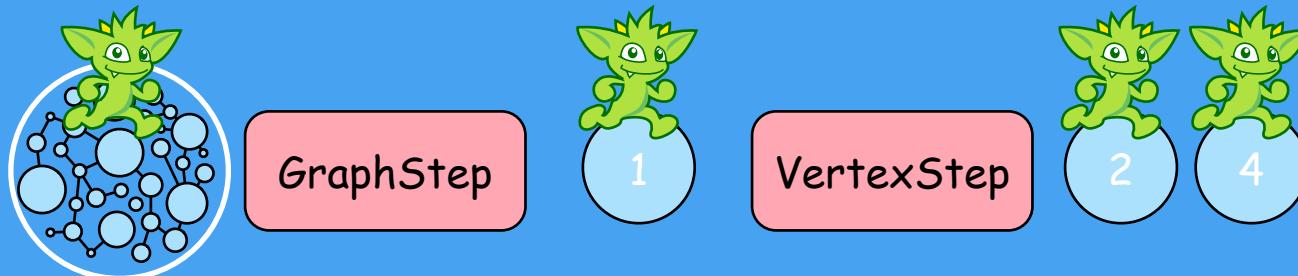
compiles to

GraphStep      VertexStep

executes as

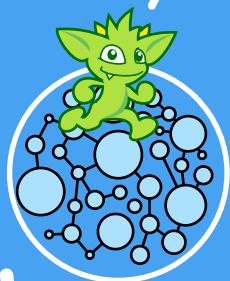


executes as



`g.V(1).out('knows')`

$\psi$



GraphStep

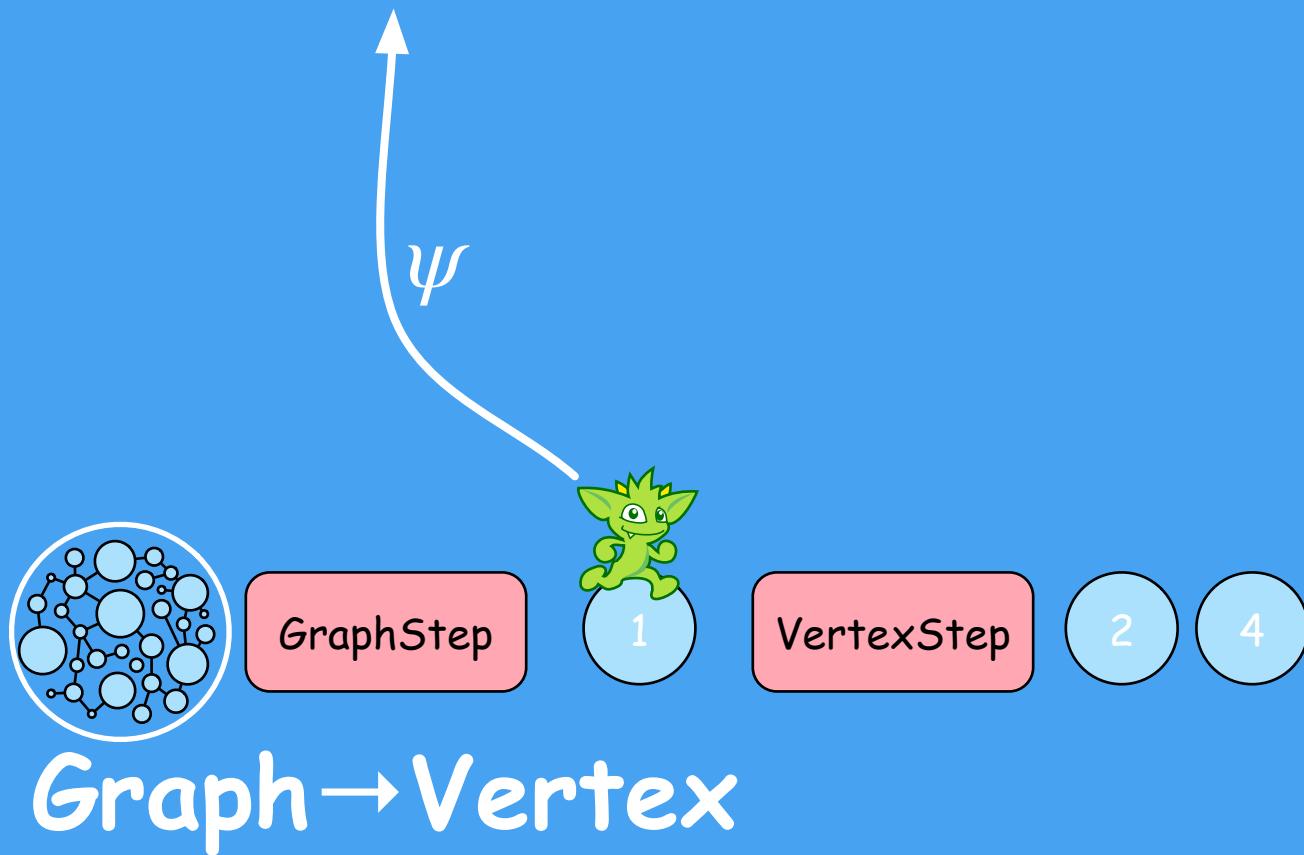
1

VertexStep

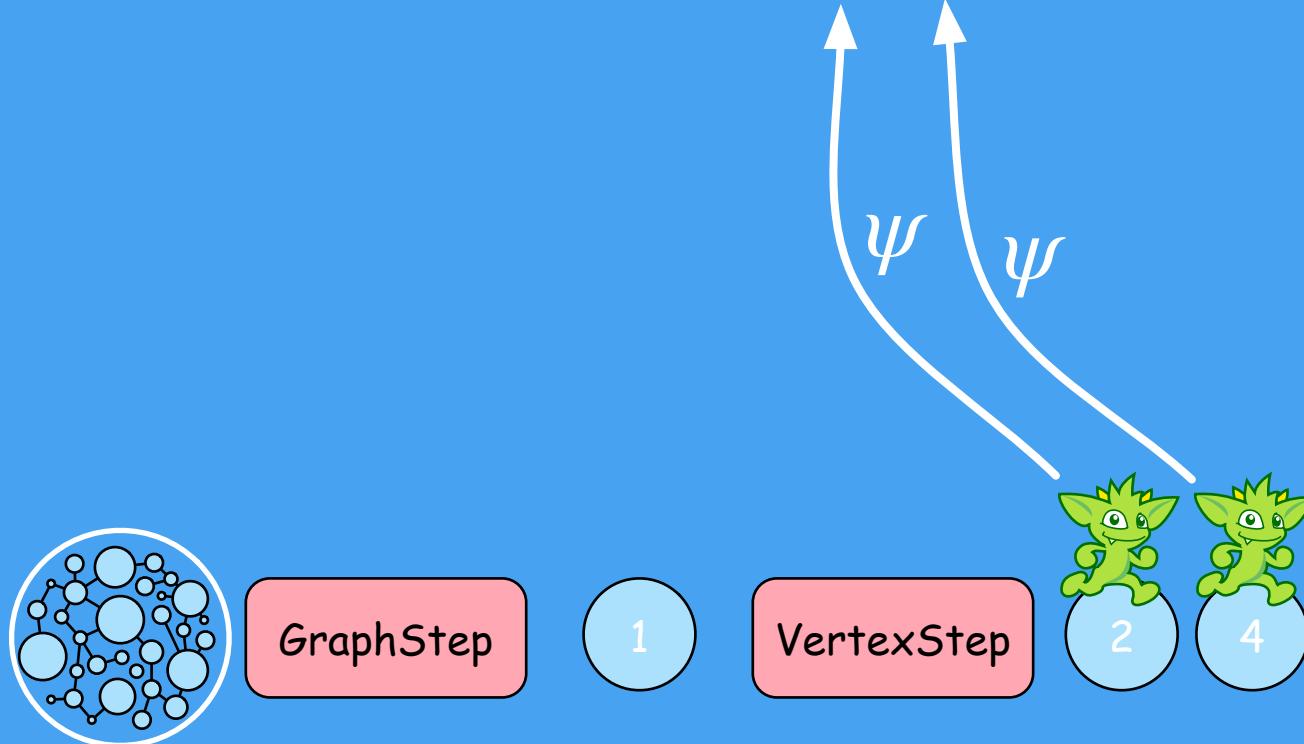
2 4

→Graph

`g.V(1).out('knows')`



`g.V(1).out('knows')`



**Vertex → Vertex**

```
g.V(1).out('knows')
```

$V_1: G \rightarrow V$

$\text{out}_{\text{knows}}: V \rightarrow V$

```
g.V(1).out('knows')
```

$V_1: G \rightarrow V^*$

$\text{out}_{\text{knows}}: V \rightarrow V^*$

`g.V(1).out('knows')`

   
 $V_1: G \rightarrow V^*$

   
 $\text{out}_{\text{knows}}: V \rightarrow V^*$

`g.V(1).out('knows')`

$V_1: G \rightarrow V^*$

$V_1(\text{green goblin} \circlearrowleft \text{blue network})$

$\text{out}_{\text{knows}}: V \rightarrow V^*$

`g.V(1).out('knows')`

   
 $V_1: G \rightarrow V^*$



   
 $\text{out}_{\text{knows}}: V \rightarrow V^*$

`g.V(1).out('knows')`

   
 $V_1: G \rightarrow V^*$

   
 $\text{out}_{\text{knows}}: V \rightarrow V^*$        $\text{out}_{\text{knows}}($    $)$

`g.V(1).out('knows')`

$V_1: G \rightarrow V^*$

$\text{out}_{\text{knows}}: V \rightarrow V^*$



```
g.V(1).out('knows')
```

$$\text{out}_{\text{knows}}(V_1(\text{ }) \text{ }) = \text{ } \text{ }$$

What are the names of the people that person 1 knows who created LinkedProcess (lop)?

```
g.V(1).out('knows').where(out('created').has('name', 'lop')).values('name')
```

What are the names of the people that person 1 knows who created LinkedProcess (lop)?

```
g.V(1).out('knows').
```

```
  where(out('created').has('name','lop'))  
  values('name')
```

global scope

locally-global scope

`g.V(1).out('knows').`

`where(out('created')).has('name', 'lop'))`

`values('name')`

local scope



```
g.V(1).out('knows').
```

```
    where(out('created').has('name', 'lop'))
```

```
    values('name')
```

global scope



```
g.V(1).out('knows').
```

```
  where(out('created').has('name', 'lop'))
```

```
    values('name')
```

global scope

```
g.V(1).out('knows').
```



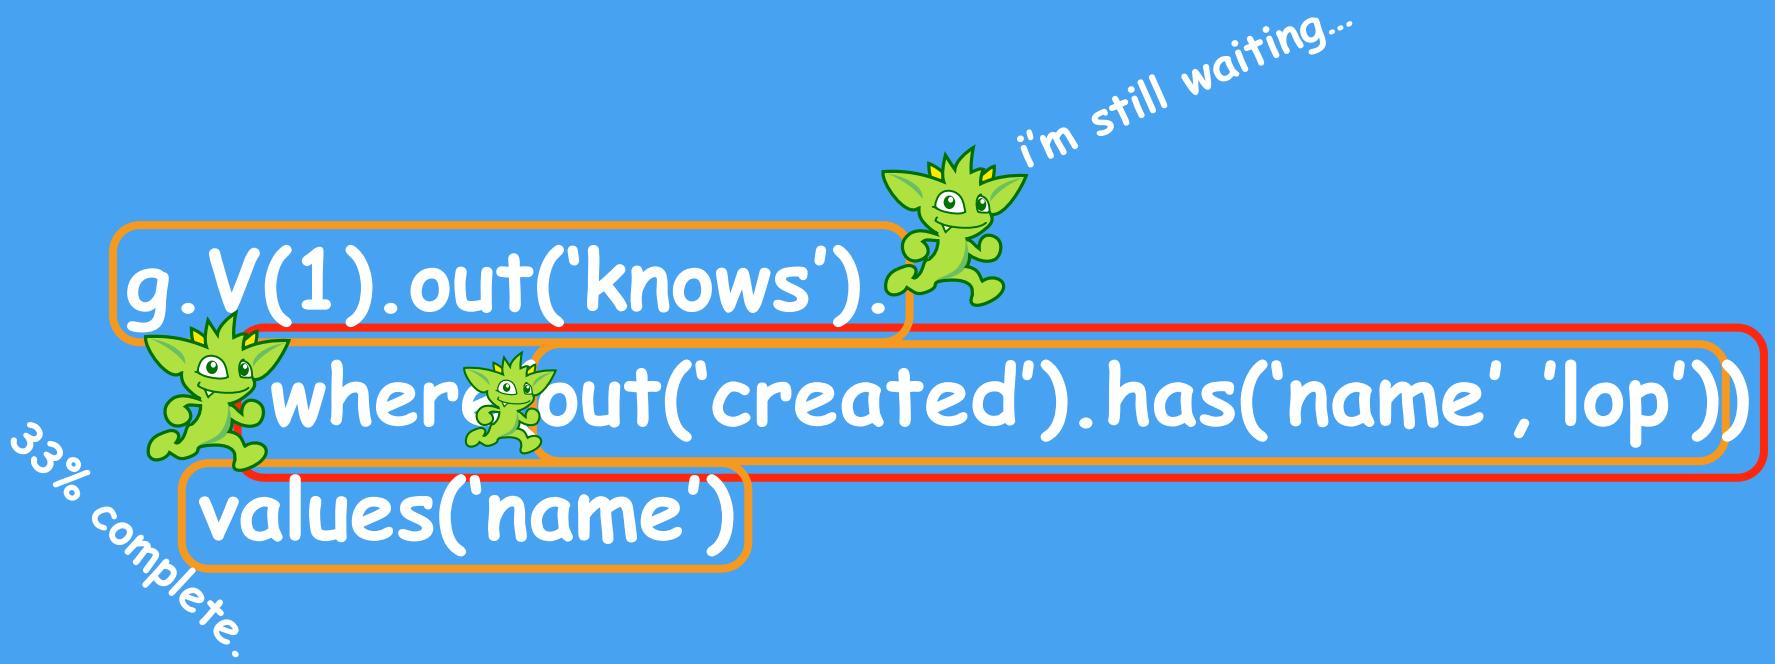
i'm waiting...

```
where(out('created').has('name','lop'))
```

```
values('name')
```

*I hope the vertex I  
reference created lop!*

local scope



locally-global scope

```
g.V(1).out('knows').
```



i'm still waiting some more...

```
where out('created').has('name', 'lop'))  
values('name')
```



uh oh!

bummer!

locally-global scope

```
g.V(1).out('knows').
```



my turn yet?

```
where(out('created').has('name', 'lop'))
```

damn!

```
values('name')
```



local scope

```
g.V(1).out('knows').
```



```
where(out('created').has('name', 'lop'))
```

```
values('name')
```

*my turn finally!*

local scope

```
g.V(1).out('knows').
```



```
where g.out('created').has('name', 'lop'))
```

```
values('name')
```

prosper  
my child.

locally-global scope

```
g.V(1).out('knows').
```



```
where out('created') as('name', 'lop')
```



```
values('name')
```

*one successful  
offspring is all I need.*

locally-global scope

`g.V(1).out('knows').`

`where(out('created') has('name', 'lolo'))  
values('name')`

woo hoo!

boo yeah!

live and learn :(

locally-global scope

```
g.V(1).out('knows').
```

```
  where(out('created').has('name', 'lop'))
```

```
    values('name')
```



I reference a vertex  
that created lop!

local scope

```
g.V(1).out('knows').
```

```
    where(out('created').has('name', 'lop'))
```

```
    values('name')
```



global scope

```
g.V(1).out('knows').
```

```
    where(out('created').has('name', 'lop'))
```

```
    values('name')
```



I reference  
a result.

global scope

# Gremlin Language Traversal

```
g.V(1).out('knows').  
  where(out('created').has('name','lop'))  
  values('name')
```

traversal = op((object | traversal)\*)\*

# Gremlin Bytecode Traversal

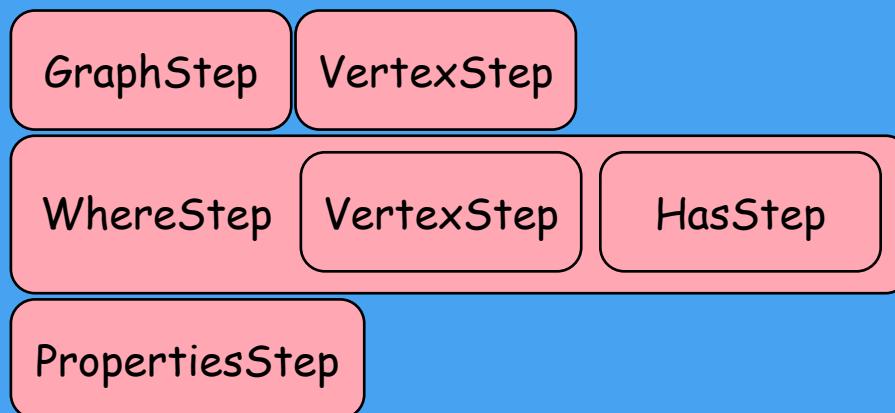
`[[V, 1][out, knows]`

`[where, [[out, created], [has, name, lop]]]`

`[values, name]]`

`traversal = [[op,(object | traversal)*]*]`

# Gremlin Machine Traversal



traversal = step[(object | traversal)\*]\*



language

```
g.V().has('name', 'marko').  
repeat(outE().identity().inV()).times(2).  
name.groupCount()
```

translate

```
[[V],[has,name,eq(marko)],  
[repeat,[[outE],[identity],[inV]]],[times,2],  
[values,name],[groupCount]]
```



bytecode

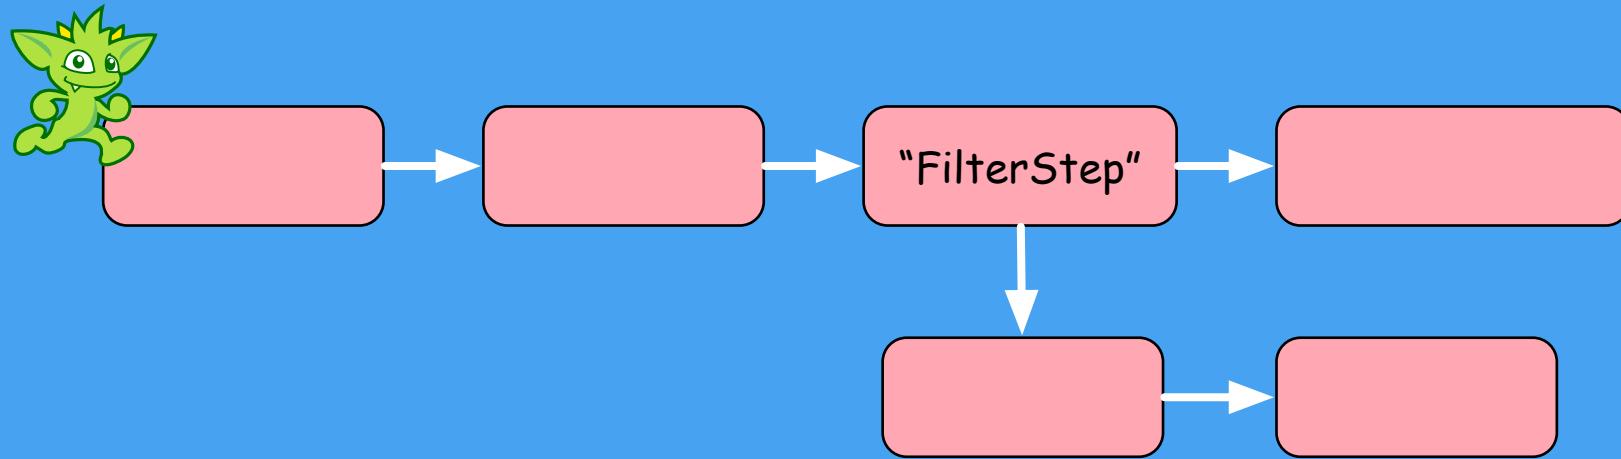
```
[GraphStep,HasStep(name,eq(marko)),  
RepeatStep([VertexStep(out,edges),IdentityStep,  
EdgeVertexStep(in)],2),  
PropertiesStep(values,name),GroupCountStep]
```

compile

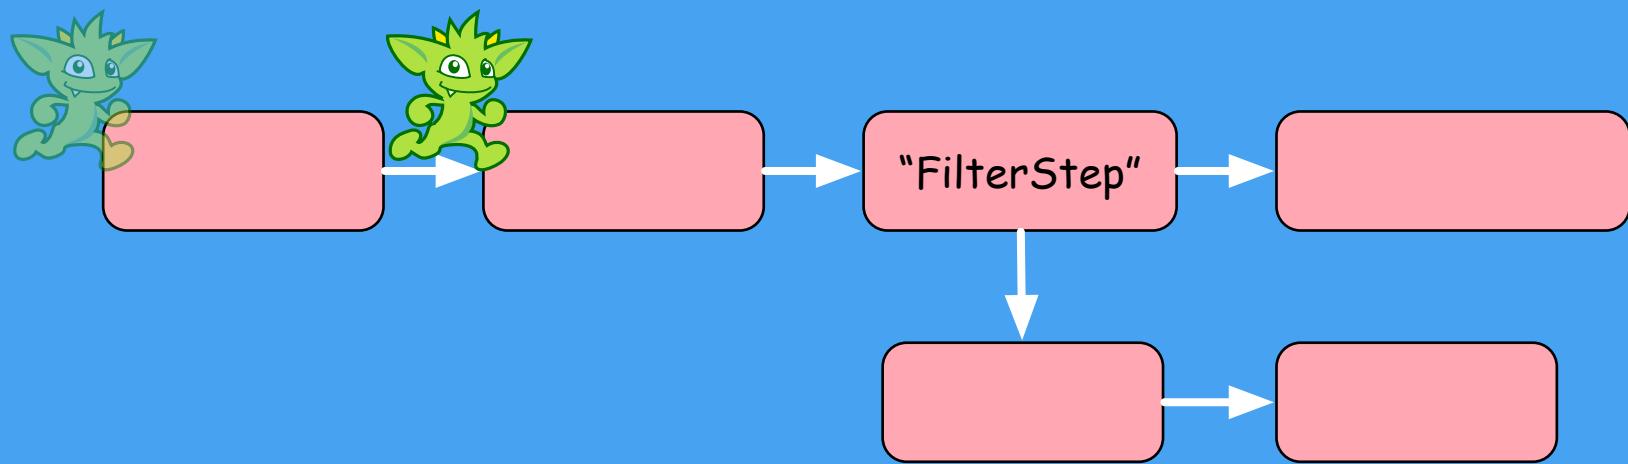
```
[ProviderGraphStep(name,eq(marko)),  
VertexStep(out,vertices),NoOpBarrierStep,  
VertexStep(out,vertices),NoOpBarrierStep,  
PropertiesStep(values,name),GroupCountStep]
```

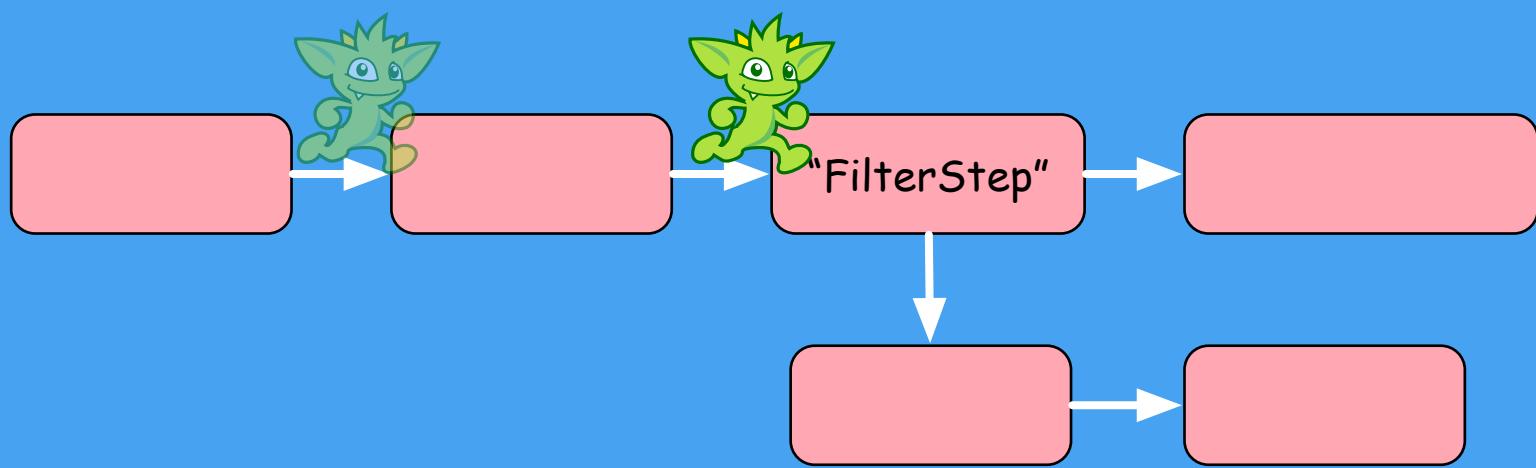
optimize

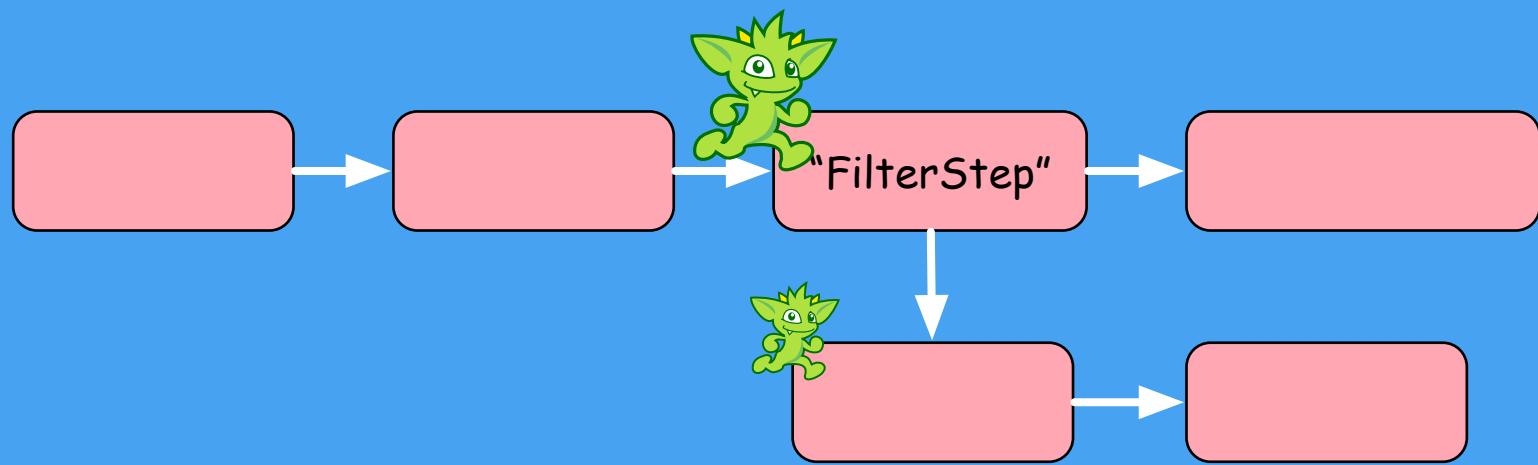


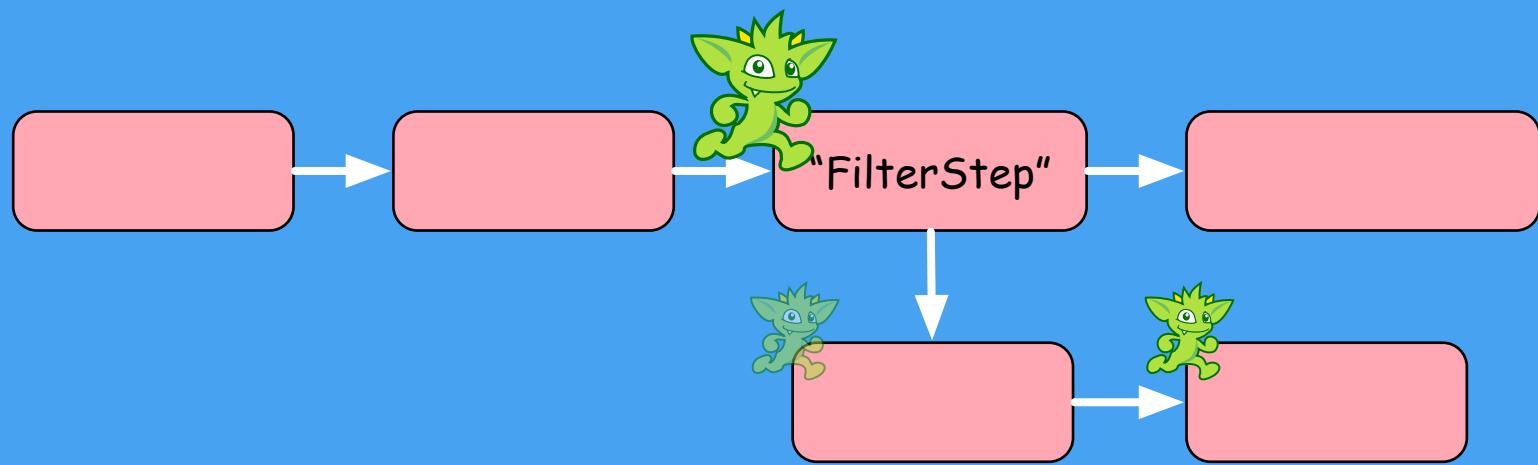


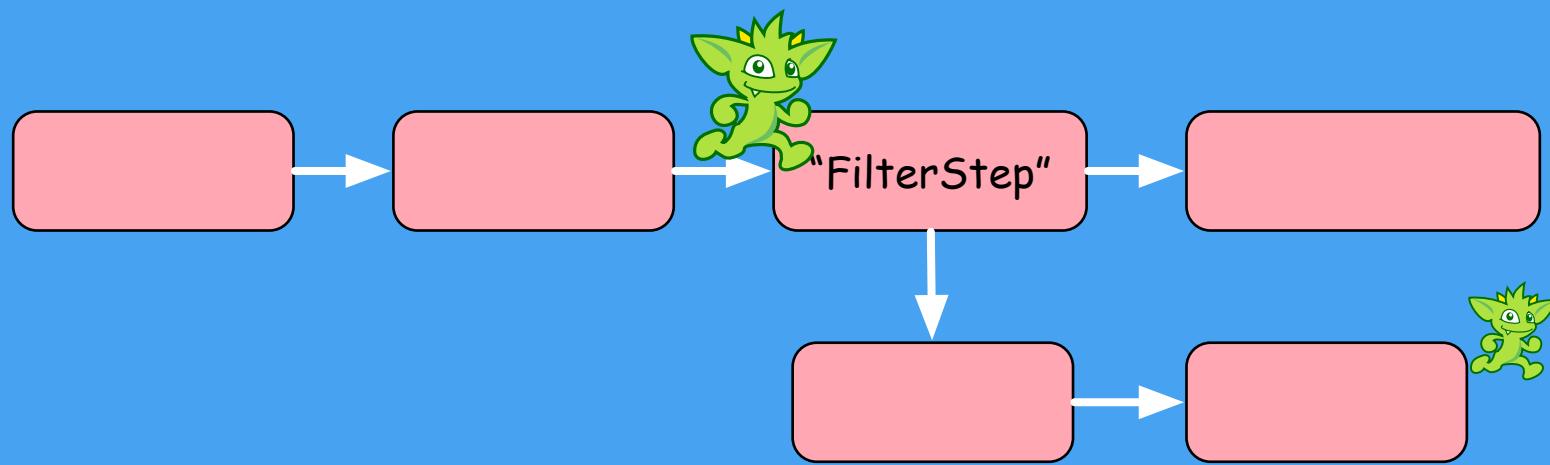
Rodriguez, M.A., "A Gremlin Implementation of the Gremlin Traversal Machine," DataStax Engineering Blog, October 2016.

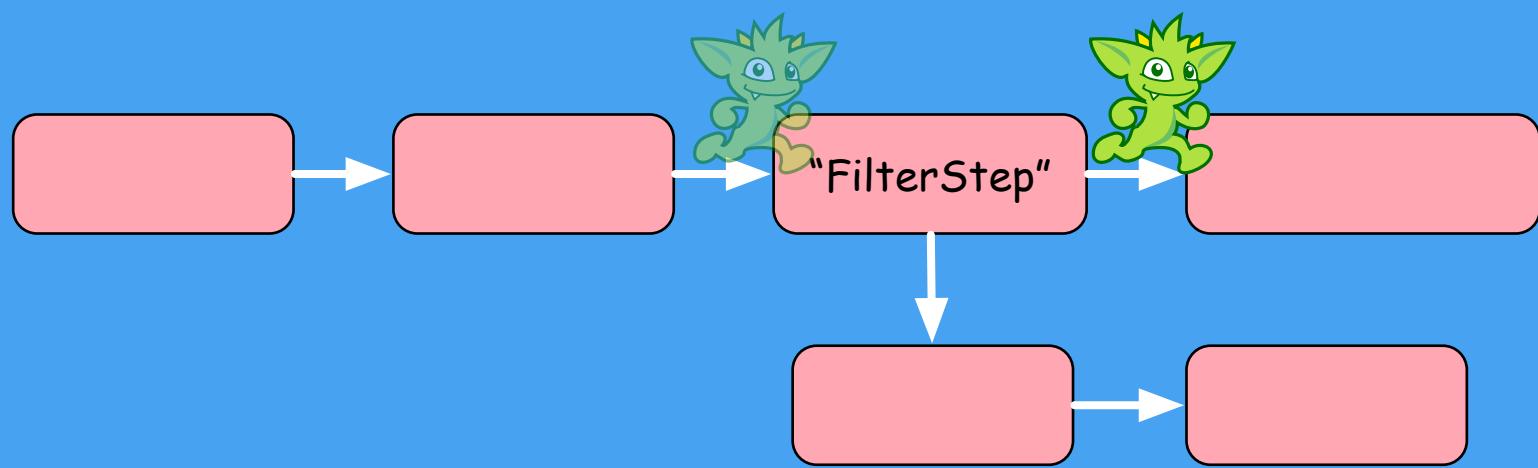


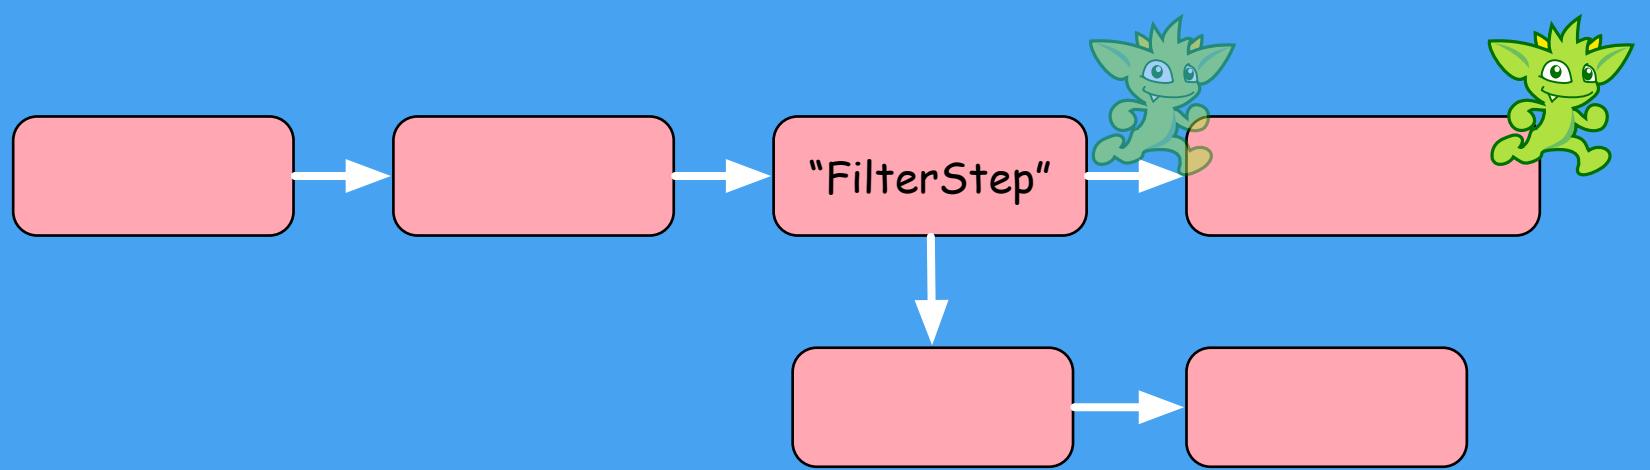


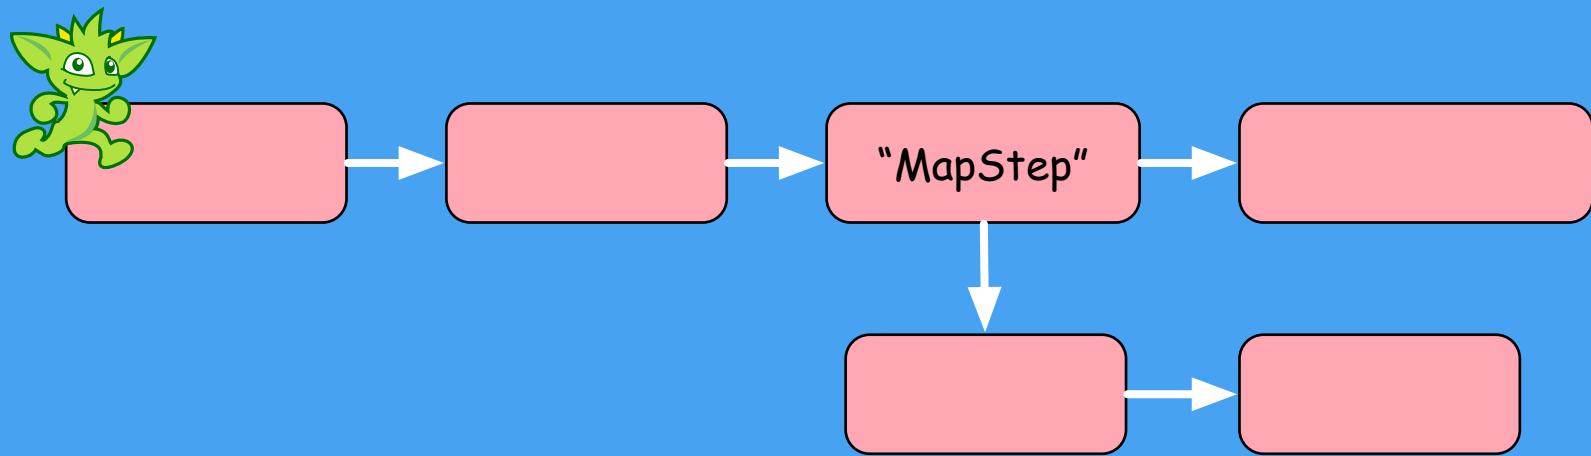


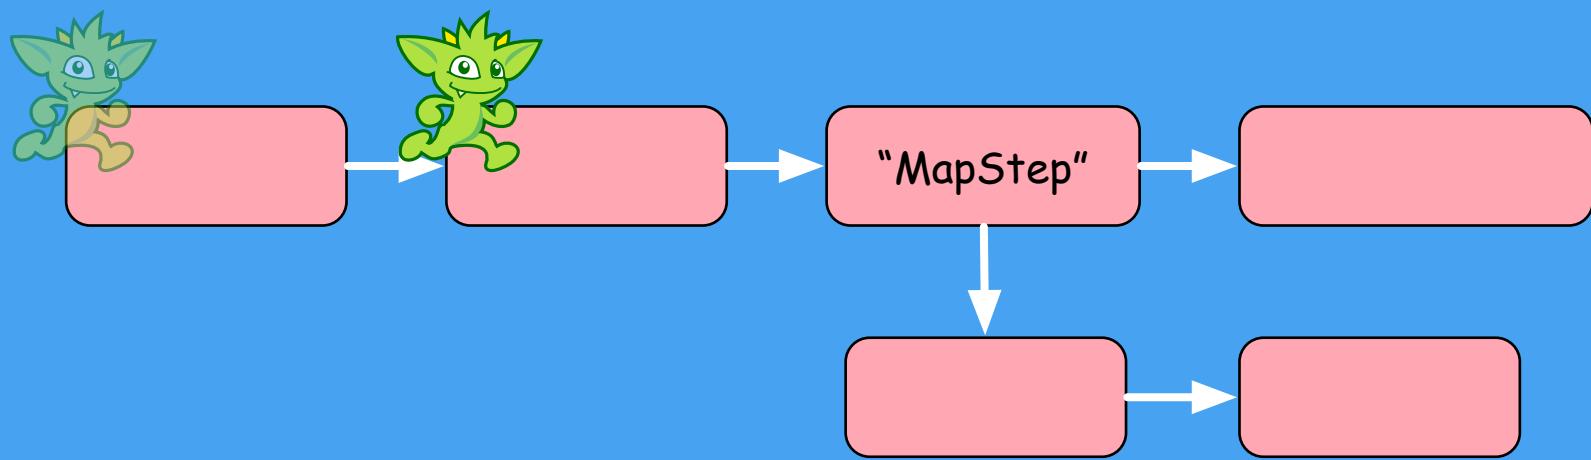


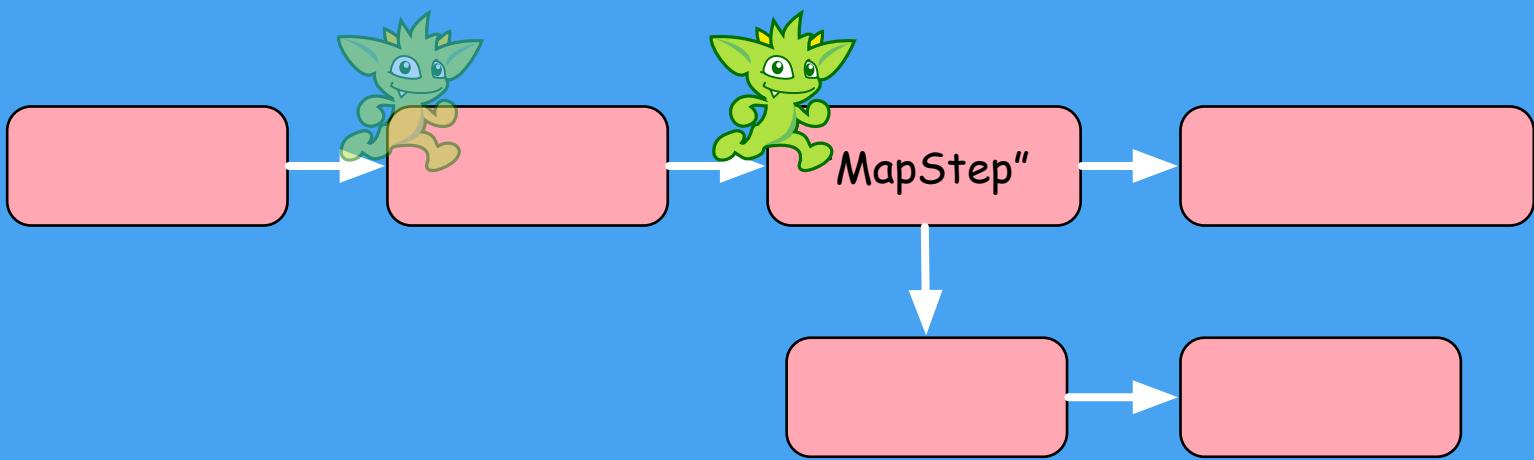


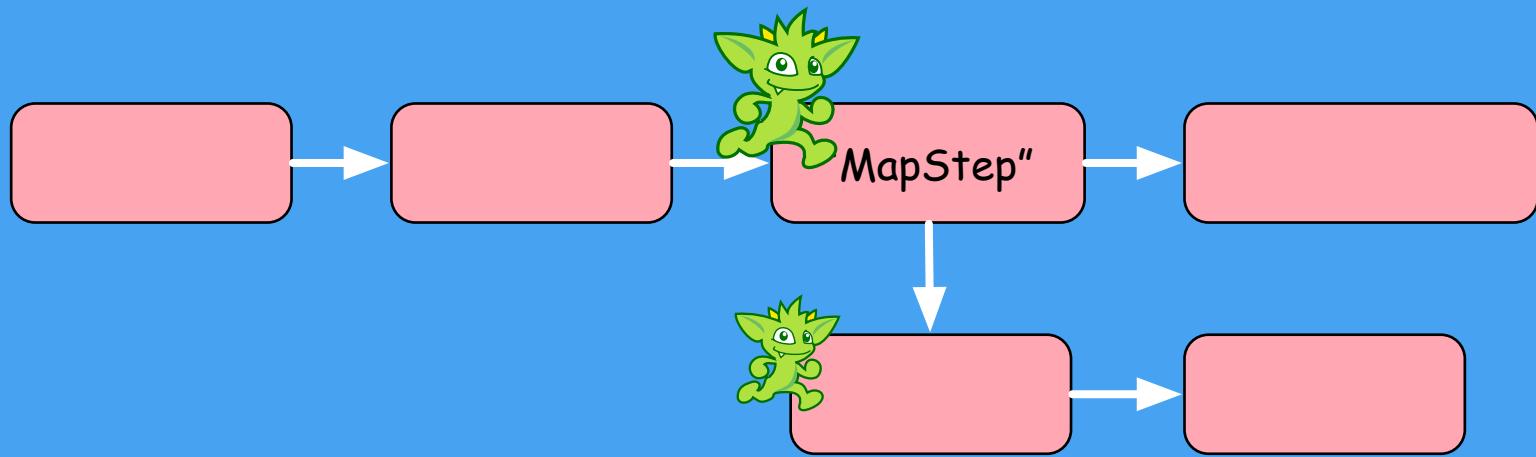


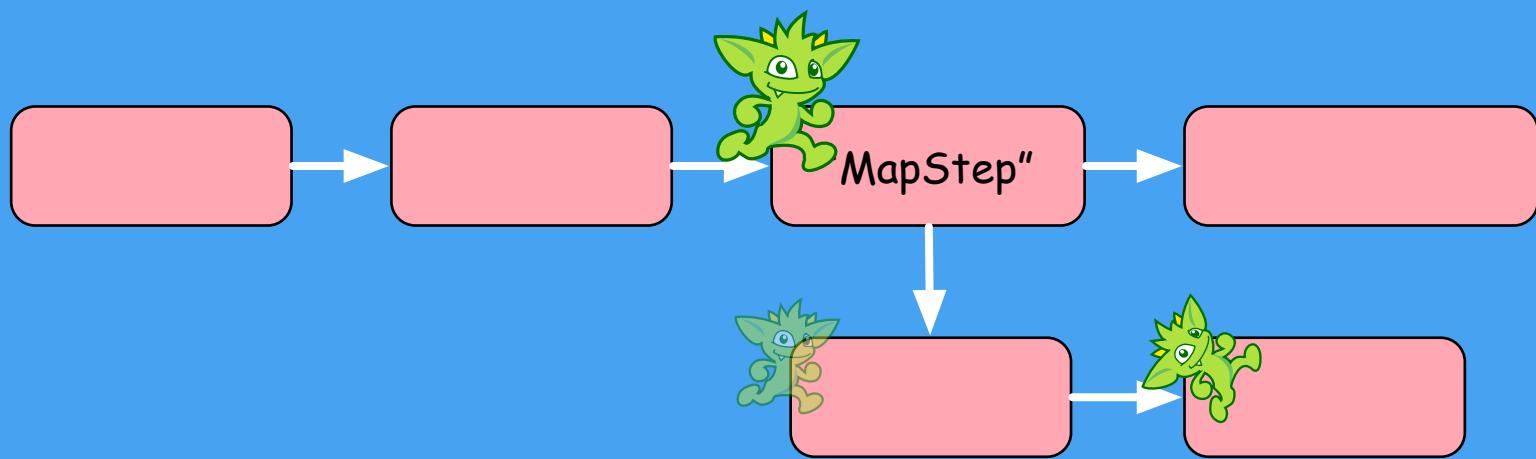


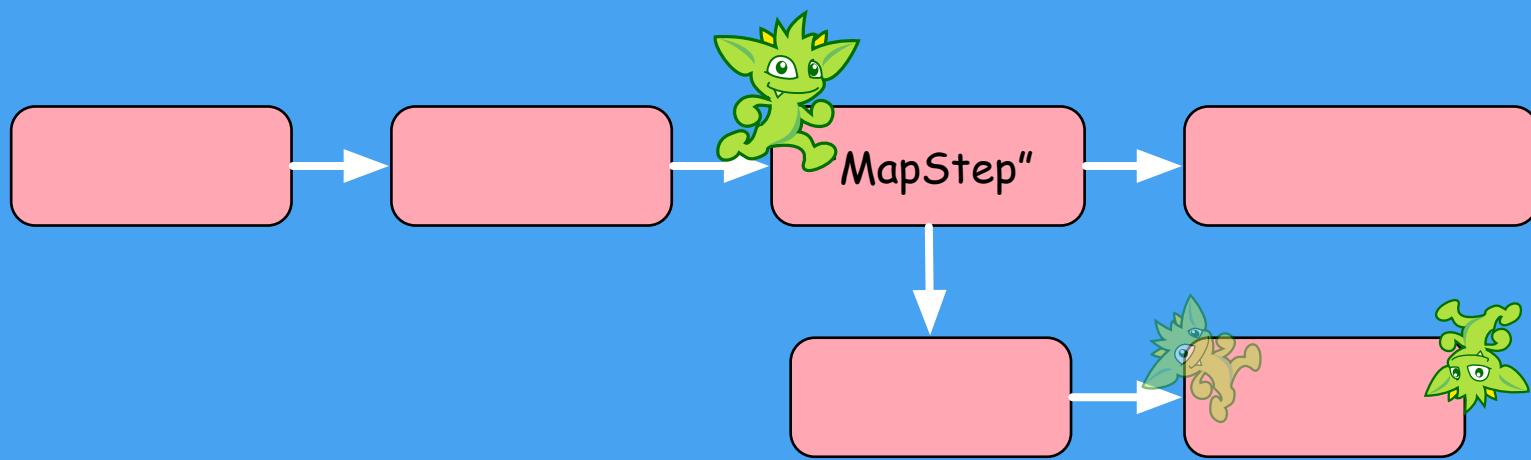


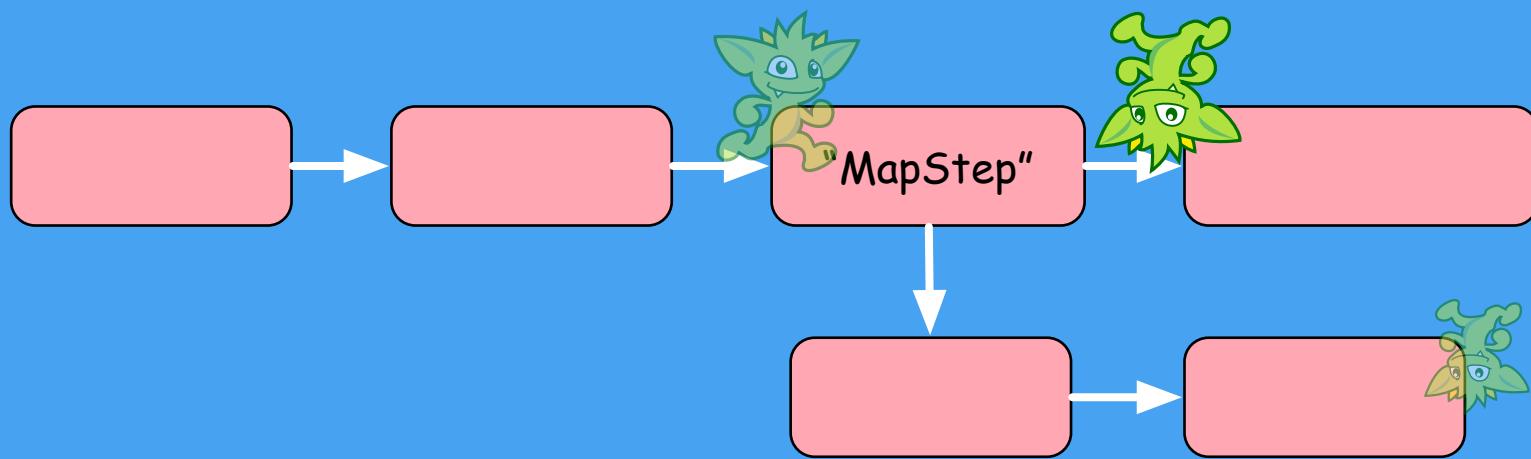


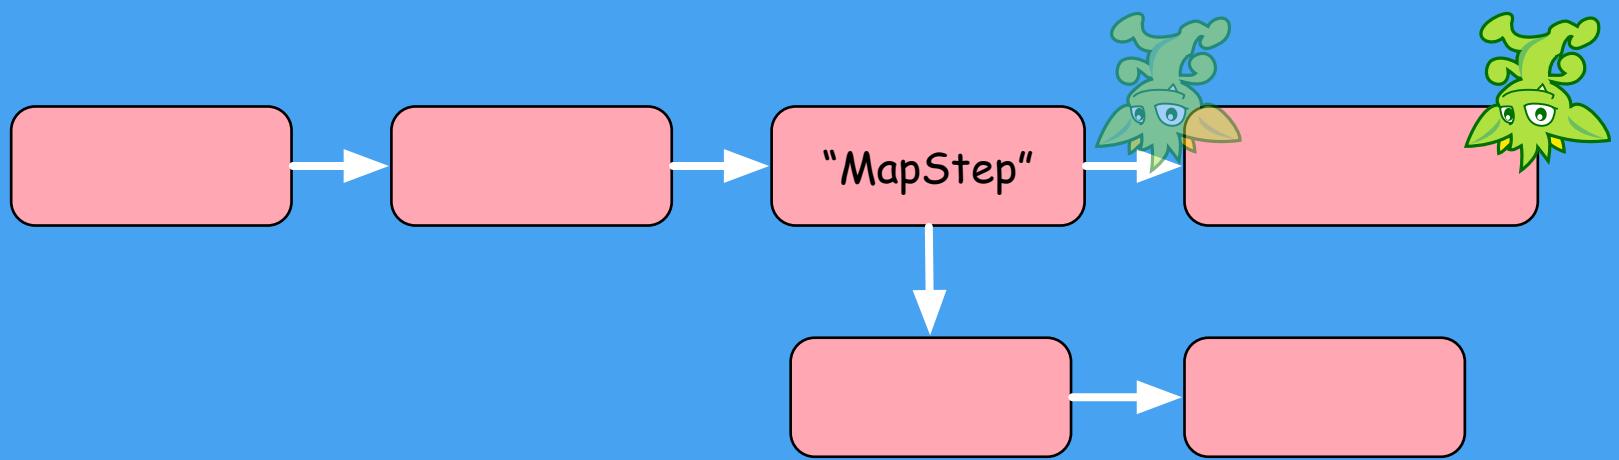












# Chained and Nested Steps



# Step Types

**Map:** one-to-one

**FlatMap:** one-to-many

**Filter:** one-to-one.or.none

**SideEffect:** one-to[?]-one

**Barrier:** many-to-many

**Reducer:** many-to-one

Map



label()  
person



Map



label()  
person



FlatMap



out('knows')



Map



label()  
person



FlatMap



out('knows')



Filter



hasLabel('android') 

Map



label()  
person



FlatMap



out('knows')



SideEffect



x=[  
1  
]



store('x')



hasLabel('android')

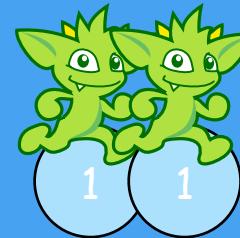
Map



label()  
person



Barrier



barrier()



FlatMap



out('knows')



Filter



hasLabel('android') 

SideEffect



x=[  
1  
]



store('x')

Map



label()  
person



FlatMap



out('knows')



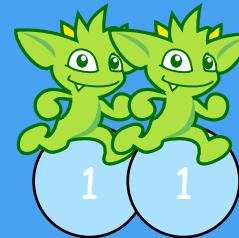
SideEffect



x=[  
  1  
]  
store('x')



Barrier



barrier()



Filter



hasLabel('android')



Reducer



count()



```
g.V().has('name','gremlin')
```



```
g.V().has('name','gremlin').  
out('bought').aggregate('stash')
```



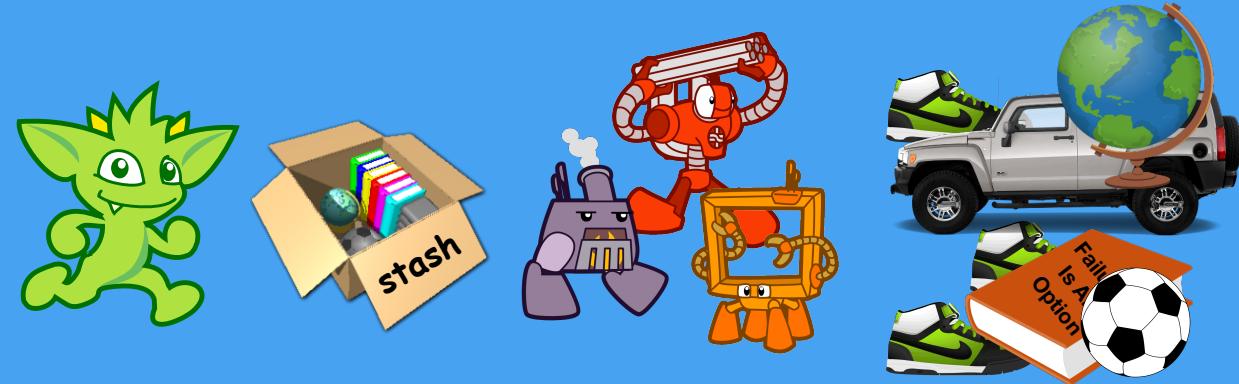
```
g.V().has('name','gremlin').  
out('bought').aggregate('stash').  
in('bought')
```



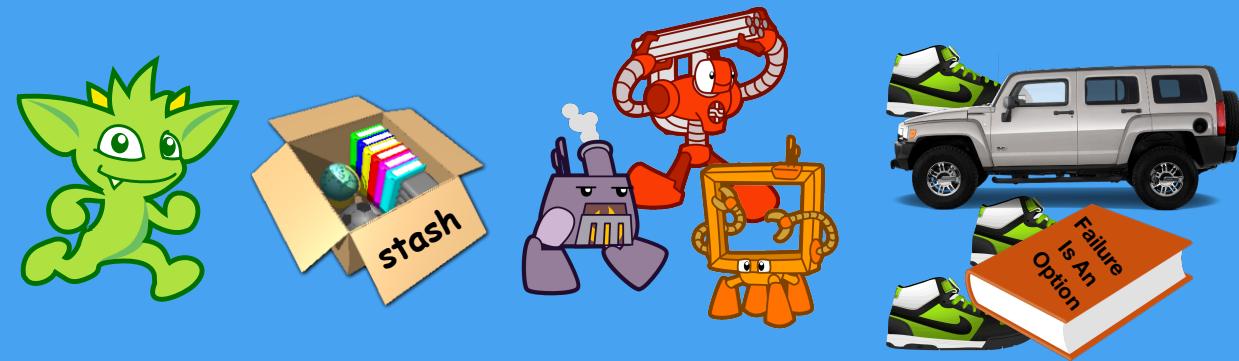
```
g.V().has('name','gremlin').  
out('bought').aggregate('stash').  
in('bought').has('name',neq('gremlin'))
```



```
g.V().has('name','gremlin').  
out('bought').aggregate('stash').  
in('bought').has('name',neq('gremlin')).  
out('bought')
```



```
g.V().has('name','gremlin').  
out('bought').aggregate('stash').  
in('bought').has('name',neq('gremlin'))).  
out('bought').not(within('stash'))
```



```
g.V().has('name','gremlin').  
out('bought').aggregate('stash').  
in('bought').has('name',neq('gremlin')).  
out('bought').not(within('stash')).  
groupCount()
```



```
g.V().has('name','gremlin').  
out('bought').aggregate('stash').  
in('bought').has('name',neq('gremlin')).  
out('bought').not(within('stash')).  
groupCount().  
order(local).by(values,desc)
```



```
g.V().has('name','gremlin').  
out('bought').aggregate('stash').  
in('bought').has('name',neq('gremlin')).  
out('bought').not(within('stash')).  
groupCount().  
order(local).by(values,descr).  
select(keys).unfold().limit(1)
```



# Turing Completeness



repeat().until()

until().repeat()

choose()

choose().option()

as()

select()

do-while

while-do

if-then-else

switch

var set

var get

# Every modern programming language supports function chaining and nesting.



`g.V(1).out("knows")`



`g.V(1).out('knows')`



`g.V(1).out('knows')`

# Quiz



out('created')

repeat(out()).times(2)

count()

id()

groupCount()

label()

where('a', eq("b"))

groupCount('m')

property('name', 'marko')

has('age')

barrier()

values('name')

`out('created')`

`flatmap: vertex → vertex*`

`repeat(out()).times(2)`

`count()`

`id()`

`groupCount()`

`label()`

`where('a', eq('b'))`

`groupCount('m')`

`property('name', 'marko')`

`has('age')`

`barrier()`

`values('name')`

**out('created')**

flatmap: vertex → vertex\*

**id()**

map: element → object

**label()**

**where('a', eq('b'))**

**property('name', 'marko')**

**barrier()**

**repeat(out()).times(2)**

flatmap: vertex → vertex\*

**count()**

reducer: object\* → long

**groupCount()**

**groupCount('m')**

**has('age')**

**values('name')**

**out('created')**

flatmap: vertex → vertex\*

**id()**

map: element → object

**label()**

map: element → string

**where('a', eq("b"))**

filter: object → object

**property('name', 'marko')**

**barrier()**

**repeat(out()).times(2)**

flatmap: vertex → vertex\*

**count()**

reducer: object\* → long

**groupCount()**

reducer: object\* → map<object, long>

**groupCount('m')**

side-effect: object →(map<object, long>) object

**has('age')**

**values('name')**

**out('created')**

flatmap: vertex → vertex\*

**id()**

map: element → object

**label()**

map: element → string

**where('a', eq('b'))**

filter: object → object

**property('name', 'marko')**

side-effect: element →(element) element

**barrier()**

barrier: object\* → object\*

**repeat(out()).times(2)**

flatmap: vertex → vertex\*

**count()**

reducer: object\* → long

**groupCount()**

reducer: object\* → map<object, long>

**groupCount('m')**

side-effect: object →(map<object, long>) object

**has('age')**

filter: element → element

**values('name')**

flatmap: element → string\*



The traversal guides the traverser down the long and winding graph.

T

The Traversers

**Traversers**

$$G \xleftarrow{\mu} t \in T \psi \xrightarrow{\Psi}$$

**Graph** **Traversal**

# Traversers

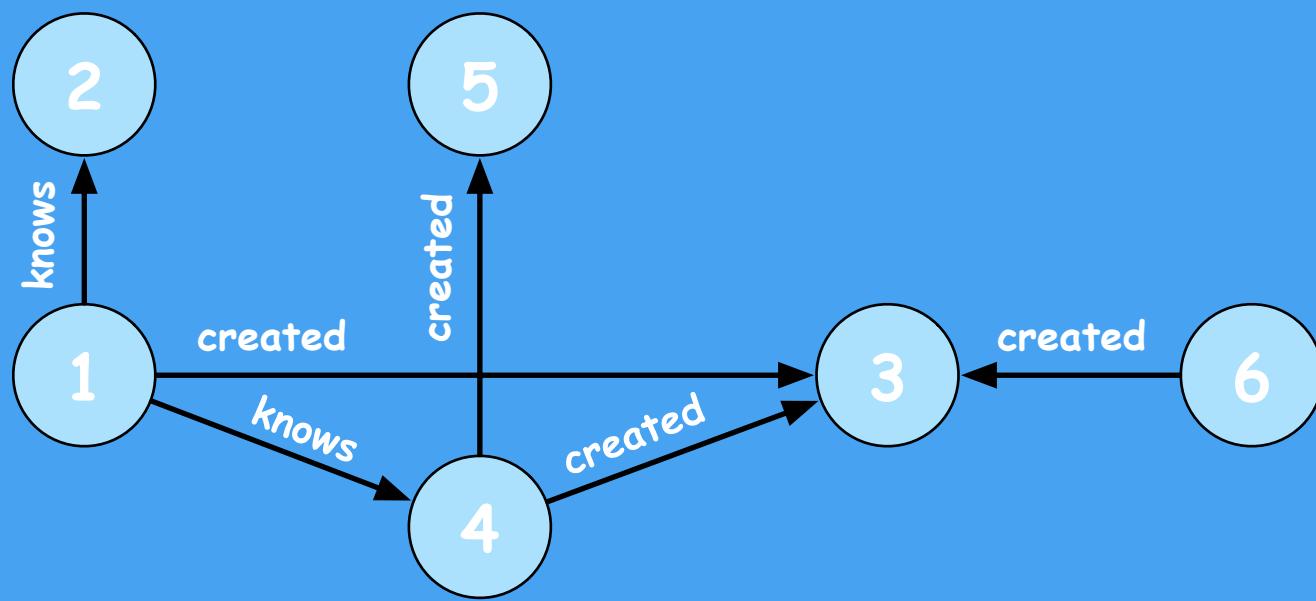
$G \xleftarrow{\mu} t \in T \xrightarrow{\psi} \Psi$

$\beta, \Delta, \zeta, \iota$

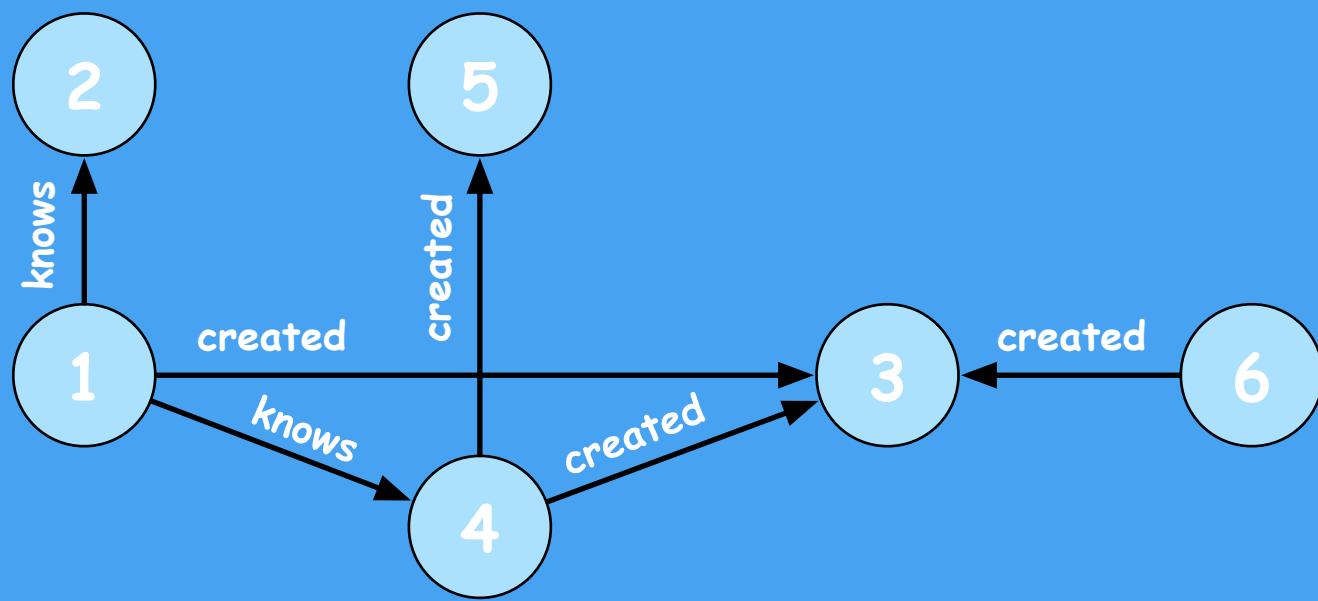


Bulk Path Sack Loops

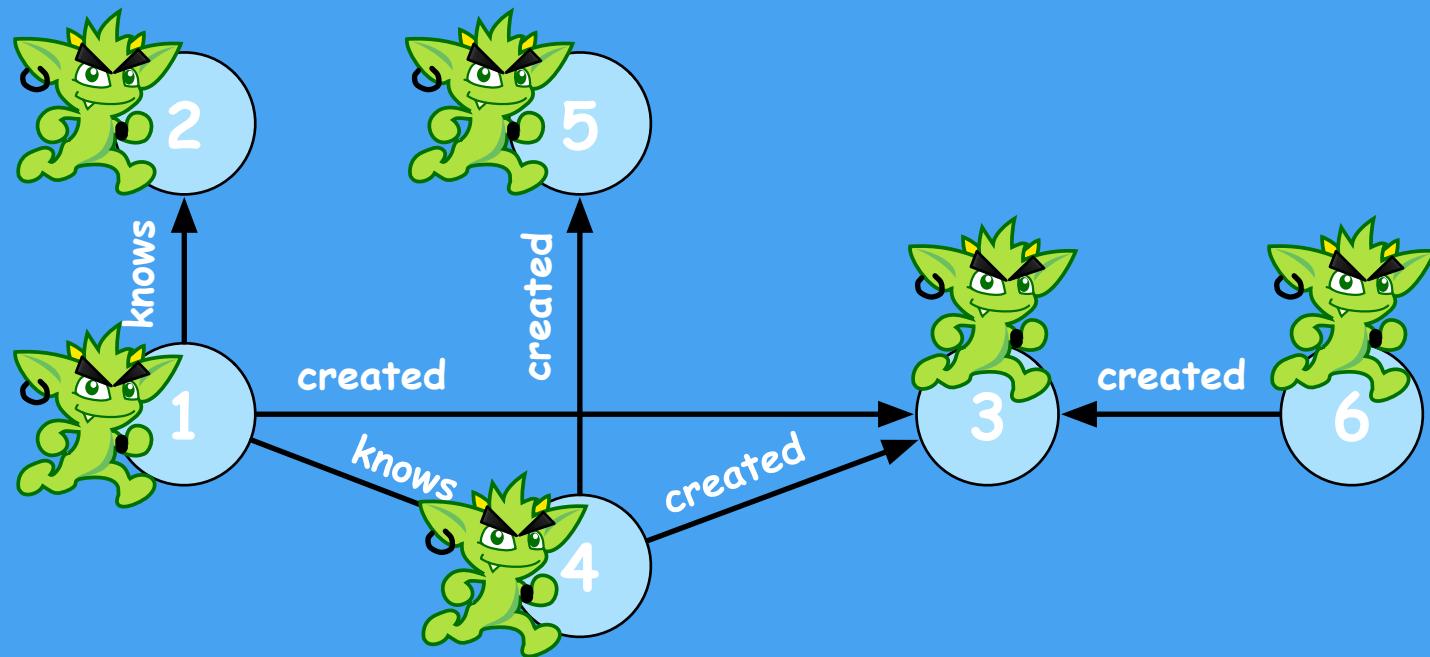
# `g.V().both().id()`



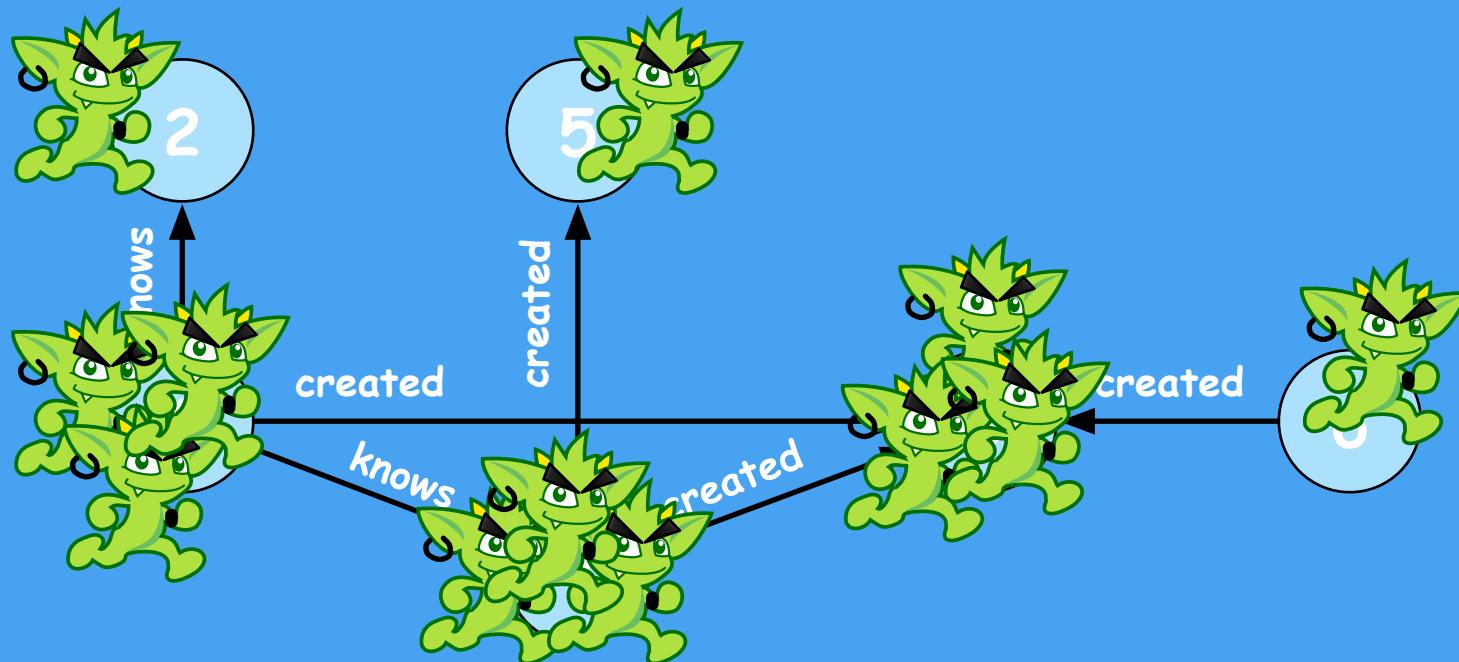
# `g.V().both().id()`



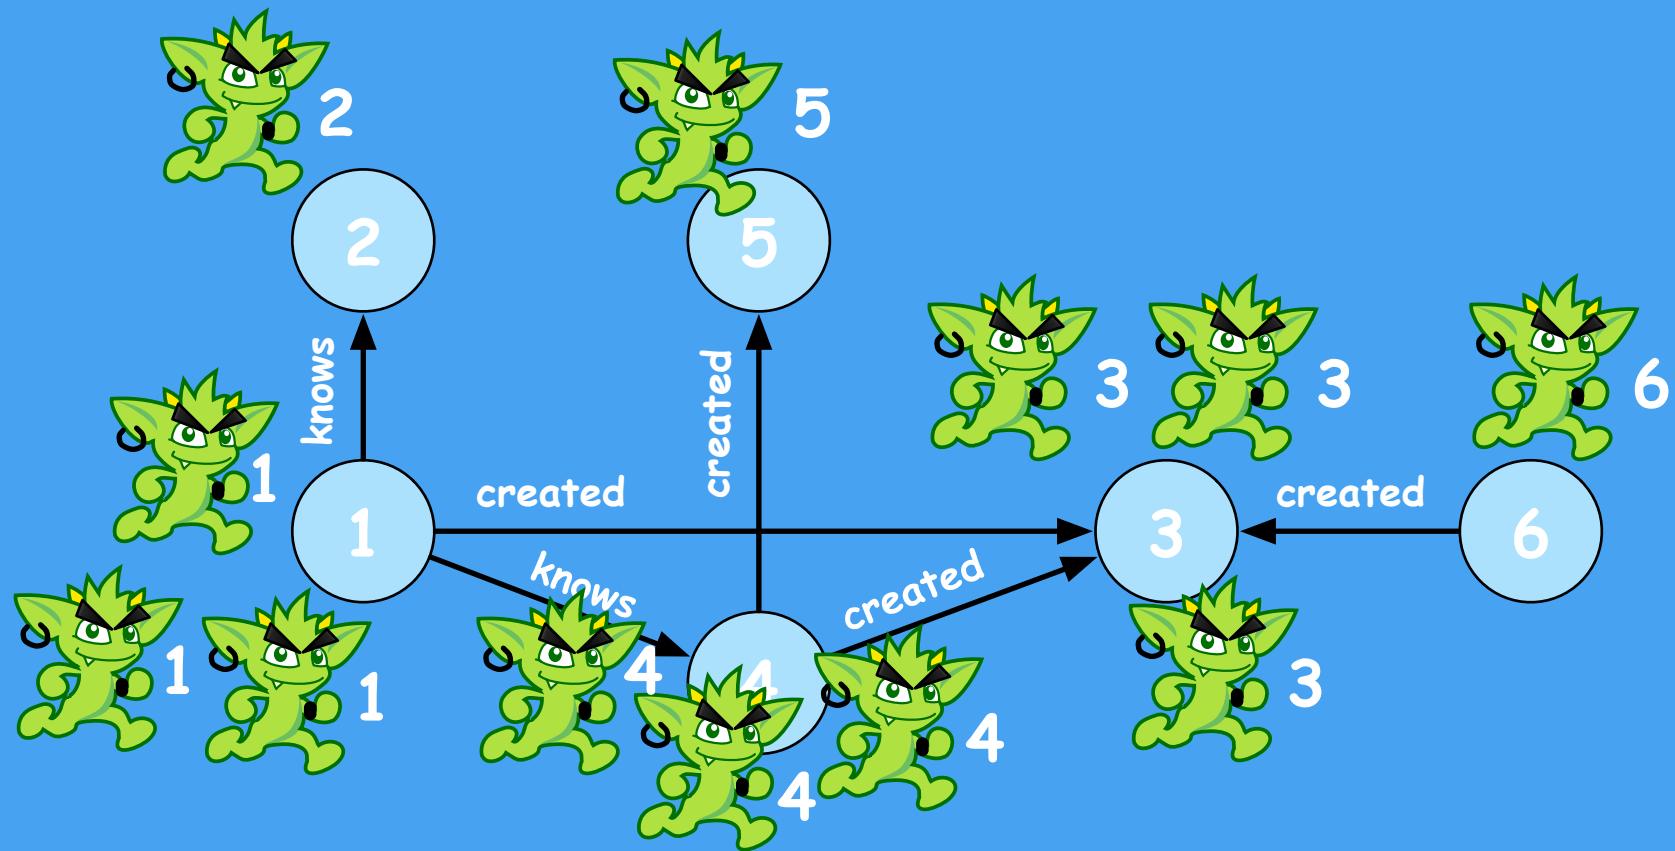
# `g.V().both().id()`



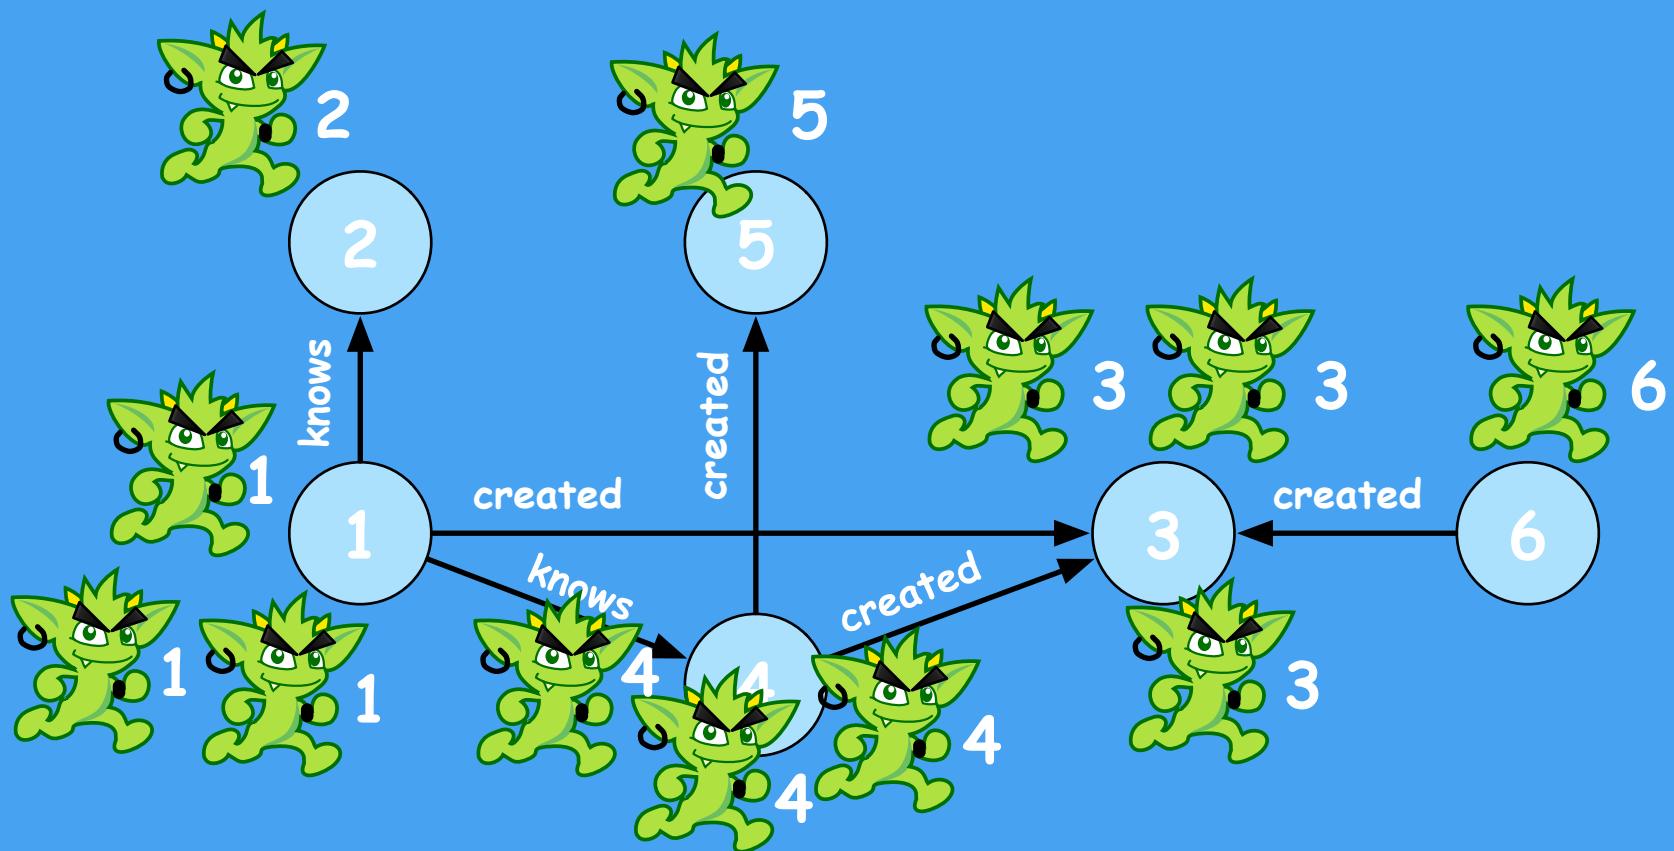
# `g.V().both().id()`



# g.V().both().id()



# `g.V().both().id()`



`==>2`

`==>5`

`==>1`

`==>1`

`==>1`

`==>4`

`==>4`

`==>4`

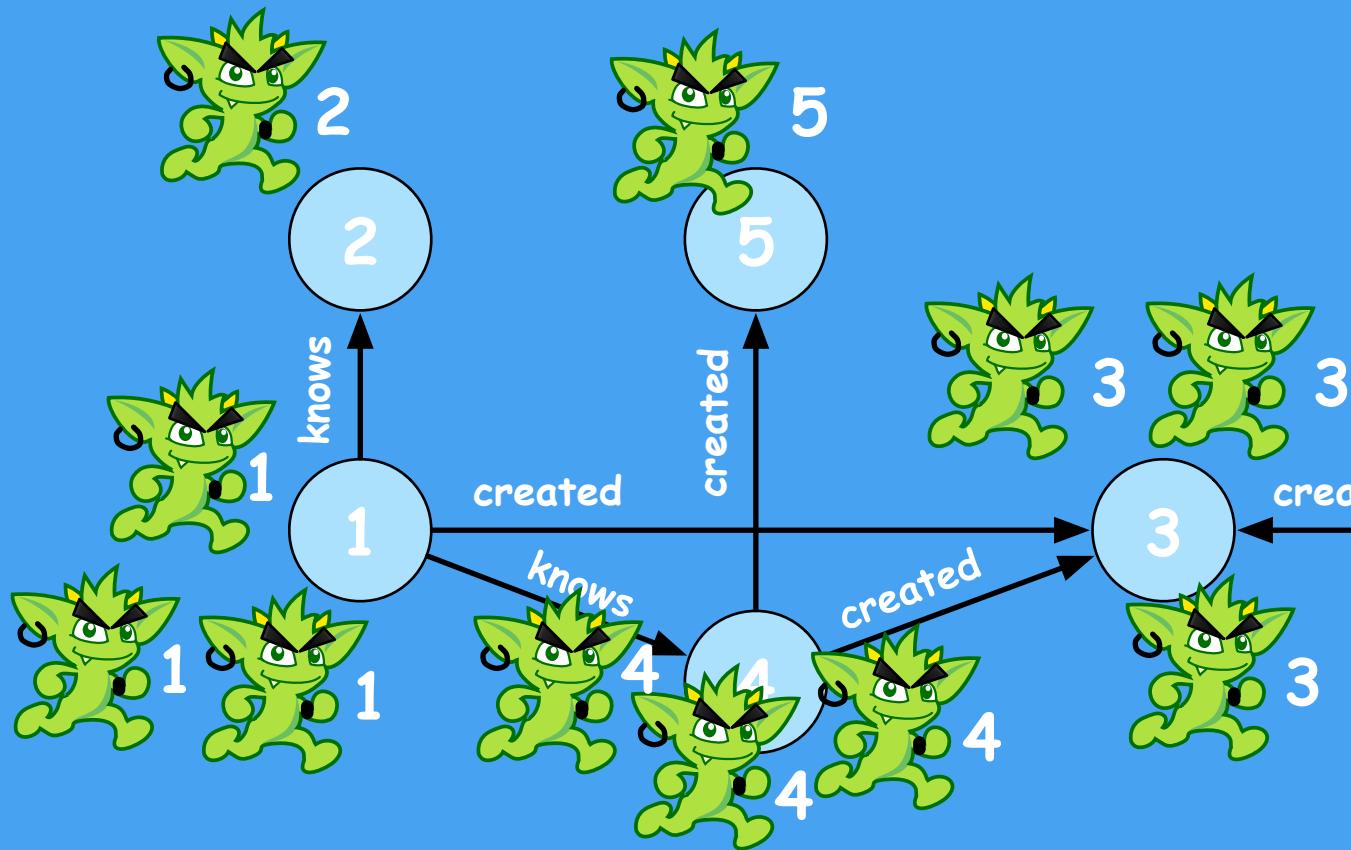
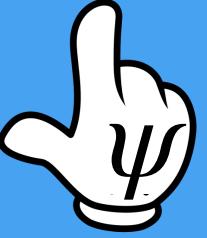
`==>3`

`==>3`

`==>3`

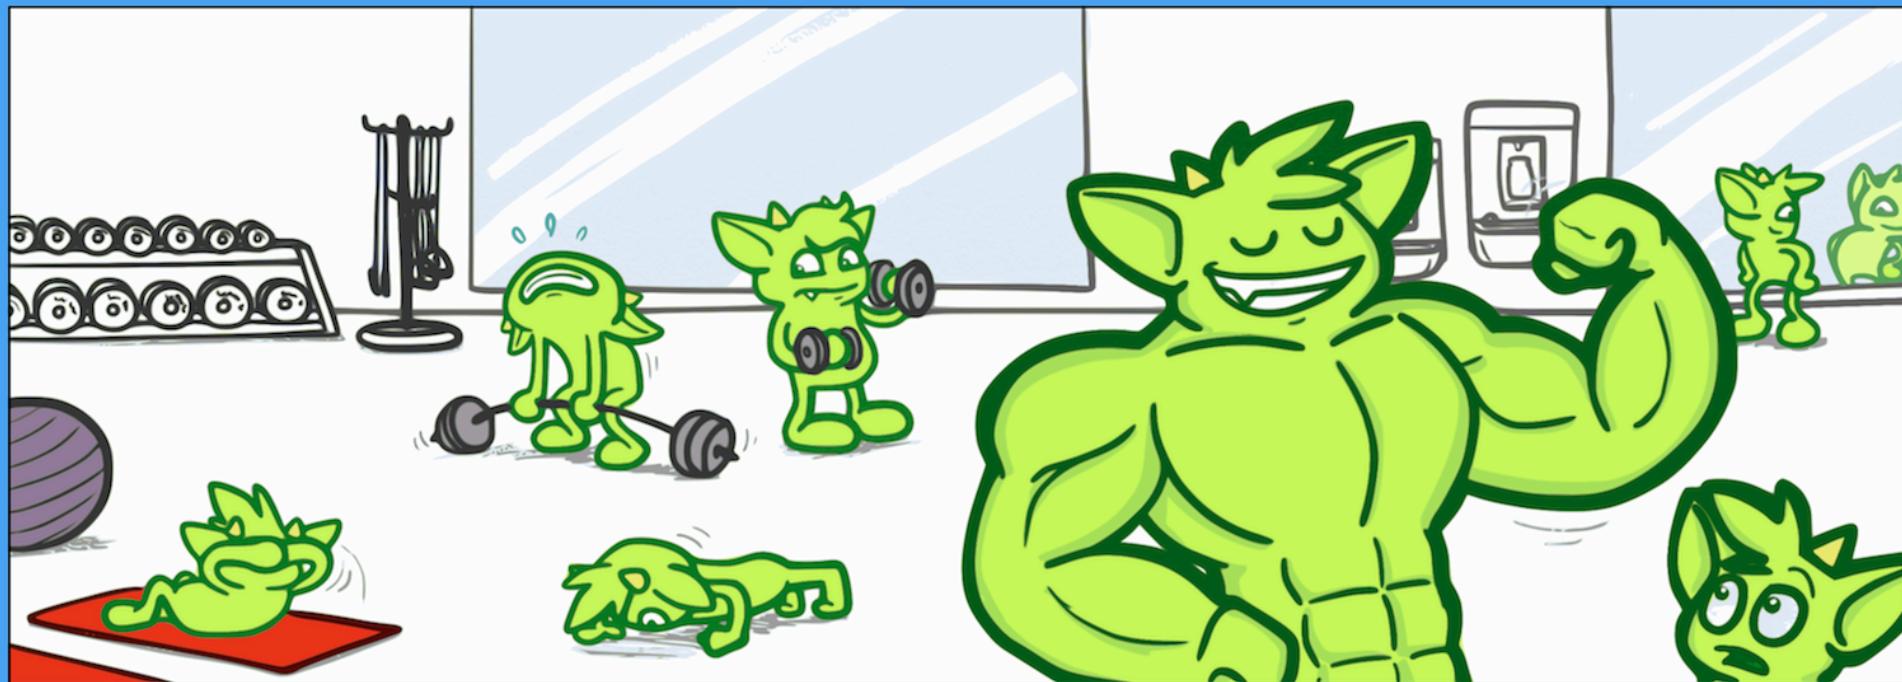
`==>6`

# g.V().both().id()



id(2) id(5)  
id(1) id(1)  
id(1) id(4)  
id(4) id(4)  
id(3) id(3)  
6 id(3)  
id(6)

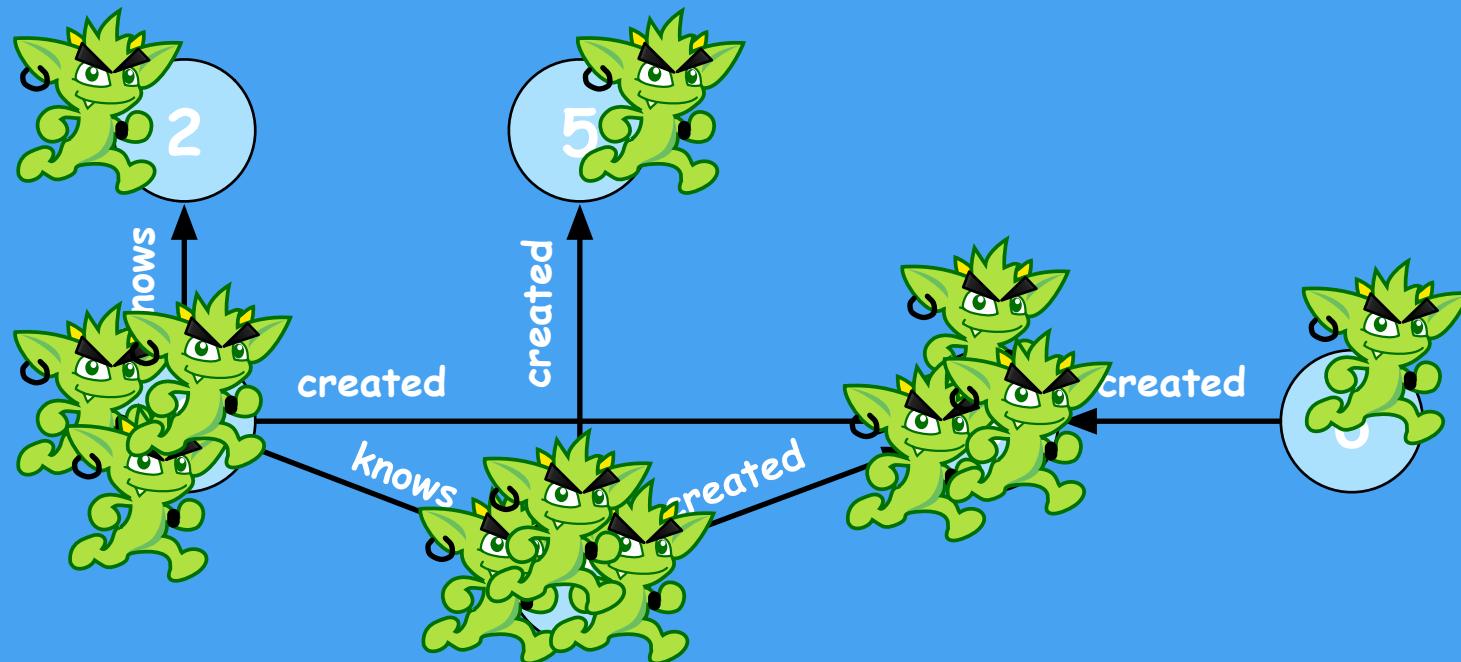
# Its time to bulk up.



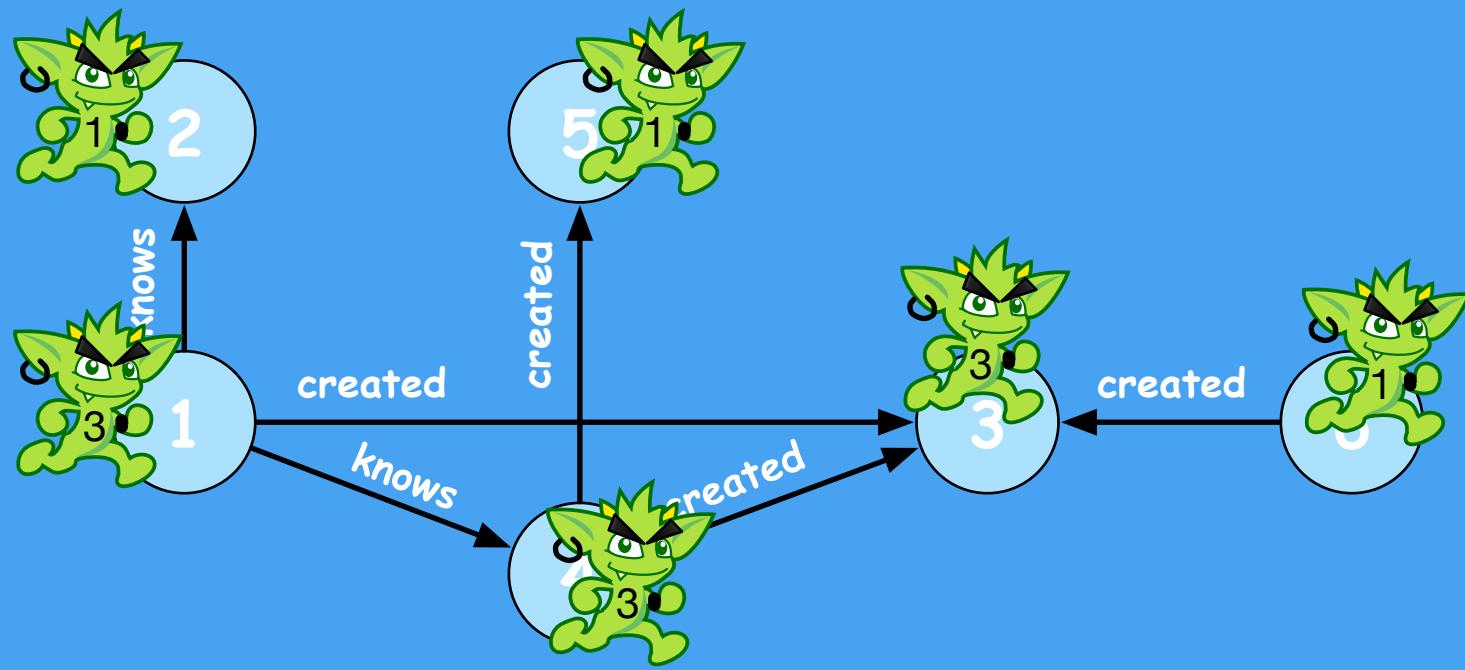
Rodriguez, M.A., "The Beauty of Bulking," Gremlin-Users Mailing List, June 2015.

<https://twitter.com/twarko/status/606513003090092035>

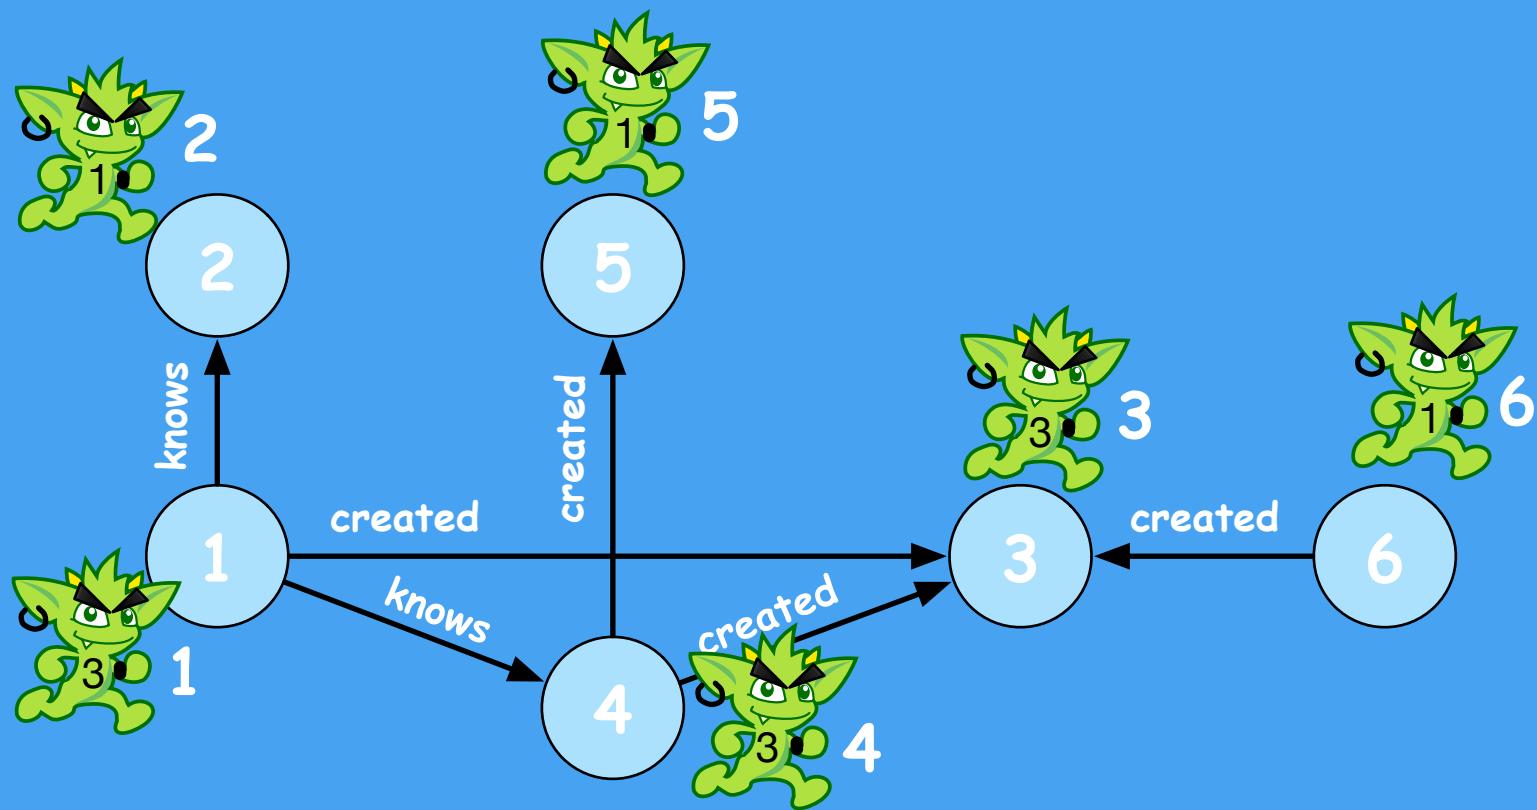
# `g.V().both().id()`



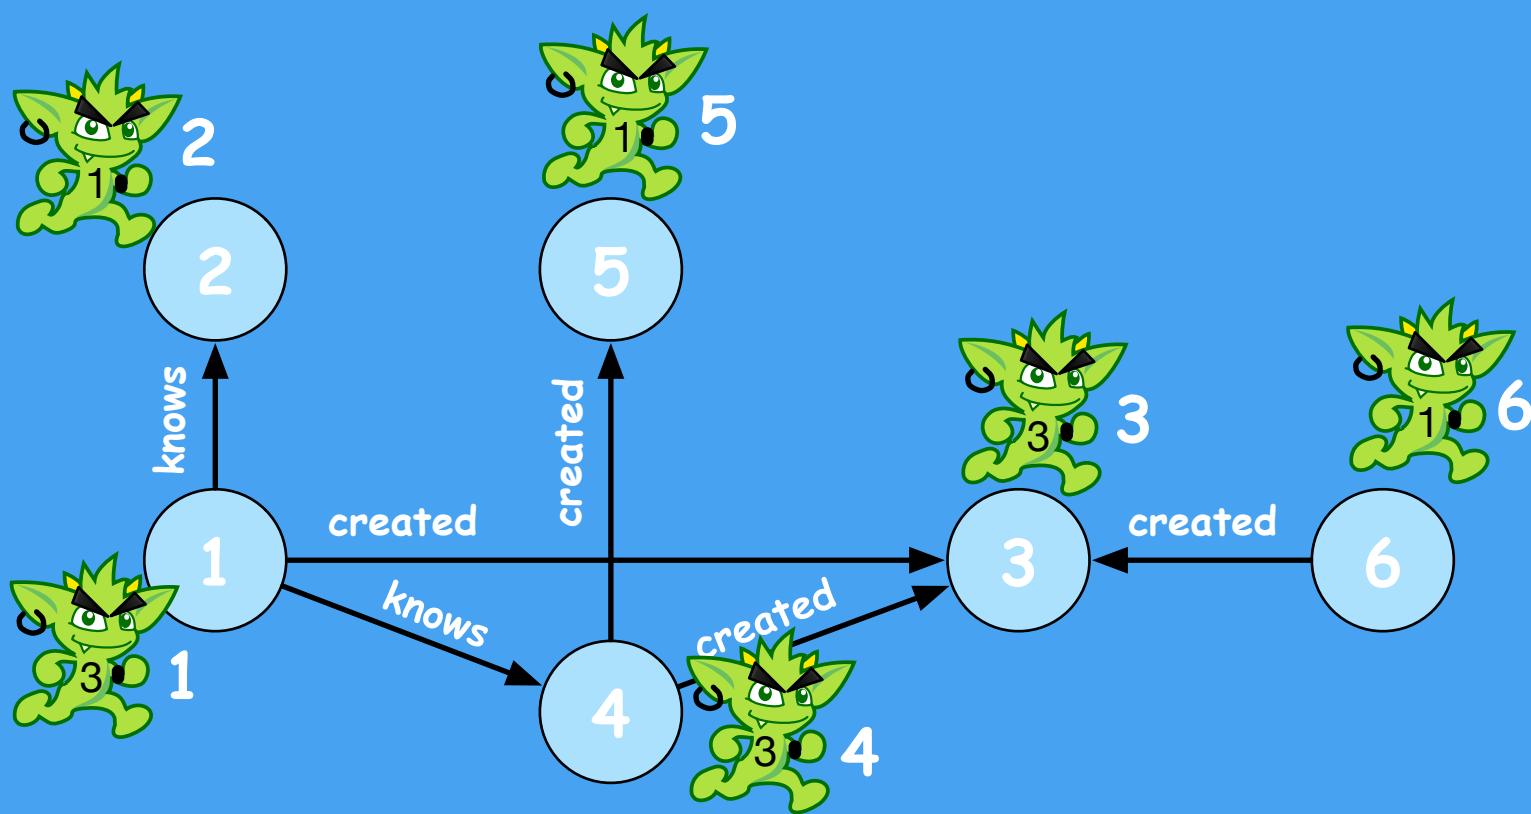
# g.V().both().id()



# g.V().both().id()



# `g.V().both().id()`



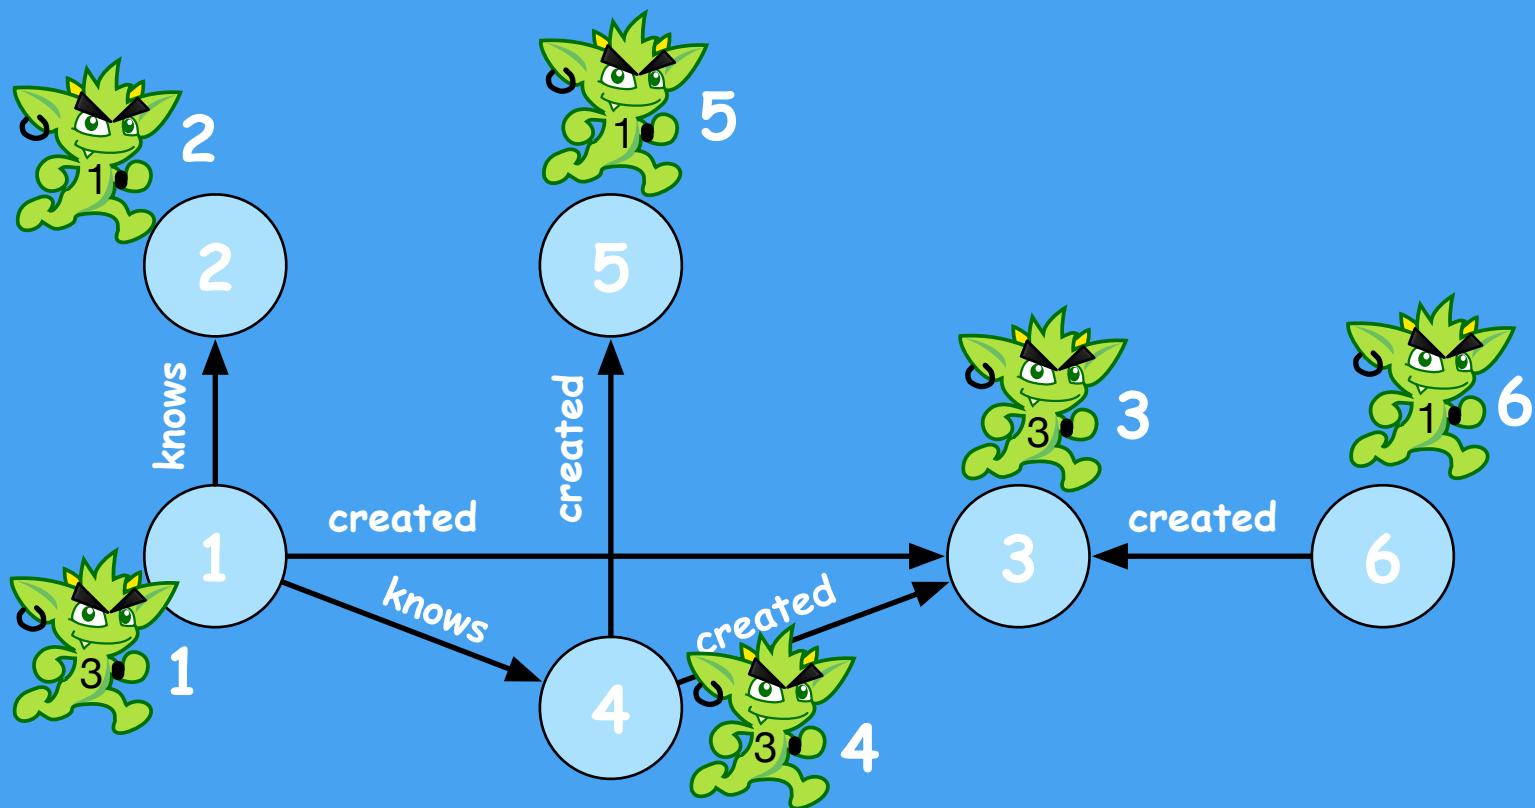
$\Rightarrow 2$   
 $\Rightarrow 5$   
 $\Rightarrow 1$   
 $\Rightarrow 1$   
 $\Rightarrow 1$   
 $\Rightarrow 4$   
 $\Rightarrow 4$   
 $\Rightarrow 4$   
 $\Rightarrow 3$   
 $\Rightarrow 3$   
 $\Rightarrow 3$   
 $\Rightarrow 6$



# g.V().both().id()



1/2 as many function calls with bulking!



# `g.V().both().id()`

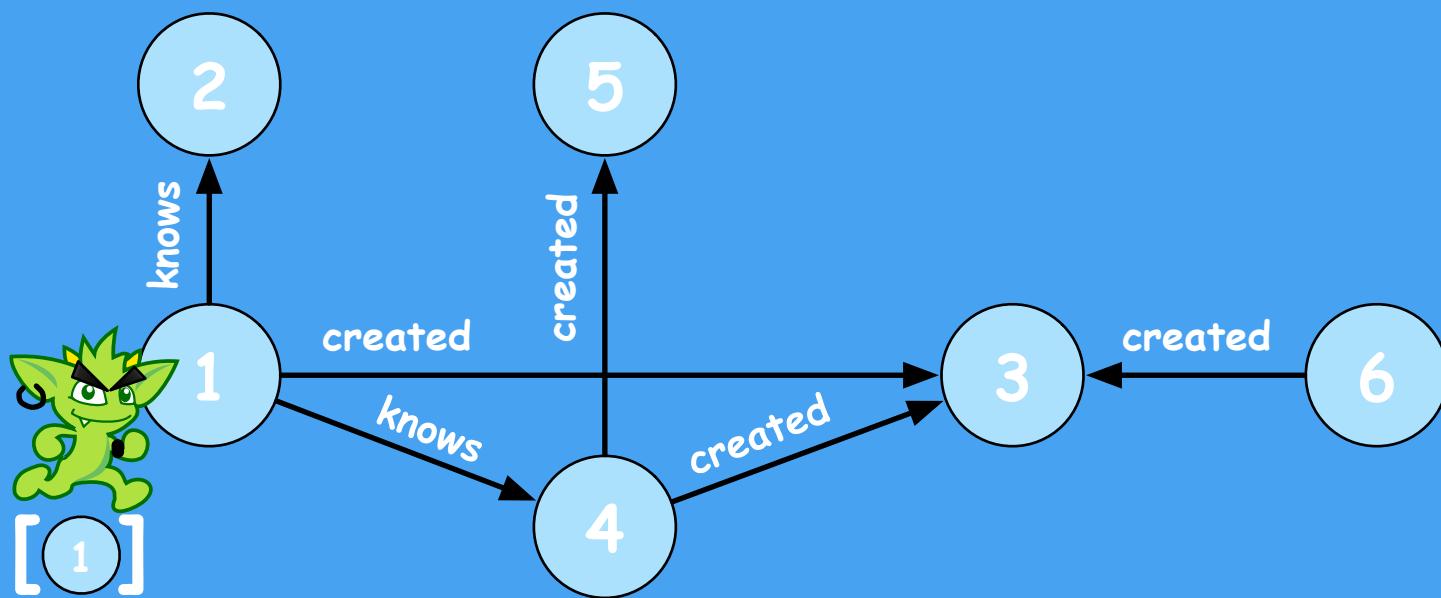
optimizes to

# `g.V().both().barrier().id()`

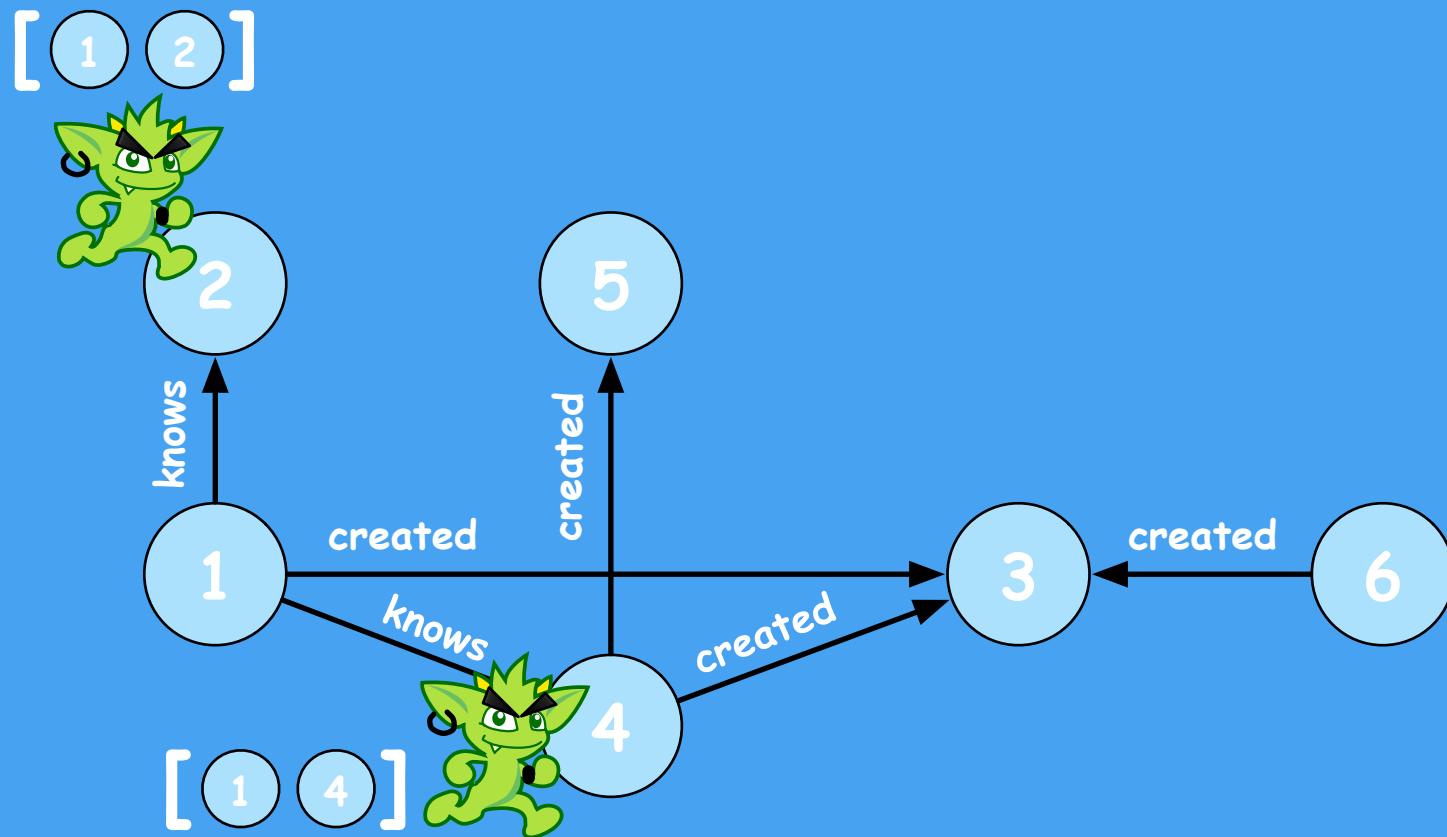
```
gremlin> g.V().both().id().explain()
==>Traversal Explanation
=====
Original Traversal [GraphStep(vertex,[]), VertexStep(BOTH,vertex), IdStep]
ConnectiveStrategy [D] [GraphStep(vertex,[]), VertexStep(BOTH,vertex), IdStep]
MatchPredicateStrategy [O] [GraphStep(vertex,[]), VertexStep(BOTH,vertex), IdStep]
FilterRankingStrategy [O] [GraphStep(vertex,[]), VertexStep(BOTH,vertex), IdStep]
InlineFilterStrategy [O] [GraphStep(vertex,[]), VertexStep(BOTH,vertex), IdStep]
RepeatUnrollStrategy [O] [GraphStep(vertex,[]), VertexStep(BOTH,vertex), IdStep]
PathRetractionStrategy [O] [GraphStep(vertex,[]), VertexStep(BOTH,vertex), IdStep]
IncidentToAdjacentStrategy [O] [GraphStep(vertex,[]), VertexStep(BOTH,vertex), IdStep]
AdjacentToIncidentStrategy [O] [GraphStep(vertex,[]), VertexStep(BOTH,vertex), IdStep]
RangeByIsCountStrategy [O] [GraphStep(vertex,[]), VertexStep(BOTH,vertex), IdStep]
LazyBarrierStrategy [O] [GraphStep(vertex,[]), VertexStep(BOTH,vertex), NoOpBarrierStep(2500), IdStep]
TinkerGraphCountStrategy [P] [GraphStep(vertex,[]), VertexStep(BOTH,vertex), NoOpBarrierStep(2500), IdStep]
TinkerGraphStepStrategy [P] [TinkerGraphStep(vertex,[]), VertexStep(BOTH,vertex), NoOpBarrierStep(2500), IdStep]
ProfileStrategy [F] [TinkerGraphStep(vertex,[]), VertexStep(BOTH,vertex), NoOpBarrierStep(2500), IdStep]
StandardVerificationStrategy [V] [TinkerGraphStep(vertex,[]), VertexStep(BOTH,vertex), NoOpBarrierStep(2500), IdStep]

Final Traversal [TinkerGraphStep(vertex,[]), VertexStep(BOTH,vertex), NoOpBarrierStep(2500), IdStep]
```

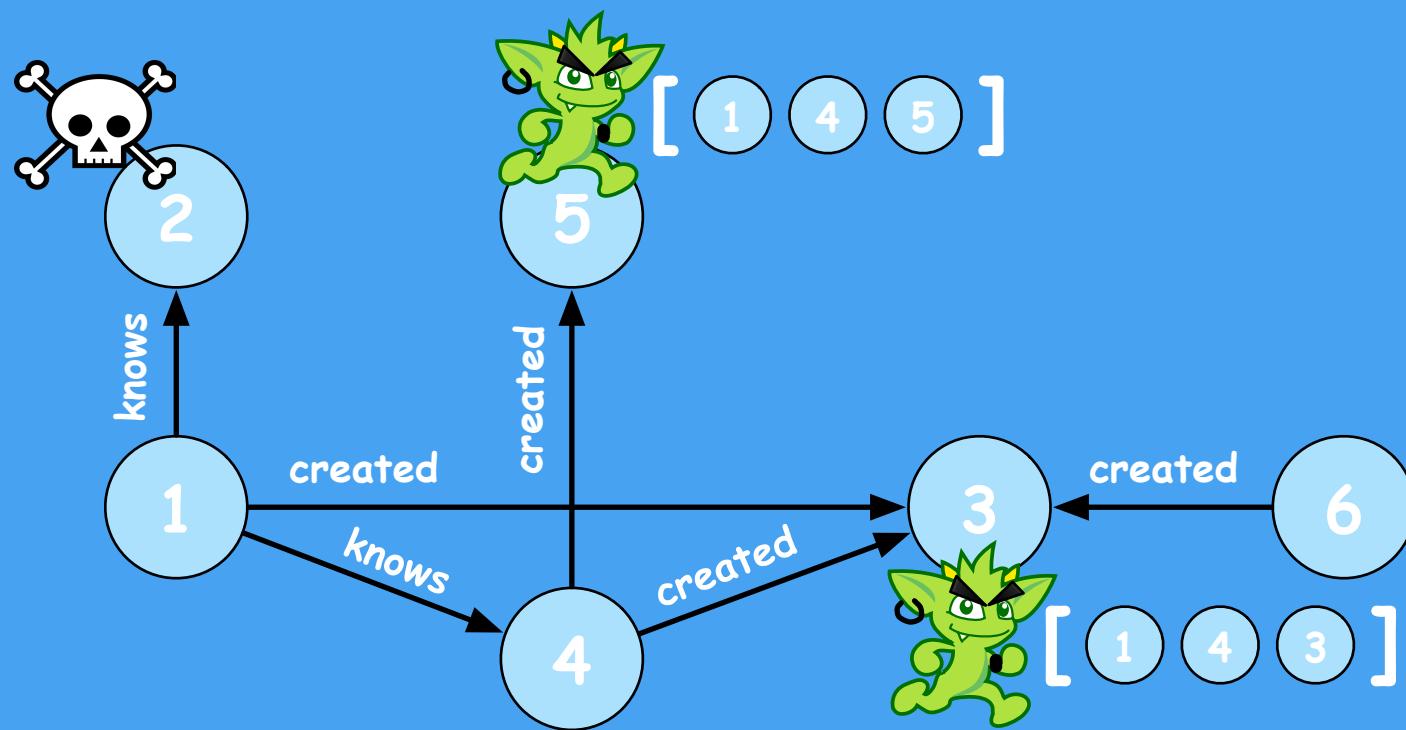
```
g.V(1).out('knows').out('created').path()
```



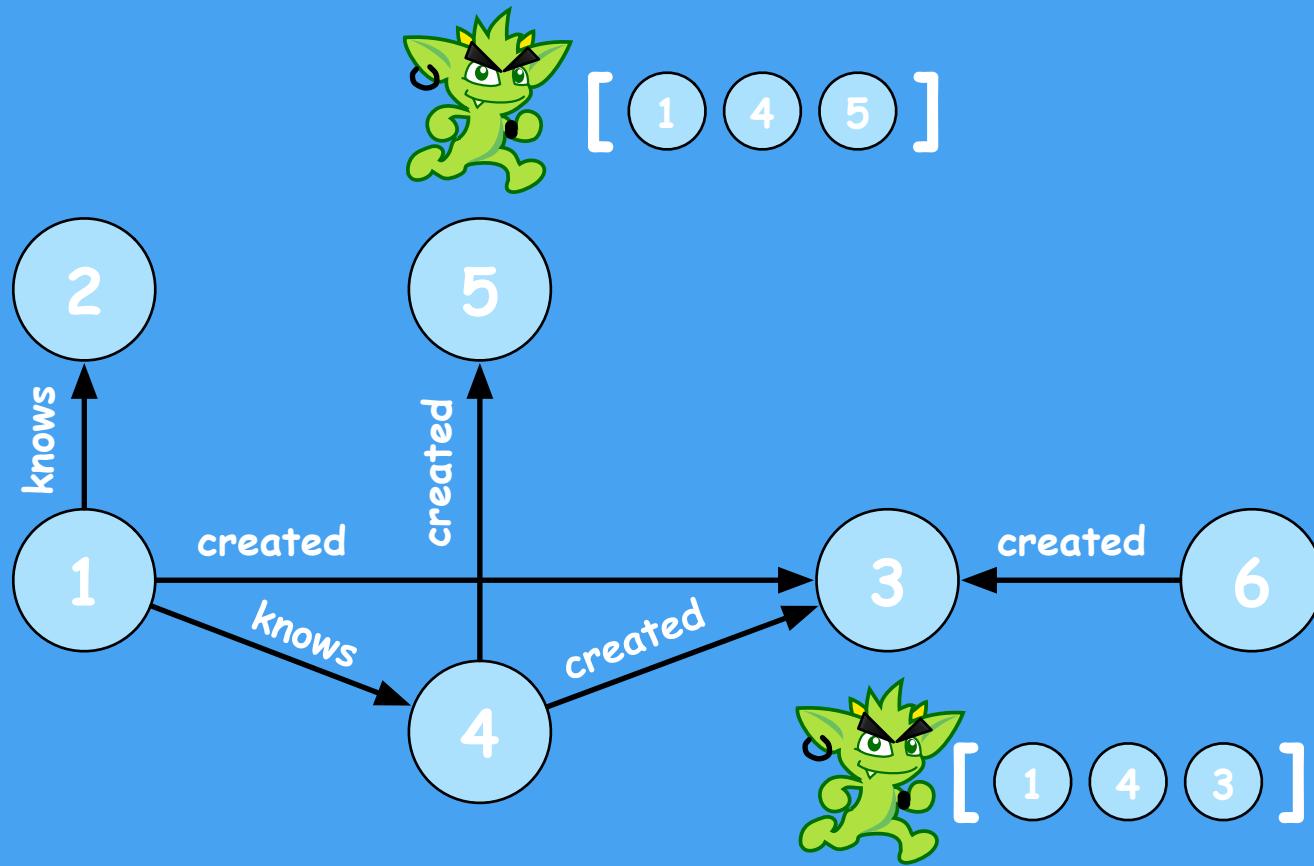
```
g.V(1).out('knows').out('created').path()
```



```
g.V(1).out('knows').out('created').path()
```



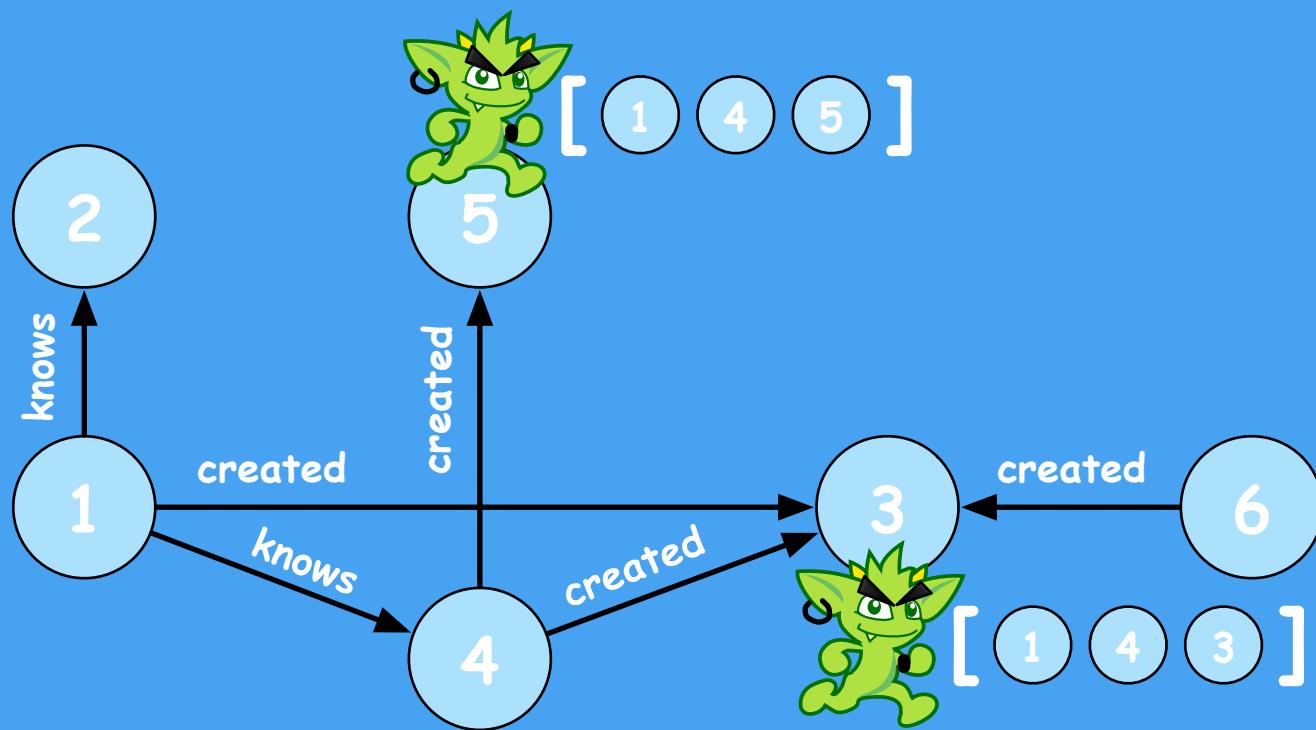
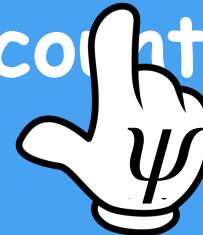
```
g.V(1).out('knows').out('created').path()
```



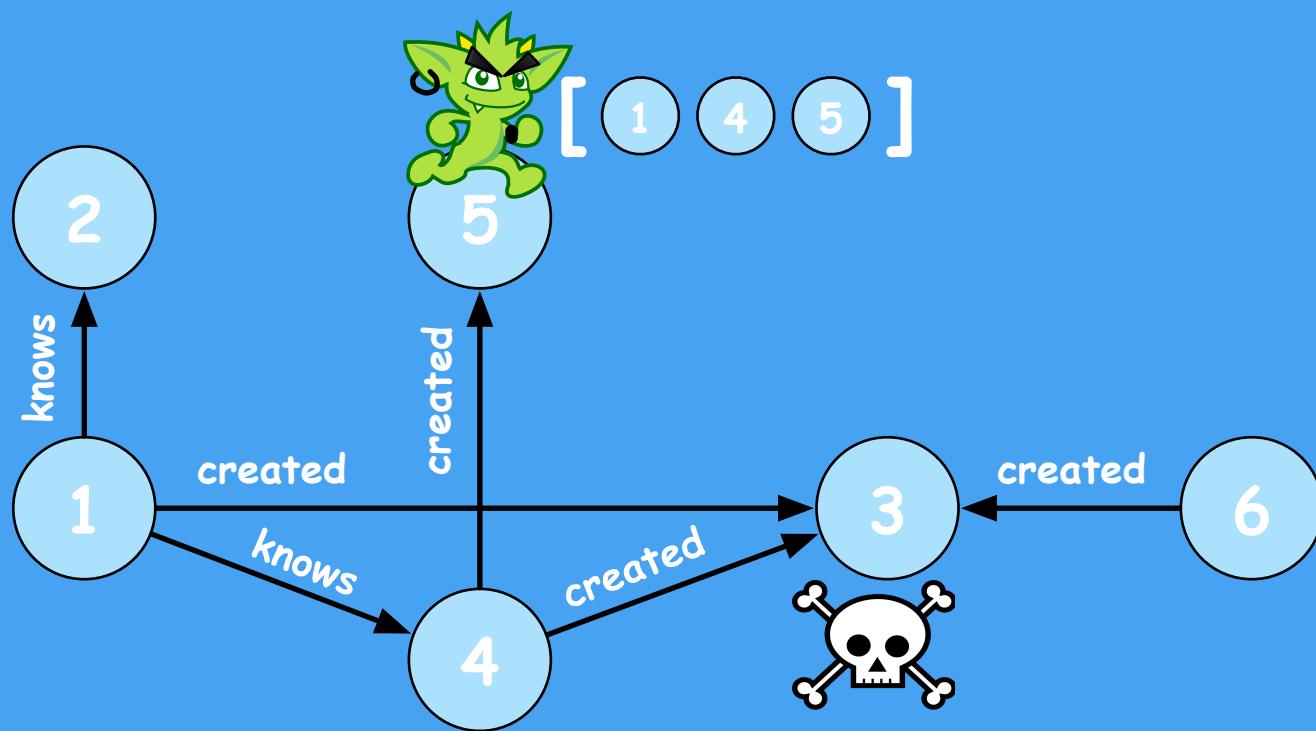
Its time to get crazy with  
locally scoped traversers.



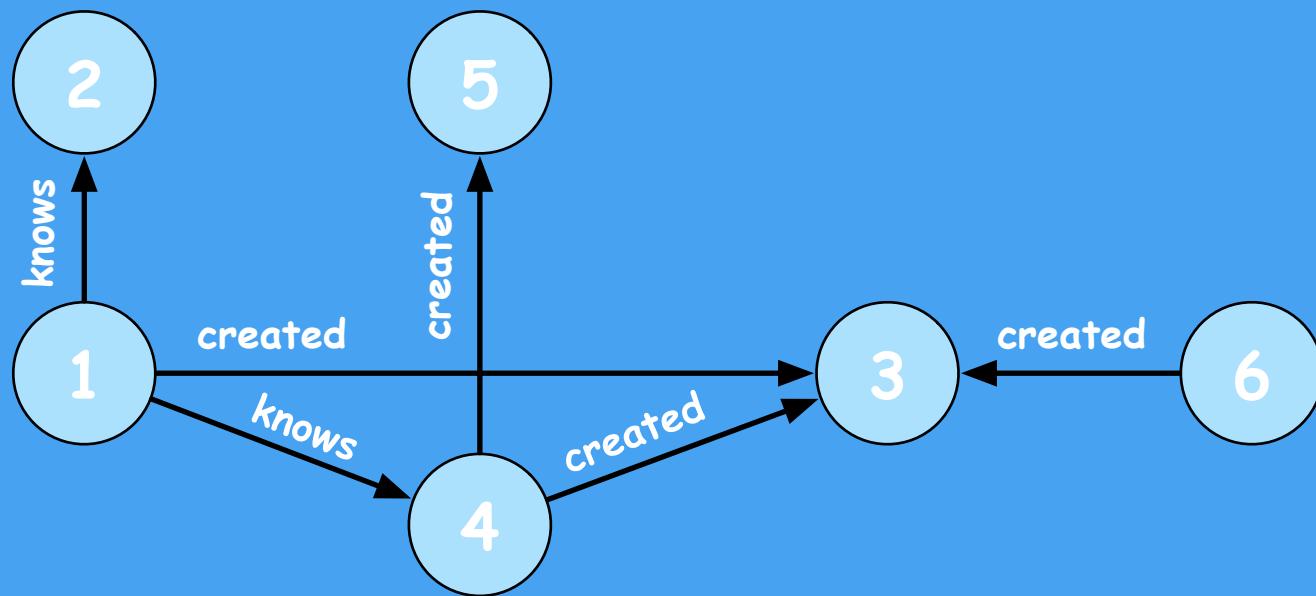
```
g.V(1).out('knows').out('created').limit(1).  
path().by(inE()).count())
```



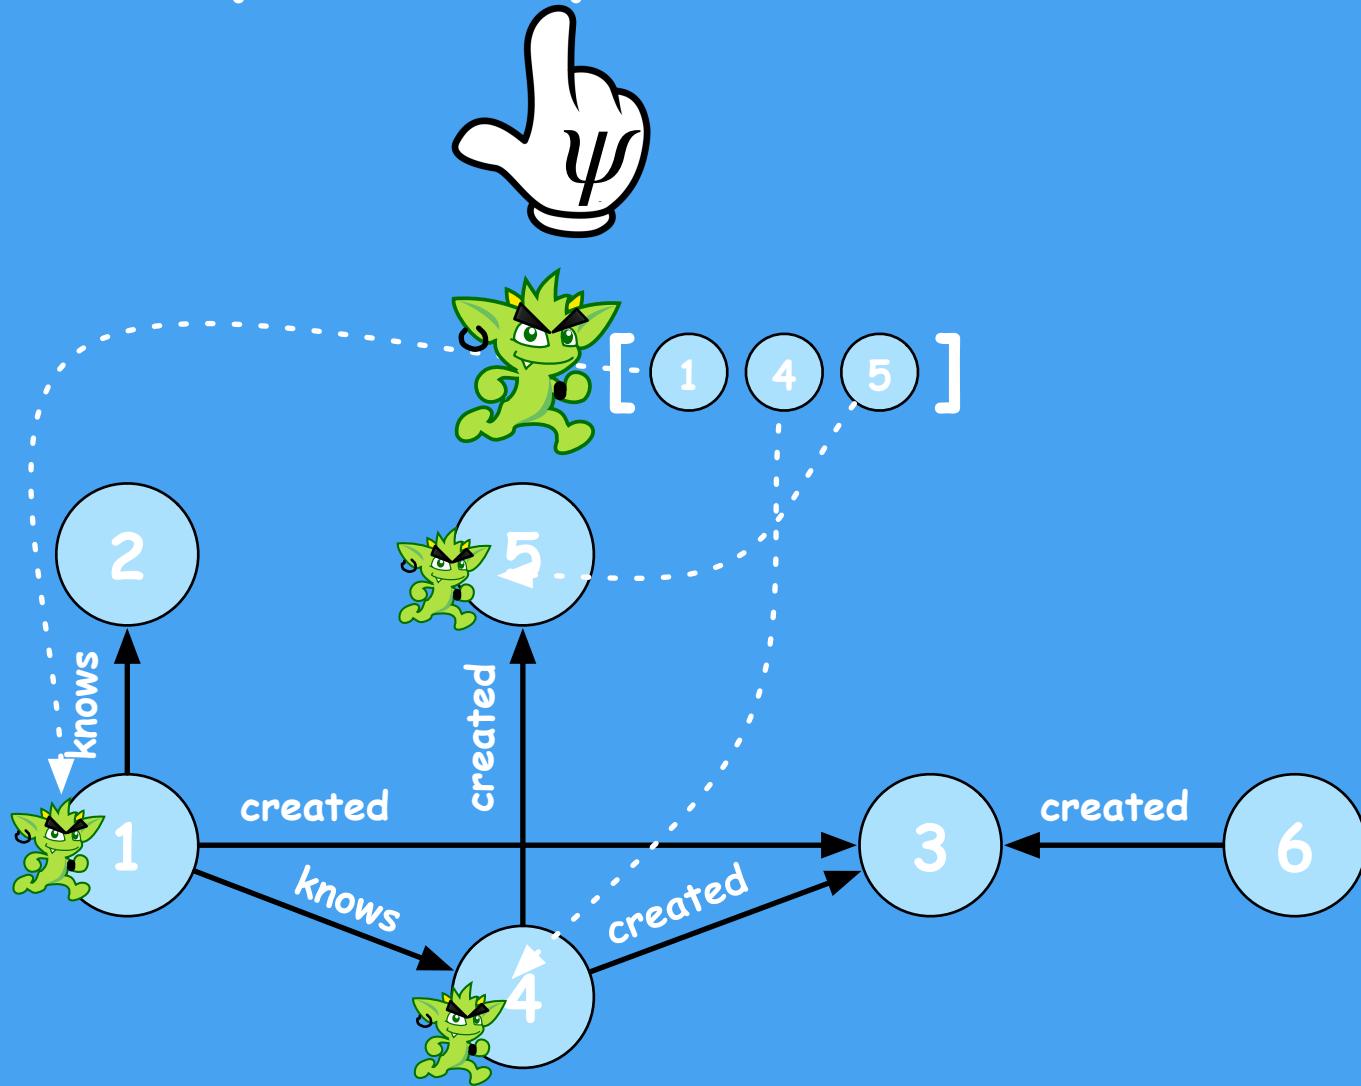
```
g.V(1).out('knows').out('created').limit(1).  
path().by(inE().count())
```



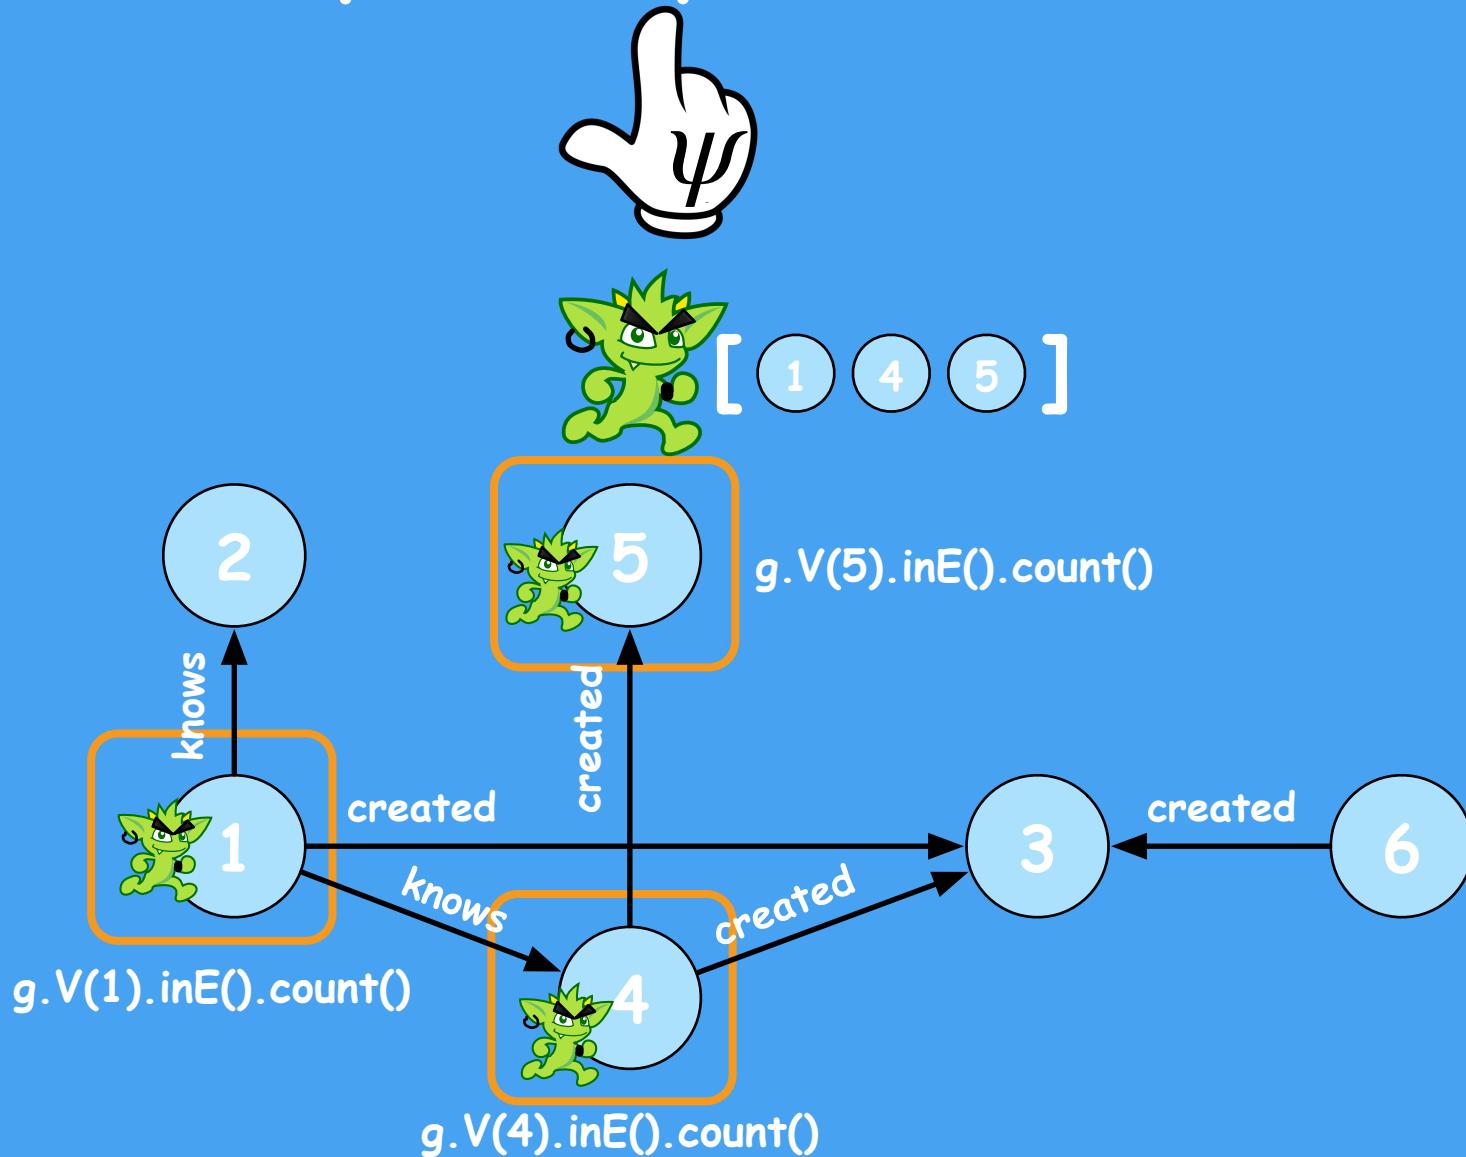
```
g.V(1).out('knows').out('created').limit(1).  
path().by(inE().count())
```



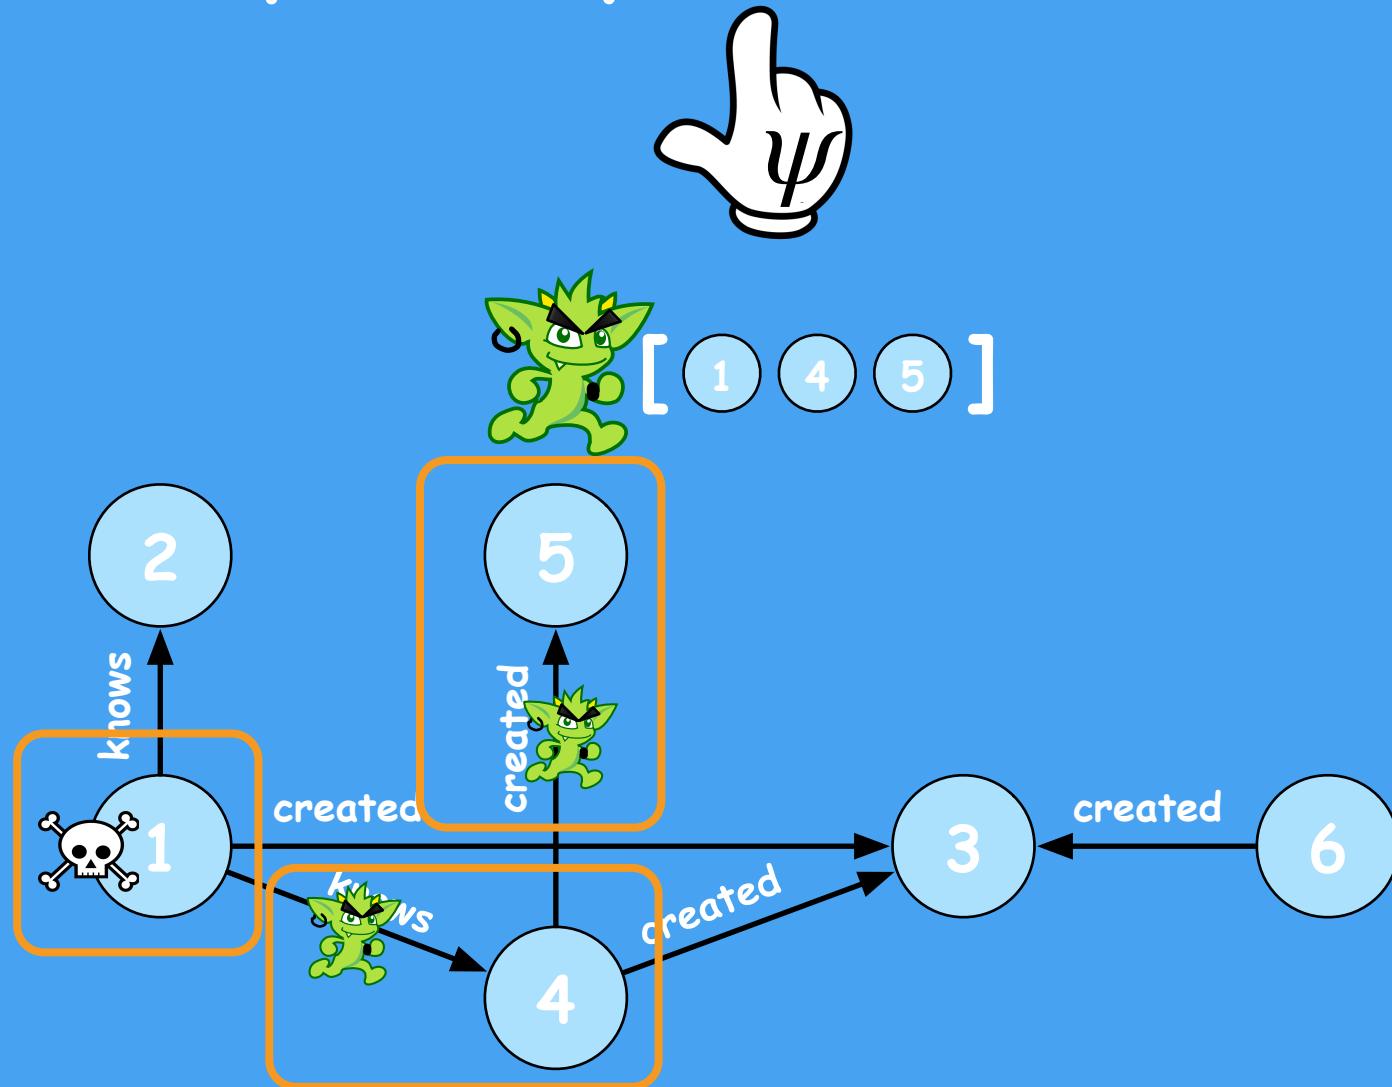
```
g.V(1).out('knows').out('created').limit(1).  
path().by(inE().count())
```



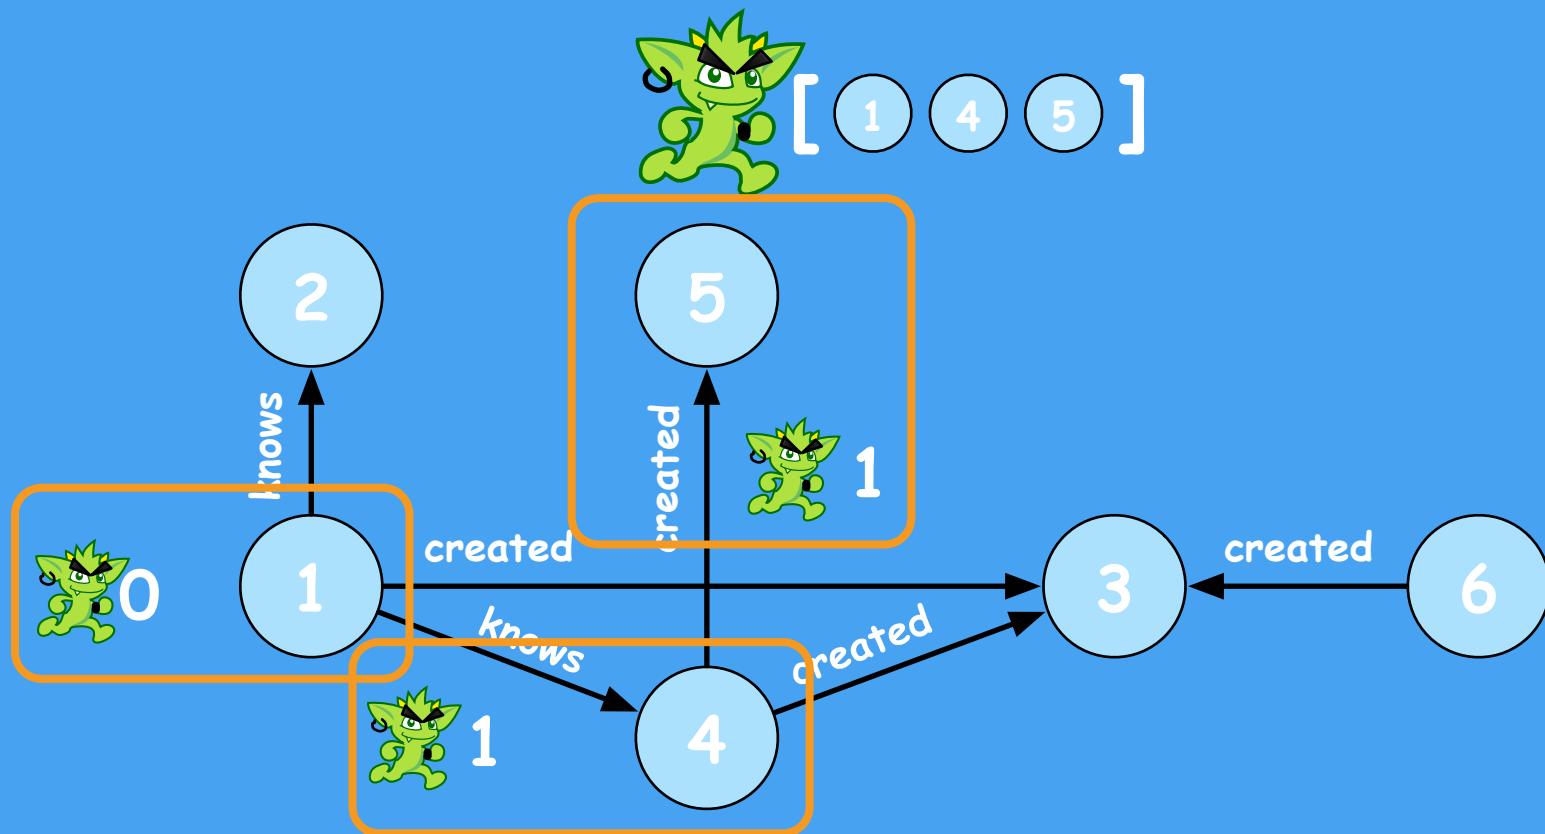
```
g.V(1).out('knows').out('created').limit(1).  
path().by(inE().count())
```



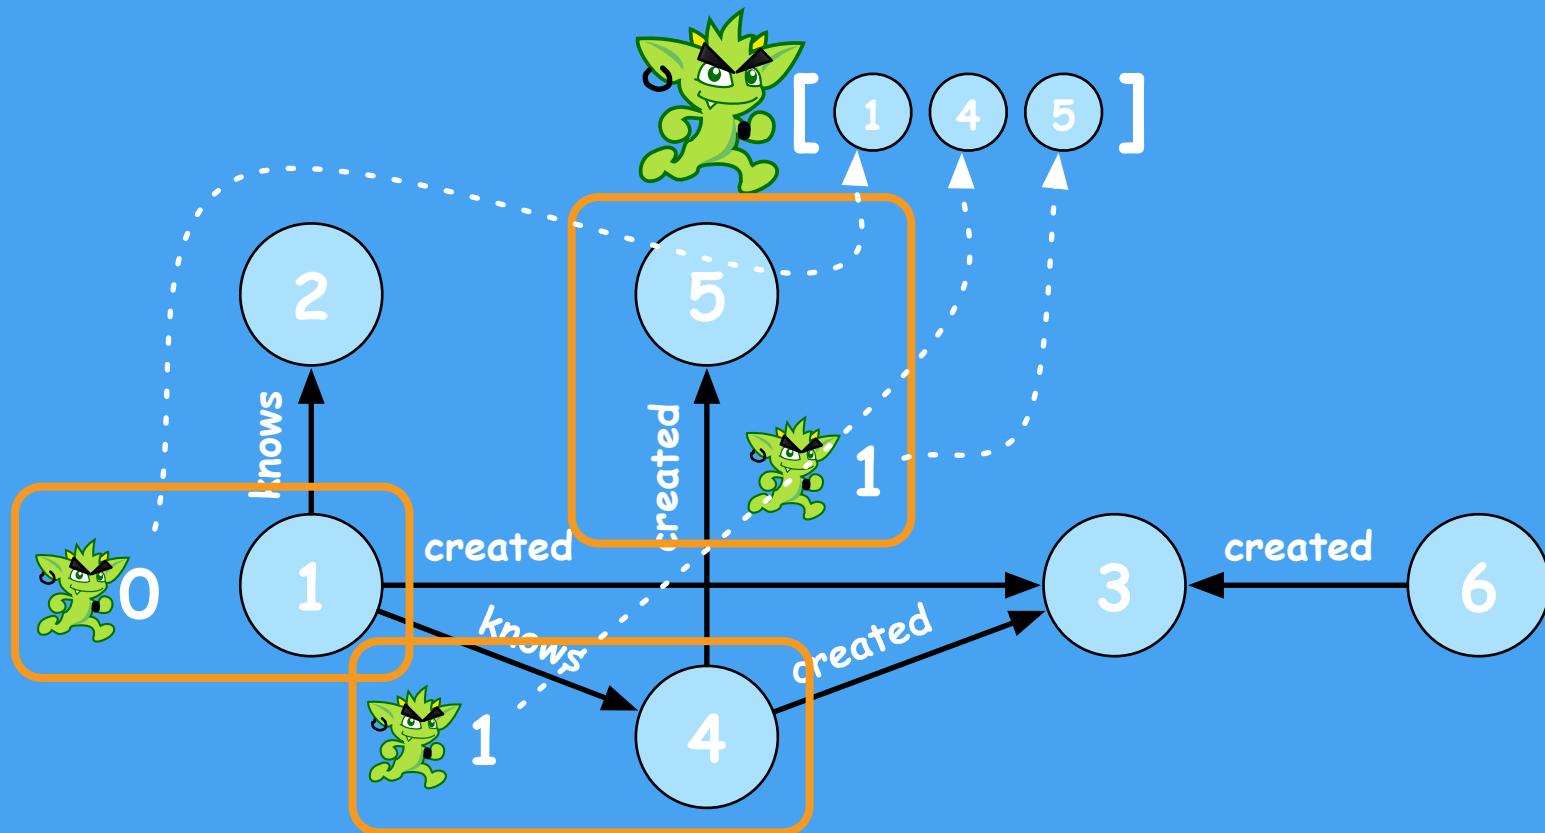
```
g.V(1).out('knows').out('created').limit(1).  
path().by(inE().count())
```



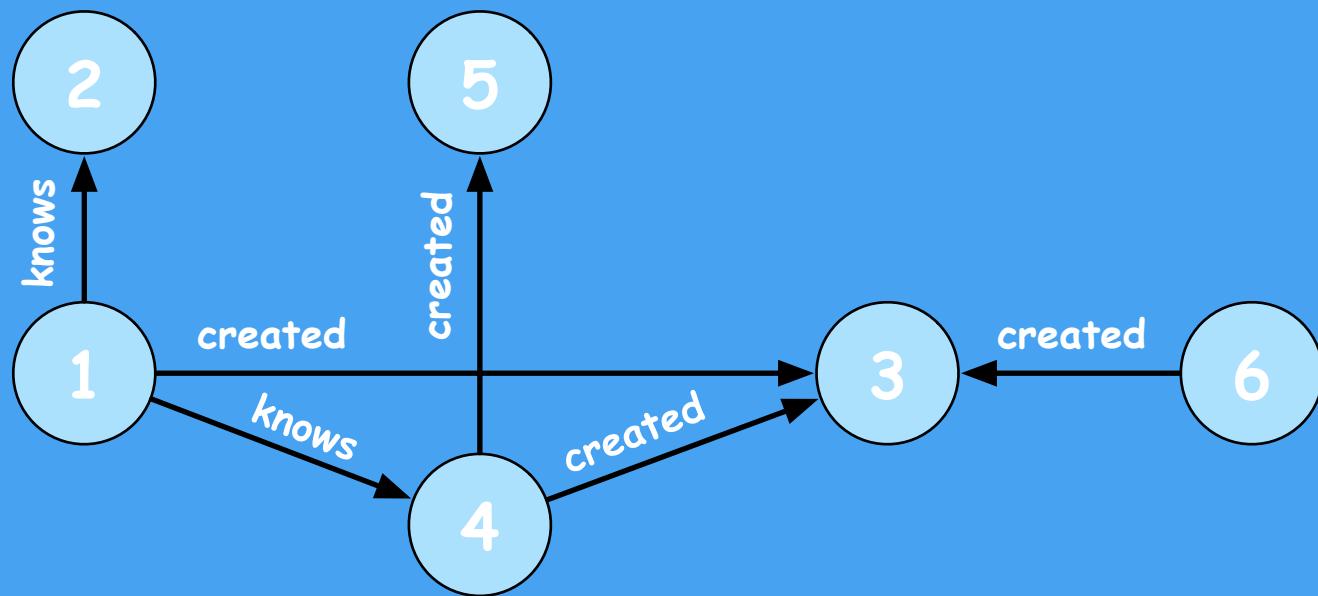
```
g.V(1).out('knows').out('created').limit(1).  
path().by(inE().count())
```



```
g.V(1).out('knows').out('created').limit(1).  
path().by(inE().count())
```

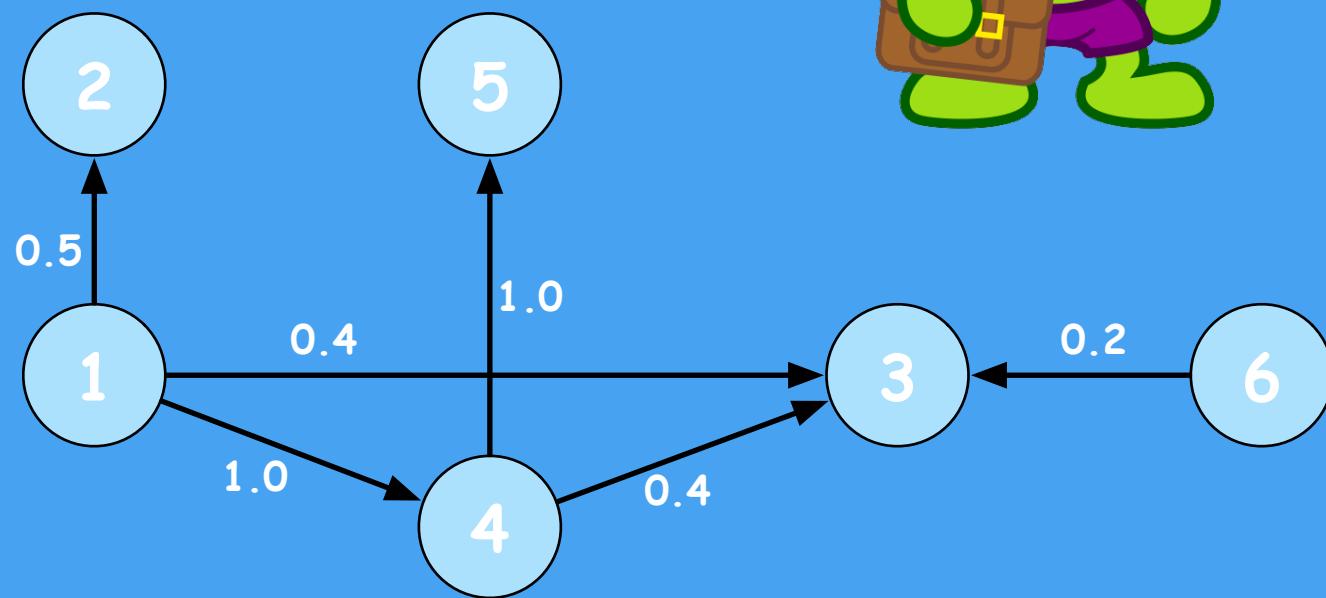


```
g.V(1).out('knows').out('created').limit(1).  
path().by(inE().count())
```



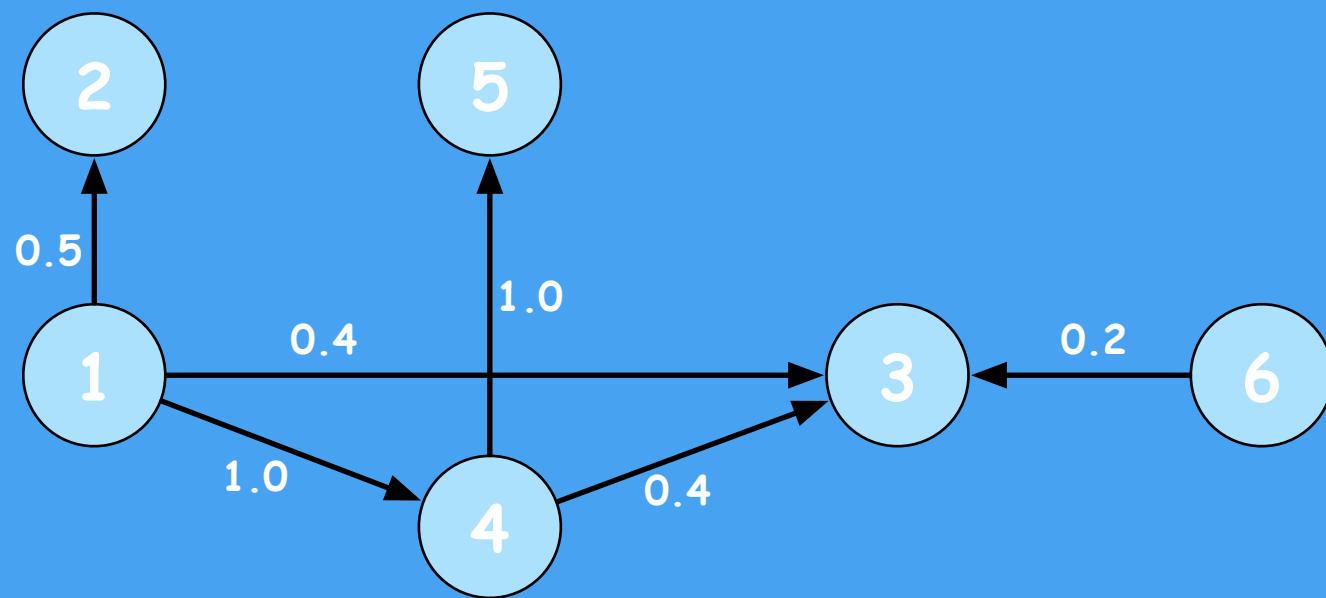
`g.withSack(1).V(1).`

`repeat(outE().sack(mult).by('weight').inV()).times(2).sack()`



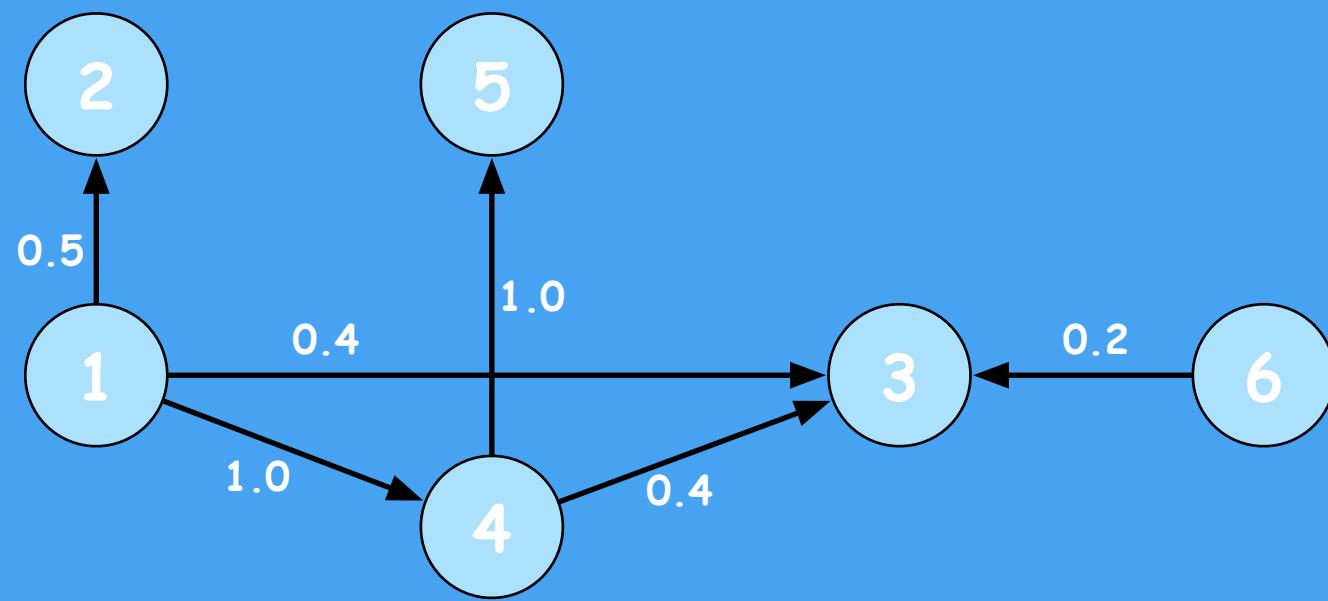
`g.withSack(1).V(1).`

 `.repeat(outE().sack(mult).by('weight').inV()).  
times(2).sack()`



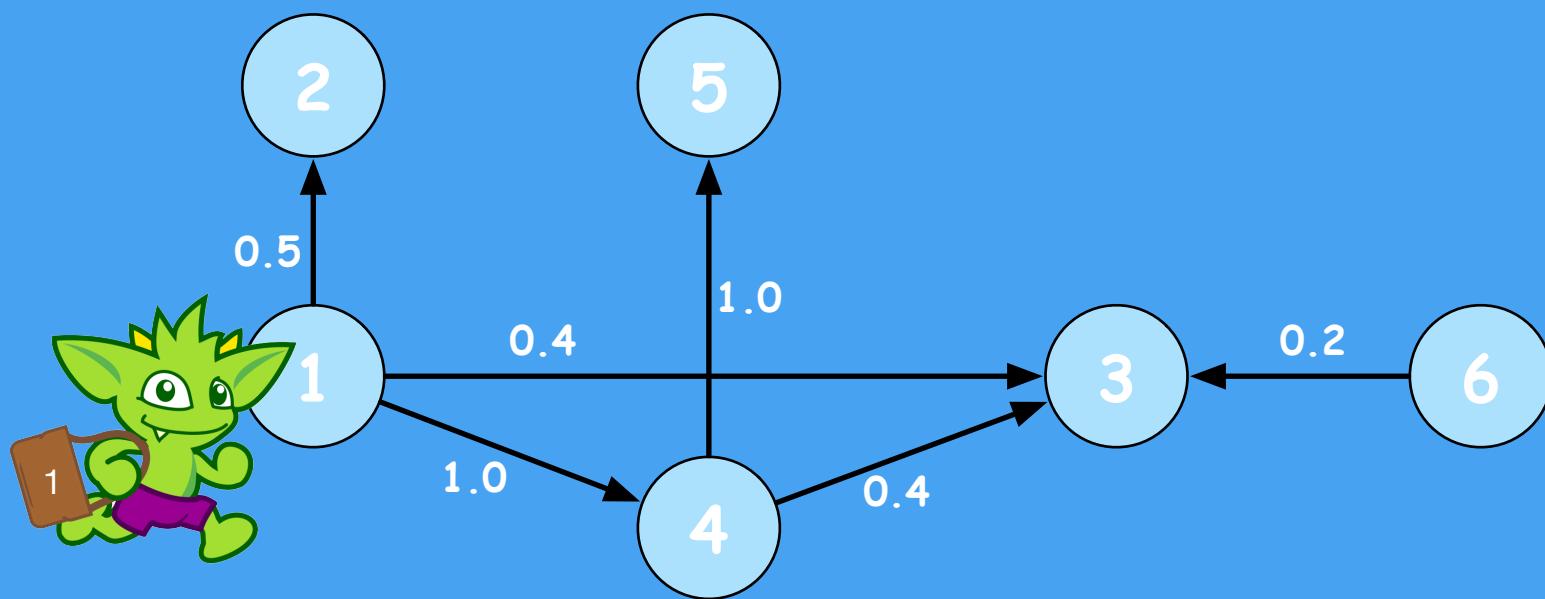
`g.withSack(1).V(1).`

`repeat(outE().sack(mult).by('weight').inV()).  
times(2).sack()`



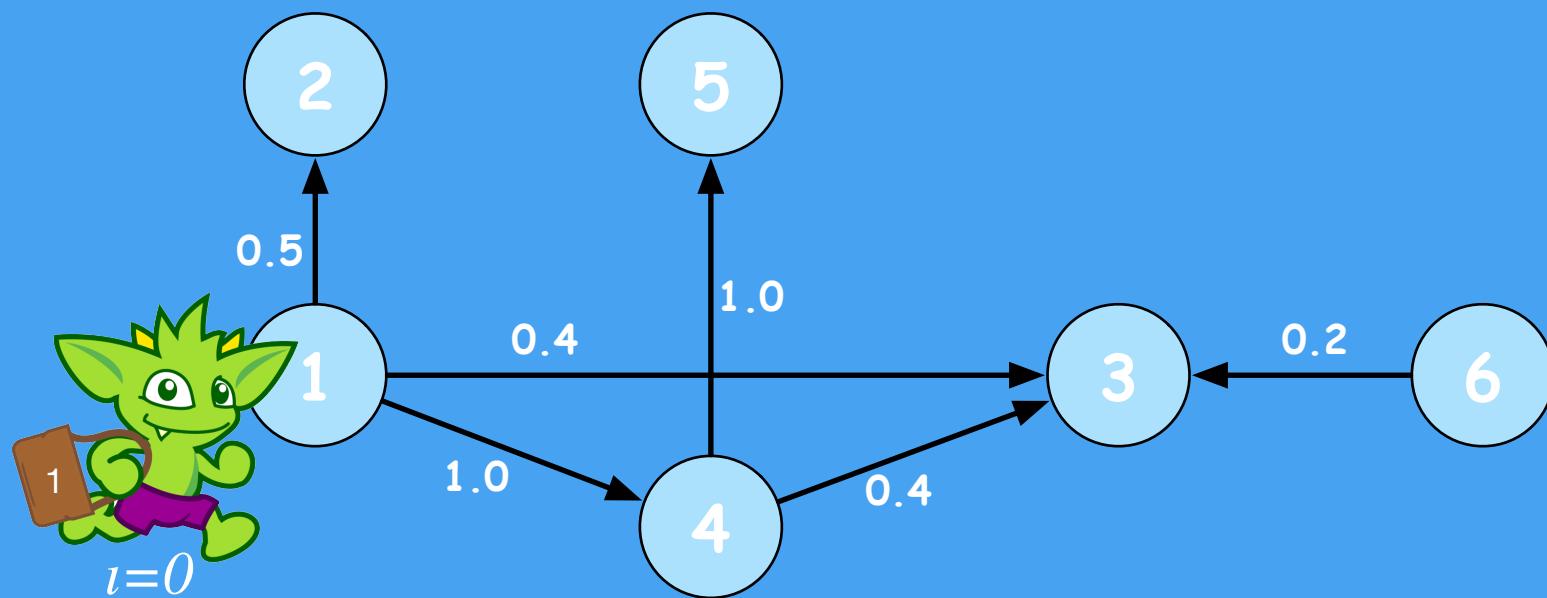
`g.withSack(1).V(1).`

`repeat(outE().sack(mult).by('weight').inV()).  
times(2).sack()`



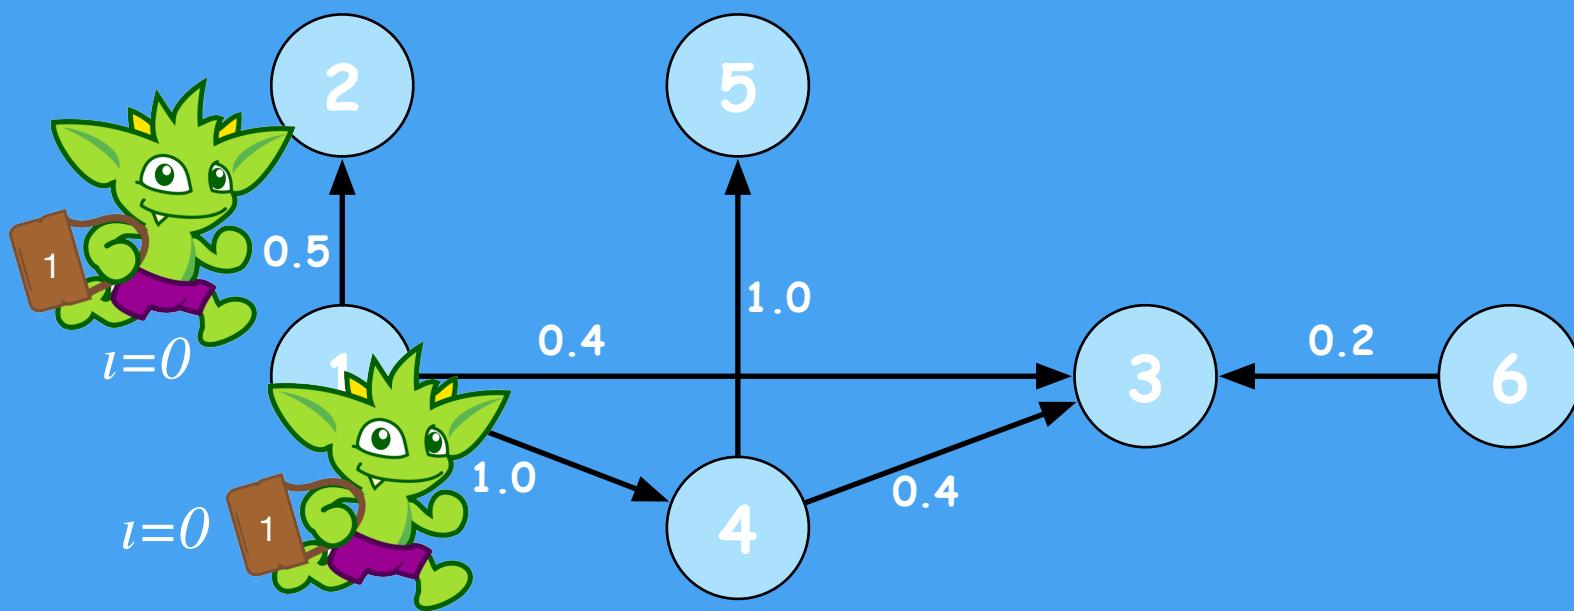
`g.withSack(1).V(1).`

`repeat(outE().sack(mult).by('weight').inV()).  
times(2).sack()`



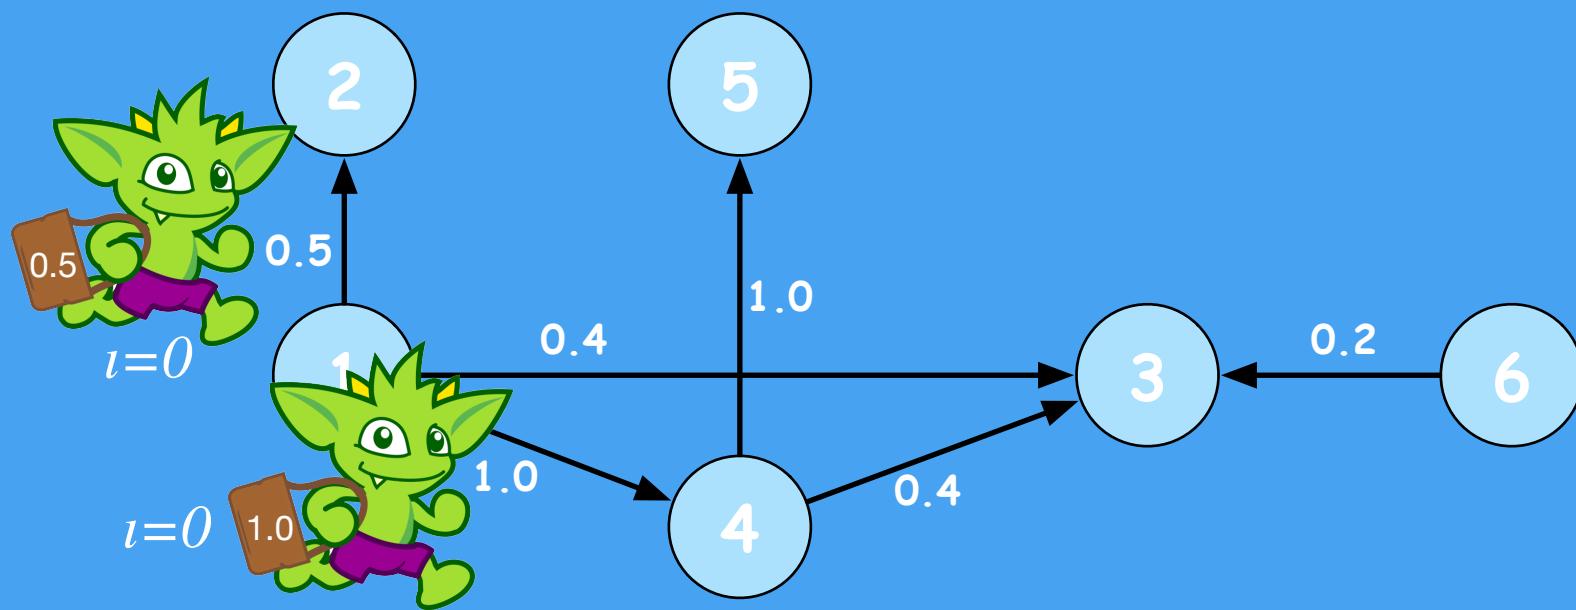
g.withSack(1).V(1).

repeat(outE().sack(mult).by('weight').inV()).  
times(2).sack()



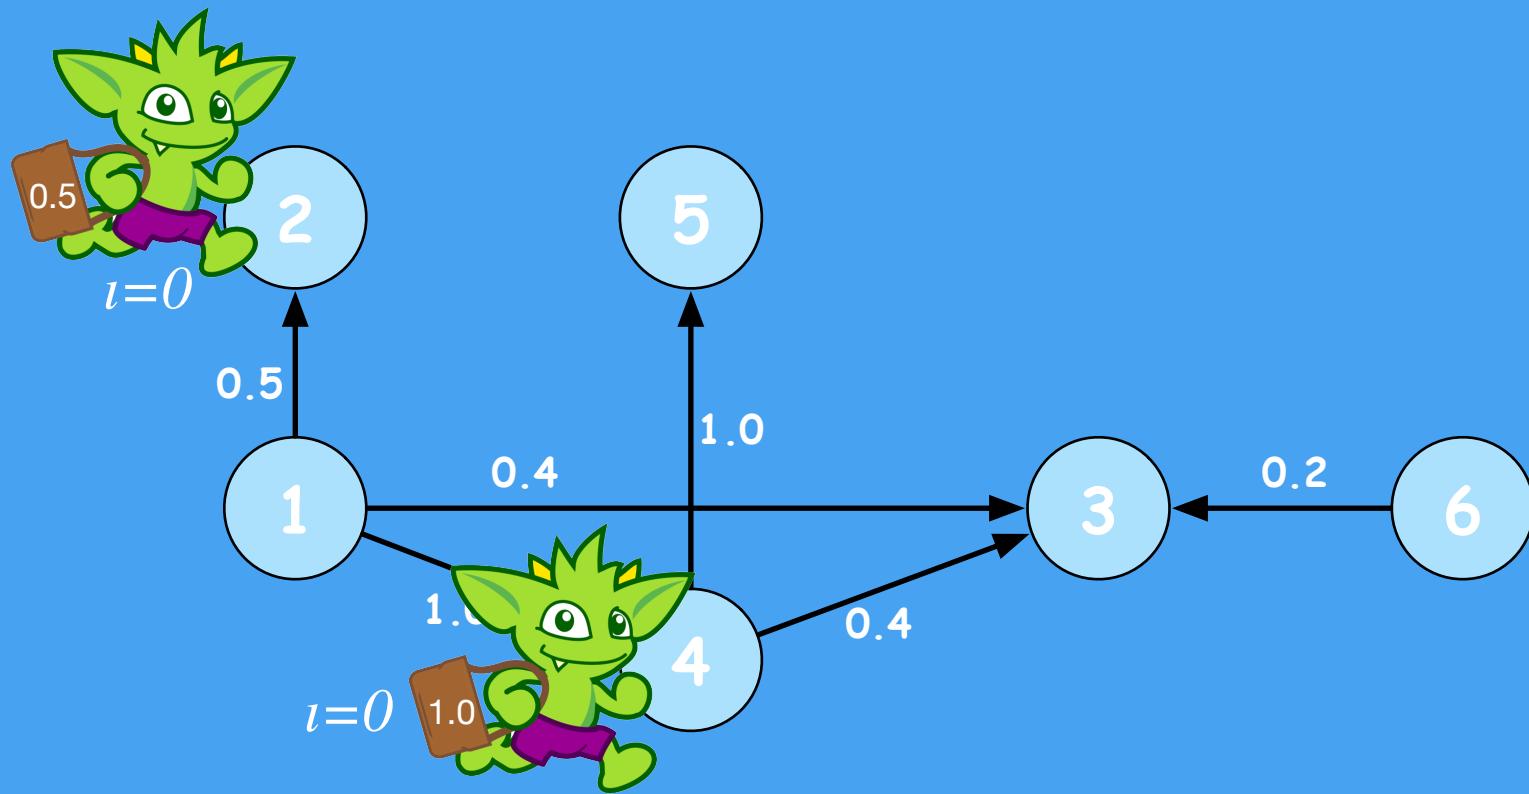
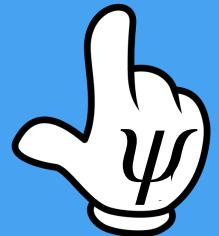
g.withSack(1).V(1).

repeat(outE().**sack(mult).by('weight')**.inV()).  
times(2).sack()



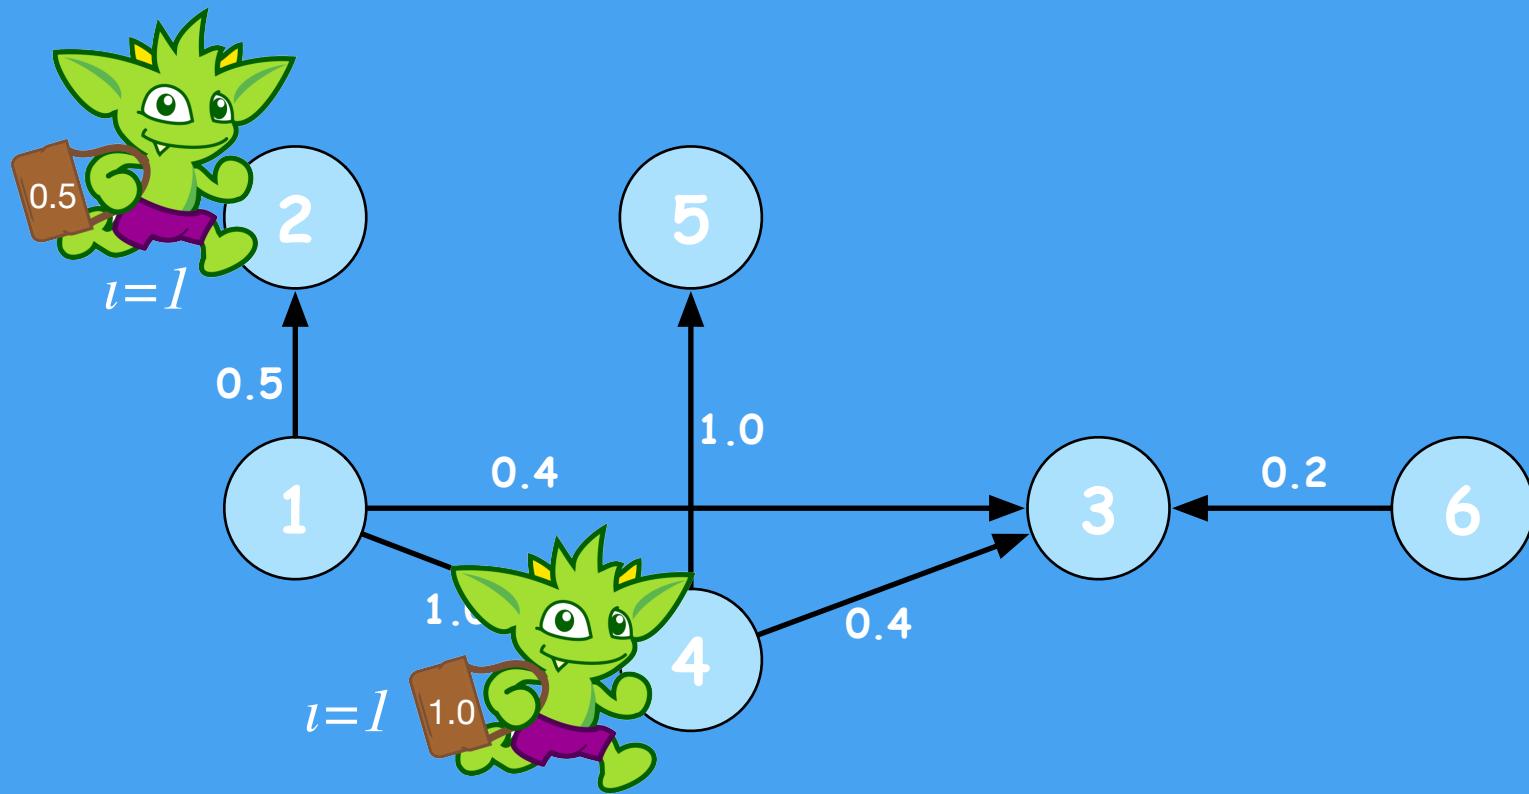
`g.withSack(1).V(1).`

`repeat(outE().sack(mult).by('weight').inV()).times(2).sack()`



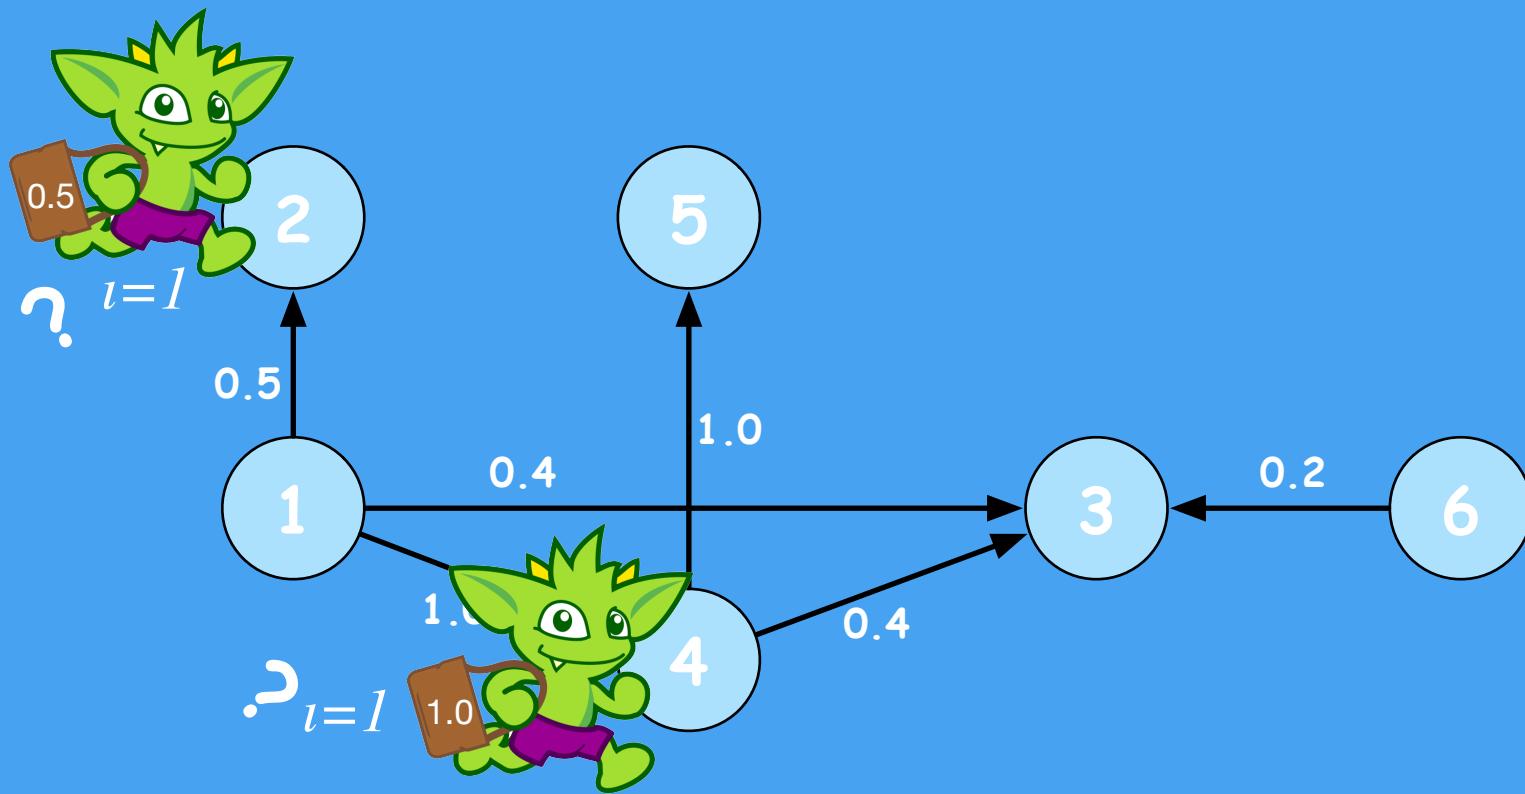
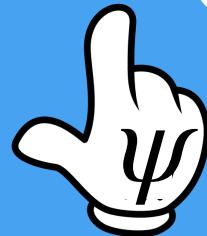
`g.withSack(1).V(1).`

`repeat(outE().sack(mult).by('weight').inV()).times(2).sack()`



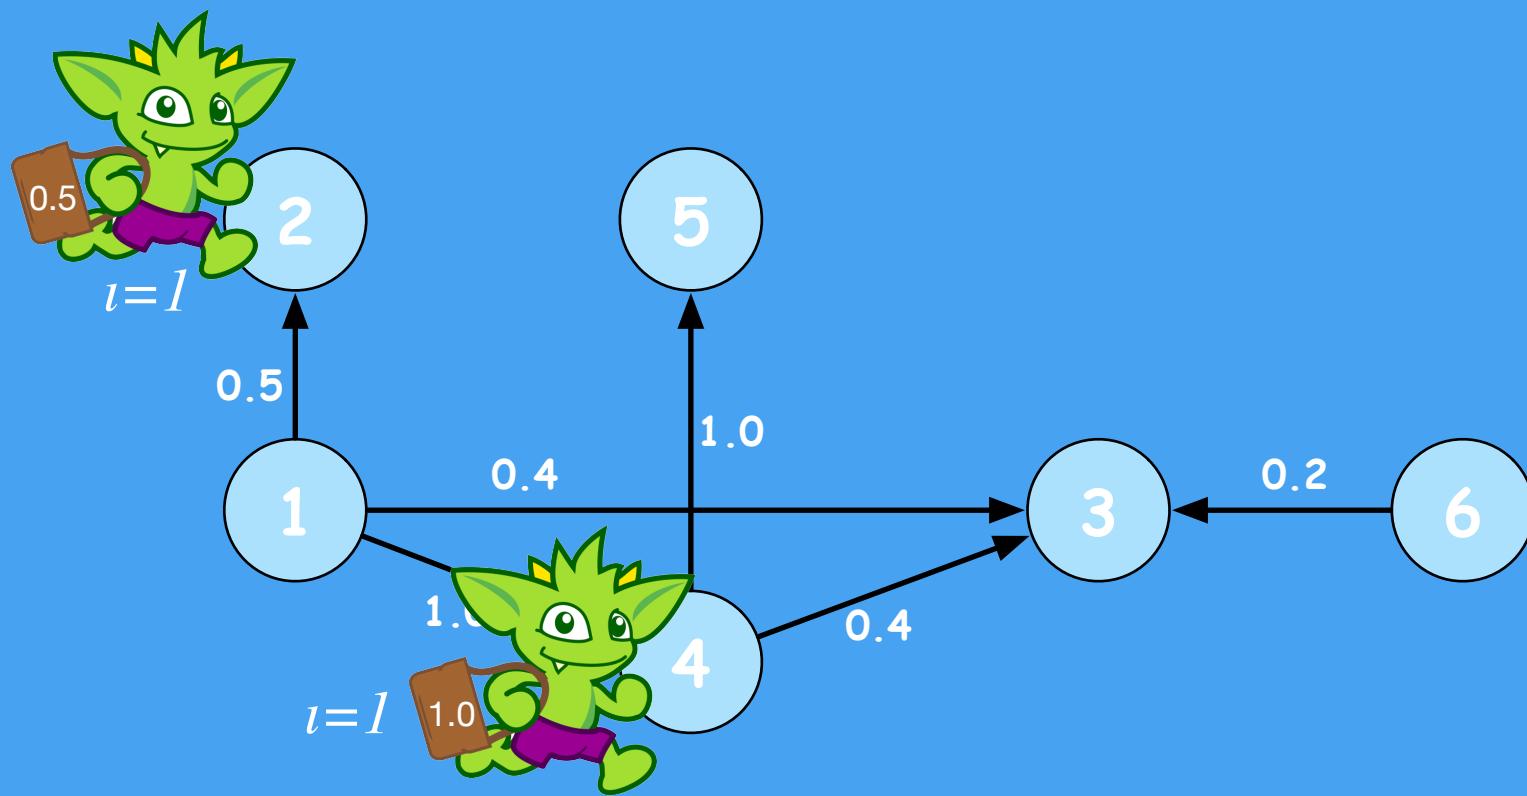
`g.withSack(1).V(1).`

`repeat(outE().sack(mult).by('weight').inV()).times(2).sack()`



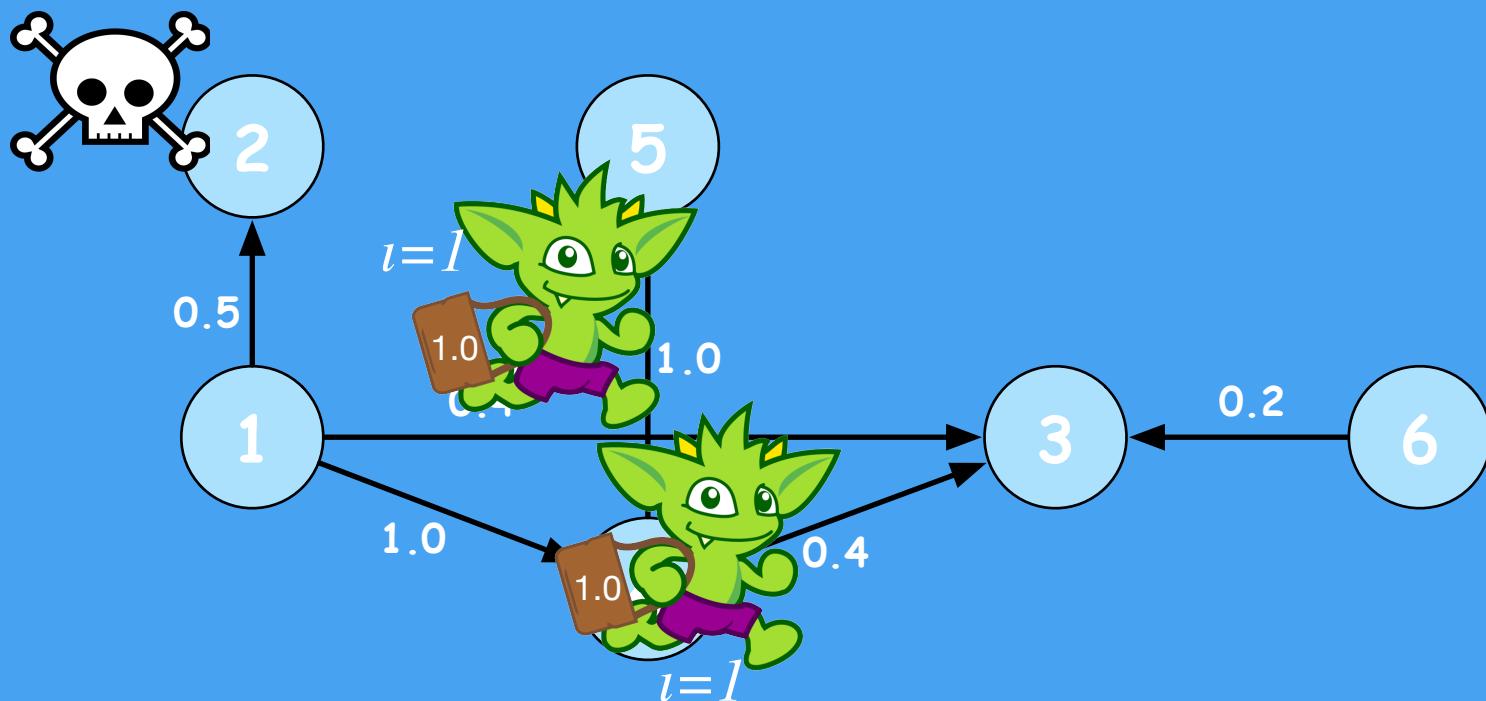
`g.withSack(1).V(1).`

`repeat(outE().sack(mult).by('weight').inV()).times(2).sack()`



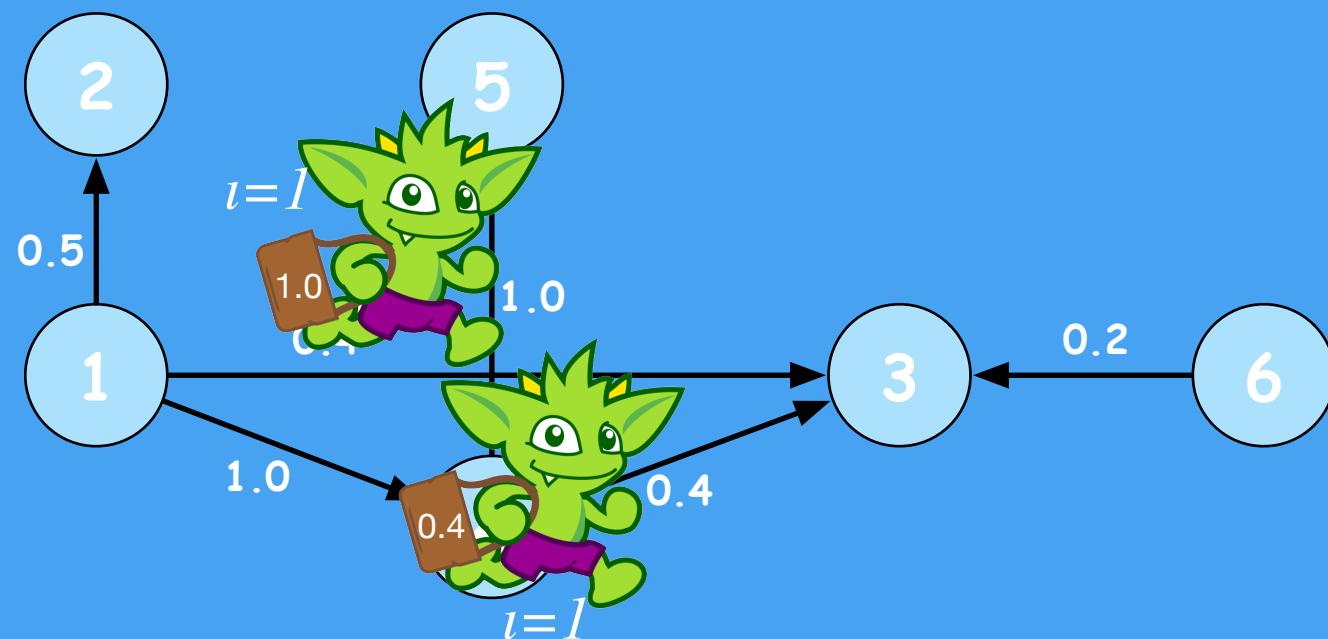
`g.withSack(1).V(1).`

`repeat(outE().sack(mult).by('weight').inV()).times(2).sack()`



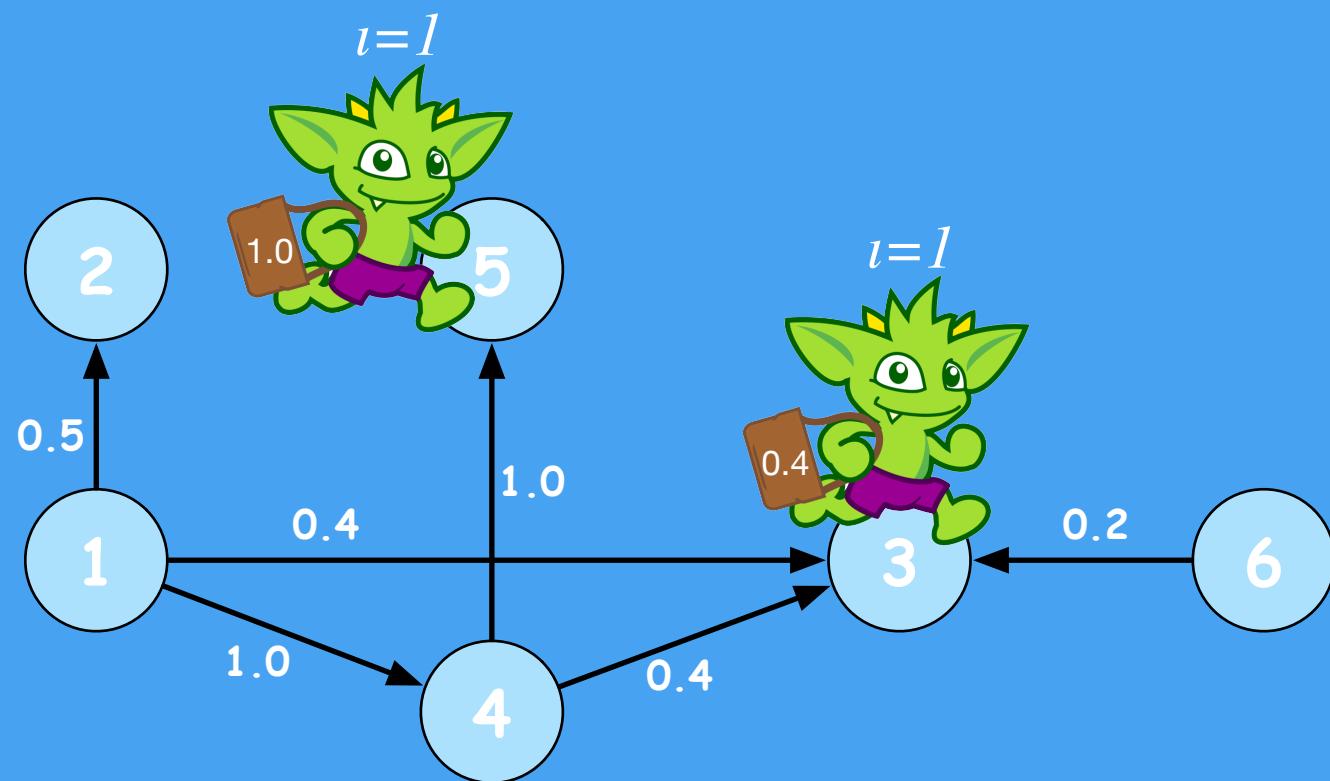
g.withSack(1).V(1).

repeat(outE().**sack(mult).by('weight')**.inV()).  
times(2).sack()



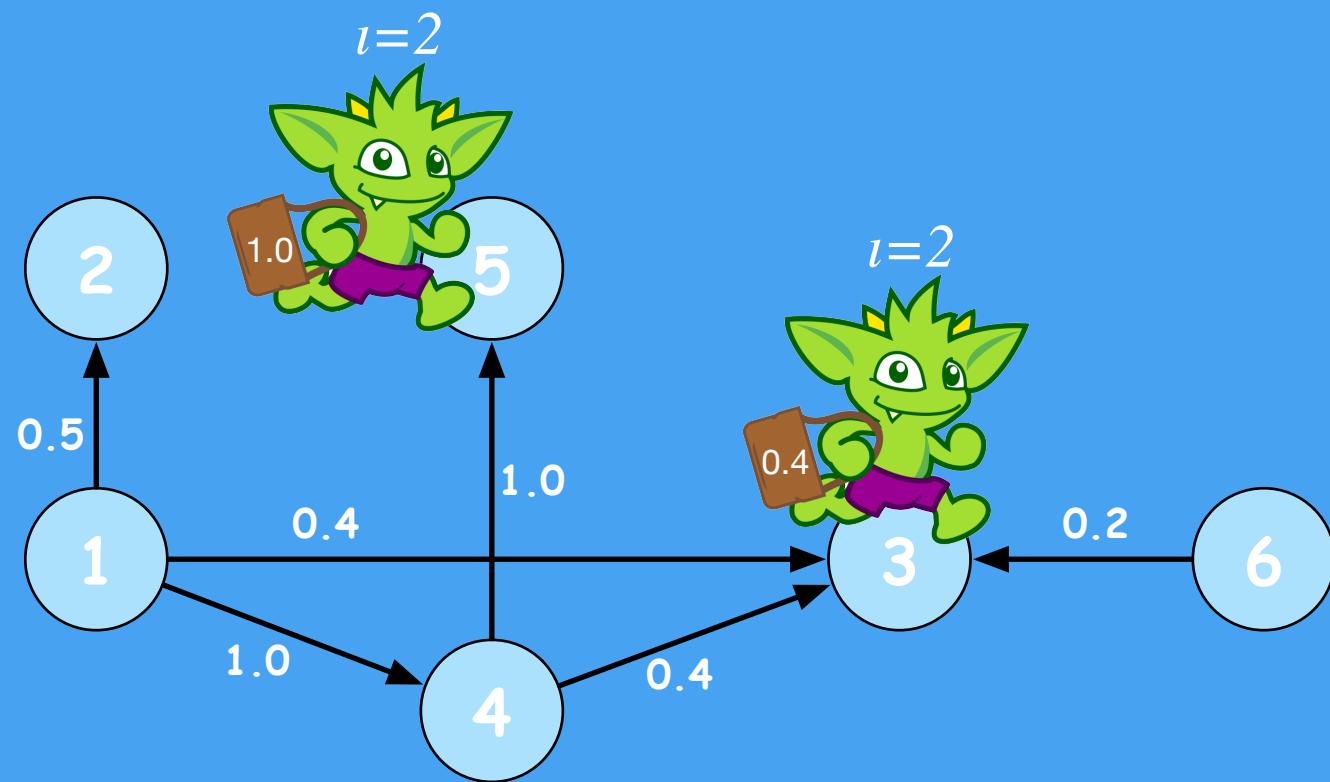
`g.withSack(1).V(1).`

`repeat(outE().sack(mult).by('weight').inV()).times(2).sack()`



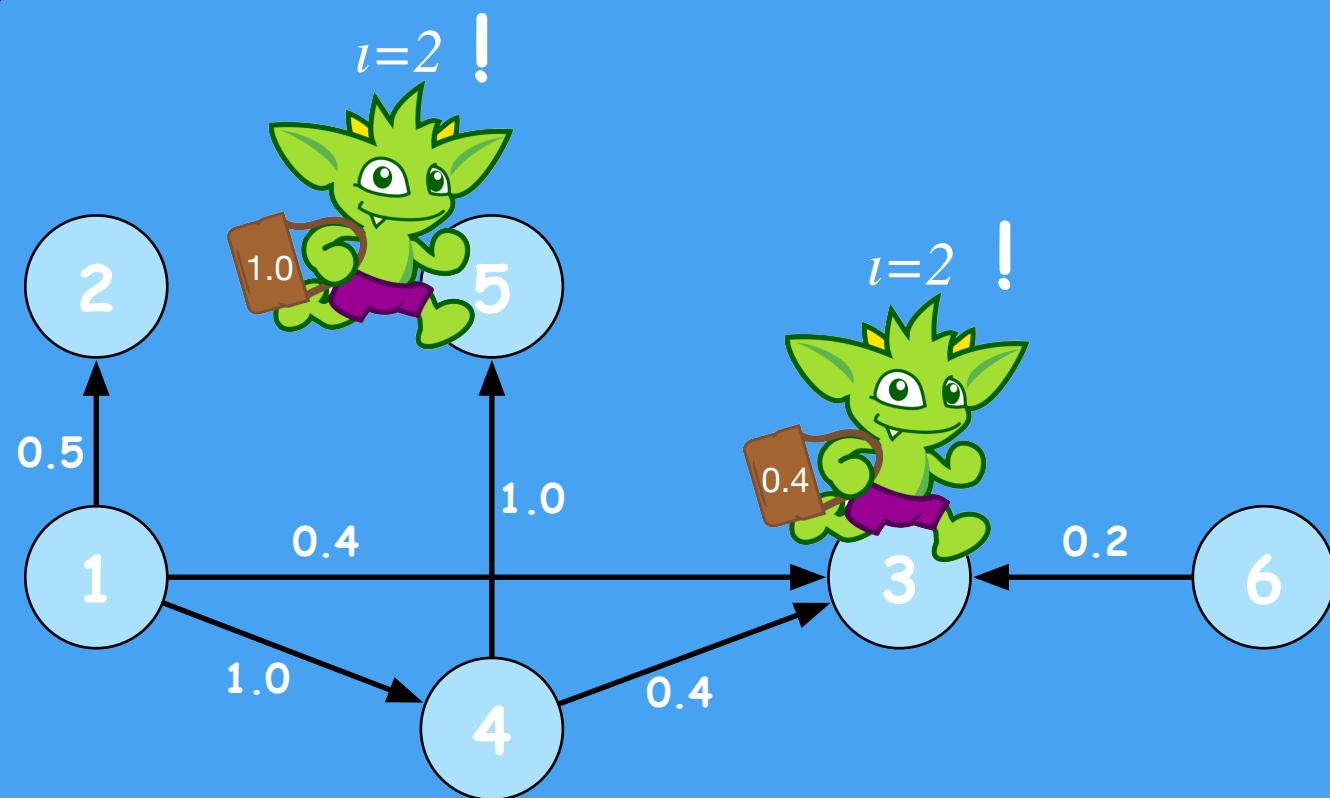
`g.withSack(1).V(1).`

`repeat(outE().sack(mult).by('weight').inV()).times(2).sack()`



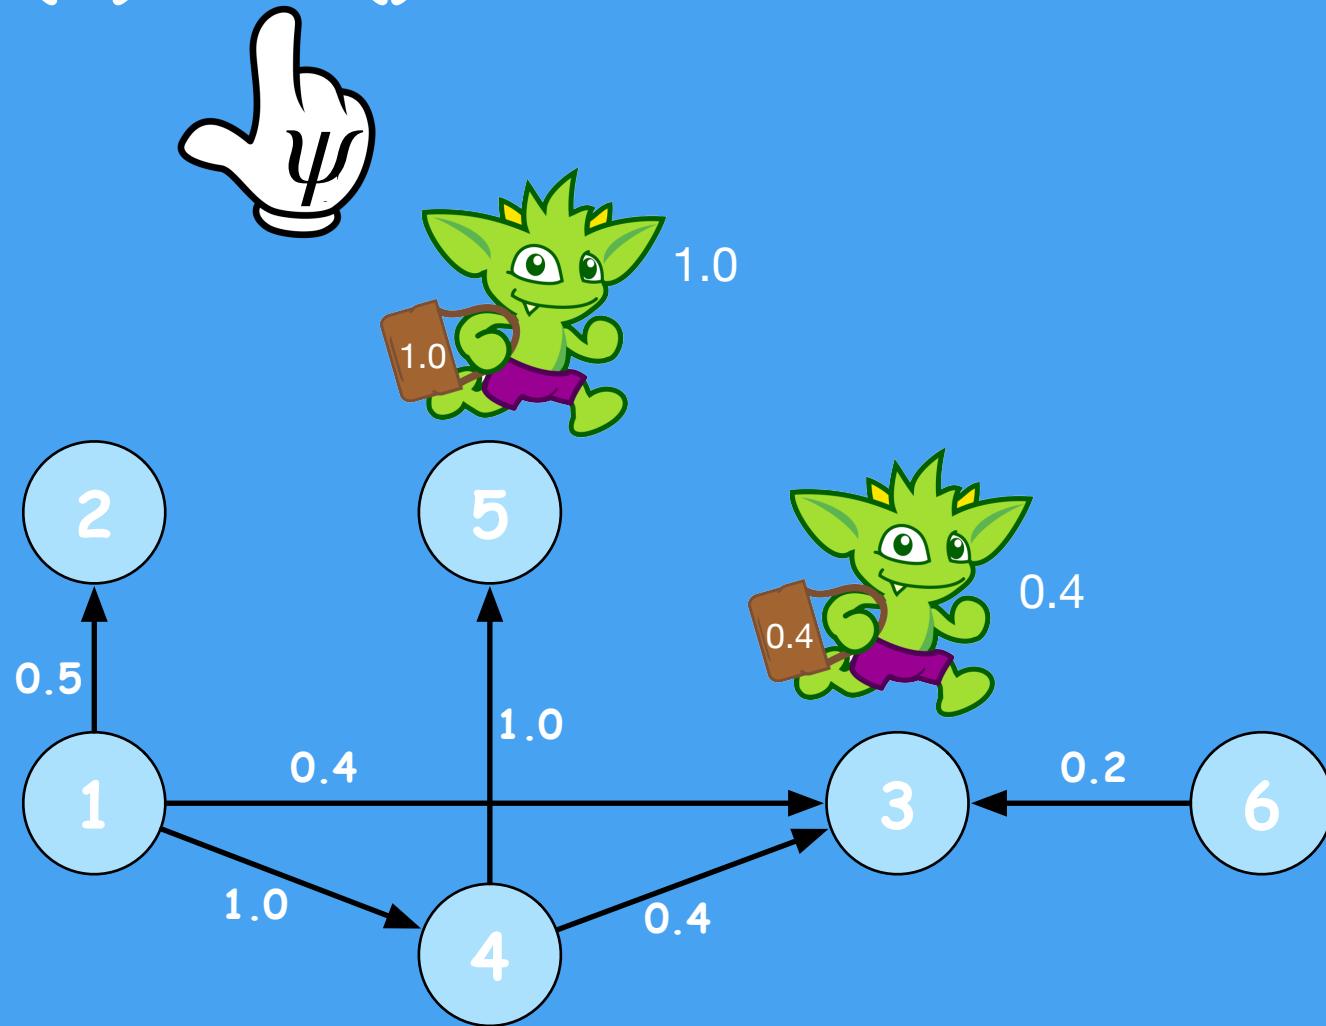
`g.withSack(1).V(1).`

`repeat(outE().sack(mult).by('weight').inV()).times(2).sack()`

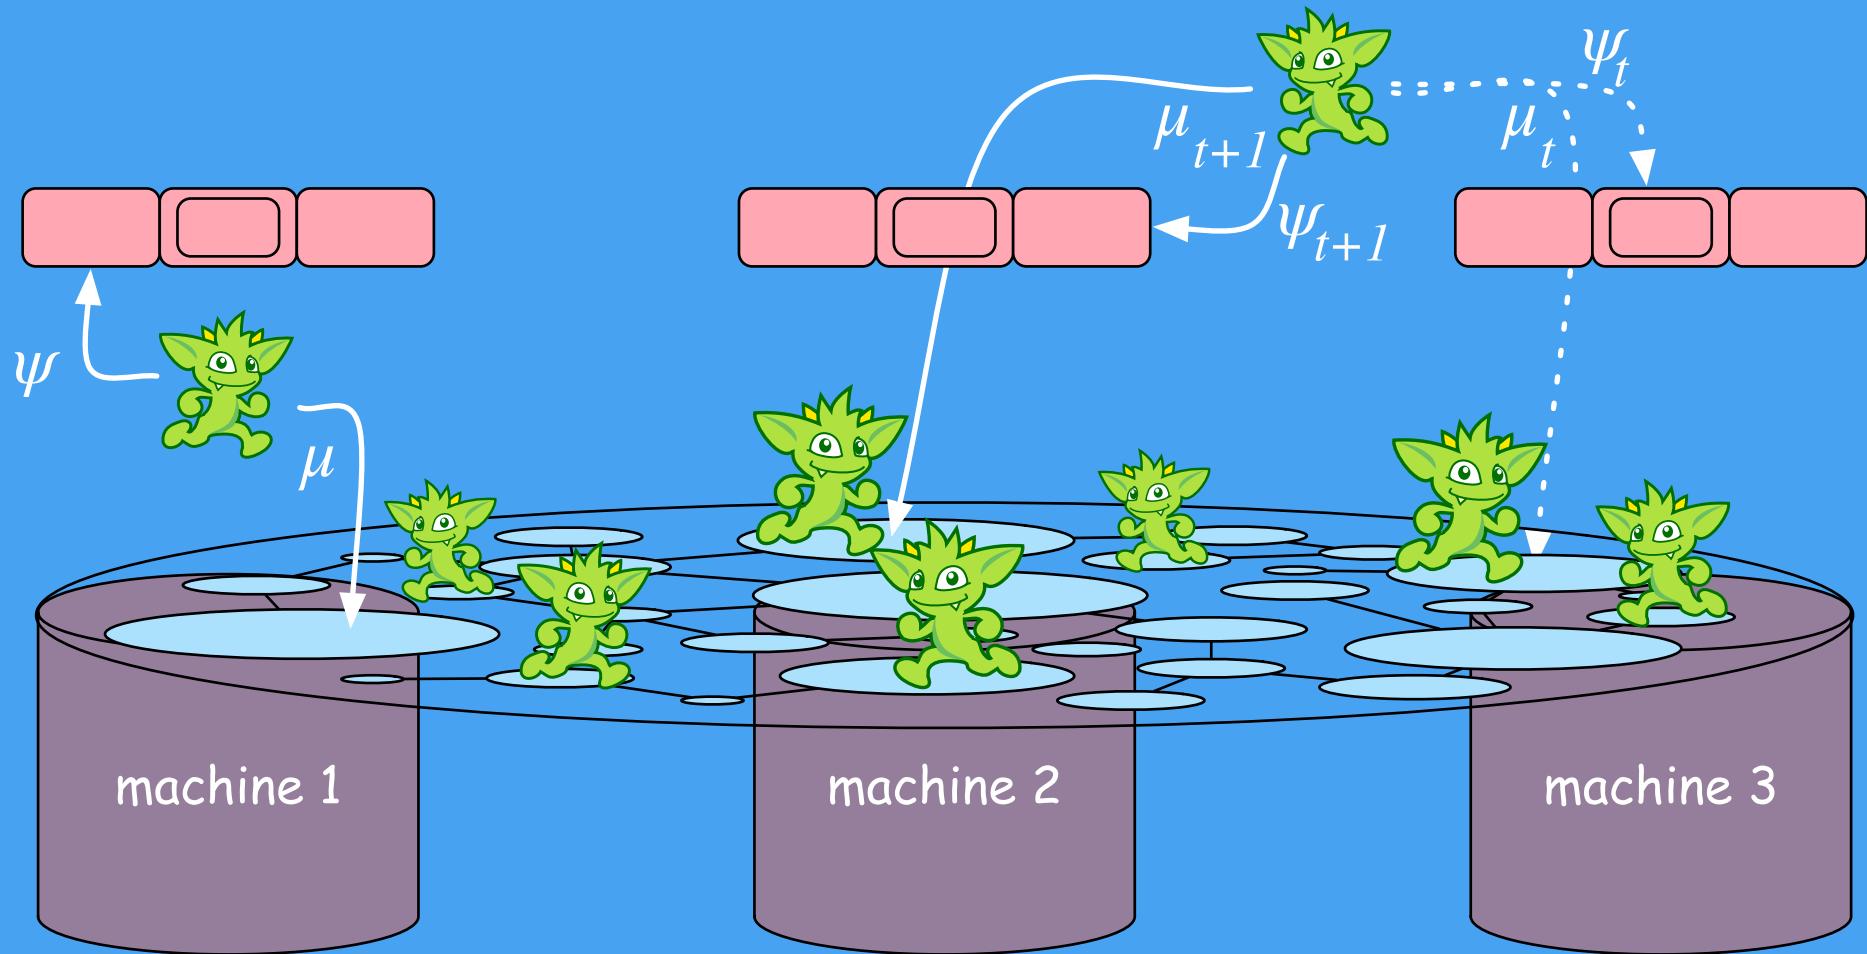


`g.withSack(1).V(1).`

`repeat(outE().sack(mult).by('weight').inV()).times(2).sack()`



The Gremlin traversal machine is a distributed virtual machine — traversers can independently move around the graph (across machines) and through the traversal.





# Quiz

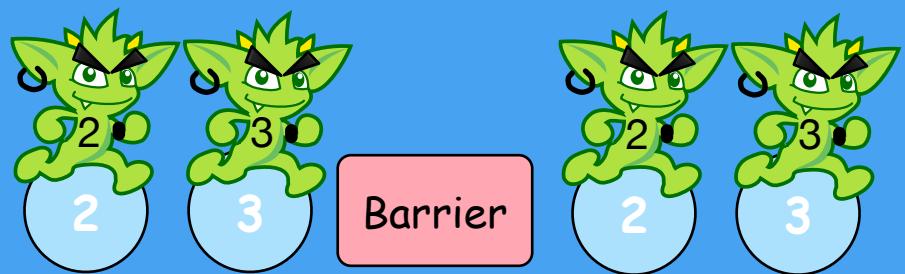
# To Bulk or Not To Bulk



# To Bulk or Not To Bulk



# To Bulk or Not To Bulk



[ 1 4 5 ] [ 1 2 5 ]



# To Bulk or Not To Bulk



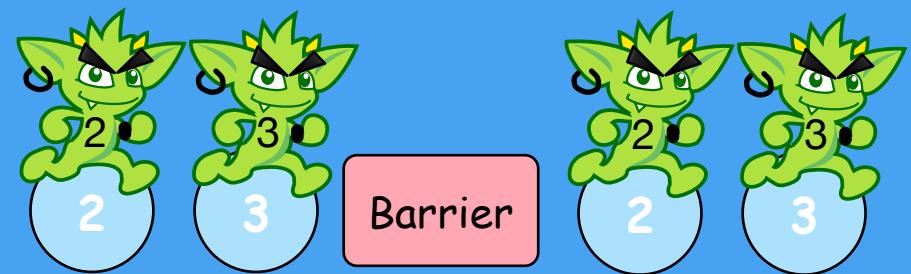
[ 1 4 5 ] [ 1 2 5 ] [ 1 4 5 ] [ 1 2 5 ]



withSack(?, sum)



# To Bulk or Not To Bulk



[ (1 4 5) ] [ (1 2 5) ] [ (1 4 5) ] [ (1 2 5) ]



withSack(?, sum)



$i=1 \quad i=2$

# To Bulk or Not To Bulk



[ 1 4 5 ] [ 1 2 5 ] [ 1 4 5 ] [ 1 2 5 ]



withSack(?, sum)



$i=1 \quad i=2$



$i=1 \quad i=2$



# To Bulk or Not To Bulk



[ 1 4 5 ] [ 1 2 5 ] [ 1 4 5 ] [ 1 2 5 ]



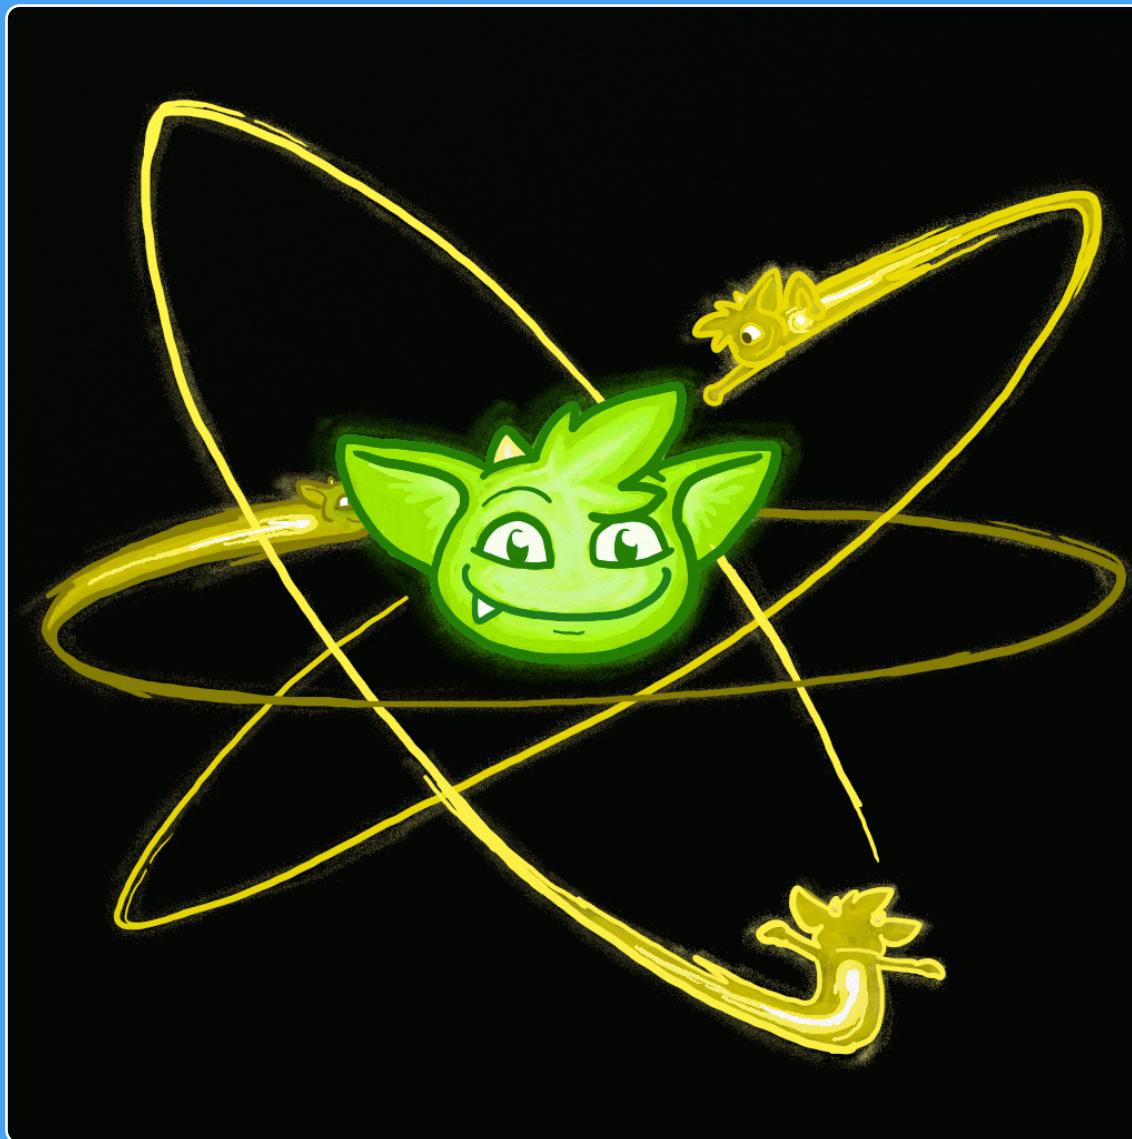
withSack(?, sum)



$i=1 \quad i=2$



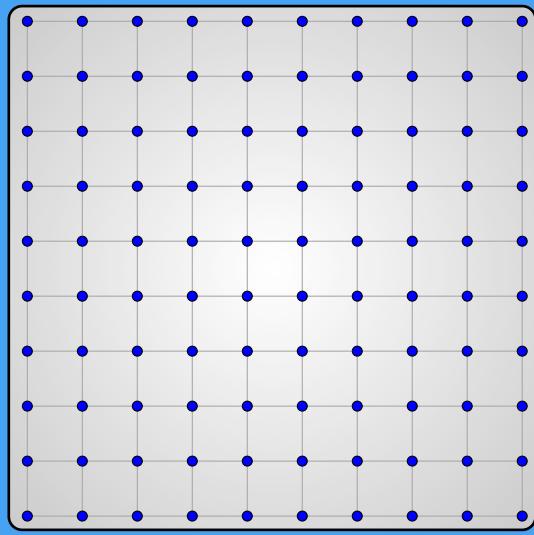
The traverser is an atomic unit of computing — moving through both the traversal and the graph.



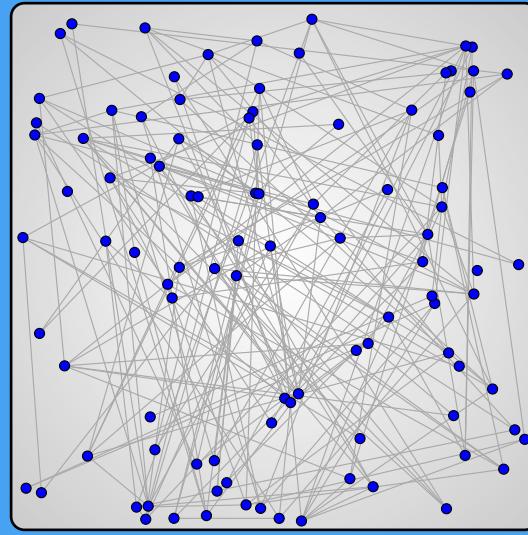
# G

## The Graph

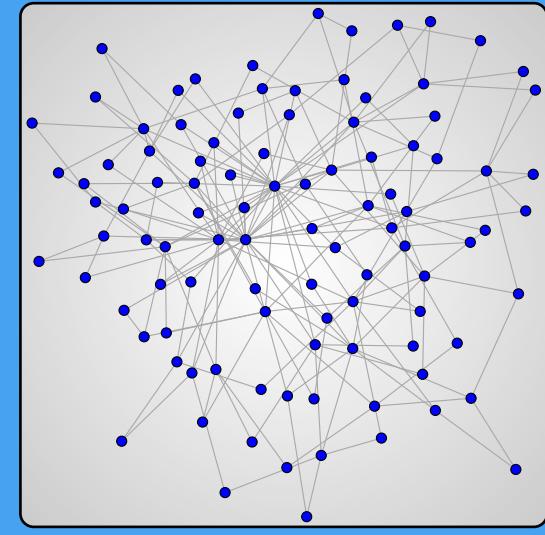
# Lattice



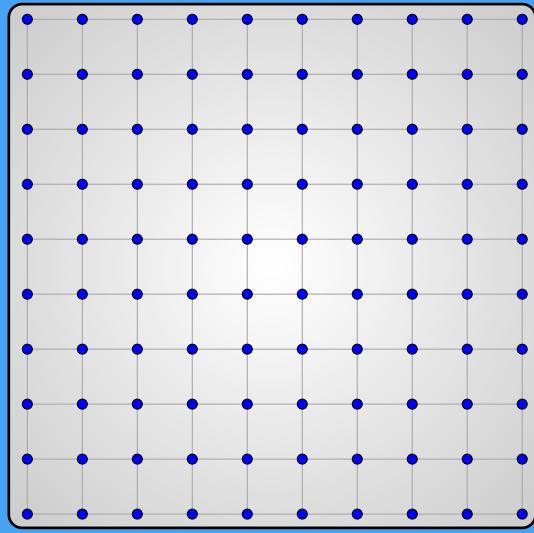
# Scale-Free



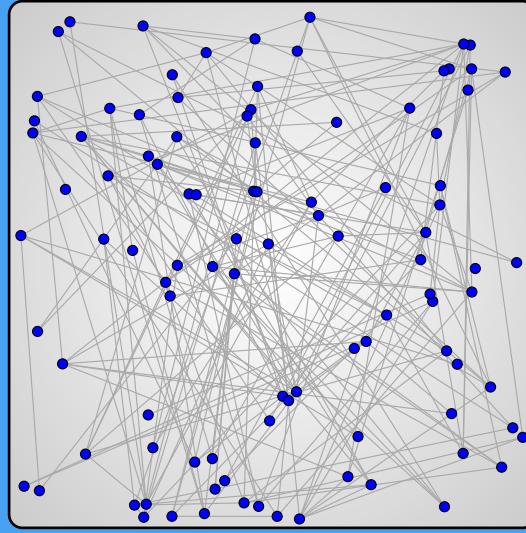
# Random



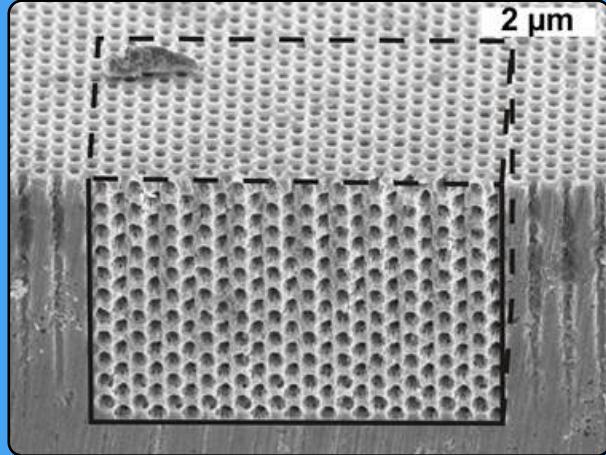
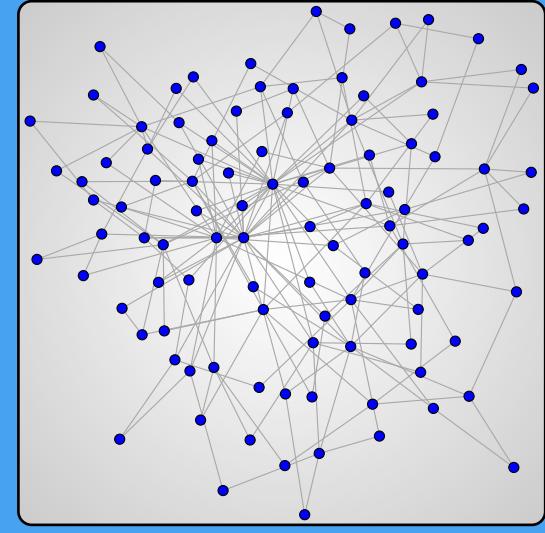
# Lattice



# Scale-Free



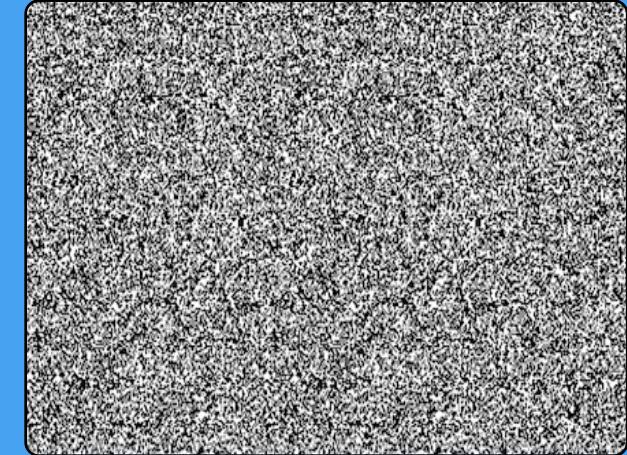
# Random



# Diamond

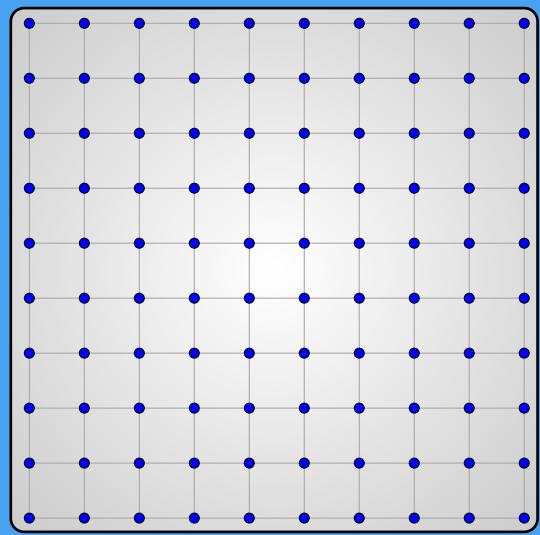


# Jackson Pollock

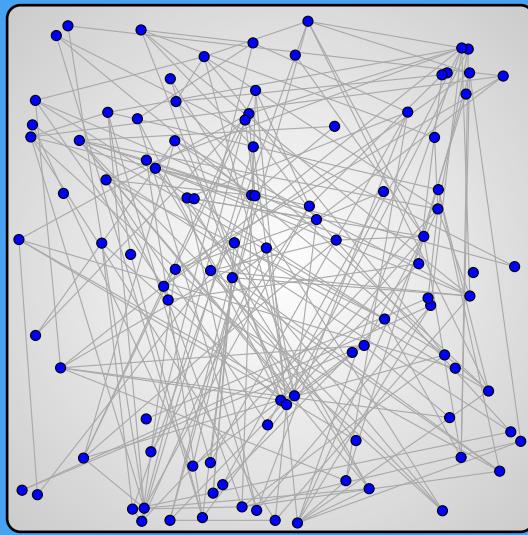


# TV Fuzz

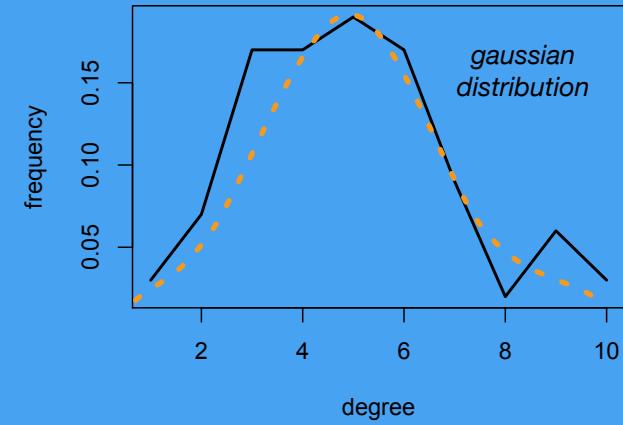
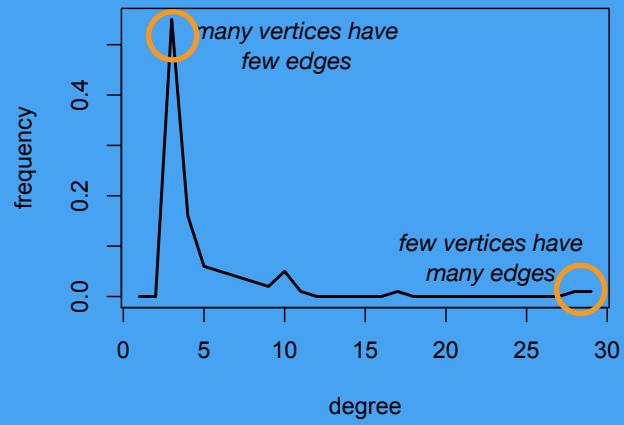
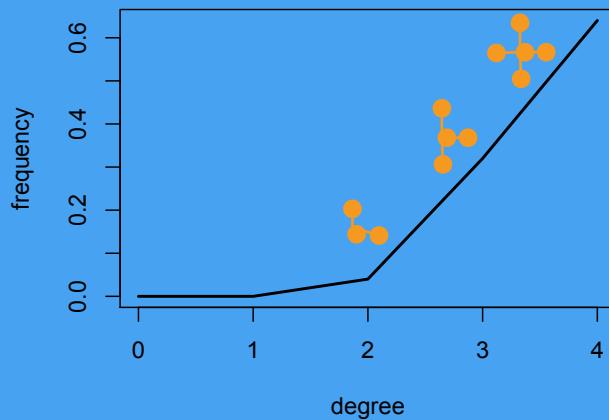
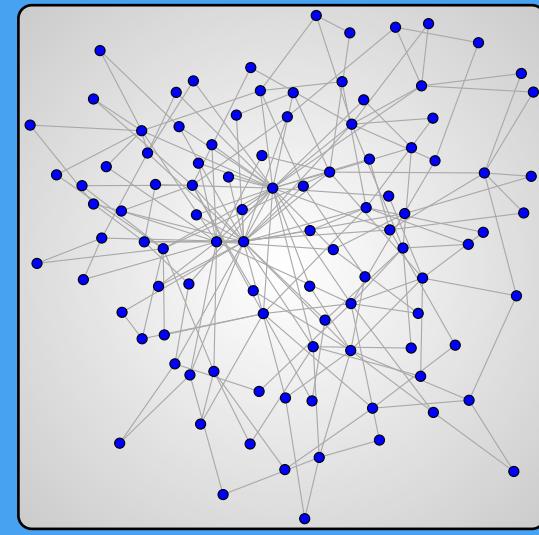
# Lattice



# Scale-Free

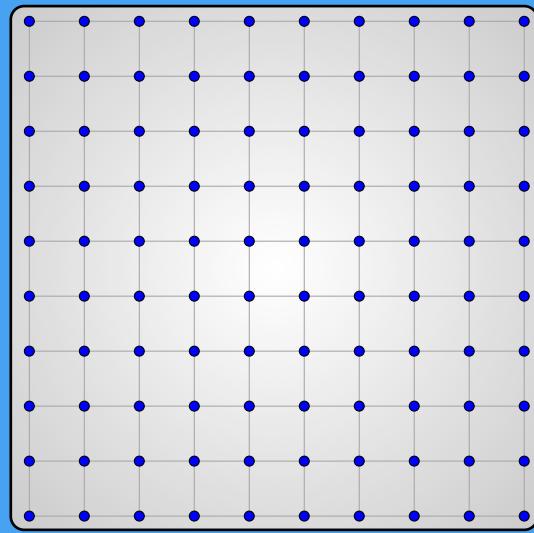


# Random

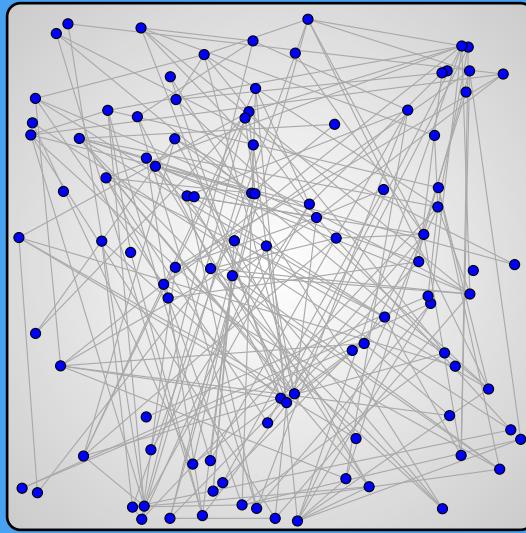


`g.V().groupCount().by(both().count())`

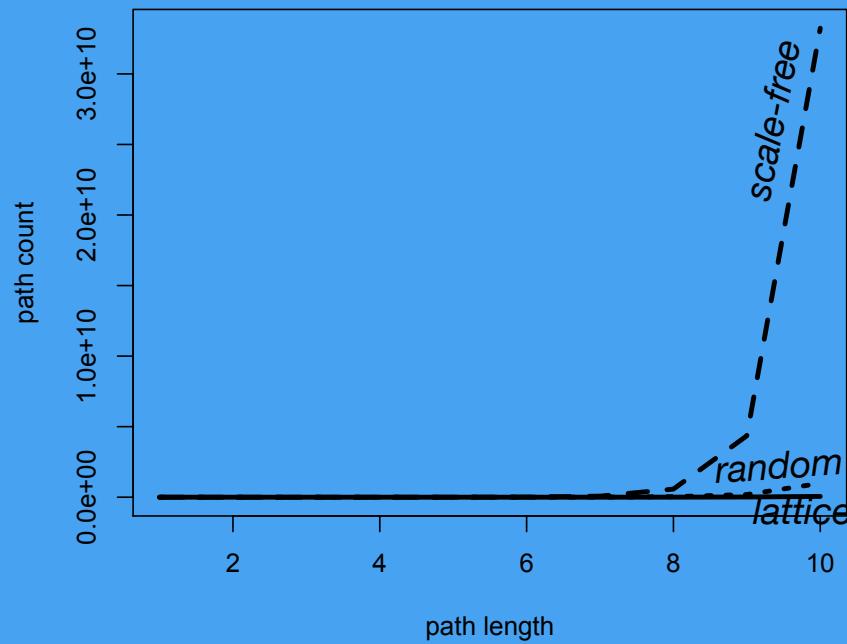
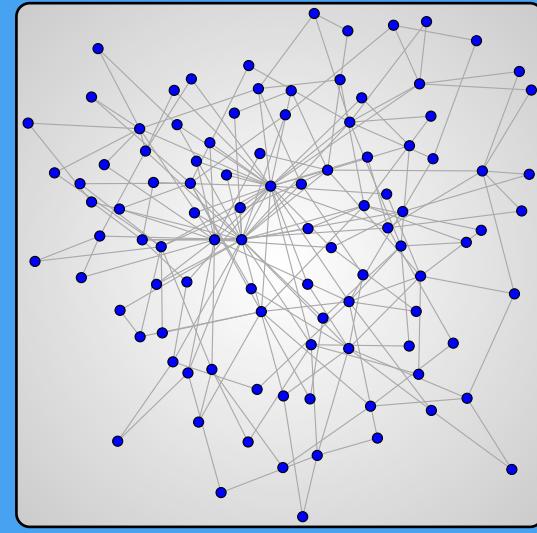
# Lattice



# Scale-Free

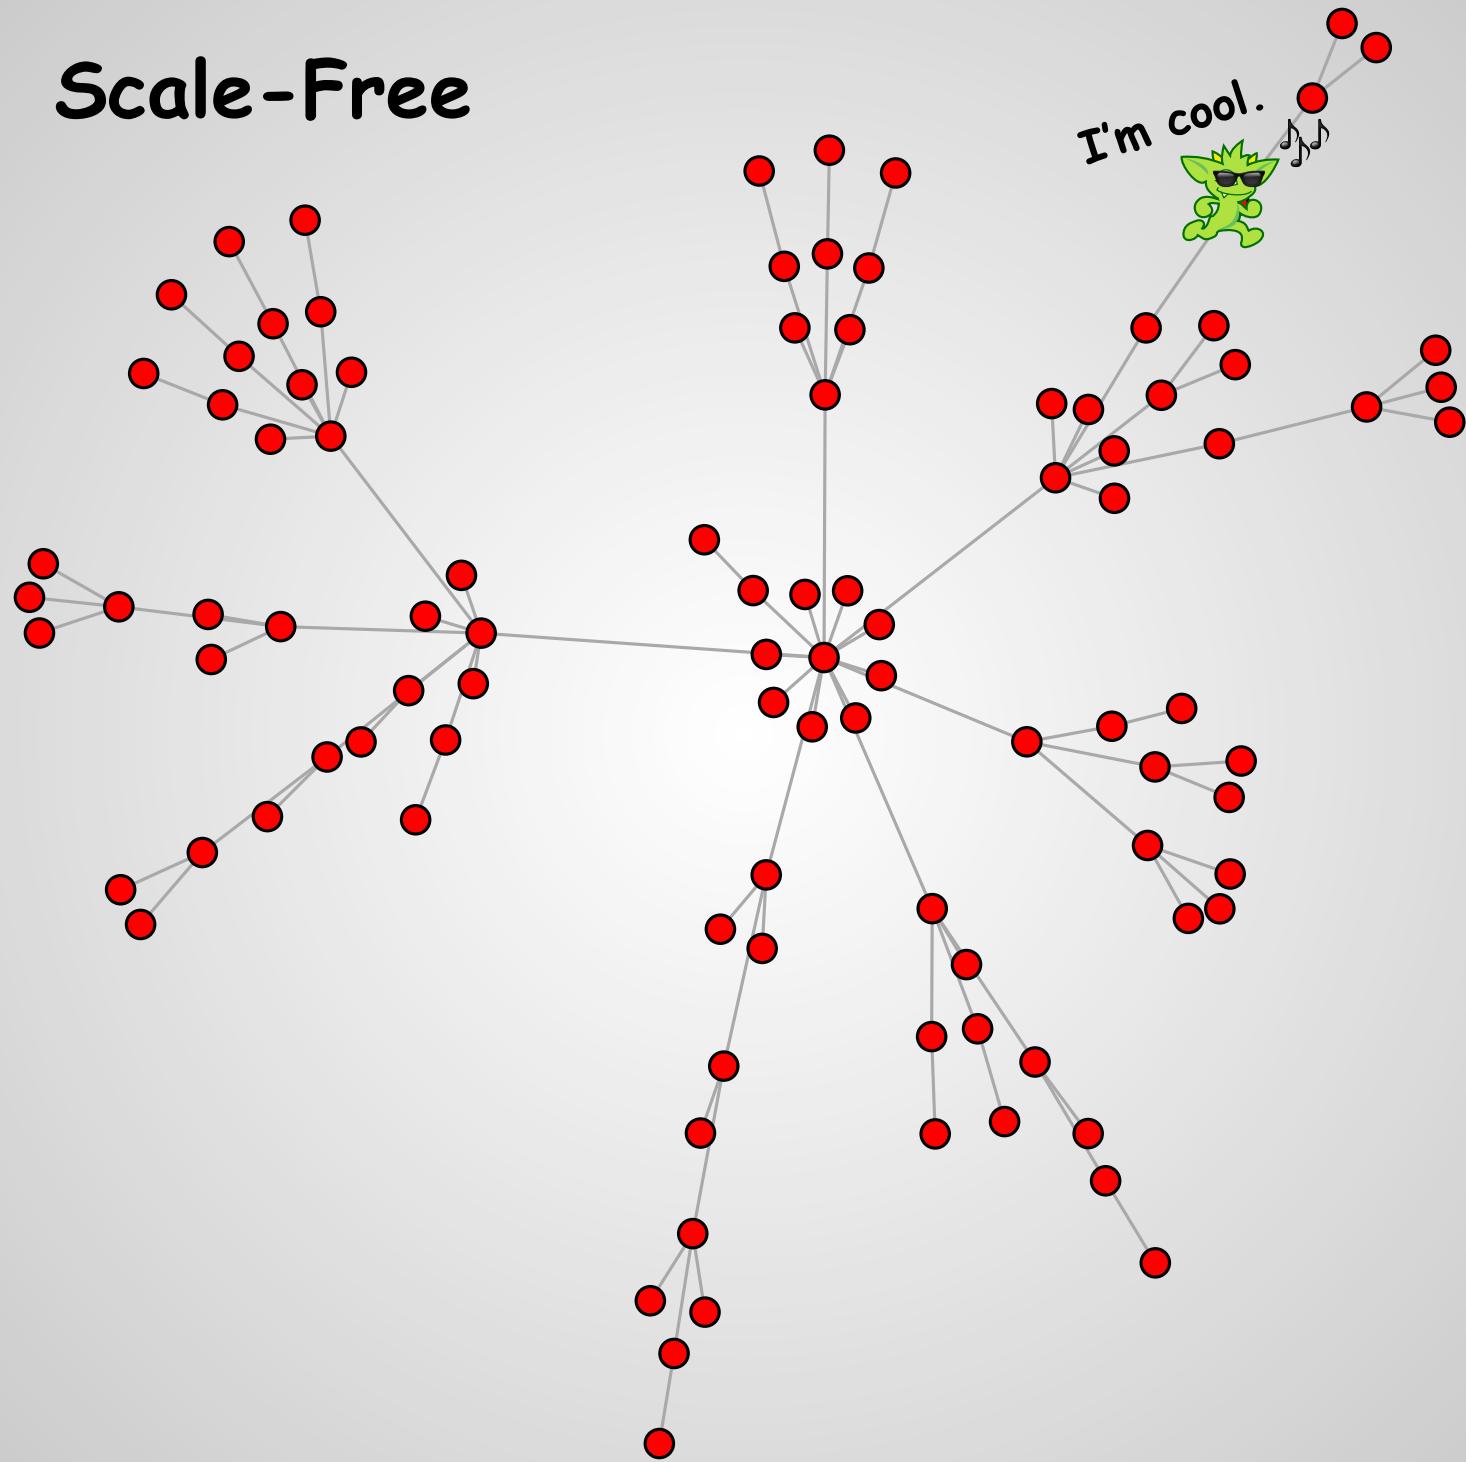


# Random

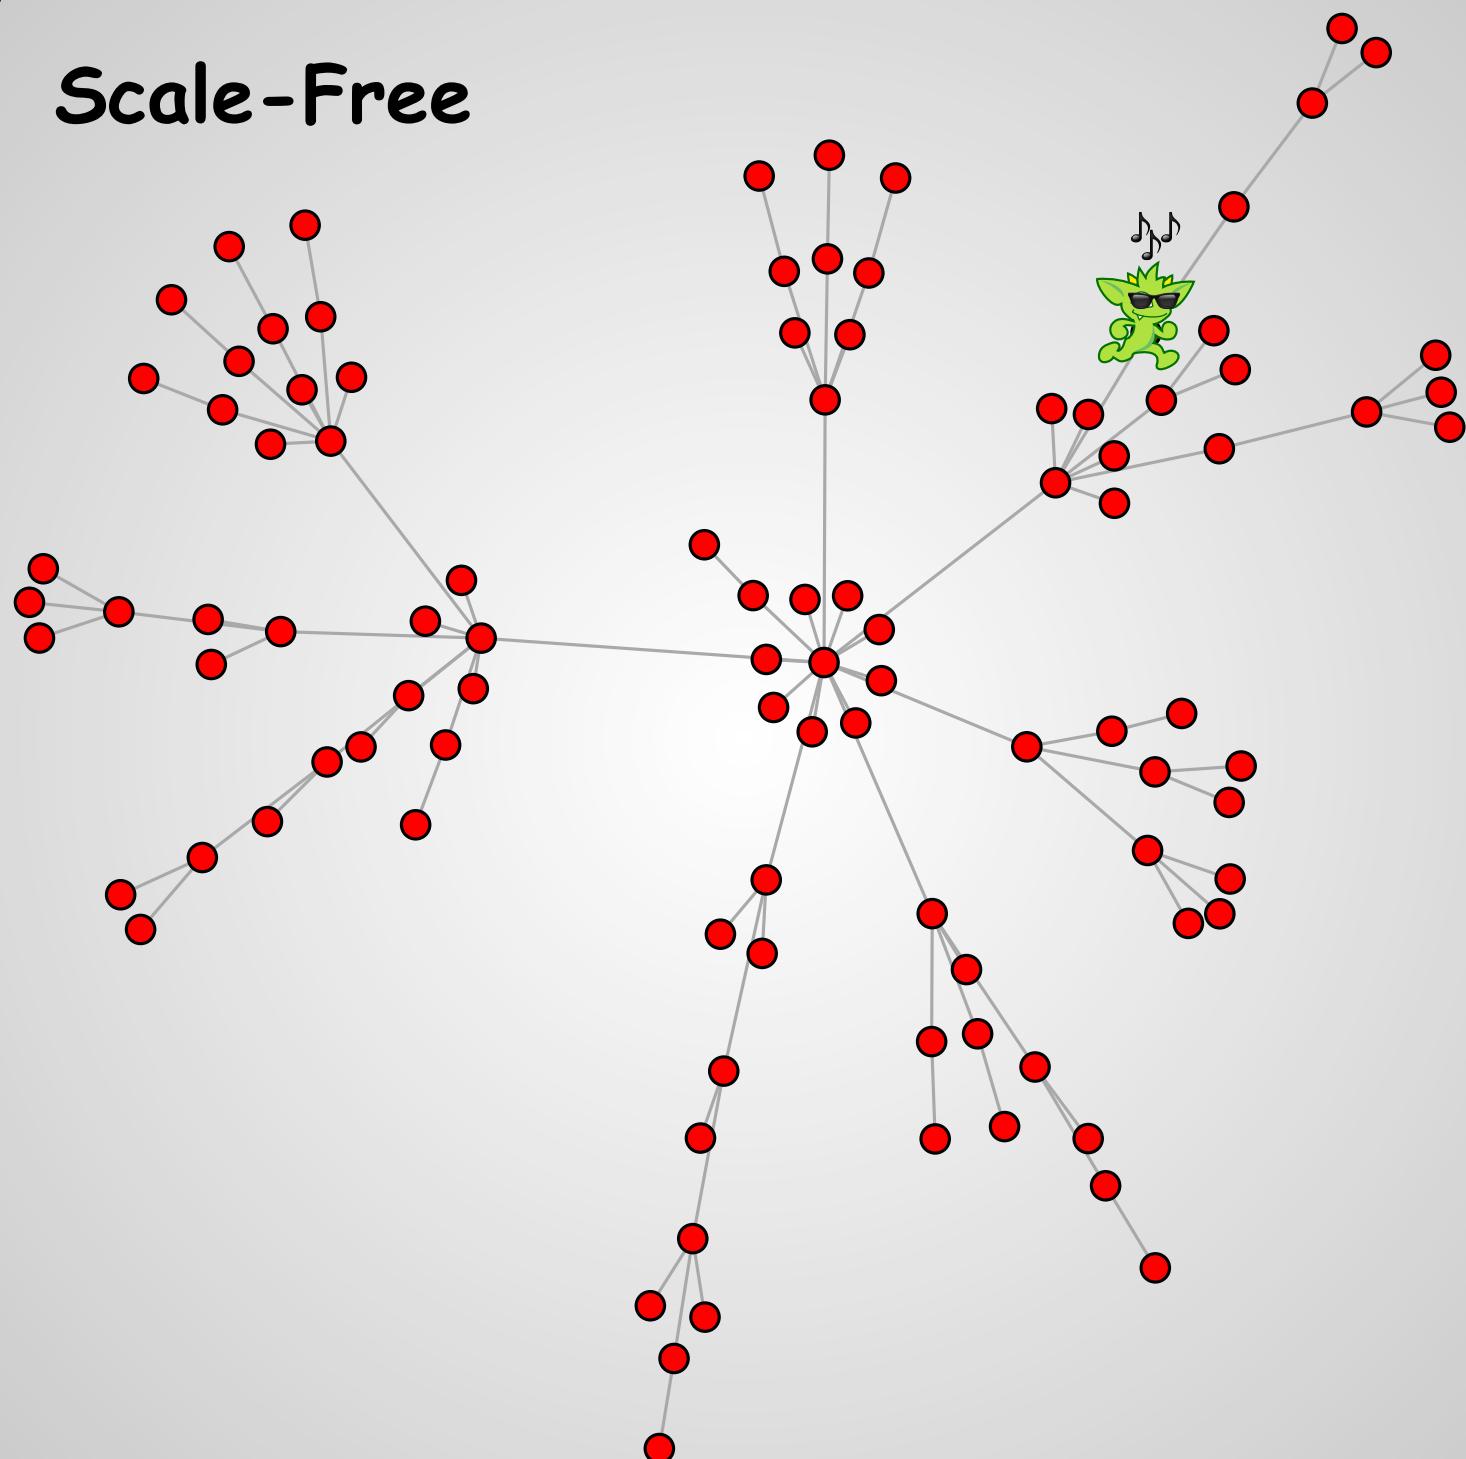


`g.V().repeat(both()).times(x).count()`

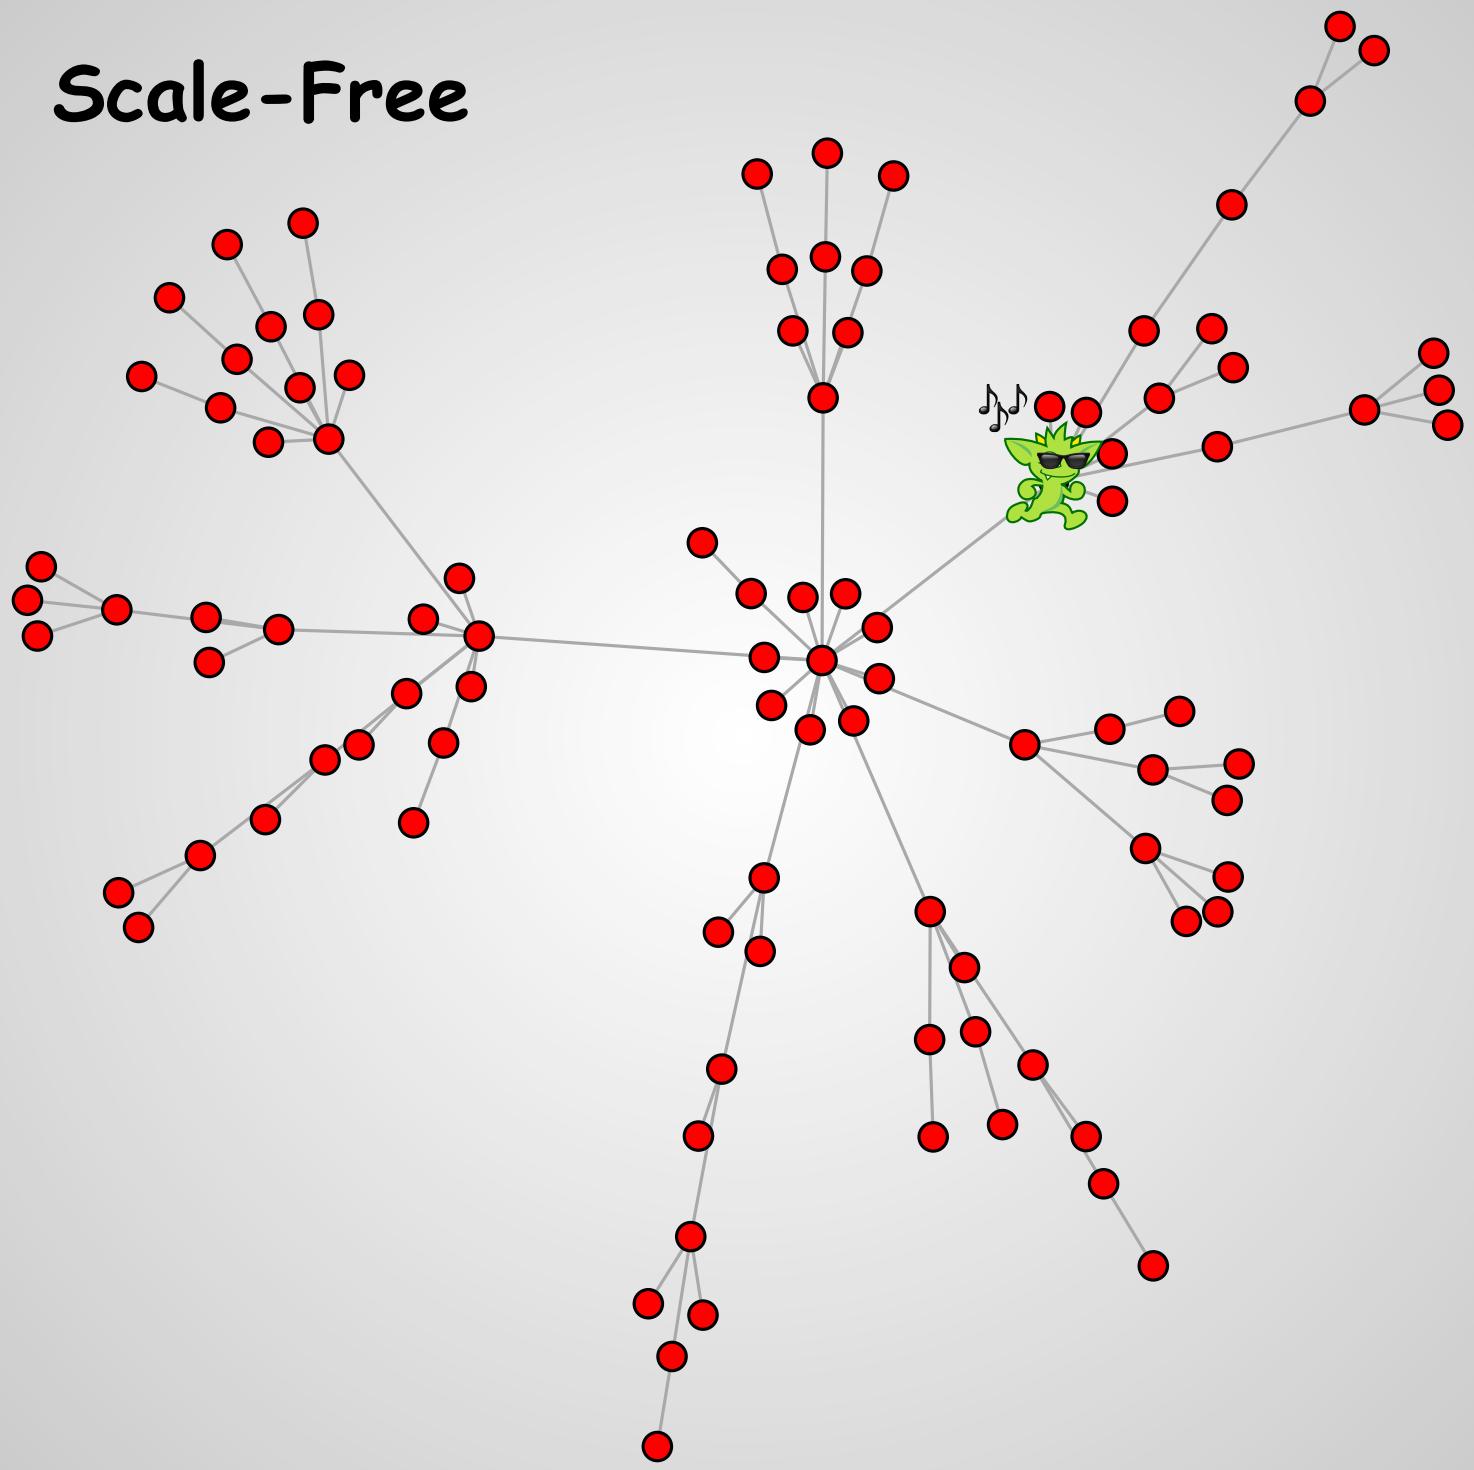
# Scale-Free



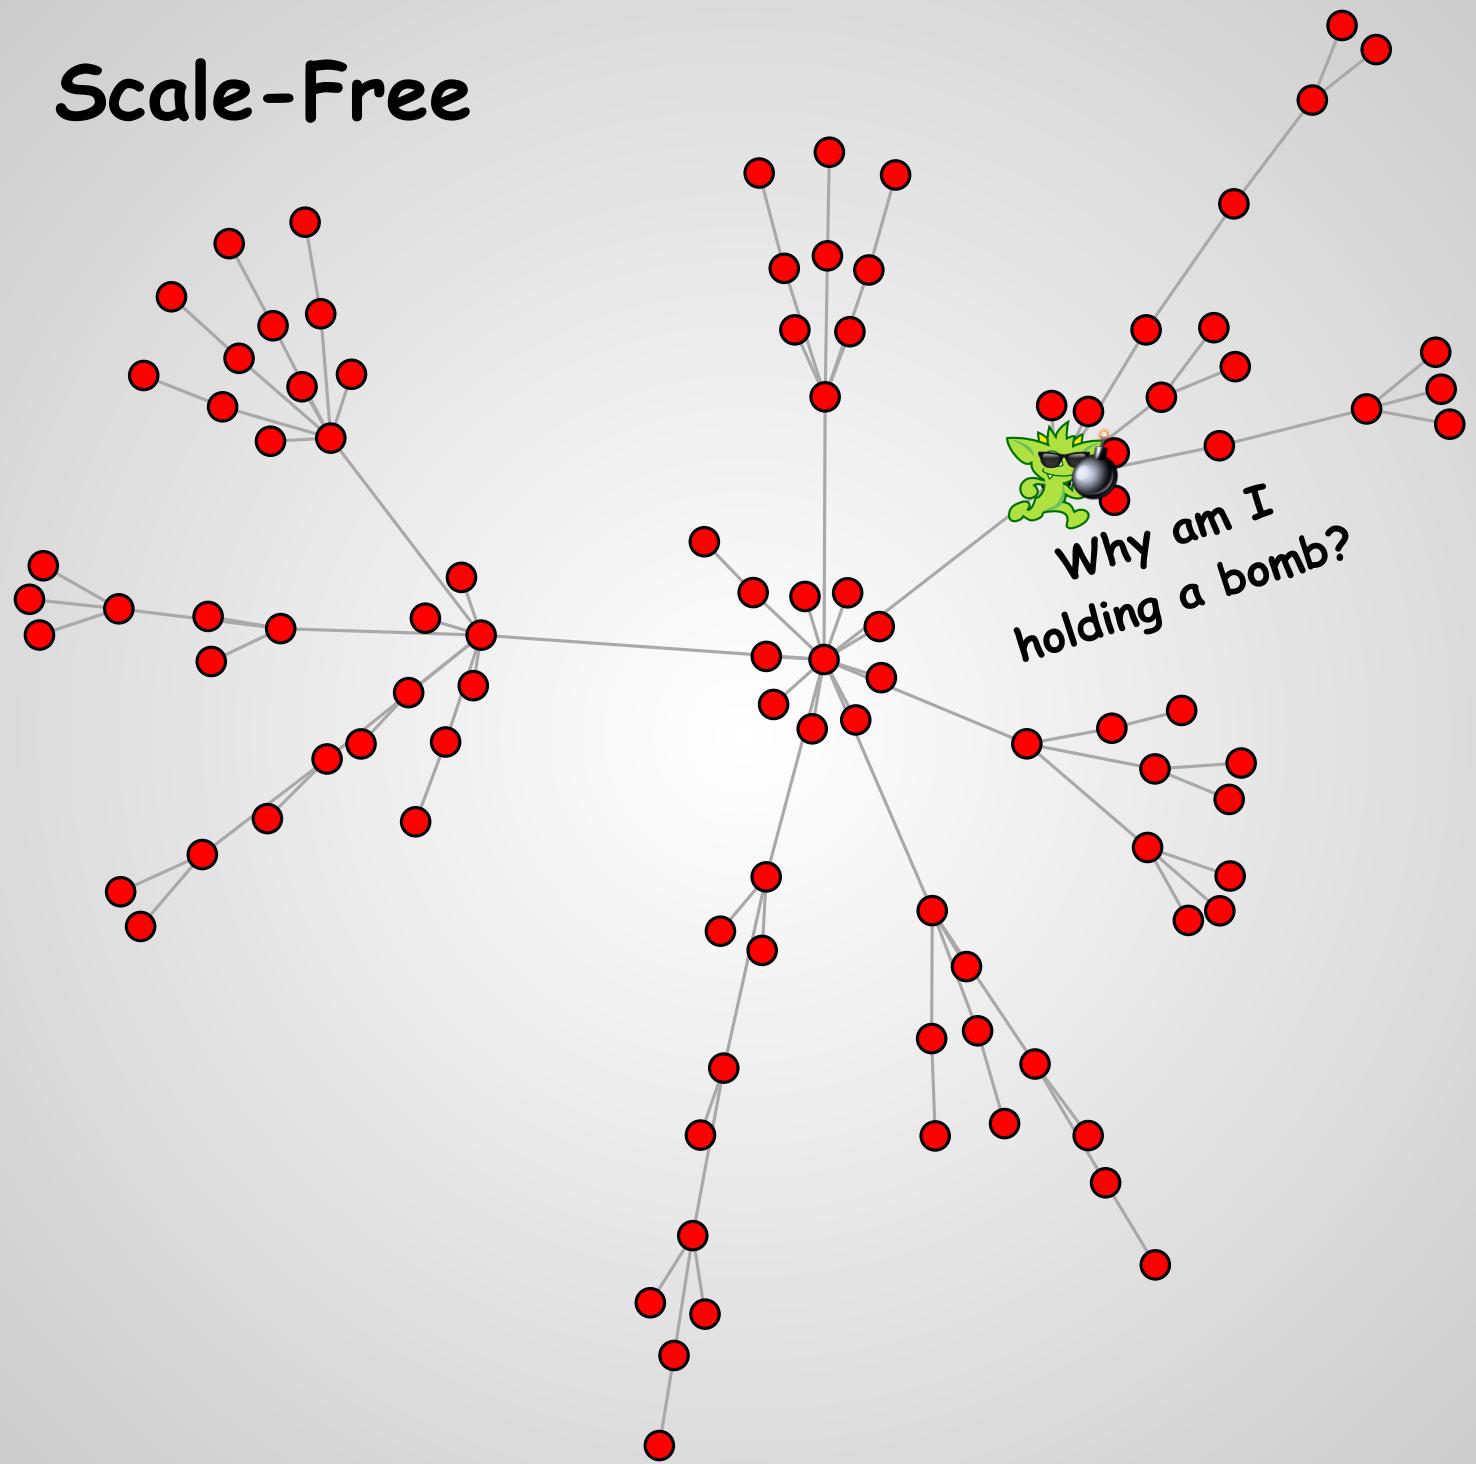
# Scale-Free



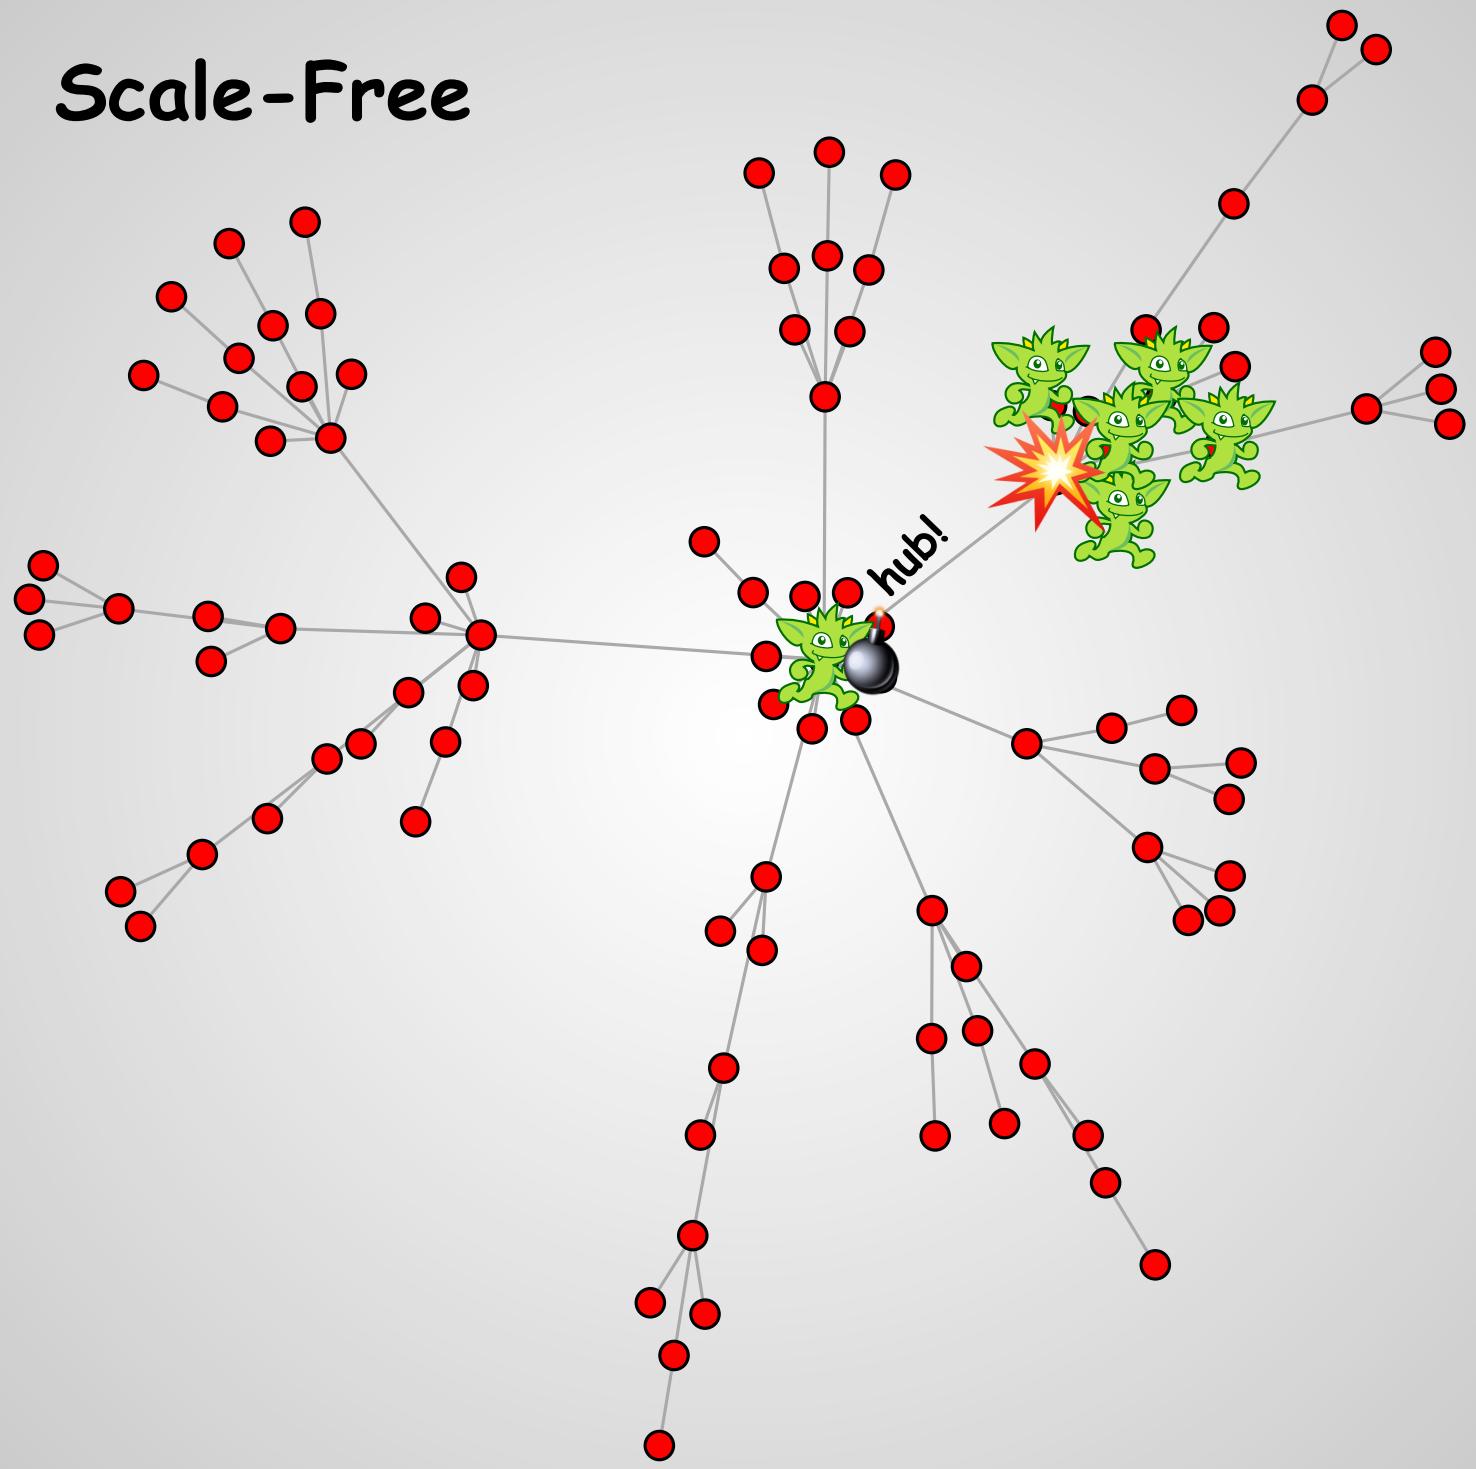
# Scale-Free



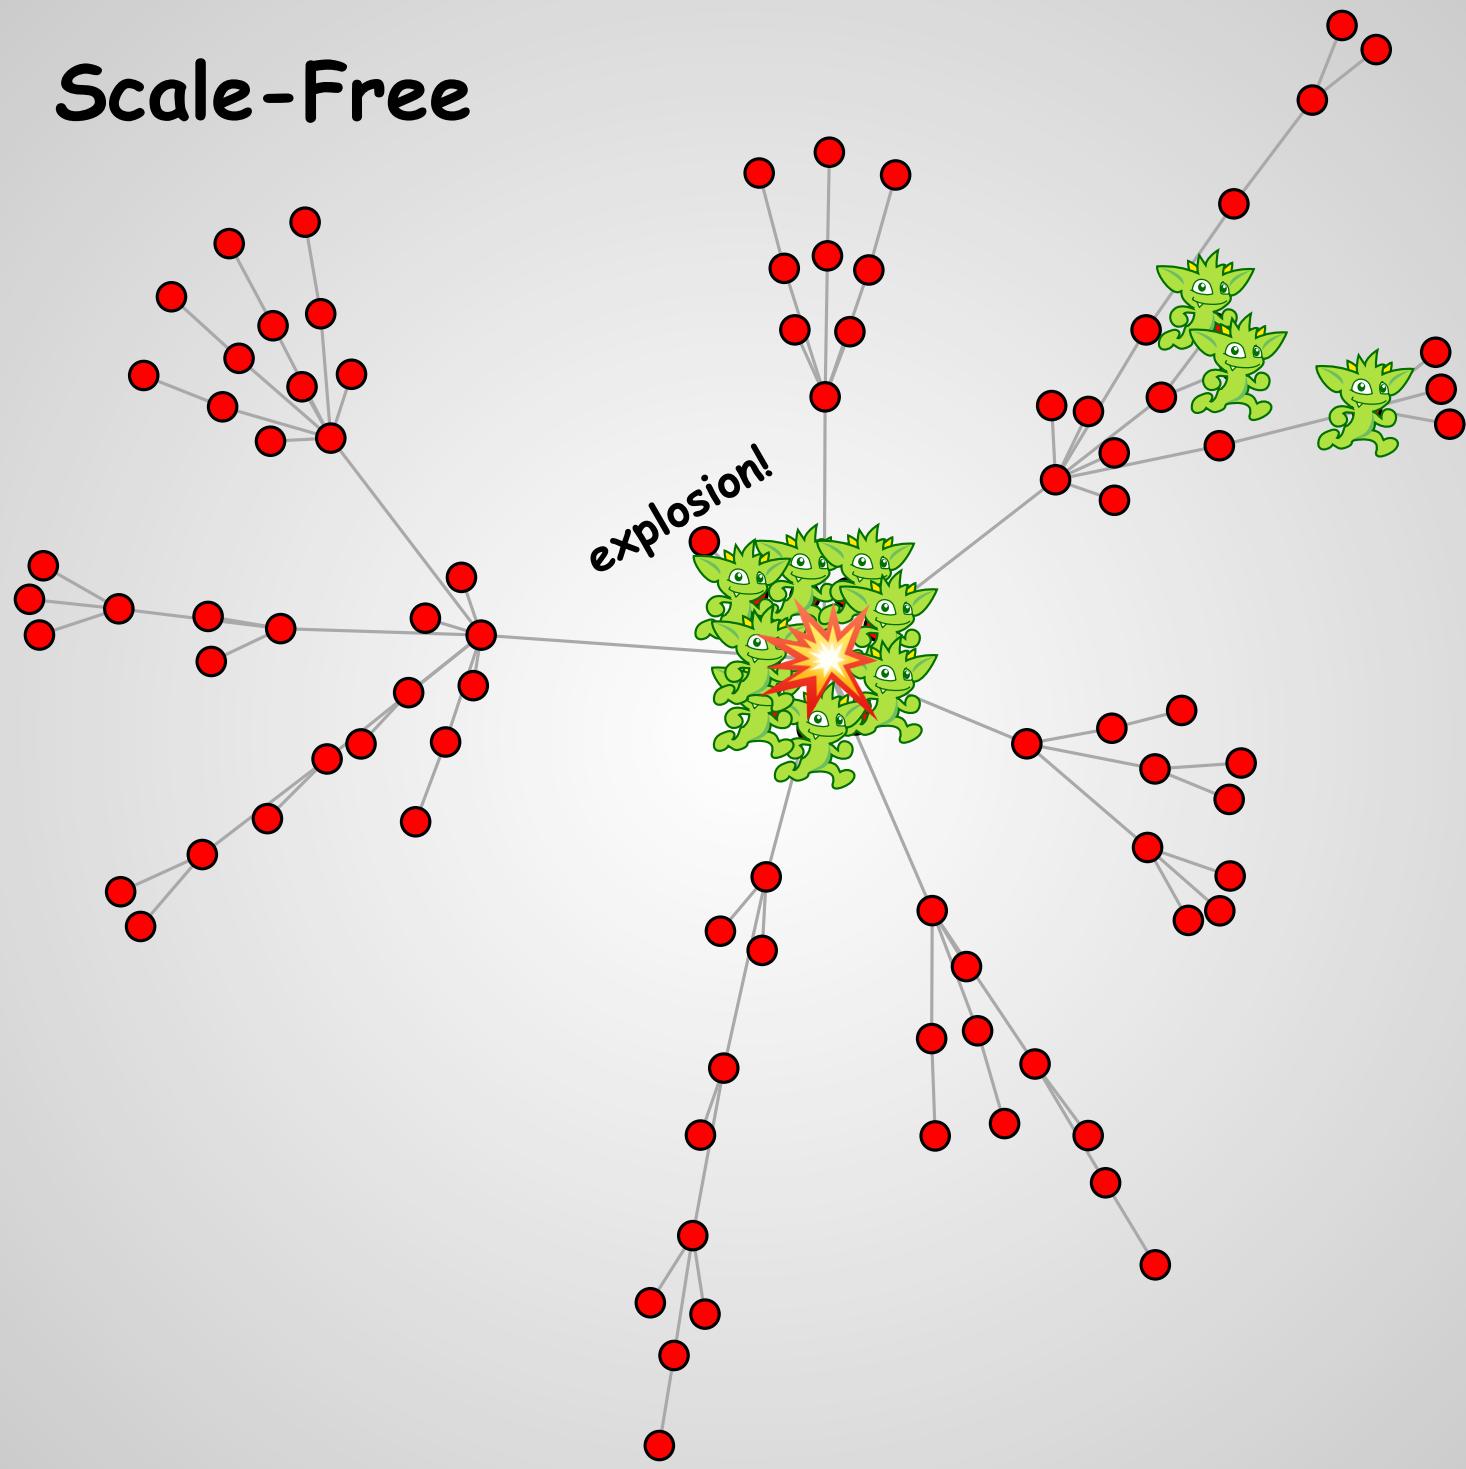
# Scale-Free



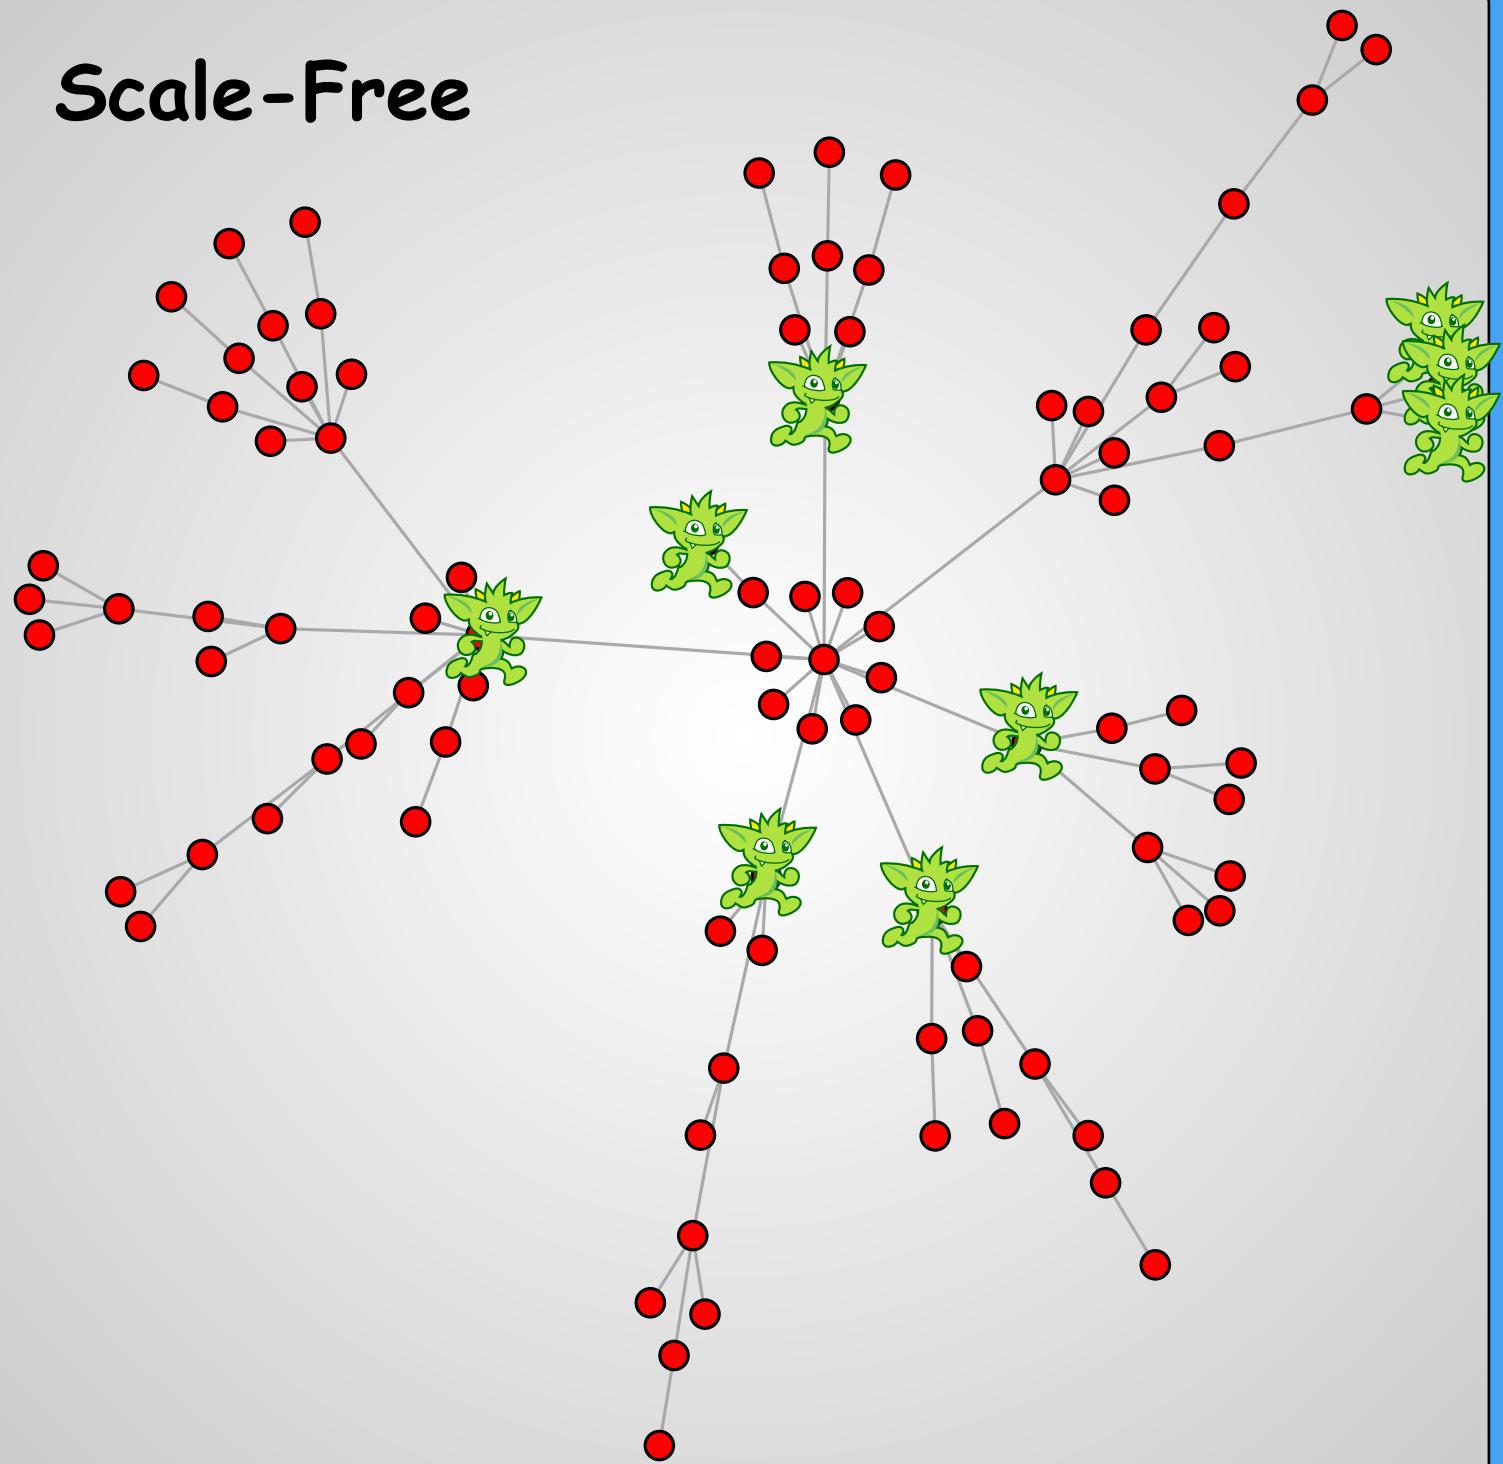
# Scale-Free



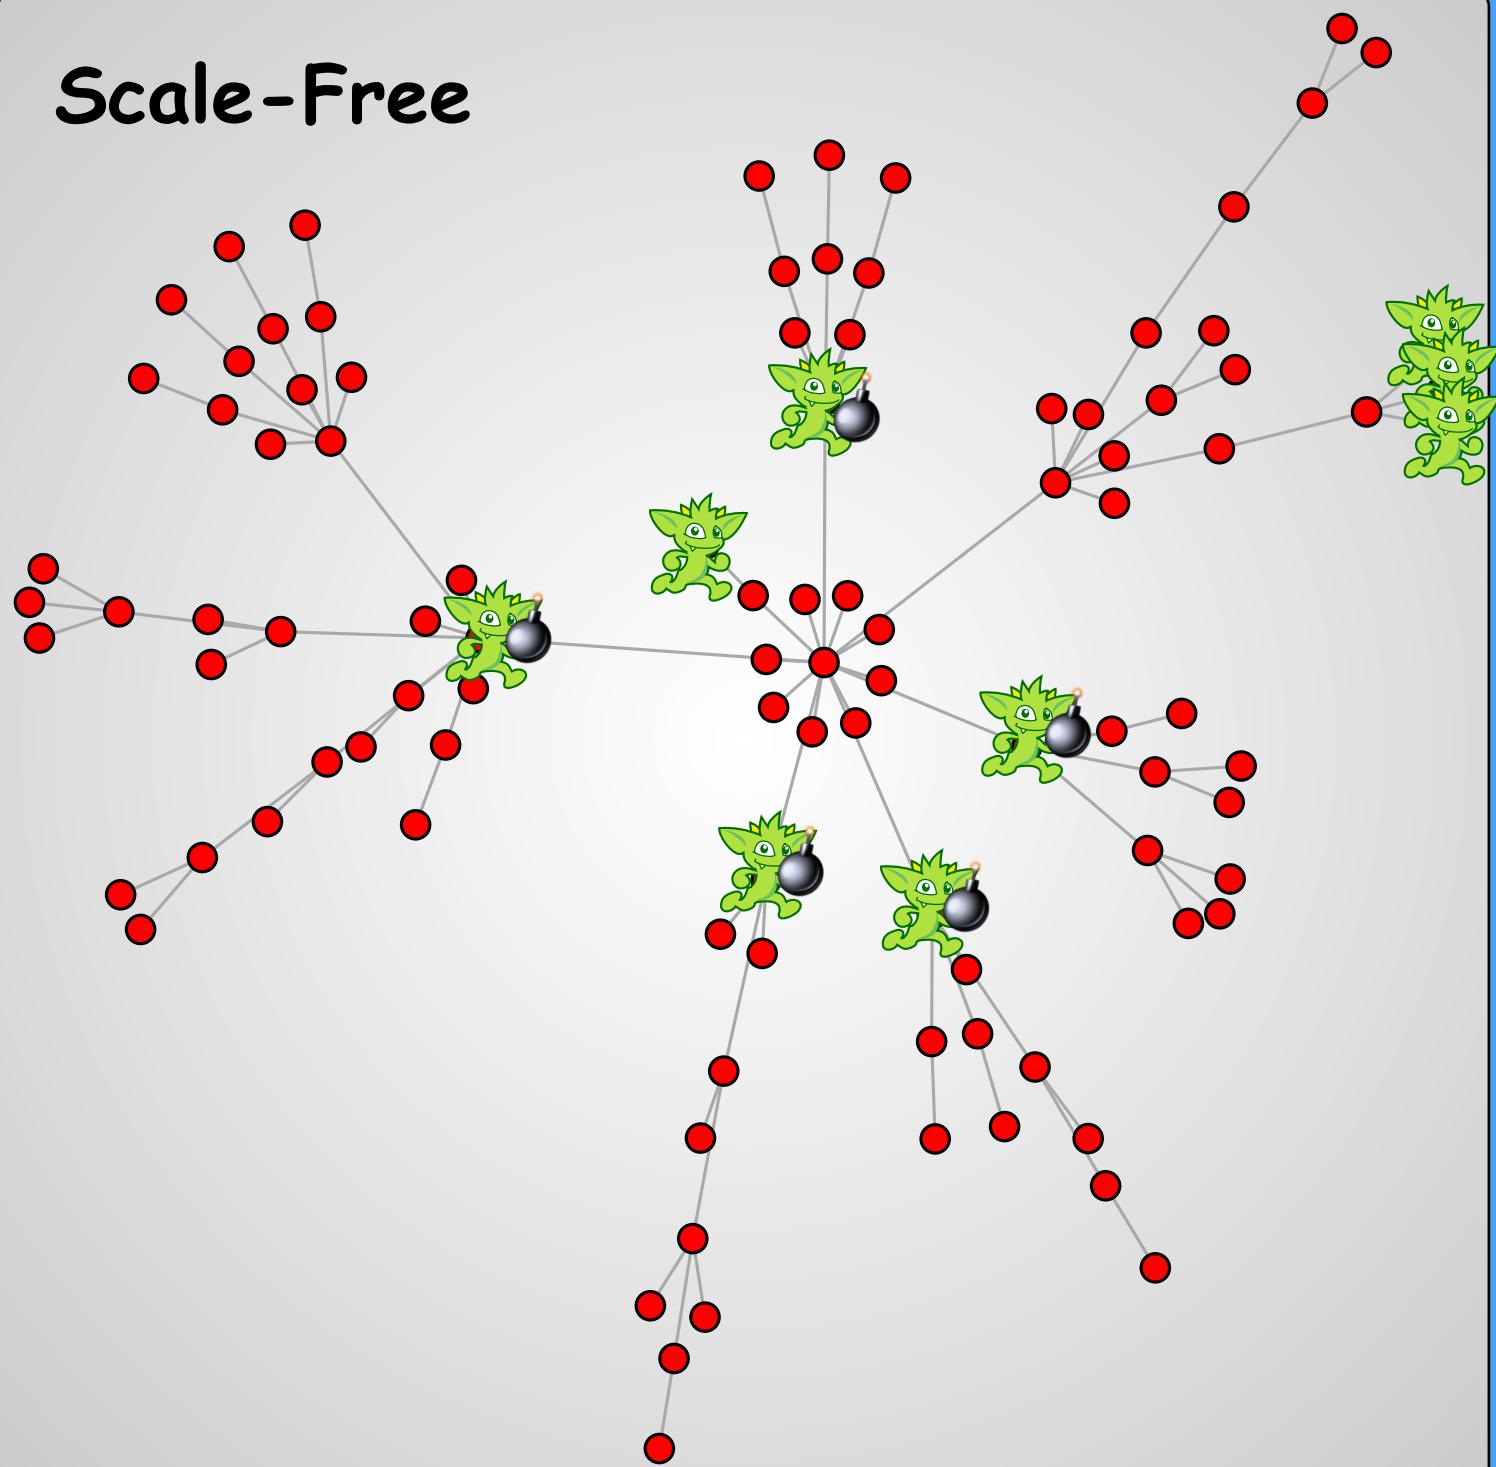
# Scale-Free



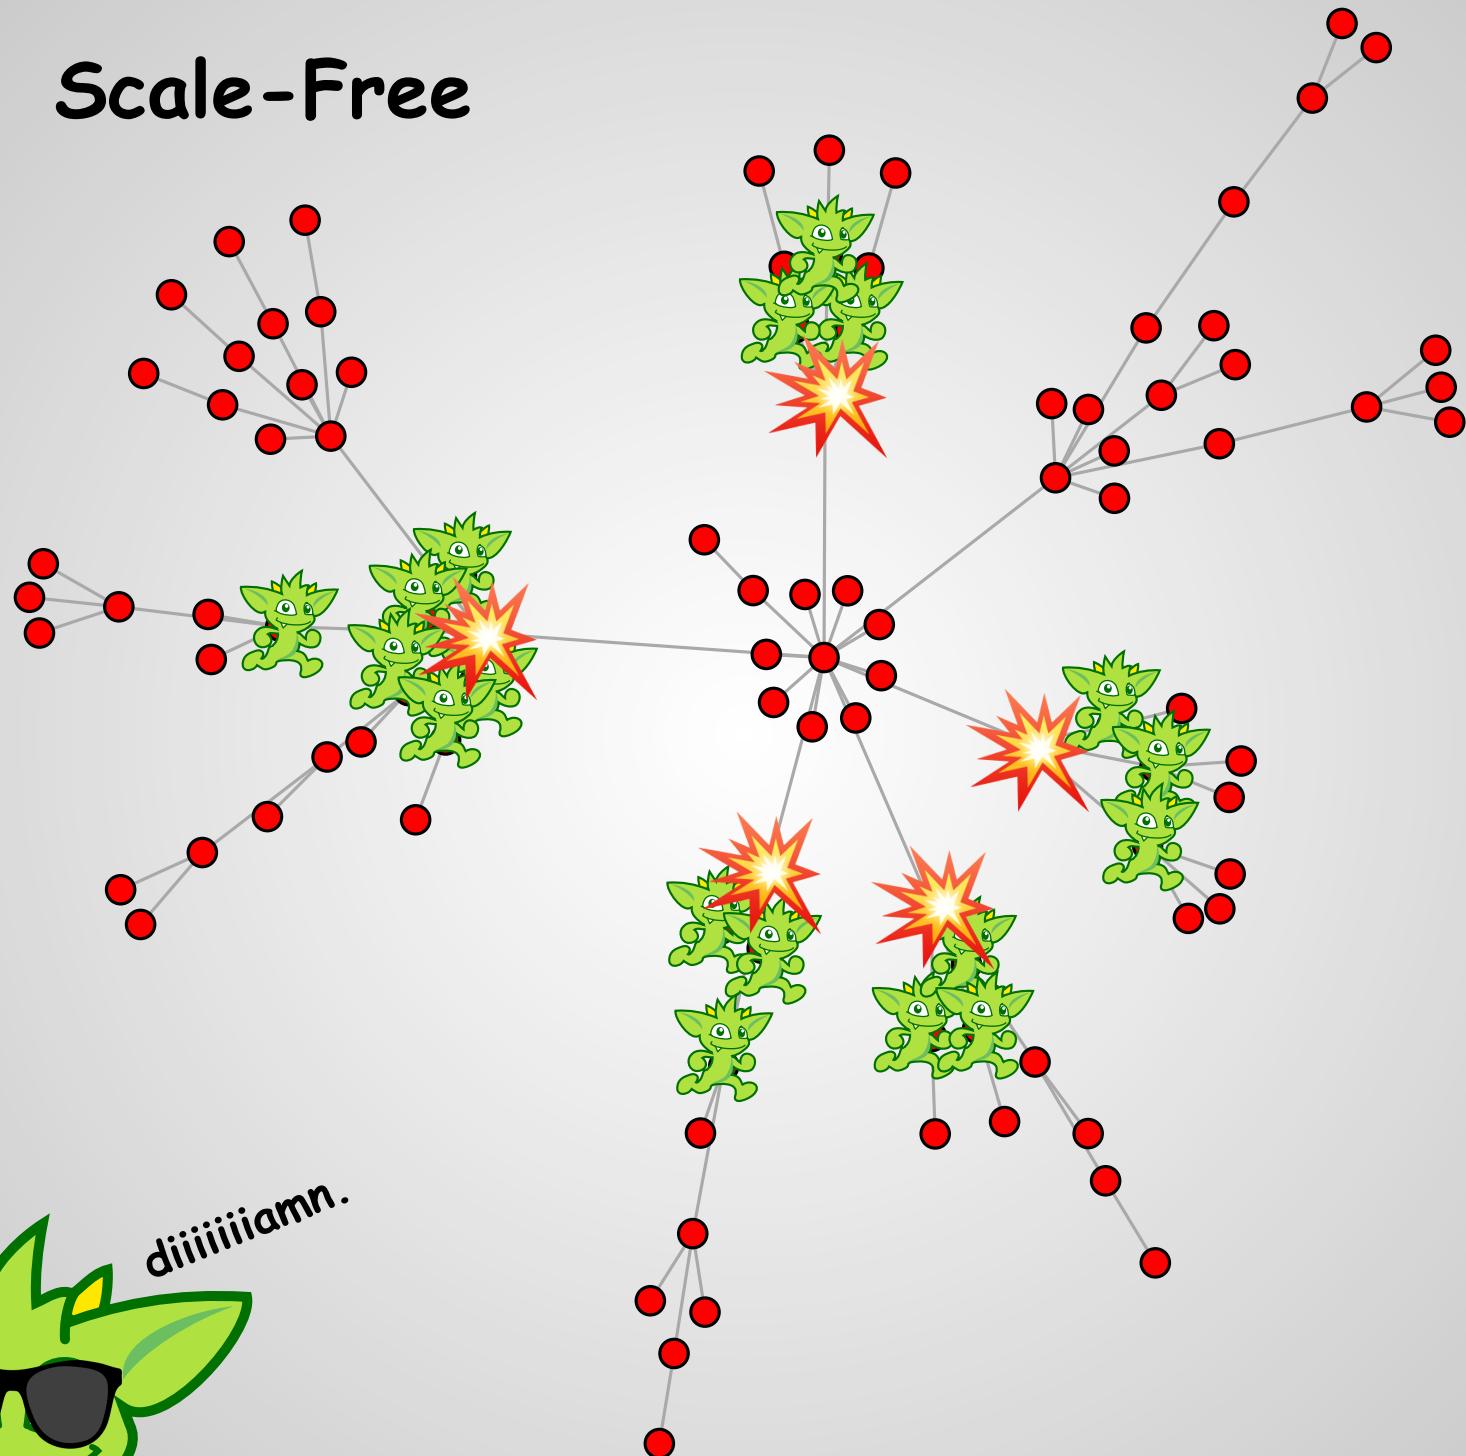
# Scale-Free



# Scale-Free



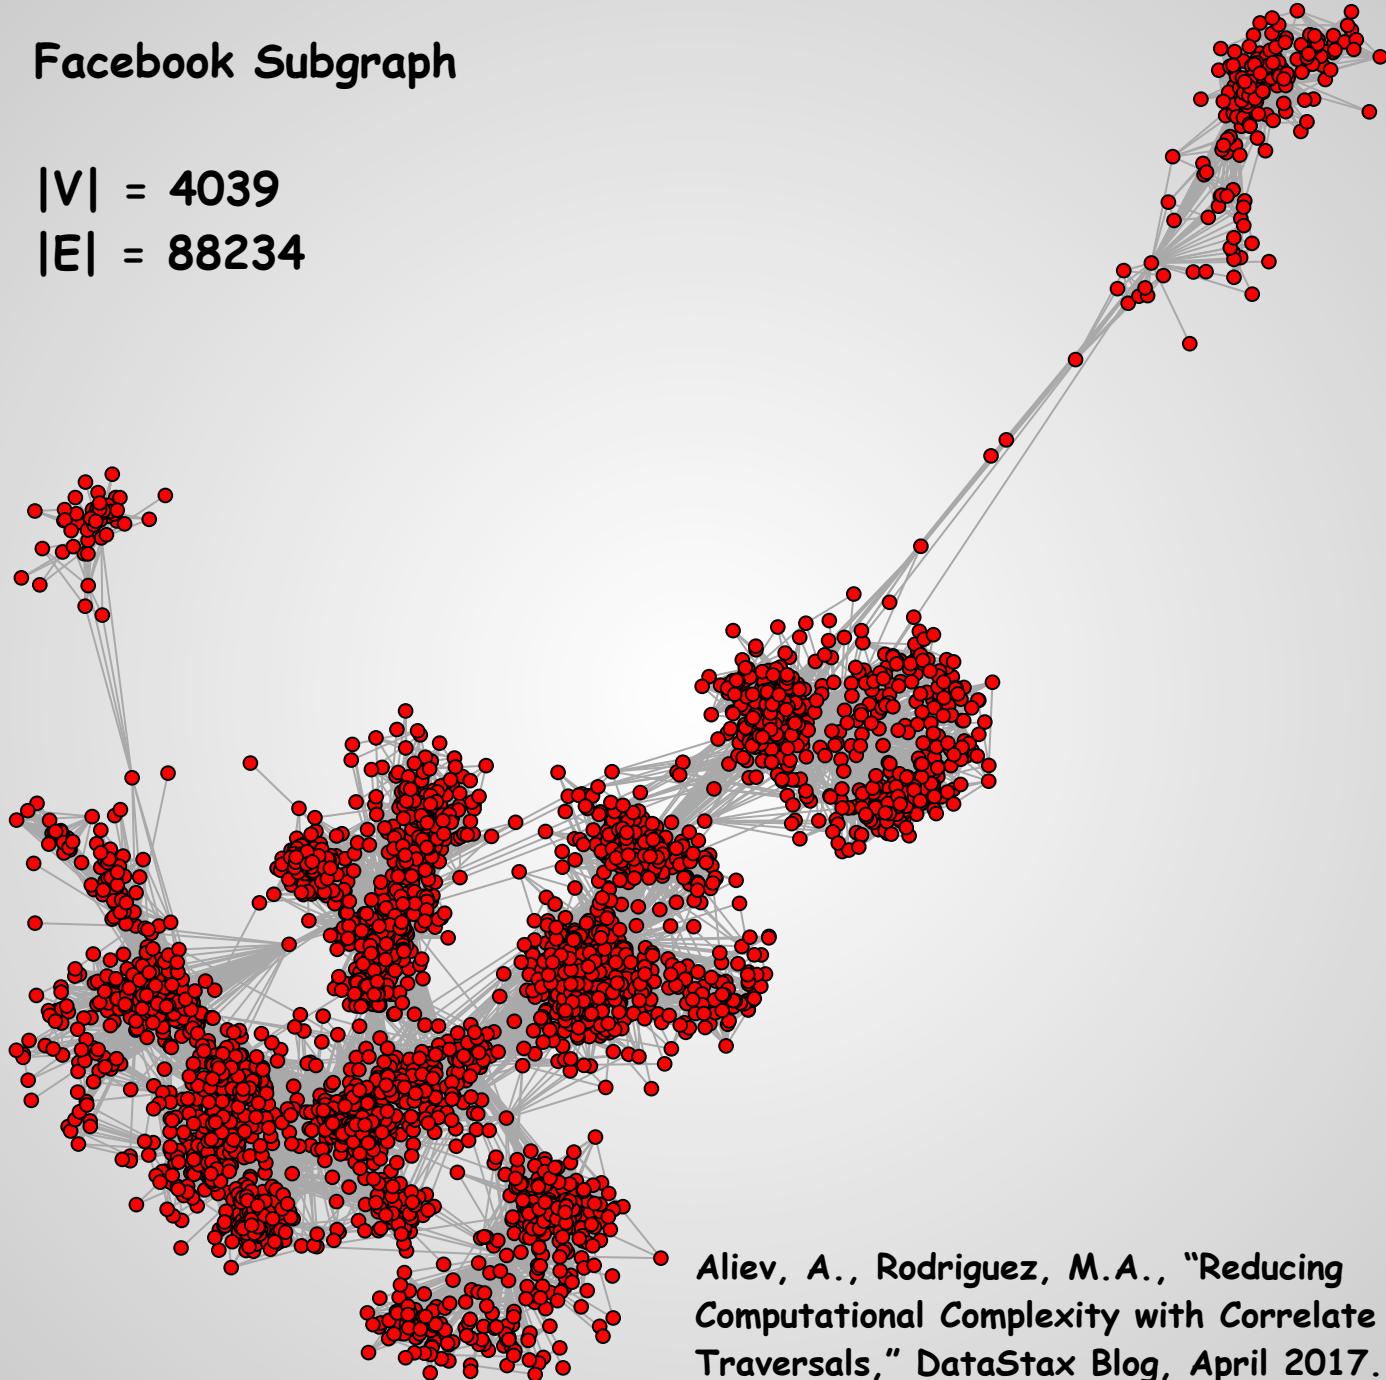
# Scale-Free



## Facebook Subgraph

$|V| = 4039$

$|E| = 88234$

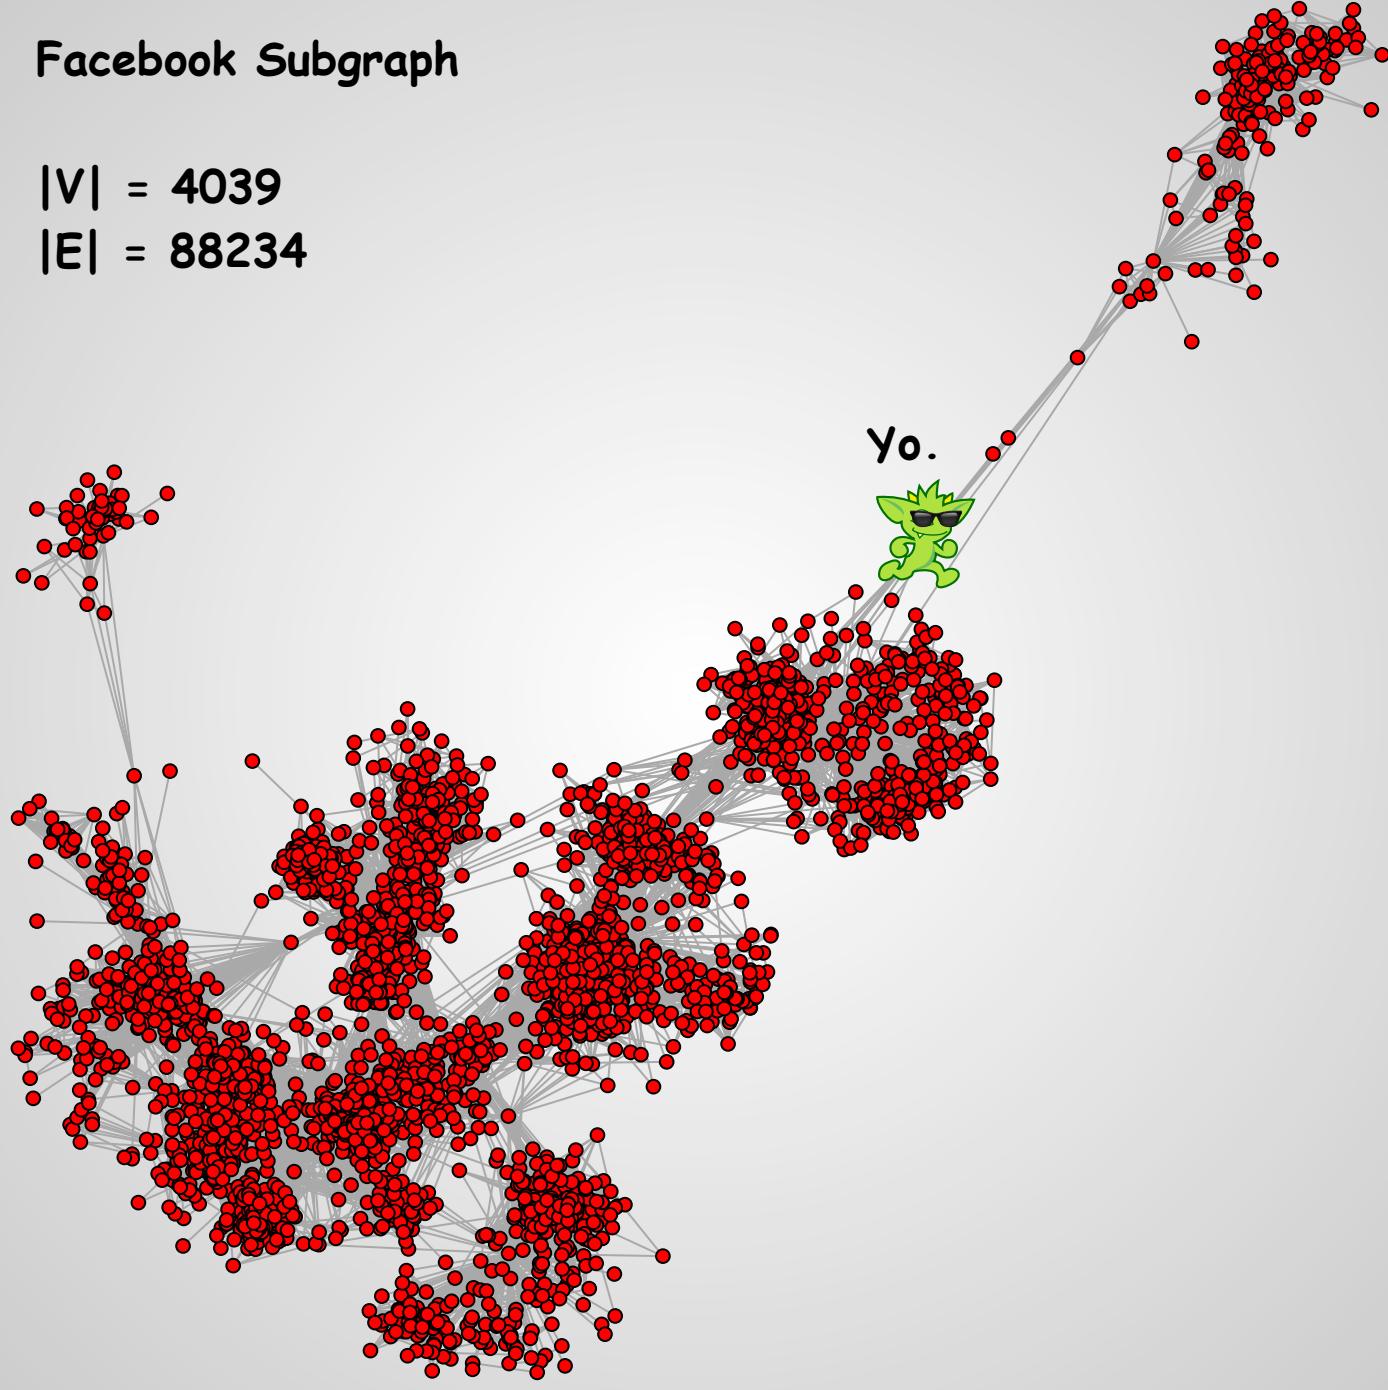


Aliev, A., Rodriguez, M.A., "Reducing Computational Complexity with Correlate Traversals," DataStax Blog, April 2017.

# Facebook Subgraph

$|V| = 4039$

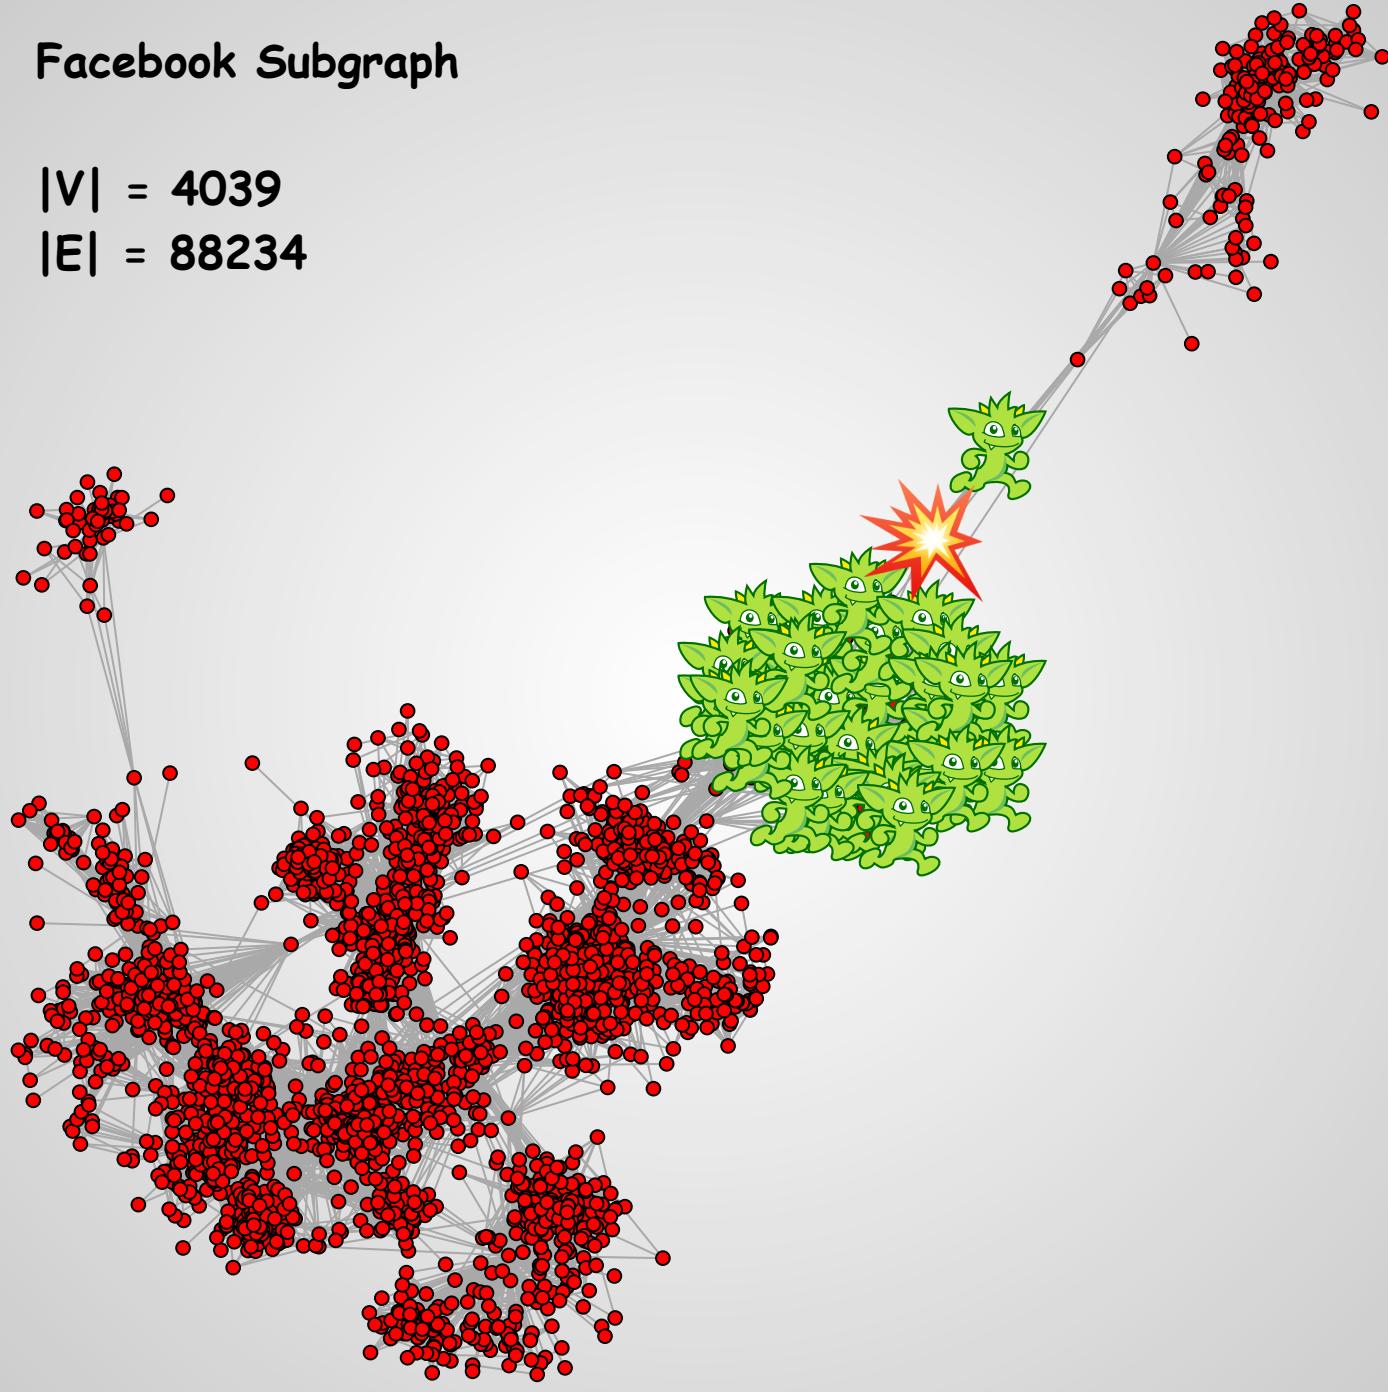
$|E| = 88234$



# Facebook Subgraph

$|V| = 4039$

$|E| = 88234$

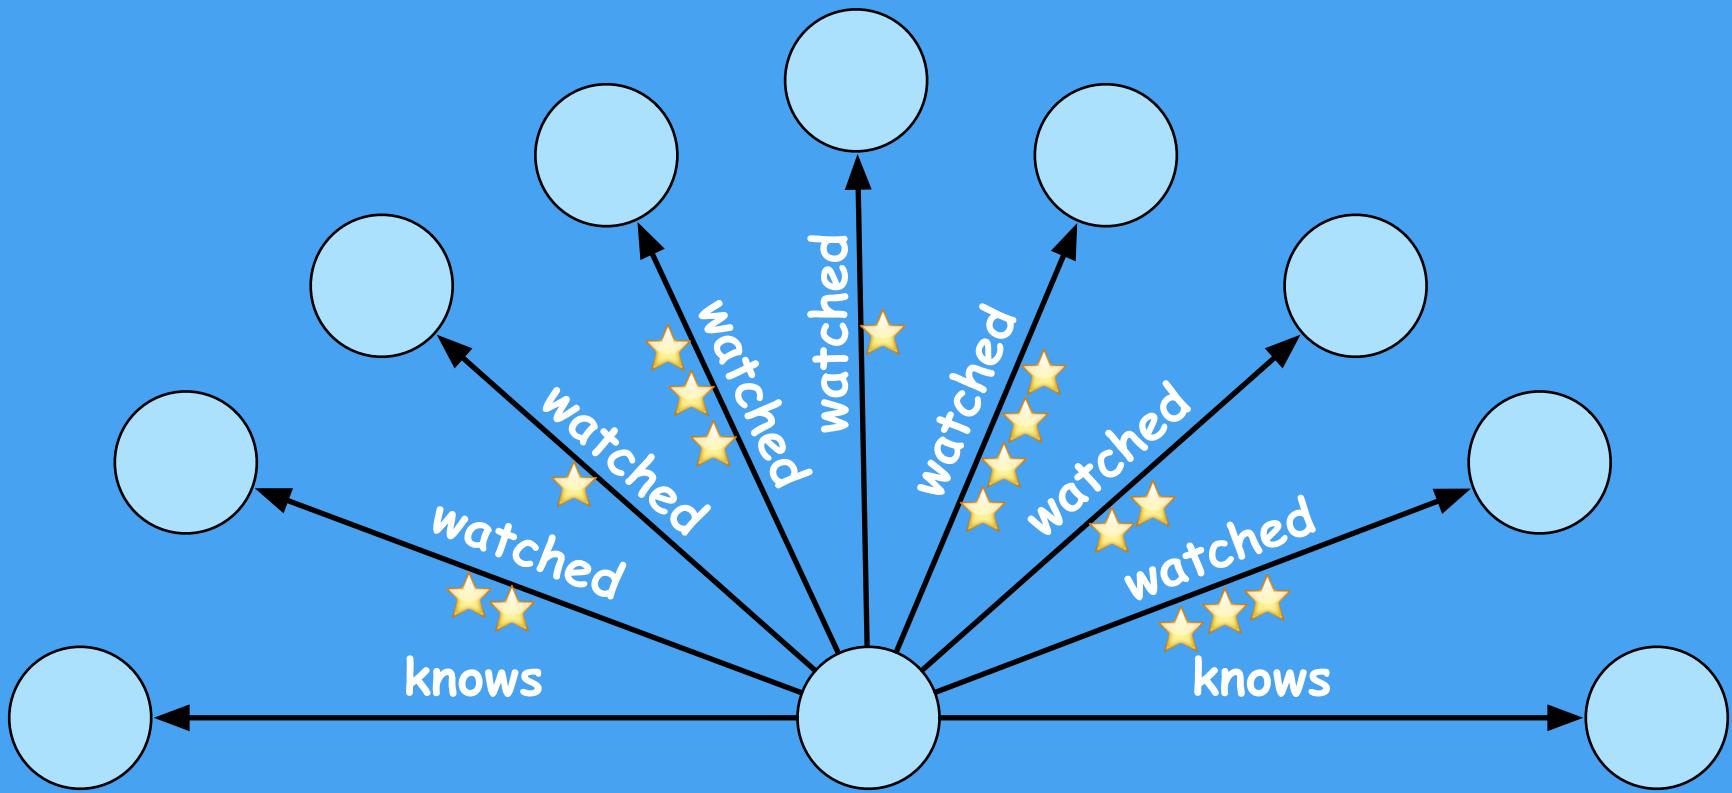


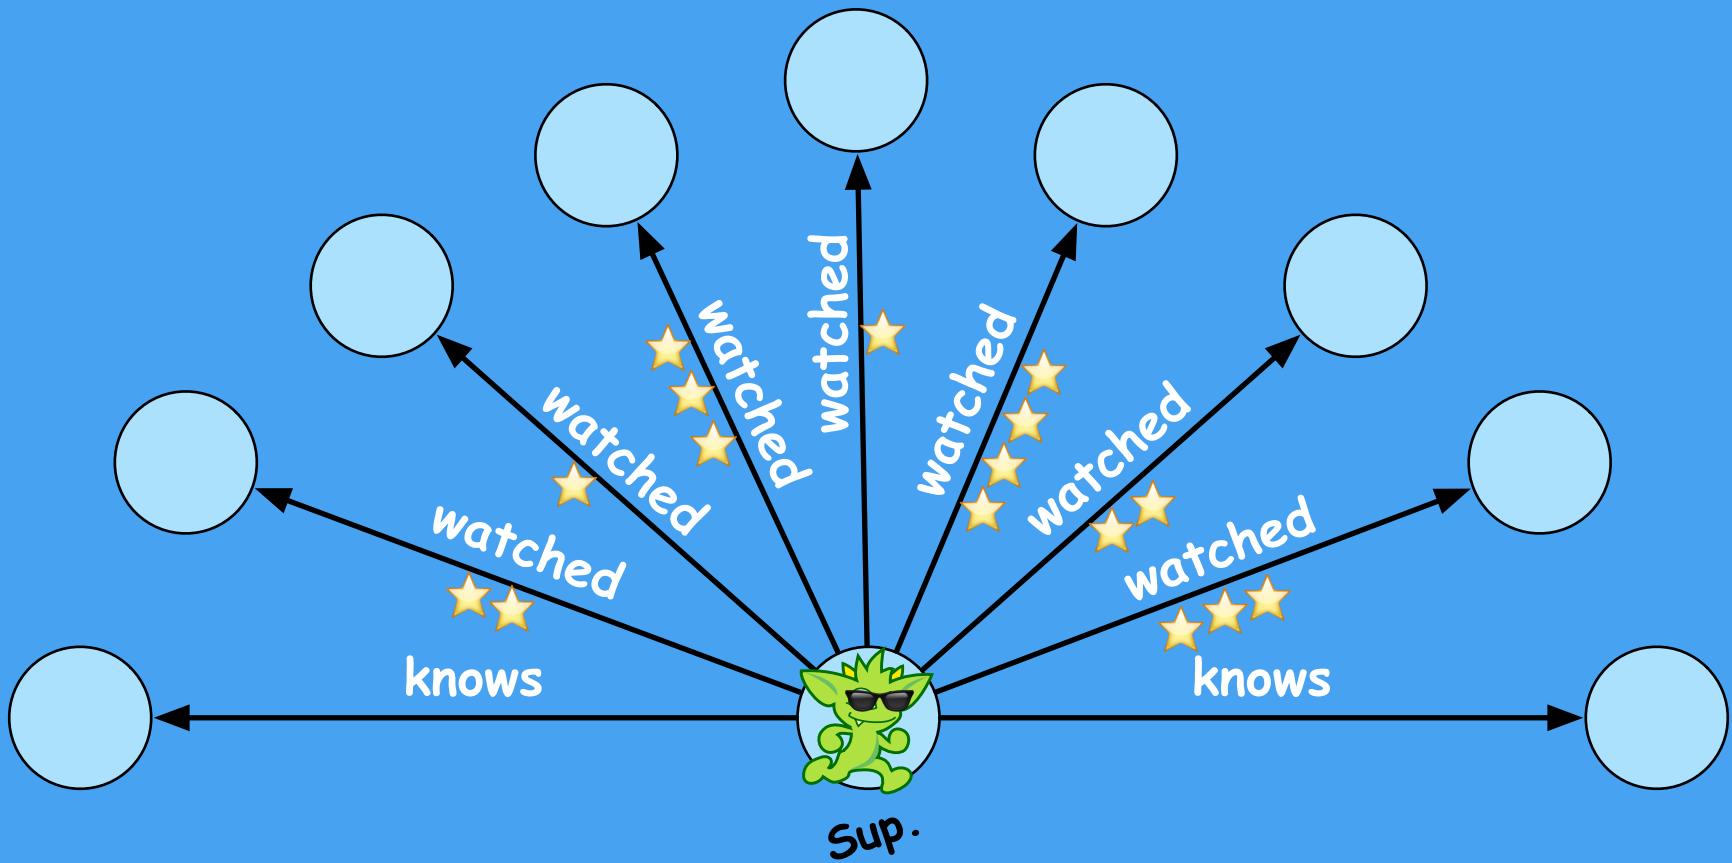
# So you want to be a graph surgeon?

- \* The same traversal on different graphs can have wildly different performance characteristics.

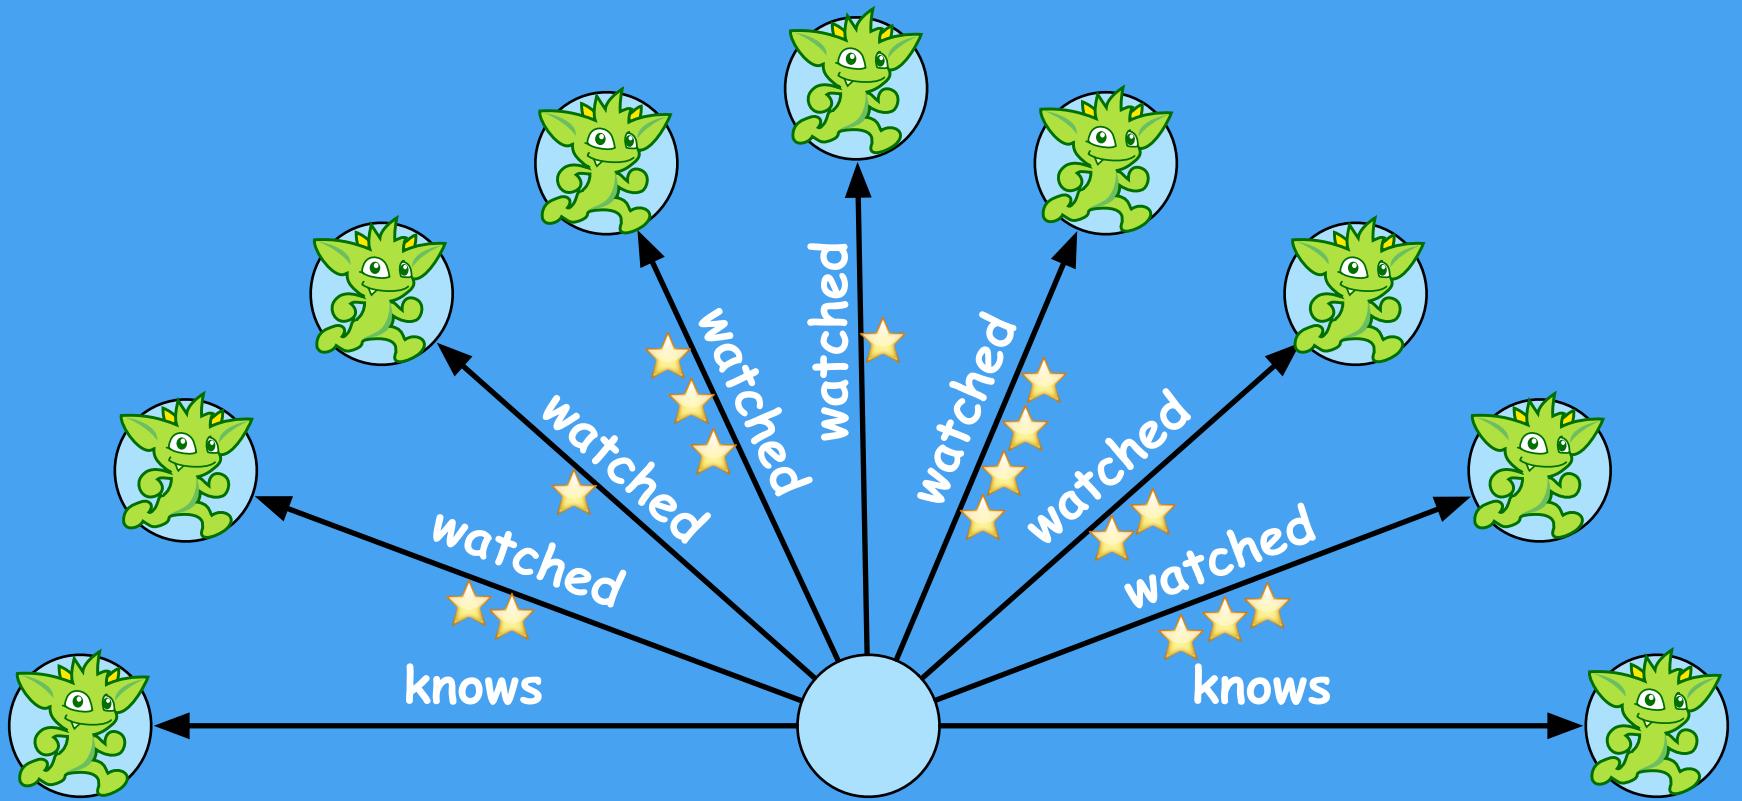


- \* Traversals on scale-free/natural graphs can easily explode due to “hubs” (super nodes).
- \* Use Gremlin to study the graph's statistics before using Gremlin to query the graph.
- \* Build your traversals up piecewise and use profile() to find bottlenecks.

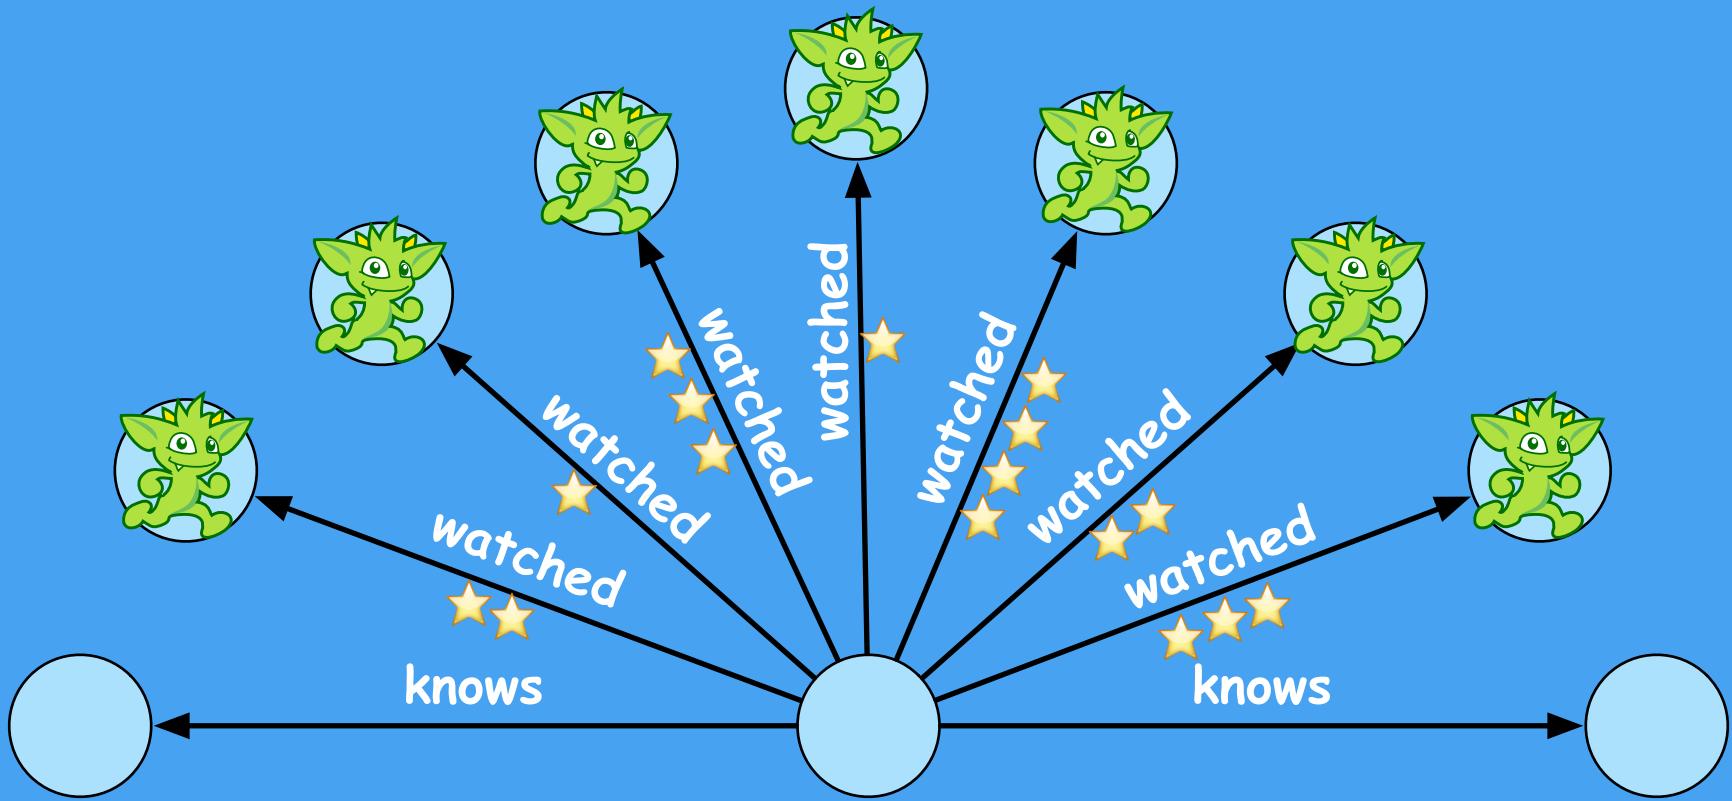




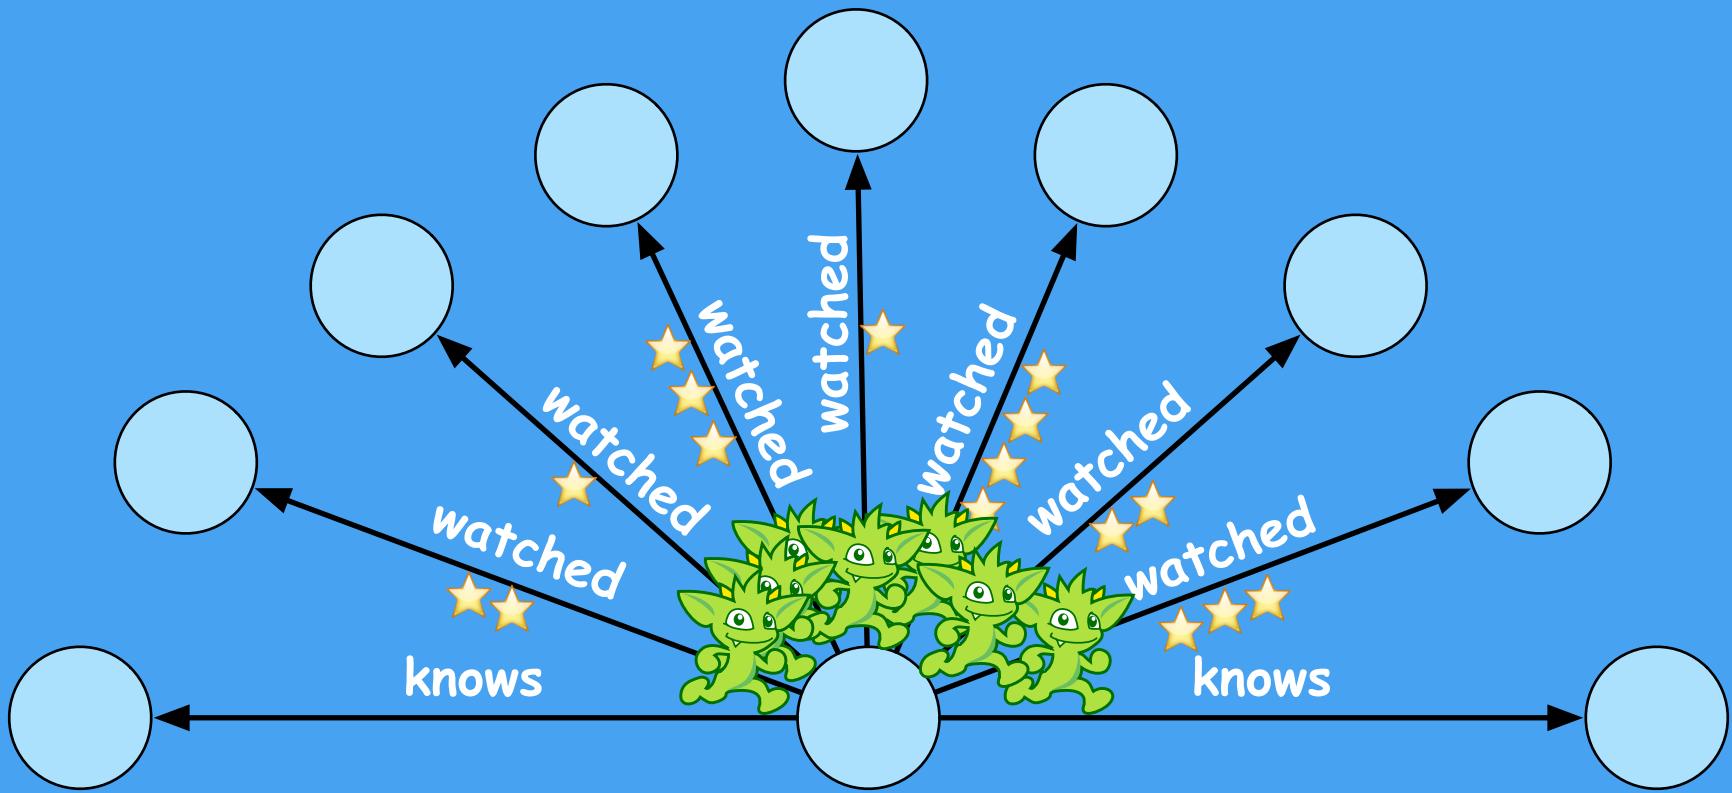
`g.V(1).count()`



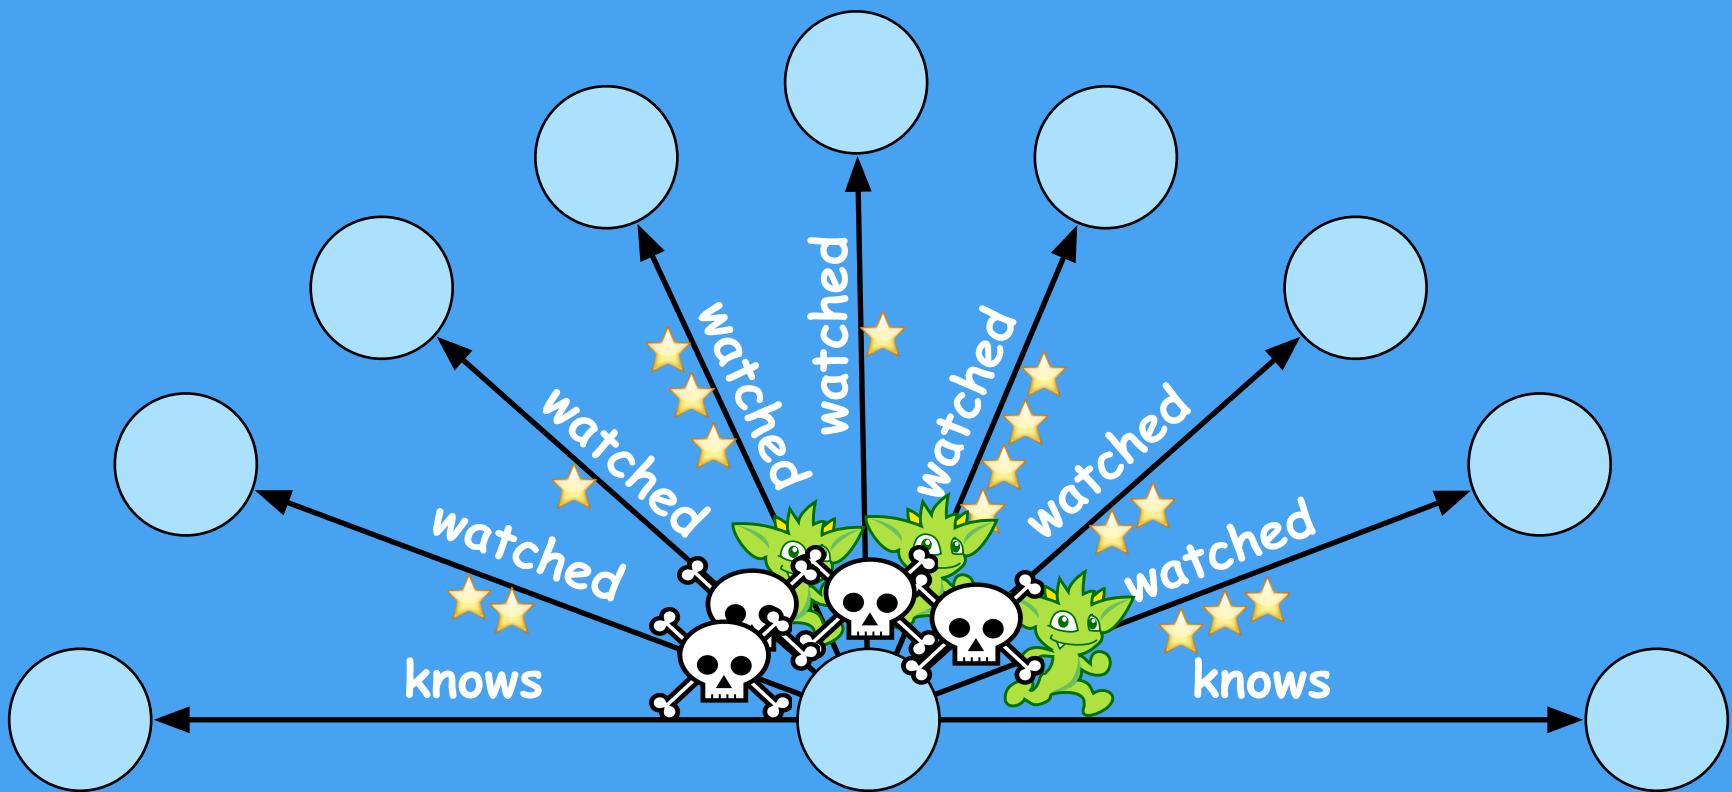
$g.V(1).out().count()$



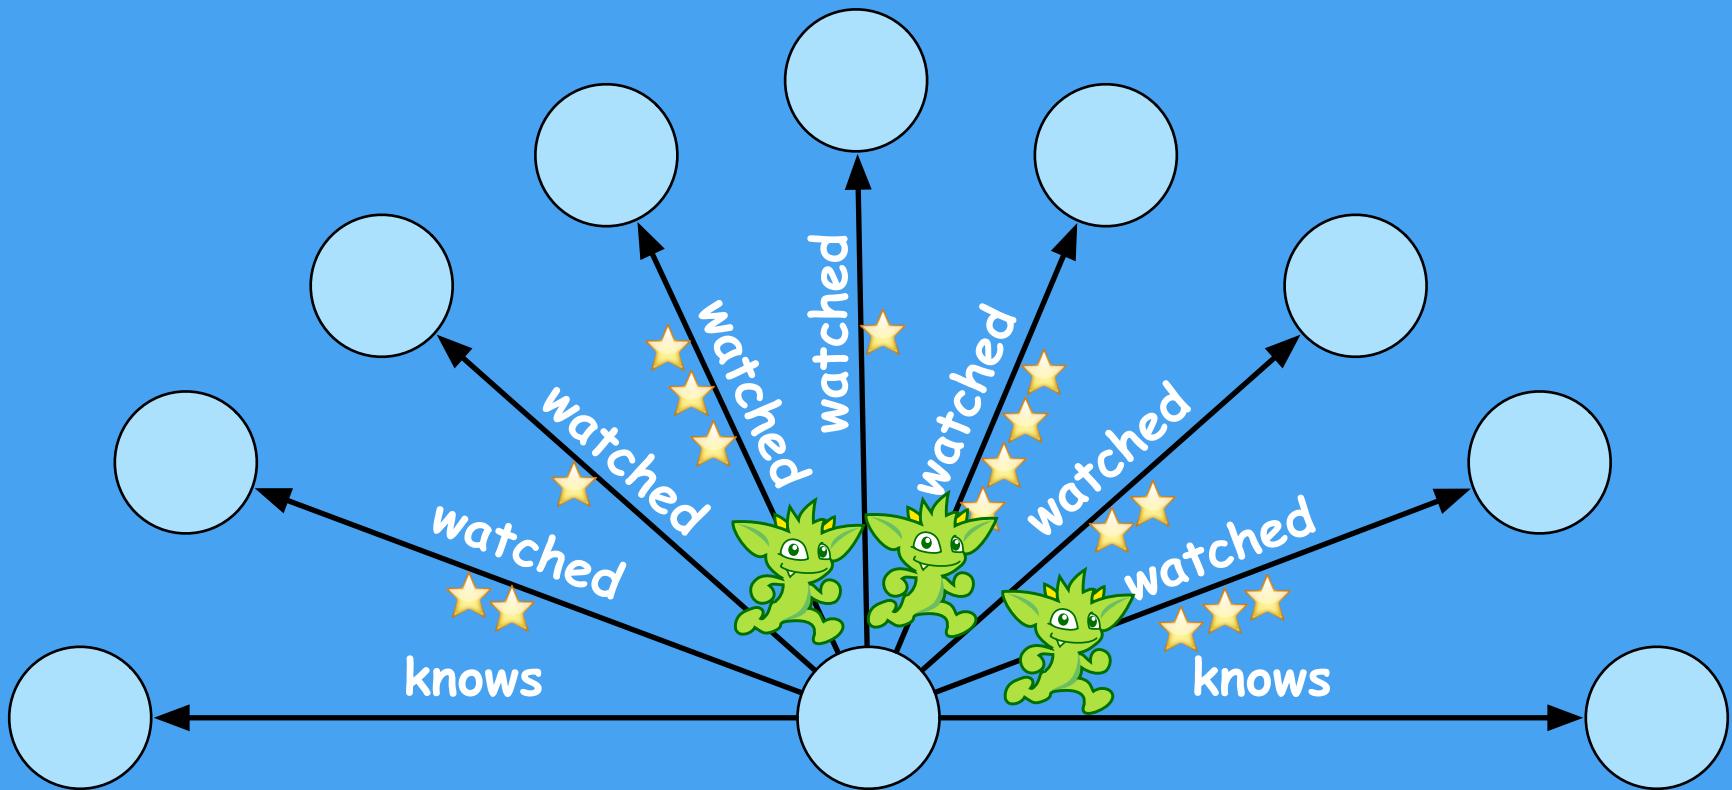
```
g.V(1).out('watched').count()
```



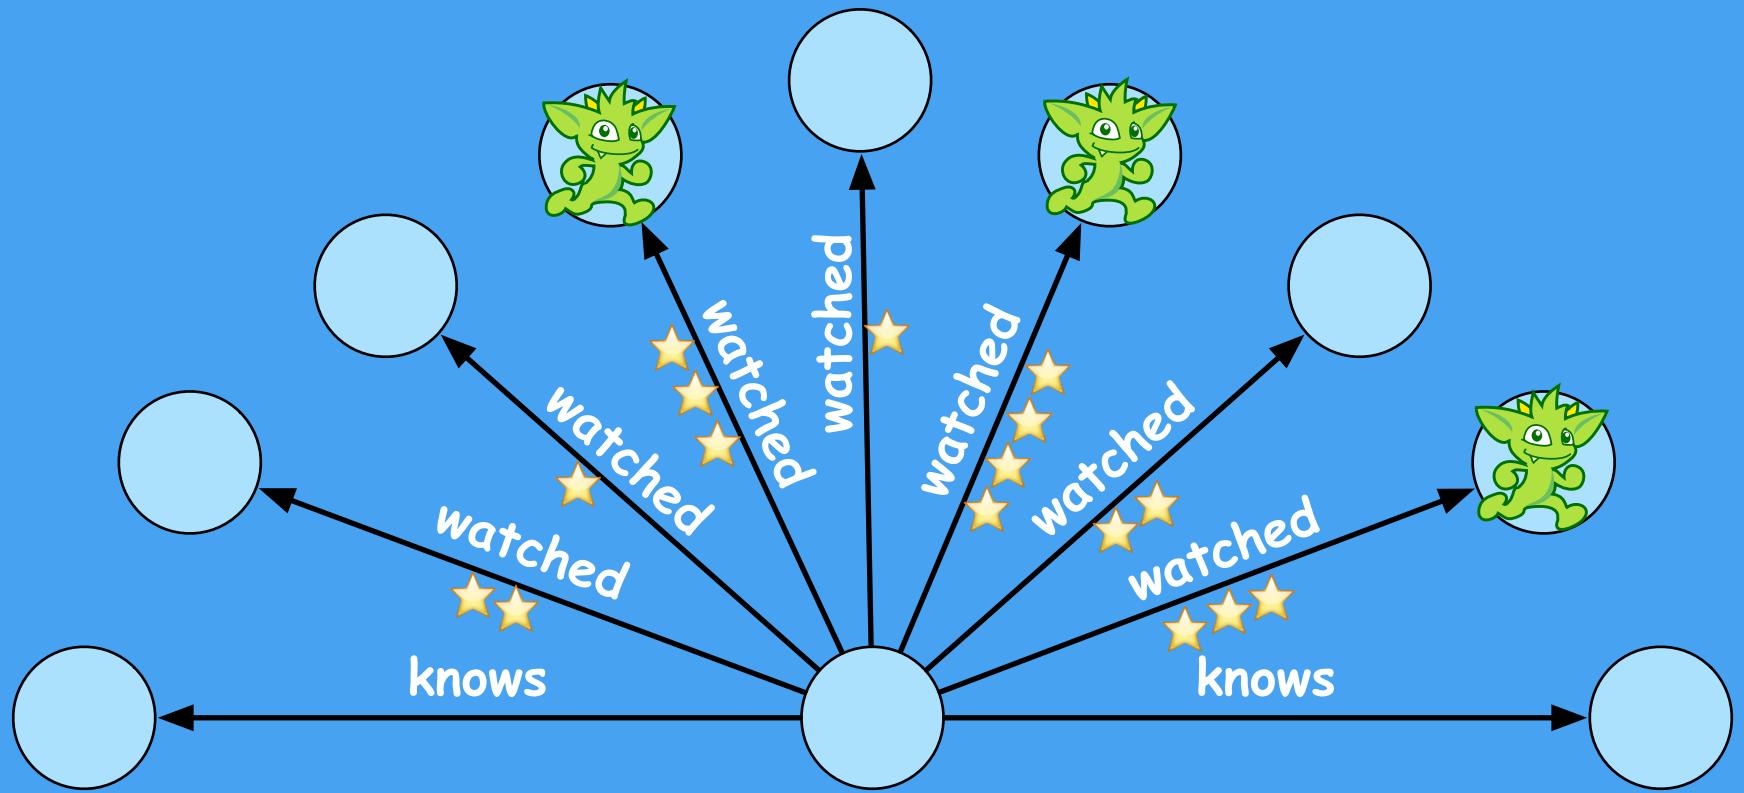
`g.V(1).outE('watched')`



```
g.V(1).outE('watched').has('stars',gte(3))
```



```
g.V(1).outE('watched').has('stars',gte(3))
```



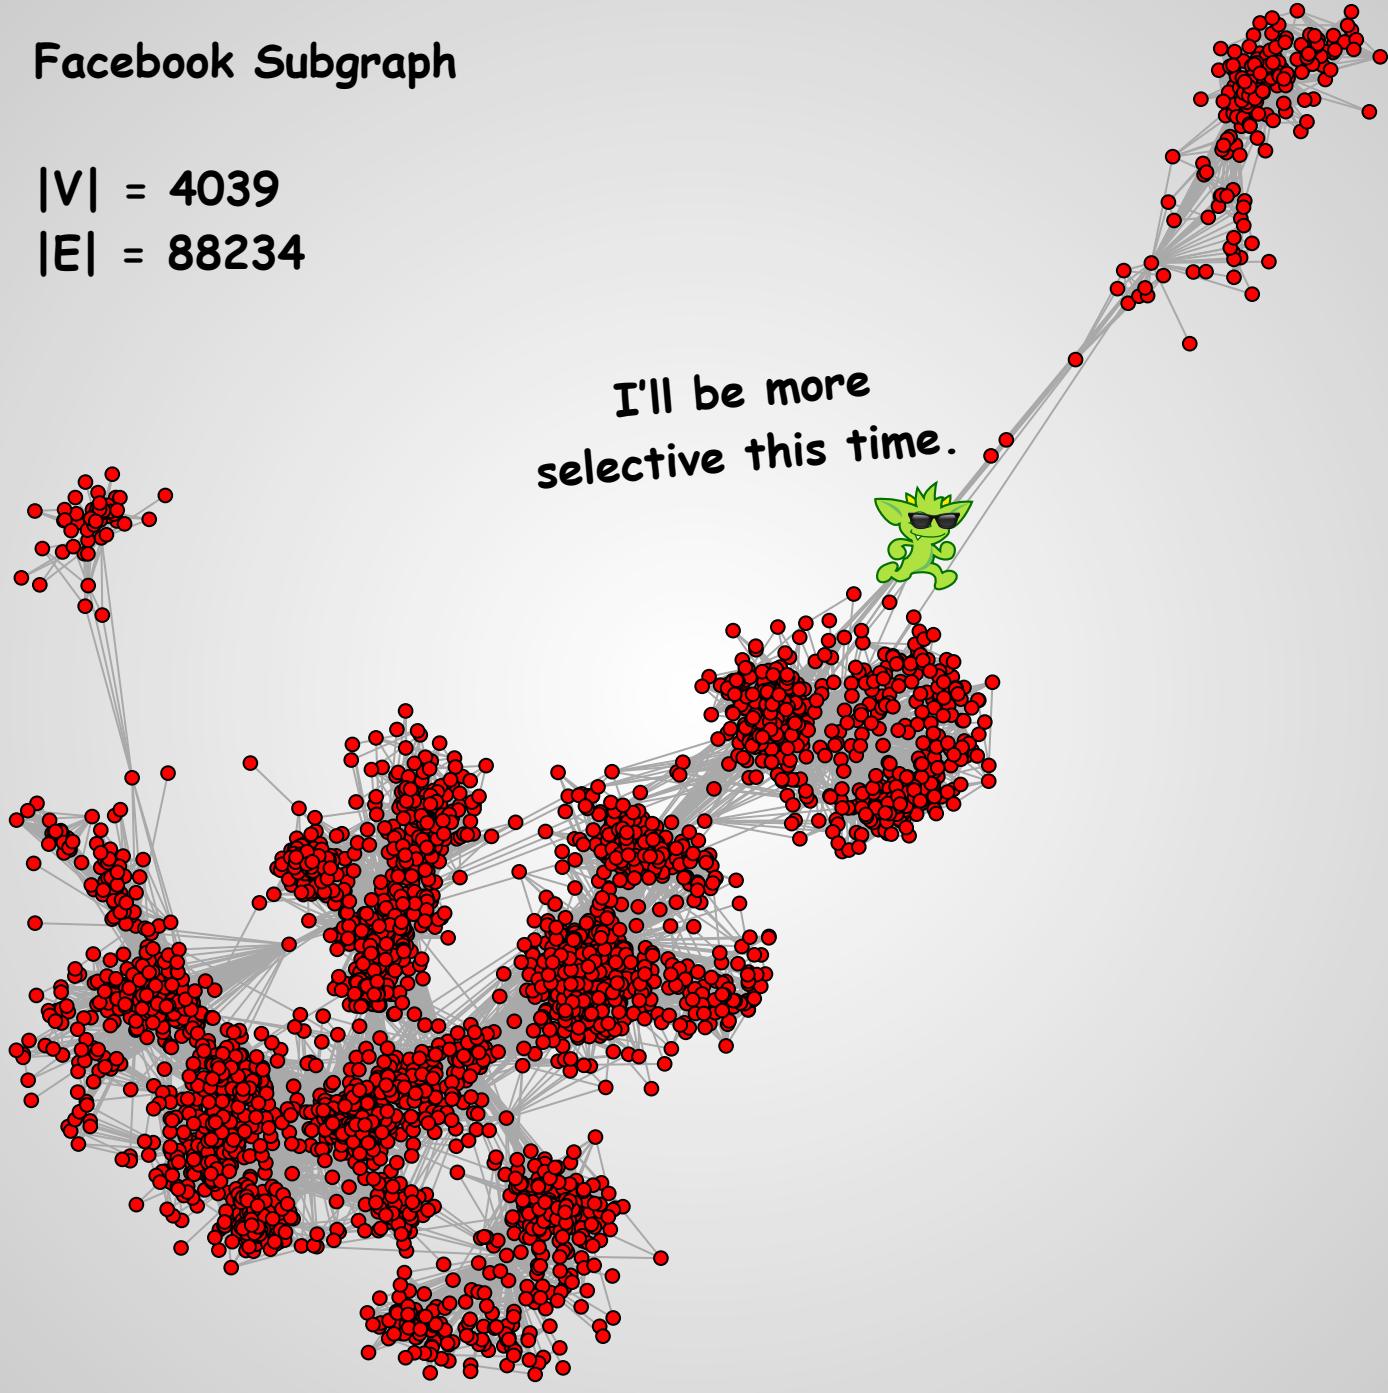
```
g.V(1).outE('watched').has('stars',gte(3)).inV().count()
```

# Facebook Subgraph

$|V| = 4039$

$|E| = 88234$

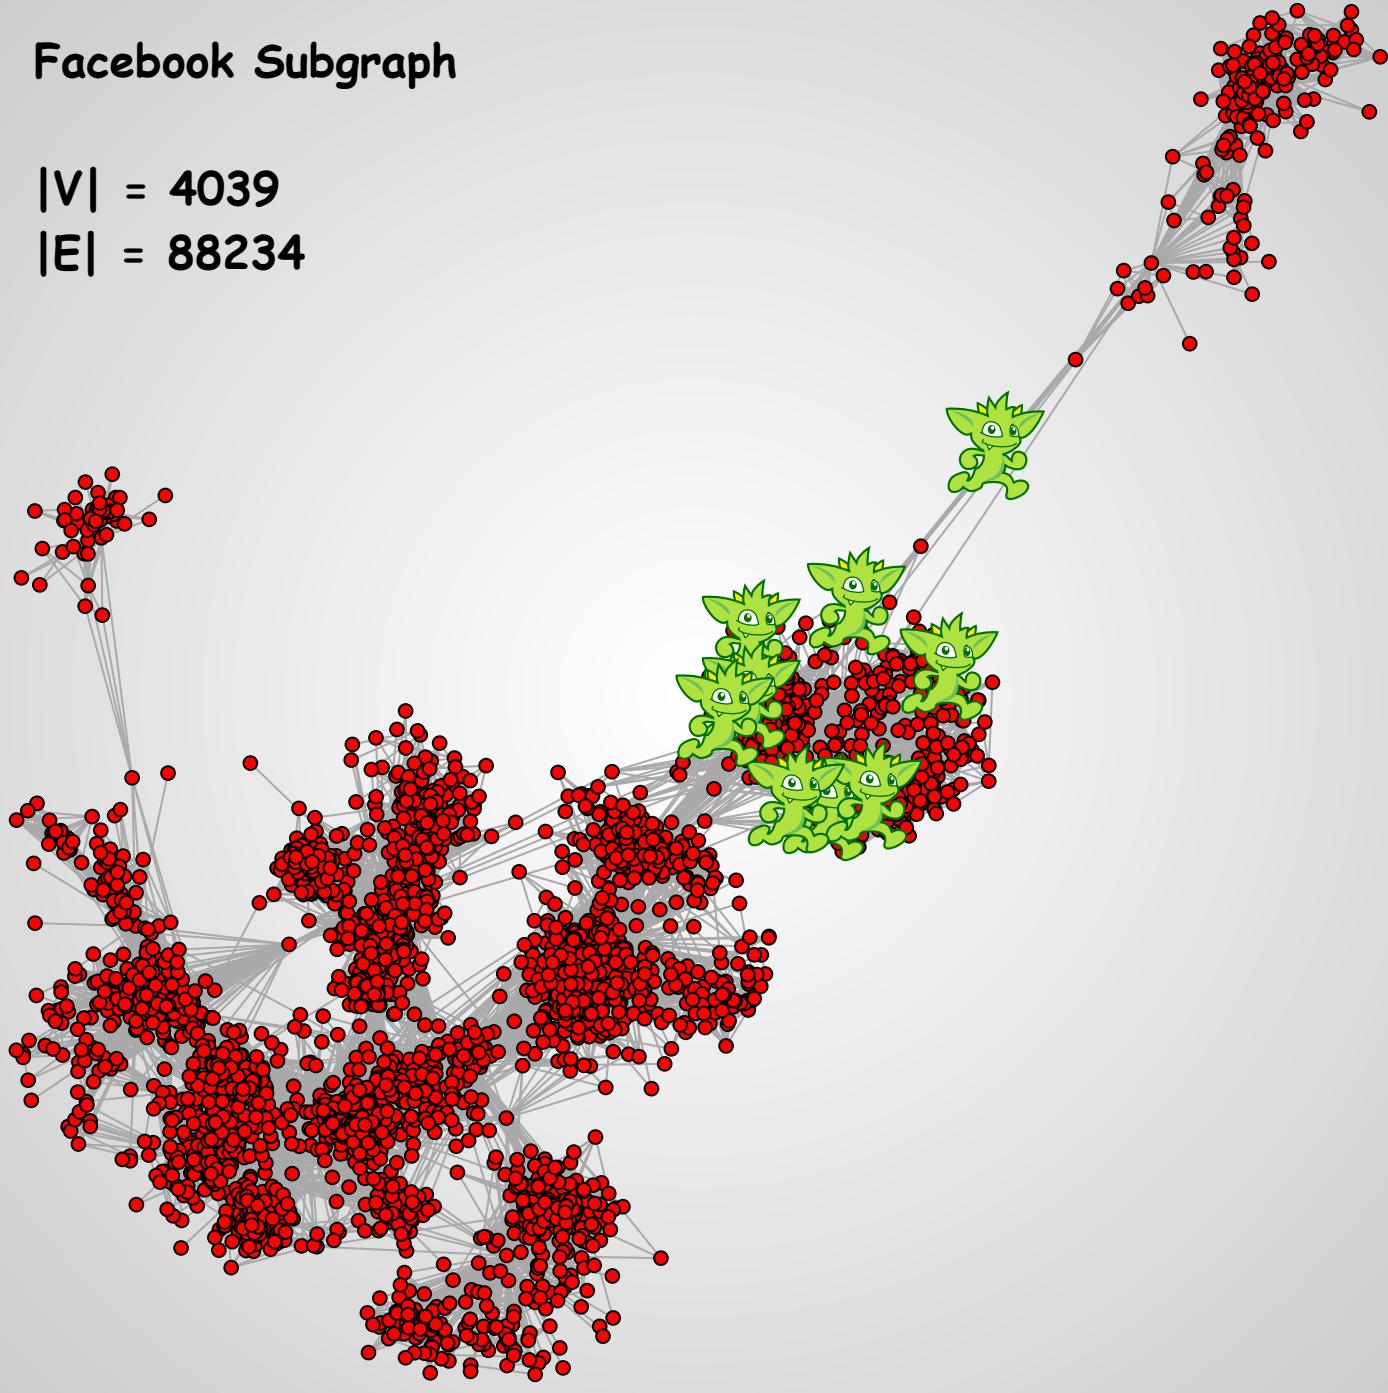
I'll be more  
selective this time.



## Facebook Subgraph

$|V| = 4039$

$|E| = 88234$





# Quiz

# Controlling Traverser Populations

out()

out('label')

outE().has('key', value).inV()

local(out().limit(10))

local(out().coin(0.5))

local(outE().sample(10).by('weight').inV())



# Controlling Traverser Populations

`out()`

*only edges in one direction*

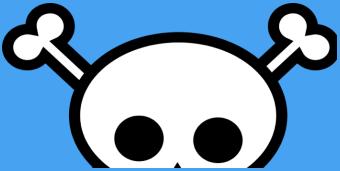
`out('label')`

`outE().has('key', value).inV()`

`local(out().limit(10))`

`local(out().coin(0.5))`

`local(outE().sample(10).by('weight').inV())`



# Controlling Traverser Populations

`out()`

only edges in one direction

`out('label')`

only edges w/ label

`outE().has('key', value).inV()`

`local(out().limit(10))`

`local(out().coin(0.5))`

`local(outE().sample(10).by('weight').inV())`



# Controlling Traverser Populations

`out()`

only edges in one direction

`out('label')`

only edges w/ label

`outE().has('key', value).inV()`

only edges w/ property condition

`local(out().limit(10))`

`local(out().coin(0.5))`

`local(outE().sample(10).by('weight').inV())`



# Controlling Traverser Populations

`out()`

only edges in one direction

`out('label')`

only edges w/ label

`outE().has('key', value).inV()`

only edges w/ property condition

`local(out().limit(10))`

only the first 10 edges

`local(out().coin(0.5))`

`local(outE().sample(10).by('weight').inV())`



# Controlling Traverser Populations

`out()`

only edges in one direction

`out('label')`

only edges w/ label

`outE().has('key', value).inV()`

only edges w/ property condition

`local(out().limit(10))`

only the first 10 edges

`local(out().coin(0.5))`

only 50% of edges

`local(outE().sample(10).by('weight').inV())`



# Controlling Traverser Populations

`out()`

only edges in one direction

`out('label')`

only edges w/ label

`outE().has('key', value).inV()`

only edges w/ property condition

`local(out().limit(10))`

only the first 10 edges

`local(out().coin(0.5))`

only 50% of edges

`local(outE().sample(10).by('weight').inV())`

only 10 edges biased by weight



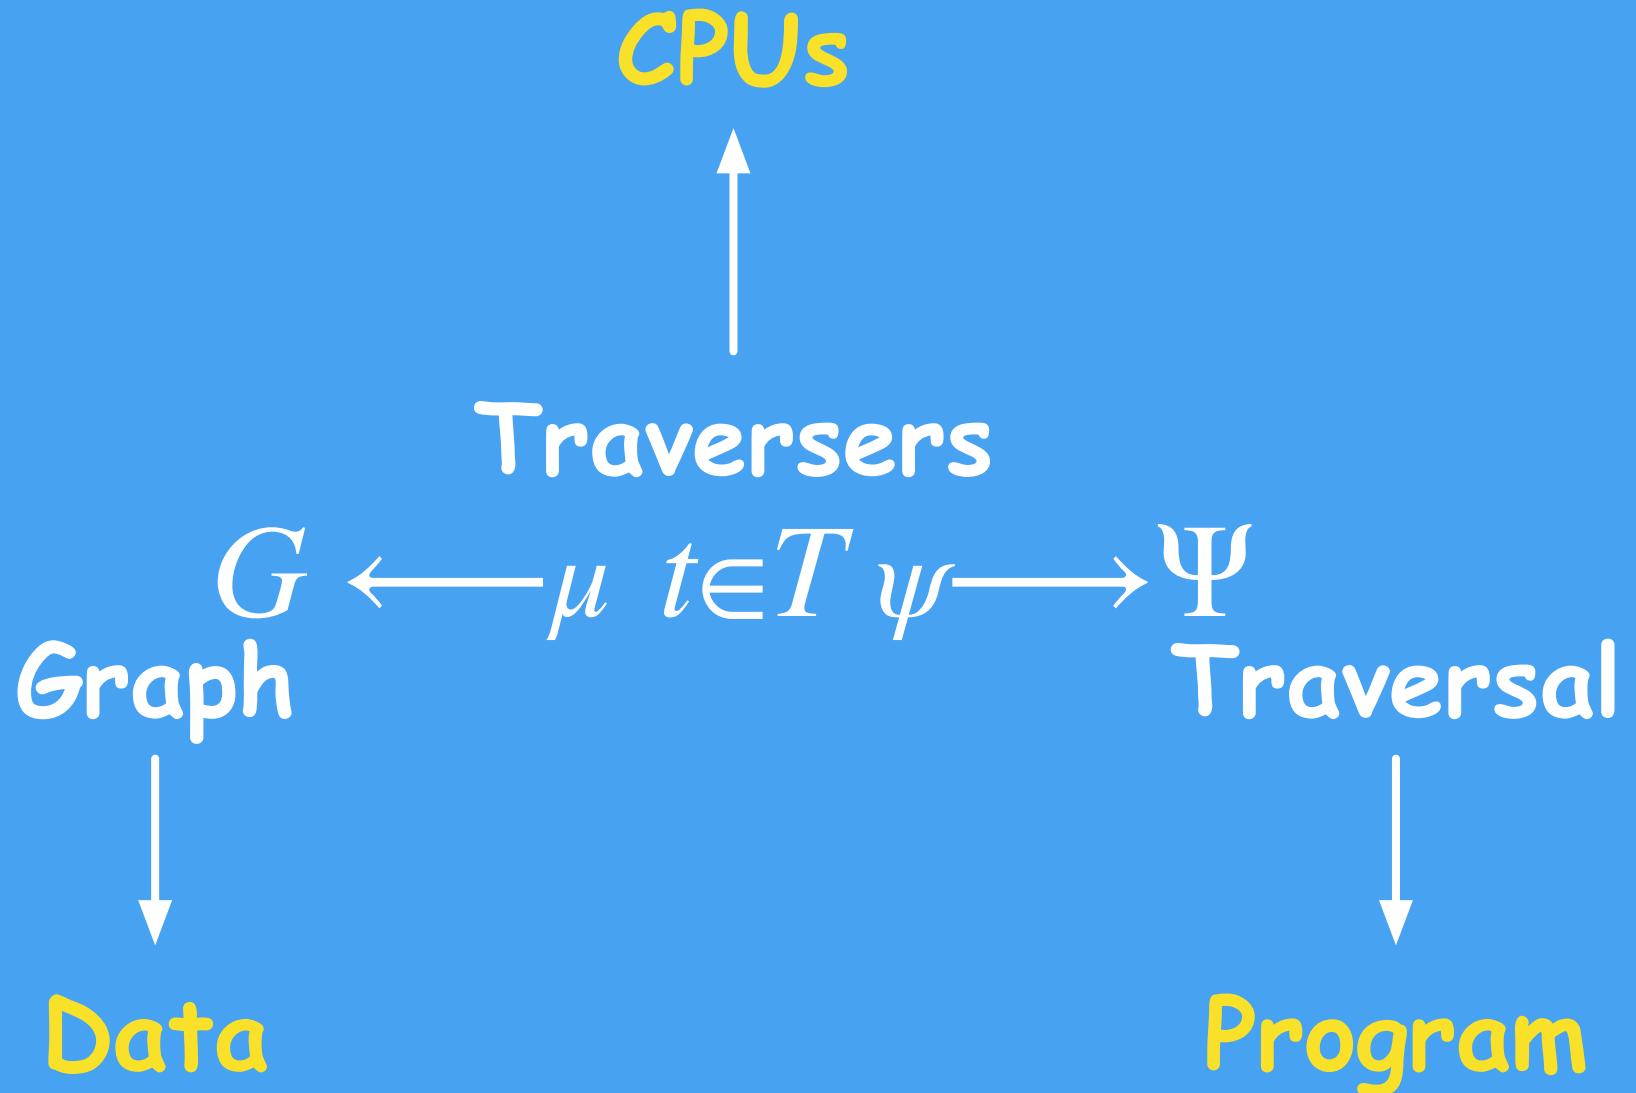
The computational complexity of the traversal is a function of the structural complexity of the graph.



**Traversers**

$$G \xleftarrow{\mu} t \in T \psi \xrightarrow{\Psi} \text{Traversal}$$

**Graph**





The show is over.  
Please tip your waitress and ~~drive~~<sup>traverse</sup> home safe!

