

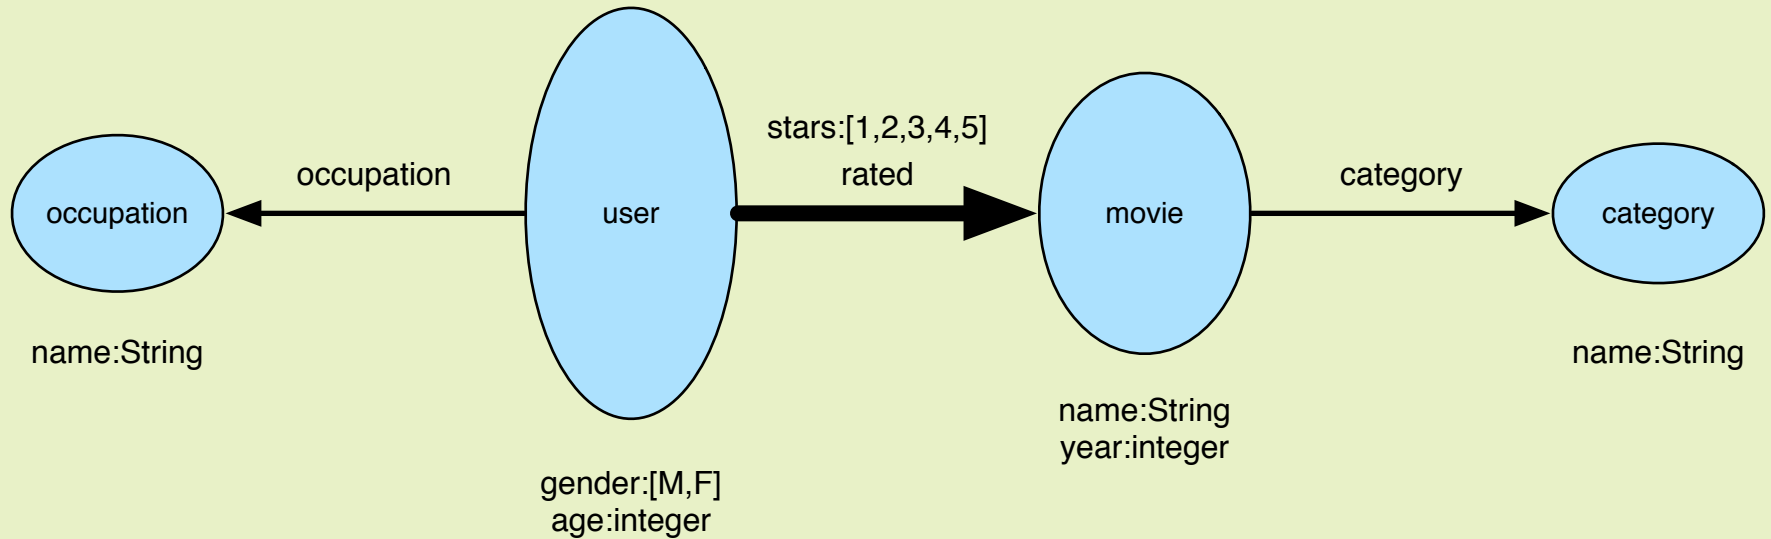
The Gremlin Graph Traversal Language



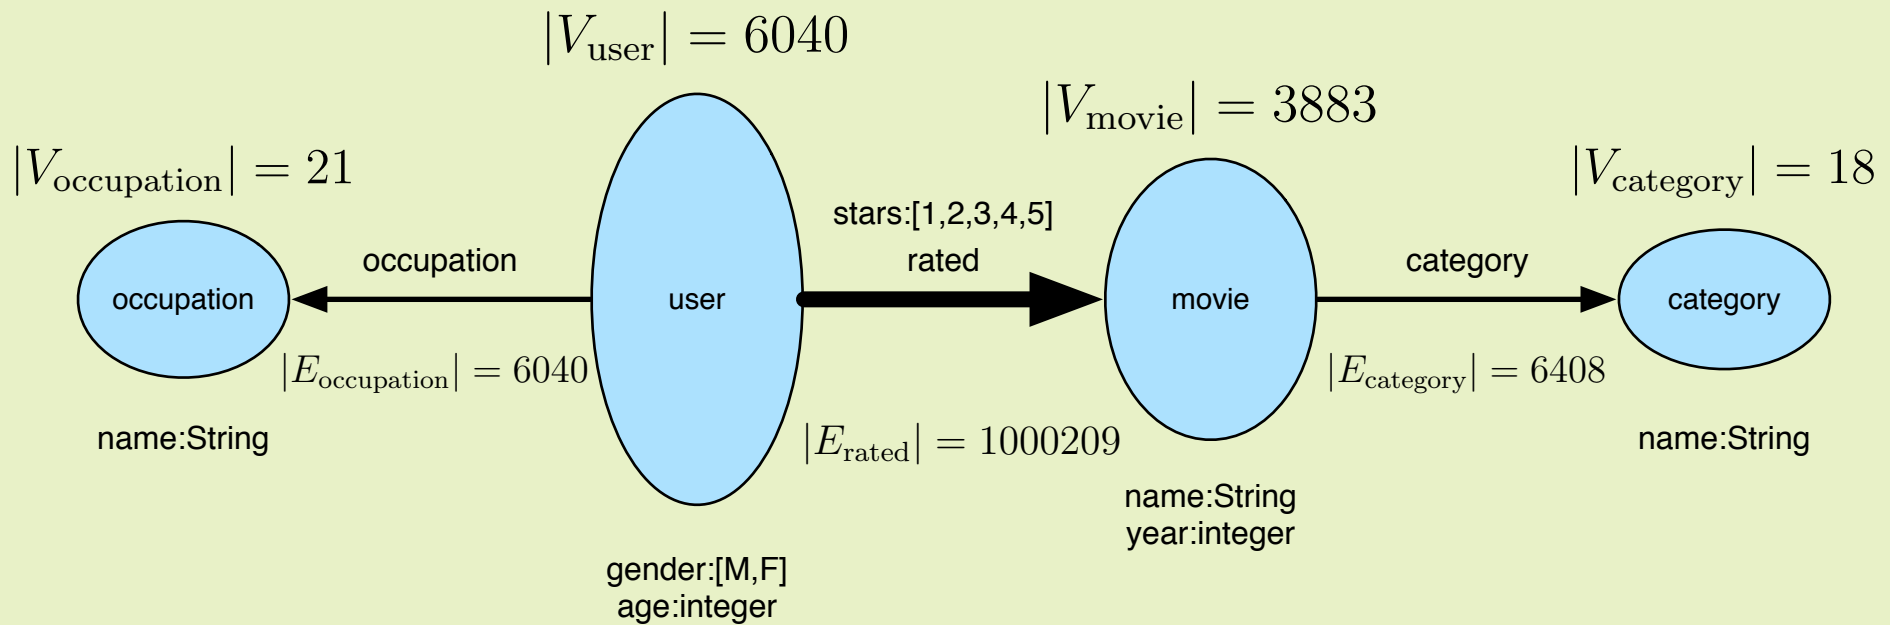
Marko A. Rodriguez and Daniel Kuppitz: 
DATASTAX

2015 NOSQL NOW!
AUG 18-20 | SAN JOSE, CA

MovieLens Dataset



MovieLens Dataset



$$G = (V, E)$$

```
~/tinkerpop3$ bin/gremlin.sh
```



```
gremlin>
```

```
~/tinkerpop3$ bin/gremlin.sh
```

```
  \,,,/
```

```
 (o o)
```

```
-----o00o-(3)-o00o-----
```

```
plugin activated: tinkerpop.server
```

```
plugin activated: tinkerpop.utilities
```

```
plugin activated: tinkerpop.tinkergraph
```

```
gremlin>
```

```
~/tinkerpop3$ bin/gremlin.sh
```

```
  \,,,/  
  (o o)
```

```
-----o00o-(3)-o00o-----
```

```
plugin activated: tinkerpop.server
```

```
plugin activated: tinkerpop.utilities
```

```
plugin activated: tinkerpop.tinkergraph
```

```
gremlin>
```

```
Gremlin-Java8  
Gremlin-Groovy*  
Gremlin-Scala  
Gremlin-Clojure  
Gremlin-JavaScript  
Gremlin-Python  
Gremlin-PHP  
...
```

```
~/tinkerpop3$ bin/gremlin.sh
```

```
  \,,,/  
  (o o)
```

```
-----o00o-(3)-o00o-----
```

```
plugin activated: tinkerpop.server
```

```
plugin activated: tinkerpop.utilities
```

```
plugin activated: tinkerpop.tinkergraph
```

```
gremlin> graph = TinkerGraph.open()
```

```
==>tinkergraph[vertices:0 edges:0]
```

```
gremlin>
```

"Create a new TinkerGraph."

```
~/tinkerpop3$ bin/gremlin.sh
```

```
  \,,,/  
  (o o)
```

```
-----o00o-(3)-o00o-----
```

```
plugin activated: tinkerpop.server
```

```
plugin activated: tinkerpop.utilities
```

```
plugin activated: tinkerpop.tinkergraph
```

```
gremlin> graph = TinkerGraph.open()
```

```
==>tinkergraph[vertices:0 edges:0]
```

```
gremlin>
```

$$G = (V = \emptyset, E = \emptyset)$$

G The graph is a set of vertices and edges

V The set of vertices in the graph

E The set of edges in the graph

\emptyset The empty set -- no elements

"Create a new TinkerGraph."


```
~/tinkerpop3$ bin/gremlin.sh
```

```
  \,,,/  
  (o o)
```

```
-----o00o-(3)-o00o-----
```

```
plugin activated: tinkerpop.server
```

```
plugin activated: tinkerpop.utilities
```

```
plugin activated: tinkerpop.tinkergraph
```

```
gremlin> graph = TinkerGraph.open()
```

```
==>tinkergraph[vertices:0 edges:0]
```

```
gremlin>
```

$$G = (V = \emptyset, E = \emptyset)$$

```
TitanGraph.open(...)  
Neo4jGraph.open(...)  
OrientGraph.open(...)  
SqlgGraph.open(...)  
HadoopGraph.open(...)  
  GiraphGraphComputer  
  SparkGraphComputer  
ElasticGraph.open(...)  
...
```

"Create a new TinkerGraph."

```
~/tinkerpop3$ bin/gremlin.sh
```

```
  \,,,/  
  (o o)
```

```
-----o00o-(3)-o00o-----
```

```
plugin activated: tinkerpop.server
```

```
plugin activated: tinkerpop.utilities
```

```
plugin activated: tinkerpop.tinkergraph
```

```
gremlin> graph = TinkerGraph.open()
```

```
==>tinkergraph[vertices:0 edges:0]
```

```
gremlin> graph.io(gryo()).readGraph('/tmp/movie-lens.kryo')
```

```
==>null
```

```
gremlin>
```

$$G = (V \neq \emptyset, E \subseteq (V \times V))$$

$A \subseteq B$ Set A is a subset of (or equal to) set B

$(V \times V)$ The set of all ordered pairs of vertices (directed binary edges)

"Load the MovieLens dataset into the newly created TinkerGraph."



```
~/tinkerpop3$ bin/gremlin.sh
```

```
  \,,,/  
  (o o)
```

```
-----o00o-(3)-o00o-----
```

```
plugin activated: tinkerpop.server
```

```
plugin activated: tinkerpop.utilities
```

```
plugin activated: tinkerpop.tinkergraph
```

```
gremlin> graph = TinkerGraph.open()
```

```
==>tinkergraph[vertices:0 edges:0]
```

```
gremlin> graph.io(gryo()).readGraph('/tmp/movie-lens.kryo')
```

```
==>null
```

```
gremlin> g = graph.traversal()
```

```
==>graphtraversalsource[tinkergraph[vertices:9962 edges:1012657], standard]
```

```
gremlin>
```

"Create a graph traversal source for spawning graph traversals over the MovieLens graph."



```
~/tinkerpop3$ bin/gremlin.sh
```

```
  \,,,/
  (o o)
-----o00o-(3)-o00o-----
plugin activated: tinkerpop.server
plugin activated: tinkerpop.utilities
plugin activated: tinkerpop.tinkergraph
gremlin> graph = TinkerGraph.open()
==>tinkergraph[vertices:0 edges:0]
gremlin> graph.io(gryo()).readGraph('/tmp/movie-lens.kryo')
==>>null
gremlin> g = graph.traversal()
==>graphtraversalsource[tinkergraph[vertices:9962 edges:1012657], standard]
gremlin> g.V().count()
==>9962
gremlin>
```

$$|V| = 9962$$

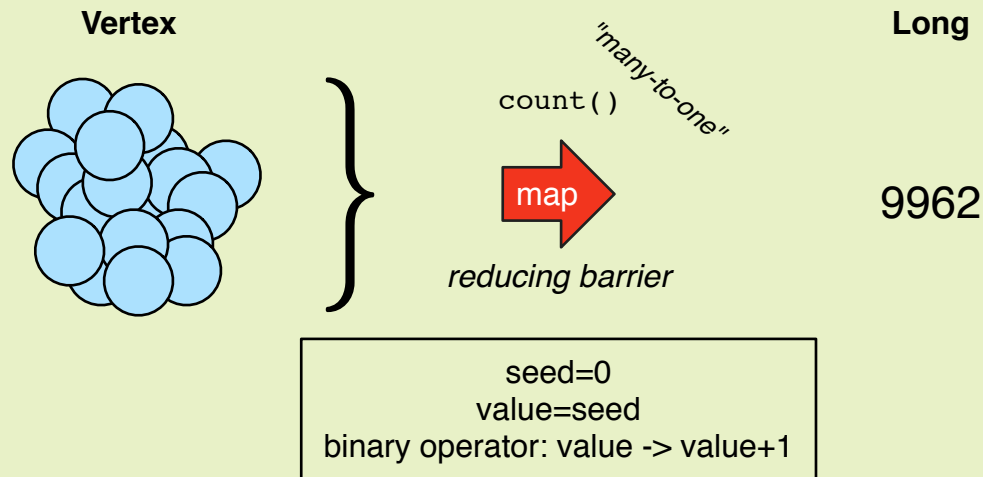
"Count the number of vertices in the graph."



```
~/tinkerpop3$ bin/gremlin.sh
```

```
  \,,,/
  (o o)
-----o00o-(3)-o00o-----
plugin activated: tinkerpop.server
plugin activated: tinkerpop.utilities
plugin activated: tinkerpop.tinkergraph
gremlin> graph = TinkerGraph.open()
==>tinkergraph[vertices:0 edges:0]
gremlin> graph.io(gryo()).readGraph('/tmp/movie-lens.kryo')
==>null
gremlin> g = graph.traversal()
==>graphtraversalsource[tinkergraph[vertices:9962 edges:1012657], standard]
gremlin> g.V().count()
==>9962
gremlin>
```

$$|V| = 9962$$




"Count the number of vertices in the graph."



```
~/tinkerpop3$ bin/gremlin.sh
```

```
  \,,,/
  (o o)
-----o00o-(3)-o00o-----
plugin activated: tinkerpop.server
plugin activated: tinkerpop.utilities
plugin activated: tinkerpop.tinkergraph
gremlin> graph = TinkerGraph.open()
==>tinkergraph[vertices:0 edges:0]
gremlin> graph.io(gryo()).readGraph('/tmp/movie-lens.kryo')
==>>null
gremlin> g = graph.traversal()
==>graphtraversalsource[tinkergraph[vertices:9962 edges:1012657], standard]
gremlin> g.V().count()
==>9962
gremlin> g.E().count()
==>1012657
gremlin>
```

$$|E| = 1012657$$



"Count the number of edges in the graph."



```
~/tinkerpop3$ bin/gremlin.sh
```

```
\\,\\,\\/  
(o o)
```

```
-----o00o-(3)-o00o-----
```

```
plugin activated: tinkerpop.server
```

```
plugin activated: tinkerpop.utilities
```

```
plugin activated: tinkerpop.tinkergraph
```

```
gremlin> graph = TinkerGraph.open()
```

```
==>tinkergraph[vertices:0 edges:0]
```

```
gremlin> graph.io(gryo()).readGraph('/tmp/movie-lens.kryo')
```

```
==>null
```

```
gremlin> g = graph.traversal()
```

```
==>graphtraversalsource[tinkergraph[vertices:9962 edges:1012657], standard]
```

```
gremlin> g.V().count()
```

```
==>9962
```

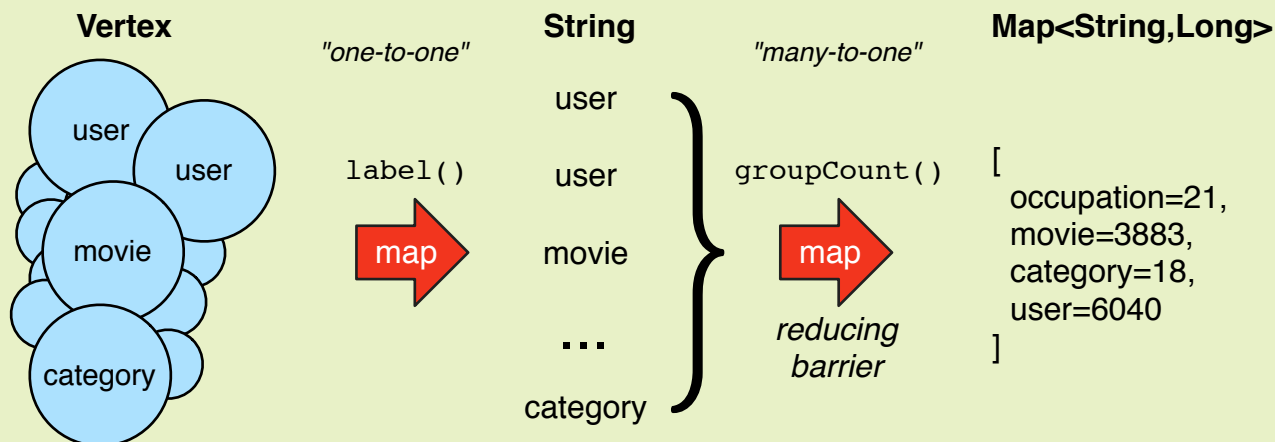
```
gremlin> g.E().count()
```

```
==>1012657
```

```
gremlin> g.V().label().groupCount()
```

```
==>[occupation:21, movie:3883, category:18, user:6040]
```

```
gremlin>
```



"For each vertex in the graph, emit its label, then group and count each distinct label."

```
~/tinkerpop3$ bin/gremlin.sh
```

```
  \,,,/  
  (o o)
```

```
-----o00o-(3)-o00o-----
```

```
plugin activated: tinkerpop.server
```

```
plugin activated: tinkerpop.utilities
```

```
plugin activated: tinkerpop.tinkergraph
```

```
gremlin> graph = TinkerGraph.open()
```

```
==>tinkergraph[vertices:0 edges:0]
```

```
gremlin> graph.io(gryo()).readGraph('/tmp/movie-lens.kryo')
```

```
==>null
```

```
gremlin> g = graph.traversal()
```

```
==>graphtraversalsource[tinkergraph[vertices:9962 edges:1012657], standard]
```

```
gremlin> g.V().count()
```

```
==>9962
```

```
gremlin> g.E().count()
```

```
==>1012657
```

```
gremlin> g.V().label().groupCount()
```

```
==>[occupation:21, movie:3883, category:18, user:6040]
```

```
gremlin> g.E().hasLabel('rated').values('stars').mean()
```

```
==>3.581564453029317
```

```
gremlin>
```

"For each rated-edge in the graph, emit its stars property value and compute the average value."


```
~/tinkerpop3$ bin/gremlin.sh
```

```
\\,\\,/
(o o)
```

```
-----o00o-(3)-o00o-----
```

```
plugin activated: tinkerpop.server
```

```
plugin activated: tinkerpop.utilities
```

```
plugin activated: tinkerpop.tinkergraph
```

```
gremlin> graph = TinkerGraph.open()
```

```
==>tinkergraph[vertices:0 edges:0]
```

```
gremlin> graph.io(gryo()).readGraph('/tmp/movie-lens.kryo')
```

```
==>null
```

```
gremlin> g = graph.traversal()
```

```
==>graphtraversalsource[tinkergraph[vertices:9962 edges:1012657], standard]
```

```
gremlin> g.V().count()
```

```
==>9962
```

```
gremlin> g.E().count()
```

```
==>1012657
```

```
gremlin> g.V().label().groupCount()
```

```
==>[occupation:21, movie:3883, category:18, user:6040]
```

```
gremlin> g.E().hasLabel('rated').values('stars').mean()
```

```
==>3.581564453029317
```

```
gremlin> g.V().hasLabel('user').map(outE('rated').count()).max()
```

```
==>2314
```

```
gremlin>
```

"What is the maximum number of movies a single user rated?"

```
~/tinkerpop3$ bin/gremlin.sh
```

```
  \,,,/  
  (o o)
```

```
-----o00o-(3)-o00o-----
```

```
plugin activated: tinkerpop.server
```

```
plugin activated: tinkerpop.utilities
```

```
plugin activated: tinkerpop.tinkergraph
```

```
gremlin> graph = TinkerGraph.open()
```

```
==>tinkergraph[vertices:0 edges:0]
```

```
gremlin> graph.io(gryo()).readGraph('/tmp/movie-lens.kryo')
```

```
==>null
```

```
gremlin> g = graph.traversal()
```

```
==>graphtraversalsource[tinkergraph[vertices:9962 edges:1012657], standard]
```

```
gremlin> g.V().count()
```

```
==>9962
```

```
gremlin> g.E().count()
```

```
==>1012657
```

```
gremlin> g.V().label().groupCount()
```

```
==>[occupation:21, movie:3883, category:18, user:6040]
```

```
gremlin> g.E().hasLabel('rated').values('stars').mean()
```

```
==>3.581564453029317
```

```
gremlin> g.V().hasLabel('user').map(outE('rated').count()).max()
```

```
==>2314
```


```
gremlin> g.V().hasLabel('movie').values('year').min()
```

```
==>1919
```

```
gremlin>
```

"What year was the oldest movie made?"

```
gremlin> g.V().hasLabel('category').values('name')
==>Animation
==>Children's
==>Comedy
==>Adventure
==>Fantasy
==>Romance
==>Drama
==>Action
==>Crime
==>Thriller
==>Horror
==>Sci-Fi
==>Documentary
==>War
==>Musical
==>Mystery
==>Film-Noir
==>Western
```



"For each vertex that is labeled 'category,' emit the name property value of that vertex."

```
gremlin> g.V().hasLabel('category').as('a','b').  
  select('a','b').  
    by('name').  
    by(inE('category').count())
```

"For each category vertex, emit a map of its name and the number of movies it represents."

```
gremlin> g.V().hasLabel('category').as('a','b').
  select('a','b').
  by('name').
  by(inE('category').count())
```

$$V : \mathbb{G} \rightarrow V^*$$

$$\text{hasLabel}_{\text{category}} : V^* \rightarrow V^*$$

$$\text{as}_{a,b} : V^* \rightarrow (V \times V)^*$$

$$\text{select}_{a,b} : (V \times V)^* \rightarrow \left[\begin{array}{c} a \\ b \end{array} \begin{array}{c} \text{values}_{\text{name}} : V^* \rightarrow \mathbb{S} \\ (\text{inE}_{\text{category}} : V^* \rightarrow E^*) \circ (\text{count} : E^* \rightarrow \mathbb{N}) \end{array} \right] \rightarrow (\mathbb{S} \times \mathbb{N})^*$$

$$f : A \rightarrow B$$

The function f maps values of type A to values of type B

$$\mathbb{G}$$

The set of all graphs

$$A^*$$

A stream of values of type A

$$(A \times B)$$

The set of all pairs of values from A and B (cross product)

$$\mathbb{N}$$

The set of all natural numbers (1, 2, 3, 4, ...)

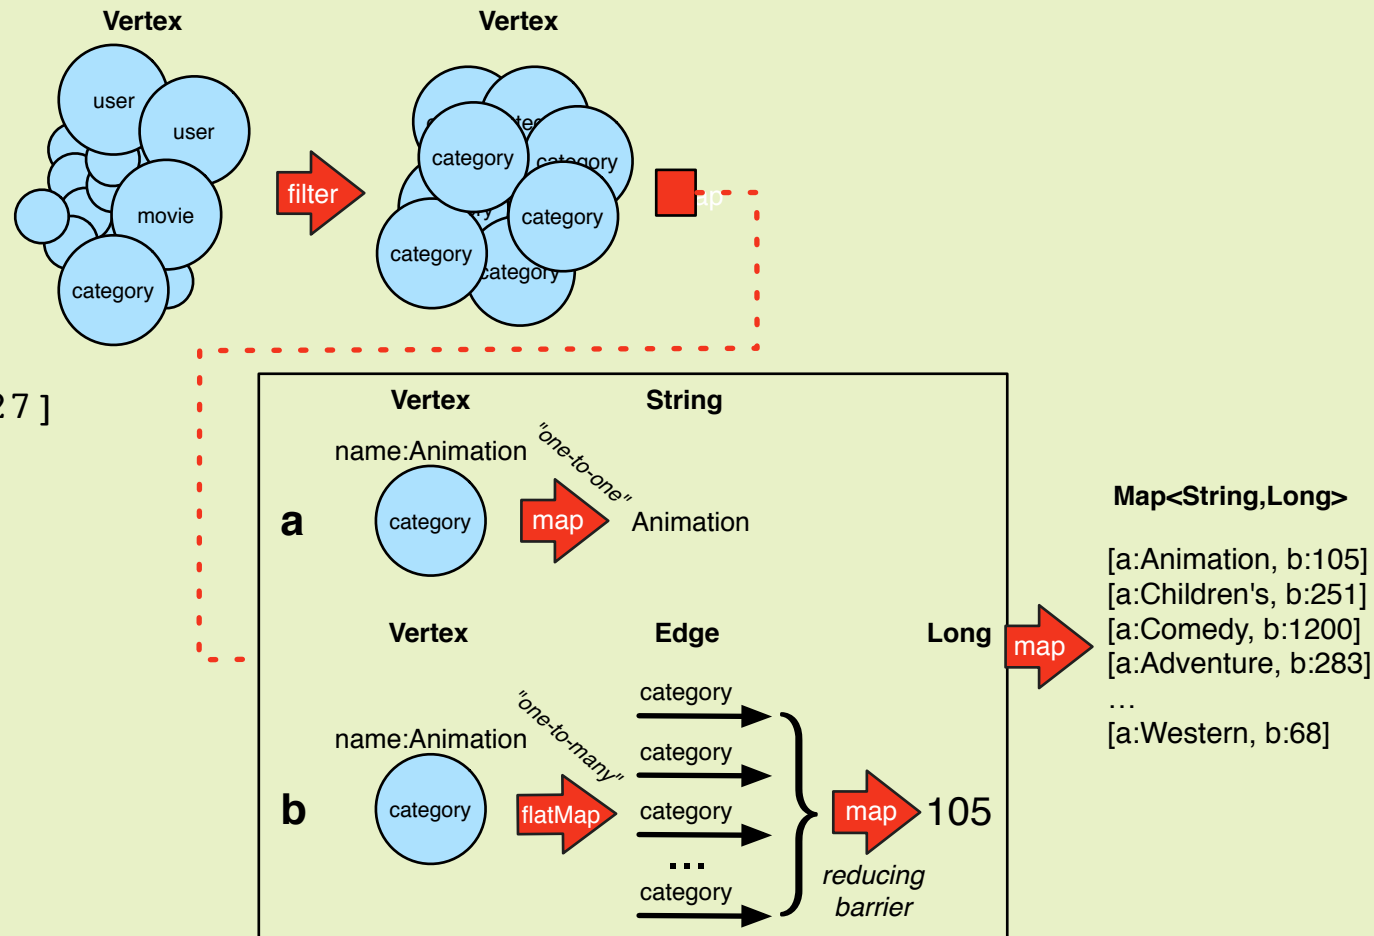
$$\mathbb{S}$$

The set of all strings (a, b, aa, ab, bb, ...) typically denoted Σ^*

"For each category vertex, emit a map of its name and the number of movies it represents."

```
gremlin> g.V().hasLabel('category').as('a','b').
  select('a','b').
  by('name').
  by(inE('category').count())
```

```
==>[a:Animation, b:105]
==>[a:Children's, b:251]
==>[a:Comedy, b:1200]
==>[a:Adventure, b:283]
==>[a:Fantasy, b:68]
==>[a:Romance, b:471]
==>[a:Drama, b:1603]
==>[a:Action, b:503]
==>[a:Crime, b:211]
==>[a:Thriller, b:492]
==>[a:Horror, b:343]
==>[a:Sci-Fi, b:276]
==>[a:Documentary, b:127]
==>[a:War, b:143]
==>[a:Musical, b:114]
==>[a:Mystery, b:106]
==>[a:Film-Noir, b:44]
==>[a:Western, b:68]
```



"For each category vertex, emit a map of its name and the number of movies it represents."

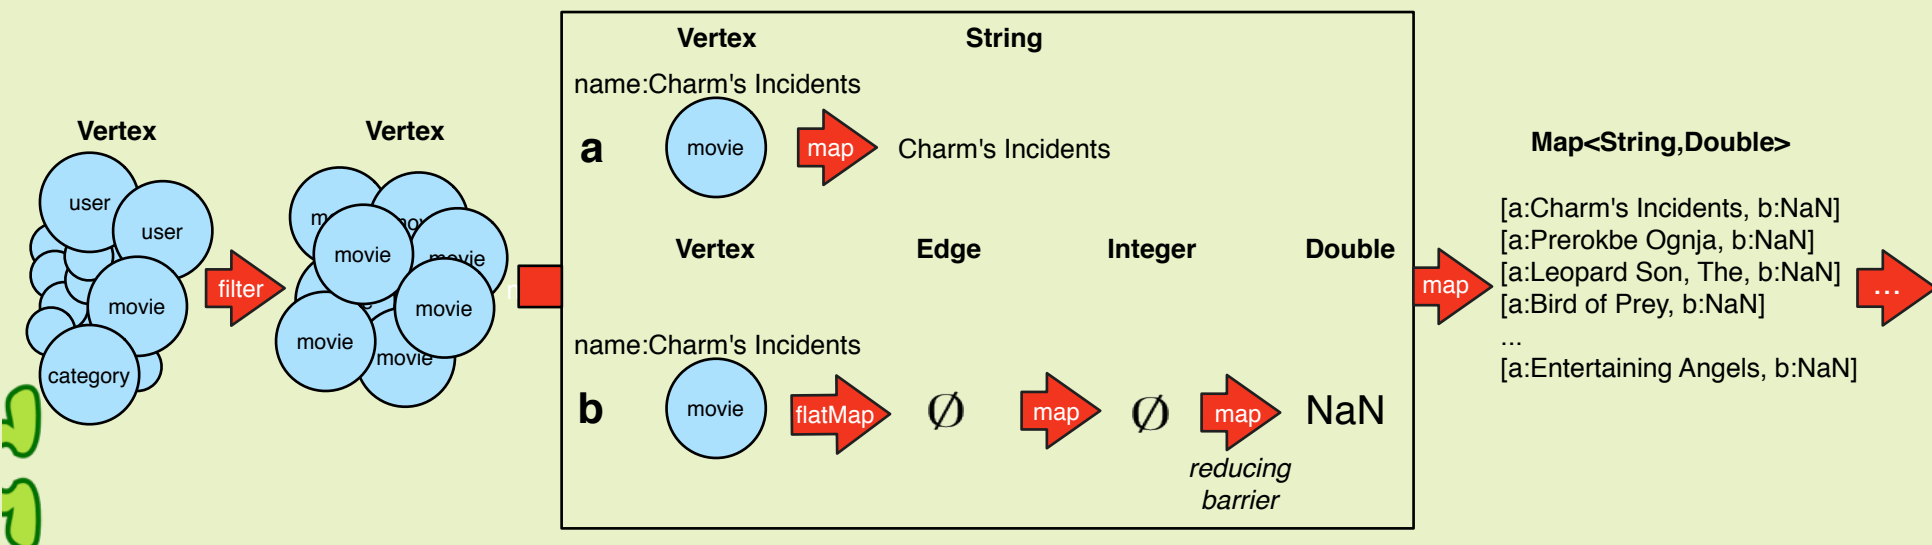
```
gremlin> g.V().hasLabel('movie').as('a','b').  
  select('a','b').  
    by('name').  
    by(inE('rated').values('stars').mean()).  
order().by(select('b'),decr).  
limit(10)
```



*"For each movie, emit a map of its name and average rating.
Sort the maps in decreasing order by their average rating. Emit the first 10 maps (i.e. top 10)."*

```
gremlin> g.V().hasLabel('movie').as('a','b').
  select('a','b').
  by('name').
  by(inE('rated').values('stars').mean()).
  order().by(select('b'),decr).
  limit(10)

==>[a:Charm's Incidents, b:NaN]
==>[a:Prerokbe Ognja, b:NaN]
==>[a:Leopard Son, The, b:NaN]
==>[a:Bird of Prey, b:NaN]
==>[a:Plutonium Circus, b:NaN]
==>[a:Hustler White, b:NaN]
==>[a:Curtis's Charm, b:NaN]
==>[a:Three Lives and Only One Death, b:NaN]
==>[a:Hoogste tijd, b:NaN]
==>[a:Entertaining Angels: The Dorothy Day Story, b:NaN]
```

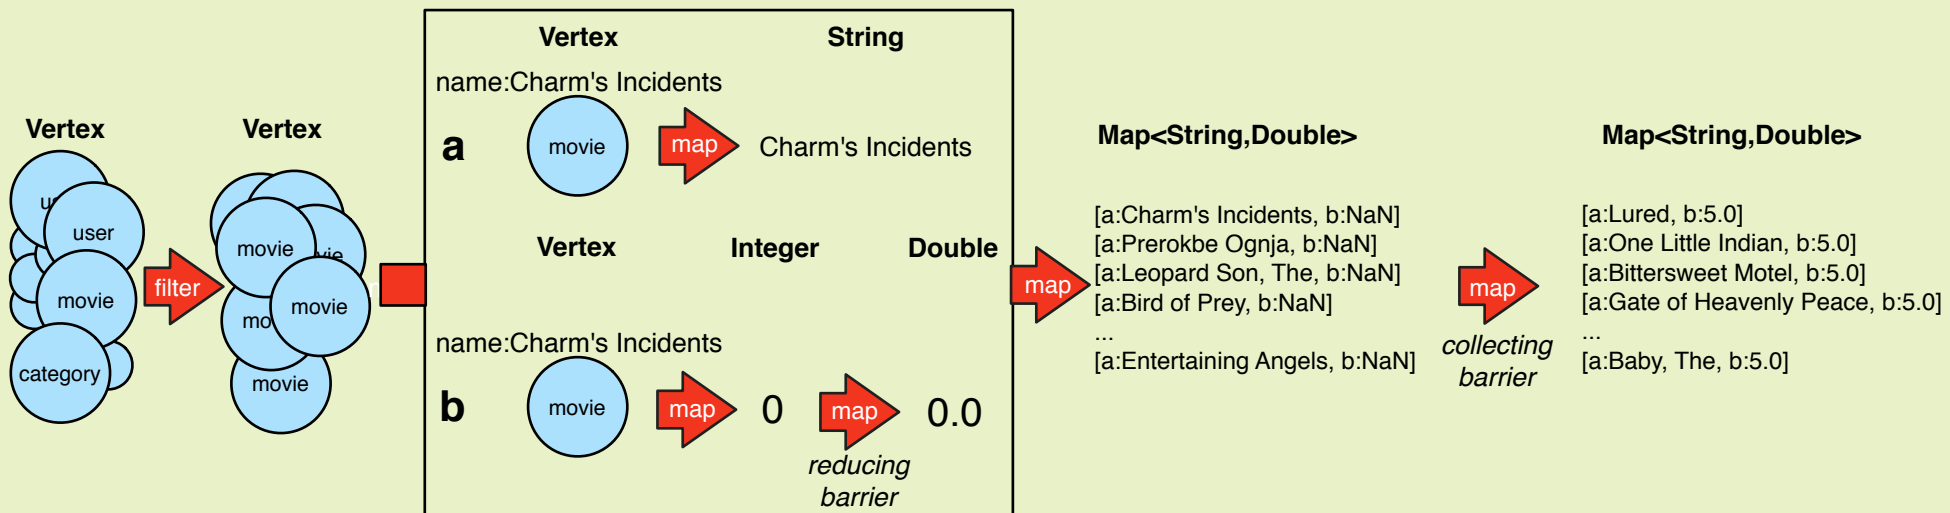


"For each movie, emit a map of its name and average rating.

Sort the maps in decreasing order by their average rating. Emit the first 10 maps (i.e. top 10)."


```
gremlin> g.V().hasLabel('movie').as('a','b').
  select('a','b').
  by('name').
  by(coalesce(
    inE('rated').values('stars'),
    constant(0)).mean()).
  order().by(select('b'),decr).
  limit(10)
```

```
==>[a:Lured, b:5.0]
==>[a:One Little Indian, b:5.0]
==>[a:Bittersweet Motel, b:5.0]
==>[a:Gate of Heavenly Peace, The, b:5.0]
==>[a:Follow the Bitch, b:5.0]
==>[a:Schlafes Bruder (Brother of Sleep), b:5.0]
==>[a:Ulysses (Ulissee), b:5.0]
==>[a:Song of Freedom, b:5.0]
==>[a:Smashing Time, b:5.0]
==>[a:Baby, The, b:5.0]
```



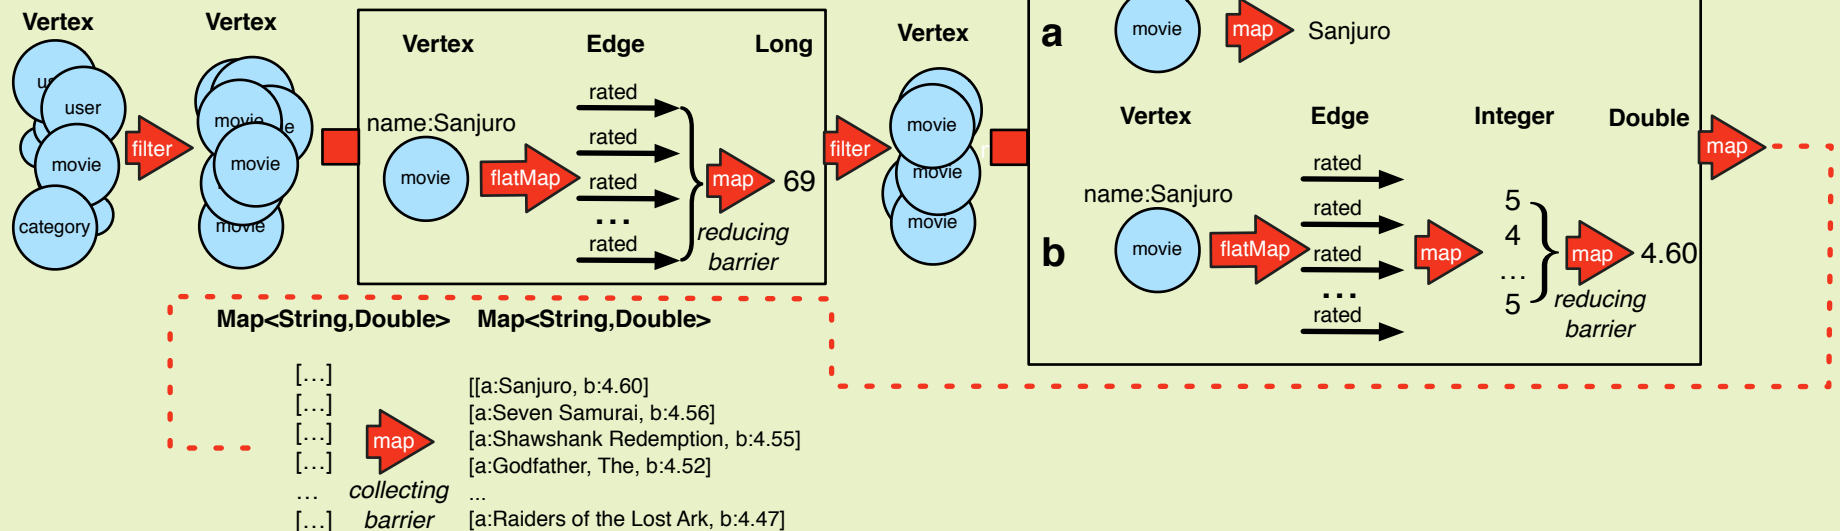
"For each movie, get its name and mean rating (or 0 if no ratings). Order by average rating and emit top 10."

```
gremlin> g.V().hasLabel('movie').as('a','b').  
  where(inE('rated').count().is(gt(10))).  
  select('a','b').  
    by('name').  
    by(inE('rated').values('stars').mean()).  
  order().by(select('b'),decr).  
  limit(10)
```

*"For each movie with at least 11 ratings, emit a map of its name and average rating.
Sort the maps in decreasing order by their average rating. Emit the first 10 maps (i.e. top 10)."*

```
gremlin> g.V().hasLabel('movie').as('a','b').
  where(inE('rated').count().is(gt(10))).
  select('a','b').
  by('name').
  by(inE('rated').values('stars').mean()).
  order().by(select('b'),decr).
  limit(10)
```

```
==>[a:Sanjuro, b:4.608695652173913]
==>[a:Seven Samurai (The Magnificent Seven), b:4.560509554140127]
==>[a:Shawshank Redemption, The, b:4.554557700942973]
==>[a:Godfather, The, b:4.524966261808367]
==>[a:Close Shave, A, b:4.52054794520548]
==>[a:Usual Suspects, The, b:4.517106001121705]
==>[a:Schindler's List, b:4.510416666666667]
==>[a:Wrong Trousers, The, b:4.507936507936508]
==>[a:Sunset Blvd. (a.k.a. Sunset Boulevard), b:4.491489361702127]
==>[a:Raiders of the Lost Ark, b:4.47772]
```



"For each movie with at least 11 ratings, emit a map of its name and average rating. Sort the maps in decreasing order by their average rating. Emit the first 10 maps (i.e. top 10)."

```

gremlin> g.V().hasLabel('movie').
    where(inE('rated').count().is(gt(10))).
    toString()
==>[GraphStep([],vertex), HasStep([~label.eq(movie)]),
    TraversalFilterStep([
    VertexStep(IN,[rated],edge),
    CountGlobalStep,
    IsStep(gt(10))] ) ]

```

$$V : \mathbb{G} \rightarrow V^*$$

$$\text{hasLabel}_{\text{movie}} : V^* \rightarrow V^*$$

$$\begin{array}{lcl}
 & \text{inE}_{\text{rated}} : V^* \rightarrow E^* & \\
 \text{where} : V^* \rightarrow & \text{count} : E^* \rightarrow \mathbb{N} & \rightarrow V^* \\
 & \text{is}_{\text{gt}(10)} : \mathbb{N} \rightarrow (\mathbb{N} \cup \emptyset) & \text{"true or false"}
 \end{array}$$

"What is the execution plan for the traversal prior to compiler optimizations being applied?"

```

gremlin> g.V().hasLabel('movie').
    where(inE('rated').count().is(gt(10))).
    iterate().toString()
==>[TinkerGraphStep(vertex,[~label.eq(movie)]),
    TraversalFilterStep([
    VertexStep(IN,[rated],edge),
    RangeGlobalStep(0,11),
    CountGlobalStep,
    IsStep(gt(10))] ) ]

```

$$V_{\text{label=movie}} : \mathbb{G} \rightarrow V^*$$

* **TinkerGraphStrategy**: Access vendor-specific vertex partition by label.

$$\text{inE}_{\text{rated}} : V^* \rightarrow E^*$$

$$\text{limit}_{11} : E^* \rightarrow E^*$$

* **RangeByIsCountStrategy**: Only iterate 1 more than required count.

$$\text{where} : V^* \rightarrow \quad \quad \quad \rightarrow V^*$$

$$\text{count} : E^* \rightarrow \mathbb{N}$$

$$\text{is}_{\text{gt}(10)} : \mathbb{N} \rightarrow (\mathbb{N} \cup \emptyset)$$

"true or false"

"What is the execution plan for the traversal after compiler optimizations have been applied?"

```
gremlin> g.getStrategies()
==>ConjunctionStrategy
    a.and().b => and(a,b)
    a.or().b => or(a,b)
    a.or().b.and().c => or(a,and(b,c))
    a.and().b.or().c => or(and(a,b),c)
==>IncidentToAdjacentStrategy
    a.outE().inV().b => a.out().b
==>AdjacentToIncidentStrategy
    a.in().count().b => a.inE().count().b
    a.where(out()).b => a.where(outE()).b
    a.and(in(),out()).b => a.and(inE(),outE()).b
==>IdentityRemovalStrategy
    a.identity().b => a.b
==>FilterRankingStrategy
    a.order().dedup().b => a.dedup().order().b
    a.and(c,d).has().b => a.has().and(c,d).b
    a.simplePath().where().b => b.where().simplePath().a
==>MatchPredicateStrategy
    a.match(c,d).where(e).b => a.match(c,d,e)
    a.match(has(),c,d).b => a.has().match(c,d).b
==>RangeByIsCountStrategy
    a.count().is(0) => a.limit(1).count().is(0)
==>TinkerGraphStepStrategy
    V.has().has().b => V[has,has].b
==>ProfileStrategy
    a.b.c.profile() => a.profile().b.profile().c.profile()
==>ComputerVerificationStrategy
    a.order.b => IllegalStateException
    a.local(out().out()).b => IllegalStateException
```

"What compilation strategies are associated with the graph traversal source?"

```
gremlin> g.V().has('movie','name','Die Hard').  
          inE('rated').values('stars').mean()  
==>4.121848739495798
```

"What is Die Hard's average rating?"

```
gremlin> g.V().has('movie','name','Die Hard').
         inE('rated').values('stars').mean()
==>4.121848739495798
```

$$V : \mathbb{G} \rightarrow V^*$$

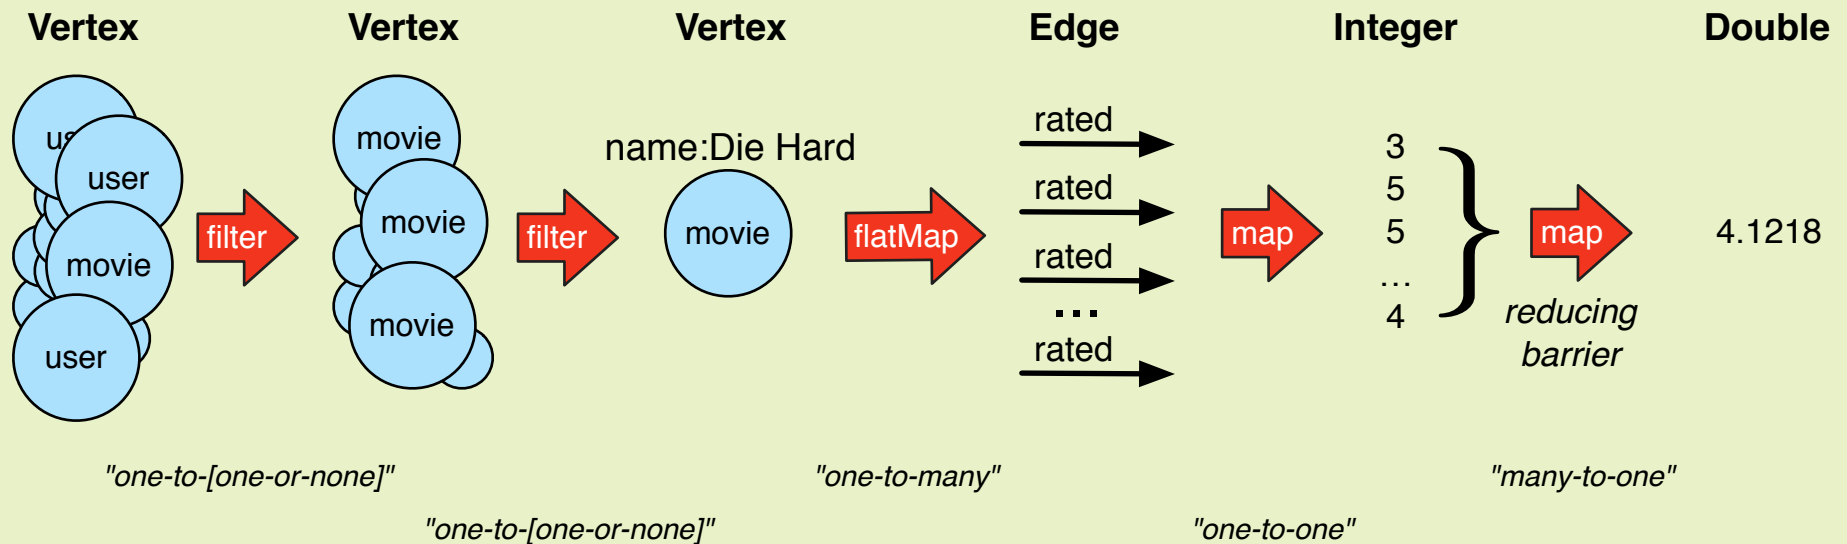
$$\text{has}_{\text{name=Die Hard}} : V^* \rightarrow V^*$$

$$\text{values}_{\text{stars}} : E^* \rightarrow \mathbb{N}^*$$

$$\text{hasLabel}_{\text{movie}} : V^* \rightarrow V^*$$

$$\text{inE}_{\text{rated}} : V^* \rightarrow E^*$$

$$\text{mean} : \mathbb{N}^* \rightarrow \mathbb{R}$$



"What is Die Hard's average rating?"

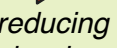

```
gremlin> g.V().has('movie','name','Die Hard').as('a').  
    inE('rated').has('stars',5).outV().  
        where(out('occupation').has('name','programmer')).  
    outE('rated').has('stars',5).inV().  
        where(neq('a')).  
    groupCount().by('name').  
        order(local).by(valueDecr).  
        limit(local,10).  
        unfold()                                // so its not printed on a single line
```

"Which programmers like Die Hard and what other movies do they like?"

Group and count the movies by their name. Sort the group count map in decreasing order by the count.

Clip the map to the top 10 entries and stream out the map's entries (for display purposes)."

```
// so its not printed on a single line
```



Group and count the movies by their name. Sort the group count map in decreasing order by the count.

Clip the map to the top 10 entries and stream out the map's entries (for display purposes)."

```

gremlin> g.V().has('movie','name','Die Hard').as('a').
    inE('rated').has('stars',5).outV().
    where(out('occupation').has('name','programmer')).
    outE('rated').has('stars',5).inV().
    where(neq('a')).
    groupCount().by('name').
    order(local).by(valueDecr).
    limit(local,10).
    unfold()                                // so its not printed on a single line
==>Raiders of the Lost Ark=36
==>Star Wars: Episode V - The Empire Strikes Back=36
==>Star Wars: Episode IV - A New Hope=34
==>Matrix, The=32
==>Terminator, The=29
==>Star Wars: Episode VI - Return of the Jedi=26
==>Sixth Sense, The=26
==>Braveheart=24
==>Aliens=23
==>Alien=22
gremlin>

```

"Which programmers like Die Hard and what other movies do they like?"

Group and count the movies by their name. Sort the group count map in decreasing order by the count.

Clip the map to the top 10 entries and stream out the map's entries (for display purposes)."

```
gremlin> g.V().
  match(
    __.as('a').hasLabel('movie'),
    __.as('a').out('category').has('name','Action'),
    __.as('a').has('year',between(1980,1990)),
    __.as('a').inE('rated').as('b'),
    __.as('b').has('stars',5),
    __.as('b').outV().as('c'),
    __.as('c').out('occupation').has('name','programmer'),
    __.as('c').has('age',between(30,40))).
select('a').groupCount().by('name').
order(local).by(valueDecr).
limit(local,10).
unfold()                                // so its not printed on a single line
```

*"What 80's action movies do 30-something programmers like?
Group count the movies by their name and sort the group count map in decreasing order by value.
Clip the map to the top 10 and emit the map entries."*

```

gremlin> g.V().
  match(
    __.as('a').hasLabel('movie'),
    __.as('a').out('category').has('name','Action'),
    __.as('a').has('year',between(1980,1990)),
    __.as('a').inE('rated').as('b'),
    __.as('b').has('stars',5),
    __.as('b').outV().as('c'),
    __.as('c').out('occupation').has('name','programmer'),
    __.as('c').has('age',between(30,40))).
  select('a').groupCount().by('name').
  order(local).by(valueDecr).
  limit(local,10).
  unfold()                                // so its not printed on a single line
==>Raiders of the Lost Ark=26
==>Star Wars: Episode V - The Empire Strikes Back=26
==>Terminator, The=23
==>Star Wars: Episode VI - Return of the Jedi=22
==>Princess Bride, The=19
==>Aliens=18
==>Boat, The (Das Boot)=11
==>Indiana Jones and the Last Crusade=11
==>Star Trek: The Wrath of Khan=10
==>Abyss, The=9
gremlin>

```

"What 80's action movies do 30-something programmers like?"

Group count the movies by their name and sort the group count map in decreasing order by value.

Clip the map to the top 10 and emit the map entries."

MatchStep

```
GraphTraversal.match(Traversal... traversalPatterns)
```

```
x.match(  
    a...b  
    a...c  
    c...  
    or(  
        a...c  
        a...b  
    )  
    c.repeat(...).b  
    not(c...a)  
    b...count().e  
    c...count().e  
).dedup(a,b).y
```

a,b,c,e : once a variable is set, it must hold equal for all patterns
c... : "predicate patterns" simply check for the existence of a result
or()/and() : nested conjunctive patterns supported
repeat(...) : recursive patterns supported
not(...) : not'ing of patterns supported
count() : barrier patterns supported
dedup(a,b) : internal de-duplication of variable values supported
x.match().y : possible to go from imperative to declarative, etc.

Plug and Play MatchAlgorithms


GreedyMatchAlgorithm :

try each pattern in the order provided by the user


CountMatchAlgorithm :

continually re-sort patterns by the cardinality of their set reductions

```
// CountMatchAlgorithm (default)
gremlin> clockWithResult(50){
    g.V().match(
        __.as('a').out('rated').as('b'),
        __.as('a').out('occupation').has('name','farmer')).
        select('a','b').count().next()}
==>66.31955294 // time in milliseconds
==>2706 // number of results
```



```
// GreedyMatchAlgorithm
gremlin> g = graph.traversal(GraphTraversalSource.build().
    with(MatchAlgorithmStrategy.build().
        algorithm(MatchStep.GreedyMatchAlgorithm).create()))
==>graphtraversalsource[tinkergraph[vertices:9962 edges:1012657], standard]
gremlin> clockWithResult(50){
    g.V().match(
        __.as('a').out('rated').as('b'),
        __.as('a').out('occupation').has('name','farmer')).
        select('a','b').count().next()}
==>1902.6290871599997 // time in milliseconds
==>2706 // number of results
```



"Which movies did each farmer rate? -- benchmark CountMatchAlgorithm vs. GreedyMatchAlgorithm."

```
gremlin> g.V().hasLabel('movie').
  where(inE('rated').count().is(gt(10))).
  group().
    by(((int)(it.value('year') / 10)) * 10).
    by().
    by(unfold().order().
      by(inE('rated').values('stars').mean(),decr).
      values('name').
      limit(1)).
  order(local).by(keyIncr).
  unfold()                                // so its not printed on a single line
```

"What is the most liked movie in each decade?"


```

gremlin> g.V().hasLabel('movie').
  where(inE('rated').count().is(gt(10))).
  group().
  by(((int)(it.value('year') / 10)) * 10).  $\lambda$ 
  by().
  by(unfold().order().
      by(inE('rated').values('stars').mean(),decr).
      values('name').
      limit(1)).
  order(local).by(keyIncr).
  unfold() // so its not printed on a single line

```

λ

Nearly every step that takes a traversal argument can also take a lambda.

It is recommended that users do not use lambdas as they are not subject to traversal strategy (i.e. compiler) optimization. However, they are useful when no provided step yields the desired computation.

"What is the most liked movie in each decade?"

```
gremlin> g.V().hasLabel('movie').
  where(inE('rated').count().is(gt(10))).
  group().
  by(((int)(it.value('year') / 10)) * 10).
  by().
  by(unfold().order().
    by(inE('rated').values('stars').mean(),decr).
    values('name').
    limit(1)).
  order(local).by(keyIncr).
  unfold() // so its not printed on a single line
==>1910=Daddy Long Legs
==>1920=General, The
==>1930=City Lights
==>1940=Third Man, The
==>1950=Seven Samurai (The Magnificent Seven)
==>1960=Sanjuro
==>1970=Godfather, The
==>1980=Raiders of the Lost Ark
==>1990=Shawshank Redemption, The
==>2000=Almost Famous
gremlin>
```

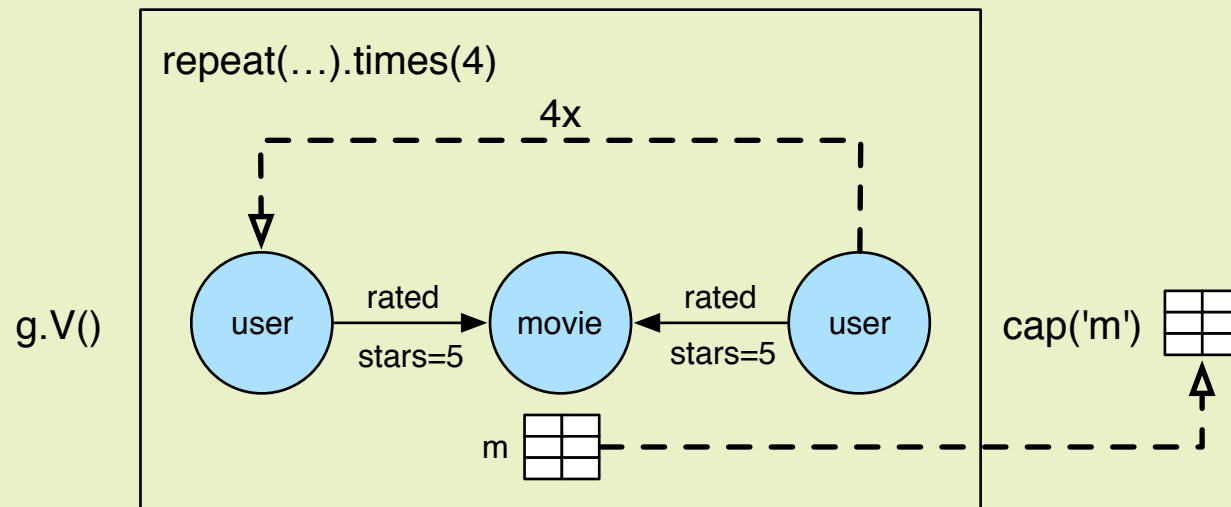
"What is the most liked movie in each decade?"

```
gremlin> graph = HadoopGraph.open('conf/hadoop/movie-lens.properties')  
==>hadoopgraph[gryoinputformat->gryooutputformat]  
gremlin> g = graph.traversal(computer(SparkGraphComputer))  
==>graphtraversalsource  
      [hadoopgraph[gryoinputformat->gryooutputformat], sparkgraphcomputer]  
gremlin>
```



"Which movies are most central in the implicit 5-stars graph?"

```
gremlin> graph = HadoopGraph.open('conf/hadoop/movie-lens.properties')
==>hadoopgraph[gryoinputformat->gryooutputformat]
gremlin> g = graph.traversal(computer(SparkGraphComputer))
==>graphtraversalsource
      [hadoopgraph[gryoinputformat->gryooutputformat], sparkgraphcomputer]
gremlin> g.V().repeat(outE('rated').has('stars', 5).inV().
      groupCount('m').by('name').
      inE('rated').has('stars', 5).outV()).
      times(4).cap('m')
```



"Which movies are most central in the implicit 5-stars graph?"

```
gremlin> graph = HadoopGraph.open('conf/hadoop/movie-lens.properties')
==>hadoopgraph[gryoinputformat->gryooutputformat]
gremlin> g = graph.traversal(computer(SparkGraphComputer))
==>graphtraversalsource
      [hadoopgraph[gryoinputformat->gryooutputformat], sparkgraphcomputer]
gremlin> g.V().repeat(outE('rated').has('stars', 5).inV().
      groupCount('m').by('name').
      inE('rated').has('stars', 5).outV()).
      times(4).cap('m')
==>Fantasia 2000=2676505178171564
==>Pale Rider=1369969000295362
==>Crucible, The=401712993698149
==>About Adam=37981148456999
==>Akira=3659939409345918
...
gremlin>
```



"Which movies are most central in the implicit 5-stars graph?"

```

gremlin> graph = HadoopGraph.open('conf/hadoop/movie-lens.properties')
==>hadoopgraph[gryoinputformat->gryooutputformat]
gremlin> g = graph.traversal(computer(SparkGraphComputer))
==>graphtraversalsource
      [hadoopgraph[gryoinputformat->gryooutputformat], sparkgraphcomputer]
gremlin> g.V().repeat(outE('rated').has('stars', 5).inV().
      groupCount('m').by('name').
      inE('rated').has('stars', 5).outV()).
      times(4).cap('m')
==>Fantasia 2000=2676505178171564
==>Pale Rider=1369969000295362
==>Crucible, The=401712993698149
==>About Adam=37981148456999
==>Akira=3659939409345918
...
gremlin> hdfs.ls('output/m')
==>rw-r--r-- daniel supergroup 0 _SUCCESS
==>rw-r--r-- daniel supergroup 245314 part-r-00000
gremlin> hdfs.head('output/m', ObjectWritable).sort {-it.value}.take(10)
==>Star Wars: Episode IV - A New Hope 35405394353105332
==>American Beauty 31943228282020585
==>Raiders of the Lost Ark 31224779793238499
==>Star Wars: Episode V - The Empire Strikes Back 30434677119726223
==>Godfather, The 30258518523013057
==>Shawshank Redemption, The 28297717387901031
==>Schindler's List 27539336654199309
==>Silence of the Lambs, The 26736276376806173
==> Fargo 26531050311325270
==>Matrix, The 26395118239203191

```



"Which movies are most central in the implicit 5-stars graph?"

```
gremlin> :plugin use tinkerpop.gephi
==>tinkerpop.gephi activated
gremlin> :remote connect tinkerpop.gephi
==>Connection to Gephi - http://localhost:8080/workspace0 with stepDelay:1000,
startRGBColor:[0.0, 1.0, 0.5], colorToFade:g, colorFadeRate:0.7, startSize:
20.0,sizeDecrementRate:0.33
gremlin>
```



```

gremlin> :plugin use tinkerpop.gephi
==>tinkerpop.gephi activated
gremlin> :remote connect tinkerpop.gephi
==>Connection to Gephi - http://localhost:8080/workspace0 with stepDelay:1000,
startRGBColor:[0.0, 1.0, 0.5], colorToFade:g, colorFadeRate:0.7, startSize:
20.0,sizeDecrementRate:0.33
gremlin> :> g.V().hasLabel('user').
    order().
        by(outE('rated').count(), decr).limit(10).as('a').
    local(outE('rated').order().
        by('stars', decr). // first by stars
        by(inV().inE('rated').count(), decr). // then by ratings
        limit(10)).
    subgraph('sg').inV().outE('category').
    subgraph('sg').select('a').outE('occupation').
    subgraph('sg').cap('sg').next()
==>tinkergraph[vertices:82 edges:233]
gremlin>

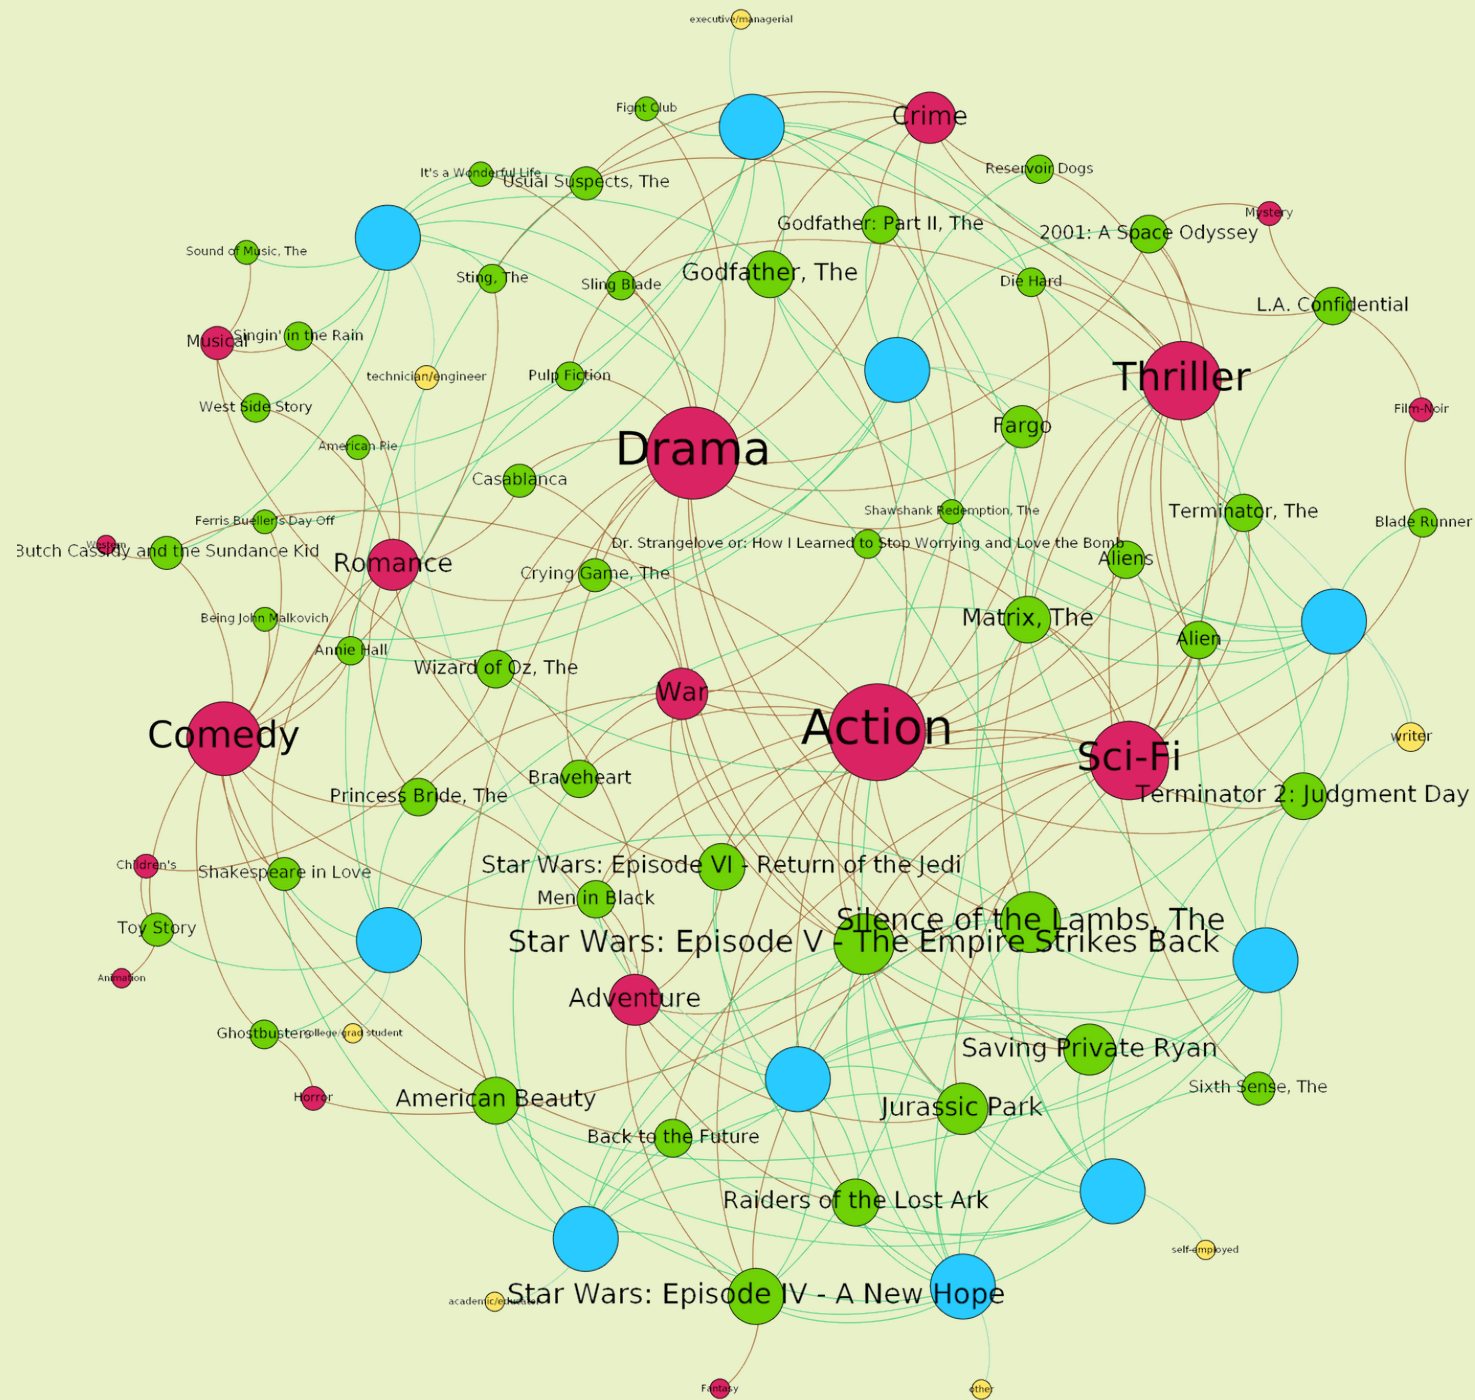
```


"movie buffs"



"Which users rated the most movies?"

For each user, display their 10 favorite movies, the categories of those movies, and their occupation.





Thanks for listening...