

THE PATH FORWARD

TITAN 1.0 + TINKERPOP3



MARKO A. RODRIGUEZ



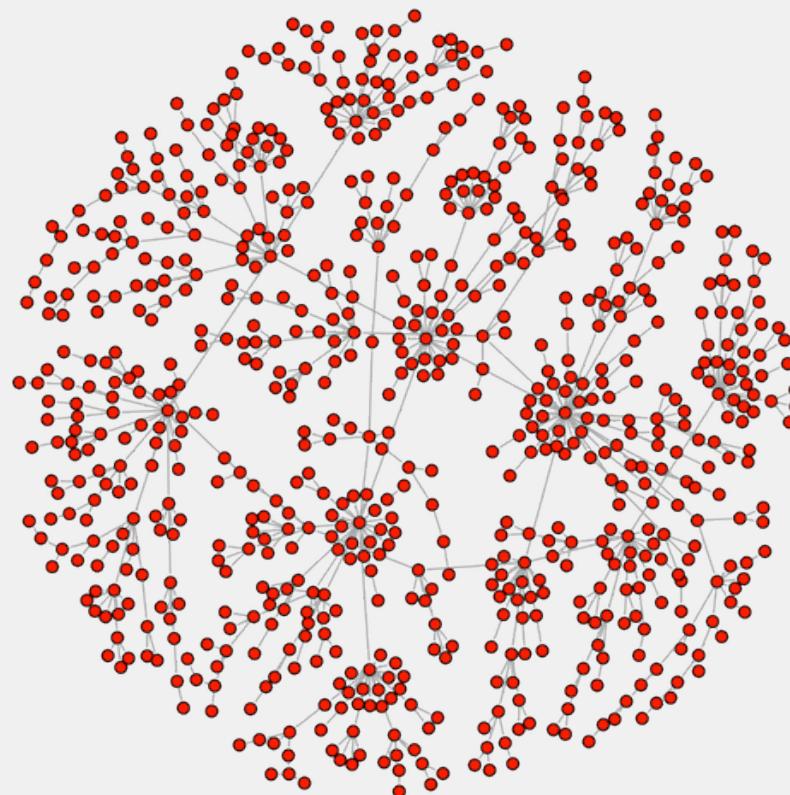
[HTTP://THINKAURELIUS.COM](http://thinkaurelius.com)

[HTTP://TINKERPOP.COM](http://tinkerpop.com)

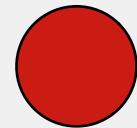


AURELIUS

PART 1: INTRODUCTION TO GRAPHS

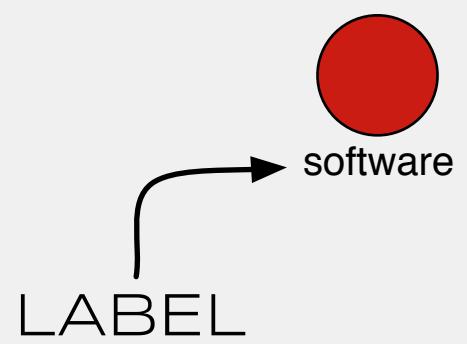


VERTEX

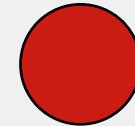




software

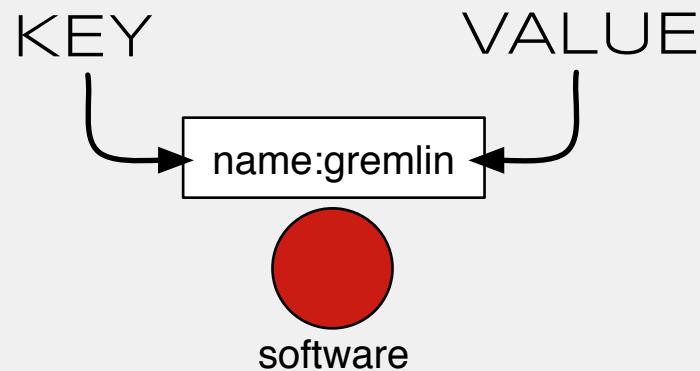


name:gremlin

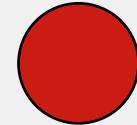


software

PROPERTY

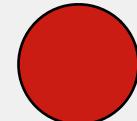


name:gremlin
use:queries



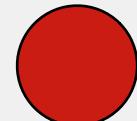
software

name:gremlin
use:queries

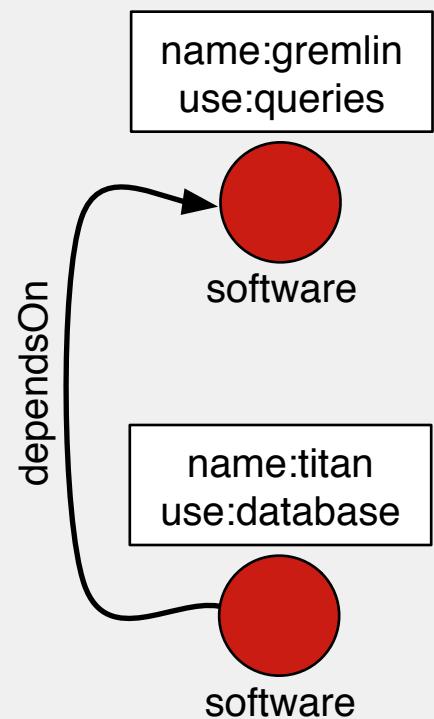


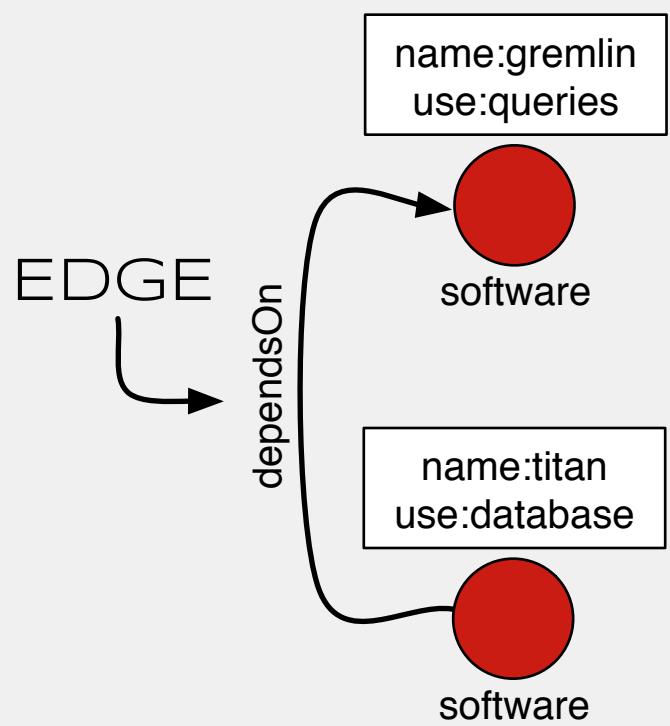
software

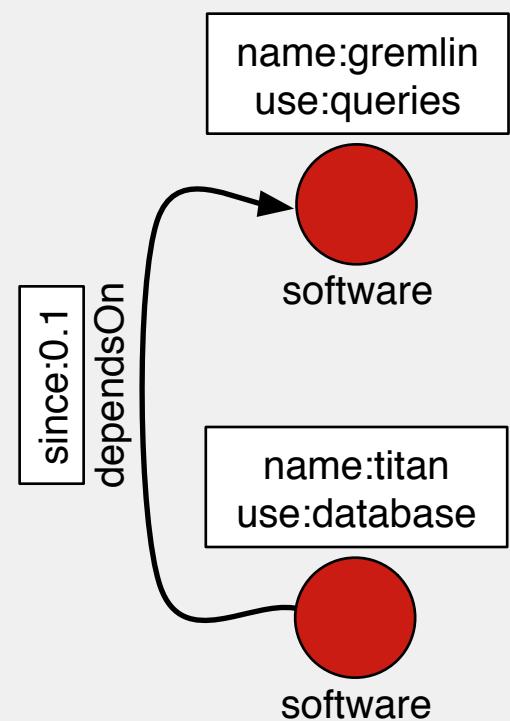
name:titan
use:database



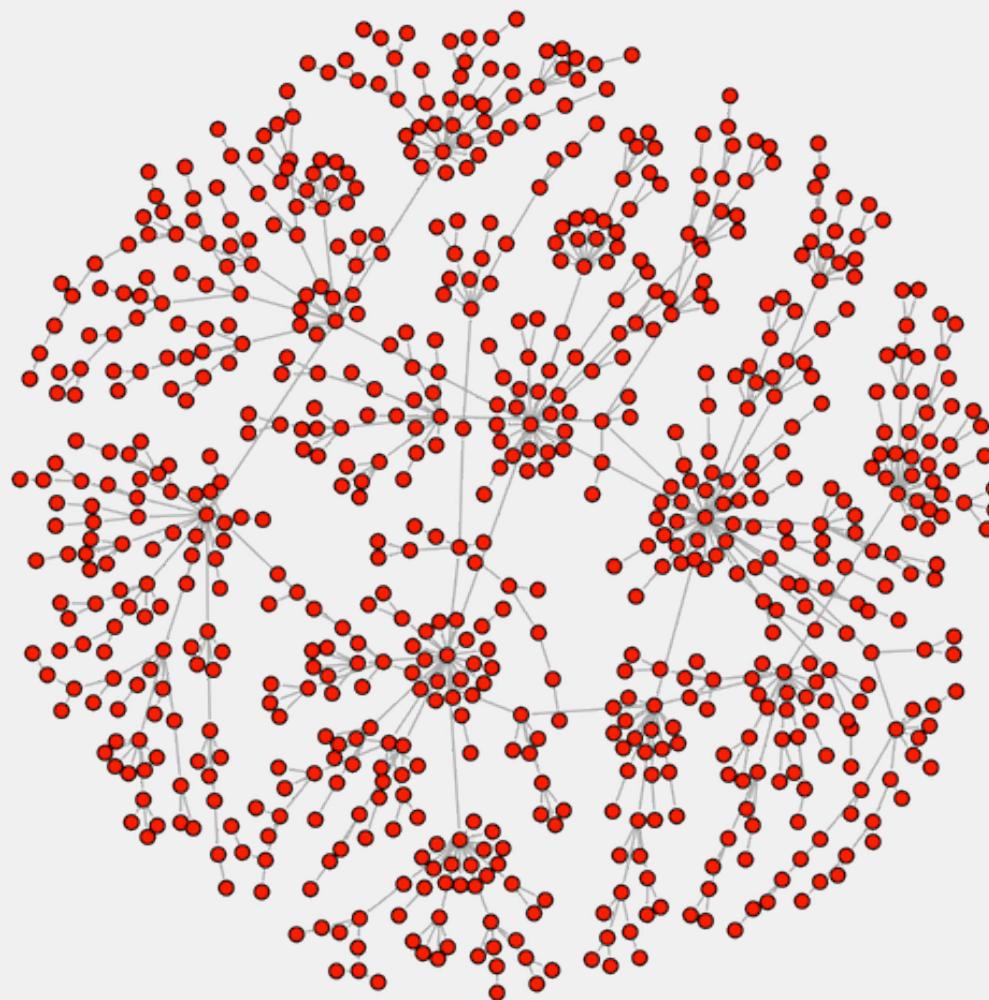
software







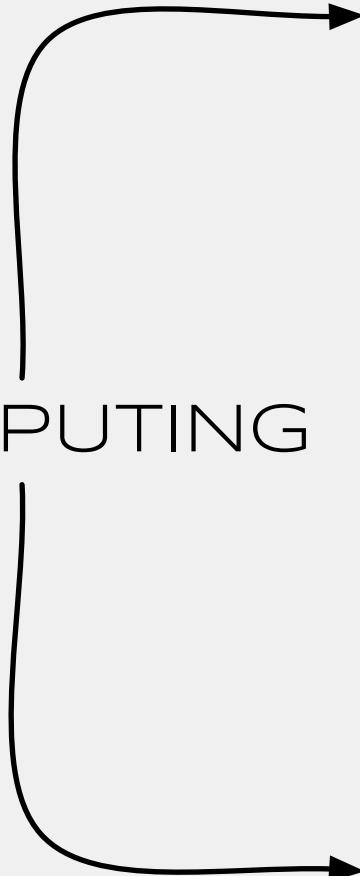
GRAPH



COMPUTING

PROCESS

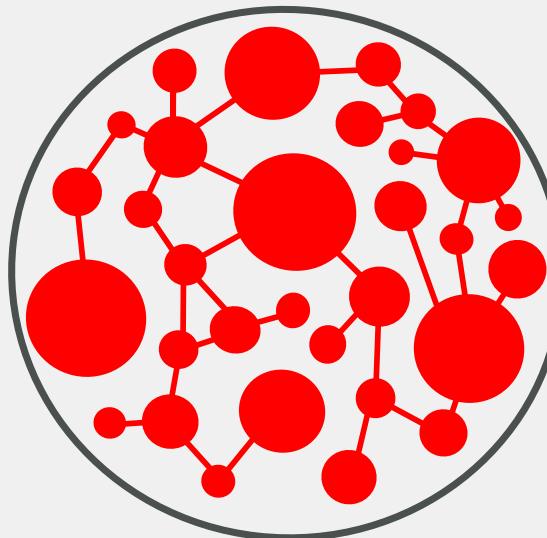
STRUCTURE



COMPUTING

PROCESS

STRUCTURE



GRAPH

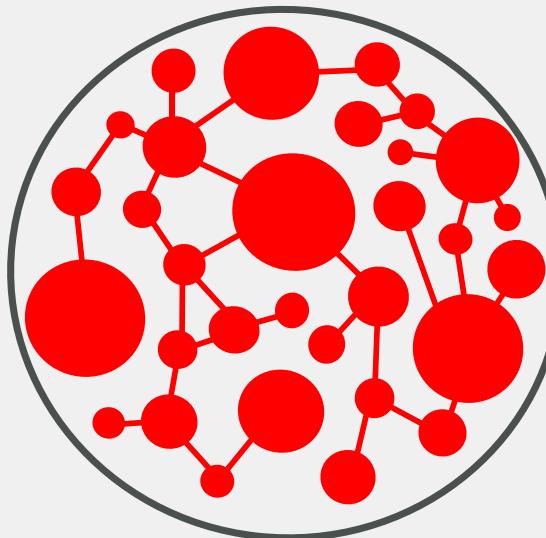
TRAVERSAL



COMPUTING

PROCESS

STRUCTURE



GRAPH

GRAPH-BASED
COMPUTING

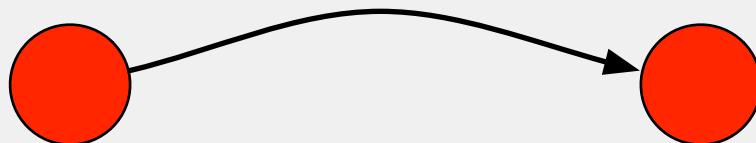
TRAVERSAL



WHY GRAPH-BASED COMPUTING?

WHY GRAPH-BASED COMPUTING?

INTUITIVE MODELING

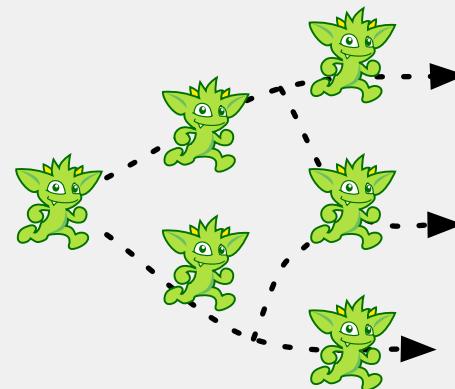


WHY GRAPH-BASED COMPUTING?

INTUITIVE MODELING



EXPRESSIVE QUERYING

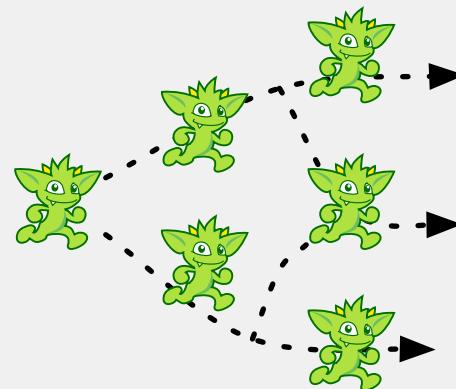


WHY GRAPH-BASED COMPUTING?

INTUITIVE MODELING



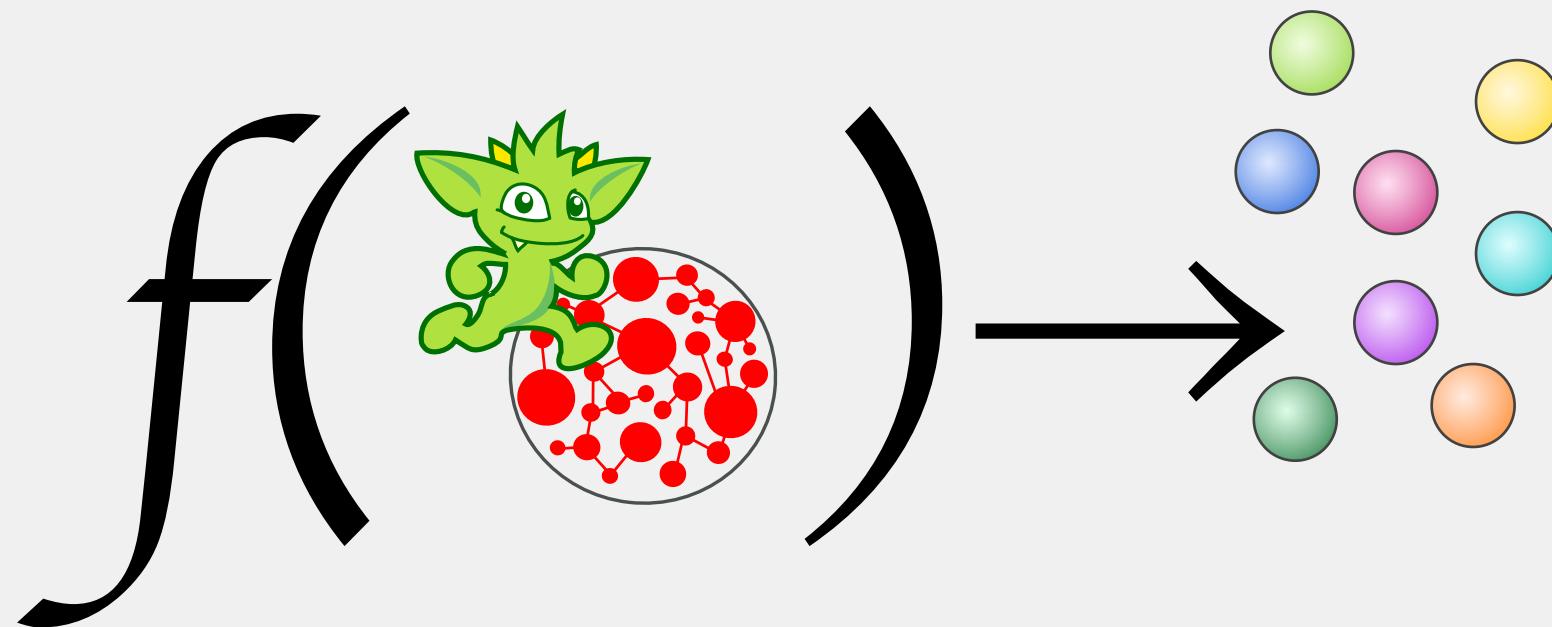
EXPRESSIVE QUERYING



NUMEROUS ANALYSES

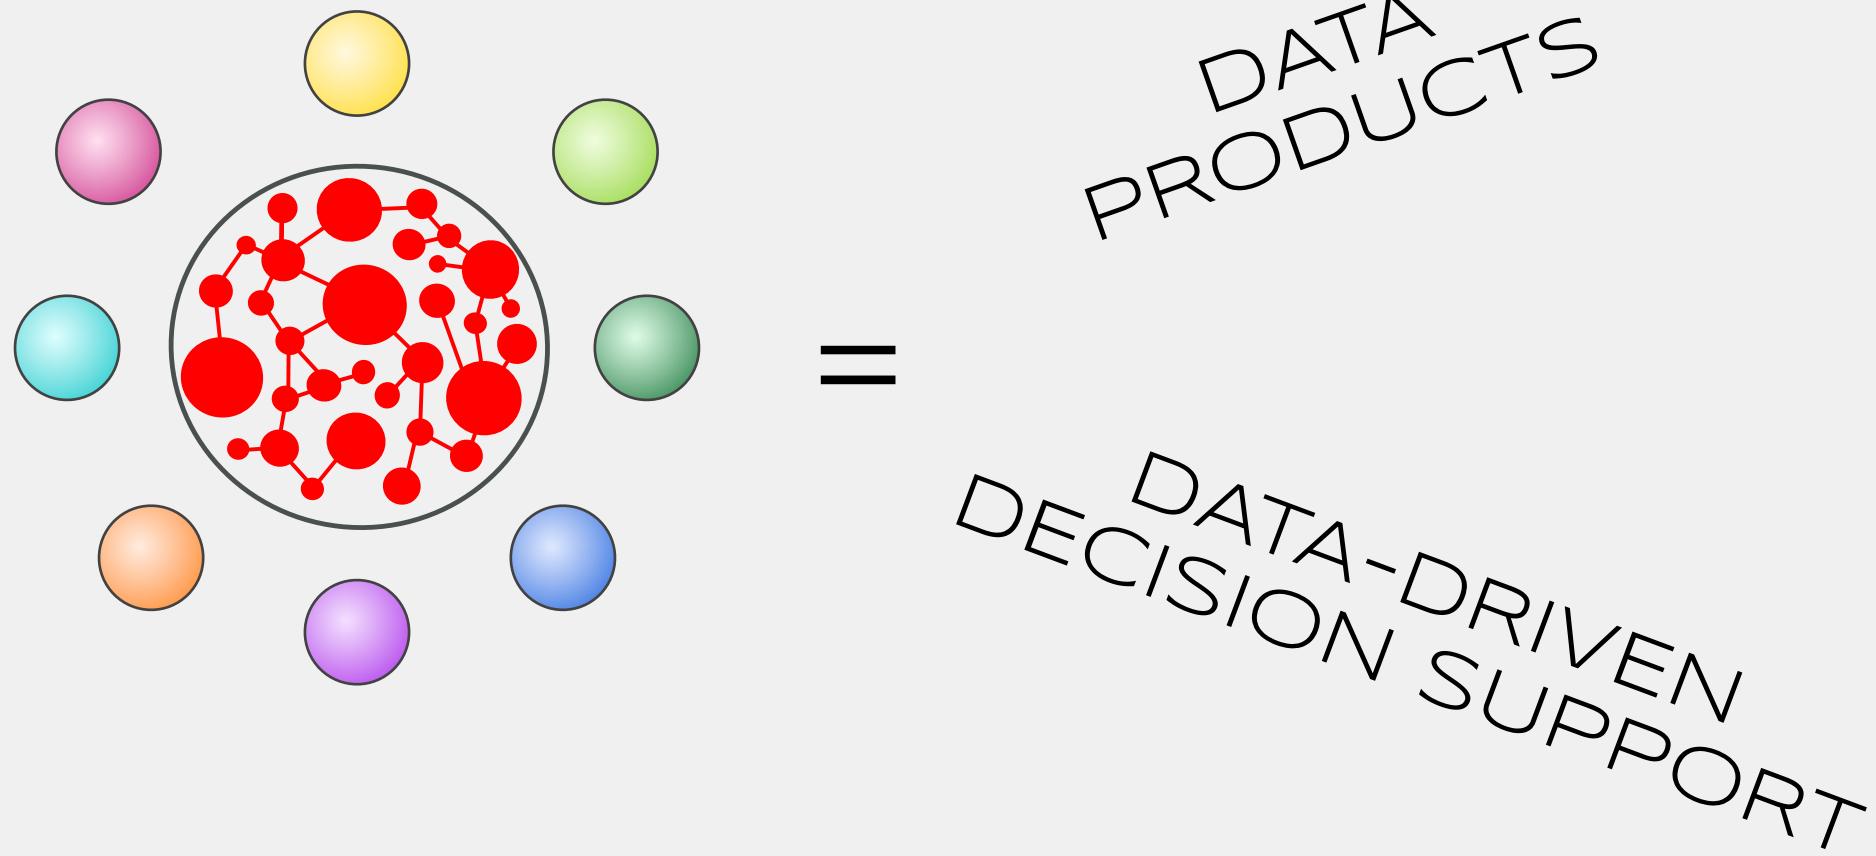
- Inference
- Scoring
- Centrality
- Mixing Patterns
- Motifs
- Path Expressions
- Geodesics
- Ranking

ANALYSES ARE THE EPIPHENOMENA OF TRAVERSAL



WHAT IS THE SIGNIFICANCE OF
GRAPH ANALYSIS?

ANALYSES YIELD INSIGHTS ABOUT THE MODEL



RECOMMENDATION



● People you may know.

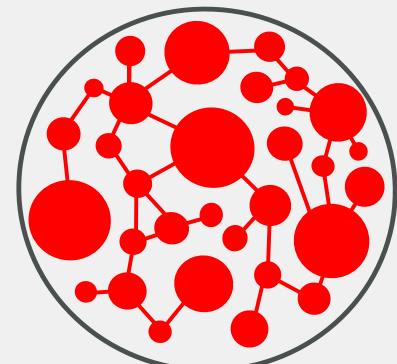
SOCIAL GRAPH

● Products you might like.

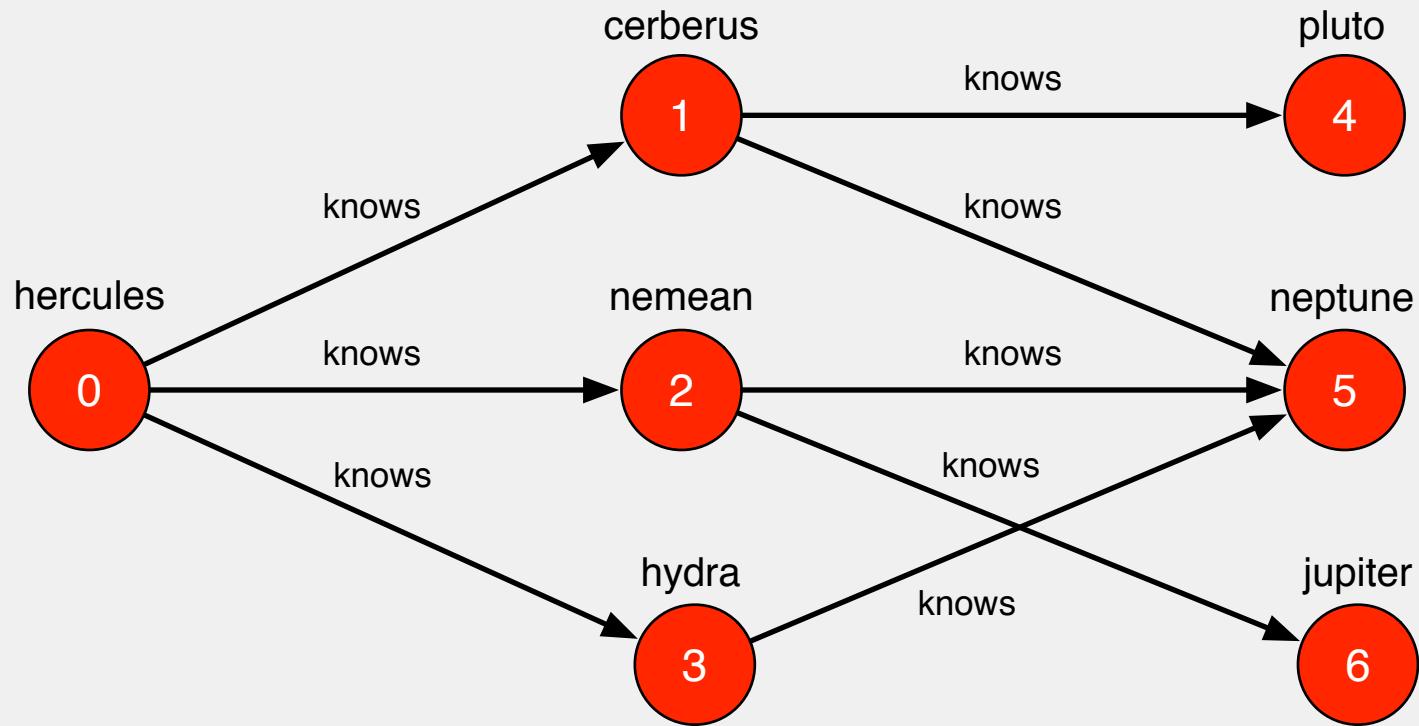
RATINGS GRAPH

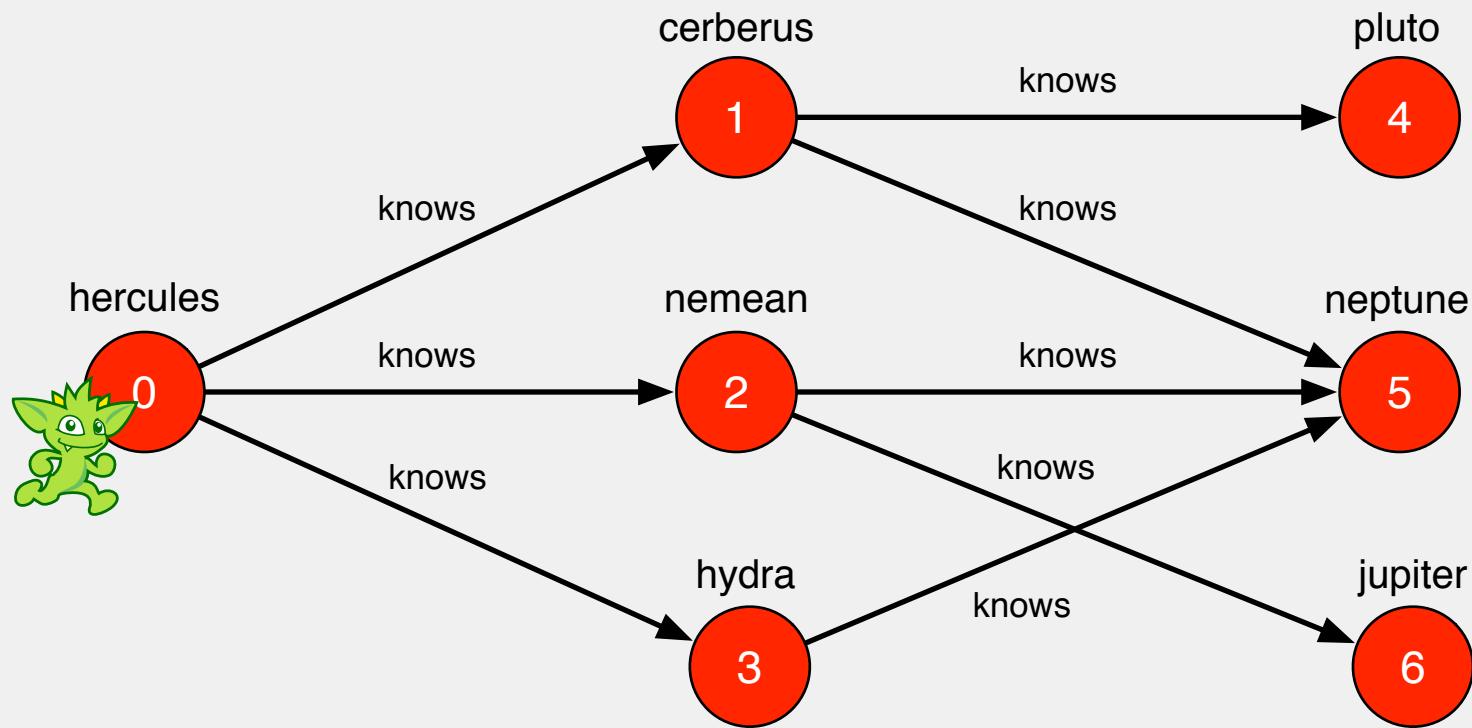
● Movies you should watch and
the friends you should watch them with.

**SOCIAL+RATINGS
GRAPH**

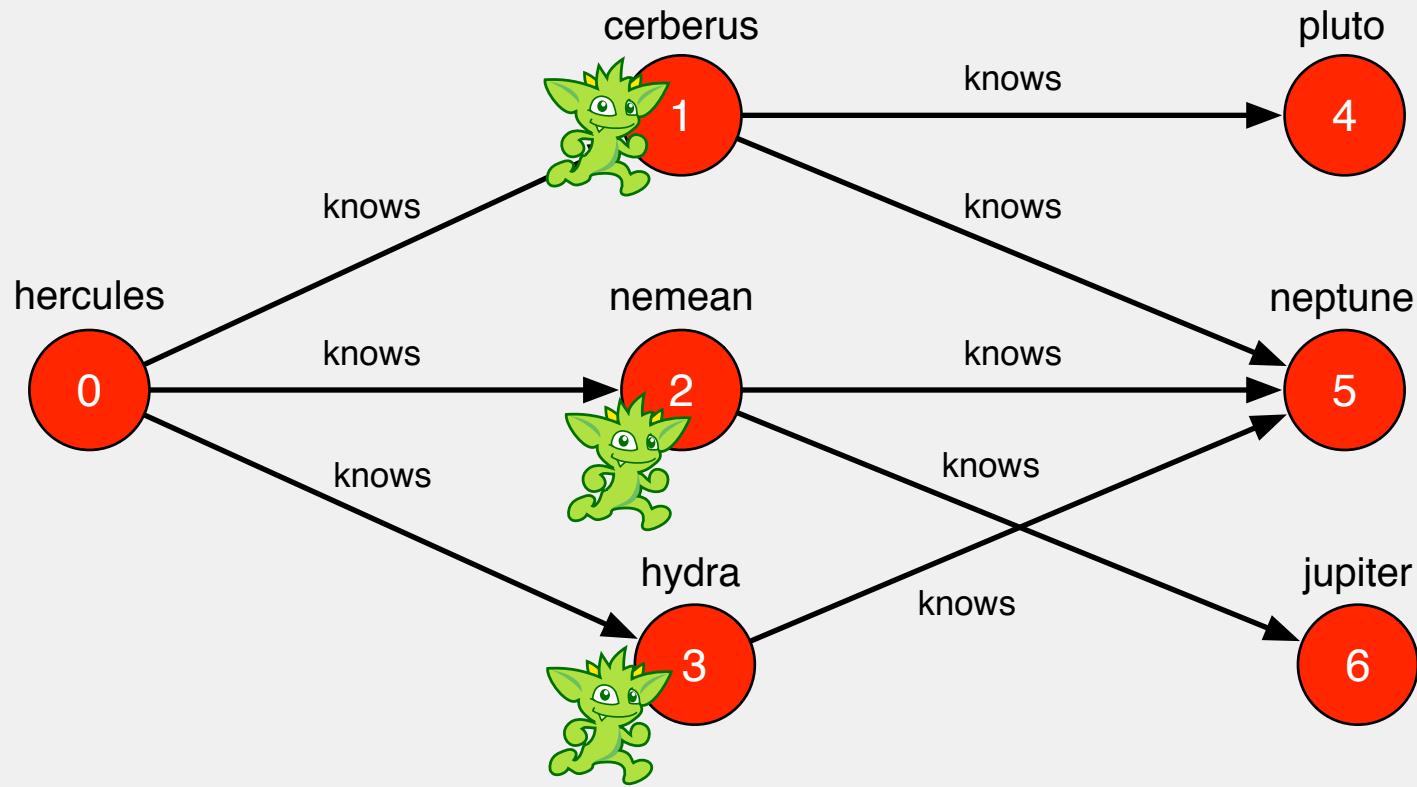


WHO ELSE MIGHT HERCULES KNOW?

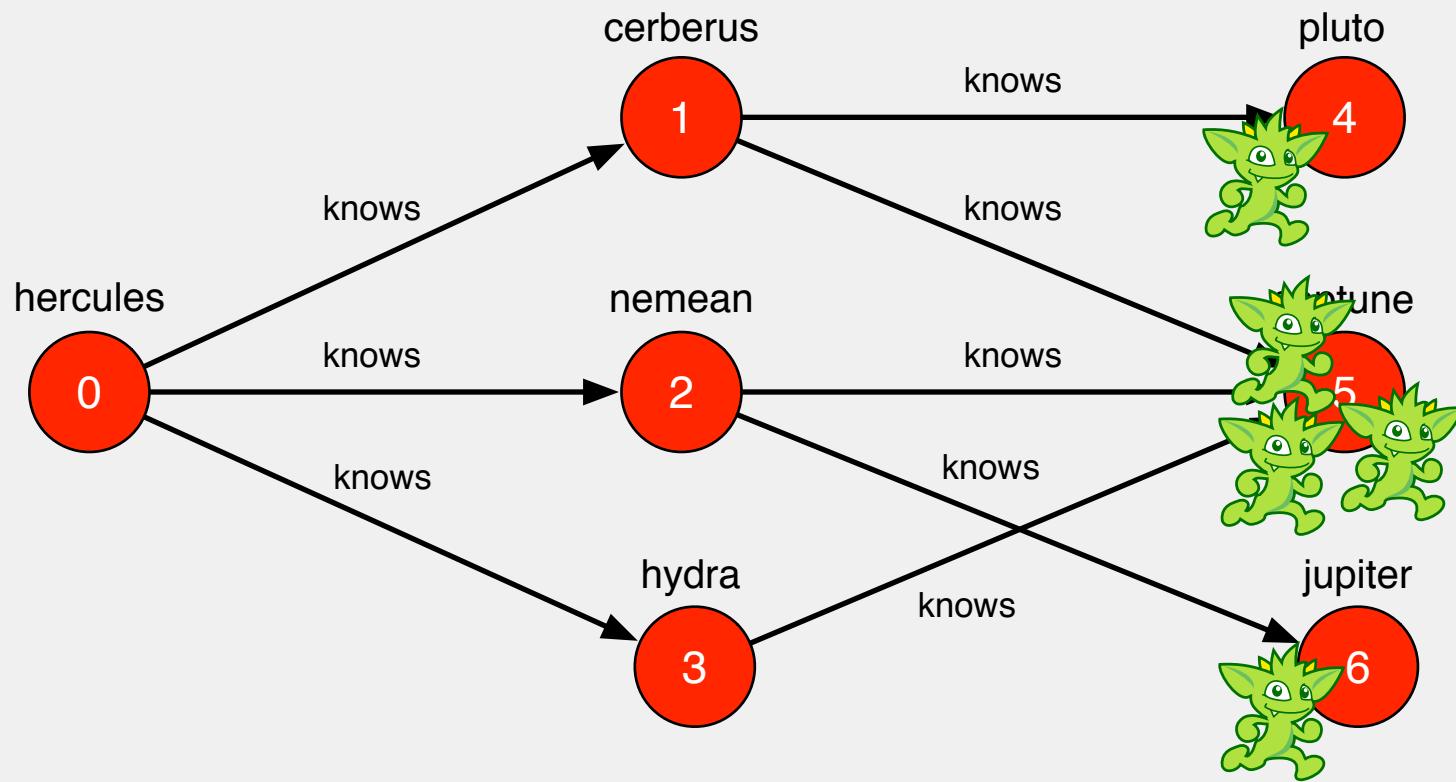




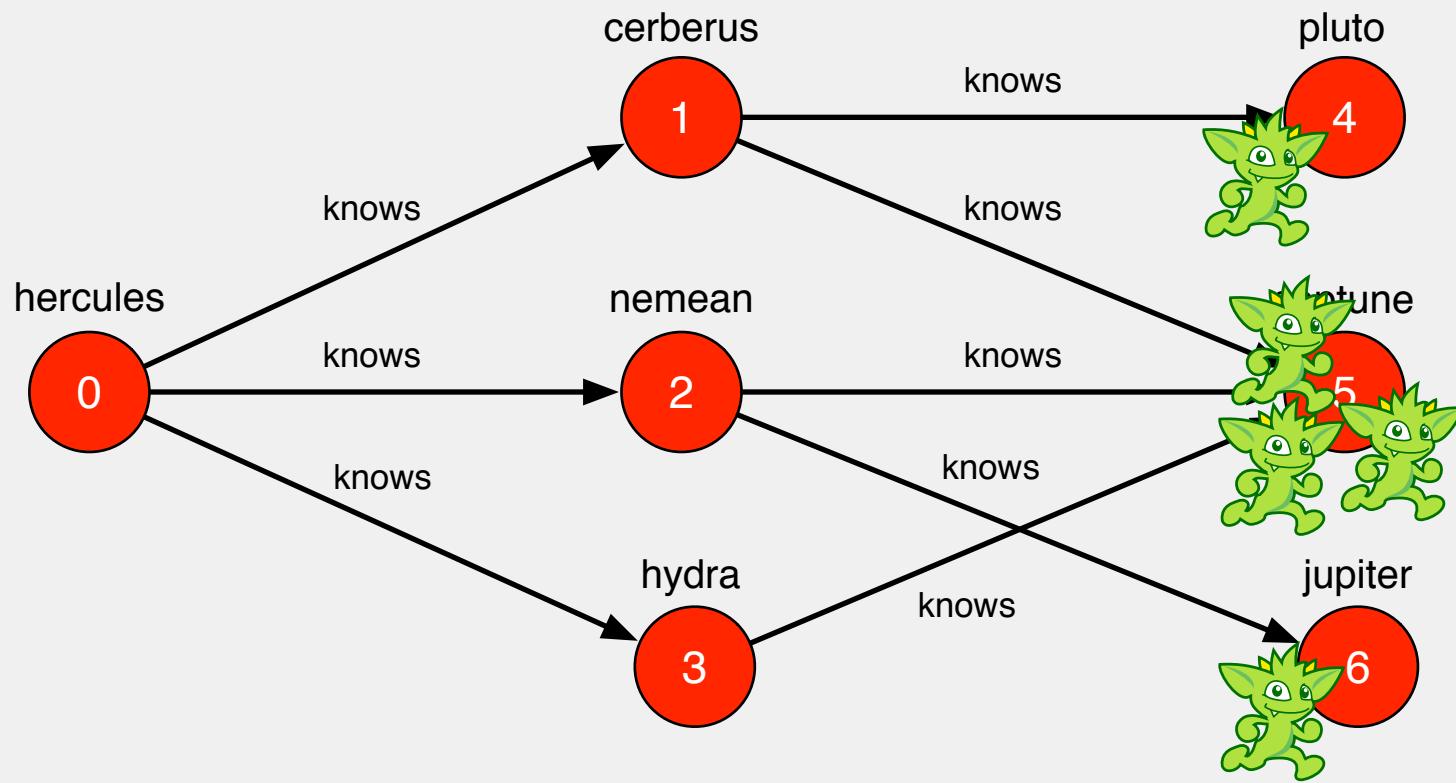
```
gremlin> hercules  
==>v[ 0 ]
```



```
gremlin> hercules.out('knows')
==>v[1]
==>v[2]
==>v[3]
```

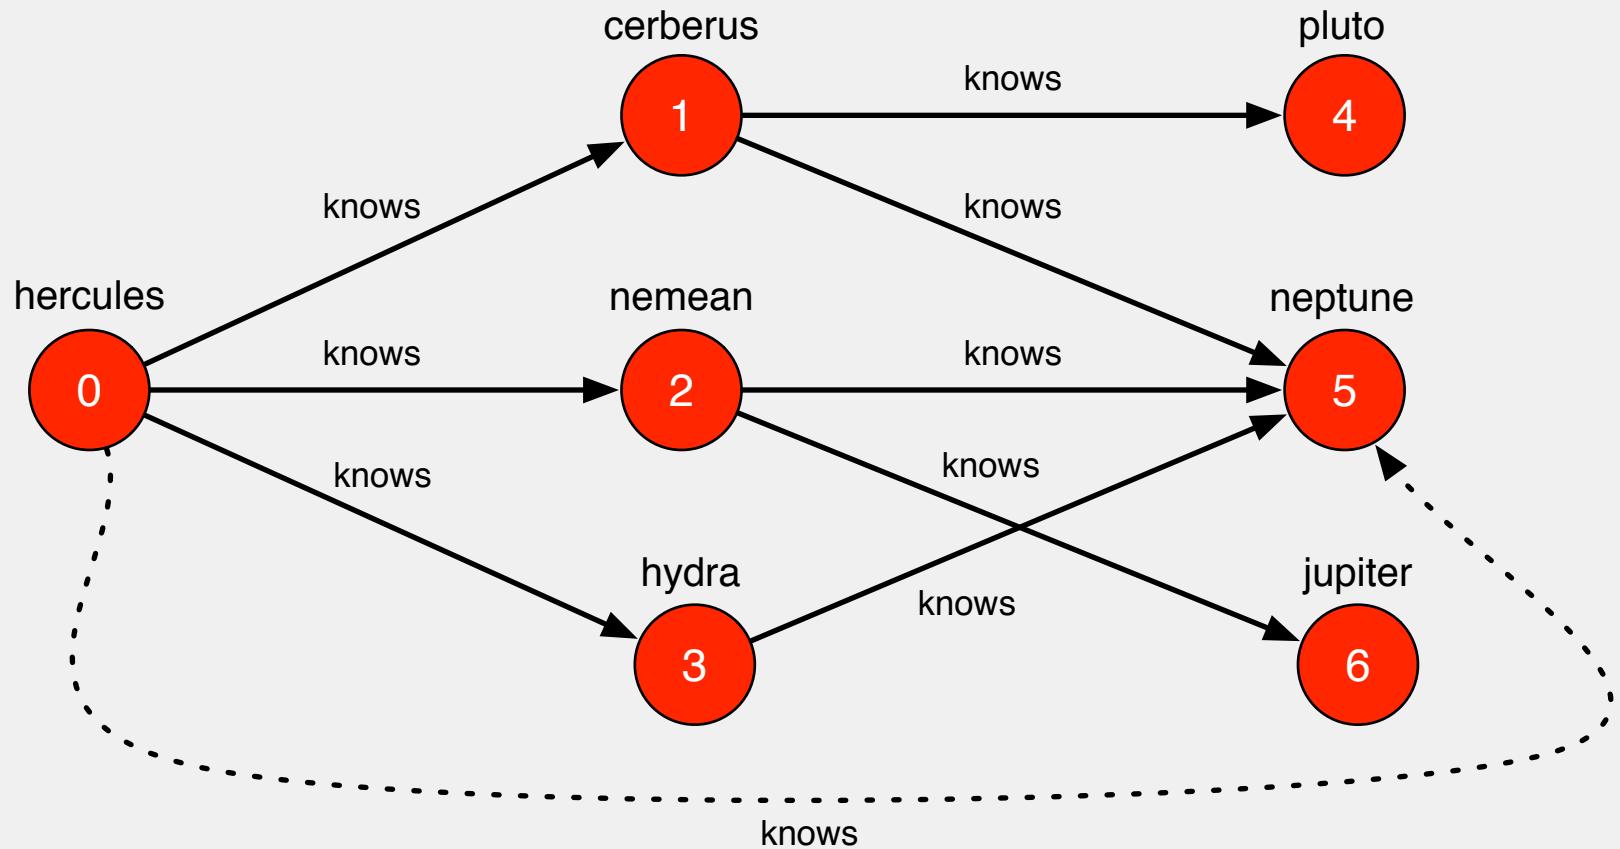


```
gremlin> hercules.out('knows').out('knows')
==>v[ 4 ]
==>v[ 5 ]
==>v[ 5 ]
==>v[ 6 ]
==>v[ 5 ]
```

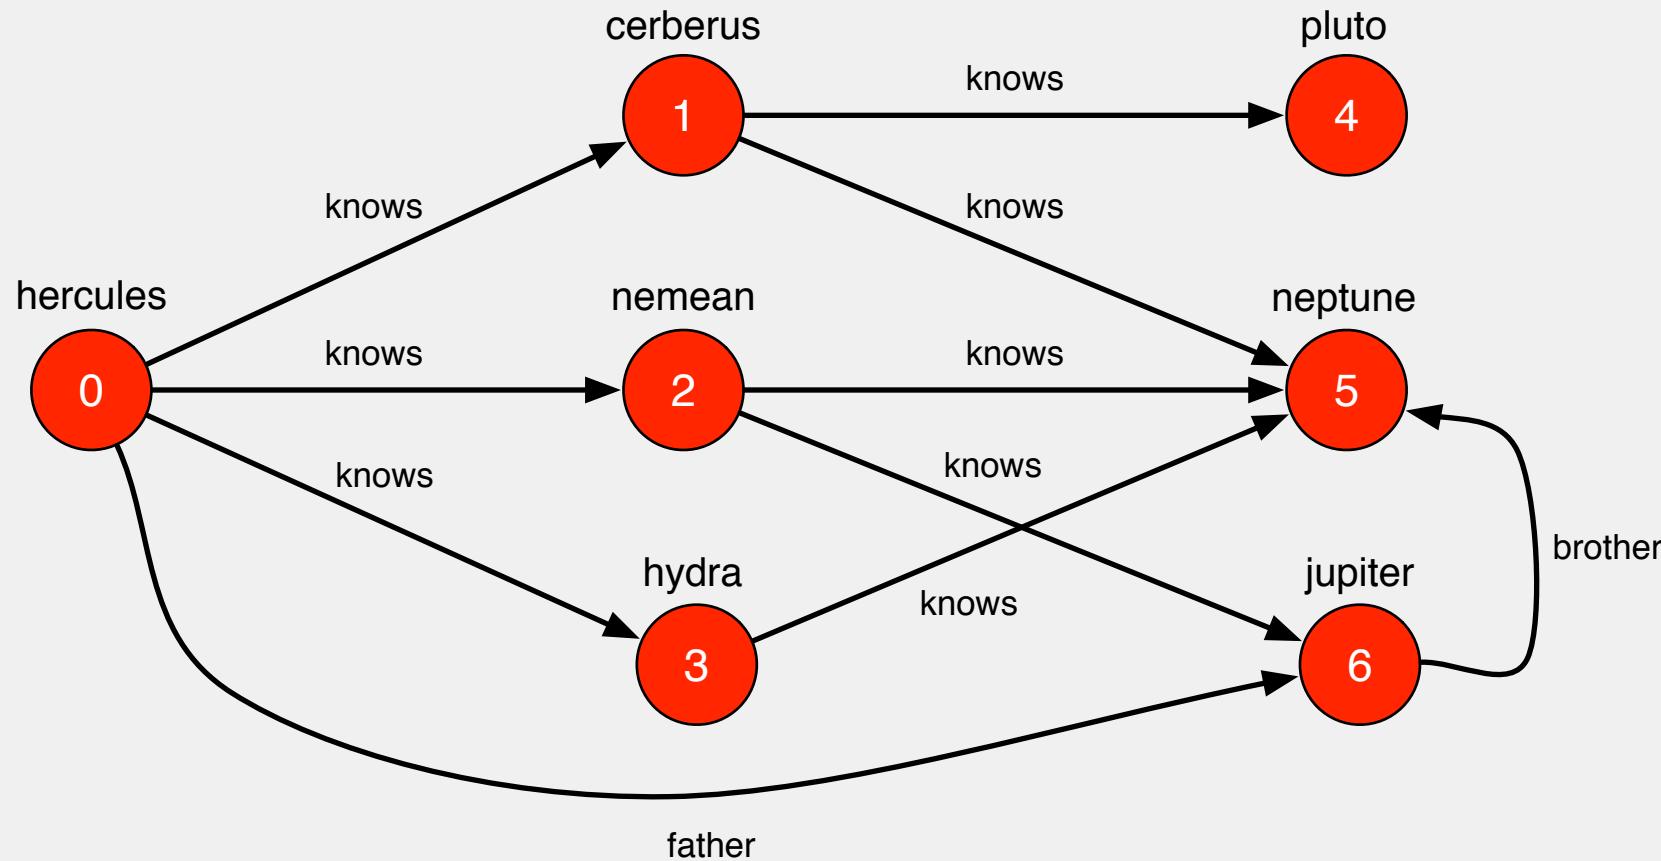


```
gremlin> hercules.out('knows').out('knows').groupCount()
==>v[4]=1
==>v[5]=3
==>v[6]=1
```

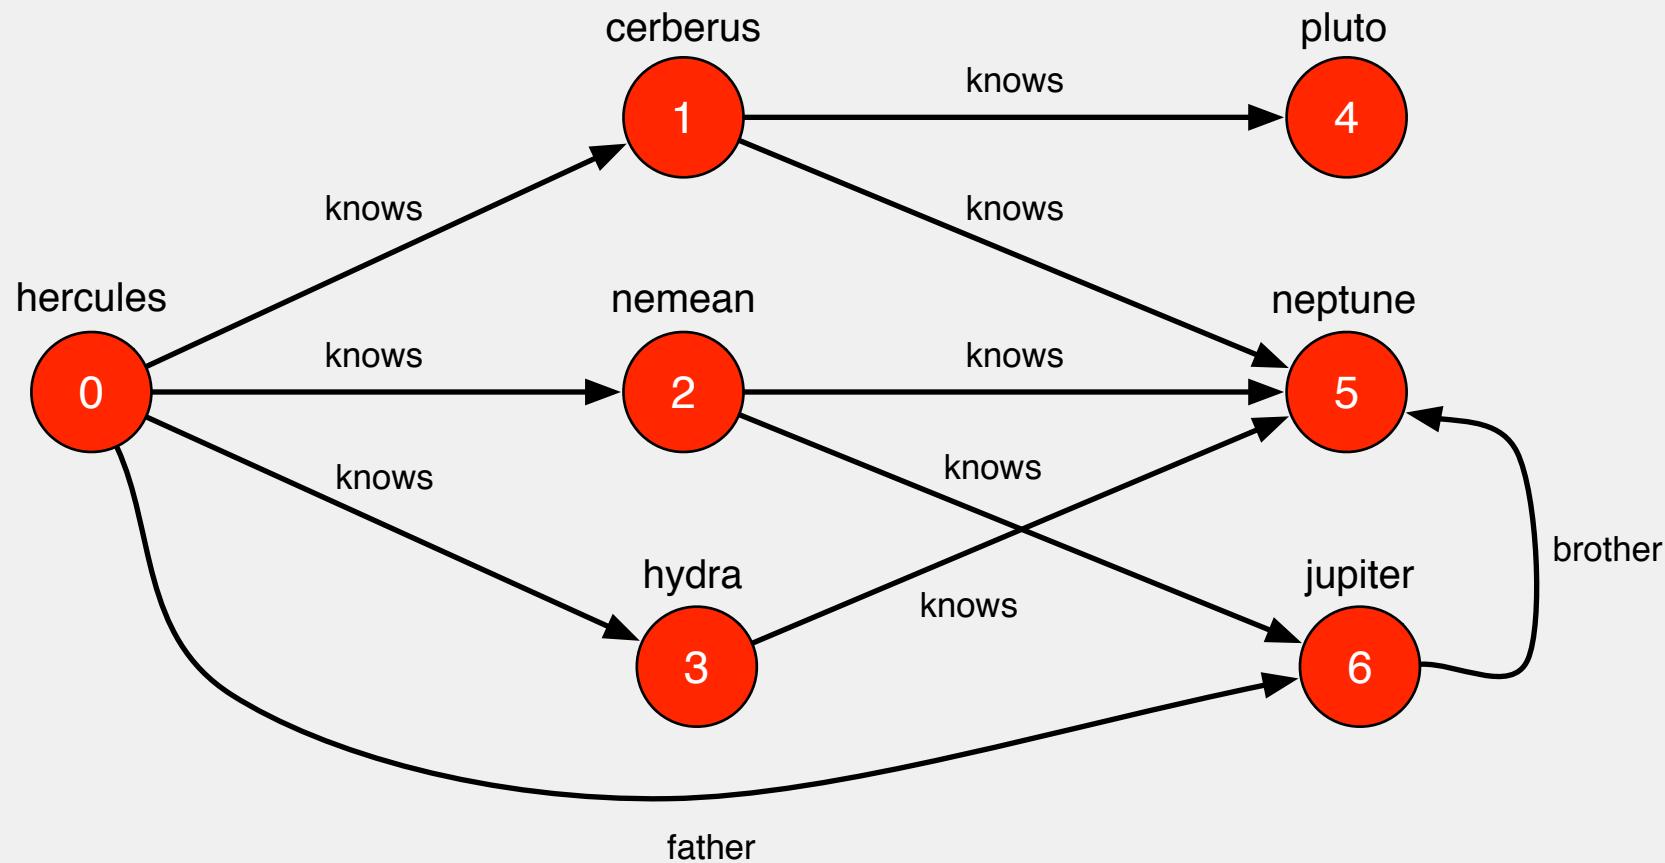
HERCULES PROBABLY KNOWS NEPTUNE

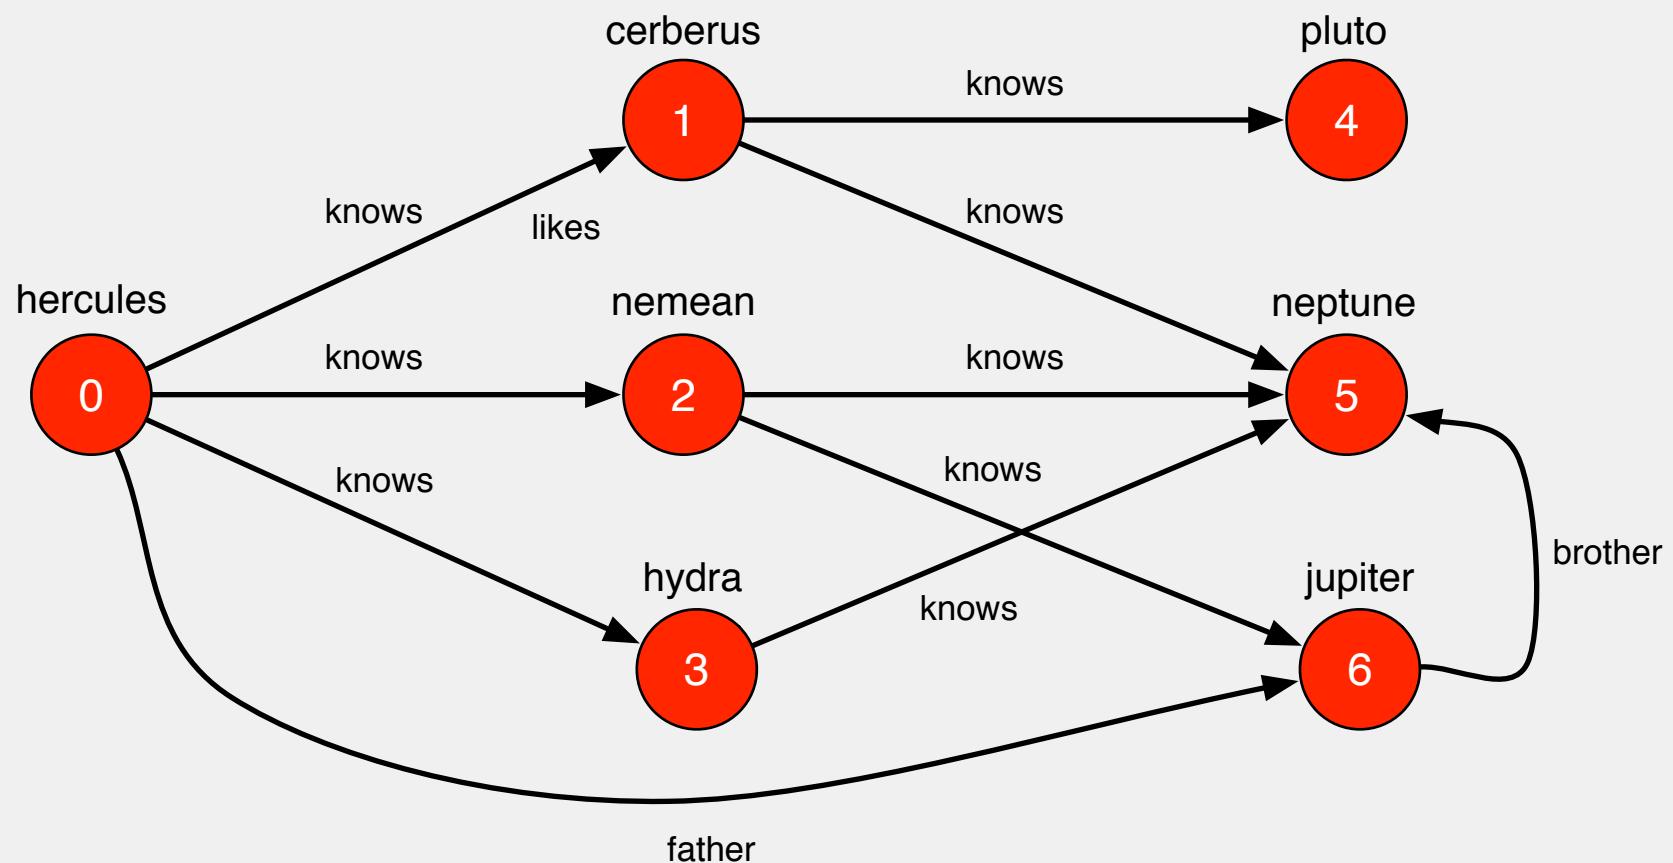


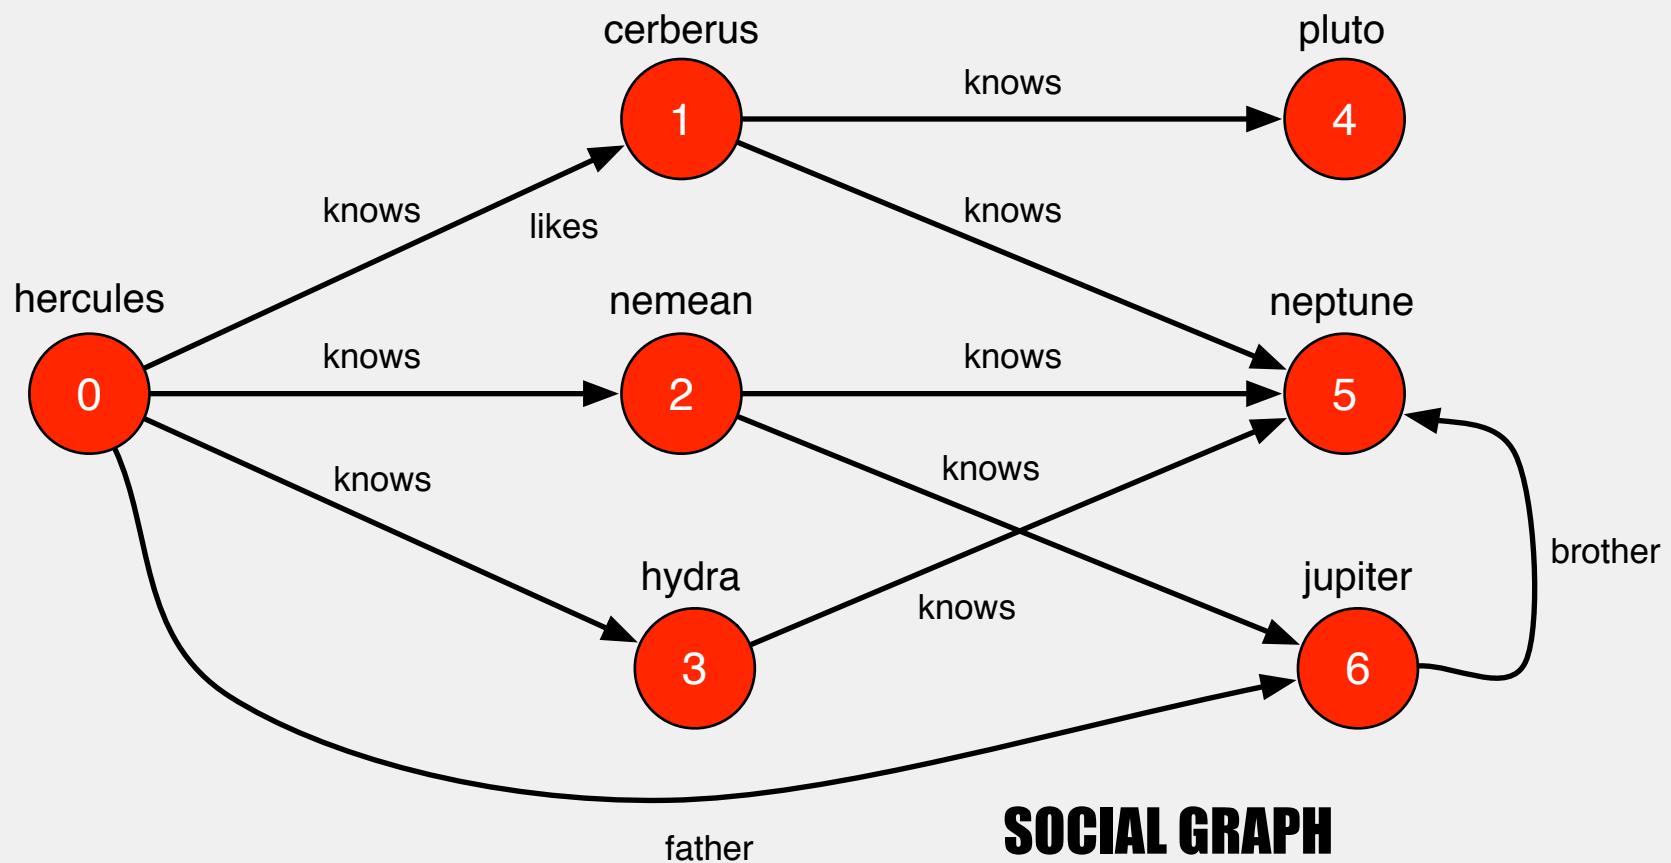
HERCULES PROBABLY KNOWS NEPTUNE

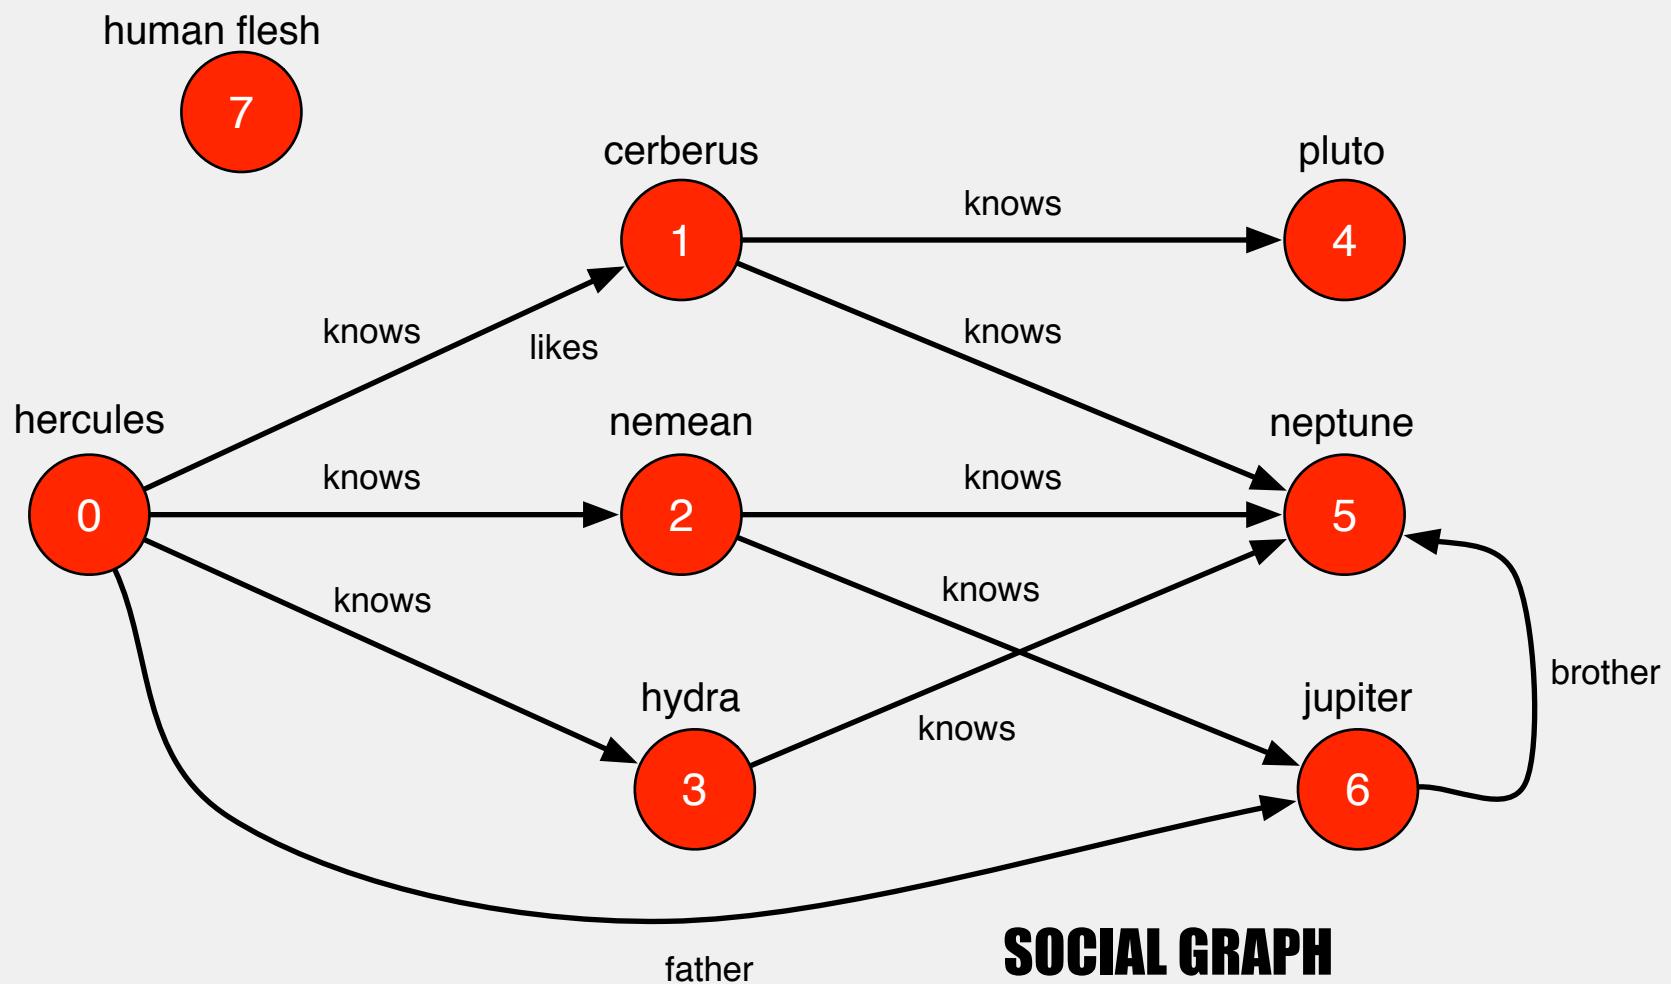


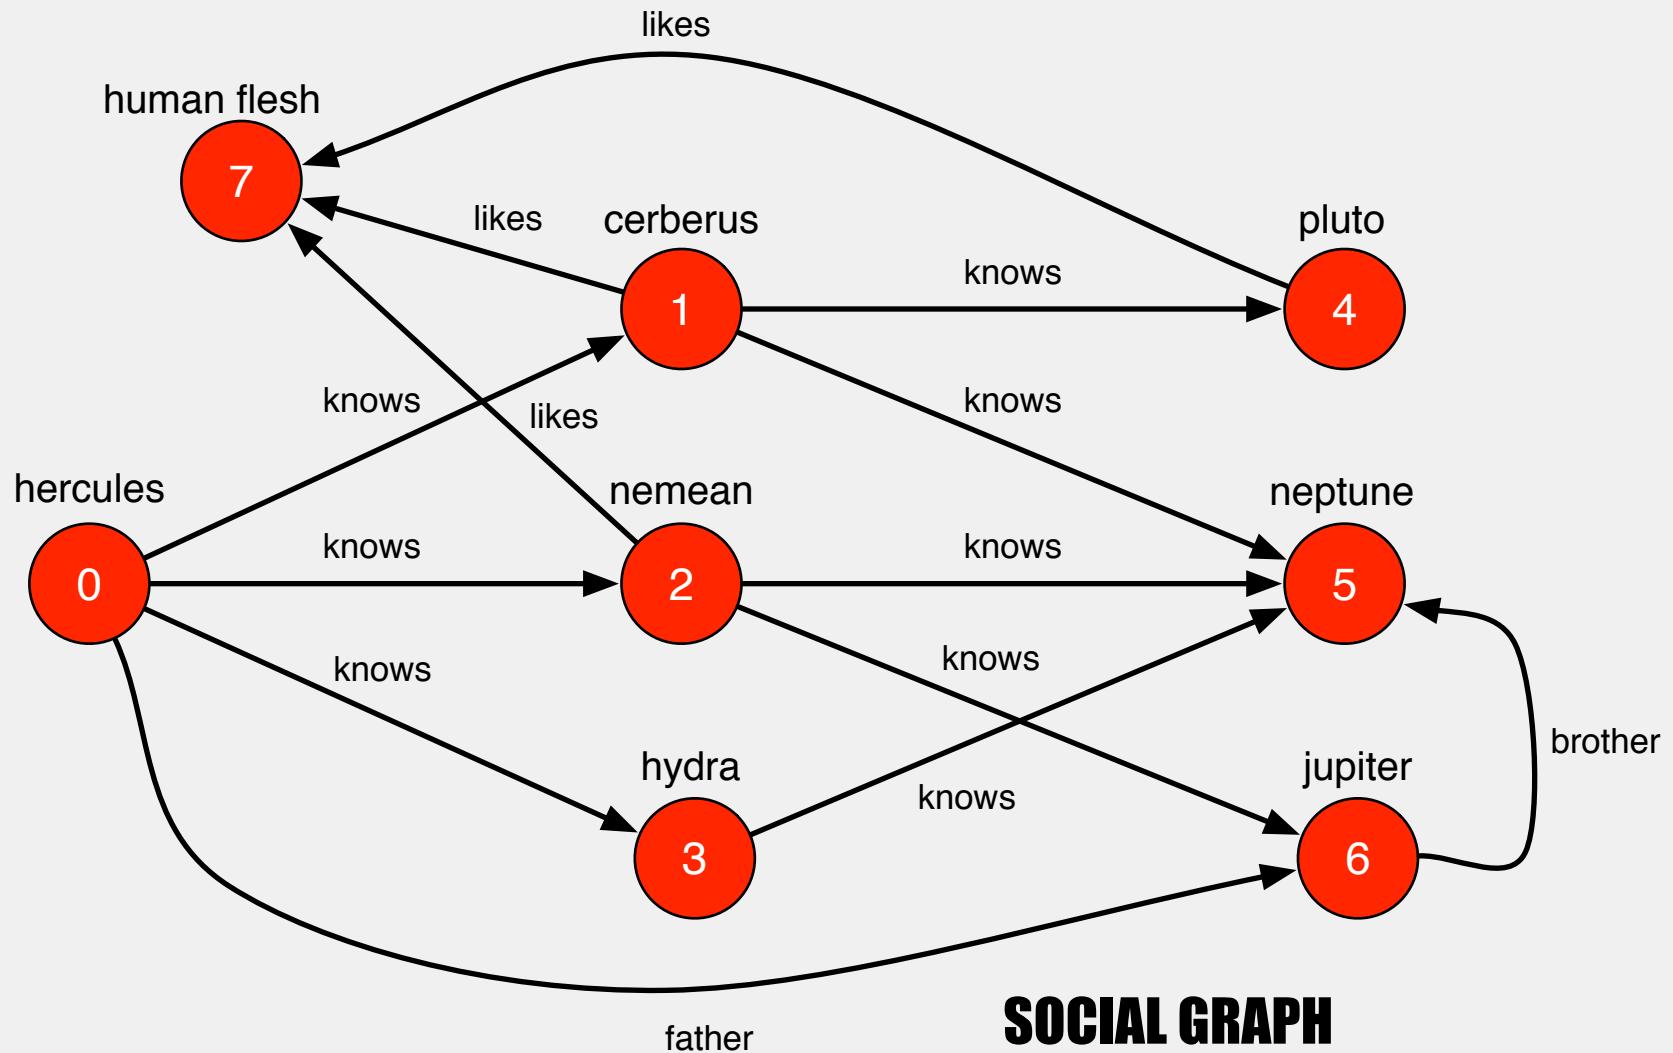
...PROBABLY MORE SO WHEN OTHER TYPES OF EDGES ARE ANALYZED

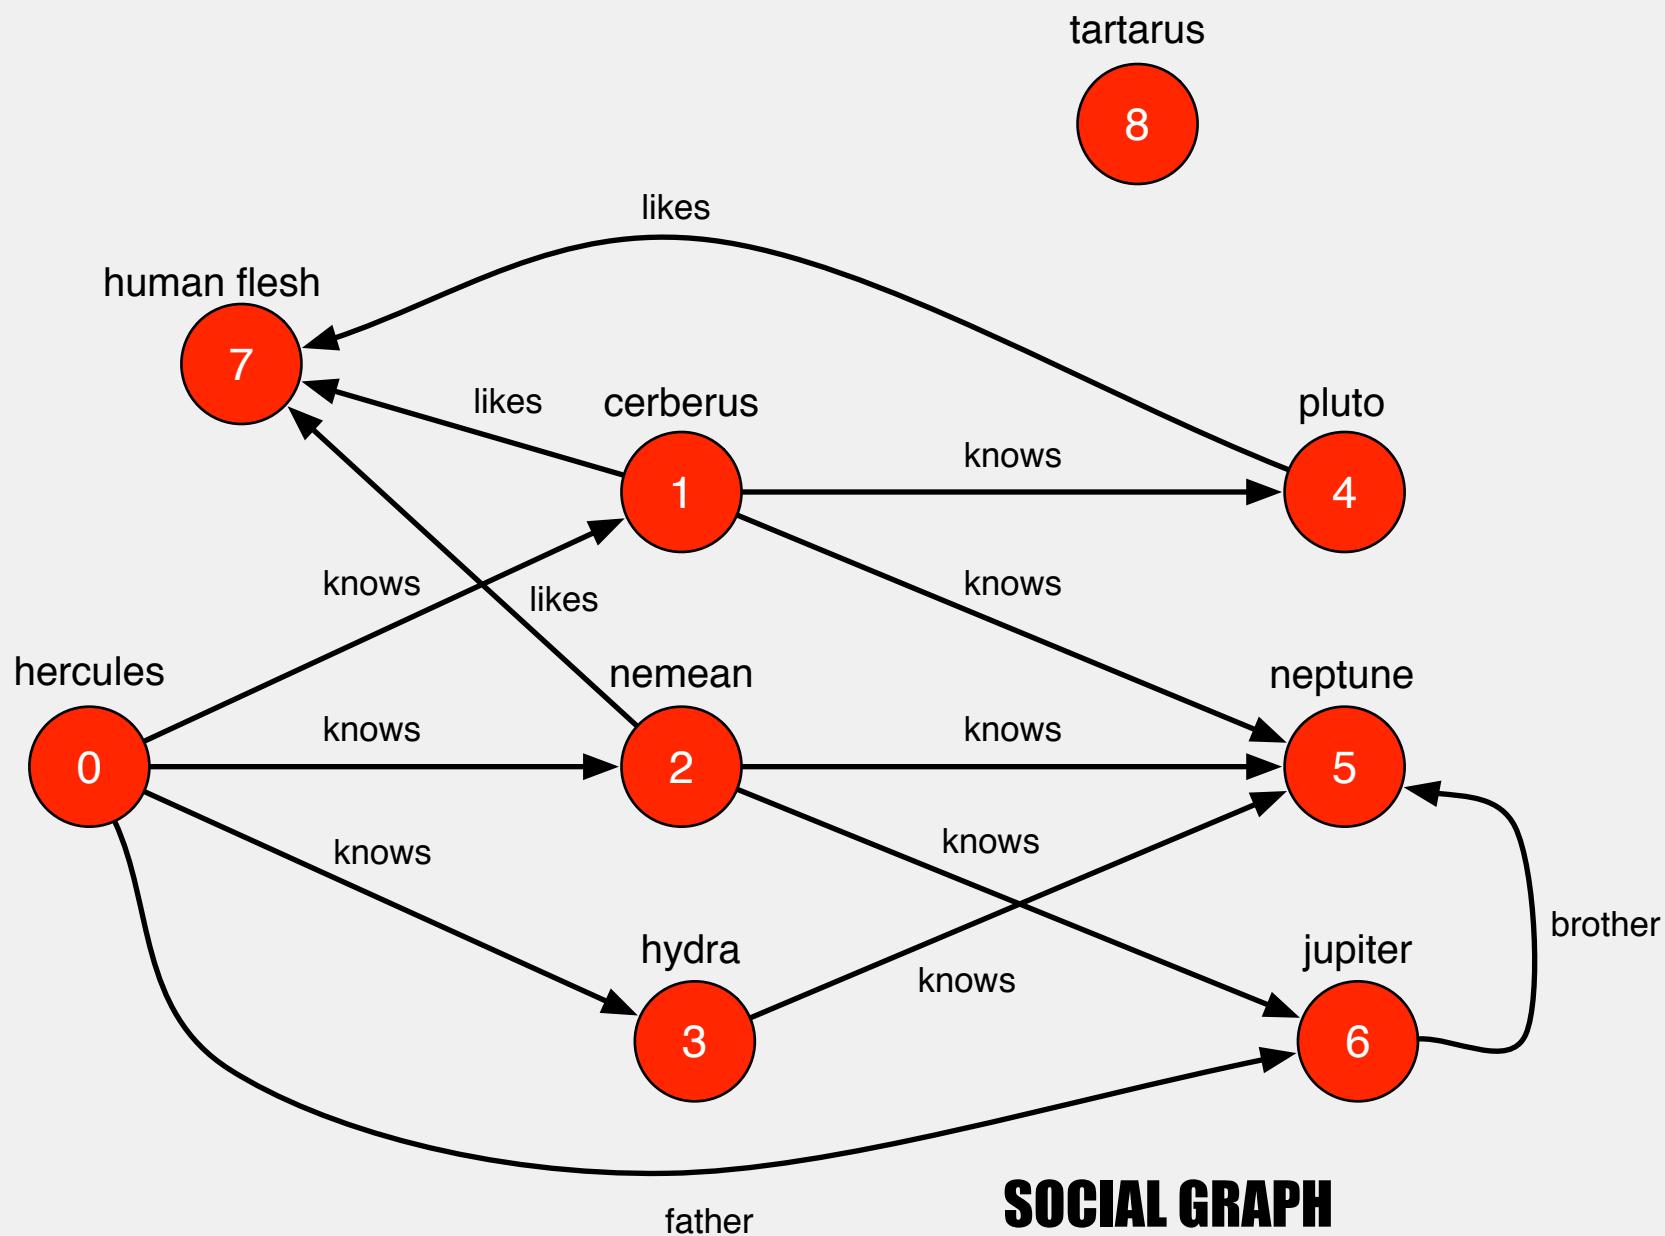


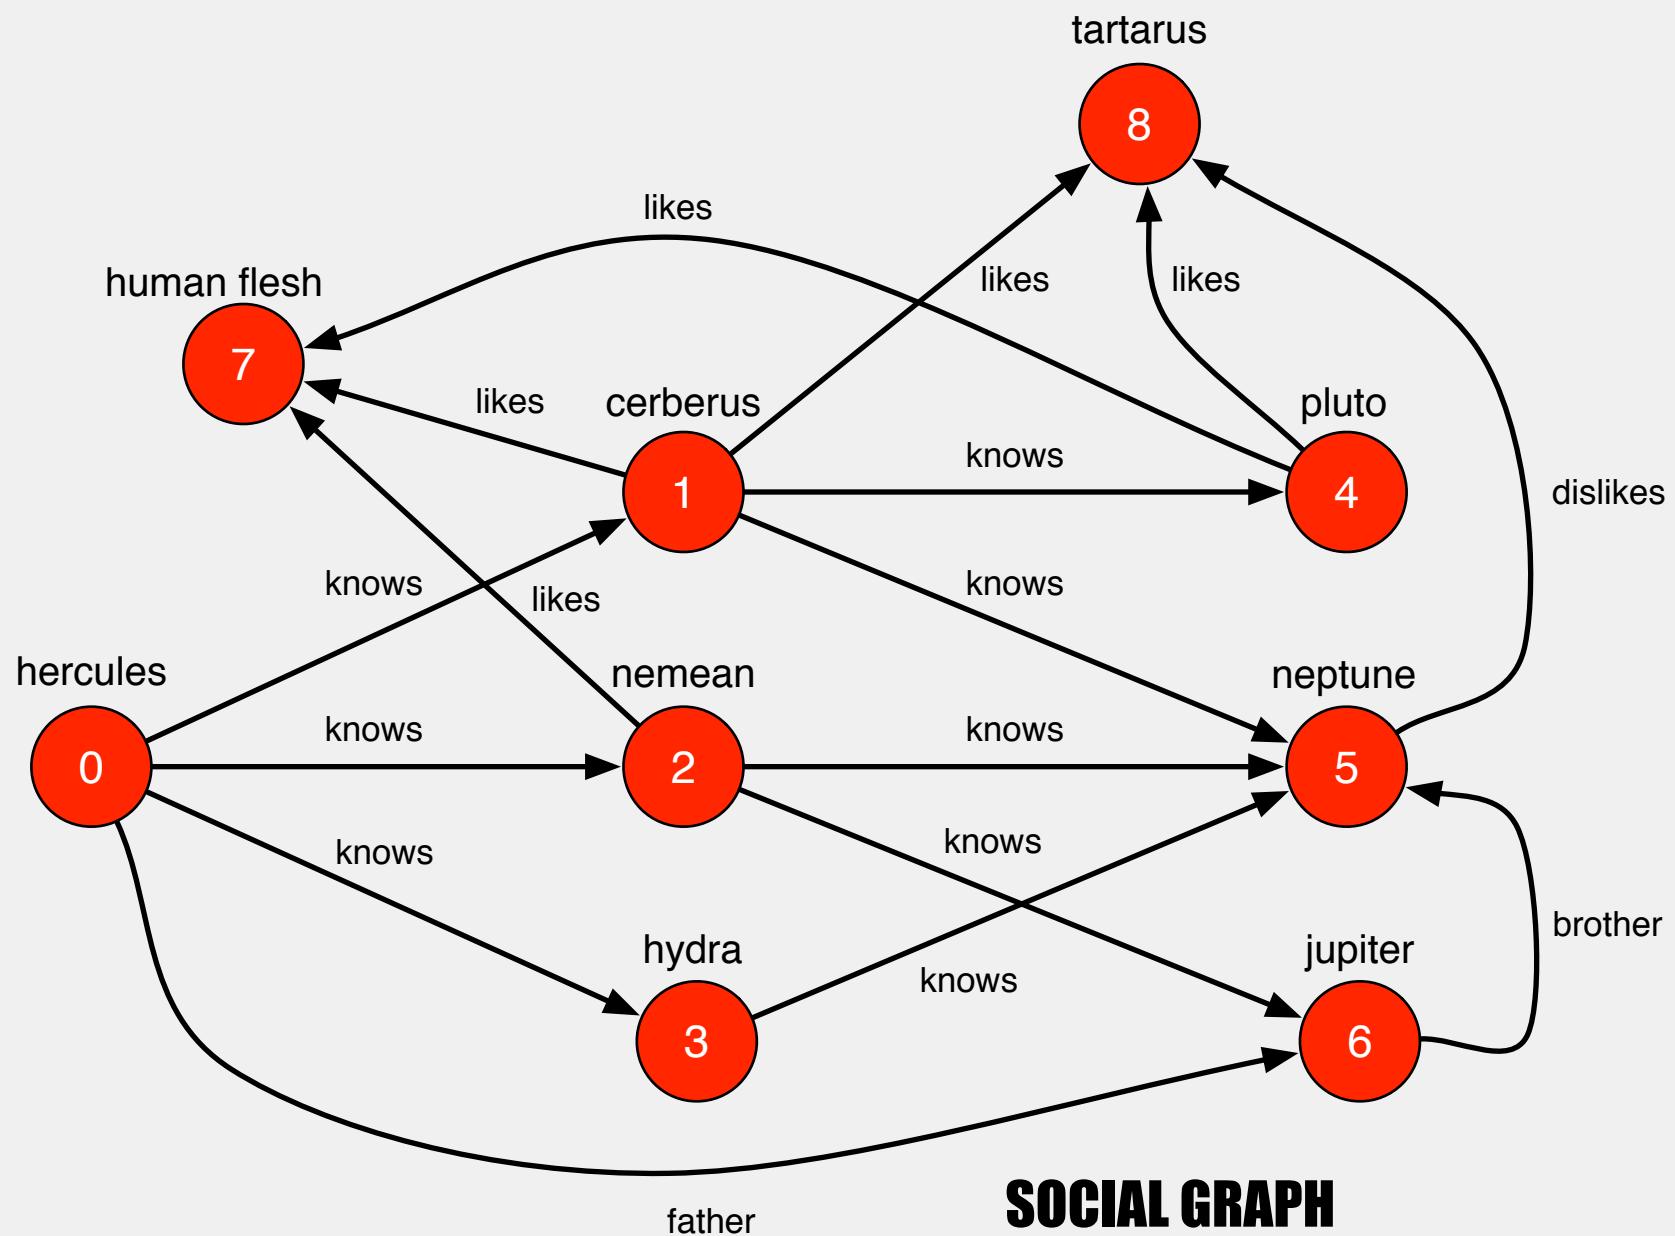




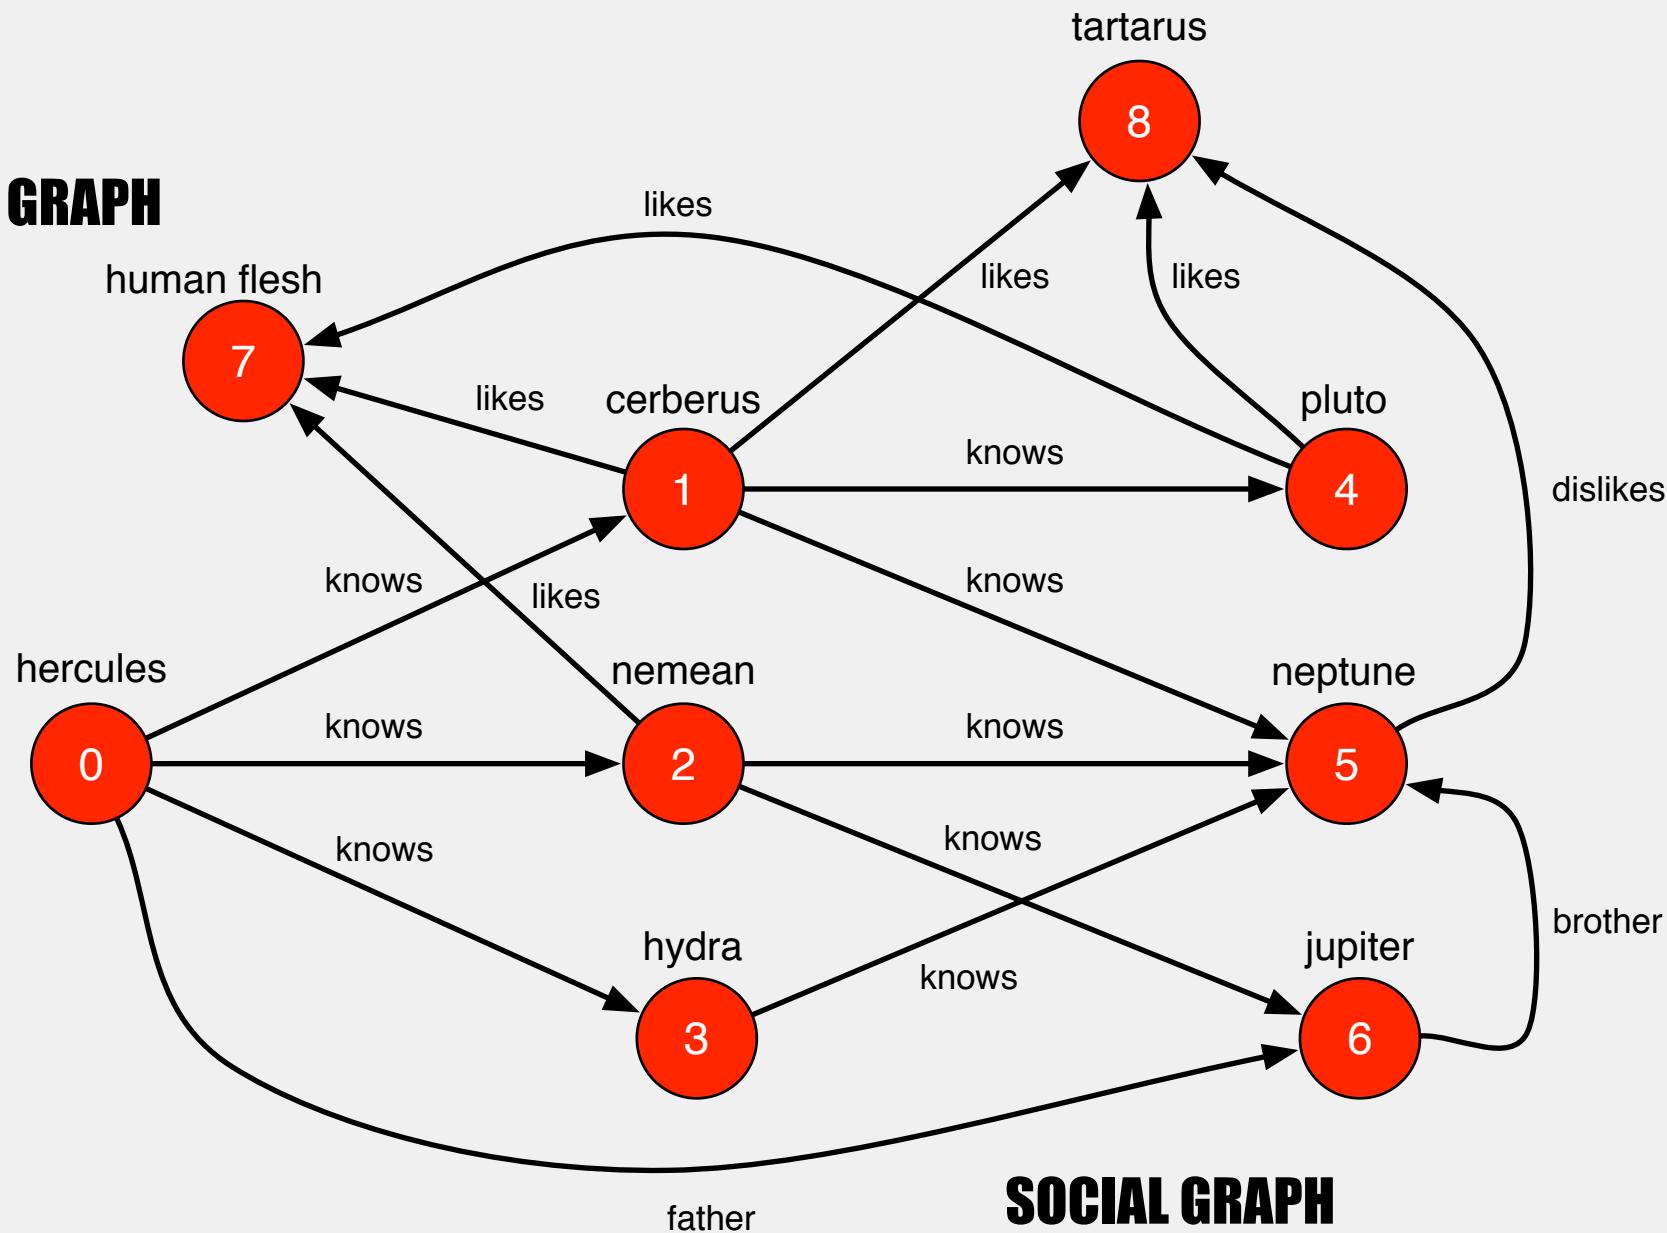






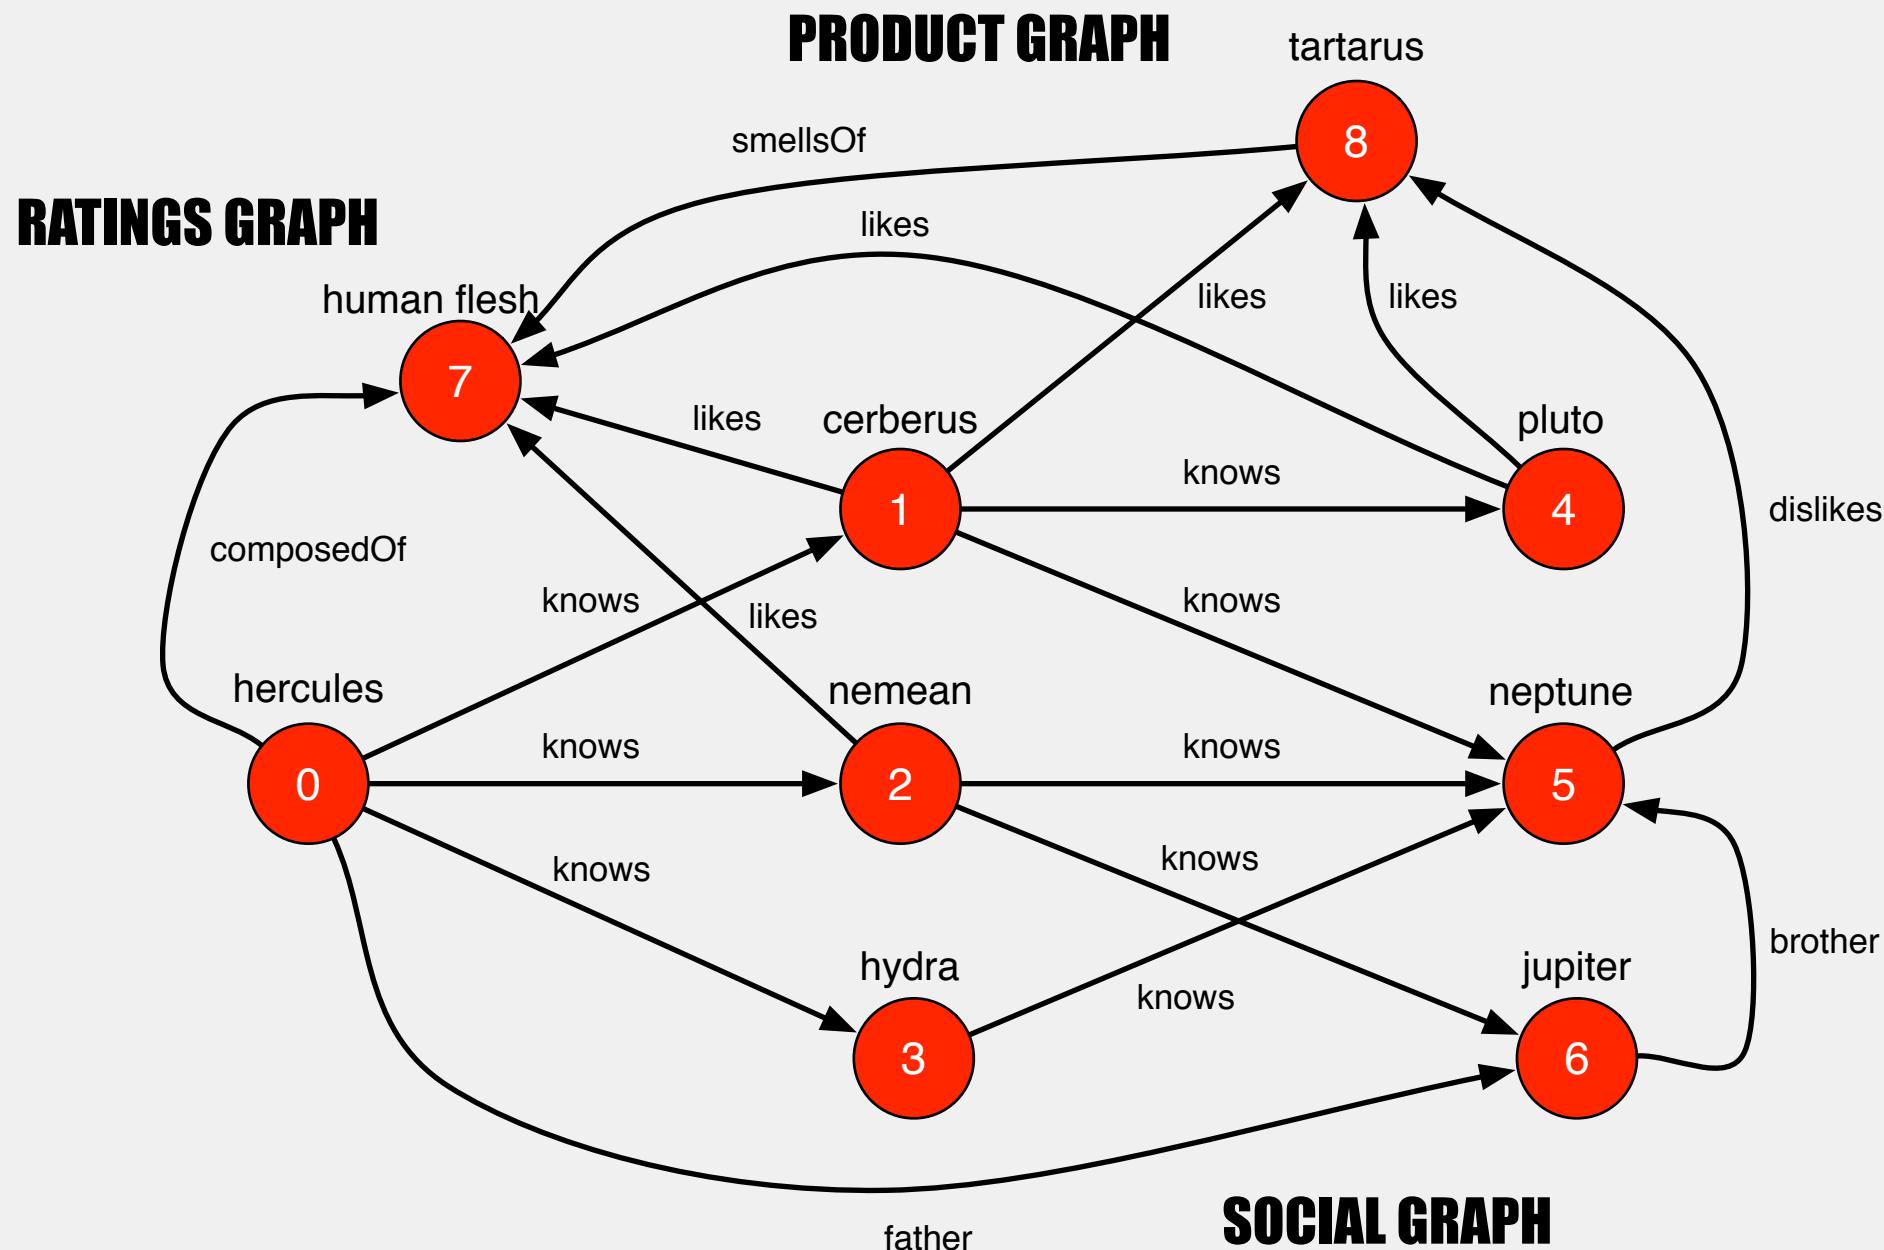


RATINGS GRAPH



SOCIAL GRAPH

Nemean Might Like Tartarus



PATH FINDING



● How is this person related to this film?

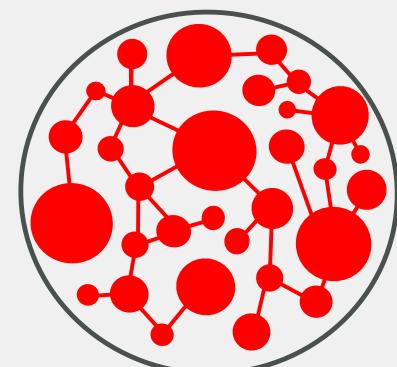
MOVIE GRAPH

● Which authors of this book also
wrote a New York Times bestseller?

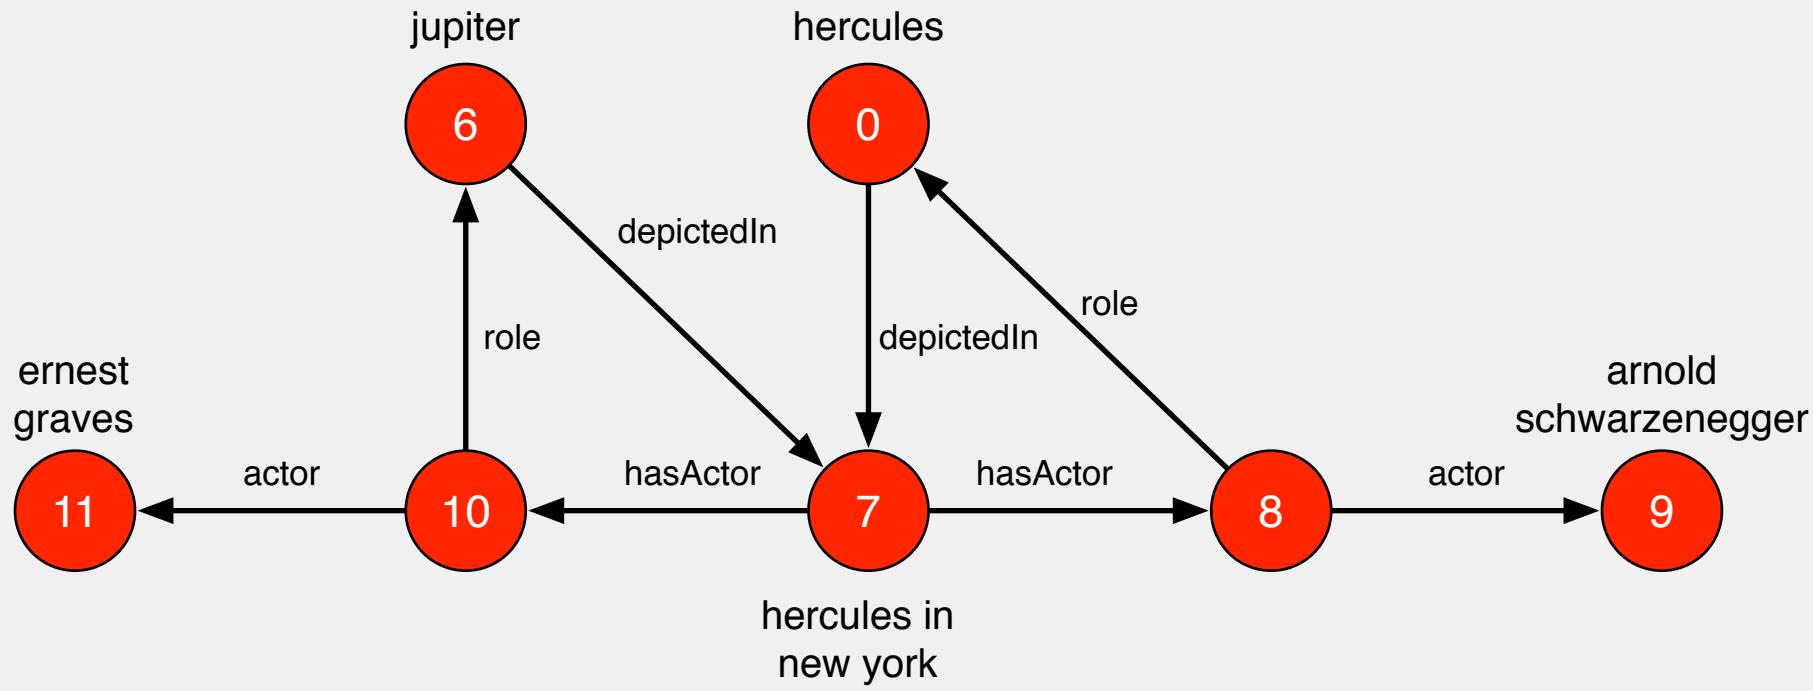
BOOK GRAPH

● Which movies are based on a book by a
New York Times bestseller?

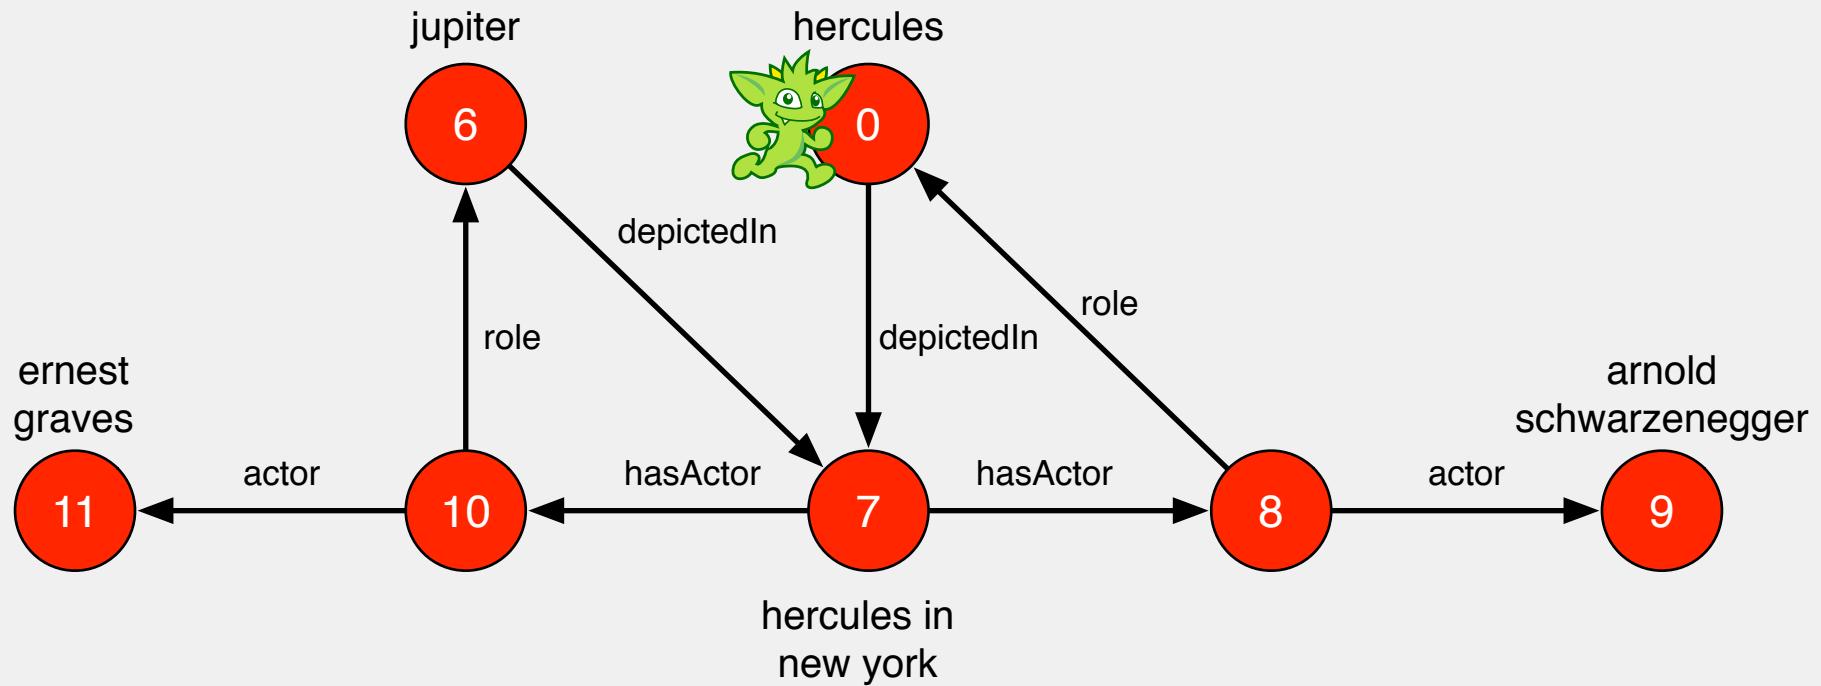
**MOVIE+BOOK
GRAPH**



WHO PLAYED HERCULES IN WHAT MOVIE?

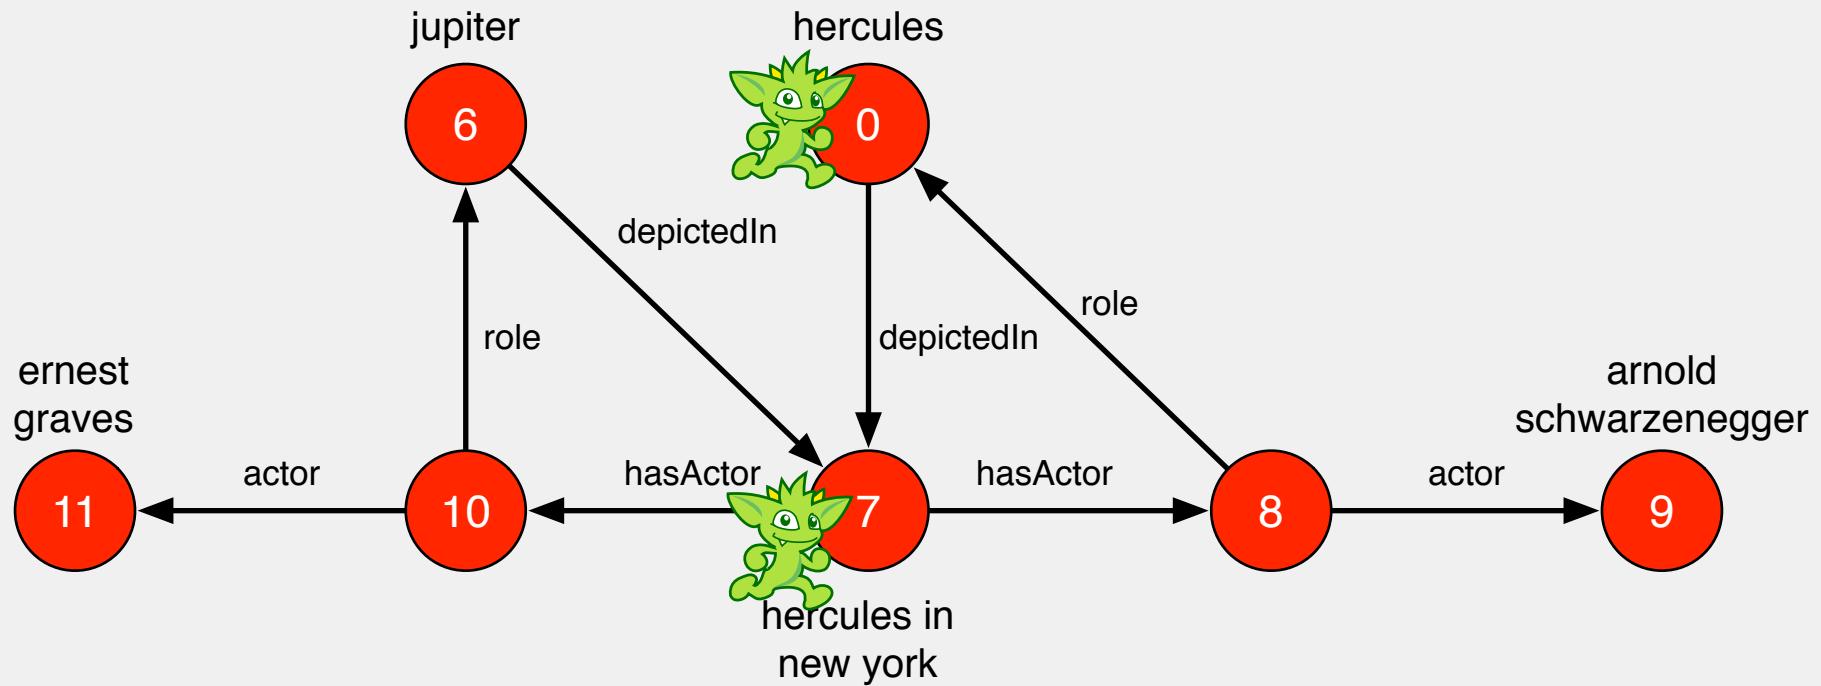


WHO PLAYED HERCULES IN WHAT MOVIE?



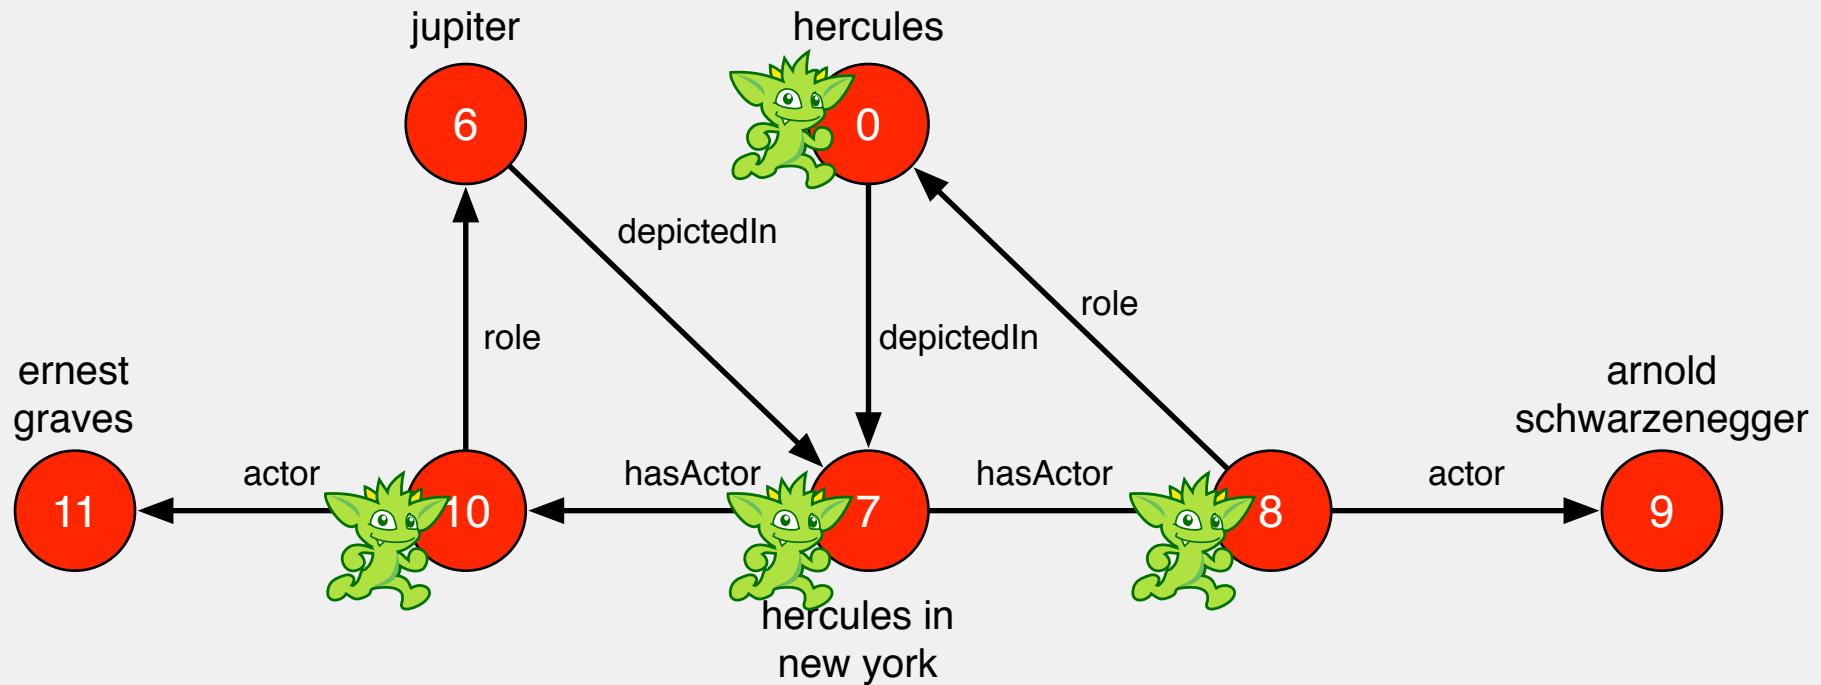
```
gremlin> hercules.match('hercules',
```

WHO PLAYED HERCULES IN WHAT MOVIE?



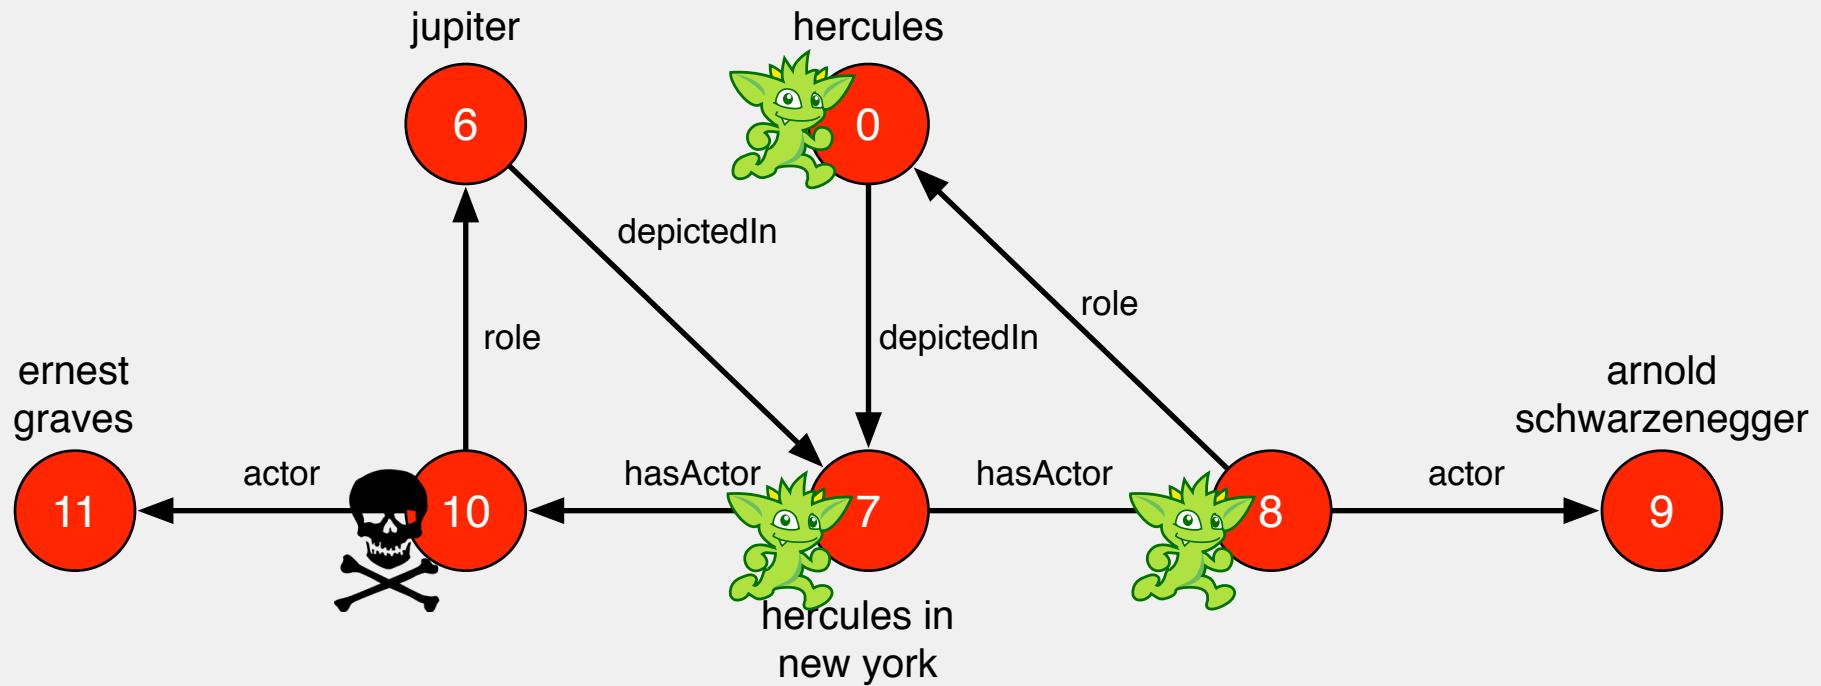
```
gremlin> hercules.match('hercules',
  g.of().as('hercules').out('depictedIn').as('movie'),
```

WHO PLAYED HERCULES IN WHAT MOVIE?



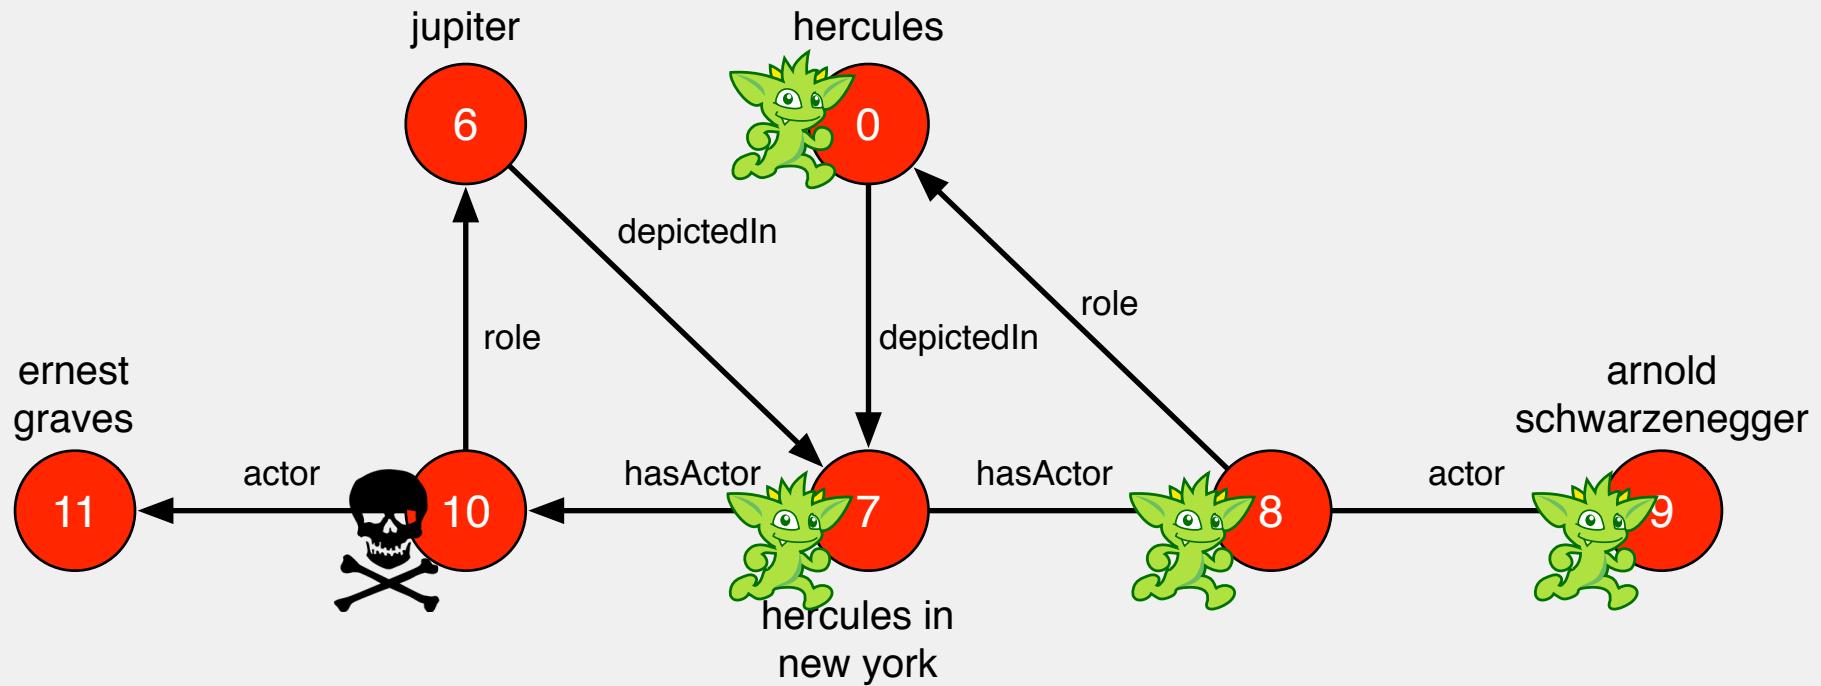
```
gremlin> hercules.match('hercules',
  g.of().as('hercules').out('depictedIn').as('movie'),
  g.of().as('movie').out('hasActor').as('actor'),
```

WHO PLAYED HERCULES IN WHAT MOVIE?



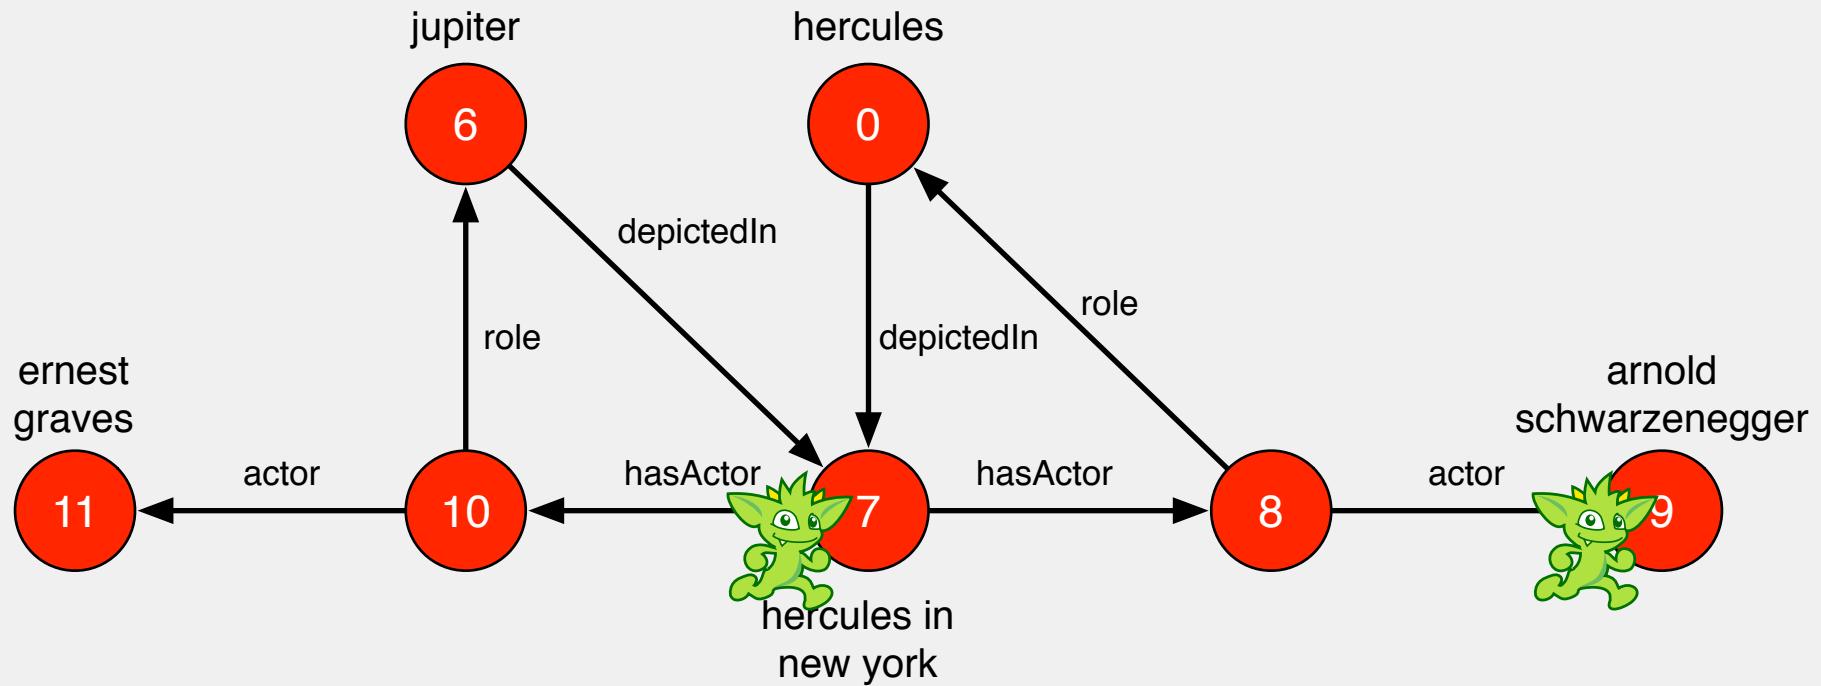
```
gremlin> hercules.match('hercules',
  g.of().as('hercules').out('depictedIn').as('movie'),
  g.of().as('movie').out('hasActor').as('actor'),
  g.of().as('actor').out('role').as('hercules'),
```

WHO PLAYED HERCULES IN WHAT MOVIE?



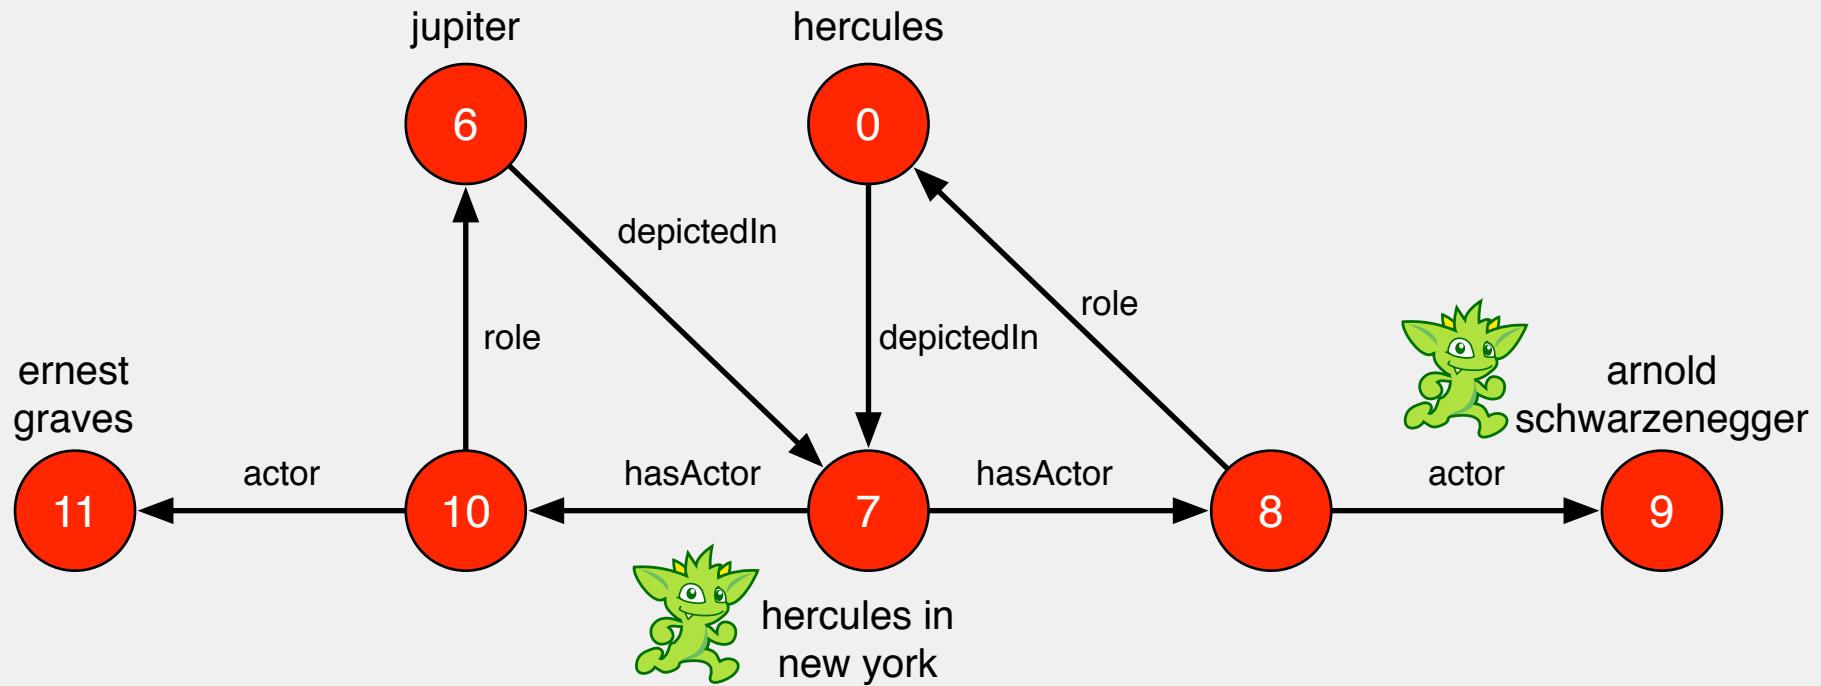
```
gremlin> hercules.match('hercules',
  g.of().as('hercules').out('depictedIn').as('movie'),
  g.of().as('movie').out('hasActor').as('actor'),
  g.of().as('actor').out('role').as('hercules'),
  g.of().as('actor').out('actor').as('star'))
```

WHO PLAYED HERCULES IN WHAT MOVIE?

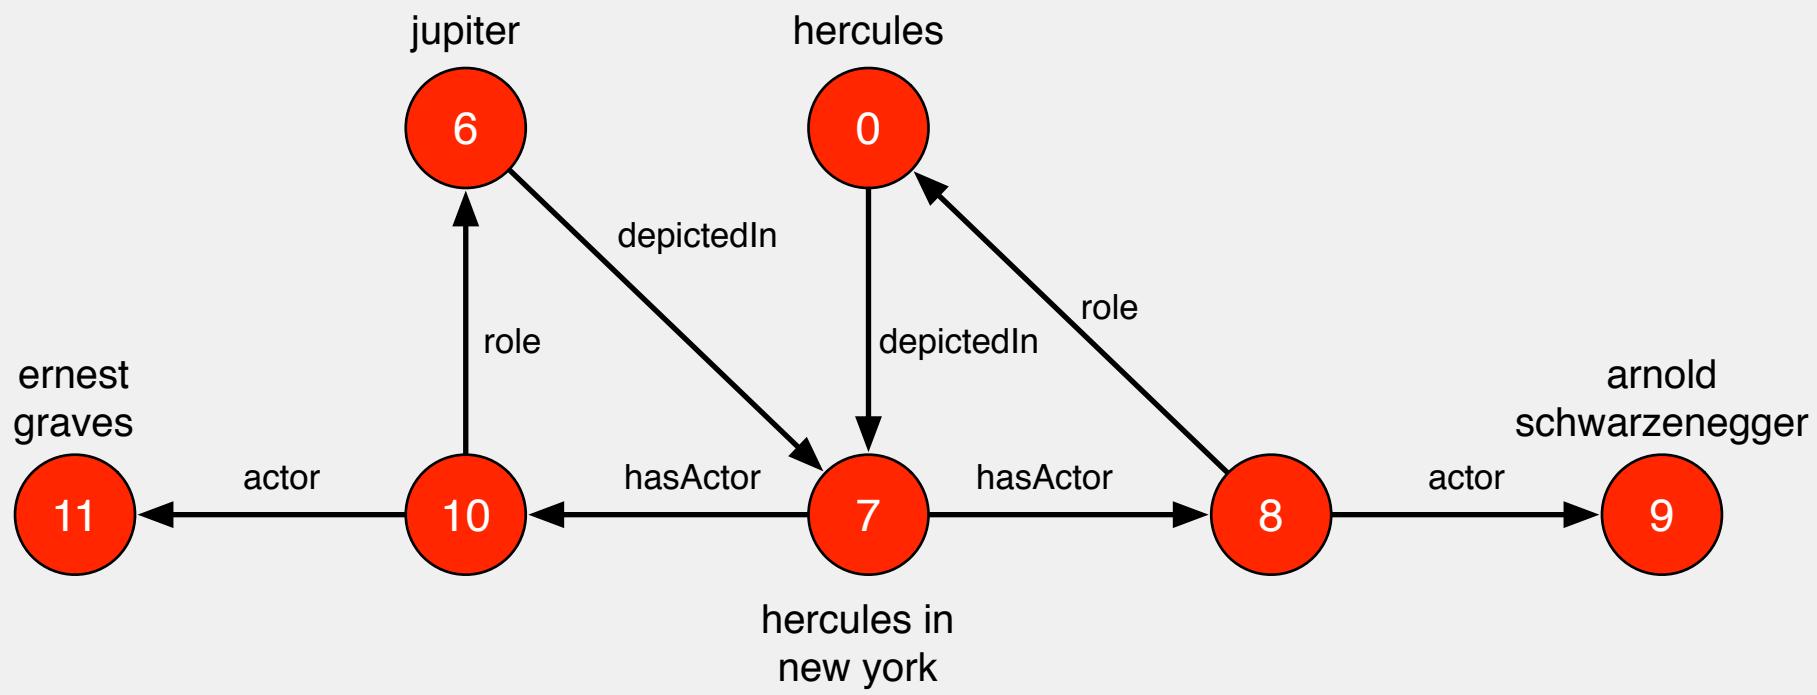


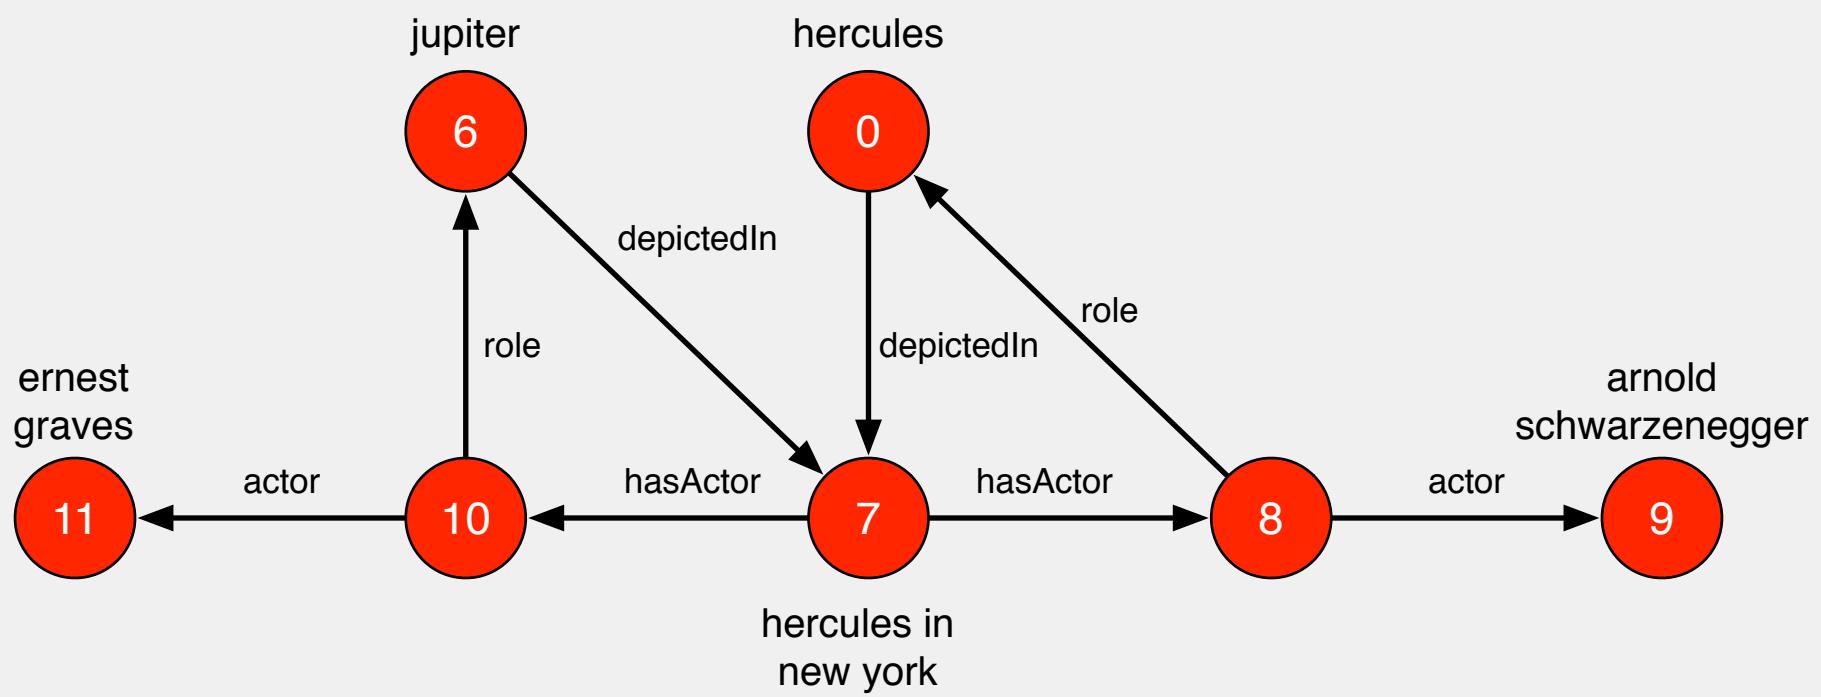
```
gremlin> hercules.match('hercules',
  g.of().as('hercules').out('depictedIn').as('movie'),
  g.of().as('movie').out('hasActor').as('actor'),
  g.of().as('actor').out('role').as('hercules'),
  g.of().as('actor').out('actor').as('star'))
.select(['movie', 'star'])
==>[movie:v[7], star:v[9]]
```

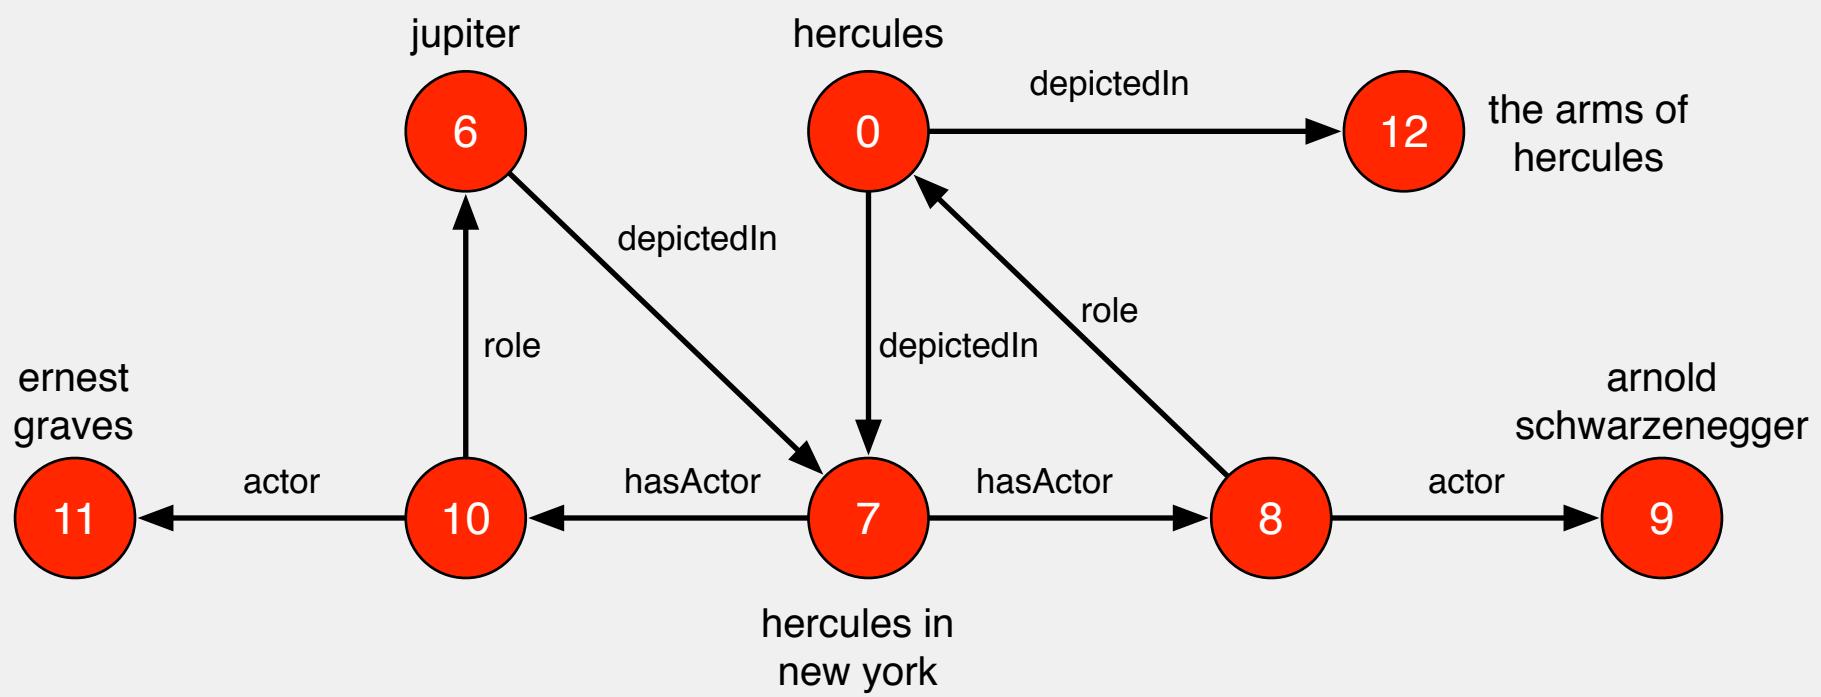
WHO PLAYED HERCULES IN WHAT MOVIE?

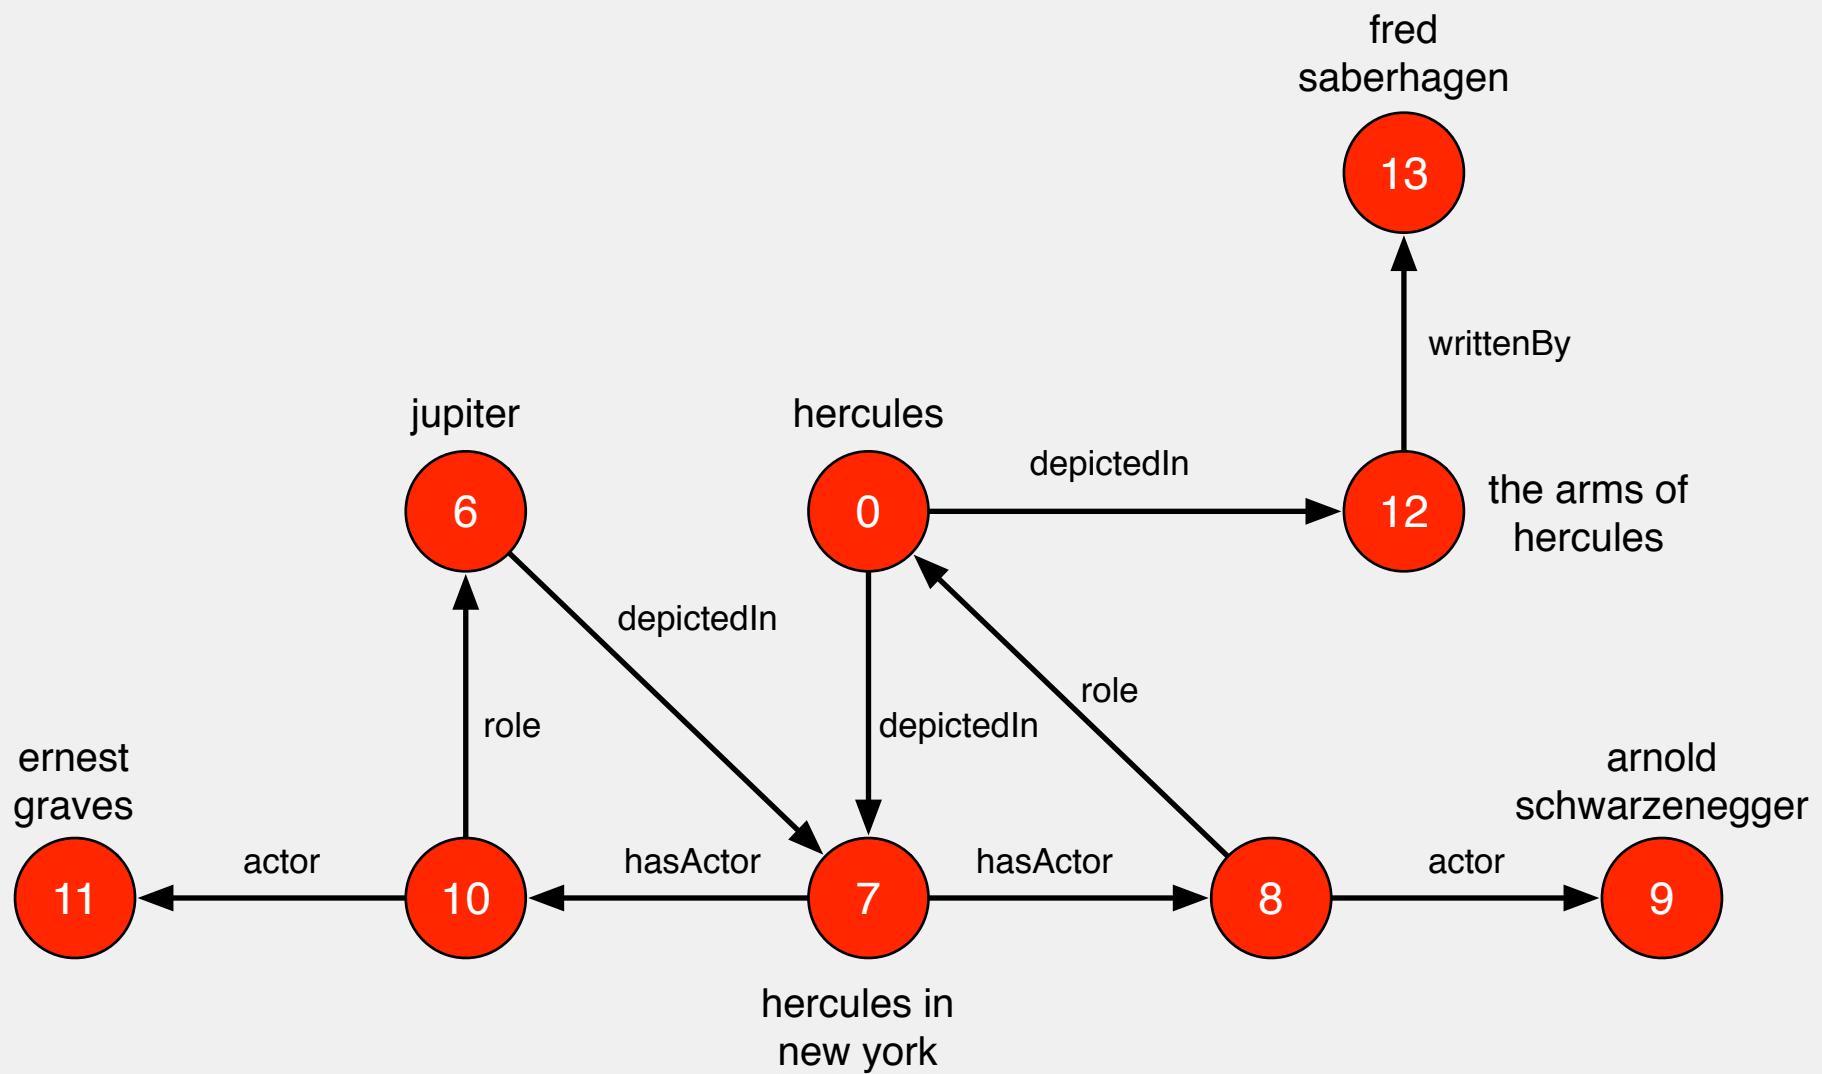


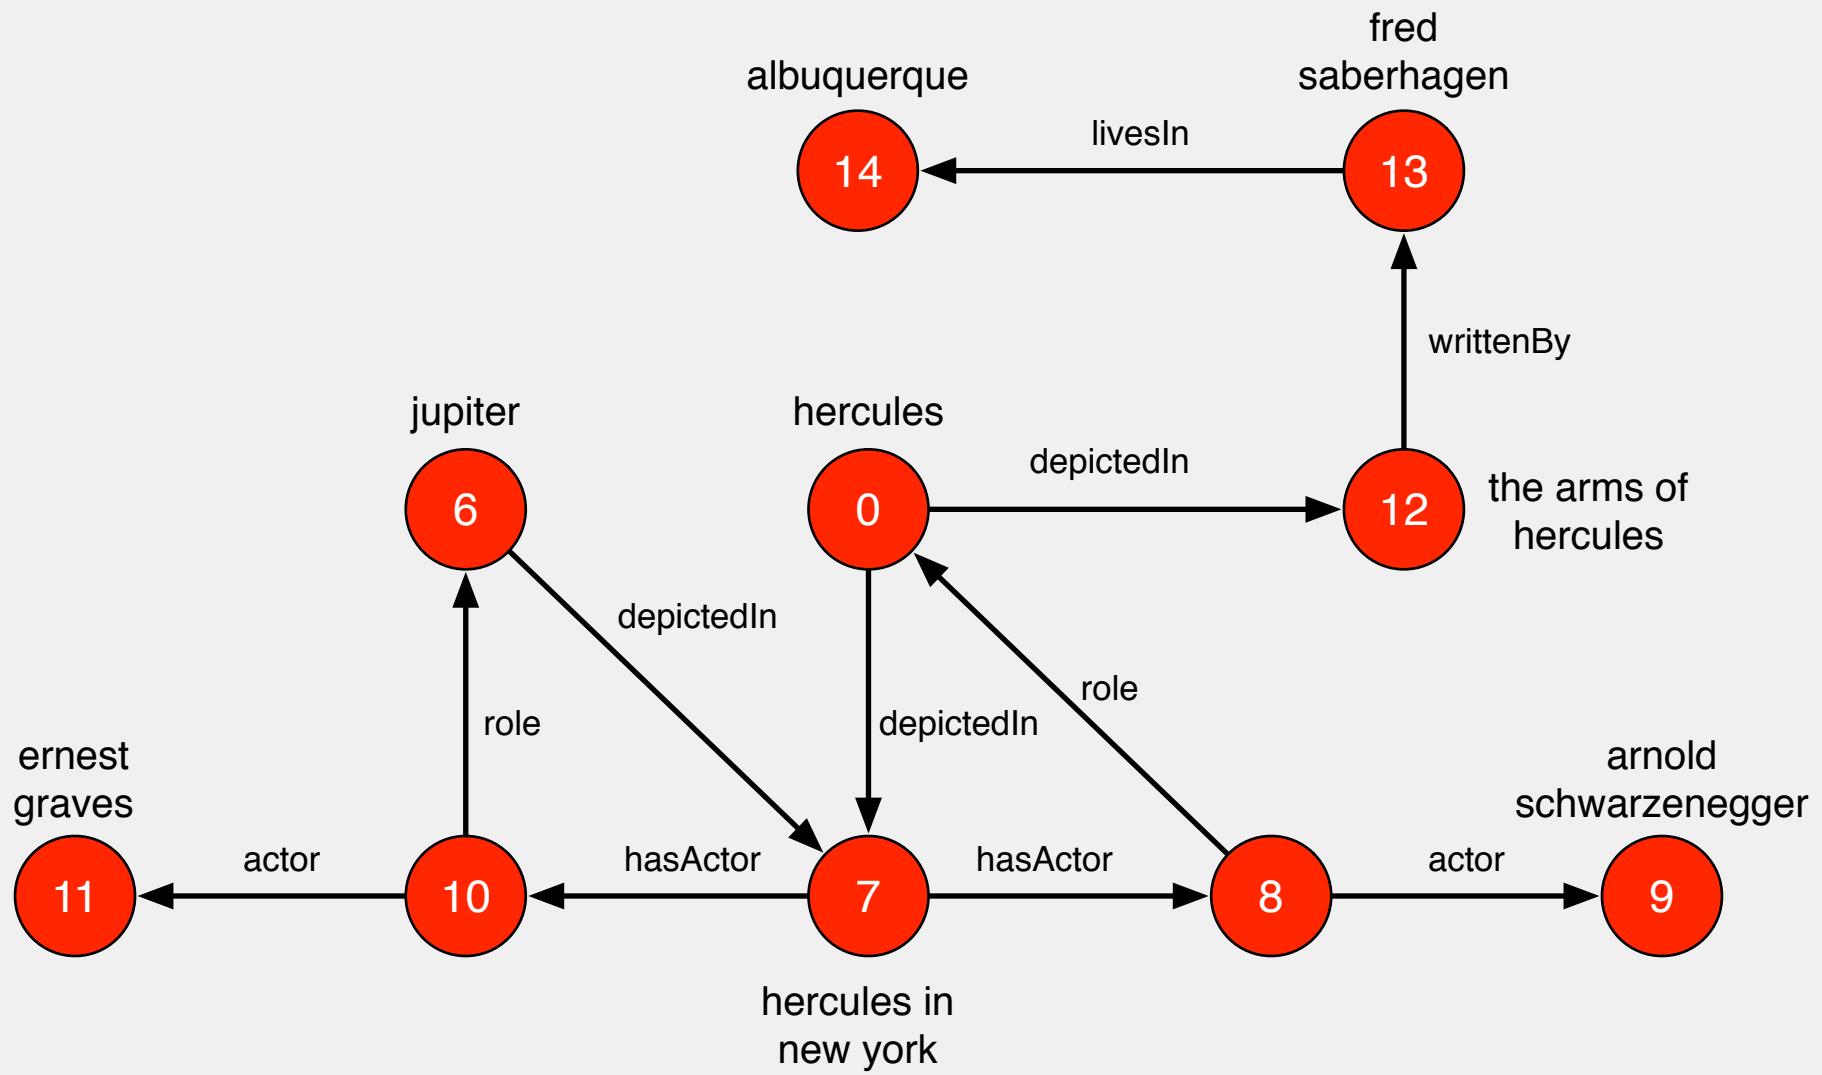
```
gremlin> hercules.match('hercules',
  g.of().as('hercules').out('depictedIn').as('movie')
  g.of().as('movie').out('hasActor').as('actor')
  g.of().as('actor').out('role').as('hercules')
  g.of().as('actor').out('actor').as('star'))
.select(['movie', 'star']).{it.name}
==>[movie:hercules in new york, star:arnold schwarzenegger]
```

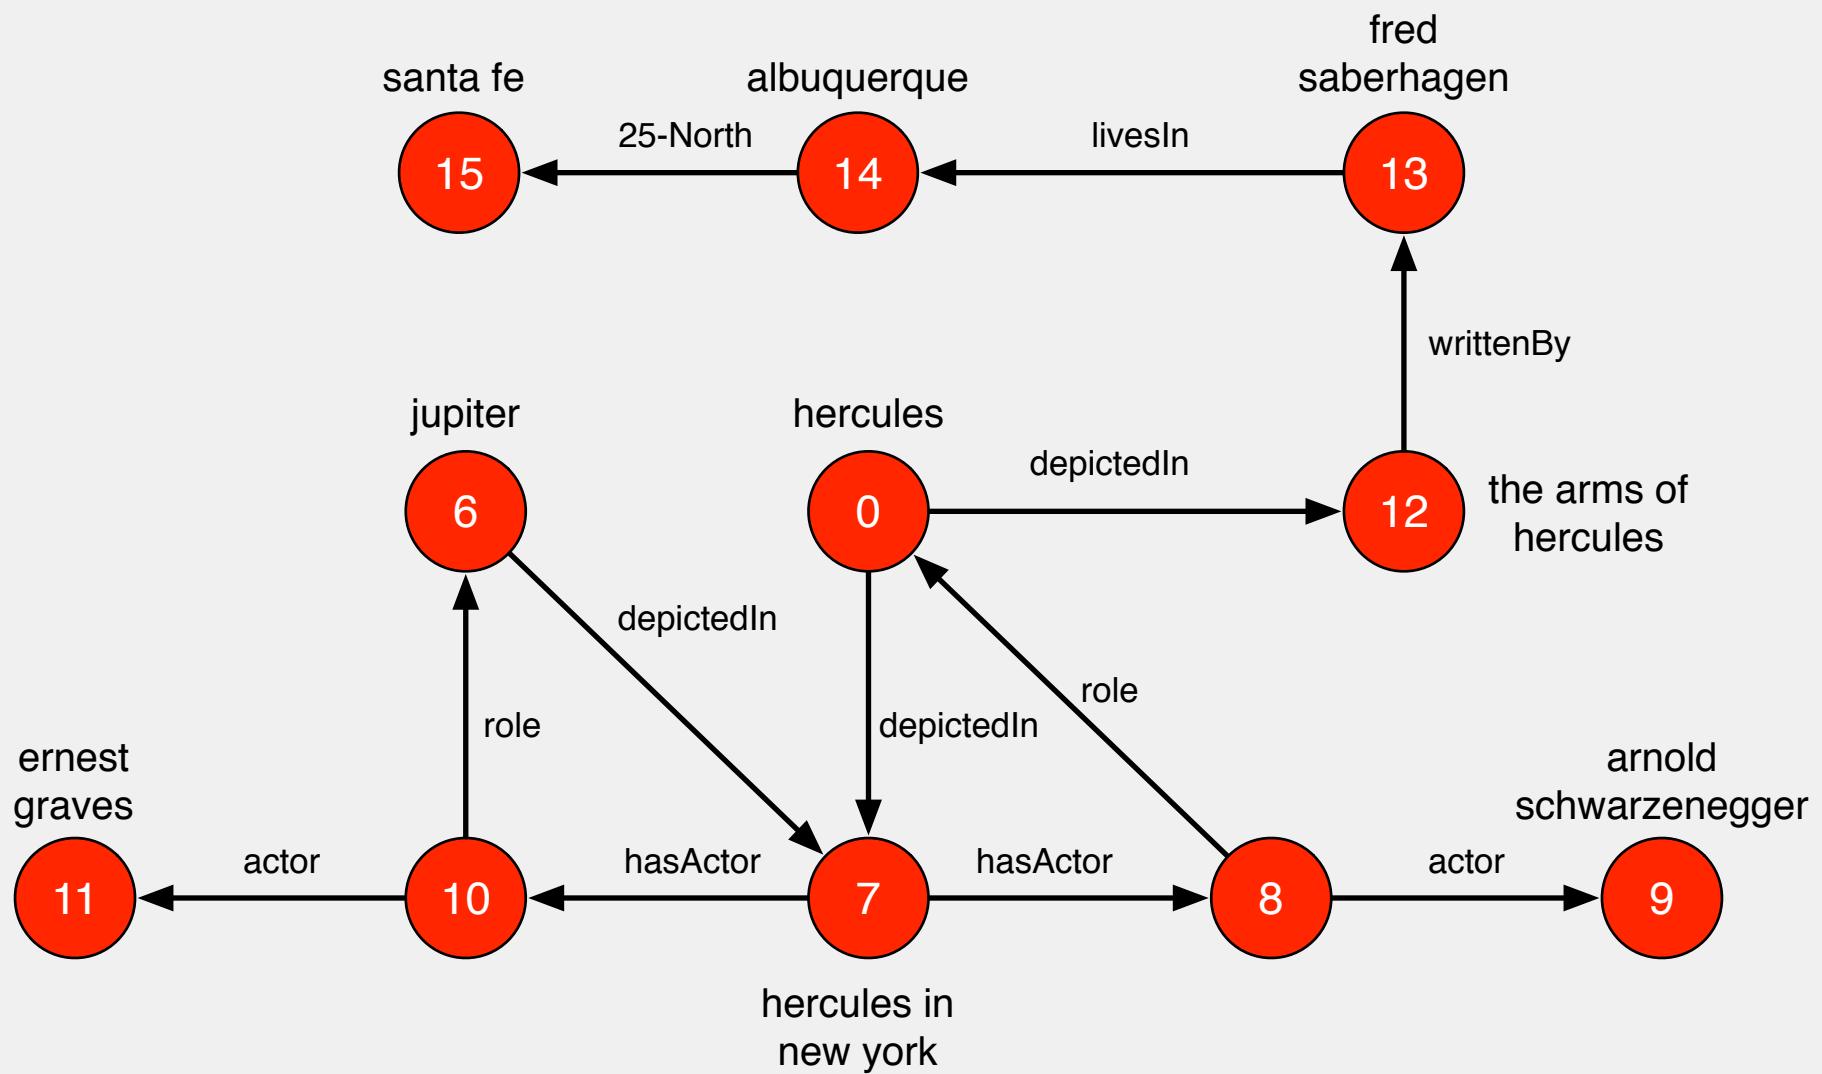


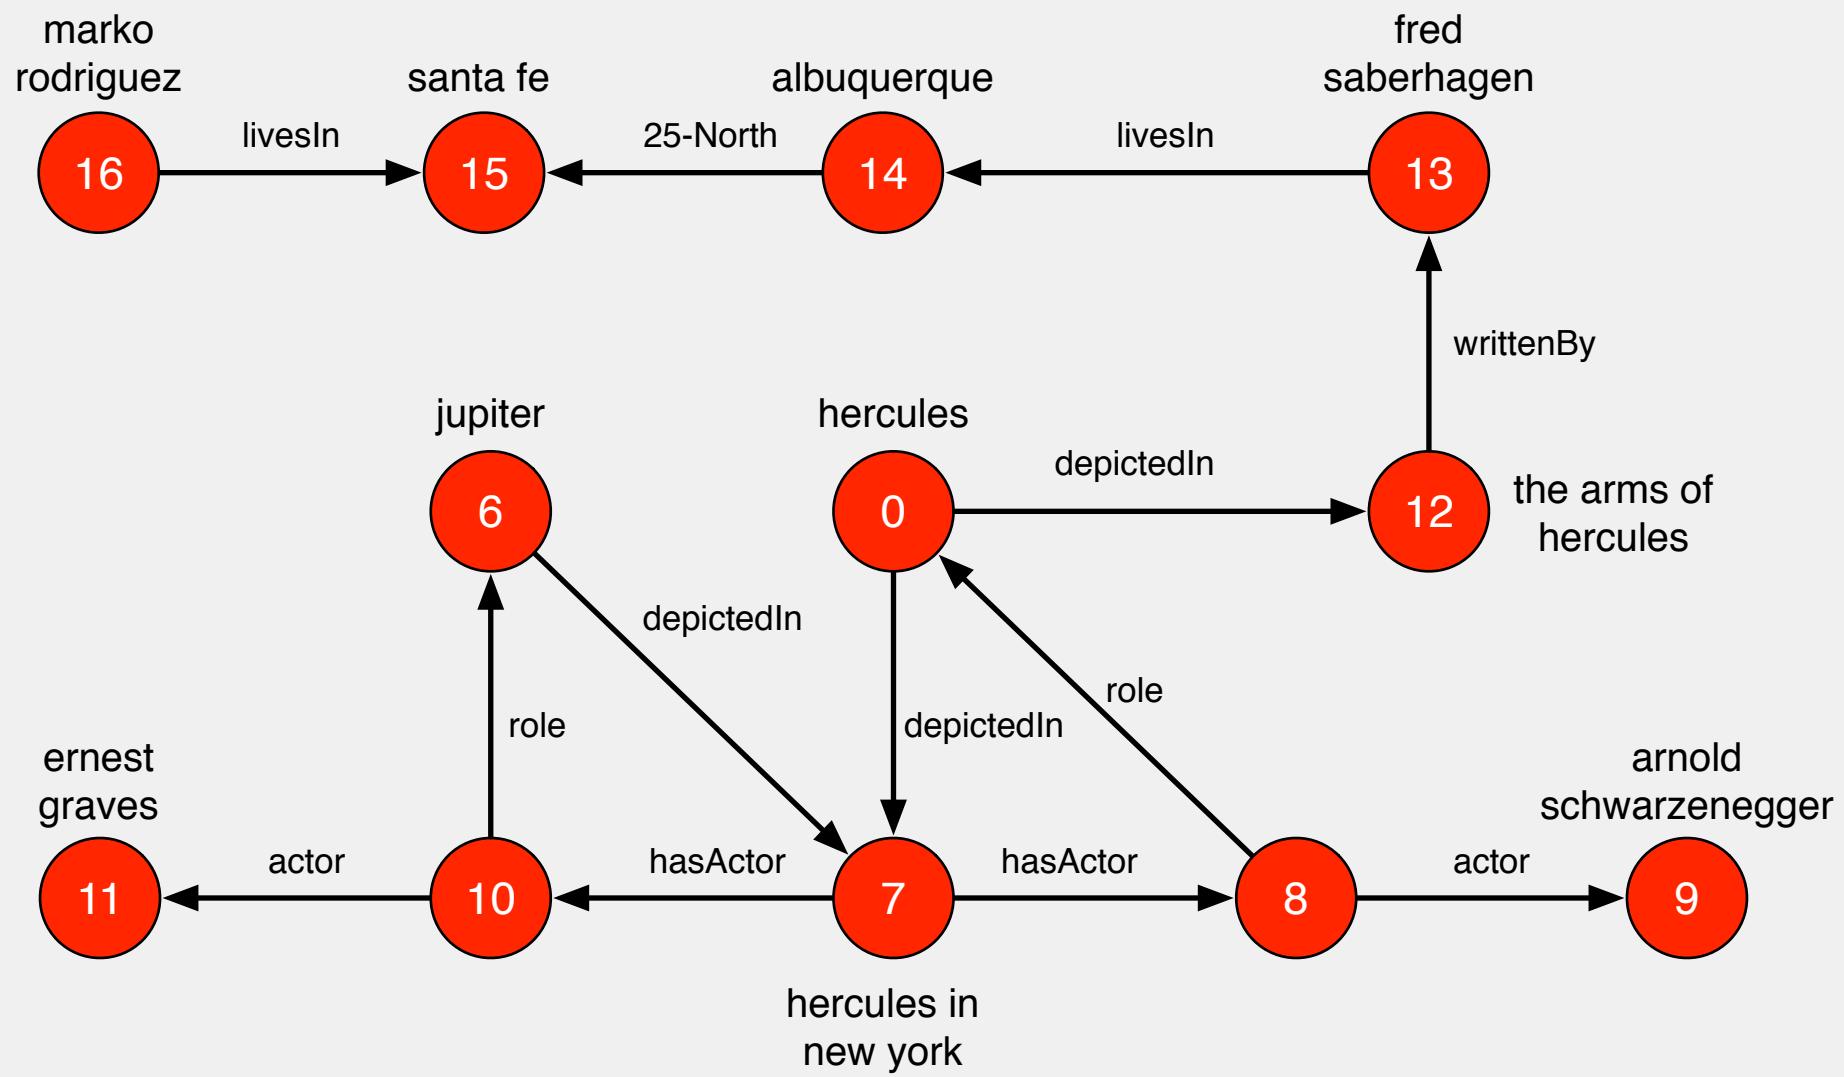


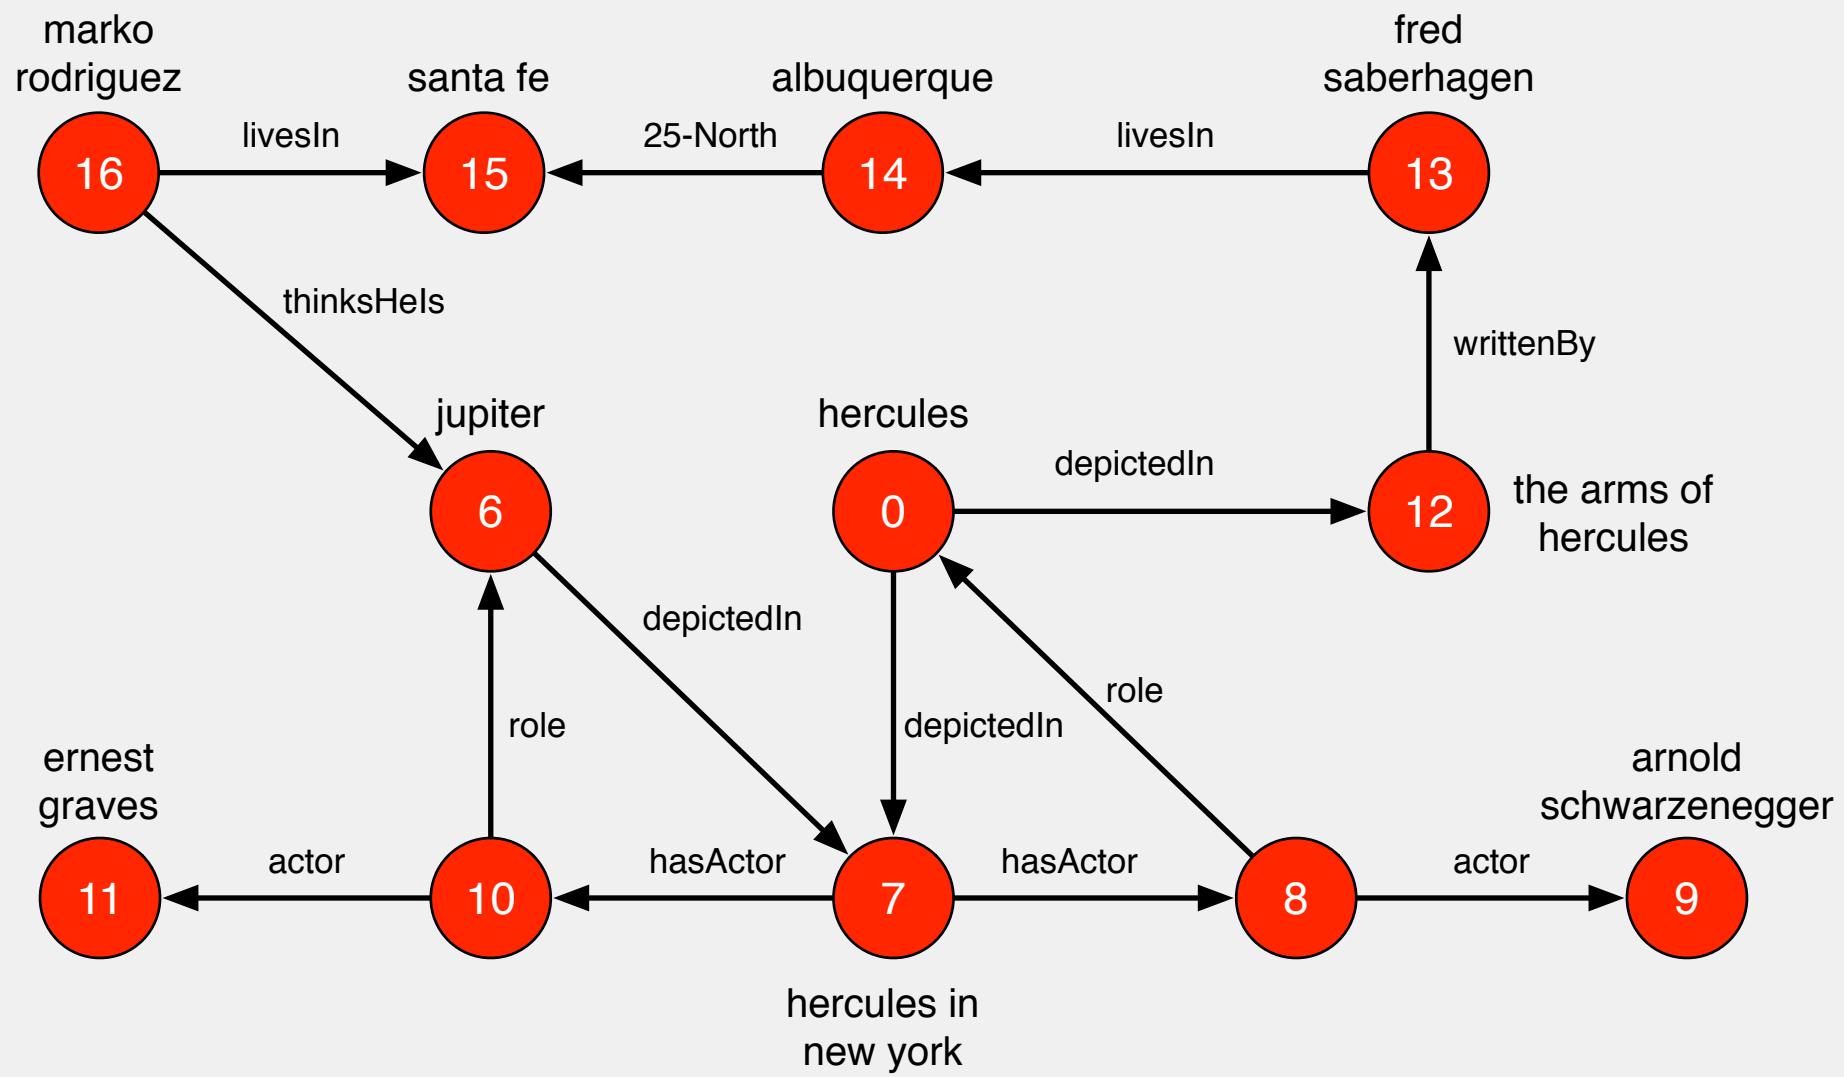




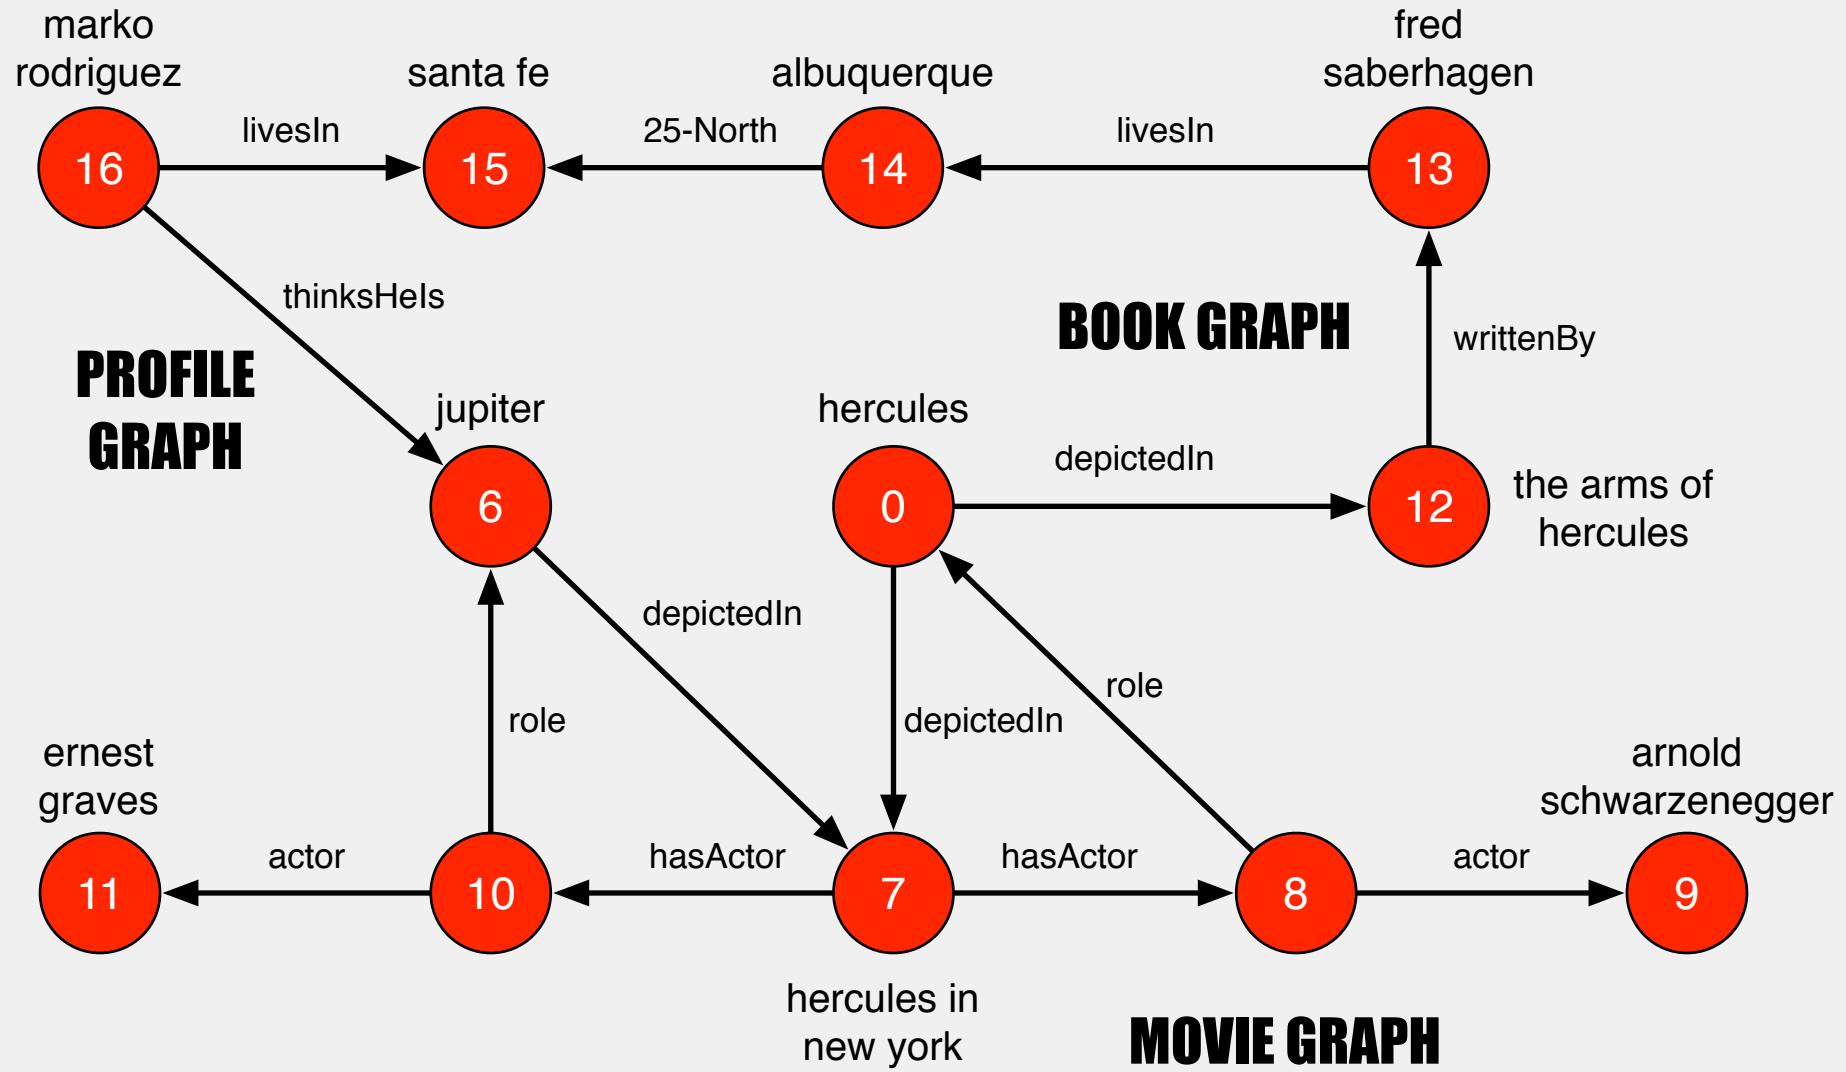








TRANSPORTATION GRAPH

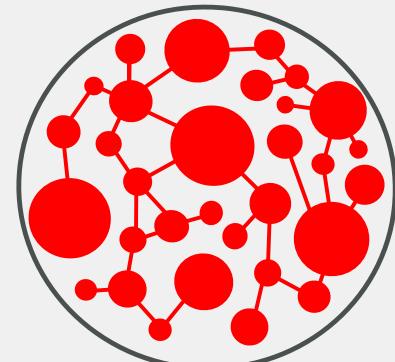


SOCIAL INFLUENCE



- Who are the most influential people in java, mathematics, art, surreal art, politics, ...?
- Which region of the social graph will propagate this advertisement this furthest?
- Which 3 experts should review this submitted article?
- Which people should I talk to at the upcoming conference and what topics should I talk to them about?

SOCIAL + COMMUNICATION + EXPERTISE + EVENT GRAPH



PATTERN IDENTIFICATION

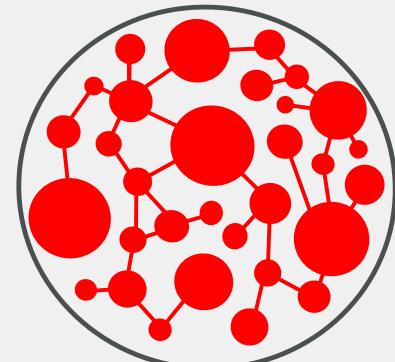


- This connectivity pattern is a sign of financial fraud.
When this motif is found, a red flag will be raised.

TRANSACTION GRAPH

- Healthy discourse is typified by a discussion board with a branch factor in this range and a concept clique score in this range.

DISCUSSION GRAPH



KNOWLEDGE DISCOVERY

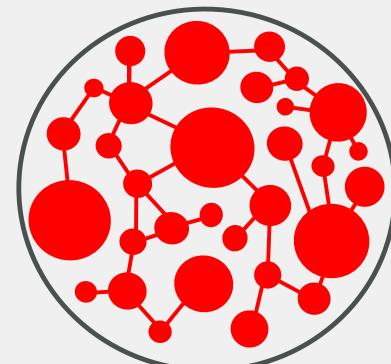


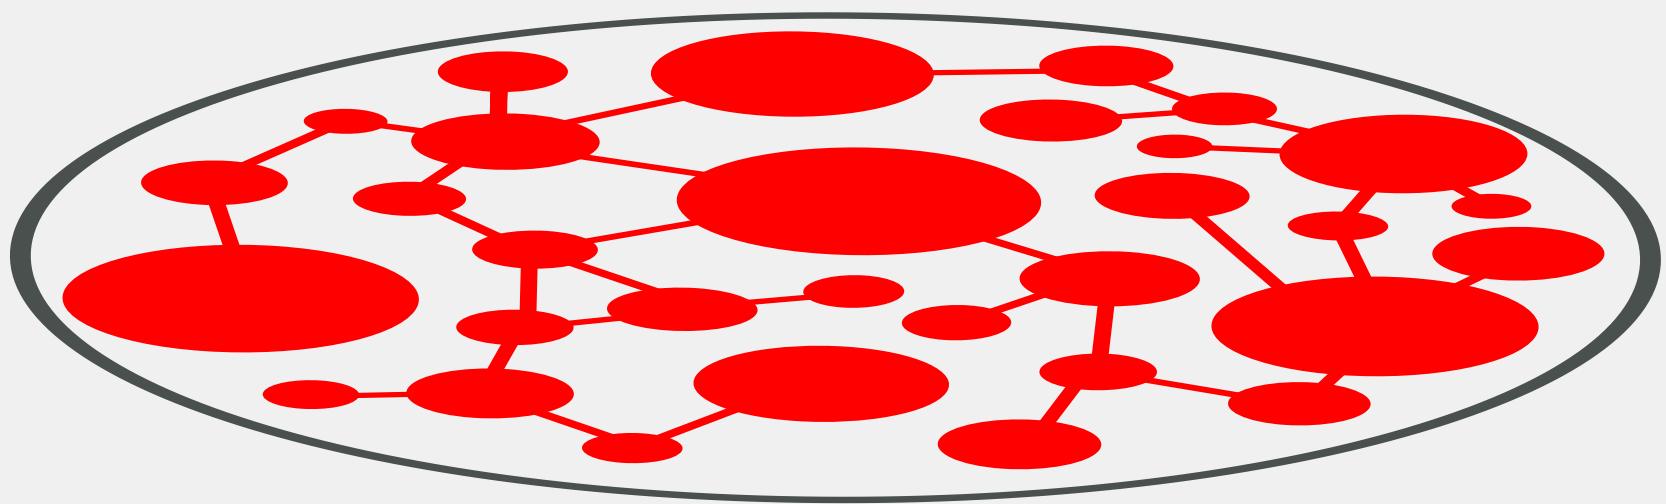
- The terms "ice", "fans", "stanley cup," are classified as "sports"

WIKIPEDIA GRAPH

- Given that all identified birds fly,
it can be deduced that all birds fly.
If contrary evidence is provided,
then this "fact" can be retracted.

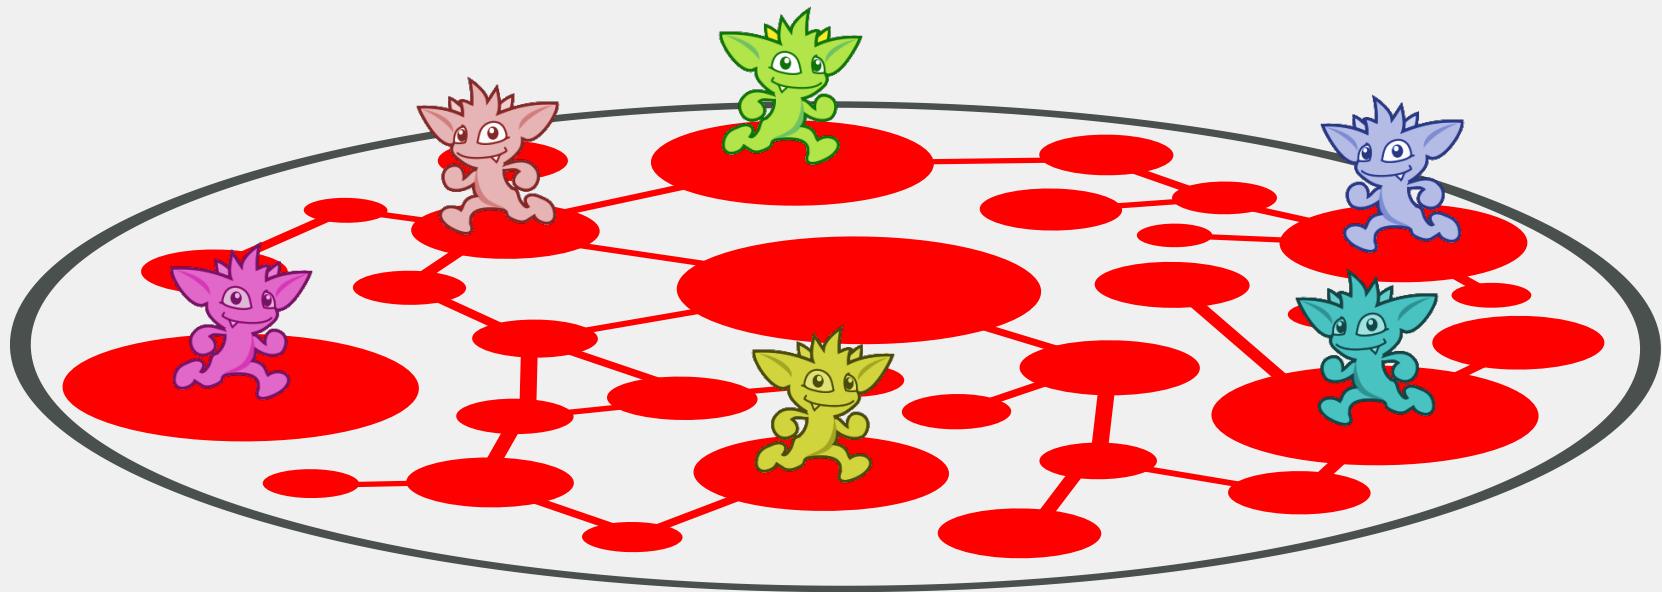
EVIDENTIAL LOGIC GRAPH





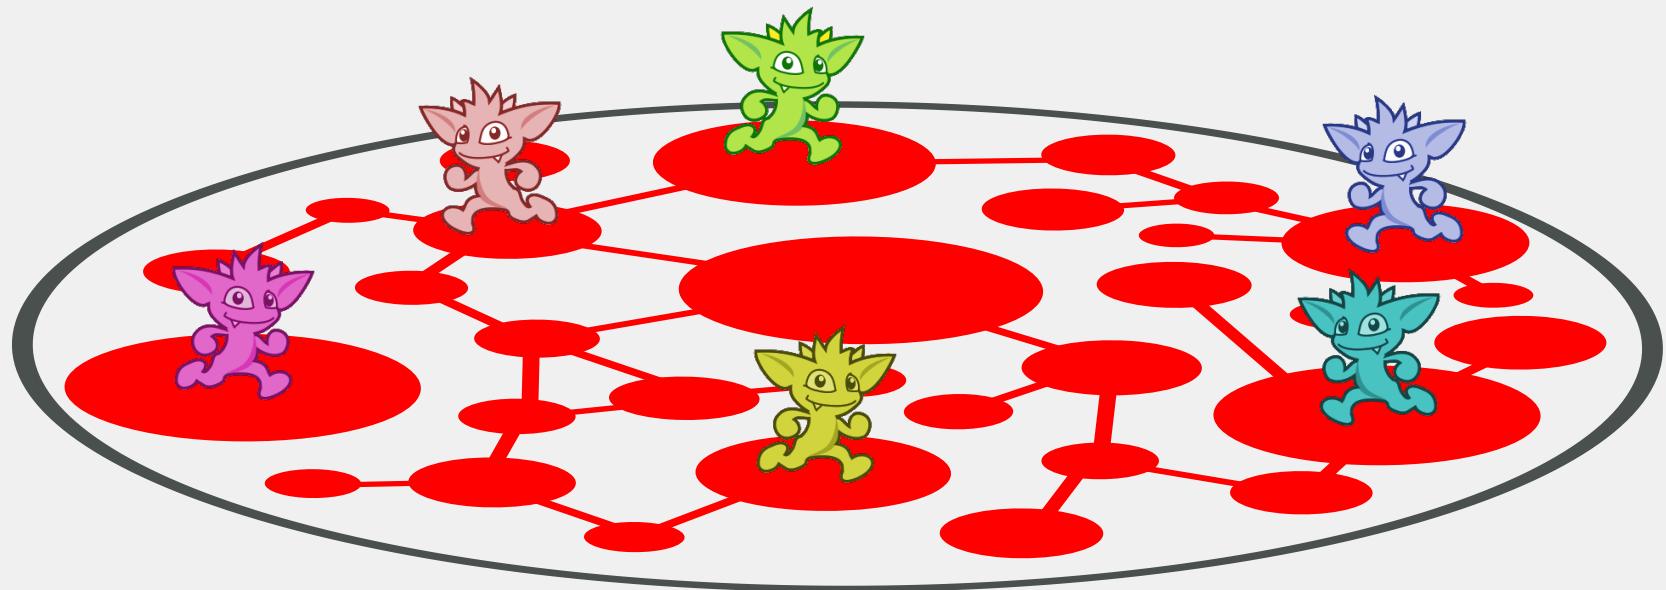
WORLD MODEL

WORLD PROCESSES



WORLD MODEL

WORLD PROCESSES



WORLD MODEL

A single world model and various types of traversers moving through that model to solve problems.

COMPUTING

PROCESS

STRUCTURE

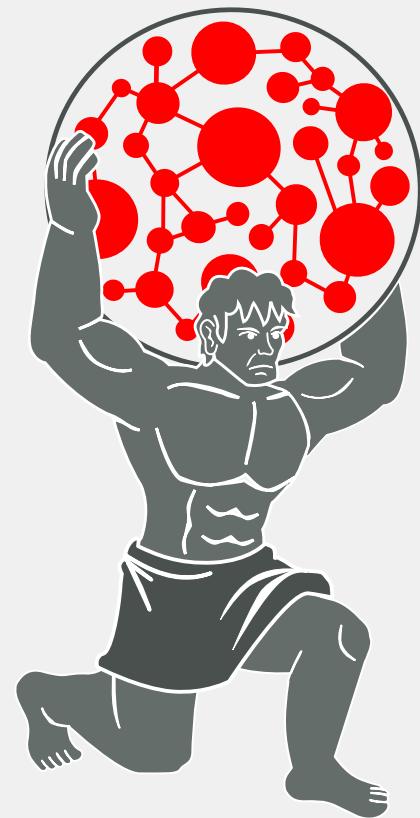


TRAVERSAL

GRAPH

GRAPH-BASED
COMPUTING

PART 2: THE PATH TO TITAN 1.0



WHY CREATE TITAN?



- A number of Aurelius' clients...
 - ...need to represent and process graphs at the 100+ billion edge scale w/ thousands of concurrent transactions.
 - ...need both local graph traversals (OLTP) and batch graph processing (OLAP).
 - ...desire a free, open source distributed graph database.

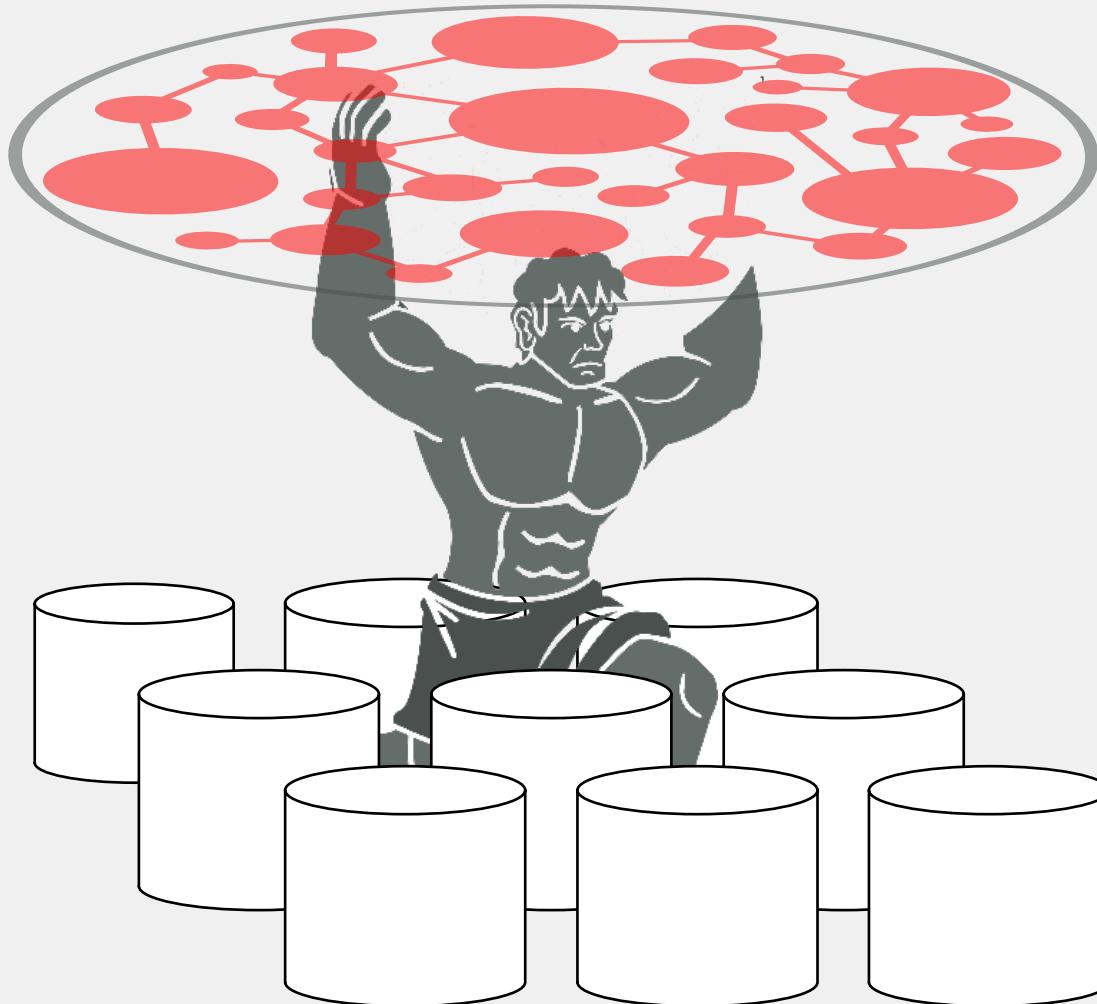
TITAN'S KEY FEATURES



Titan provides...

- ..."infinite size" graphs and "unlimited" users by means of a distributed storage engine.
- ...real-time local traversals (OLTP) and support for global batch processing via Hadoop (OLAP).
- ...distribution via the liberal, free, open source Apache2 license.

BACKEND AGNOSTIC



Cassandra

-OR-

A P A C H E
HBASE

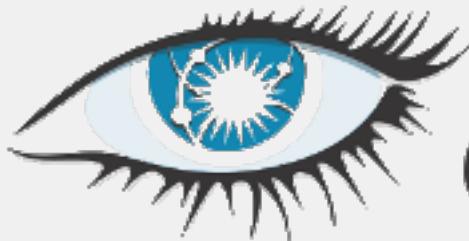
TITAN DISTRIBUTED VIA CASSANDRA



```
titan$ bin/gremlin.sh
```

```
\,,,/
( o o )
-----o00o-(_)-o00o-----
gremlin> conf = new BaseConfiguration();
==>org.apache.commons.configuration.BaseConfiguration@763861e6
gremlin> conf.setProperty("storage.backend", "cassandra");
gremlin> conf.setProperty("storage.hostname", "77.77.77.77");
gremlin> g = TitanFactory.open(conf);
==>titangraph[cassandra:77.77.77.77]
gremlin>
```

INHERITED FEATURES



Cassandra

- Continuously available with no single point of failure.
- No write bottlenecks to the graph as there is no master/slave architecture.
- Built-in replication ensures data is available during machine failure.
- Caching layer ensures that continuously accessed data is available in memory.
- Elastic scalability allows for the introduction and removal of machines.

TITAN DISTRIBUTED VIA HBASE



```
titan$ bin/gremlin.sh
```

\,,,/

(o o)

-----o00o-(_)-o00o-----

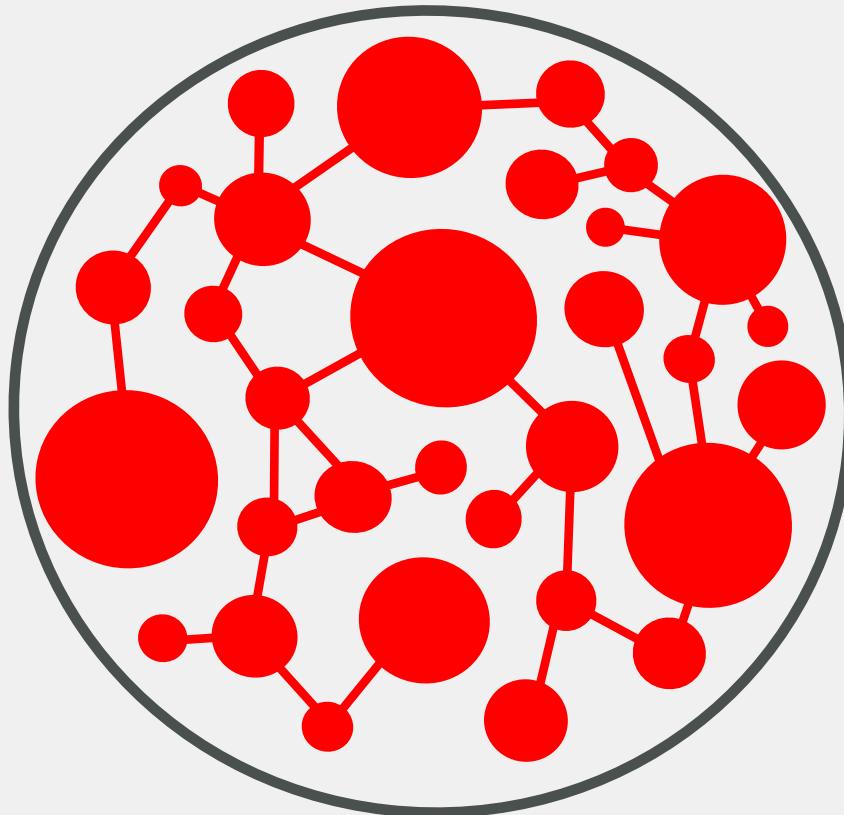
```
gremlin> conf = new BaseConfiguration();
==>org.apache.commons.configuration.BaseConfiguration@763861e6
gremlin> conf.setProperty("storage.backend", "hbase");
gremlin> conf.setProperty("storage.hostname", "77.77.77.77");
gremlin> g = TitanFactory.open(conf);
==>titangraph[hbase:77.77.77.77]
gremlin>
```

INHERITED FEATURES

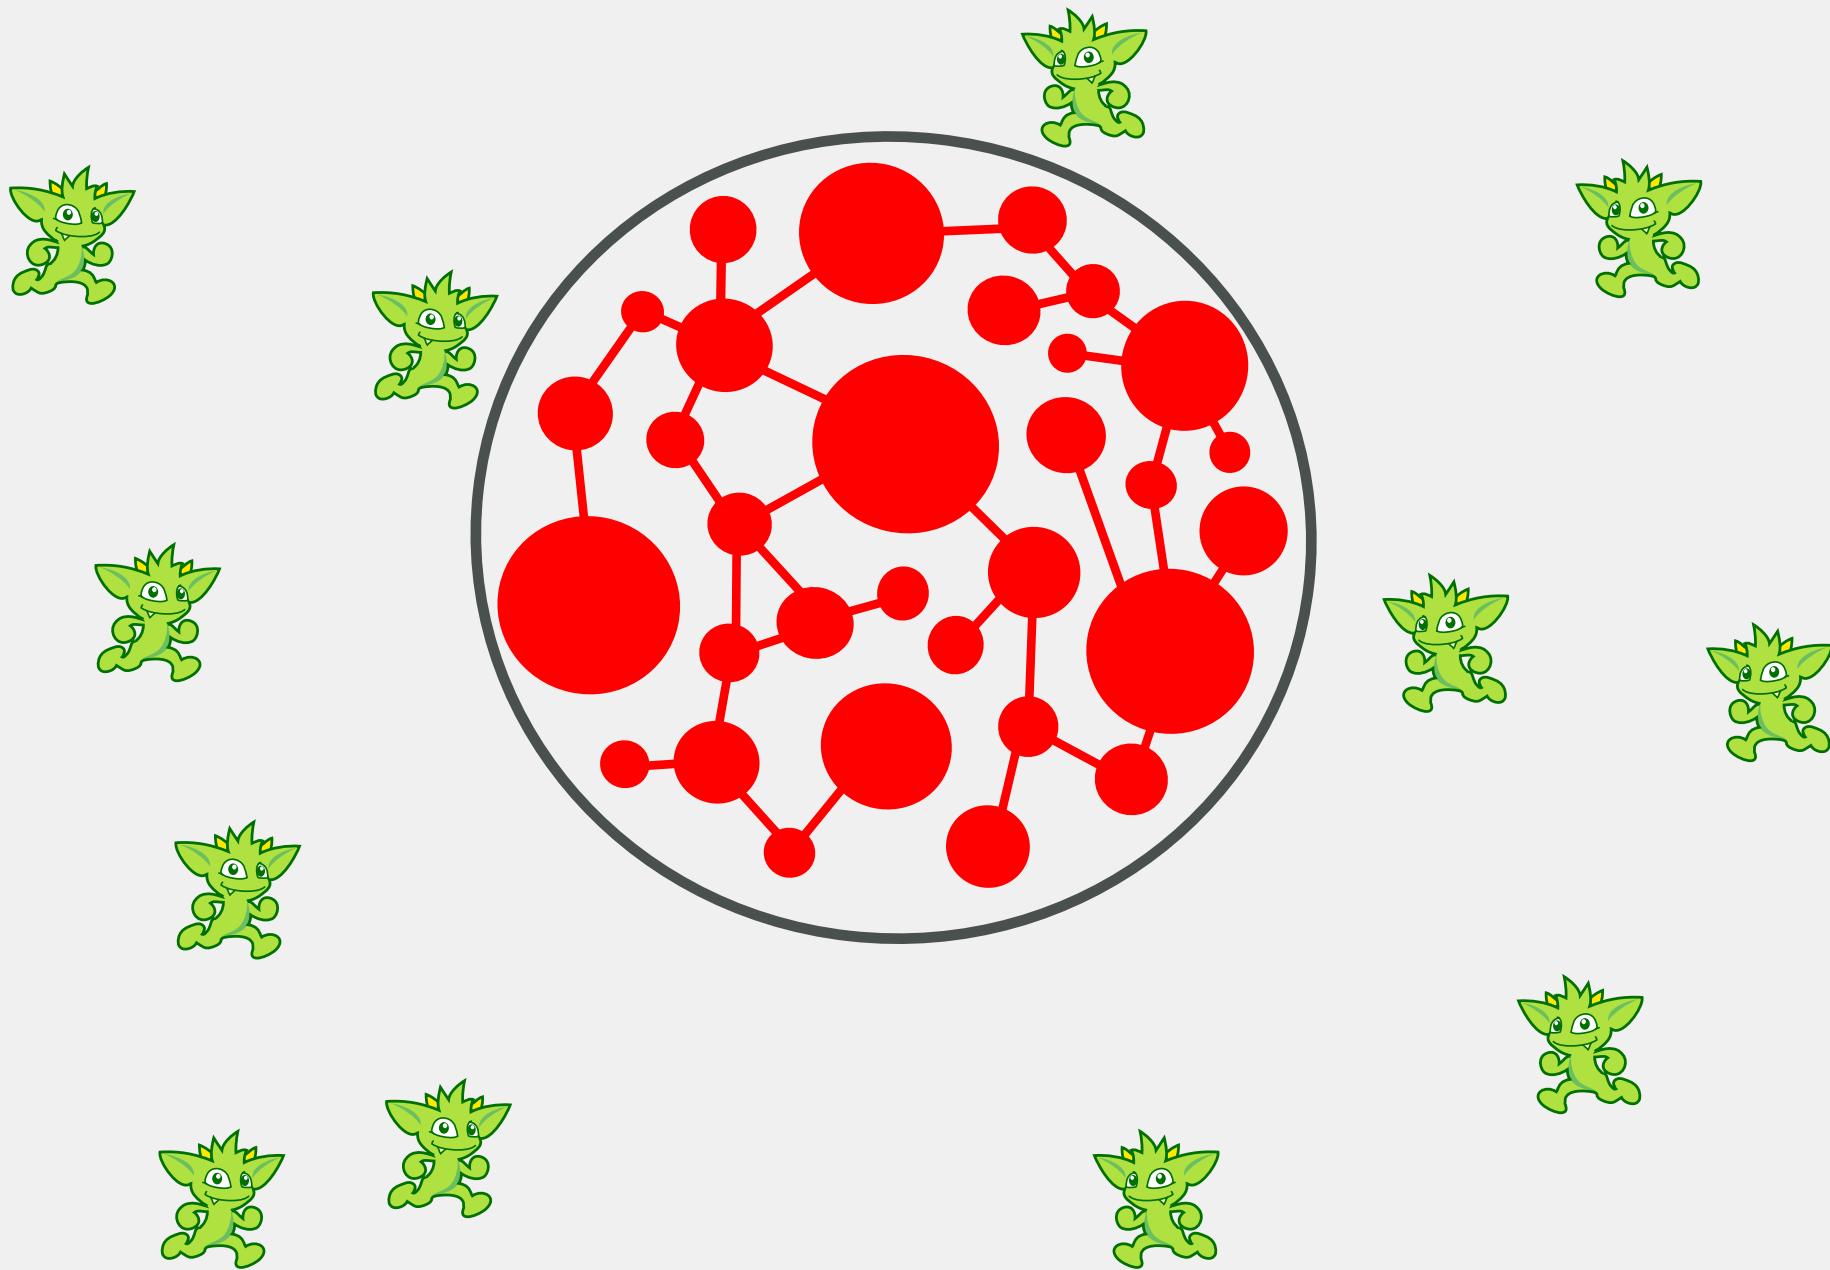


- Strictly consistent reads and writes.
- Linear scalability with the addition of machines.
- Base classes for backing Hadoop MapReduce jobs with HBase tables.
- HDFS-based data replication.
- Generally good integration with the tools in the Hadoop ecosystem.

Titan is all about ...



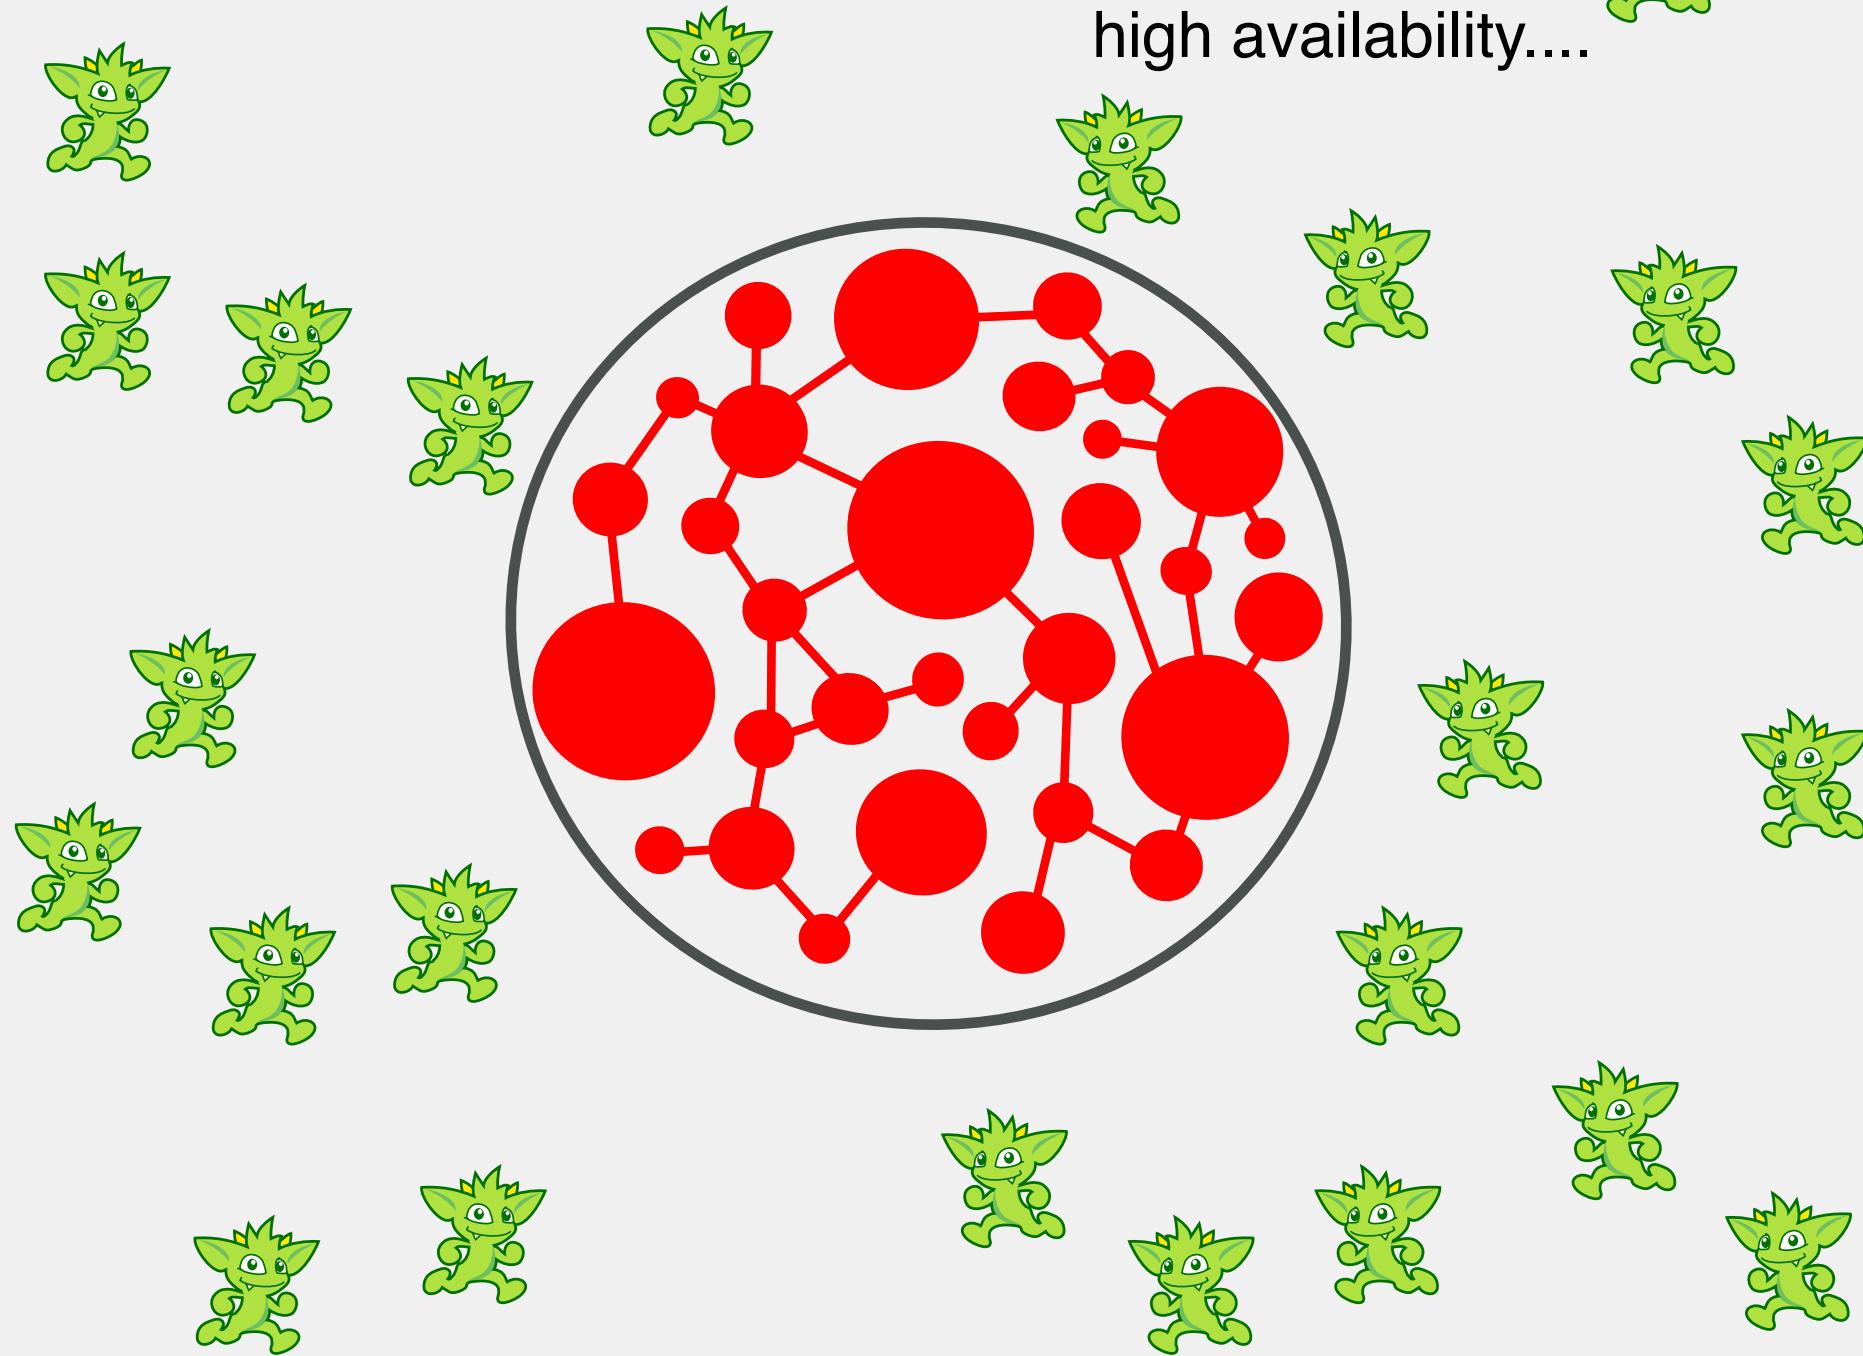
Titan is all about numerous concurrent users...



Titan is all about numerous concurrent users...



high availability....

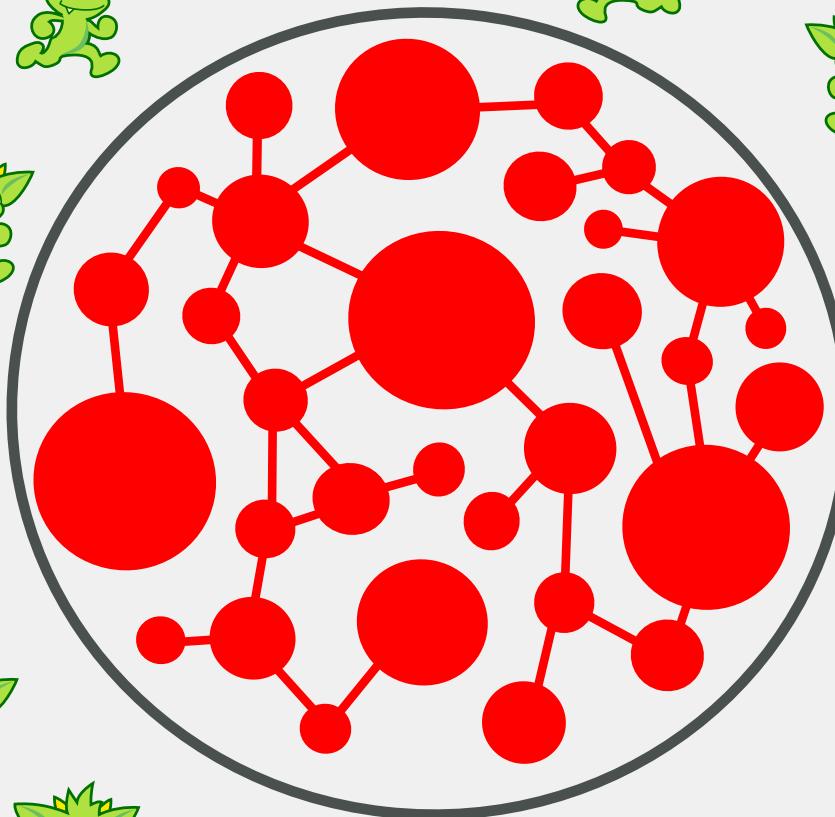


Titan is all about numerous concurrent users...



high availability....

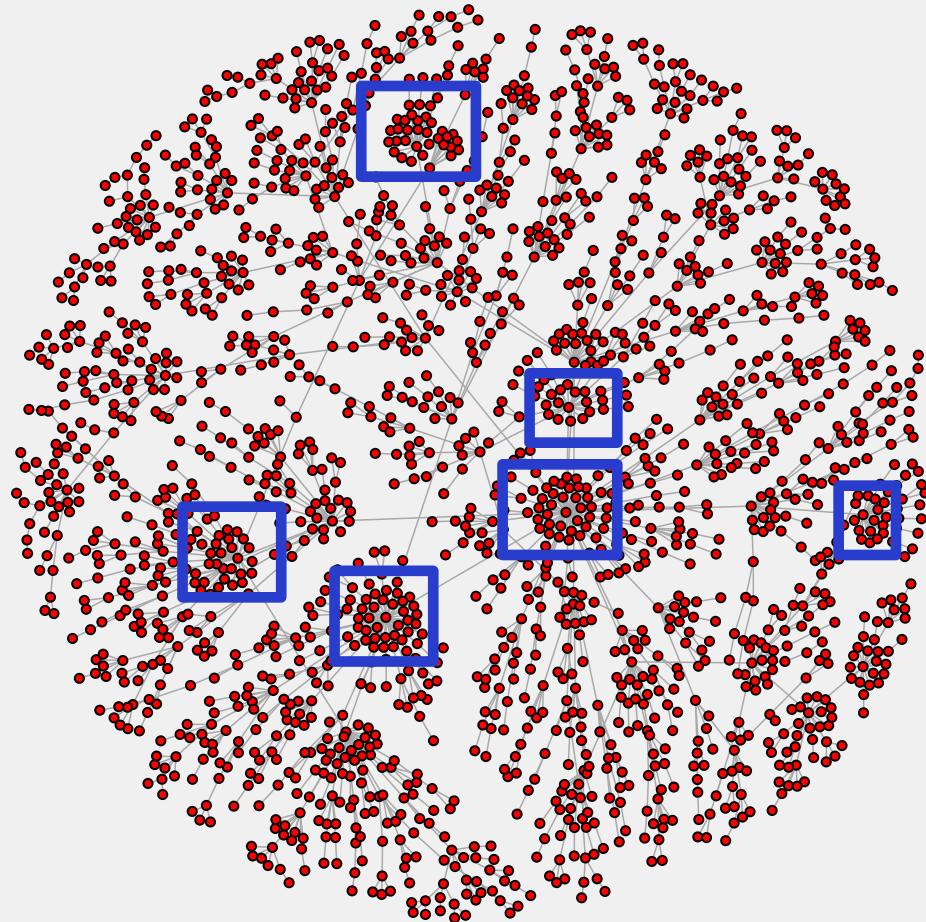
dynamic scalability...



VERTEX-CENTRIC INDICES

THE SUPER NODE PROBLEM

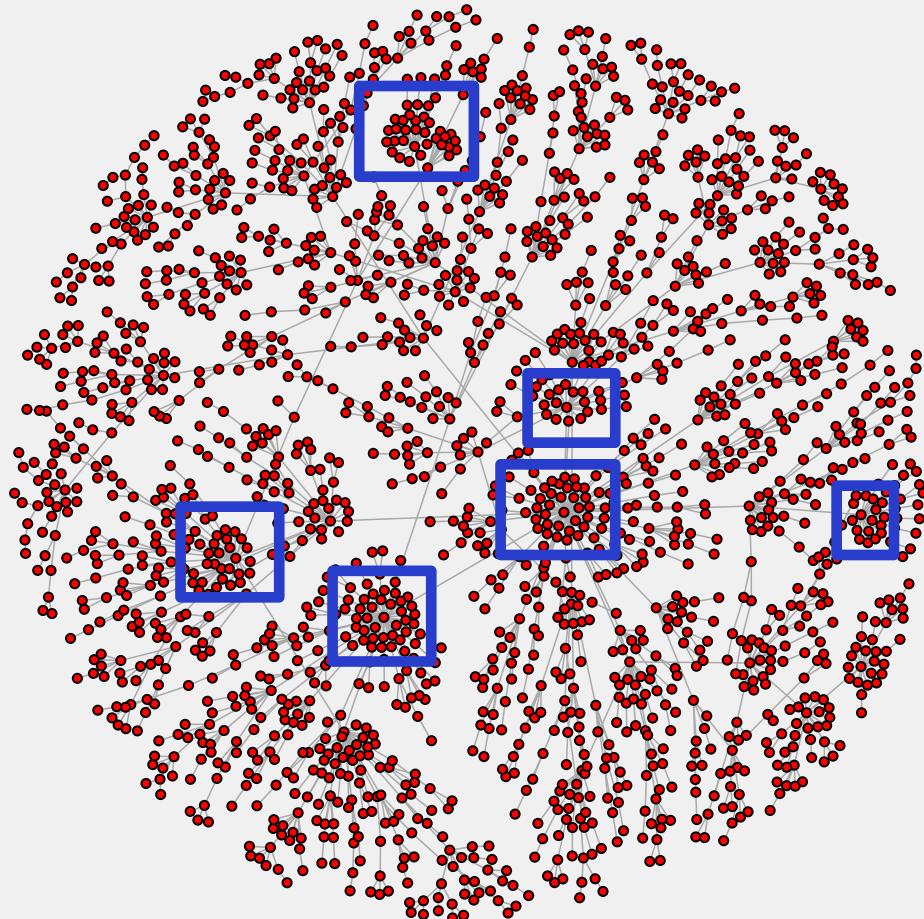
- Natural, real-world graphs contain vertices of high degree.
- Even if rare, their degree ensures that they exist on many paths.
- Traversing a high degree vertex means touching numerous incident edges and potentially touching most of the graph in only a few steps.



VERTEX-CENTRIC INDICES

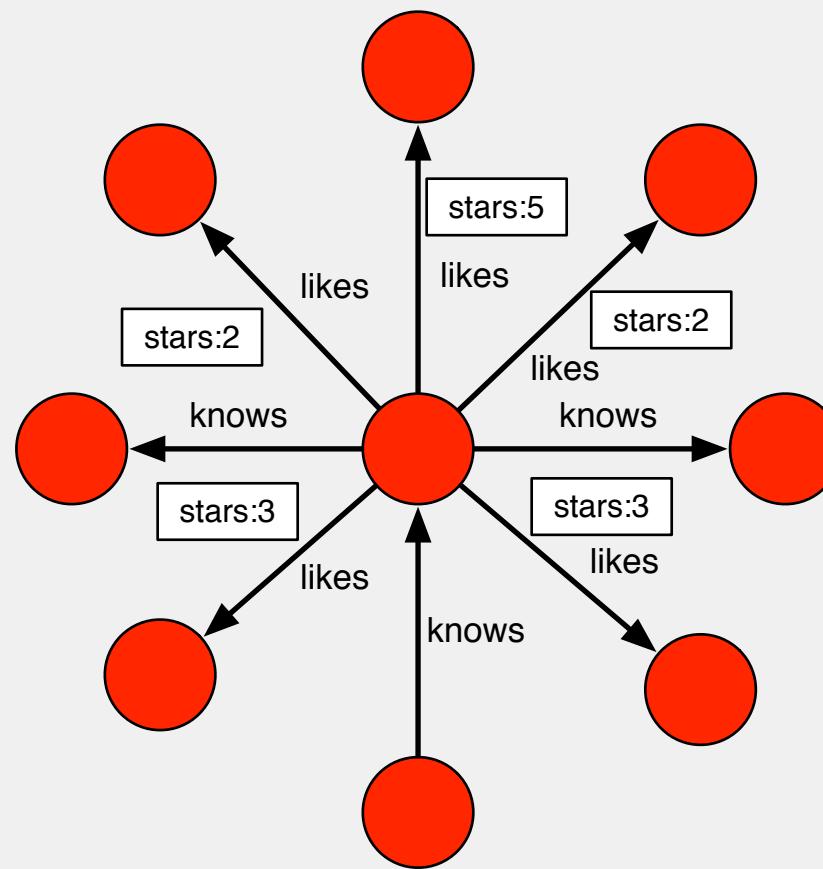
A SUPER NODE SOLUTION

- A "super node" only exists from the vantage point of classic "textbook style" graphs.
- In the world of property graphs, intelligent disk-level filtering can interpret a "super node" as a more manageable low-degree vertex.
- Vertex-centric querying utilizes sort orders for speedy lookup of incident edges with particular qualities.



VERTEX-CENTRIC INDICES PUSHDOWN PREDICATES

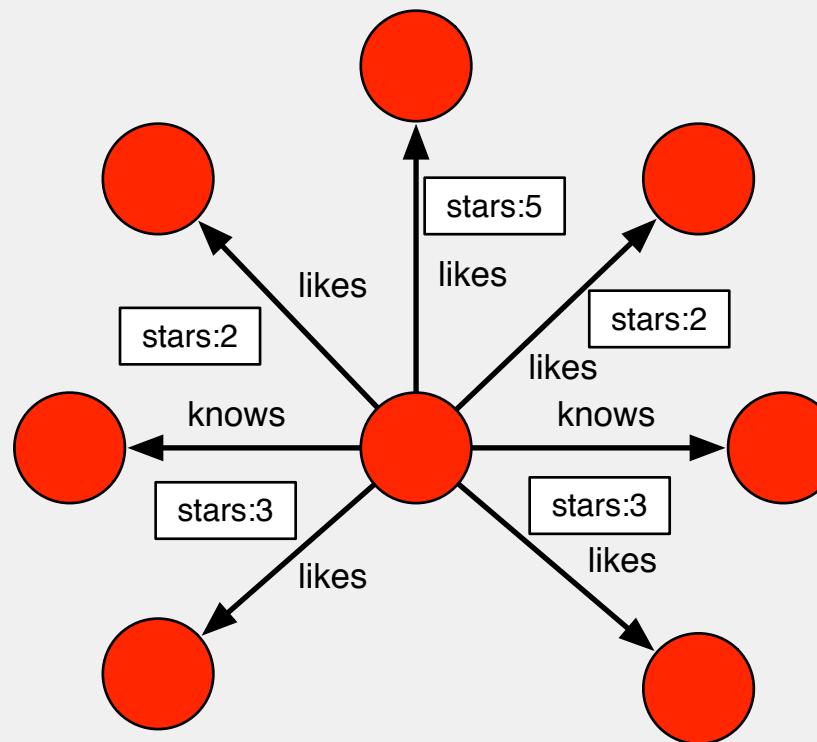
`vertex.bothE()`



8 edges

VERTEX-CENTRIC INDICES PUSHDOWN PREDICATES

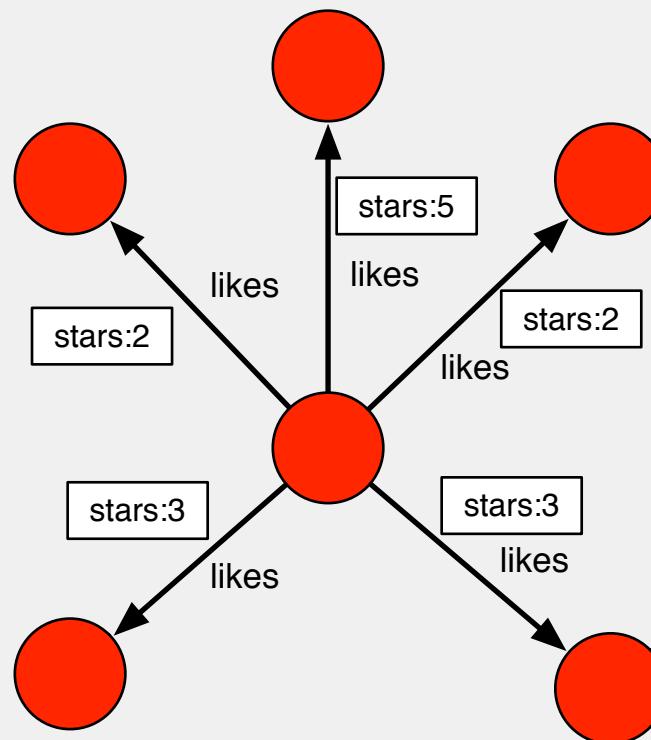
`vertex.outE()`



7 edges

VERTEX-CENTRIC INDICES PUSHDOWN PREDICATES

```
vertex.outE("likes")
```

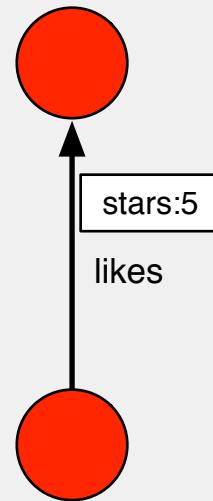


5 edges

VERTEX-CENTRIC INDICES

PUSHDOWN PREDICATES

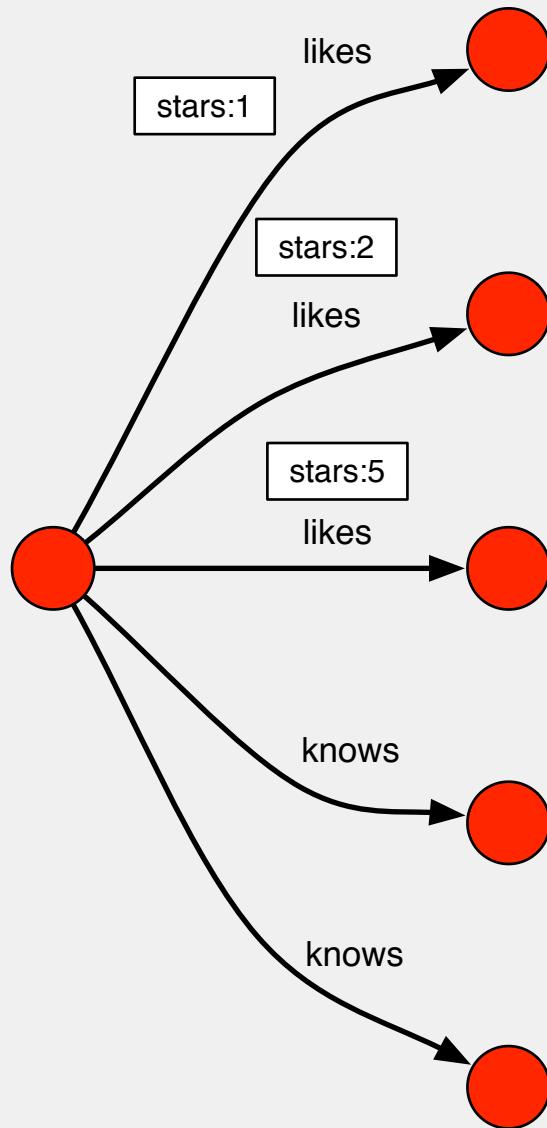
```
vertex.outE("likes").has("stars", 5)
```



1 edge

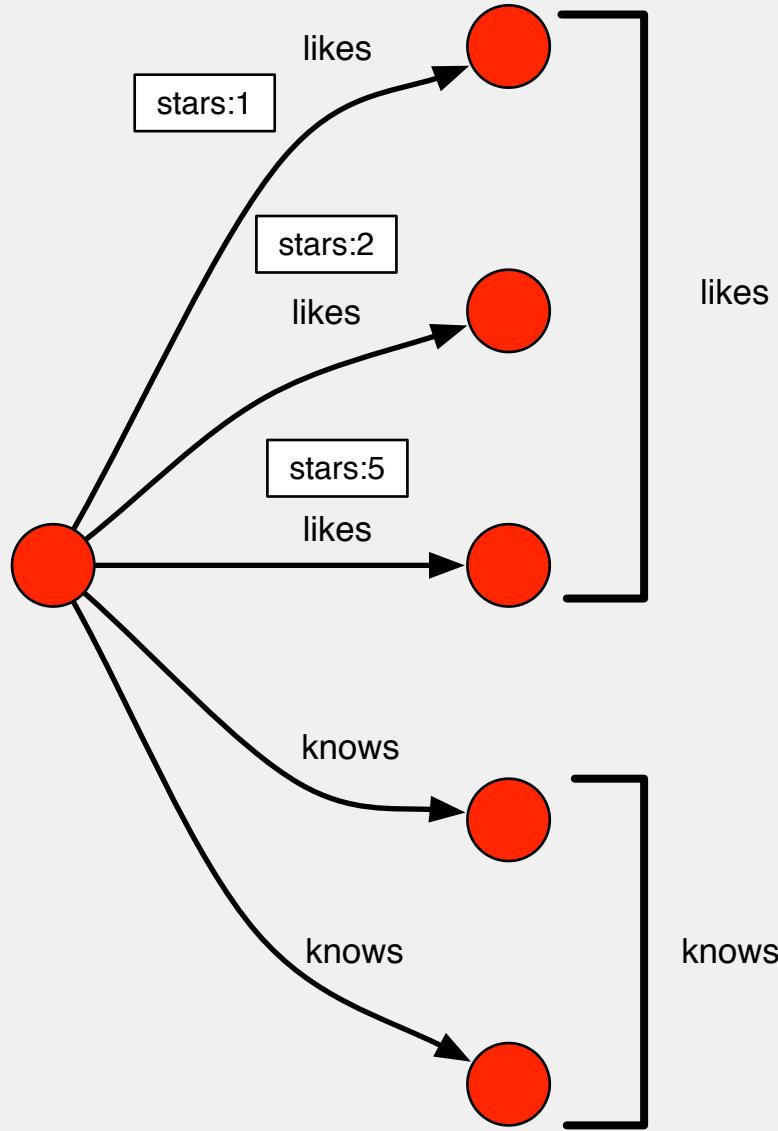
VERTEX-CENTRIC INDICES

DISK-LEVEL SORTING/INDEXING



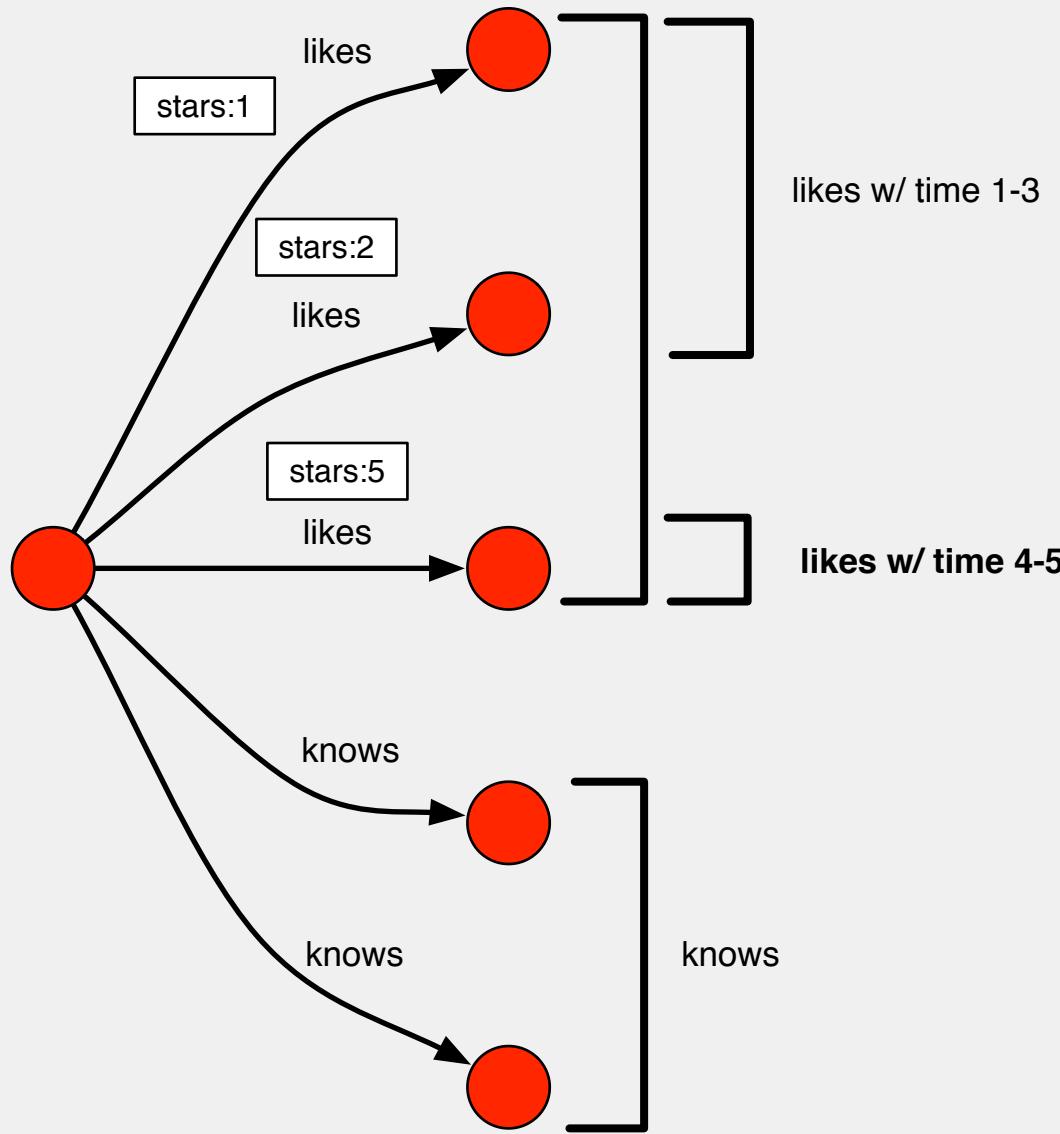
VERTEX-CENTRIC INDICES

DISK-LEVEL SORTING/INDEXING



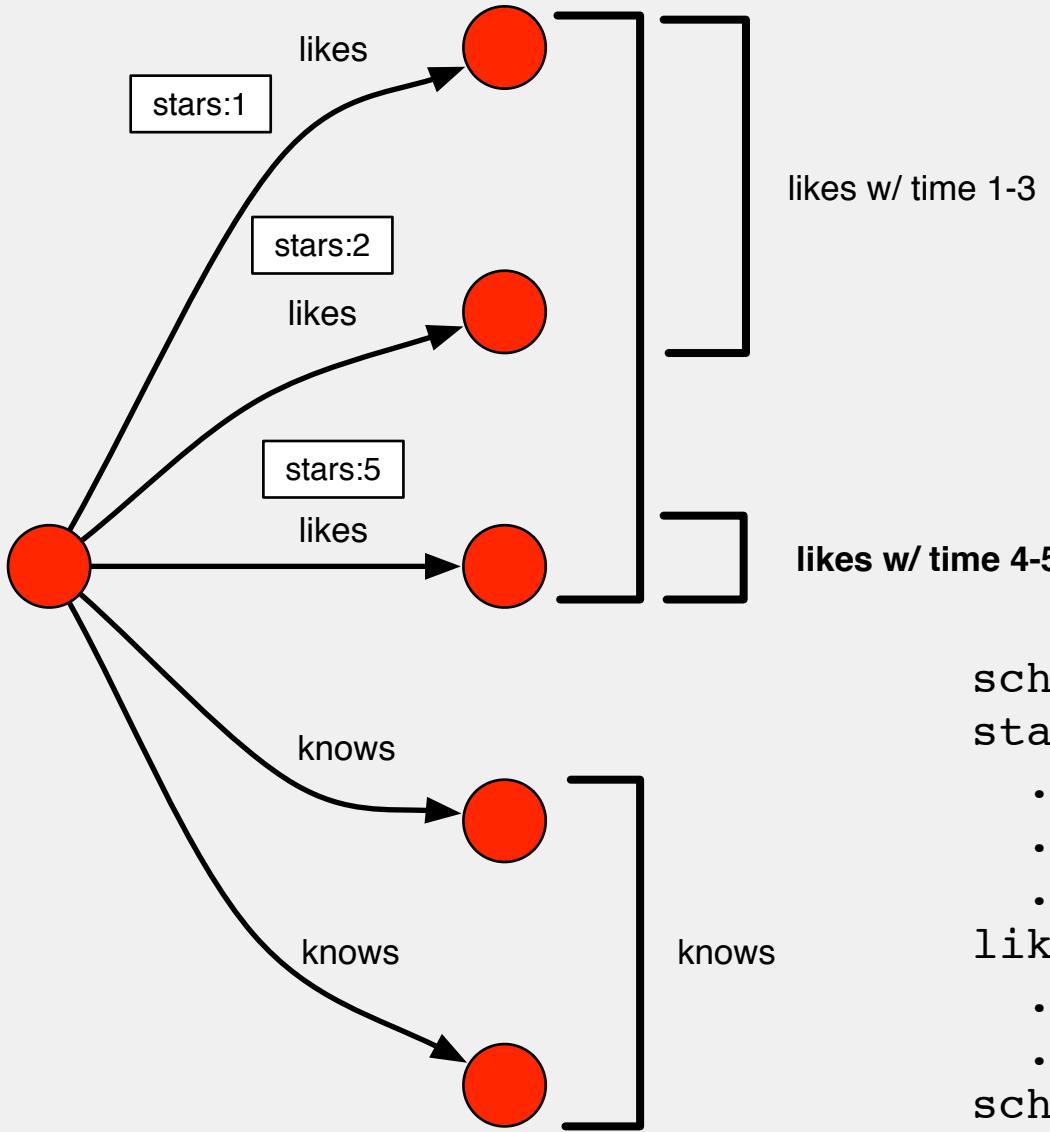
VERTEX-CENTRIC INDICES

DISK-LEVEL SORTING/INDEXING



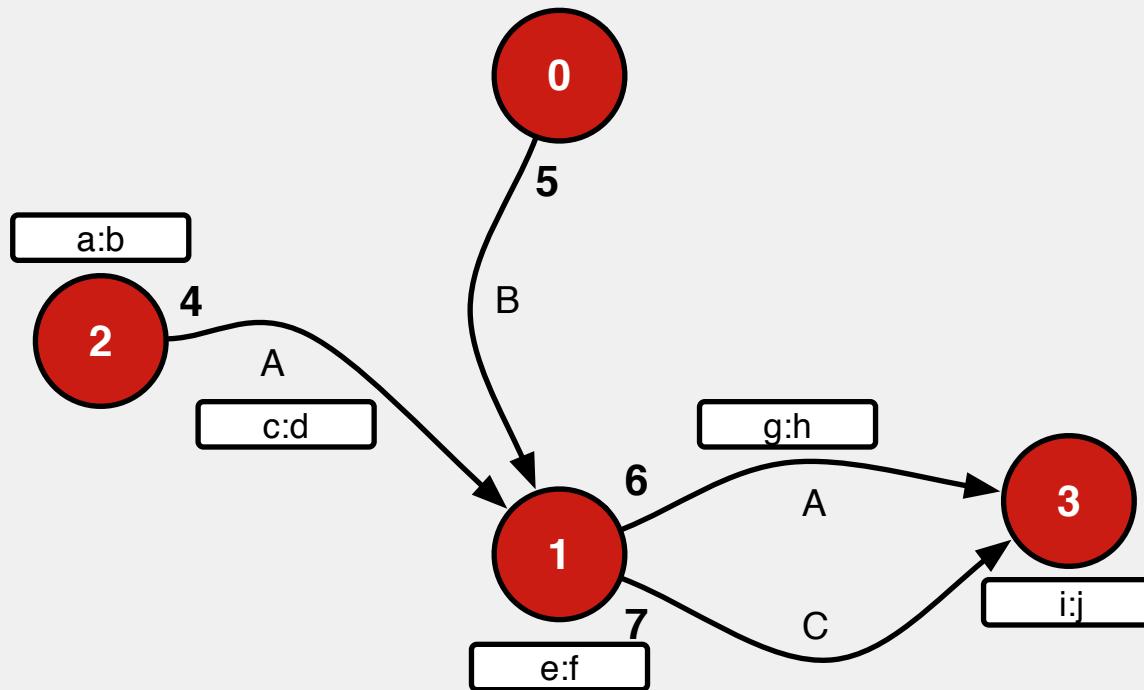
VERTEX-CENTRIC INDICES

DISK-LEVEL SORTING/INDEXING



```
schema = g.openManagementSystem()
stars = schema
    .makePropertyKey("stars")
    .dataType(Integer.class)
    .make()
likes = schema
    .makeEdgeLabel('likes')
    .make()
schema.buildEdgeIndex(likes,
    'timeIdx', BOTH, DESC, stars)
```

THE TITAN OLAP STORY



incoming edges

vertex	edge	edge	edge	edge
id props	id props label id	id label id	id props label id	id label id

1 e:f

4 c:d A 2

5 B 0

6 g:h A 3

7 C 3

AN ADJACENCY LIST

0

1

3

4

5

6

7

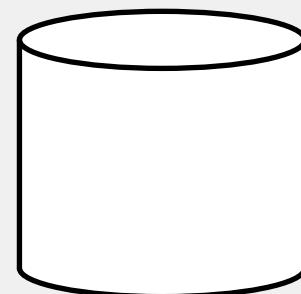
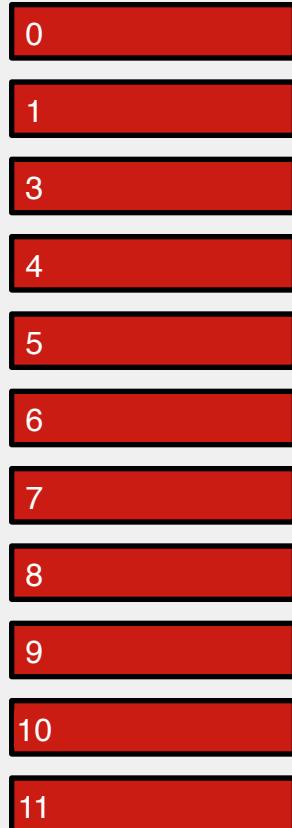
8

9

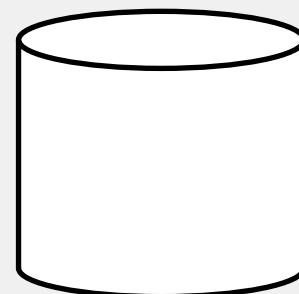
10

11

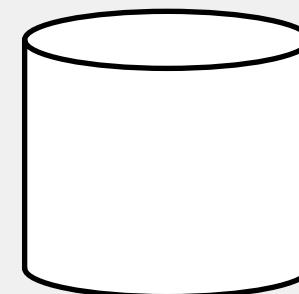
AN ADJACENCY LIST + CLUSTER



127.0.0.2

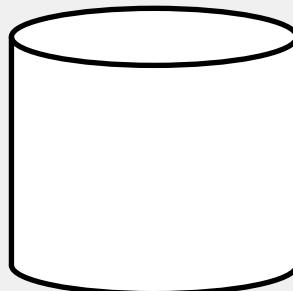
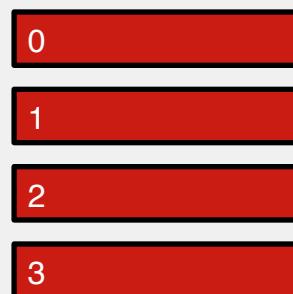


127.0.0.3

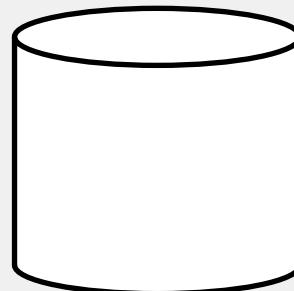


127.0.0.4

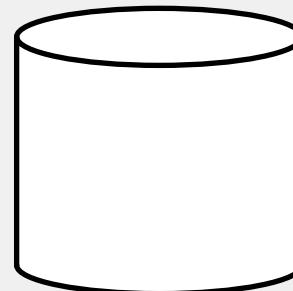
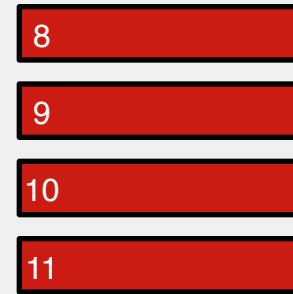
A DISTRIBUTED ADJACENCY LIST



127.0.0.2



127.0.0.3



127.0.0.4

HADOOP

Hadoop is a distributed computing platform composed of two key components:

HDFS:

A distributed file system that stores arbitrarily large files within a cluster.

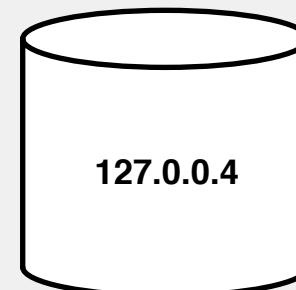
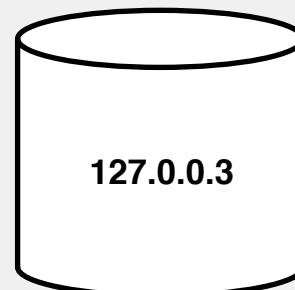
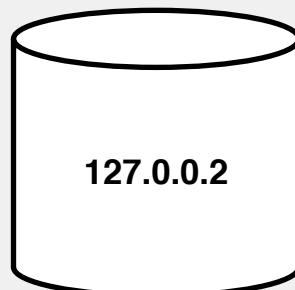
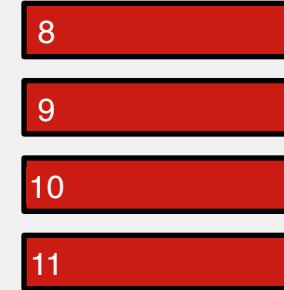
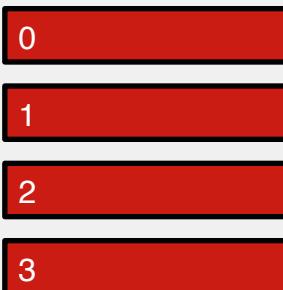
MapReduce:

A parallel functional computing model for key/value pair data.



FAUNUS AND HADOOP

Process



Structure



Faunus provides graph input/output formats (structure) and a traversal language for graphs (process).

THE TITAN OLAP STORY



Serial Key/Value Data Structure



0	██████████
1	██████████
3	██████████
4	██████████
5	██████████
6	██████████
7	██████████
8	██████████
9	██████████
10	██████████
11	██████████



Indexed Key/Indexed Value Data Structure

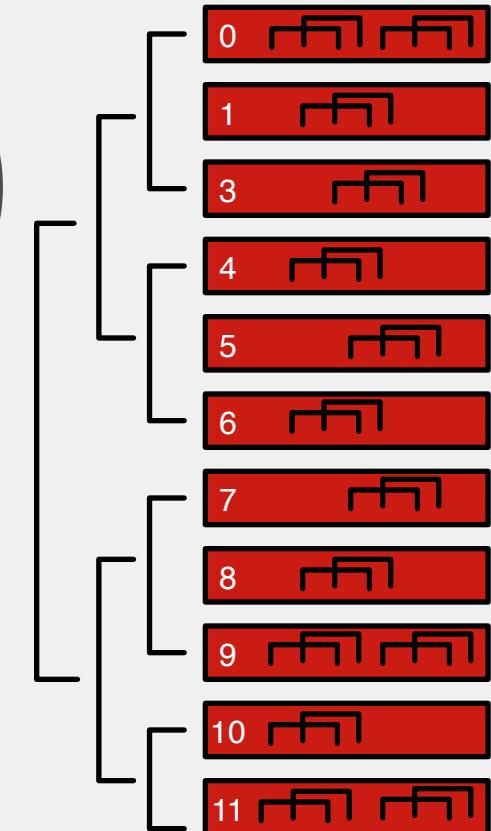
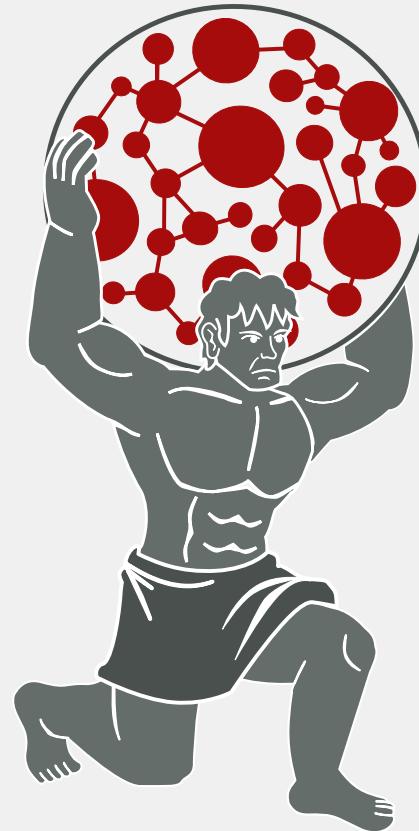
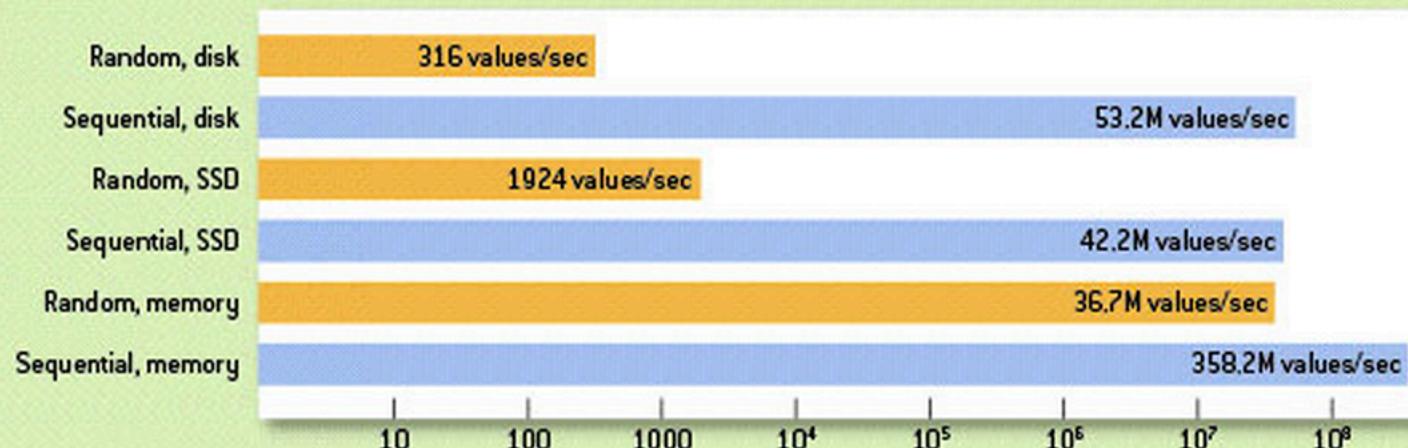


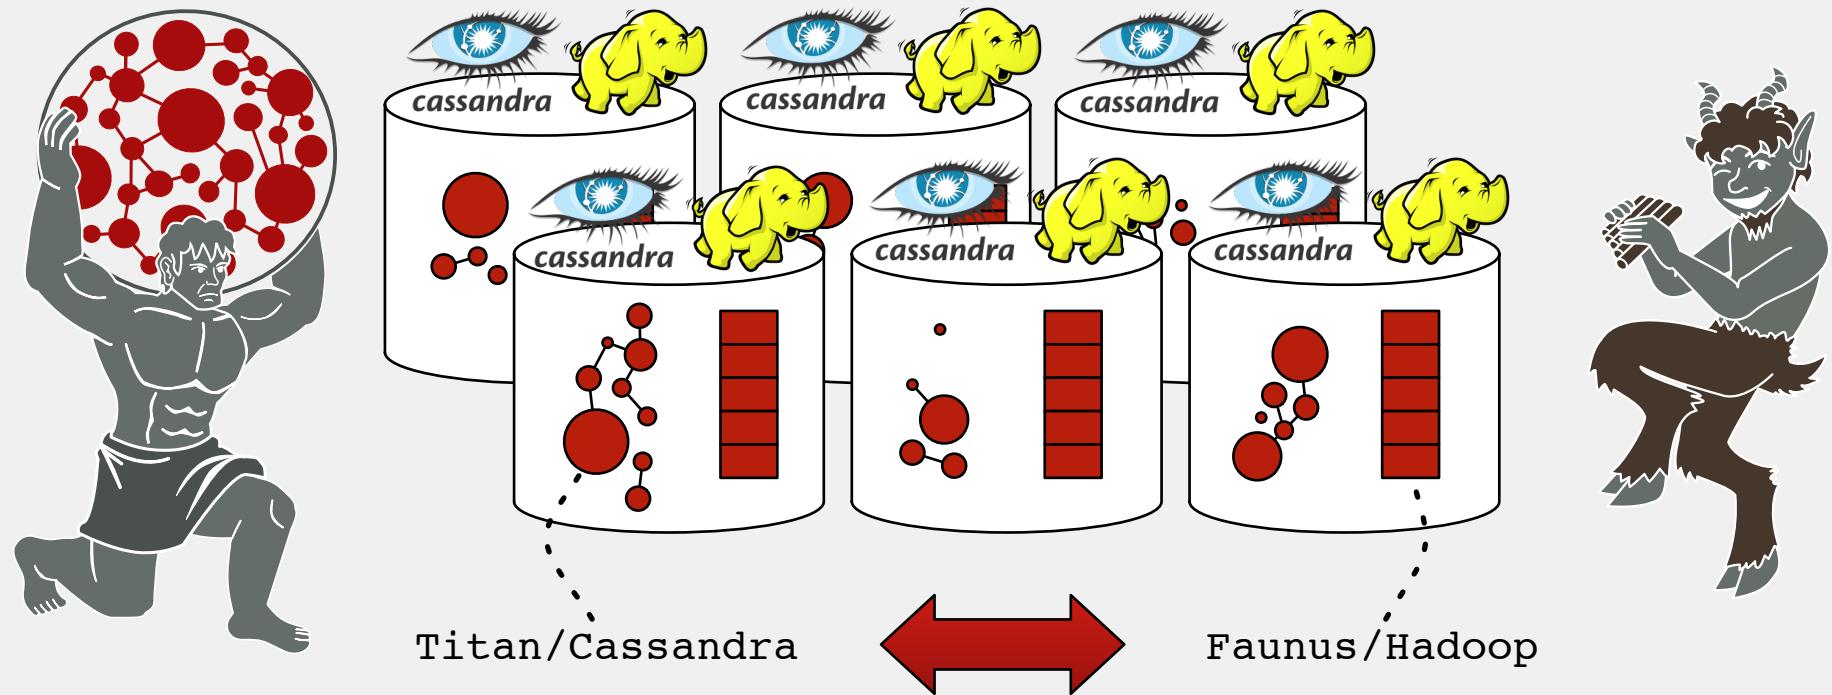
FIGURE
3

Comparing Random and Sequential Access in Disk and Memory



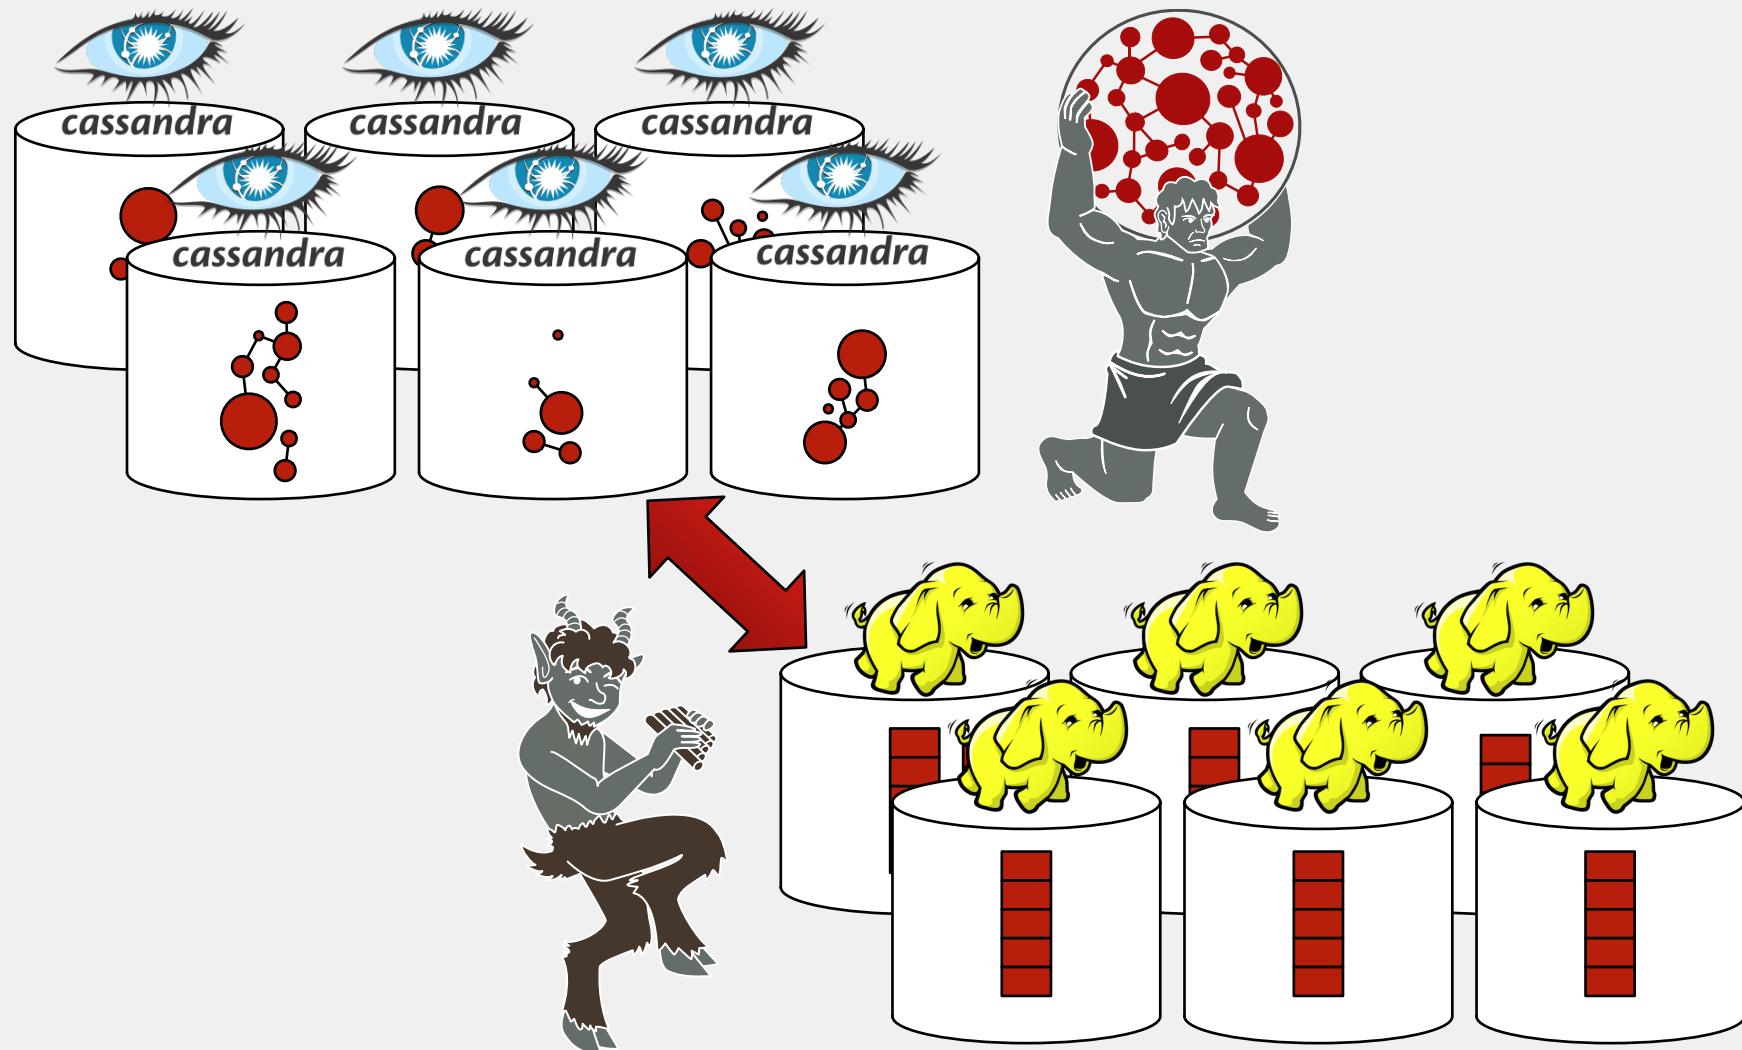
Note: Disk tests were carried out on a freshly booted machine (a Windows 2003 server with 64-GB RAM and eight 15,000-RPM SAS disks in RAID5 configuration) to eliminate the effect of operating-system disk caching. SSD test used a latest-generation Intel high-performance SATA SSD.

INTRA-CLUSTER CONFIGURATION



- Data is processed on the machine where it is located.
- Limited network communication.

INTER-CLUSTER CONFIGURATION

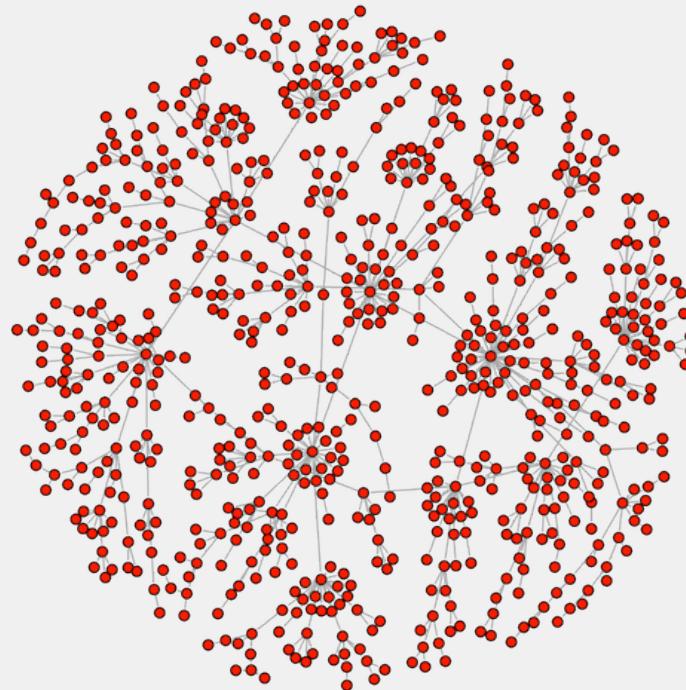


- Graph data is offloaded to another cluster.
- Repeated analysis does not interfere with production graph database.

THE TITAN OLAP STORY



Titan currently provides a Hadoop OLAP engine
Titan-Hadoop (nicknamed **Faunus**)

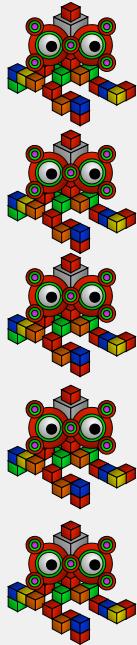


Titan 1.0 will also provide an in-memory, off head, single-machine OLAP engine.
Titan-Memory (nicknamed **Fulgora**)

PART 3: THE PATH TO TINKERPOP3



TINKERPOP3 FEATURES



- A standard, vendor-agnostic graph query language.
- Both OLTP and OLAP engines for evaluating queries.
- A way for developers to create custom, domain specific variants.
- A server system for pushing queries to the server for evaluation.
- A vertex-centric computing framework for distributed graph algorithms.

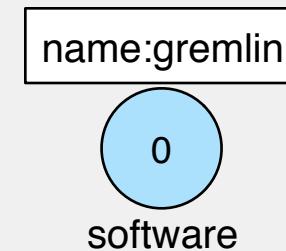




gremlin>

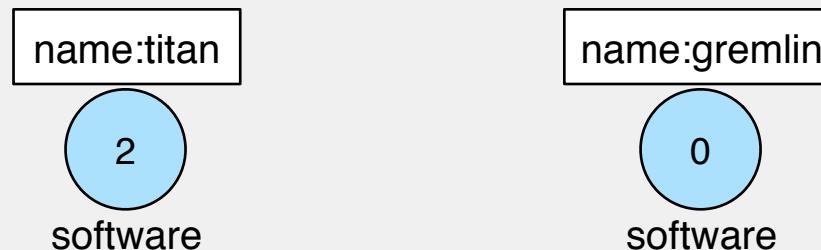


```
gremlin> gremlin = g.addVertex(label,'software','name','gremlin')  
==>v[0]
```



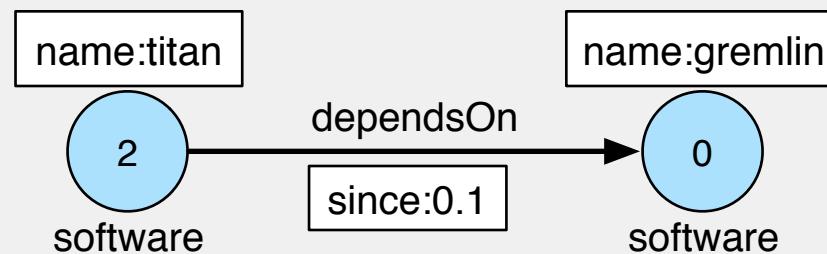


```
gremlin> gremlin = g.addVertex(label,'software','name','gremlin')
==>v[0]
gremlin> titan = g.addVertex(label,'software','name','titan')
==>v[2]
```





```
gremlin> gremlin = g.addVertex(label,'software','name','gremlin')
==>v[0]
gremlin> titan = g.addVertex(label,'software','name','titan')
==>v[2]
gremlin> titan.addEdge('dependsOn',gremlin,'since','0.1')
==>e[4][2-dependson->0]
```





Gremlin is a style of graph traversal that is heavily influenced by functional programming.

```
gremlin> gremlin = g.addVertex(label,'software','name','gremlin')
==>v[0]
gremlin> titan = g.addVertex(label,'software','name','titan')
==>v[2]
gremlin> titan.addEdge('dependsOn',gremlin,'since','0.1')
==>e[4][2-dependson->0]
```

//////////////////////////////

```
gremlin> g.V
==>v[0]
==>v[2]
gremlin> g.V.has('name','titan')
==>v[2]
gremlin> g.V.has('name','titan').outE
==>e[4][2-dependson->0]
gremlin> g.V.has('name','titan').outE.since
==>0.1
gremlin> g.V.has('name','titan').out.name
==>gremlin
```

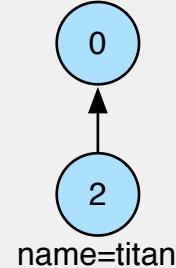
```
g.V.has('name','titan').out.name
```

What are the names of the vertices outgoing adjacent to Titan?



```
g.V.has('name','titan').out.name
```

name=gremlin



For the graph 'g'...



Traversals are defined using fluent, function composition.

```
g.V.has('name','titan').out.name
```

name=gremlin



name=titan

...get all the vertices...



Objects flow through the traversal in a lazy evaluation manner.

```
g.V.has('name','titan').out.name
```



...filter all vertices that don't have the name 'titan'...



The general functions are map(), flatMap(), filter(), sideEffect(), and branch().

```
g.V.has('name','titan').out.name
```



...get all outgoing adjacent vertices...



Any programming language can provide a Gremlin implementation.

```
g.V.has('name','titan').out.name
```

gremlin

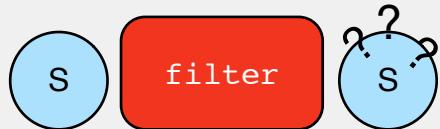
...get the names of those vertices...



Any graph system vendor can provide a binding to Gremlin.

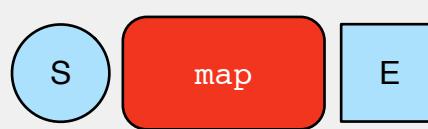
THE GENERIC STEPS

$$S \rightarrow S \cup \emptyset$$



has()
retain()
except()
interval()
simplePath()
dedup()
timeLimit()
...

$$S \rightarrow E$$



back()
fold()
valueMap()
match()
orderBy()
shuffle()
...

$$S \rightarrow E^*$$



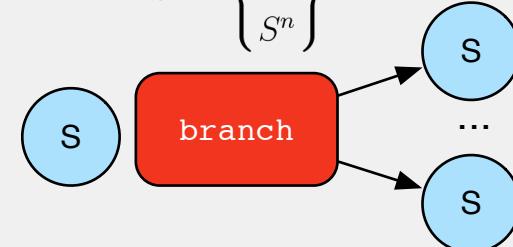
out()
outE()
inV()
both()
values()
unfold()
...

$$S \xrightarrow{A} S$$



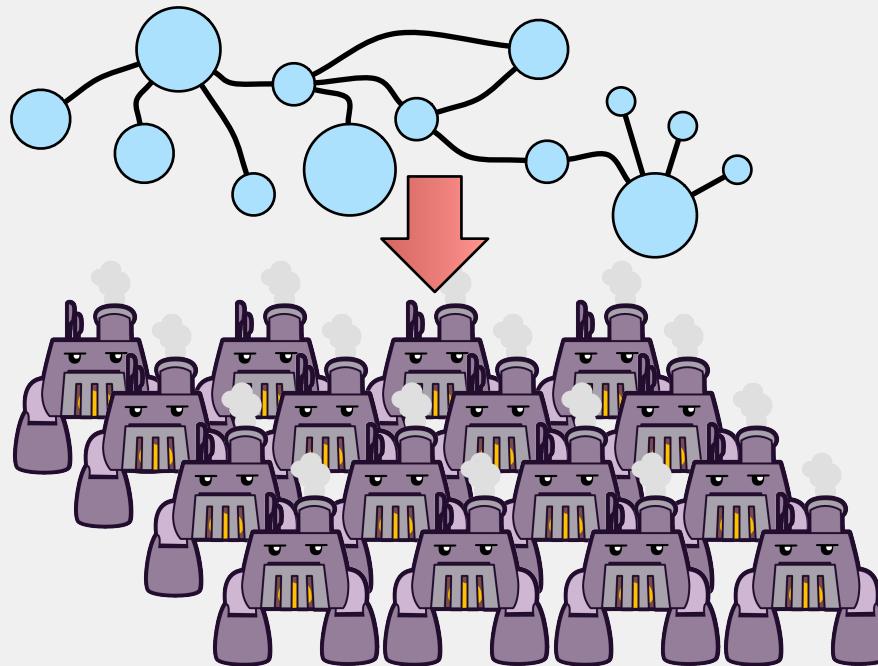
aggregate()
store()
groupCount()
groupBy()
tree()
count()
subgraph()
...

$$S \rightarrow \left\{ \begin{matrix} S^1 \\ \dots \\ S^n \end{matrix} \right\}$$



jump()
until()
choose()
union()
...

TINKERPOP'S OLAP STORY

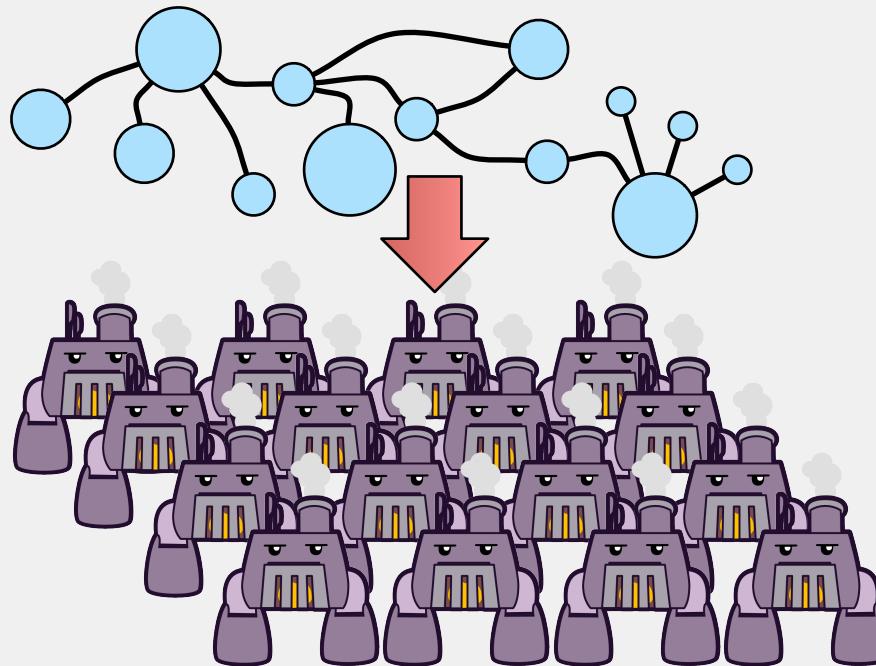


Execute parallel, distributed Gremlin traversals.

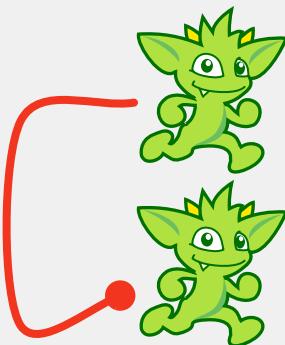


Execute vertex-centric message passing algorithms.

TINKERPOP'S OLAP STORY

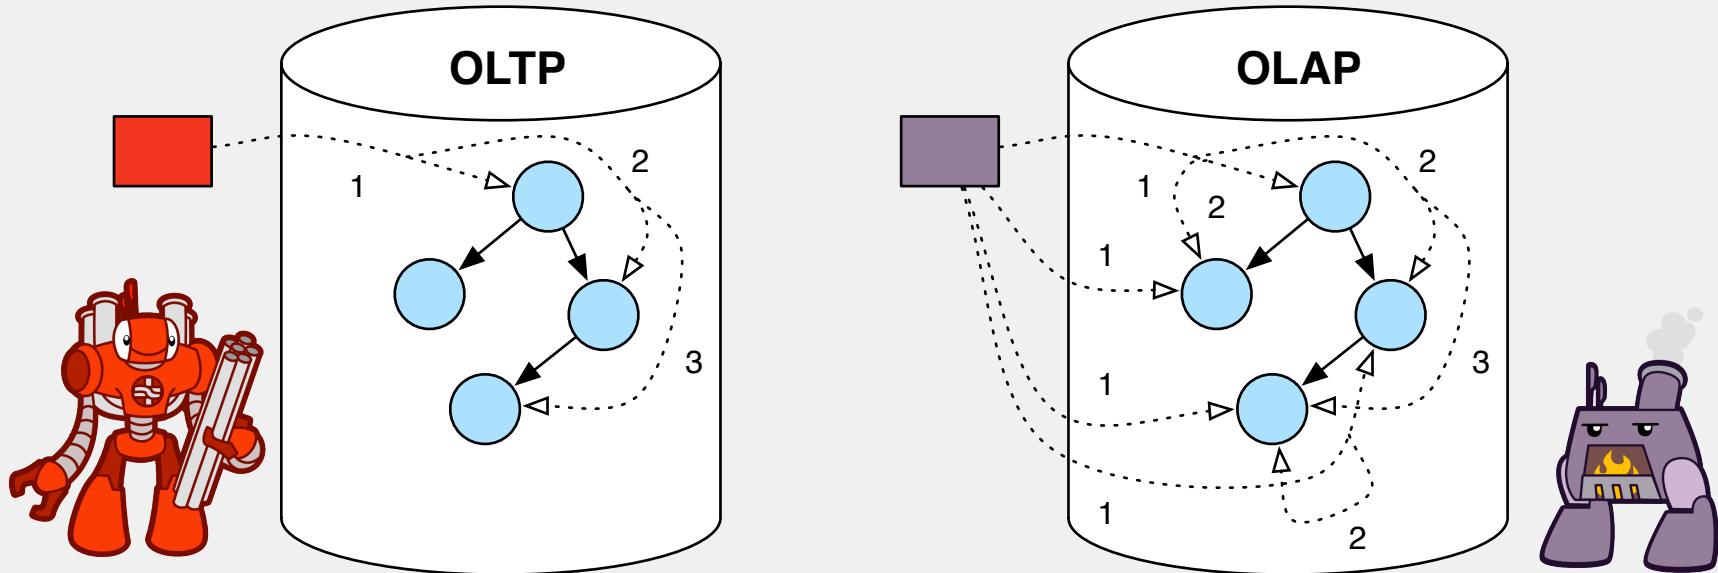


traversers are the messages!



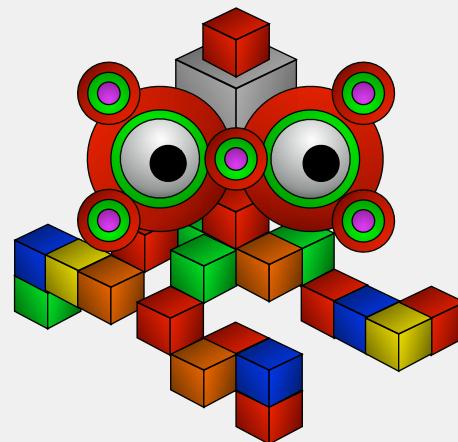
Execute parallel, distributed Gremlin traversals.

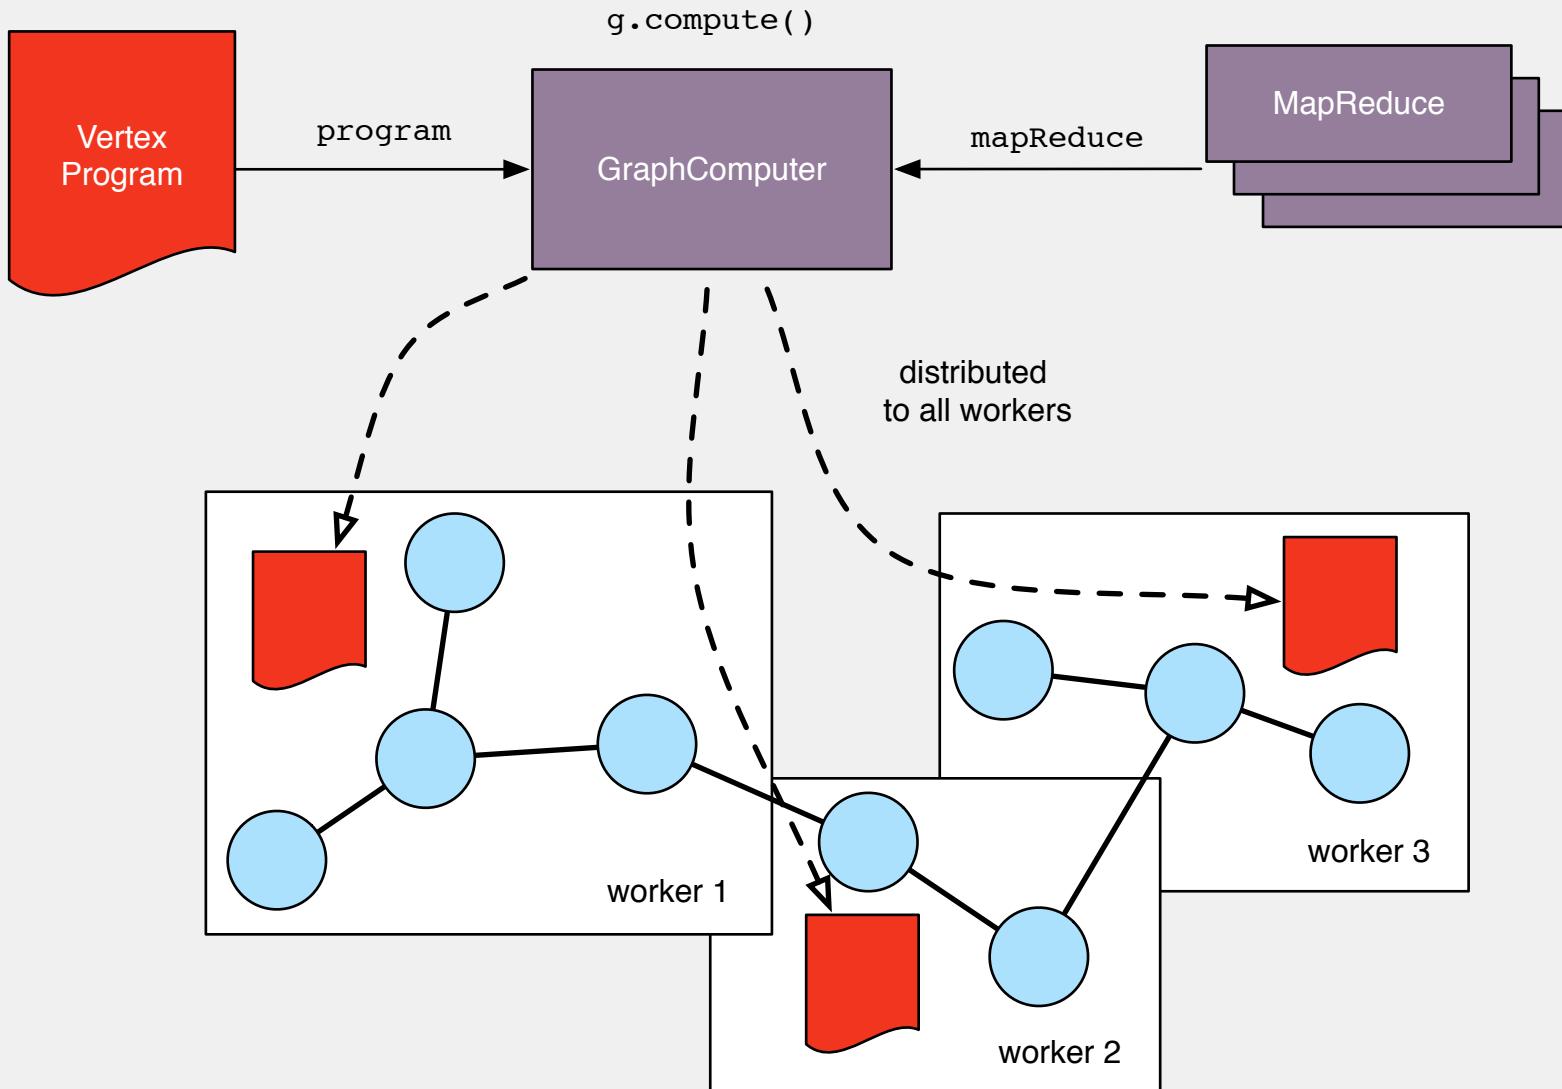
Execute vertex-centric message passing algorithms.



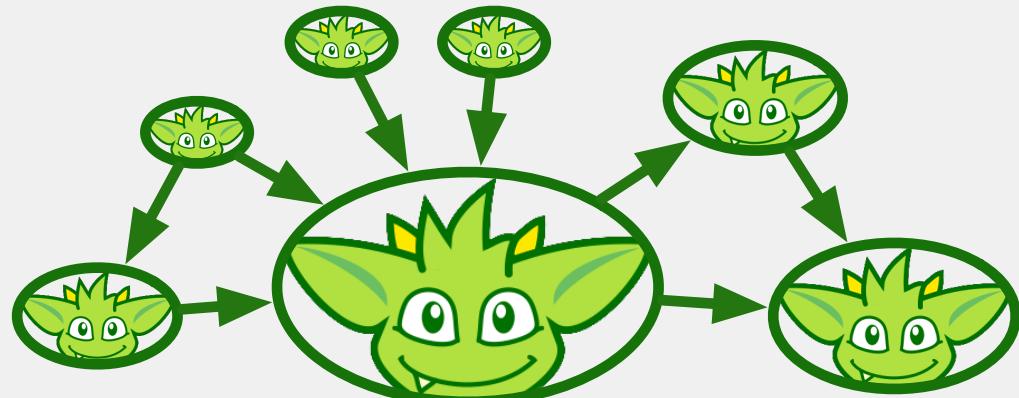
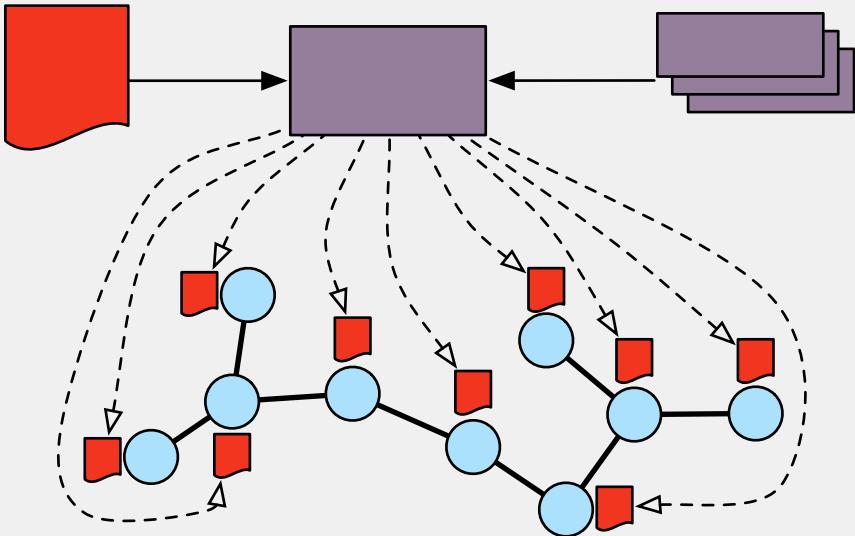
- ✓ Real-time queries.
- ✓ Touching a subset of the graph.
- ✓ Numerous concurrent queries.
- ✓ Good for local graph mutations.

- ✓ Long running queries.
- ✓ Touching the entire of the graph.
- ✓ Single user.
- ✓ Good for bulk loading/mutations.





1. Logically distribute the vertex program to all the vertices in the graph.
2. The vertices execute the vertex program sending and receiving messages on each iteration.
3. Any number of MapReduce jobs can follow to aggregate the computed data in the graph.

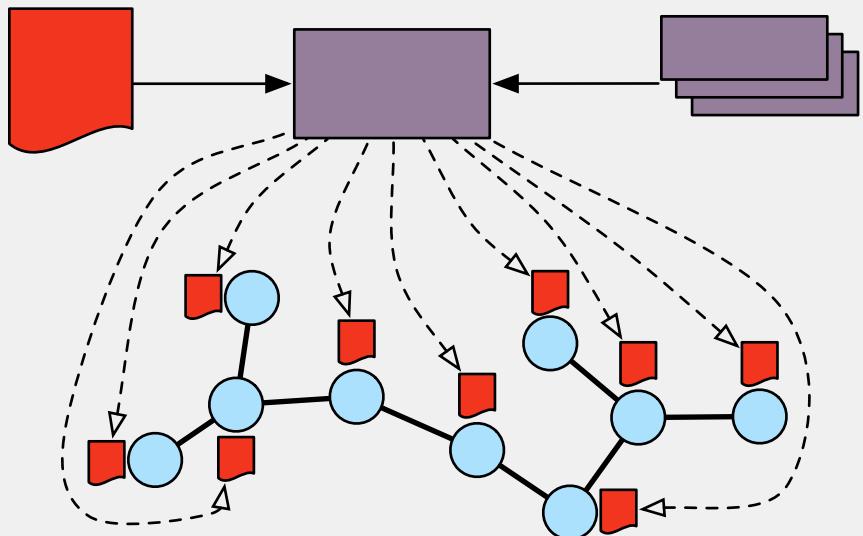


```

gremlin> result = g.compute().program(PageRankVertexProgram).submit()
==>result[tinkergraph[vertices:6 edges:6],memory[size:0]]
gremlin> result.memory.iteration
==>30
gremlin> result.graph
==>tinkergraph[vertices:6 edges:6]
gremlin> result.graph.v.map{ [it.id,it.pageRank]}
==>[1, 0.1500000000000002]
==>[2, 0.1925000000000003]
==>[3, 0.4018125]
==>[4, 0.1925000000000003]
==>[5, 0.2318125000000003]
==>[6, 0.1500000000000002]

```

* Minor tweaks to example to make it fit nicely on a single line.



```
gremlin> g.V.both.both.name.groupCount()
==>[ripple:3, peter:3, vadas:3, josh:7, lop:7, marko:7]
```

OLTP

```
gremlin> g.V.both.both.name.groupCount().submit(g.compute())
==>[ripple:3, peter:3, vadas:3, josh:7, lop:7, marko:7]
```

OLAP

MULTIPLE GRAPH COMPUTERS FOR THE SAME GRAPH

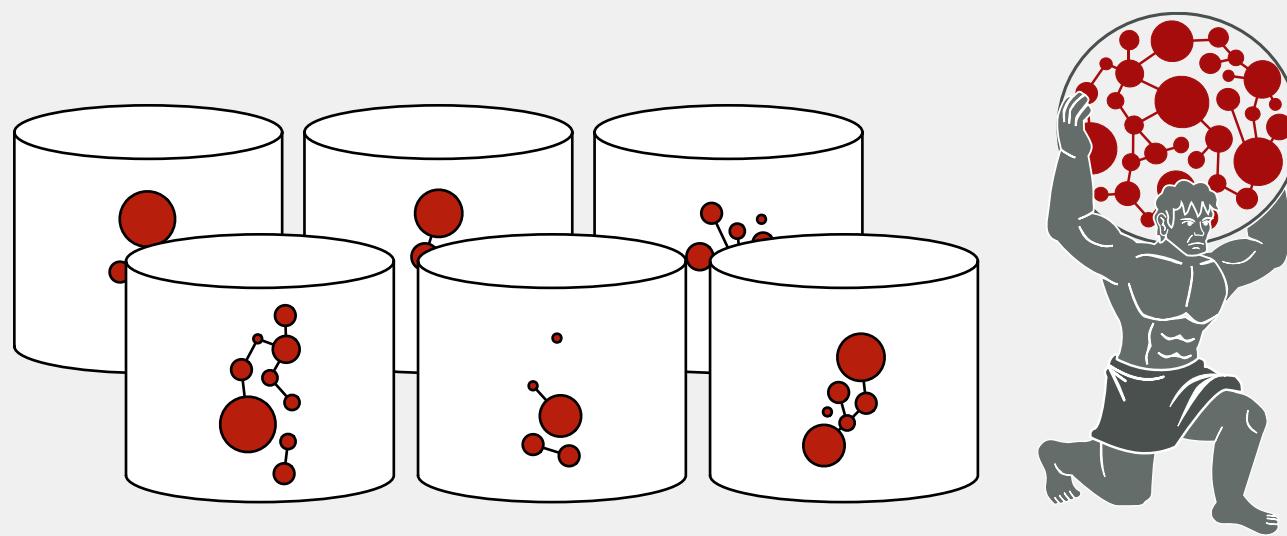


```
gremlin> g = TitanFactory.open(conf)
==>titangraph[66.77.88.99]
gremlin> g.compute(FaunusGraphComputer)
          .program(PageRankVertexProgram)
          .submit()
==>result[faunusgraph,memory[size:0]]
```

```
gremlin> g = TitanFactory.open(conf)
==>titangraph[66.77.88.99]
gremlin> g.compute(FulgoraGraphComputer)
          .program(PageRankVertexProgram)
          .submit()
==>result[fulgoragraph,memory[size:0]]
```



GREMLIN SERVER

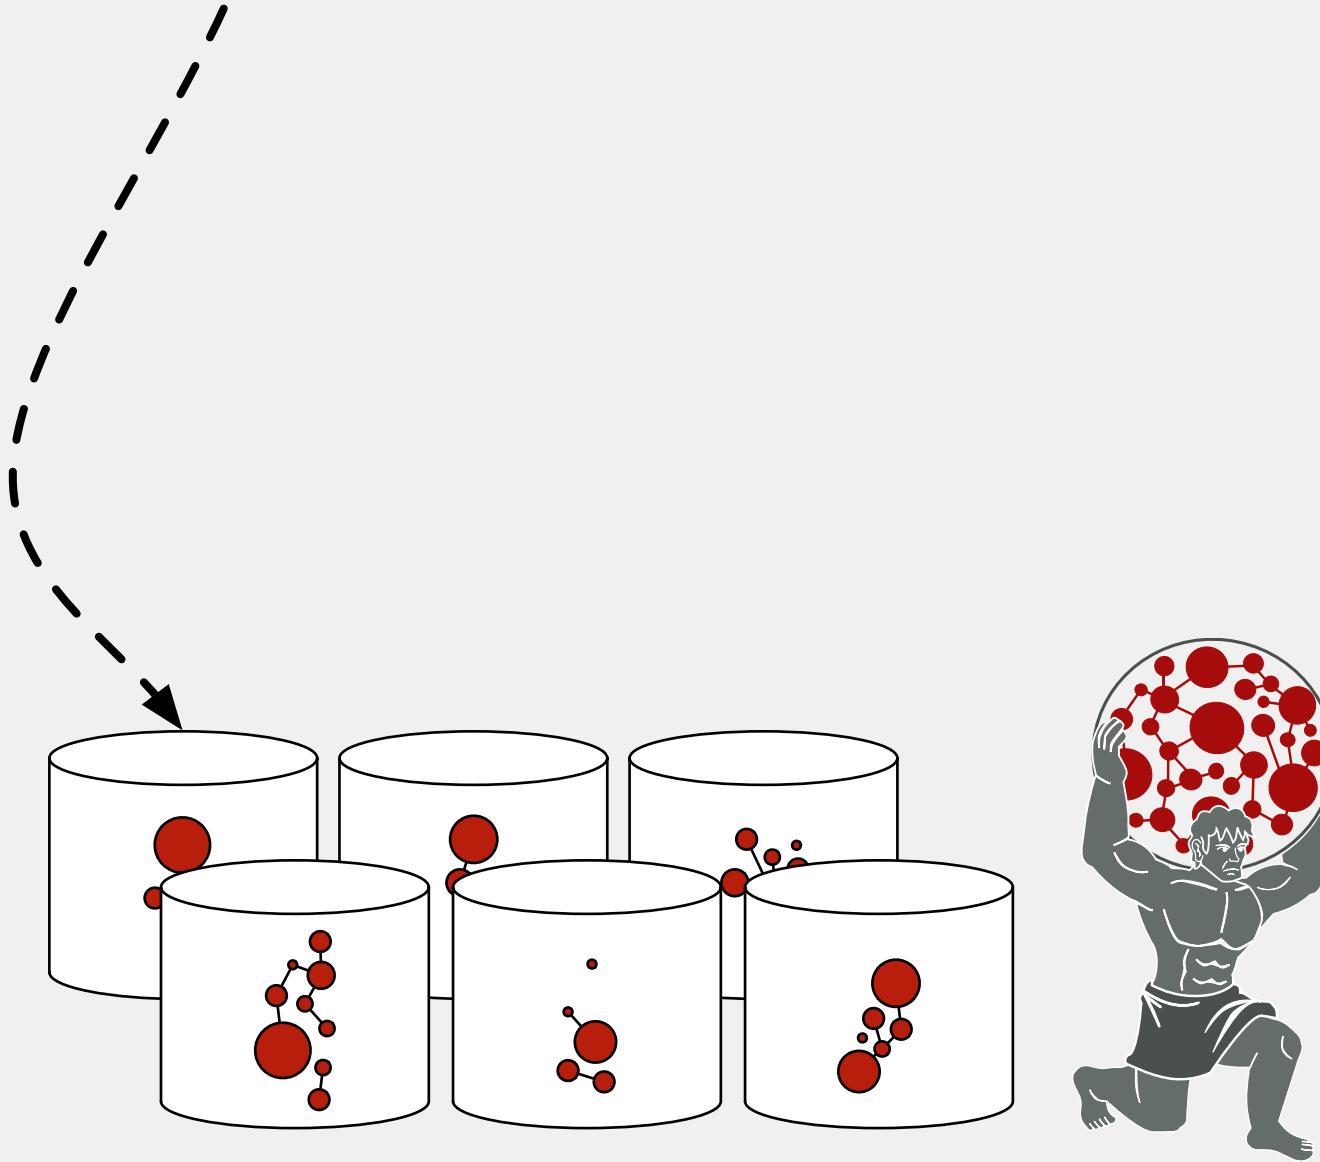


GREMLIN SERVER

```
gremlin> g = TitanFactory.open(conf)
==>titangraph[66.77.88.99]
```



localhost

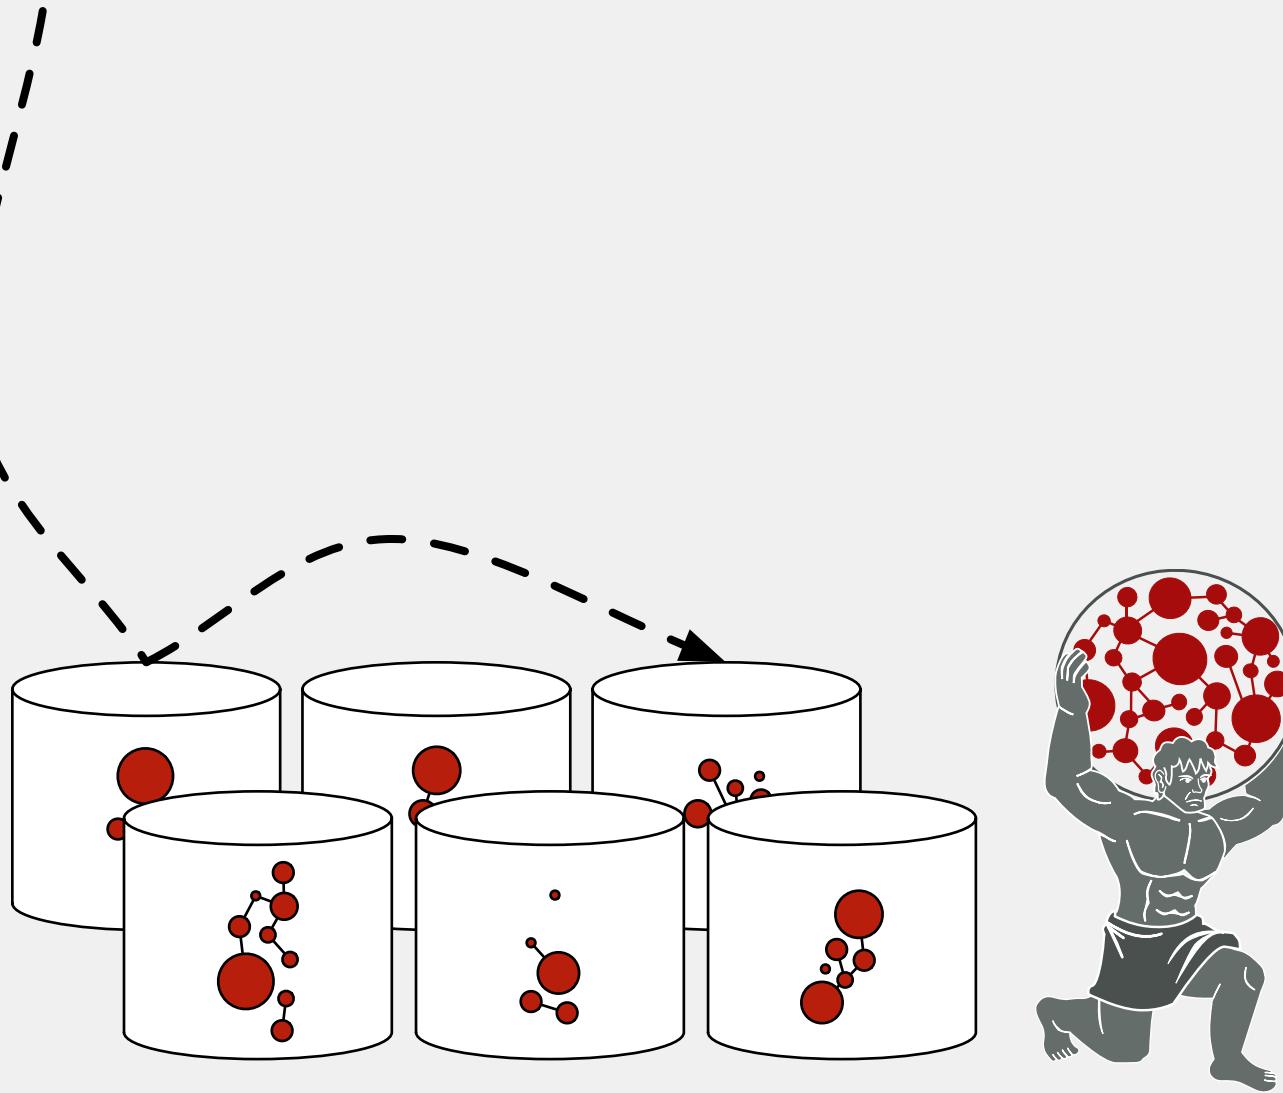


GREMLIN SERVER

```
gremlin> g = TitanFactory.open(conf)
==>titangraph[66.77.88.99]
gremlin> g.V.has('name','titan')
==>v[2]
```



localhost



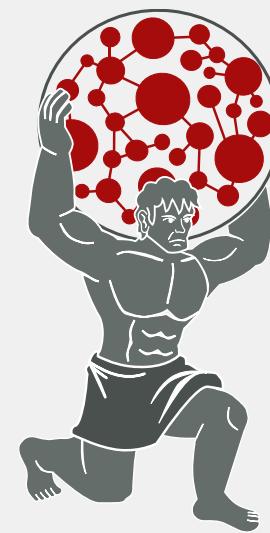
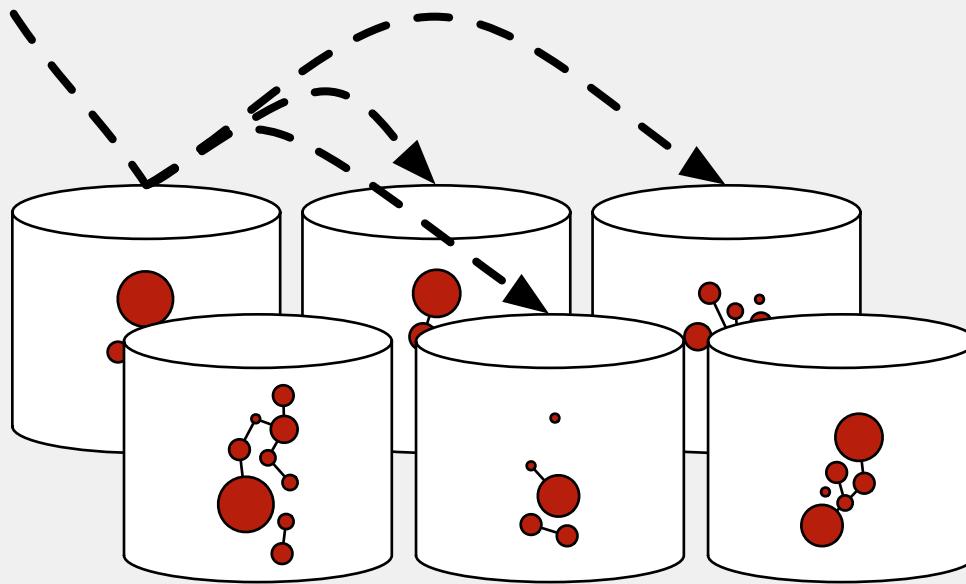
GREMLIN SERVER



localhost

```
gremlin> g = TitanFactory.open(conf)
==>titangraph[66.77.88.99]
gremlin> g.V.has('name','titan')
==>v[2]
gremlin> g.V.has('name','titan').out
```

1. get the titan vertex
2. get titan's adjacents



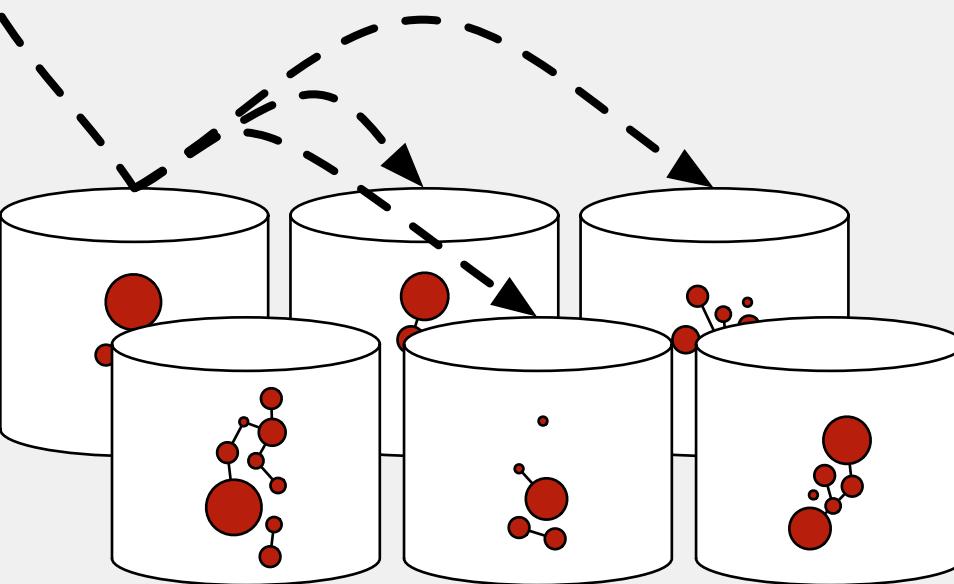
GREMLIN SERVER



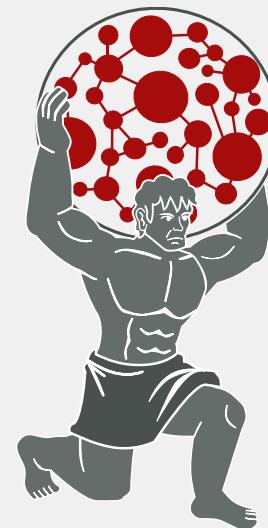
localhost

```
gremlin> g = TitanFactory.open(conf)
==>titangraph[66.77.88.99]
gremlin> g.V.has('name','titan')
==>v[2]
gremlin> g.V.has('name','titan').out
```

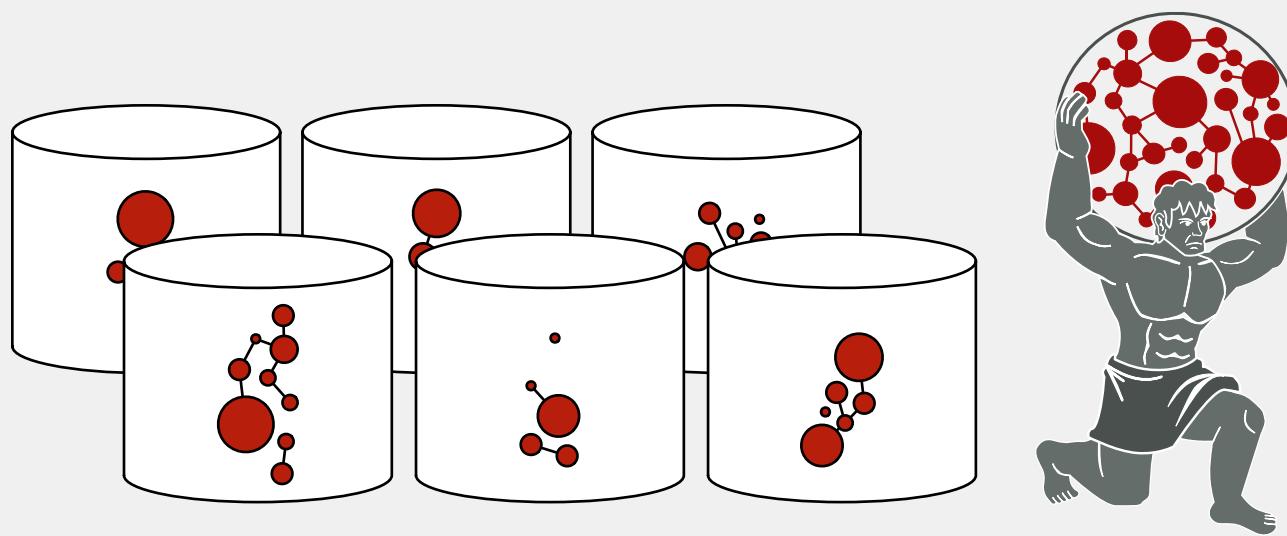
1. get the titan vertex
2. get titan's adjacents



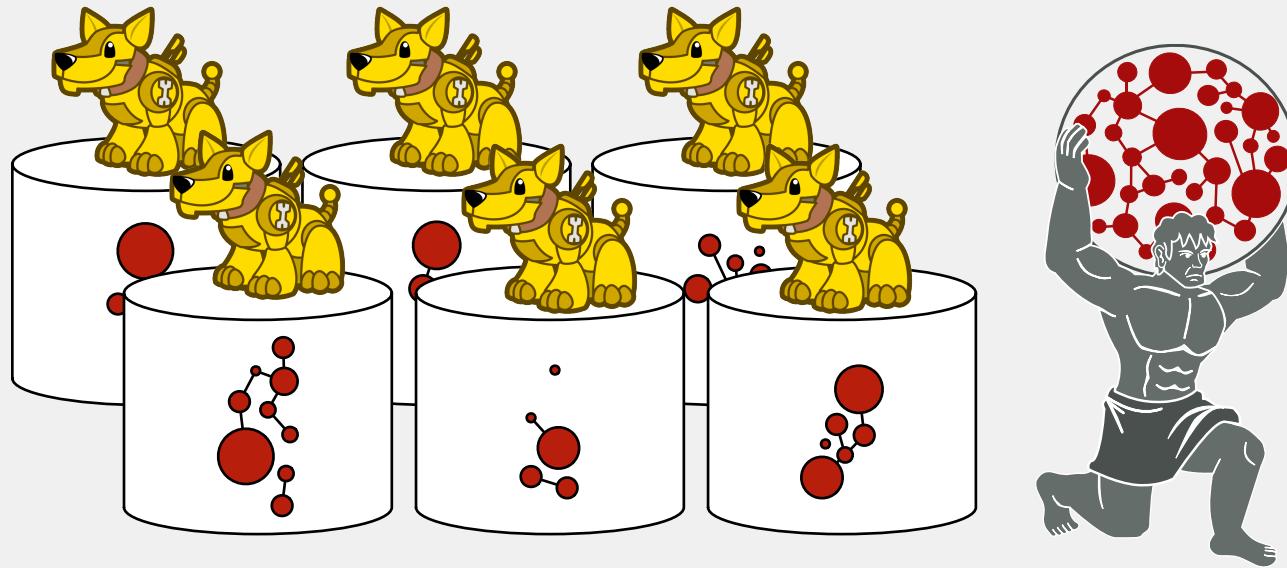
CHATTY
Each step in the traversal is a round trip



GREMLIN SERVER



GREMLIN SERVER

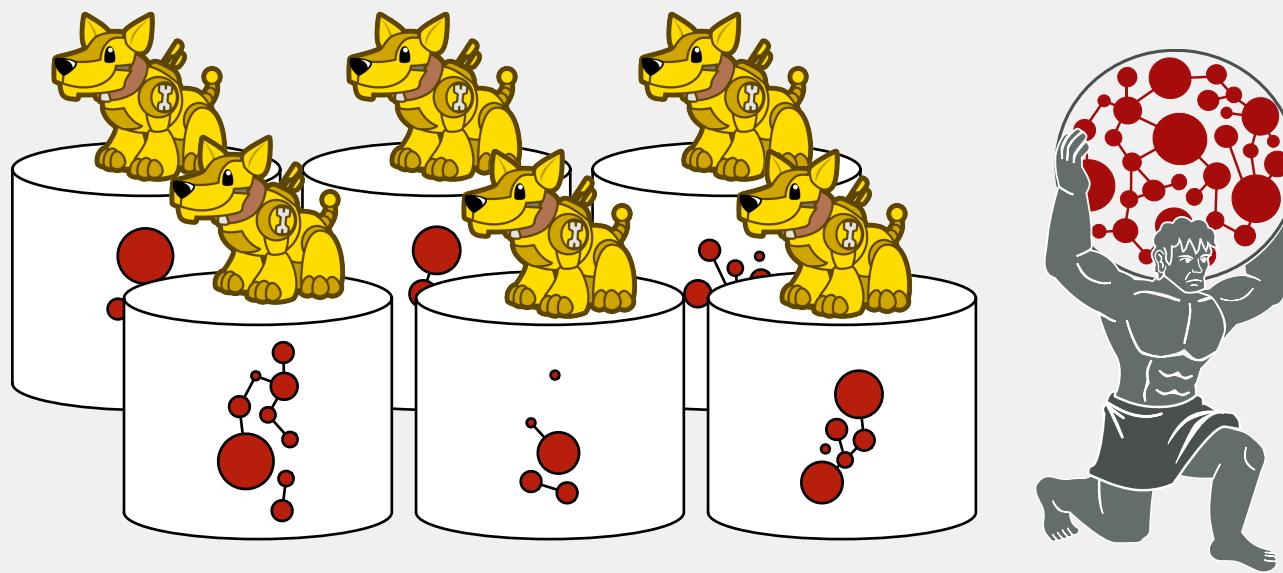


GREMLIN SERVER

```
gremlin> :remote connect tinkerpop.server titan.yaml  
==>connected - 66.77.88.99:8182,11.22.33.44:8182,...
```



localhost

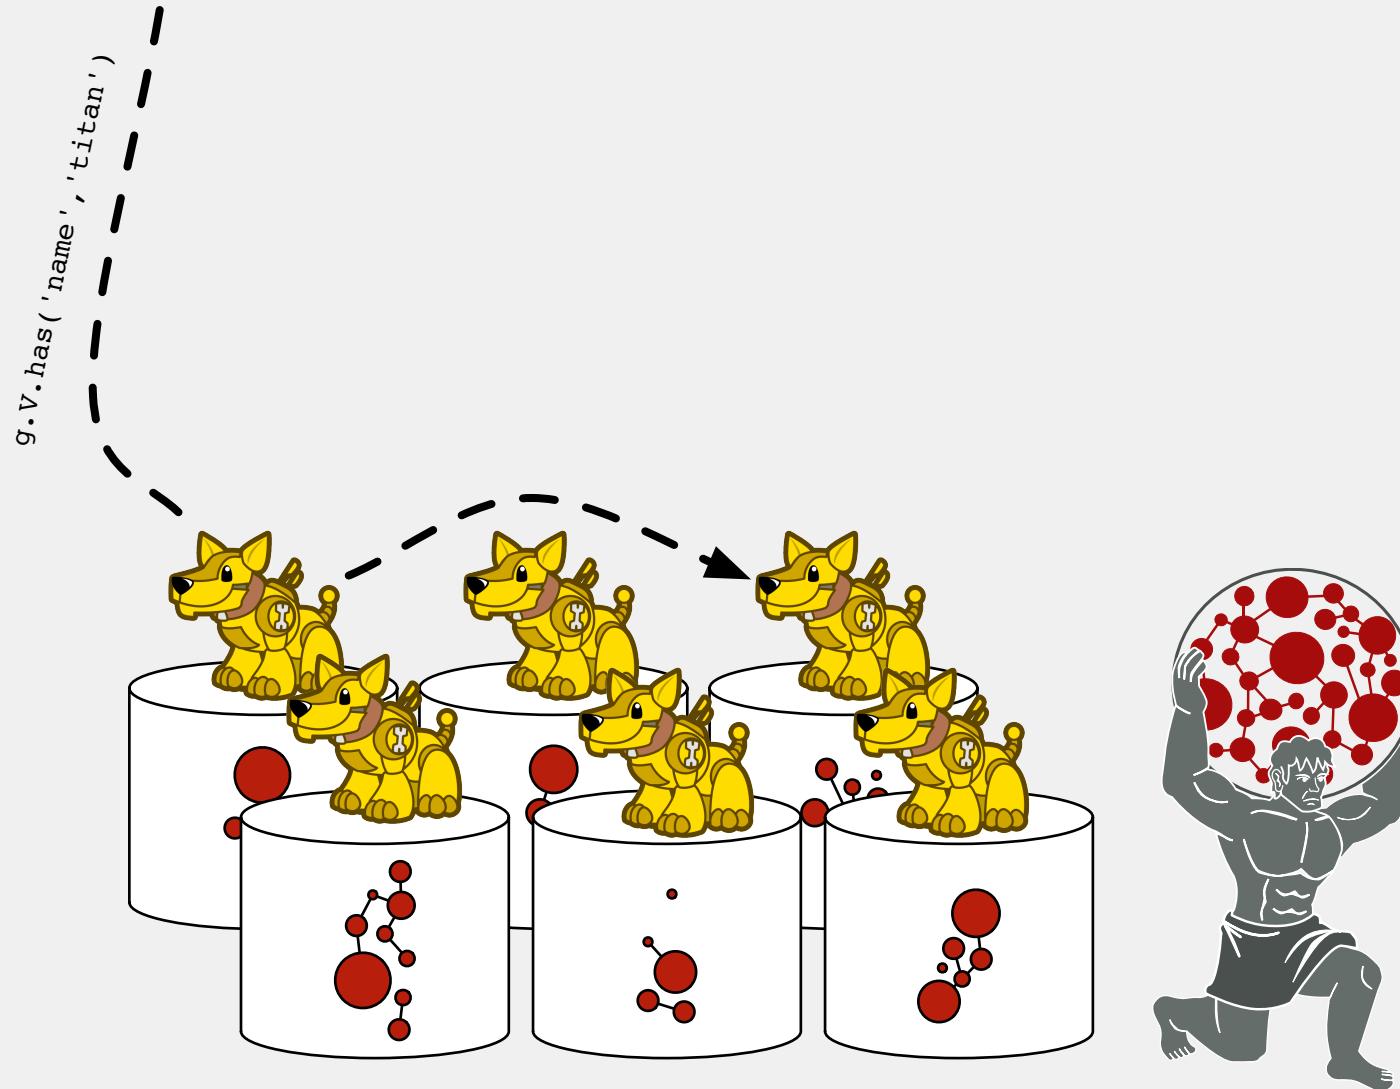


GREMLIN SERVER



localhost

```
gremlin> :remote connect tinkerpop.server titan.yaml  
==>connected - 66.77.88.99:8182,11.22.33.44:8182,...  
gremlin> :> g.v.has('name','titan')  
==>v[2]
```

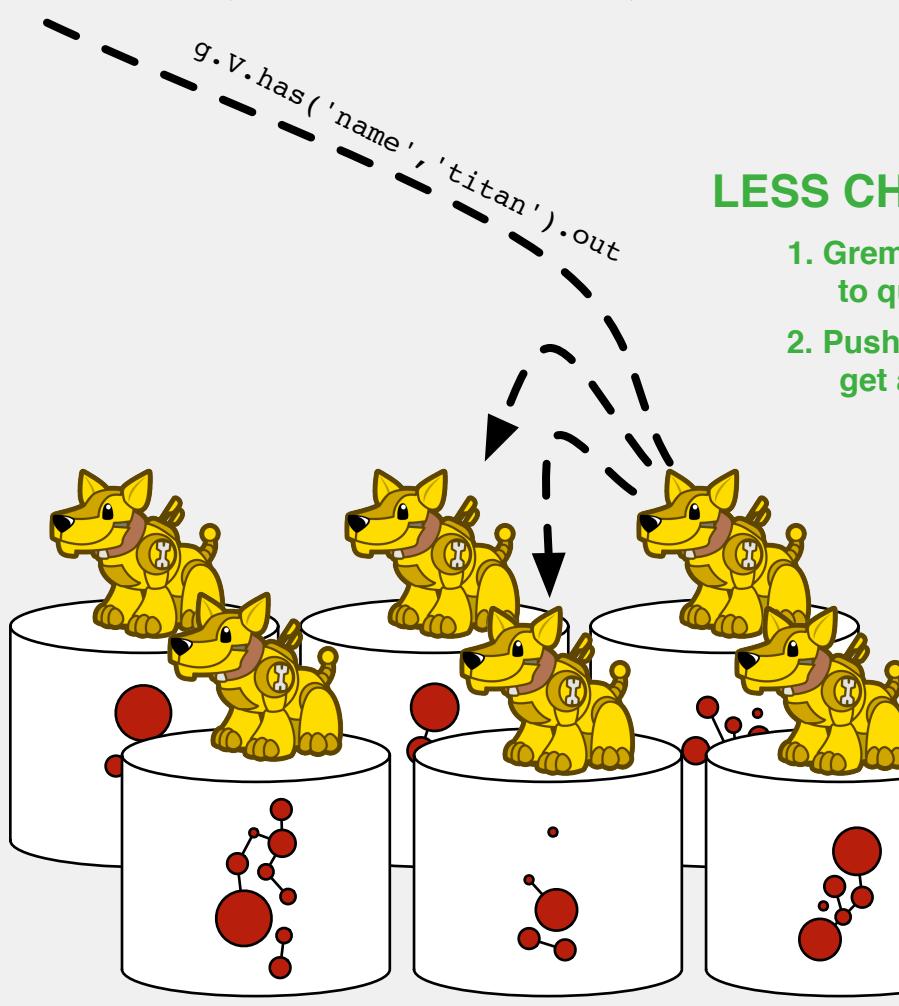


GREMLIN SERVER



localhost

```
gremlin> :remote connect tinkerpop.server titan.yaml  
==>connected - 66.77.88.99:8182,11.22.33.44:8182,...  
gremlin> :> g.V.has('name','titan')  
==>v[2]  
gremlin> :> g.V.has('name','titan').out
```



LESS CHATTY

1. Gremlin Server client knows the hash function to query the machine with the Titan vertex
2. Push the traversal to the server and get a WebSockets iterator back

GREMLIN SERVER



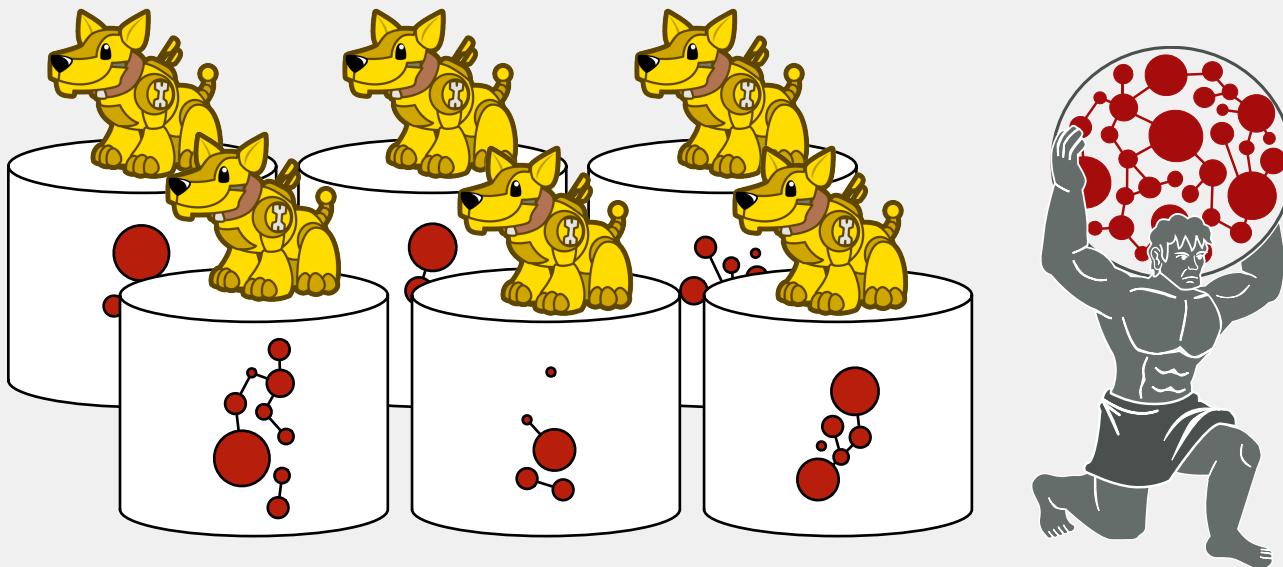
Allows non-JVM languages to interact with a remote graph.

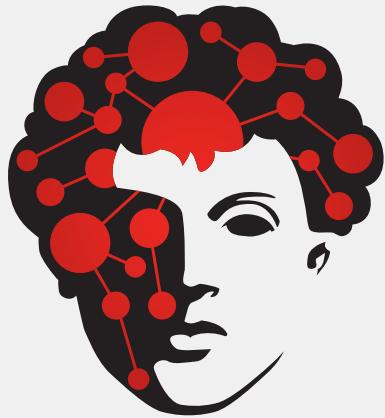


Supports stored procedures -- MyAlgorithms.class



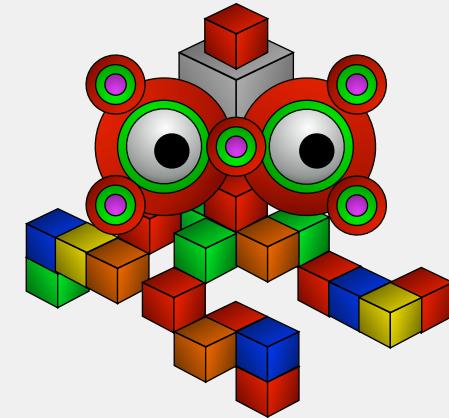
Supports WebSockets, REST, and various serialization mechanisms.





CREDITS

PRESENTED BY
MARKO A. RODRIGUEZ



MANY THANKS TO
MATTHIAS BRÖCHELER
STEPHEN MALLETTÉ
DAN LAROCQUE
DANIEL KUPPITZ
BOB BRIODY
BRYN COOKE
JOSH SHINAVIER
PAVEL YASKEVICH
AURELIUS COMMUNITY
TINKERPOP COMMUNITY
KETRINA YIM